



# Kubernetes

In Practice





Somkiat Puisungnoen

Update Info 1 View Activity Log 10+ ...

Timeline About Friends 3,138 Photos More

When did you work at Opendream? X

... 22 Pending Items

Post Photo/Video Live Video Life Event

What's on your mind?

Public Post

Intro

Software Craftsmanship

Software Practitioner at สยามช่างนาฏกิจ พ.ศ. 2556

Agile Practitioner and Technical at SPRINT3r

Somkiat Puisungnoen 15 mins · Bangkok · ...

Java and Bigdata

 Kubernetes in practice

© 2017 - 2018 Siam Chamnankit Company Limited. All rights reserved.

Facebook somkiat.cc

Page Messages Notifications 3 Insights Publishing Tools Settings Help ▾

somkiat.cc  
@somkiat.cc

Home Posts Videos Photos

Liked Following Share ... + Add a Button

Help people take action on this Page. ×



# Agenda Day 1

1. Cloud Native Application
2. Kubernetes architecture
3. Key-features
4. Pods and Containers
5. Service
6. Replication Controller (RC)
7. Deployment and ReplicaSet (RS)
8. Volume



# Agenda Day 2

1. Resource management
2. Horizontal Pods Autoscaler (HPA)
3. ConfigMap and Secret
4. Log and monitoring
5. Ingress network
6. Working with Persistence storage

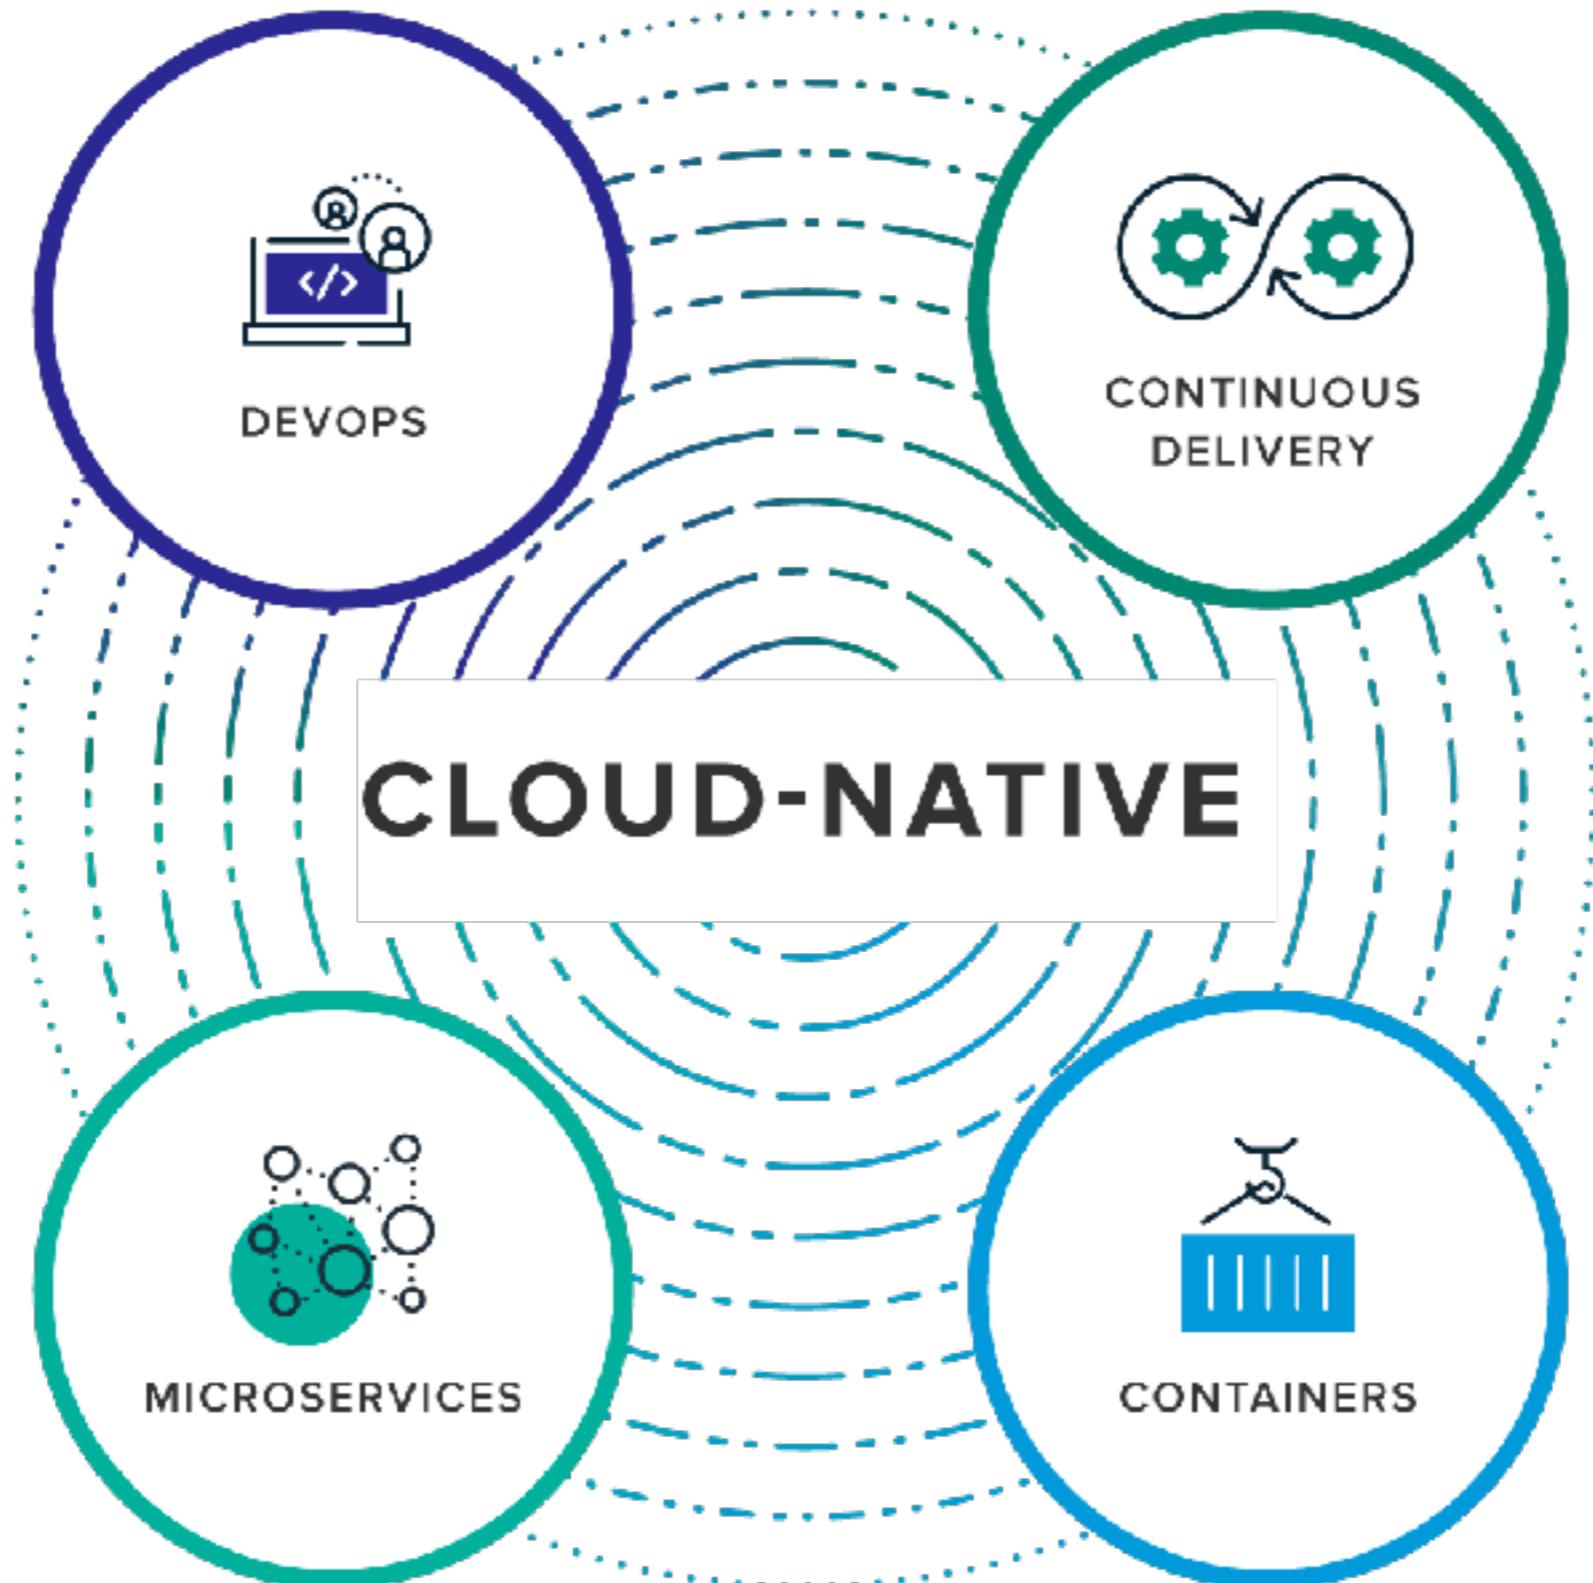


**<https://github.com/up1/course-kubernetes-in-practice>**



# Cloud Native Application





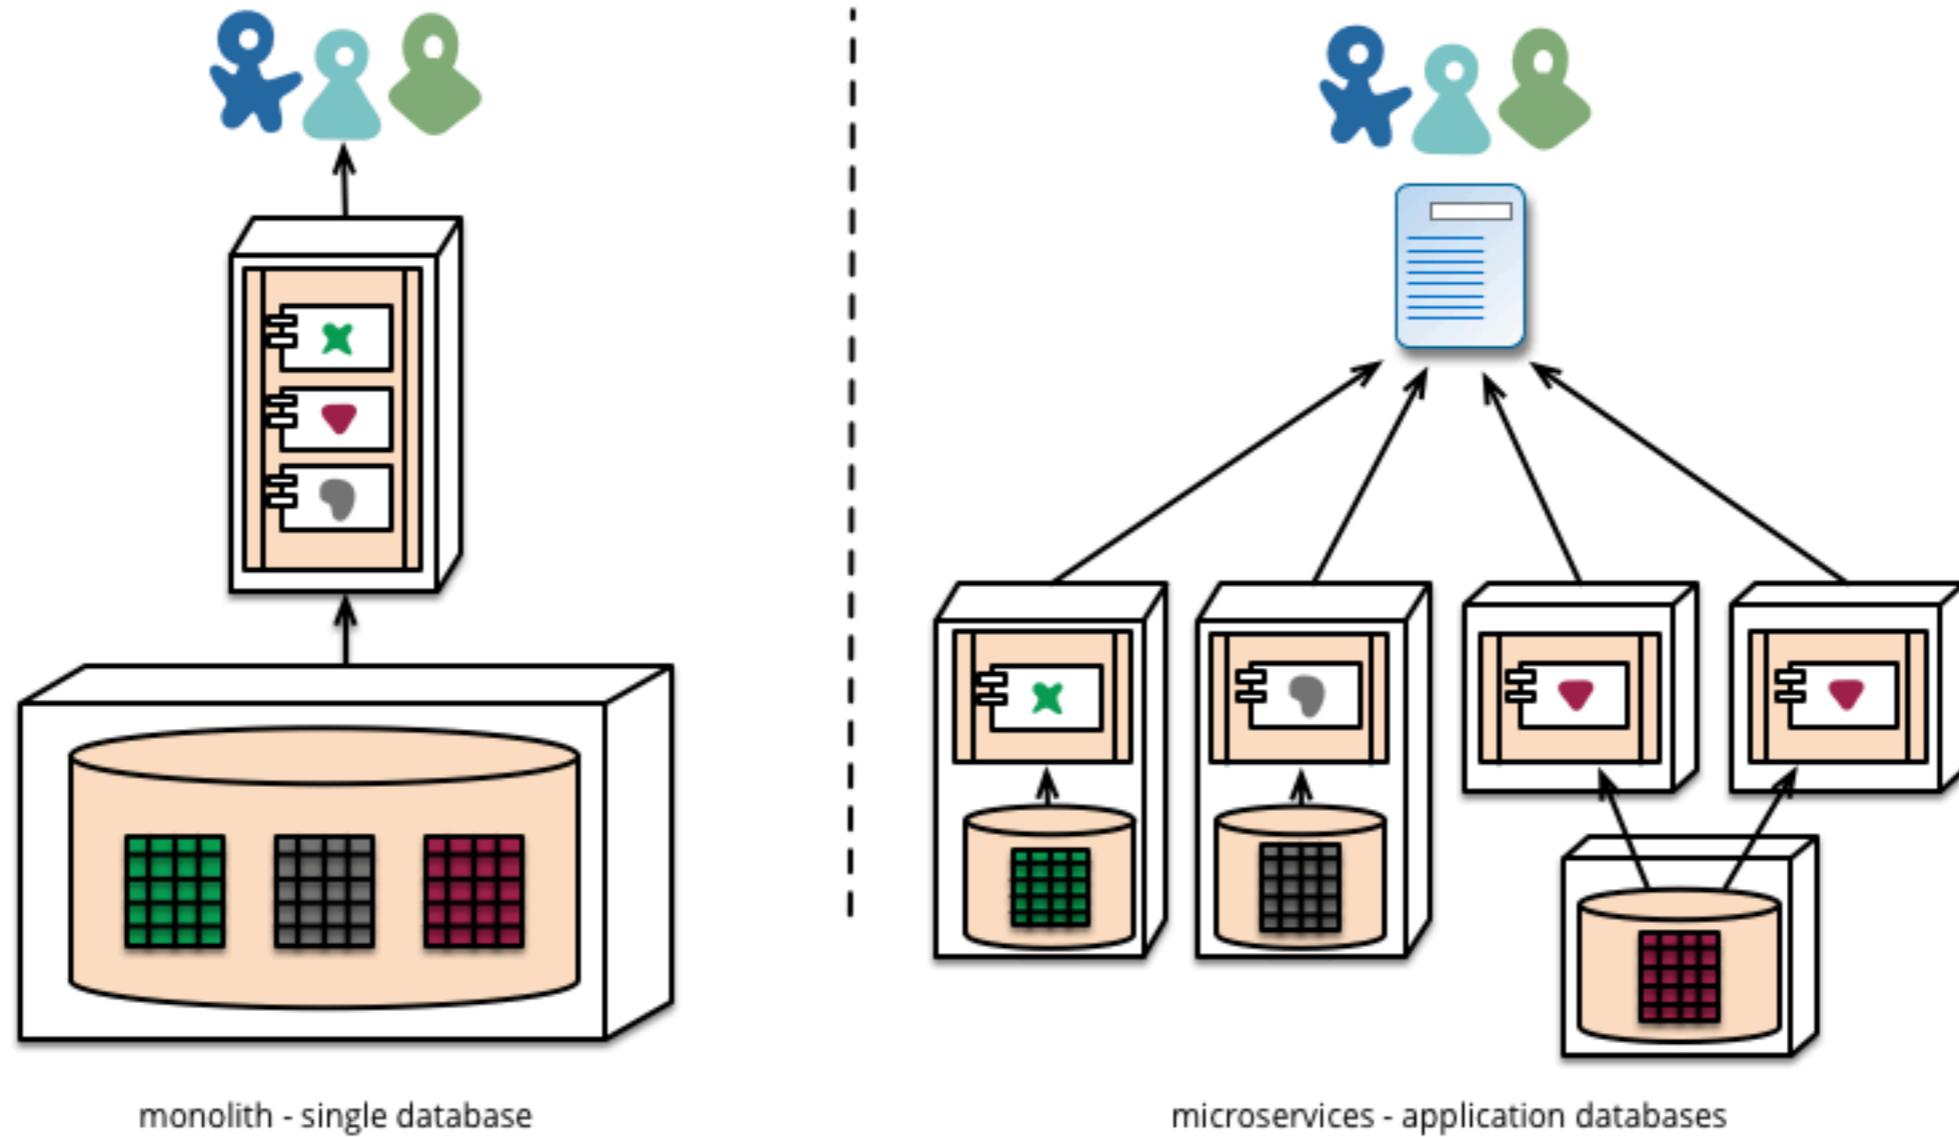
<https://pivotal.io/cloud-native>



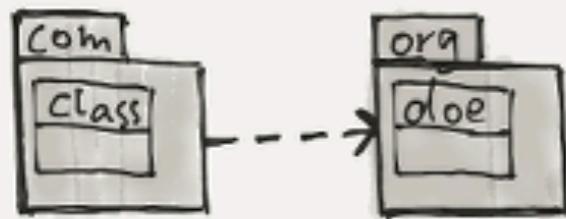
# Evolution of Architecture



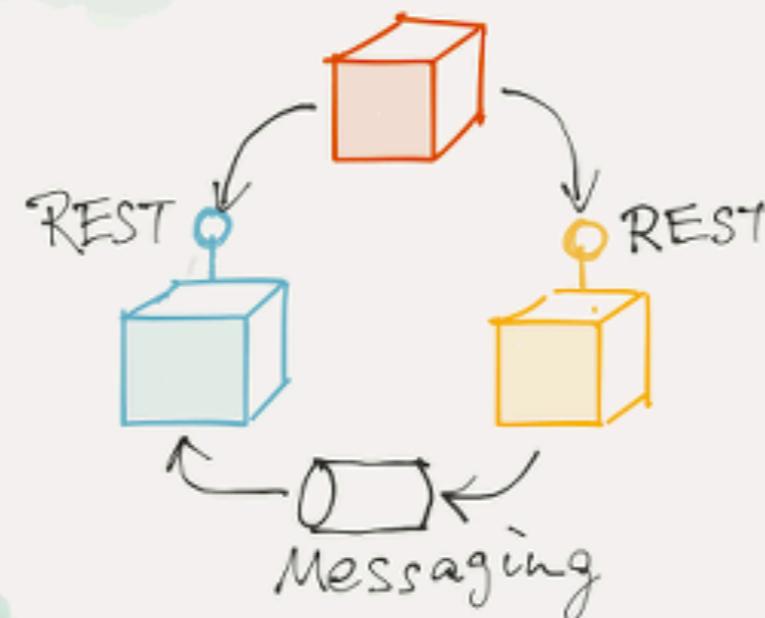
# Microservices



# Architecture



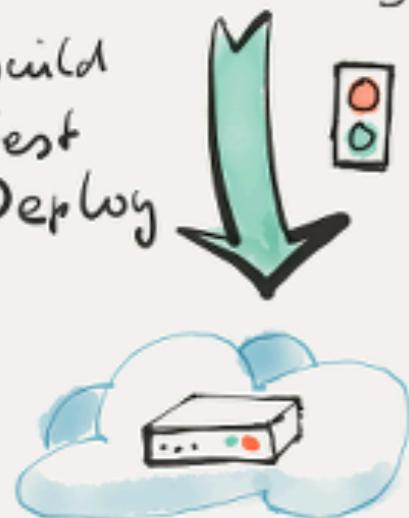
# Microservices



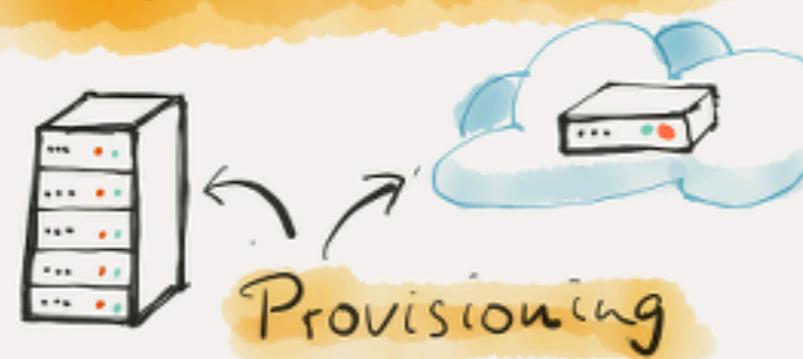
# Deployment

## Continuous Delivery

`{ var i=1; }`  
Build  
Test  
Deploy



# Infrastructure

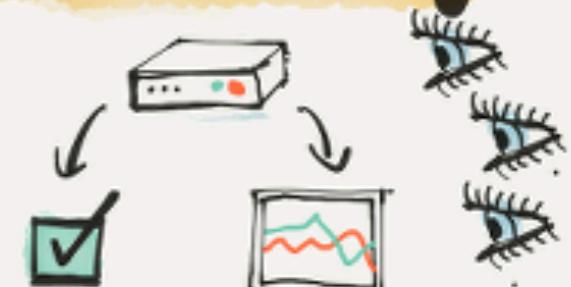


# People & Teams



Communication  
Collaboration

# Monitoring

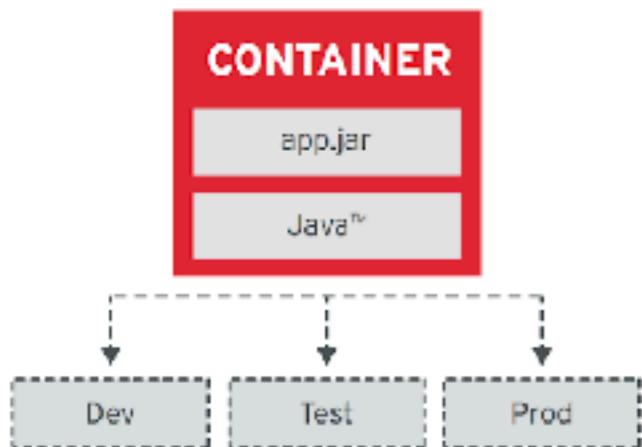


Features & Technology



# Container design principles

Image Immutability Principle



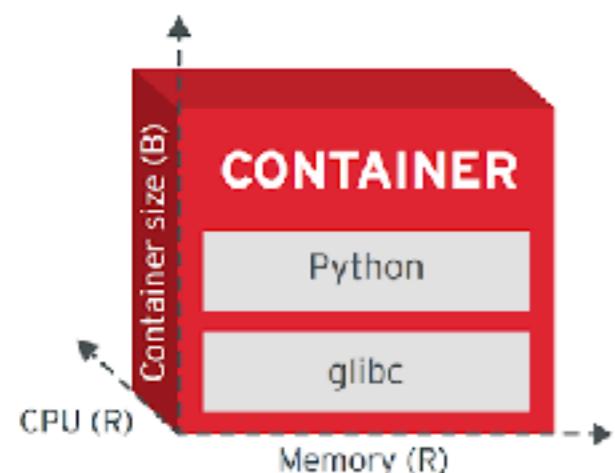
High Observability Principle



Process Disposability Principle



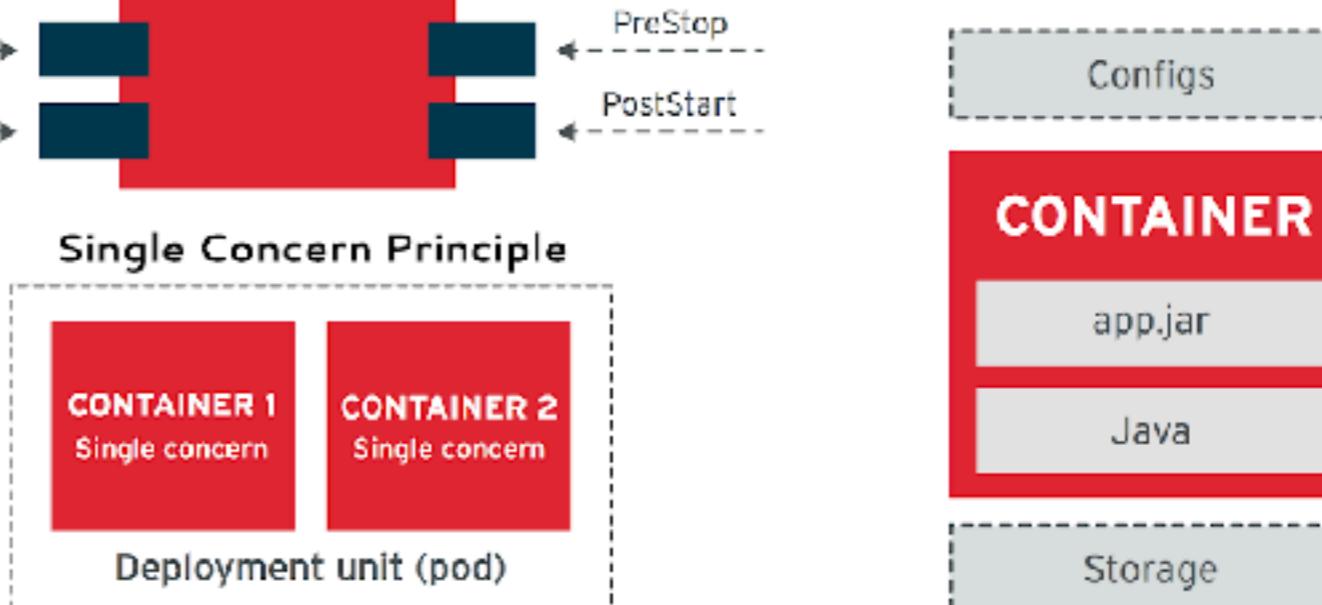
Runtime Confinement Principle



Lifecycle Conformance Principle



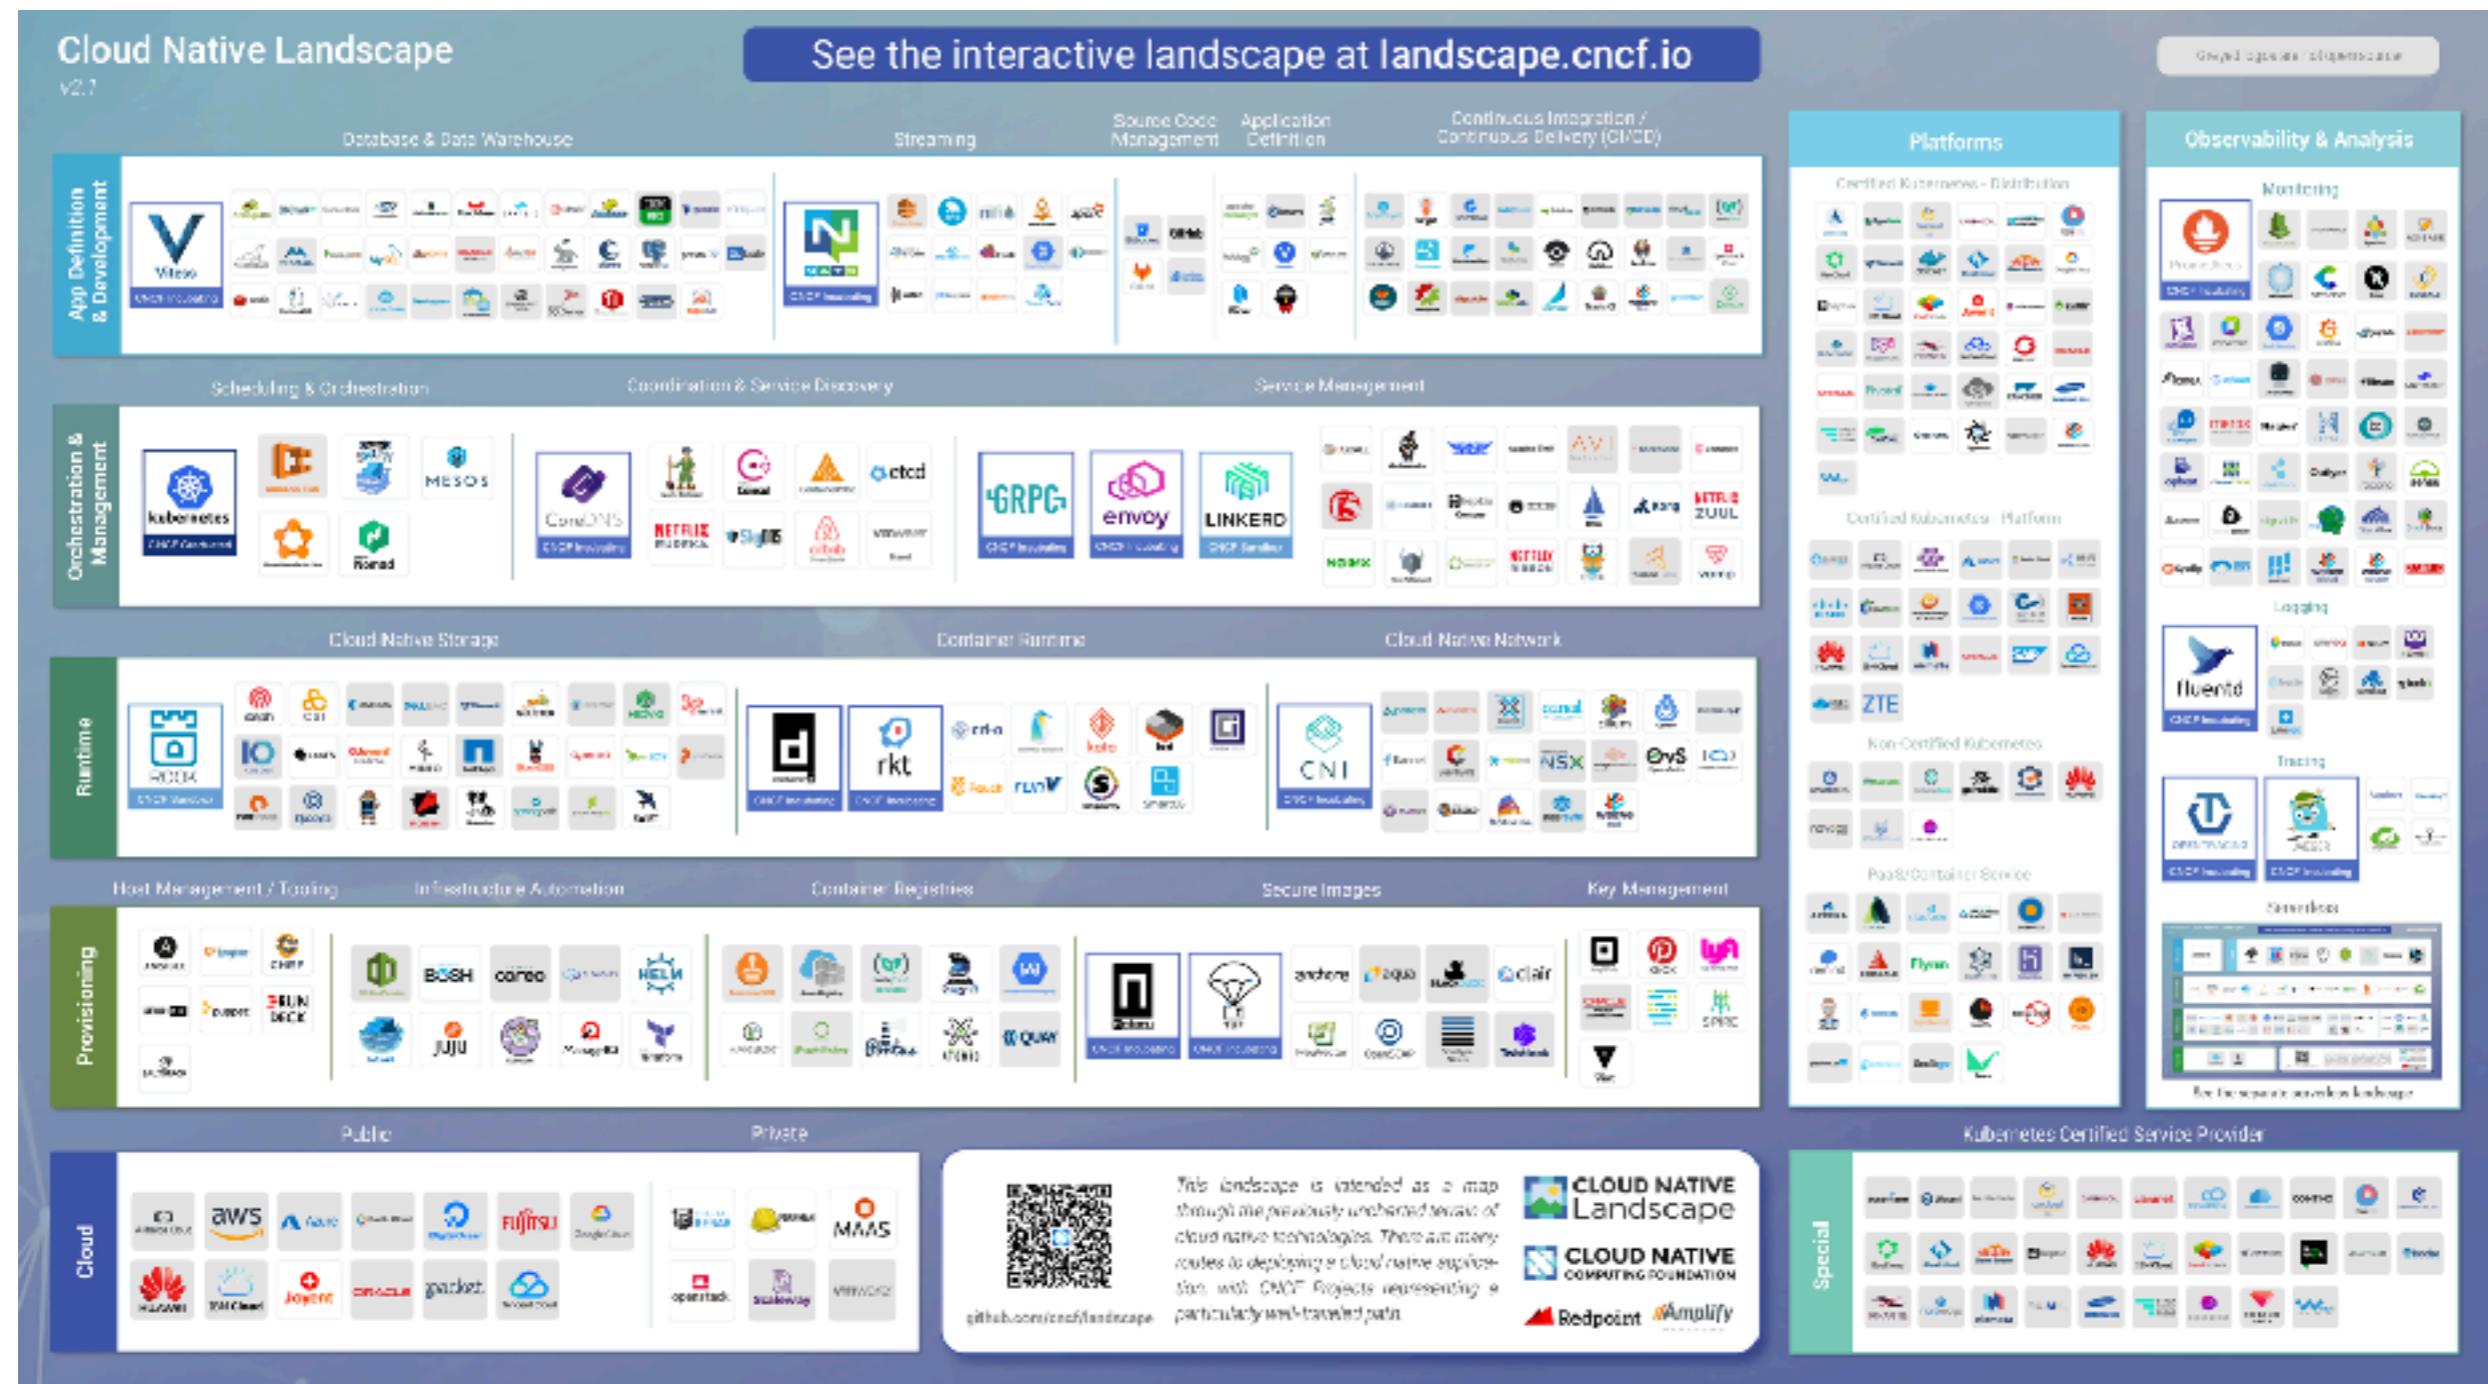
Self-Containment Principle



Single Concern Principle



# Cloud native landscape



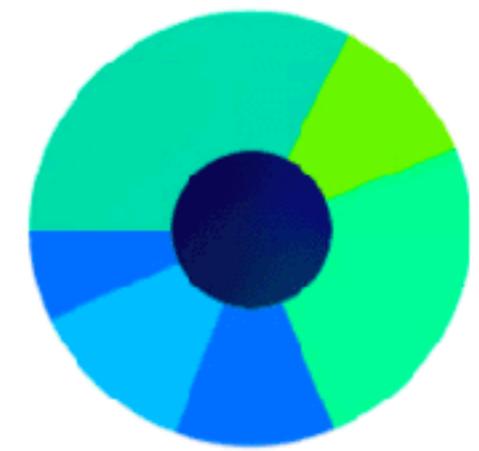
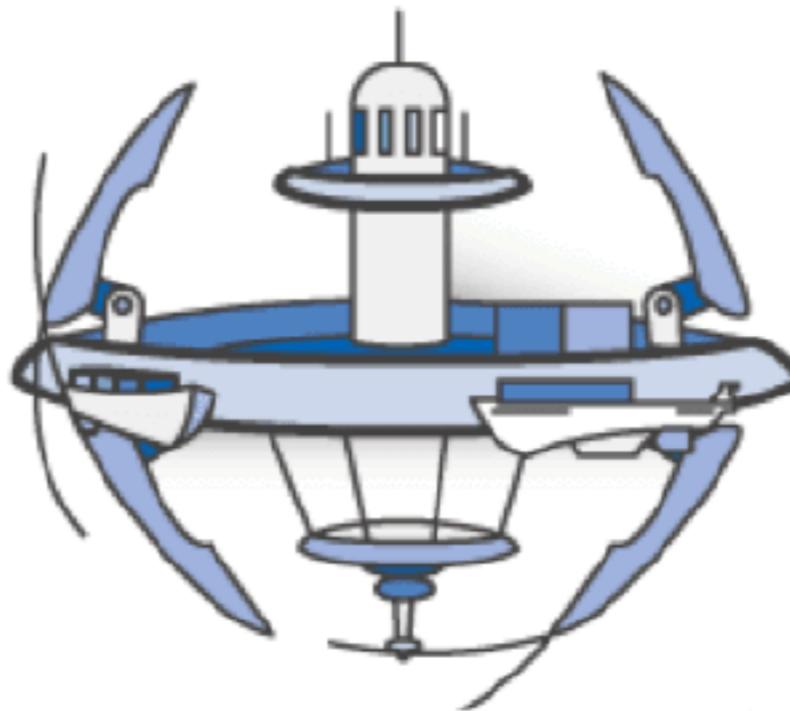
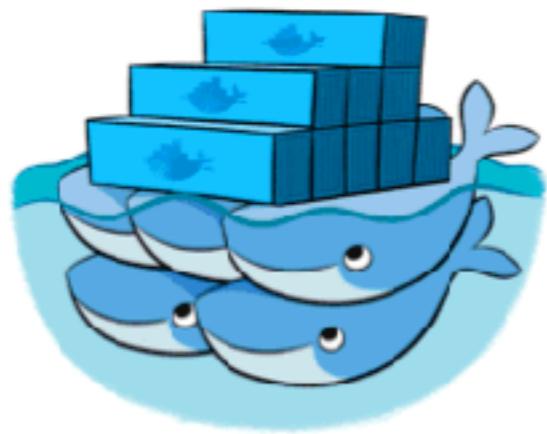
<https://github.com/cncf/landscape>







MESOS



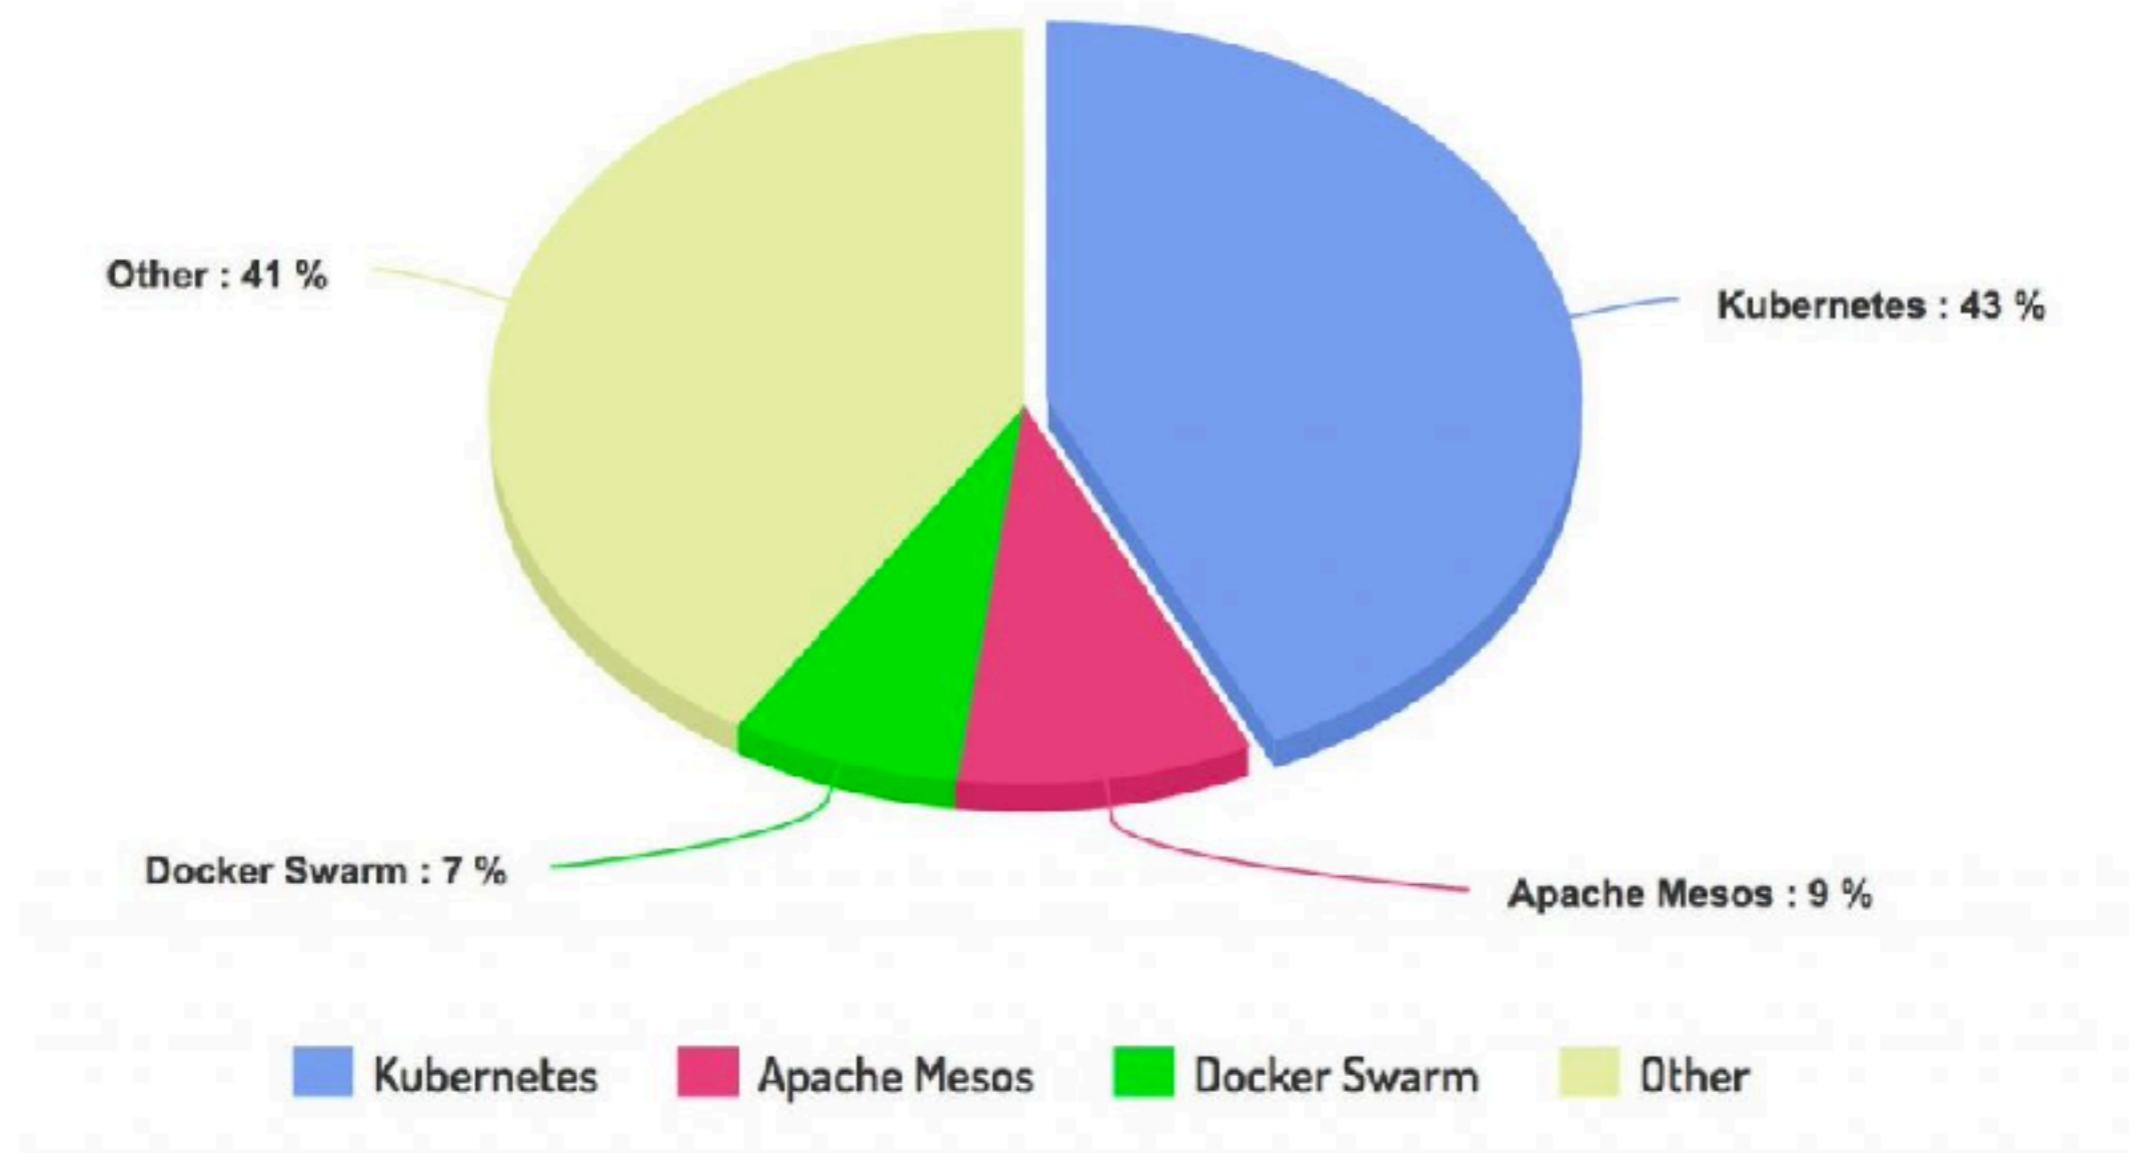
Marathon



kubernetes  
by Google



# Container Orchestrators in Sysdig's 2017 Docker Usage Report



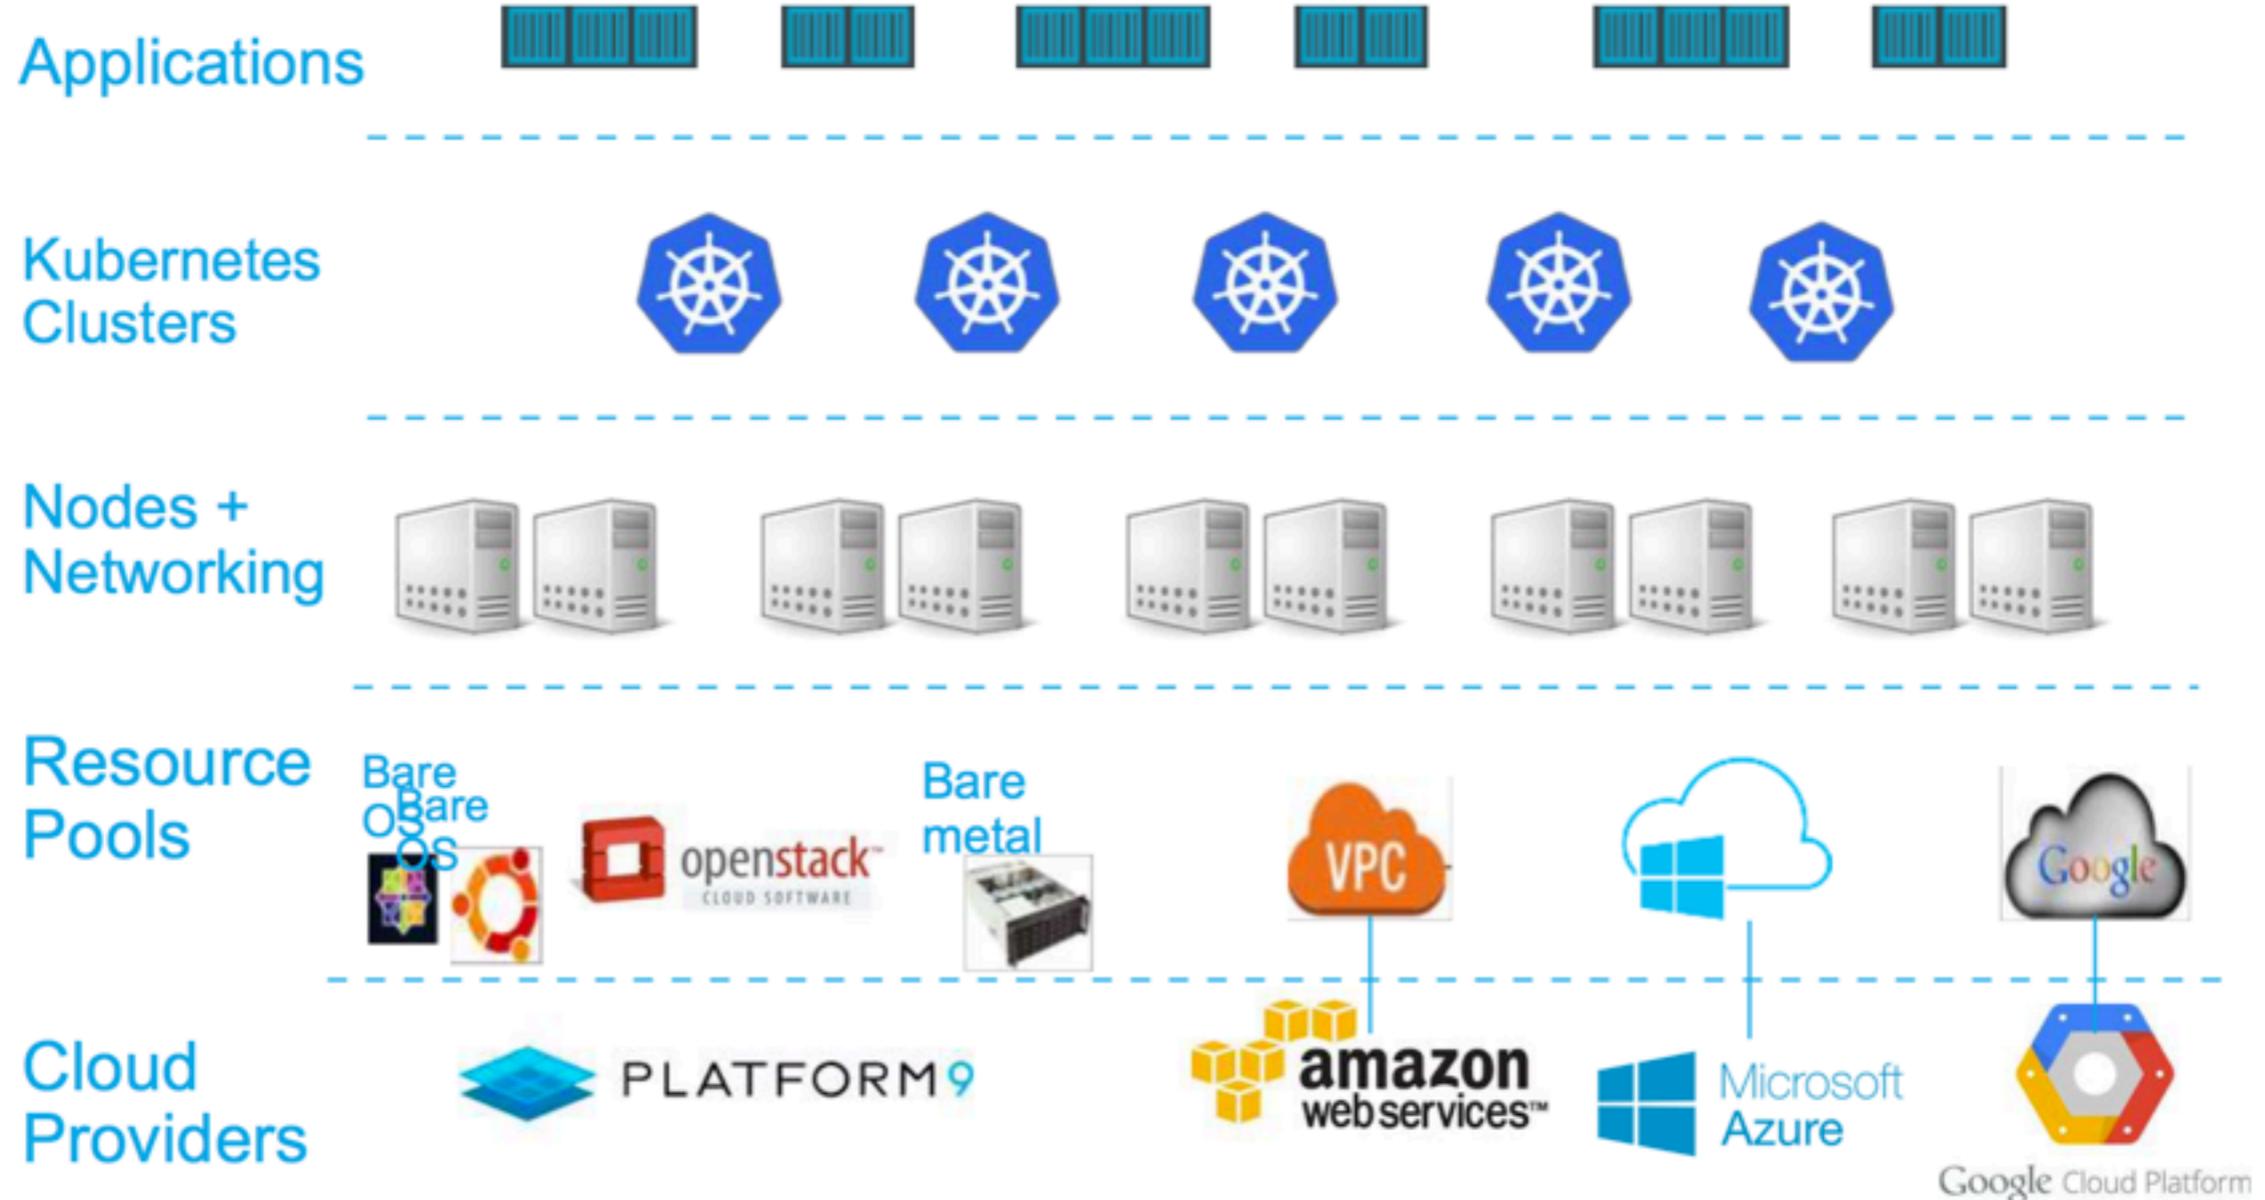
<https://sysdig.com/blog/sysdig-docker-usage-report-2017/>



# Why Kubernetes ?



# Write once, run anywhere



# Write once, run anywhere

Eliminate infrastructure lock-in

Use containers

Provides management for containers



# Modular app design

Monolithic app makes everything worse

- Larger teams slow thing down

- Spaghetti dependencies

Lack of ownership for sharing components

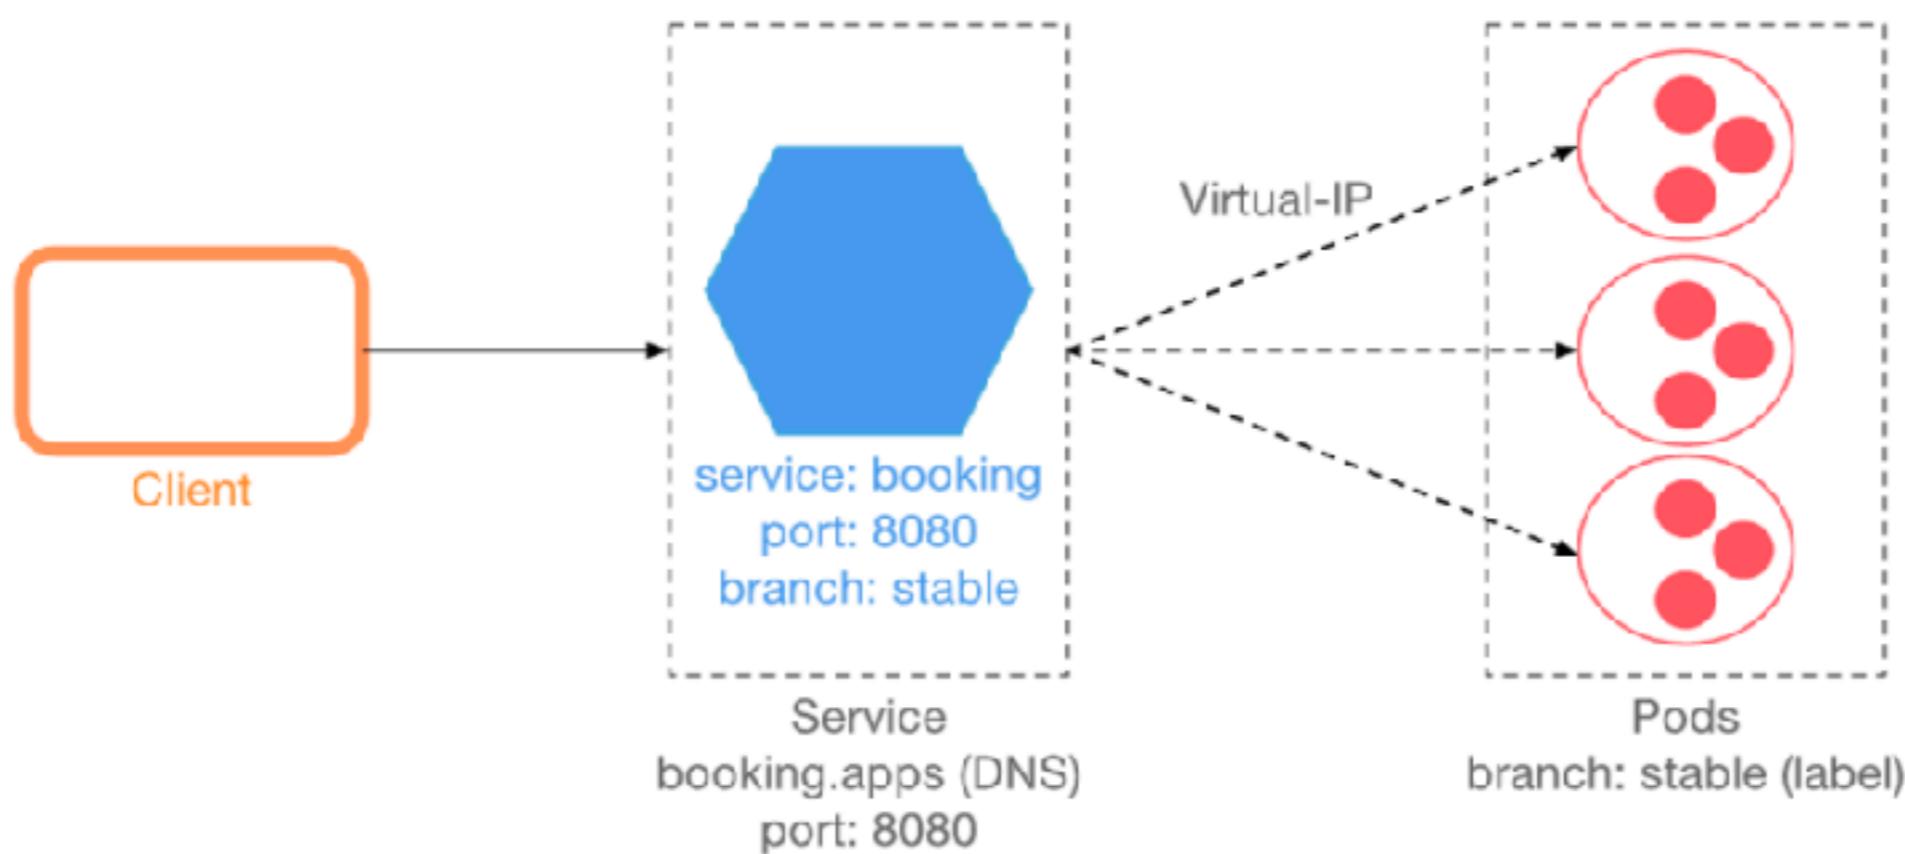
- Complexity to testing

- Slow building



# Modular app design

Container/Image boundary similar to class  
How to use/manage a collection of container ?



# Fault-tolerant by design

Design for failure

Infrastructure provisioning/re-provisioning

Configuration networking and load balance

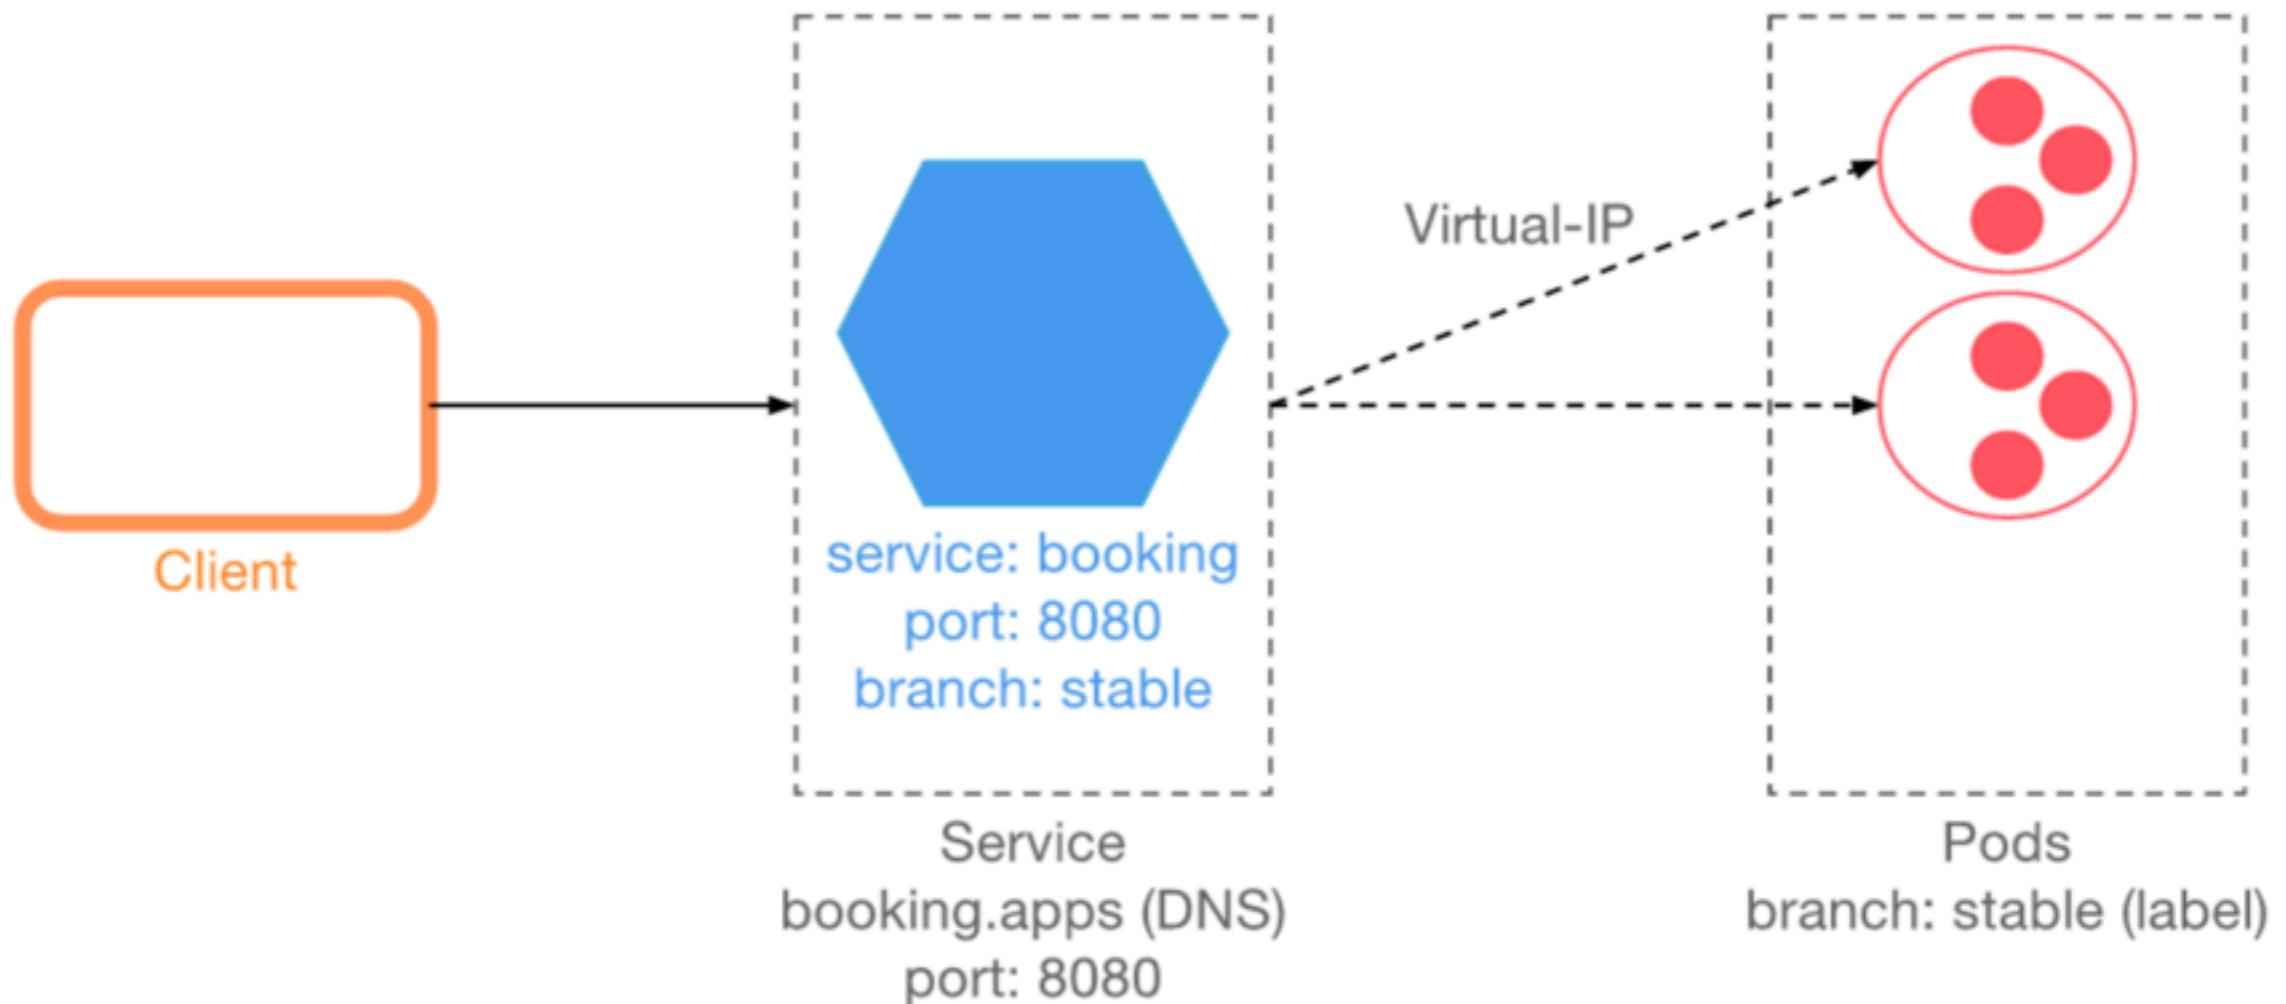
Redundancy (scale-out)

Lifecycle management (Software update)



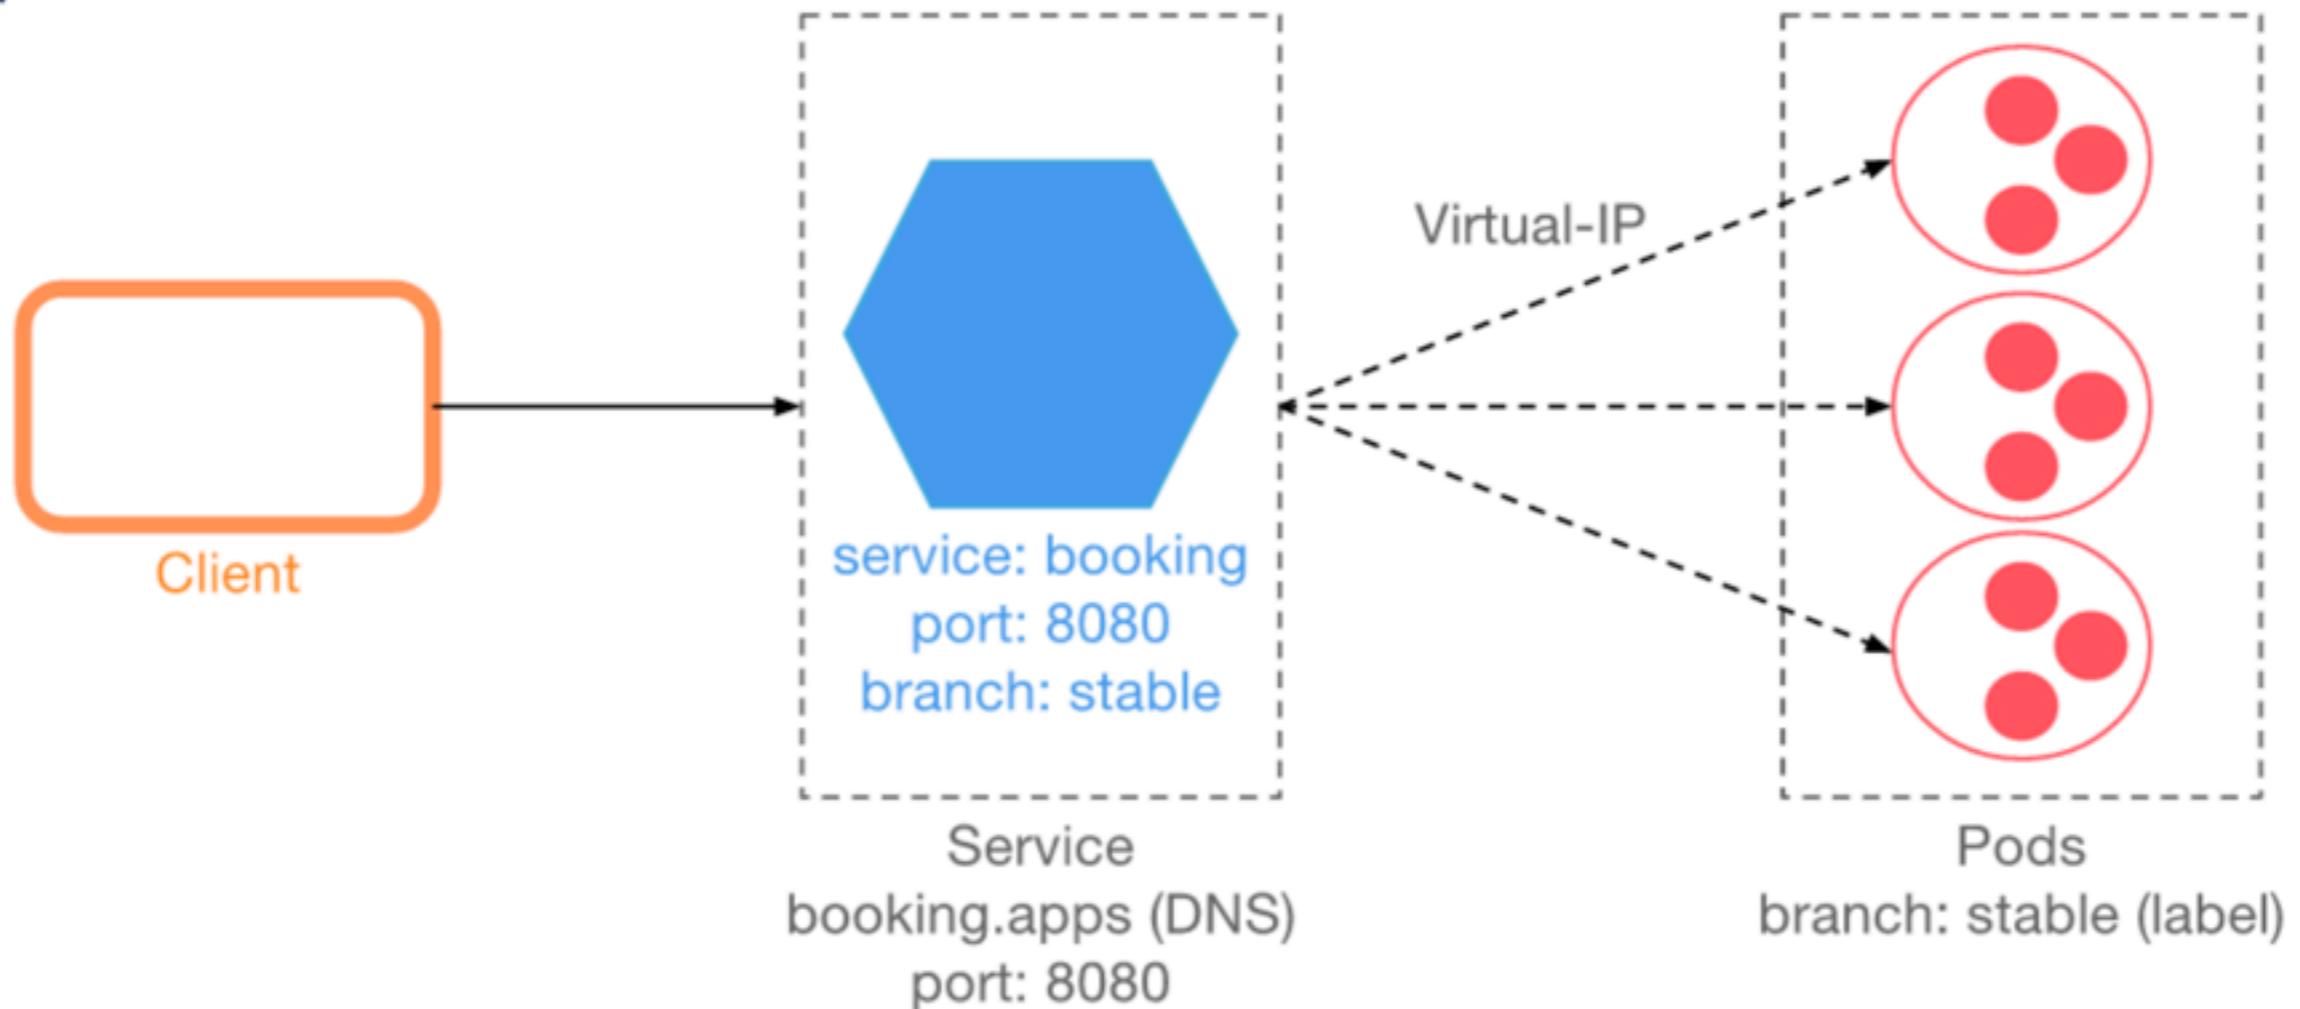
# Fault-tolerant by design

replicas = 2



# Fault-tolerant by design

replicas = 3



# Deployment, not Infrastructure

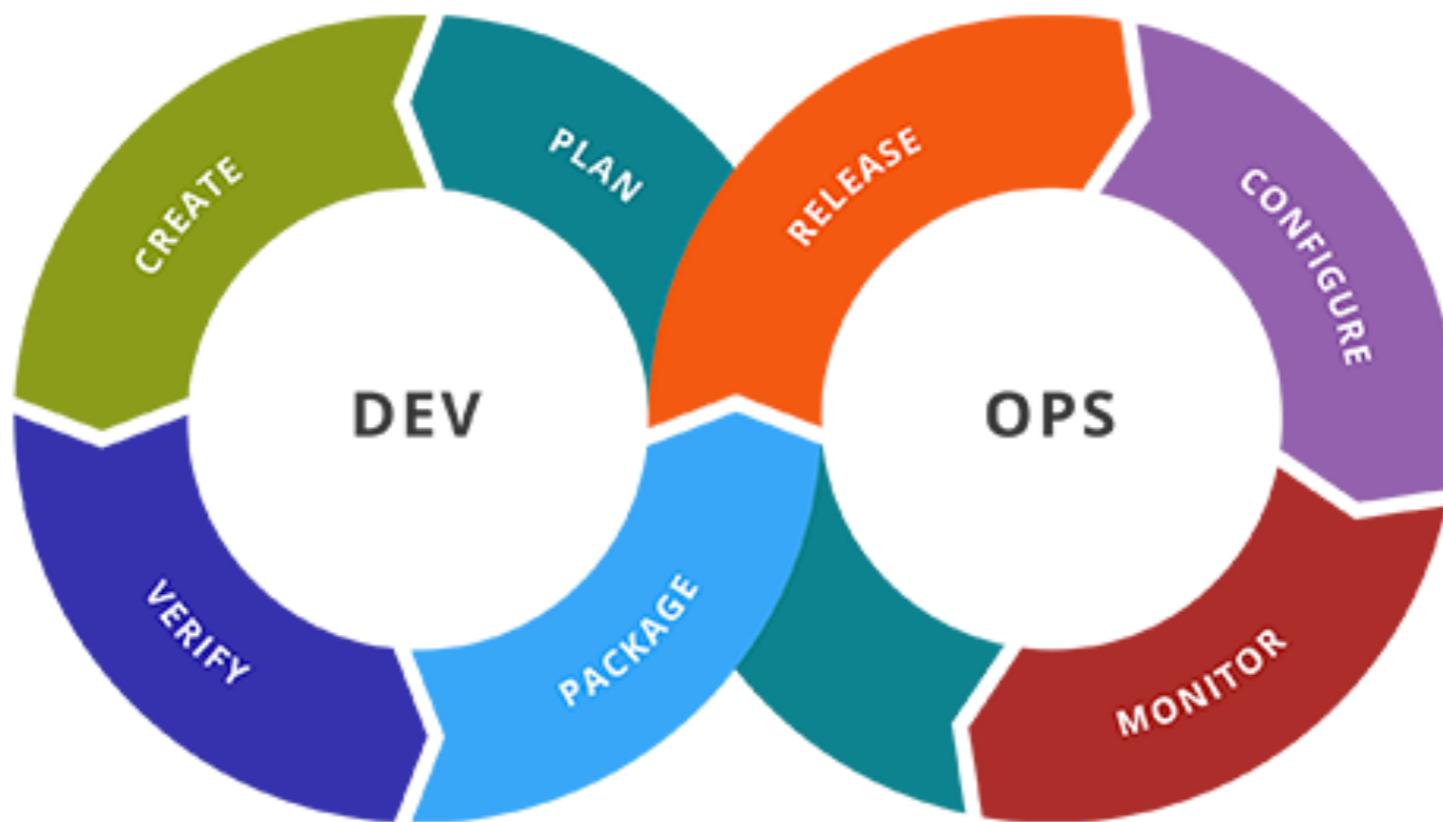
Software deployment is hard  
Infrastructure provisioning/re-provisioning  
Configuration networking and load balance  
Redundancy (scale-out)  
Lifecycle management (Software update)



# Deployment, not Infrastructure

Kubernetes support for deployment  
Controllers are in focus  
Scale-out service  
Rolling update for new version  
Rollback to a previous version  
Pause and resume a deployment  
Horizontal auto-scaling  
Canary deployment





# Let's start !!



# Installation



# Software requirement

minkube

Kubernetes command-line tool

<https://github.com/kubernetes/minikube>



# Starting minikube

```
$minikube start  
$minikube status
```



# Version of Kubernetes

```
$kubectl get nodes -o yaml
```

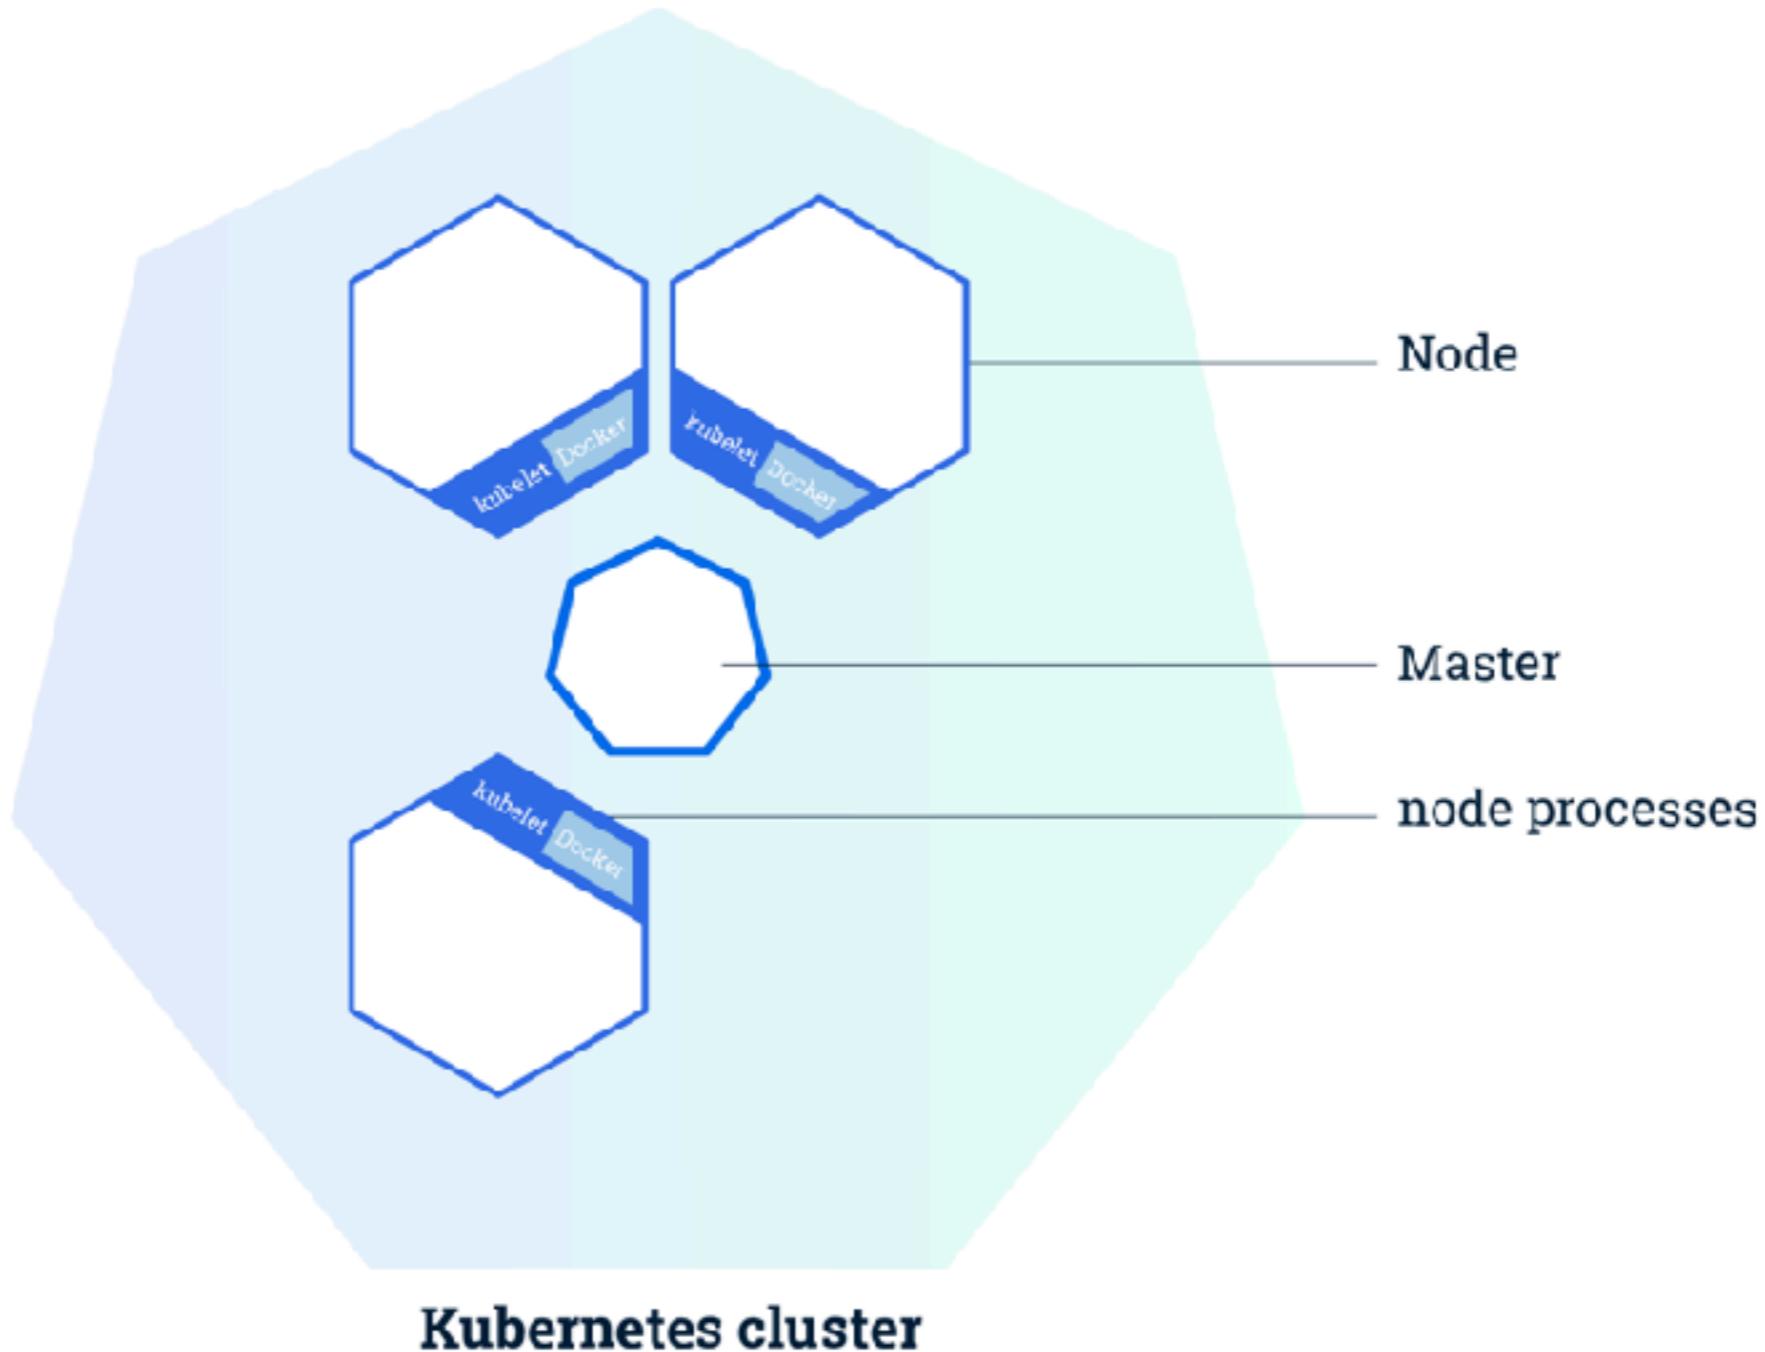
```
nodeInfo:  
  architecture: amd64  
  bootID: 7c769465-27e5-4dd8-a89f-319ba1b8ef57  
  containerRuntimeVersion: docker://17.9.0  
  kernelVersion: 4.9.64  
  kubeProxyVersion: v1.9.4  
  kubeletVersion: v1.9.4  
  machineID: edffd4ca8bf24253a736ea86d2448185  
  operatingSystem: linux  
  osImage: Buildroot 2017.11  
  systemUUID: 98BAE0EF-8C9E-45DF-9AC7-F0E05806B189
```



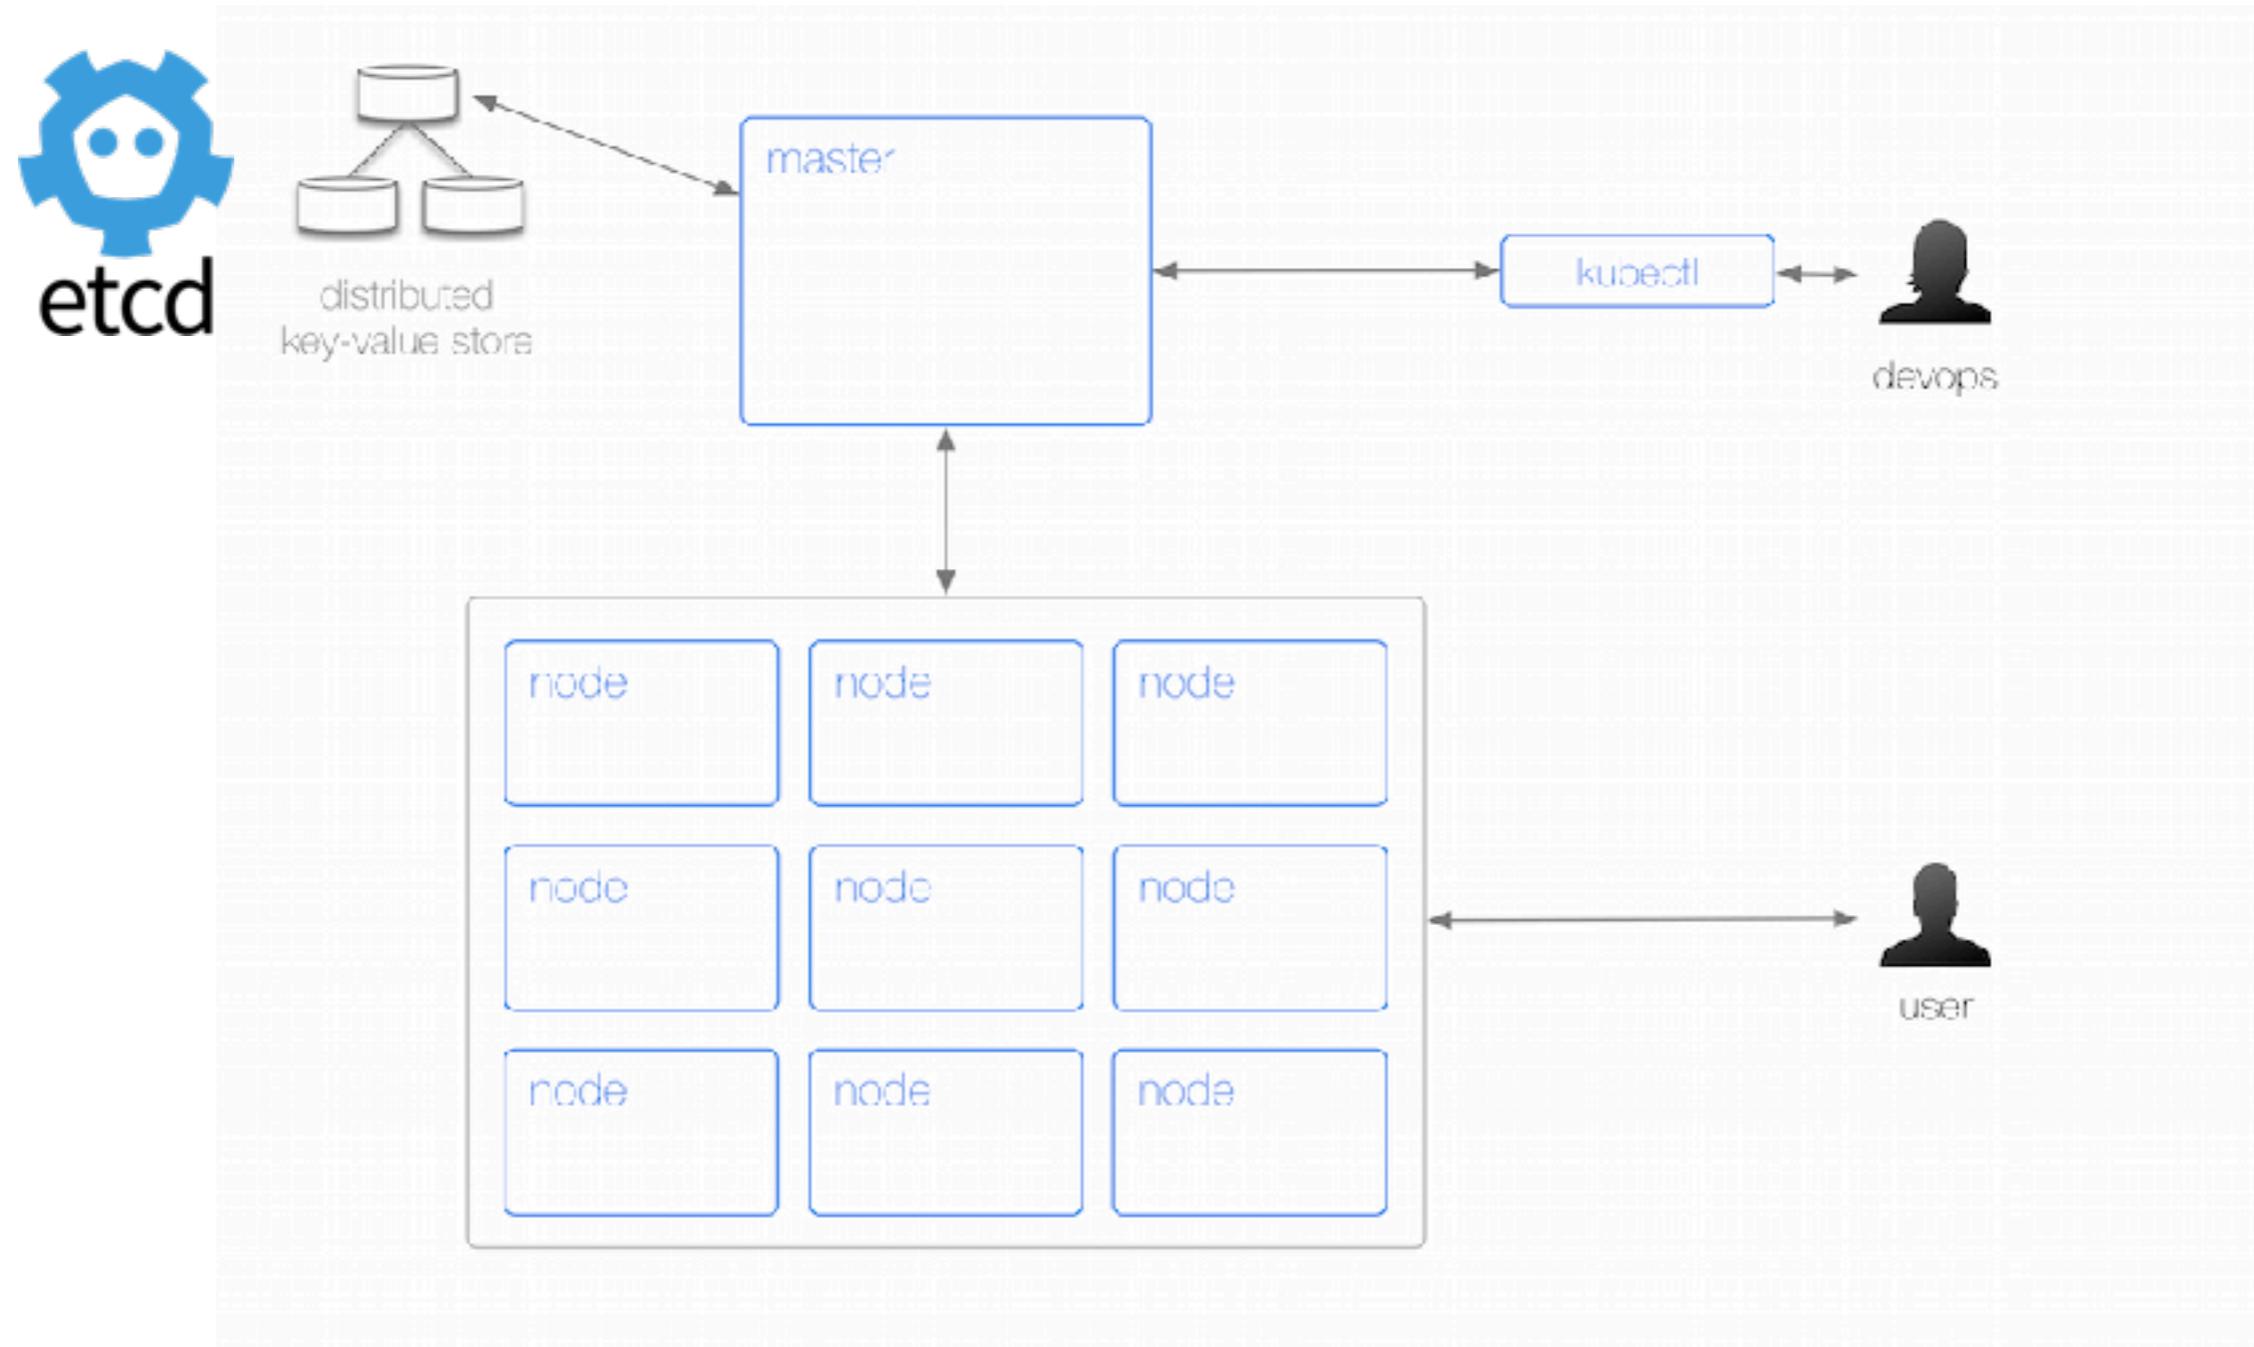
# Basic of Kubernetes



# Kubernetes cluster



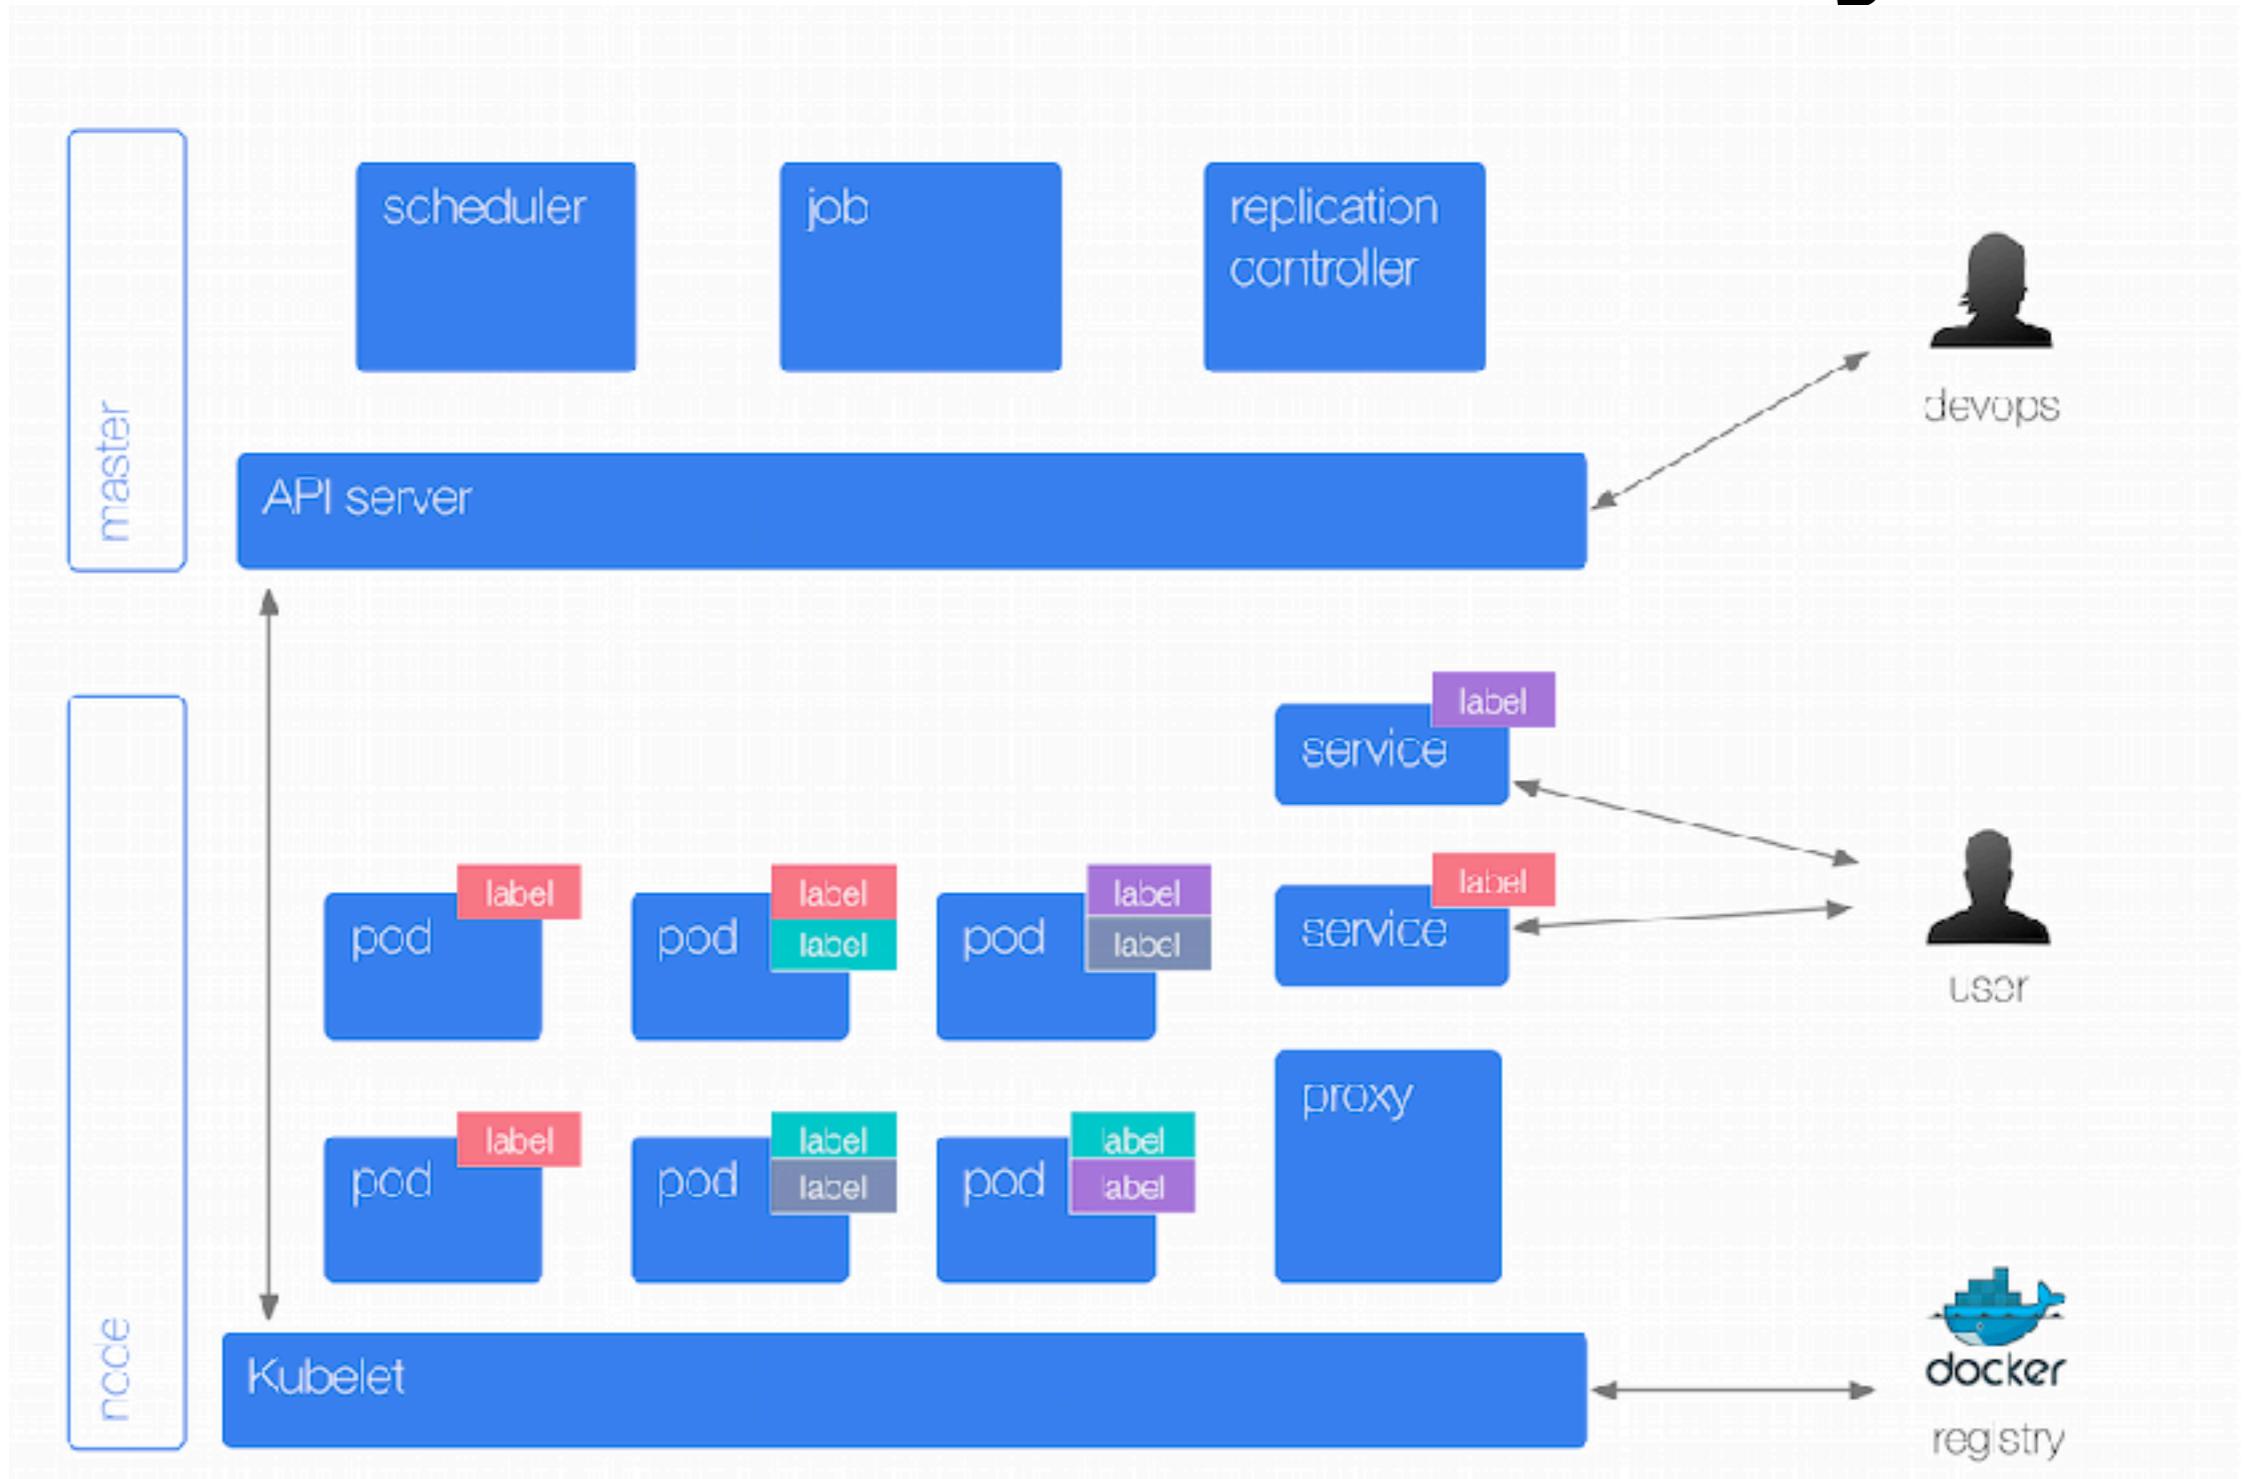
# Kubernetes physical layout



<http://k8s.info/cs.html>



# Kubernetes abstraction layout



<http://k8s.info/cs.html>



# Hello with Kubernetes

file /01-hello/instruction.txt



# Hello with Kubernetes

file /01-hello/instruction2.txt



# Hello Kubernetes

Working with **kubectl** in command line  
Try to run image from Docker hub



# Try to use somkiat/hello

The screenshot shows the Docker Hub interface for the repository `somkiat/hello`. At the top, there's a dark header bar with a search bar, a dashboard icon, and navigation links for Dashboard, Explore, Organizations, Create, and a user profile for `somkiat`. Below the header, the repository name `somkiat/hello` is displayed with a star icon, indicating it's a public repository. A note says "Last pushed: 2 months ago". There are tabs for Repo Info, Tags, Collaborators, Webhooks, and Settings, with the Tags tab currently selected. A table lists three tags: `v3`, `v2`, and `latest`, each with a compressed size of 5 MB and a last updated date of 2 months ago. Each row has a delete icon.

Tag Name	Compressed Size	Last Updated	
v3	5 MB	2 months ago	
v2	5 MB	2 months ago	
latest	5 MB	2 months ago	

<https://hub.docker.com/r/somkiat/hello>



Kubernetes in practice

© 2017 - 2018 Siam Chamnankit Company Limited. All rights reserved.

# Create a container

```
$kubectl run hello --image=somkiat/hello  
--port=8080 --generator=run/v1
```



# Create a container

```
$kubectl run hello --image=somkiat/hello  
--port=8080 --generator=run/v1
```

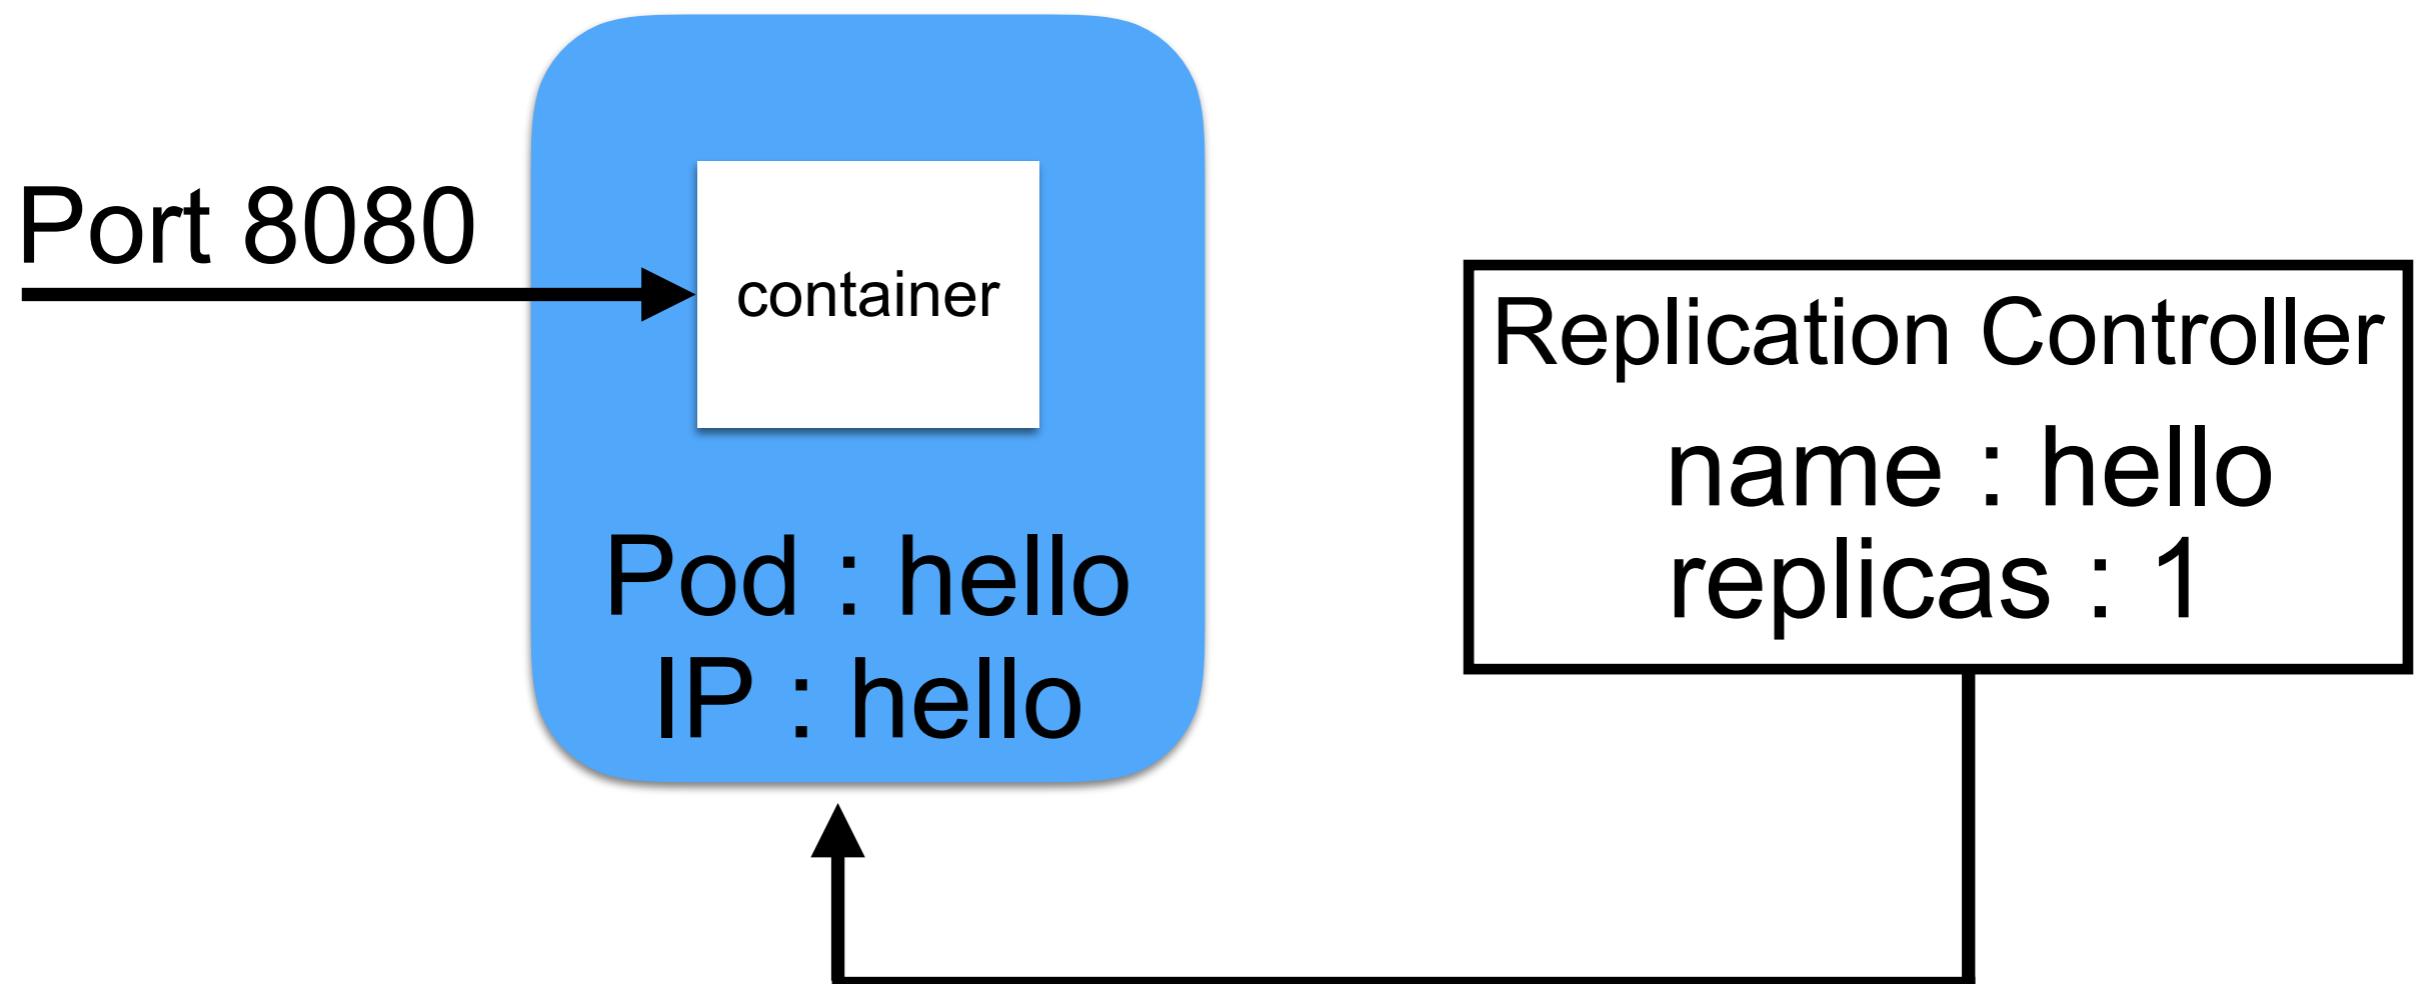


*Create Replication Controller (RC)*

<https://kubernetes.io/docs/reference/kubectl/conventions/#generators>



# Create a container

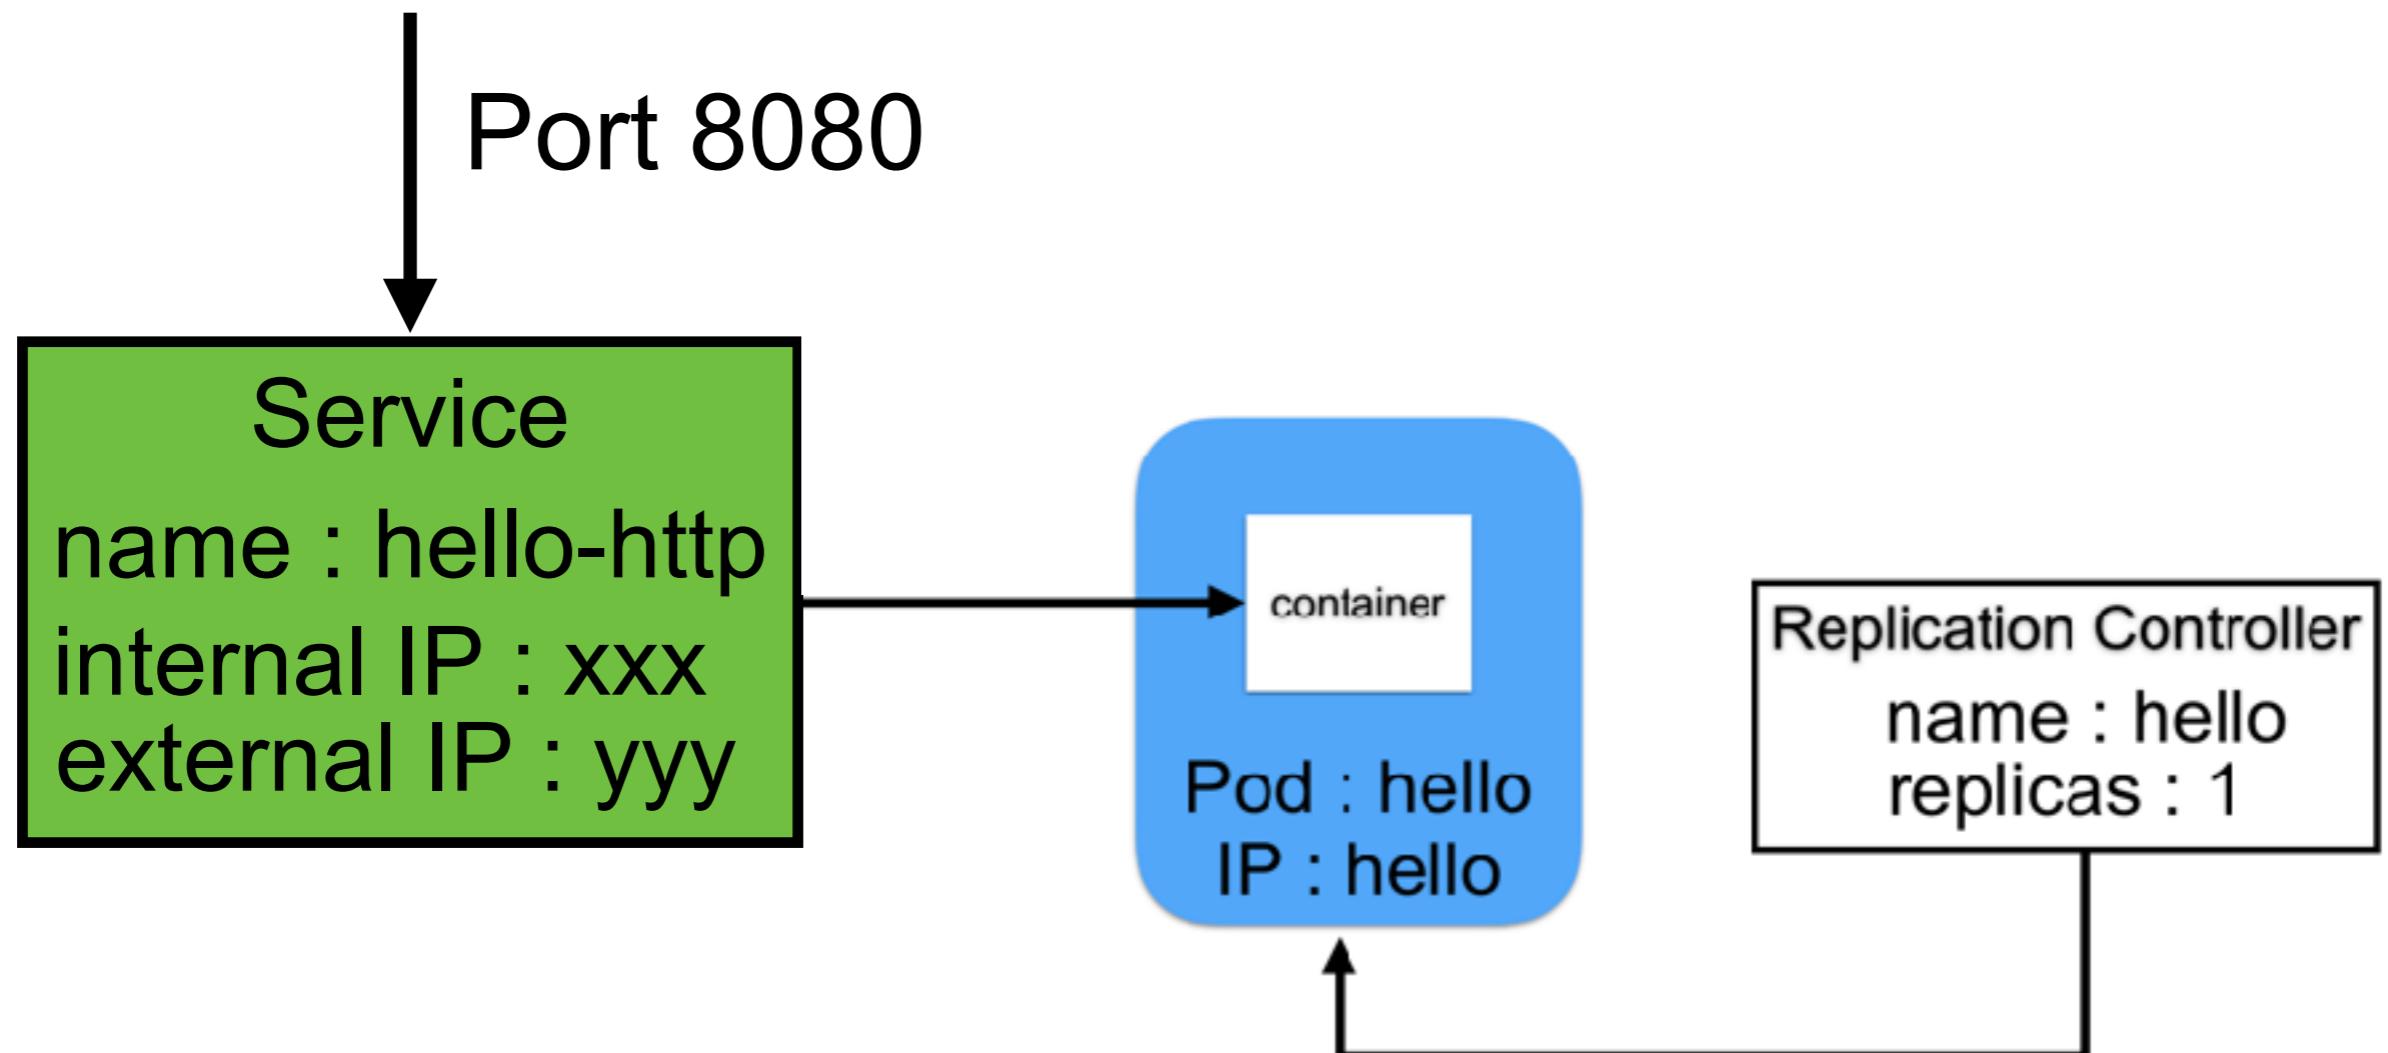


# Expose RC with service

```
$kubectl expose rc hello  
--type=LoadBalancer --name hello-http
```



# Expose RC with service



# Access service with minikube

```
$minikube service hello-http
```

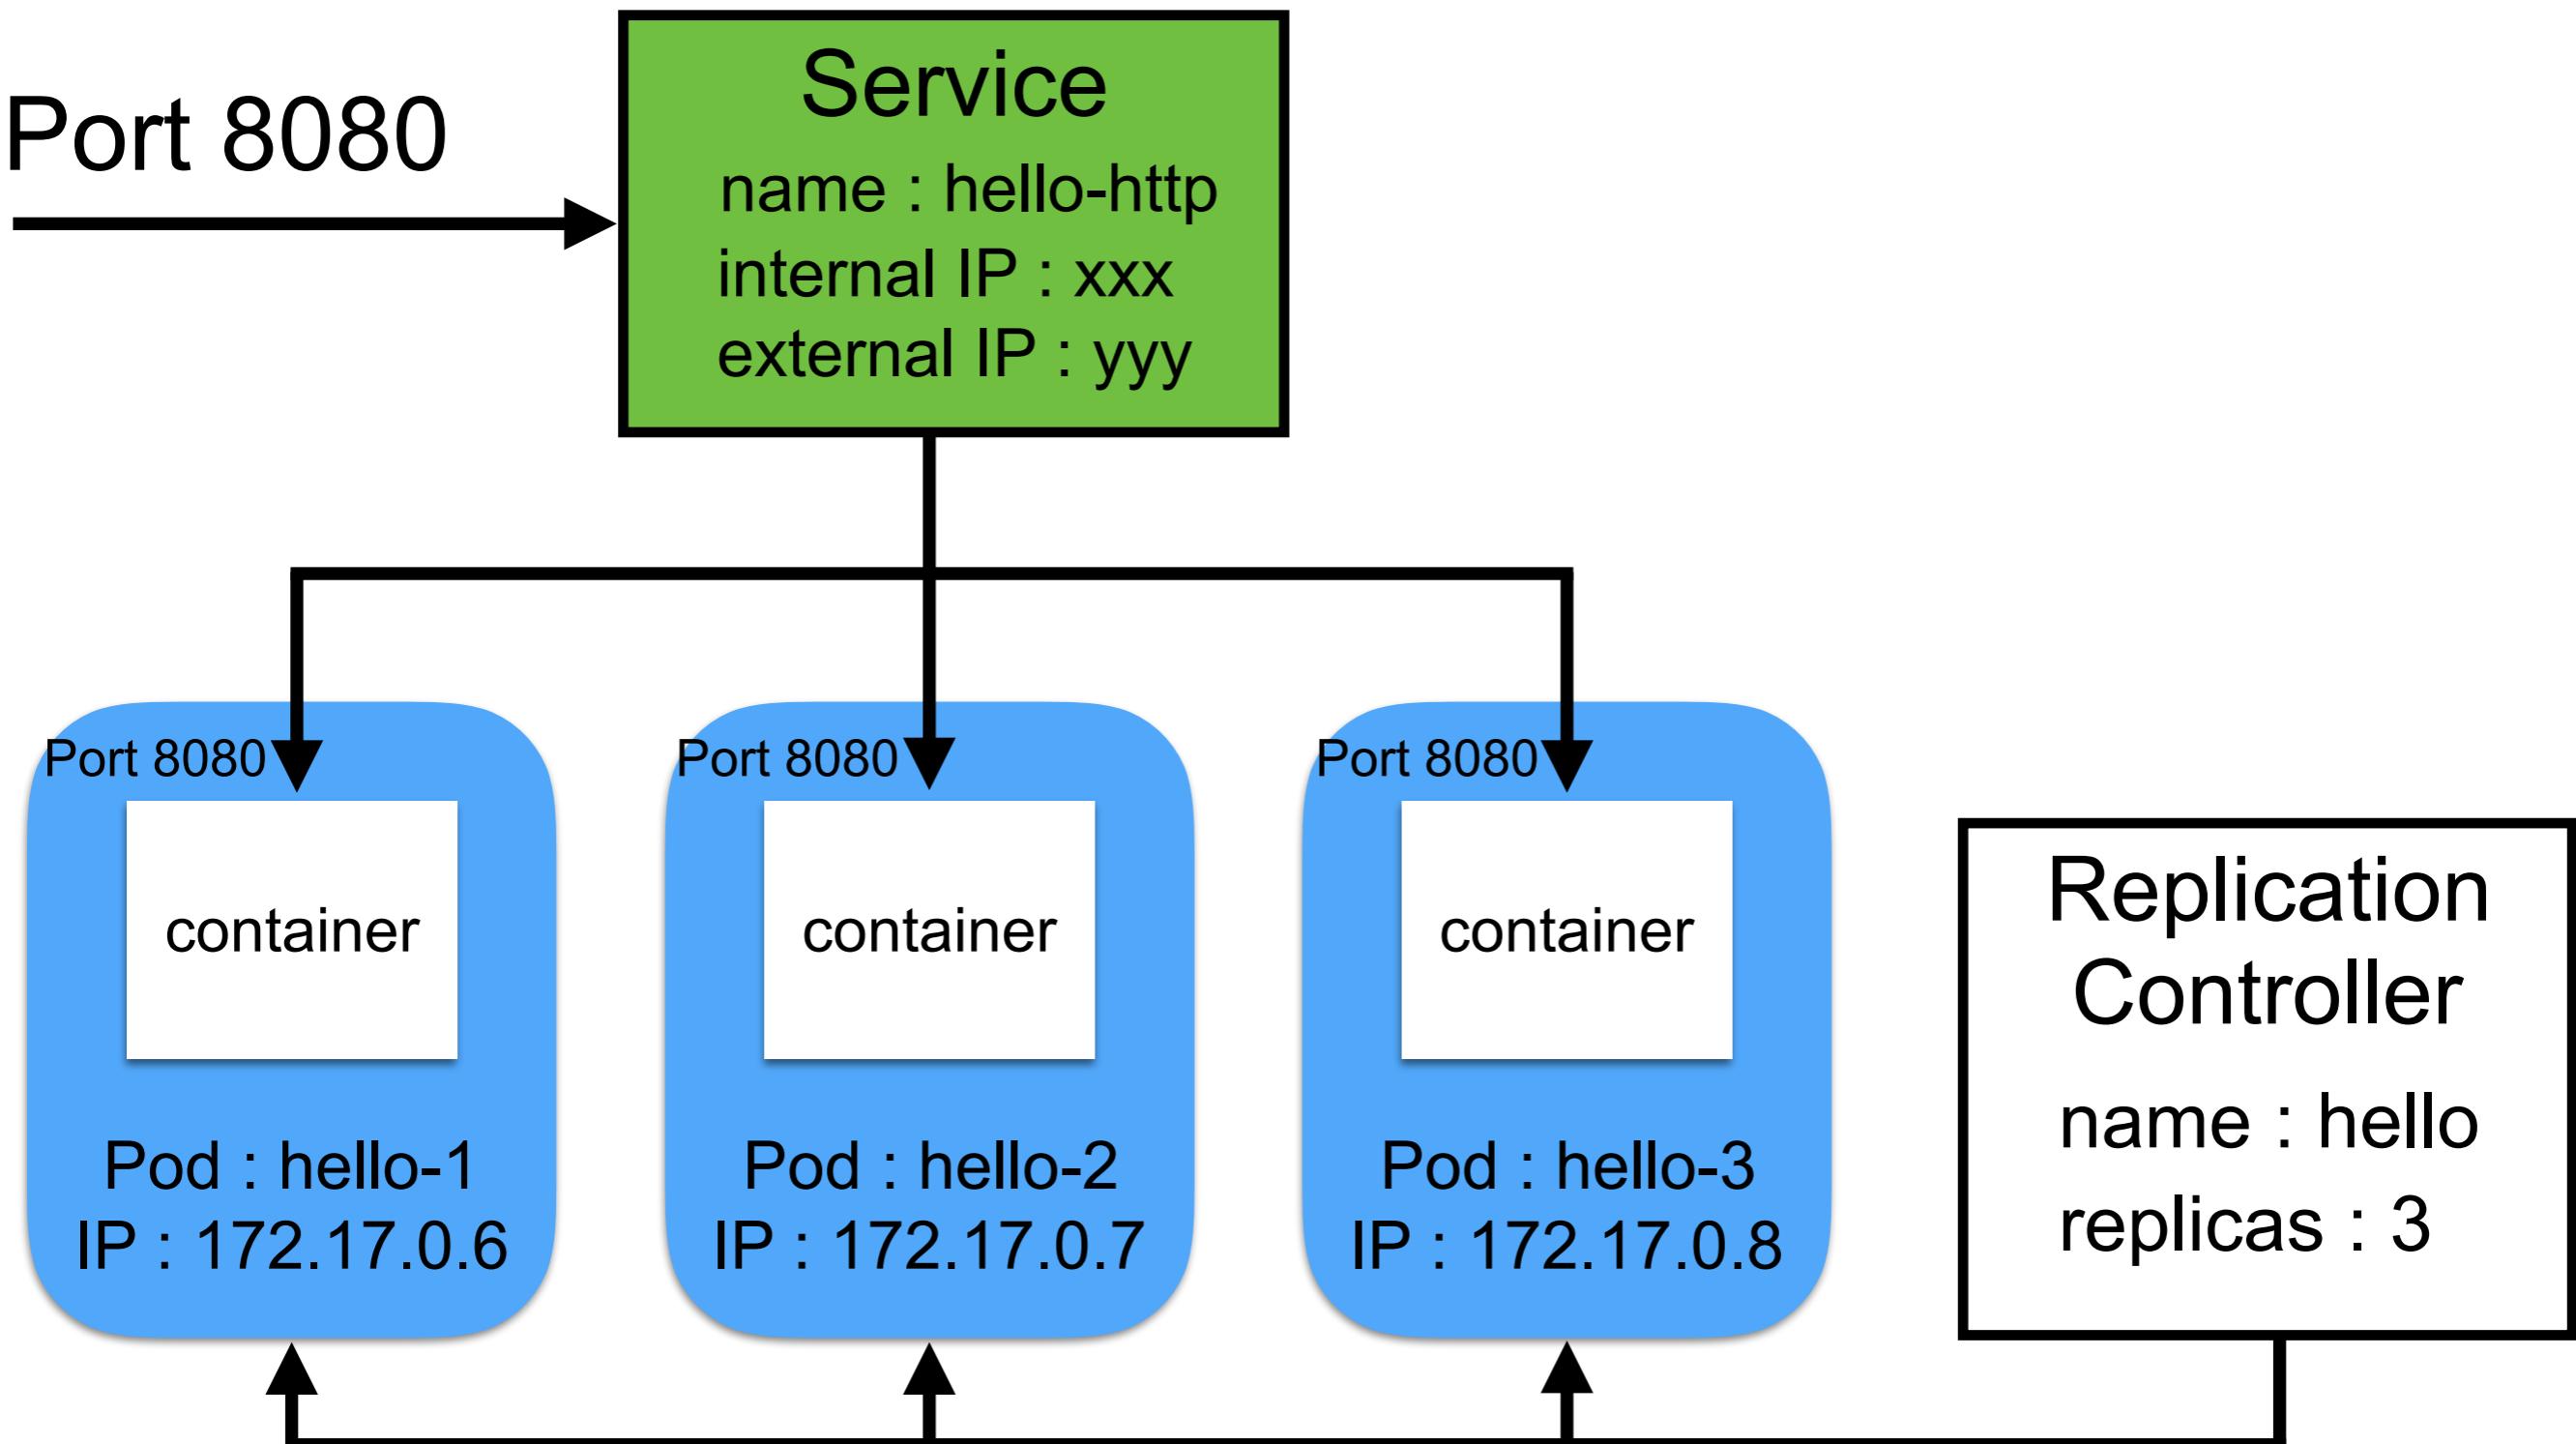


# Scaling the application

```
$ kubectl scale rc hello --replicas=3
```



# Scaling the application



# List of RC

\$kubectl get rc

NAME	DESIRED	CURRENT	READY	AGE
hello	3	3	3	35m



# List of Pods

\$kubectl get pod -o wide

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
hello-lw2p7	1/1	Running	0	29m	172.17.0.8	minikube
hello-pqnpz	1/1	Running	0	36m	172.17.0.6	minikube
hello-wwsnh	1/1	Running	0	29m	172.17.0.7	minikube



# Testing

```
export SERVICE=192.168.99.100:32614
```

```
$curl http://$SERVICE  
Hello, "/" on hello-wwsnh
```

```
$curl http://$SERVICE  
Hello, "/" on hello-pqnzp
```

```
$curl http://$SERVICE  
Hello, "/" on hello-lw2p7
```



# Kubernetes dashboard

## \$minikube dashboard

The screenshot shows the Kubernetes dashboard interface. At the top, there is a navigation bar with a Kubernetes logo, a search bar, and a '+ CREATE' button. Below the navigation bar, a blue header bar displays the text 'Overview'. On the left side, there is a sidebar with 'Cluster' and 'Namespaces' sections. Under 'Namespaces', 'default' is selected. Below the namespaces section, there is a 'Workloads' section with sub-options: 'Cron Jobs', 'Daemon Sets', 'Deployments', 'Jobs', and 'Pods'. The main content area is titled 'Workloads Statuses' and contains two green circular charts, both showing '100.00%' for 'Pods' and 'Replication Controllers'. Below this, there is a table titled 'Pods' with the following data:

Name	Node	Status	Restarts	Age	⋮
hello-lw2p7	minikube	Running	0	53 minutes	⋮
hello-wwsnh	minikube	Running	0	53 minutes	⋮
hello-pqnpz	minikube	Running	0	59 minutes	⋮



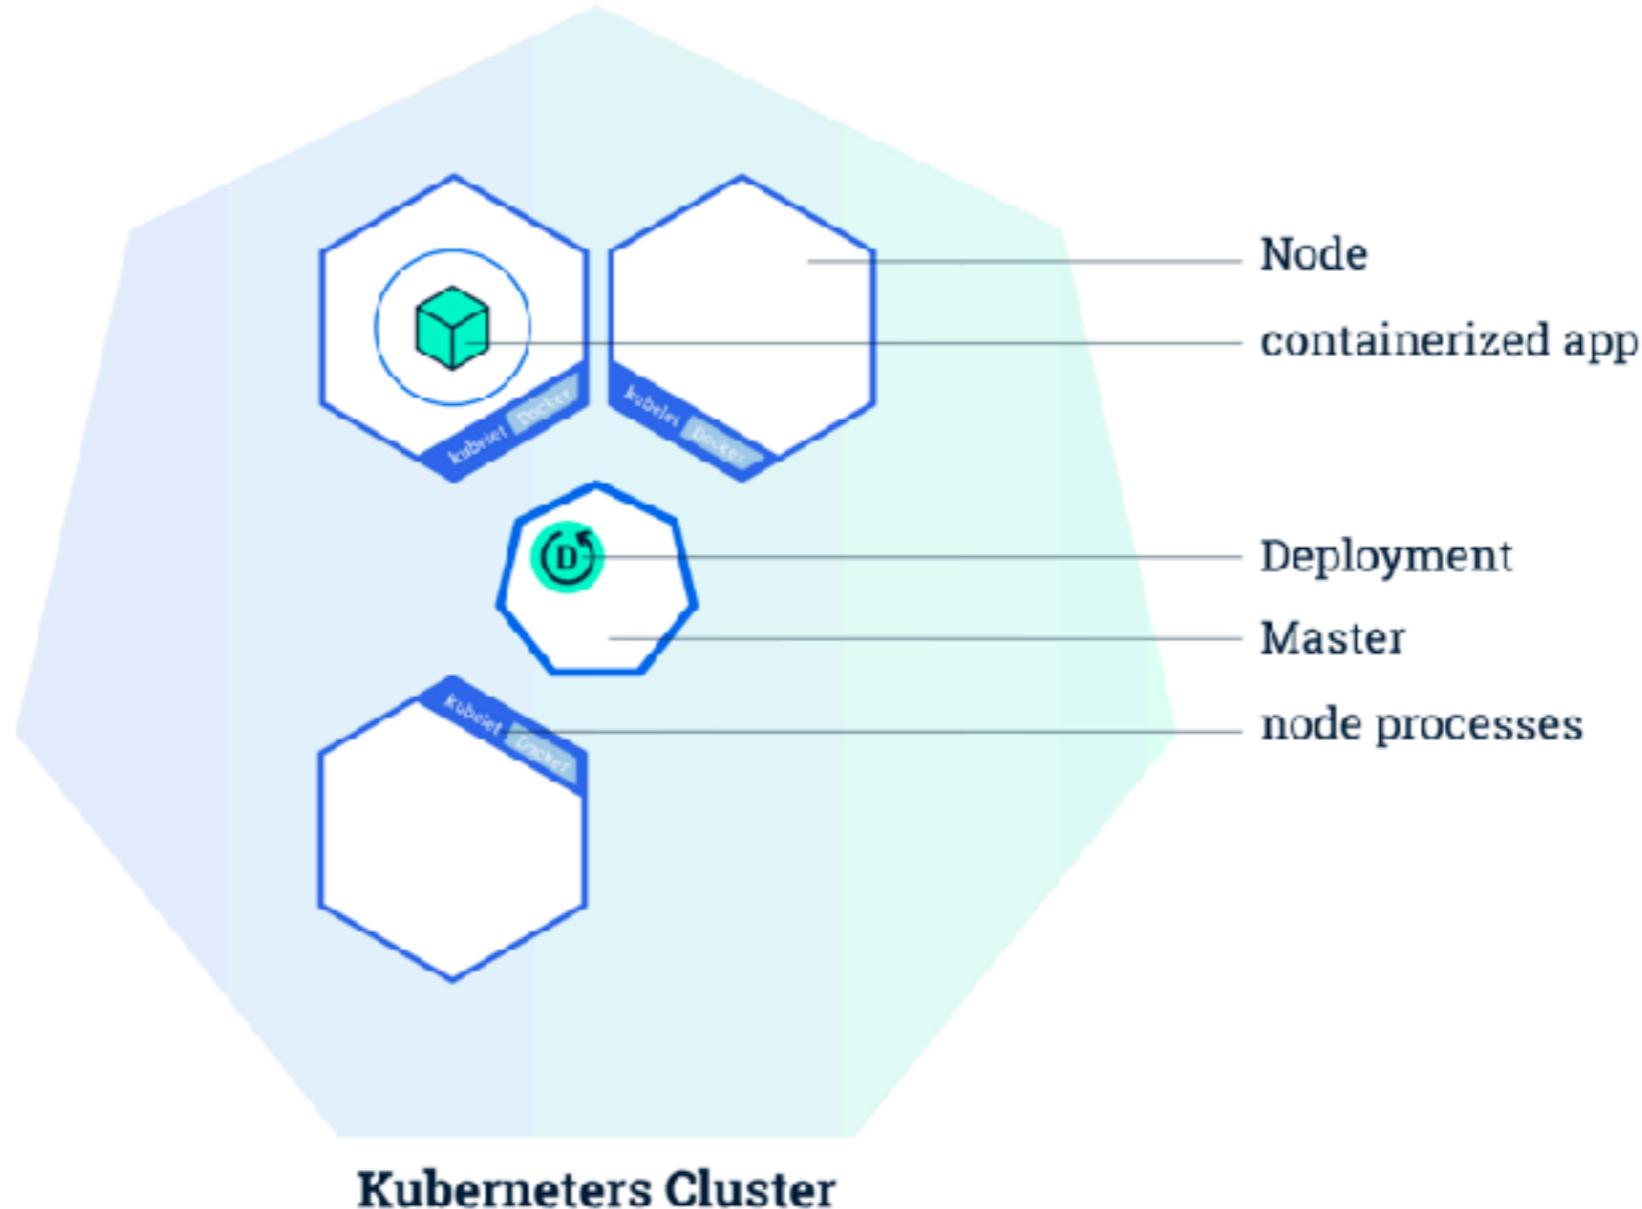
# Welcome to Kubernetes



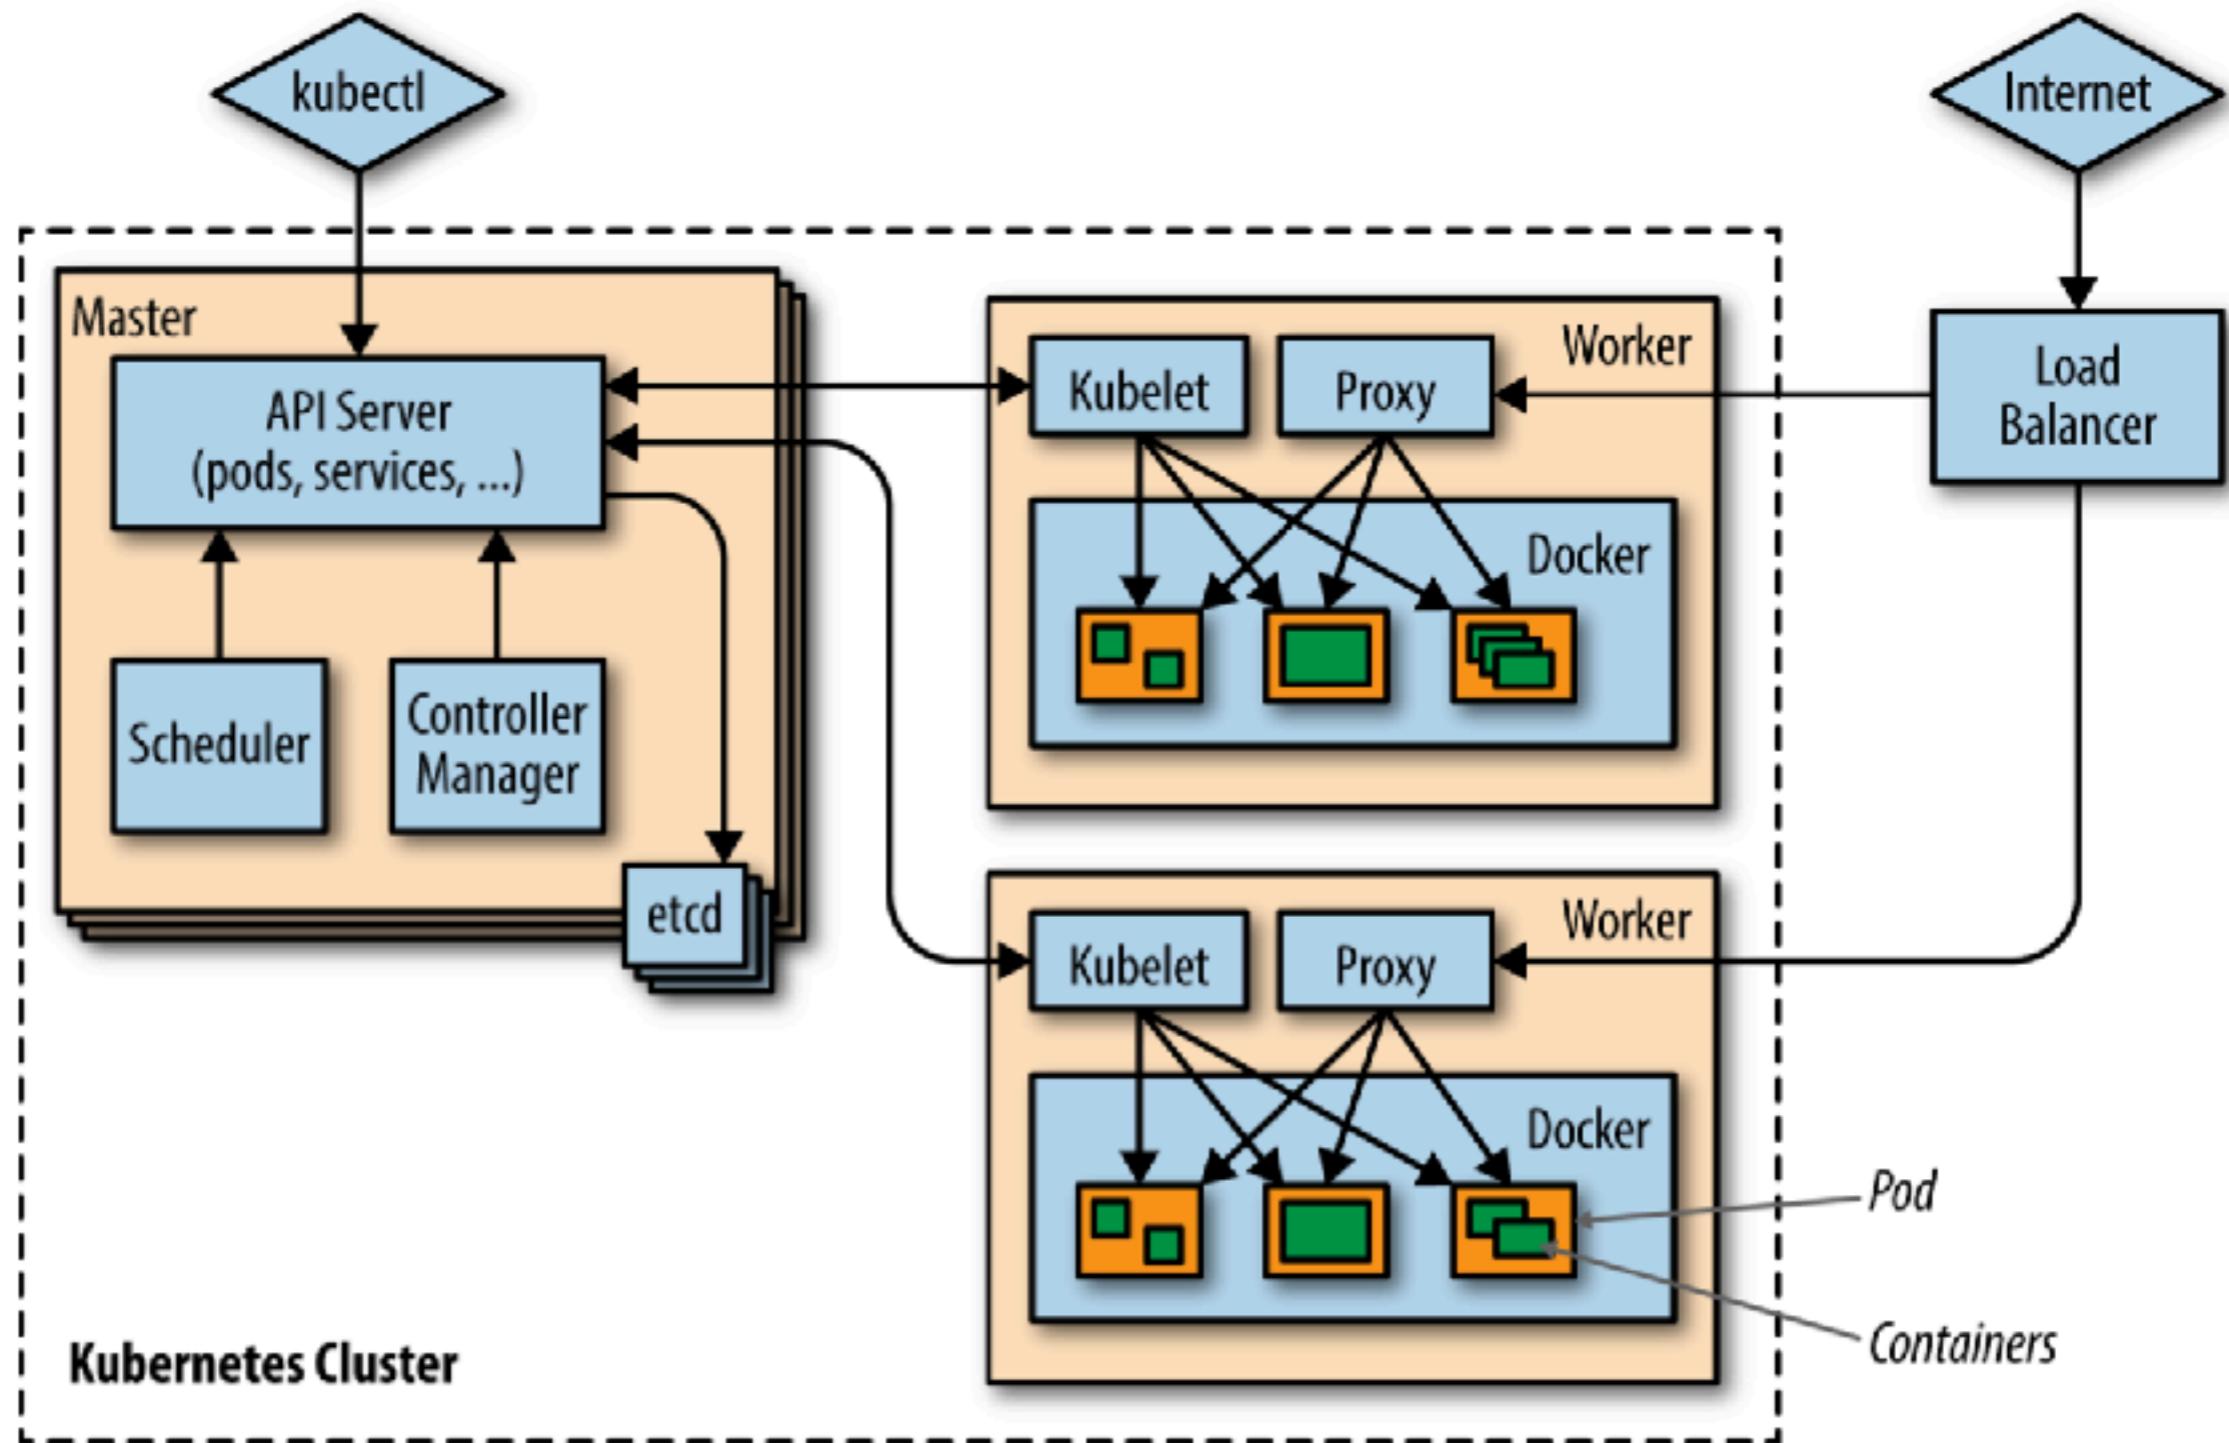
# Kubernetes Architecture



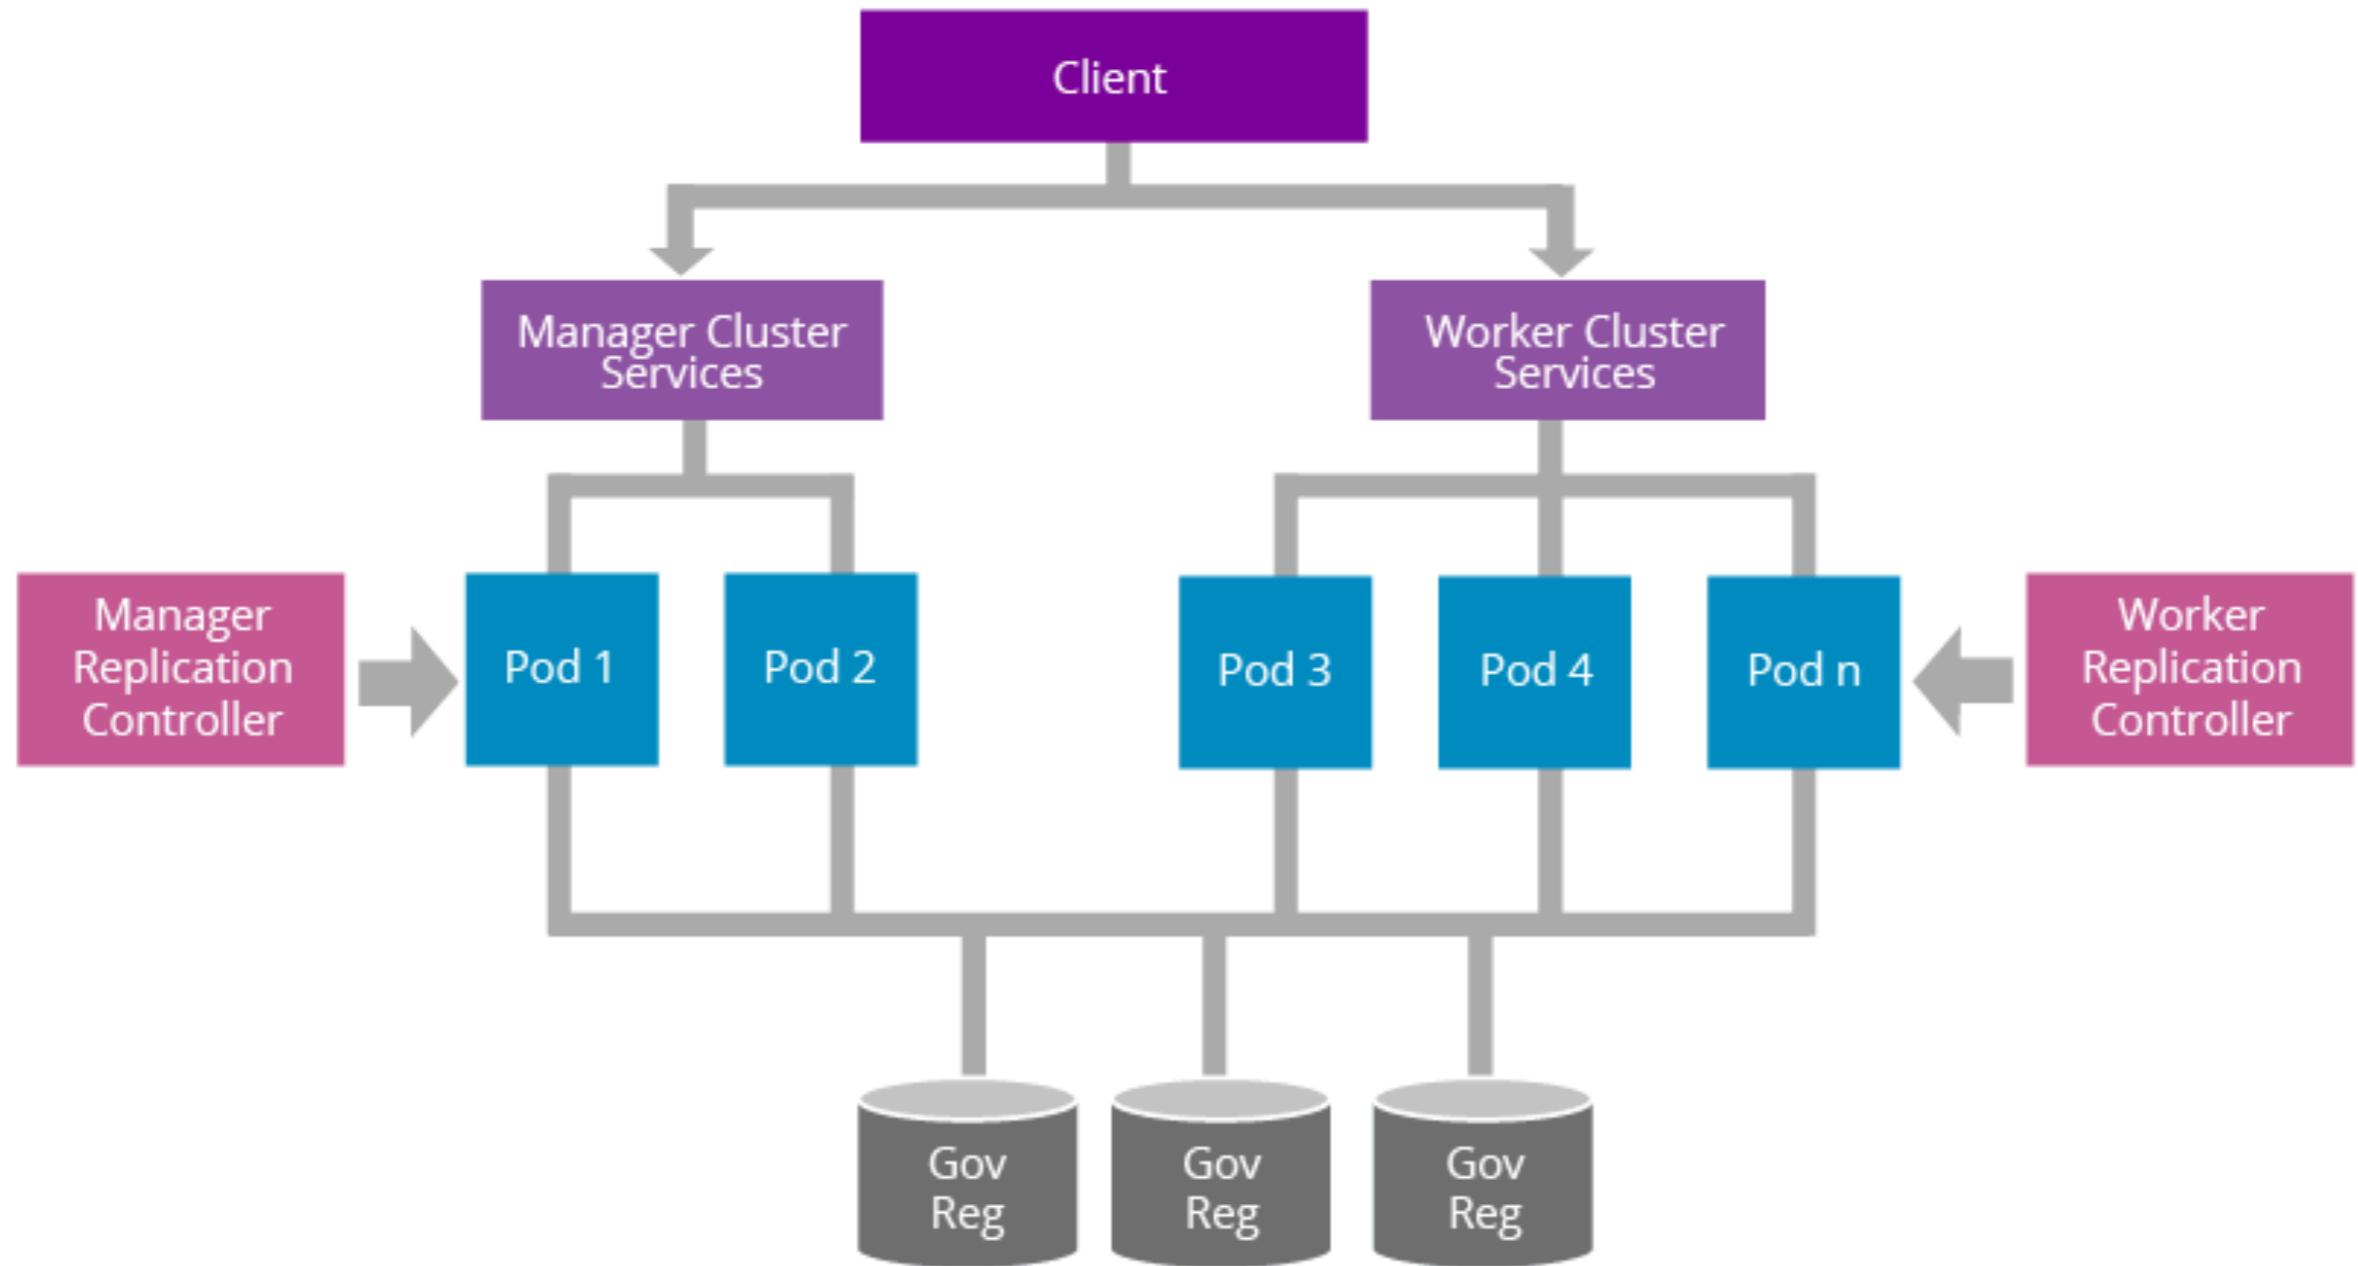
# Kubernetes cluster



# Kubernetes architecture



# Kubernetes architecture



<https://wso2.com/whitepapers/a-reference-architecture-for-deploying-wso2-middleware-on-kubernetes/>



# Key features

<https://kubernetes.io/>



# Key features

- Automatic binpacking
- Horizontal Pod Autoscaler (HPA)
- Automated rollouts and rollbacks
- Storage orchestration



# Key features

Self-healing

Service discovery and Load Balancing (LB)

Secret and config management

Batch execution



# Automatic binpacking



# Automatic binpacking

Limit of CPU/Memory can define on Pods  
Schedule will select a node that have enough resources



# Automatic binpacking

When reach the Memory limit ?

1. Kill current Pod
2. If Pod have restart flag, try to create in other node



# Automatic binpacking

When reach the CPU limit ?

1. Schedule not kill Pod
2. Schedule waiting it back to normal state



# Automatic binpacking

Try to check resources of node

**\$kubectl describe node**

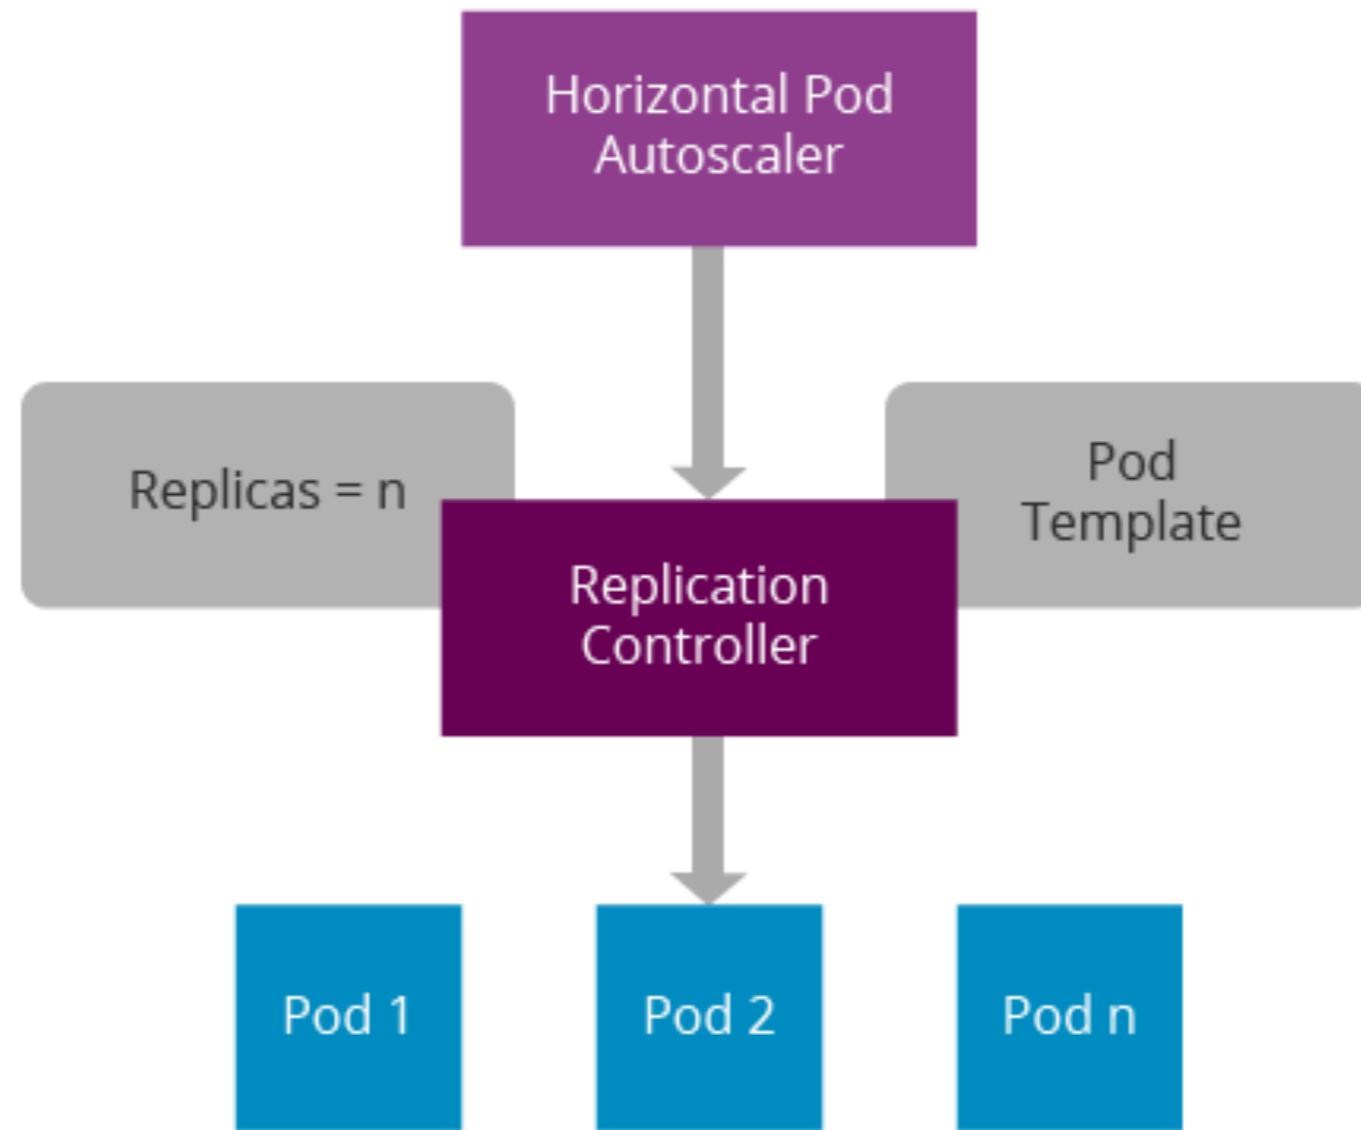
CPU Requests	CPU Limits	Memory Requests	Memory Limits
5m (0%)	0 (0%)	50Mi (2%)	0 (0%)
260m (13%)	0 (0%)	110Mi (5%)	170Mi (8%)
0 (0%)	0 (0%)	0 (0%)	0 (0%)
0 (0%)	0 (0%)	0 (0%)	0 (0%)



# **Horizontal Pod Autoscaler (HPA)**



# Horizontal Pod Autoscaler



# Horizontal Pod Autoscaler

Monitor workload on Pods (based on CPU)  
and automatic scaling-up application

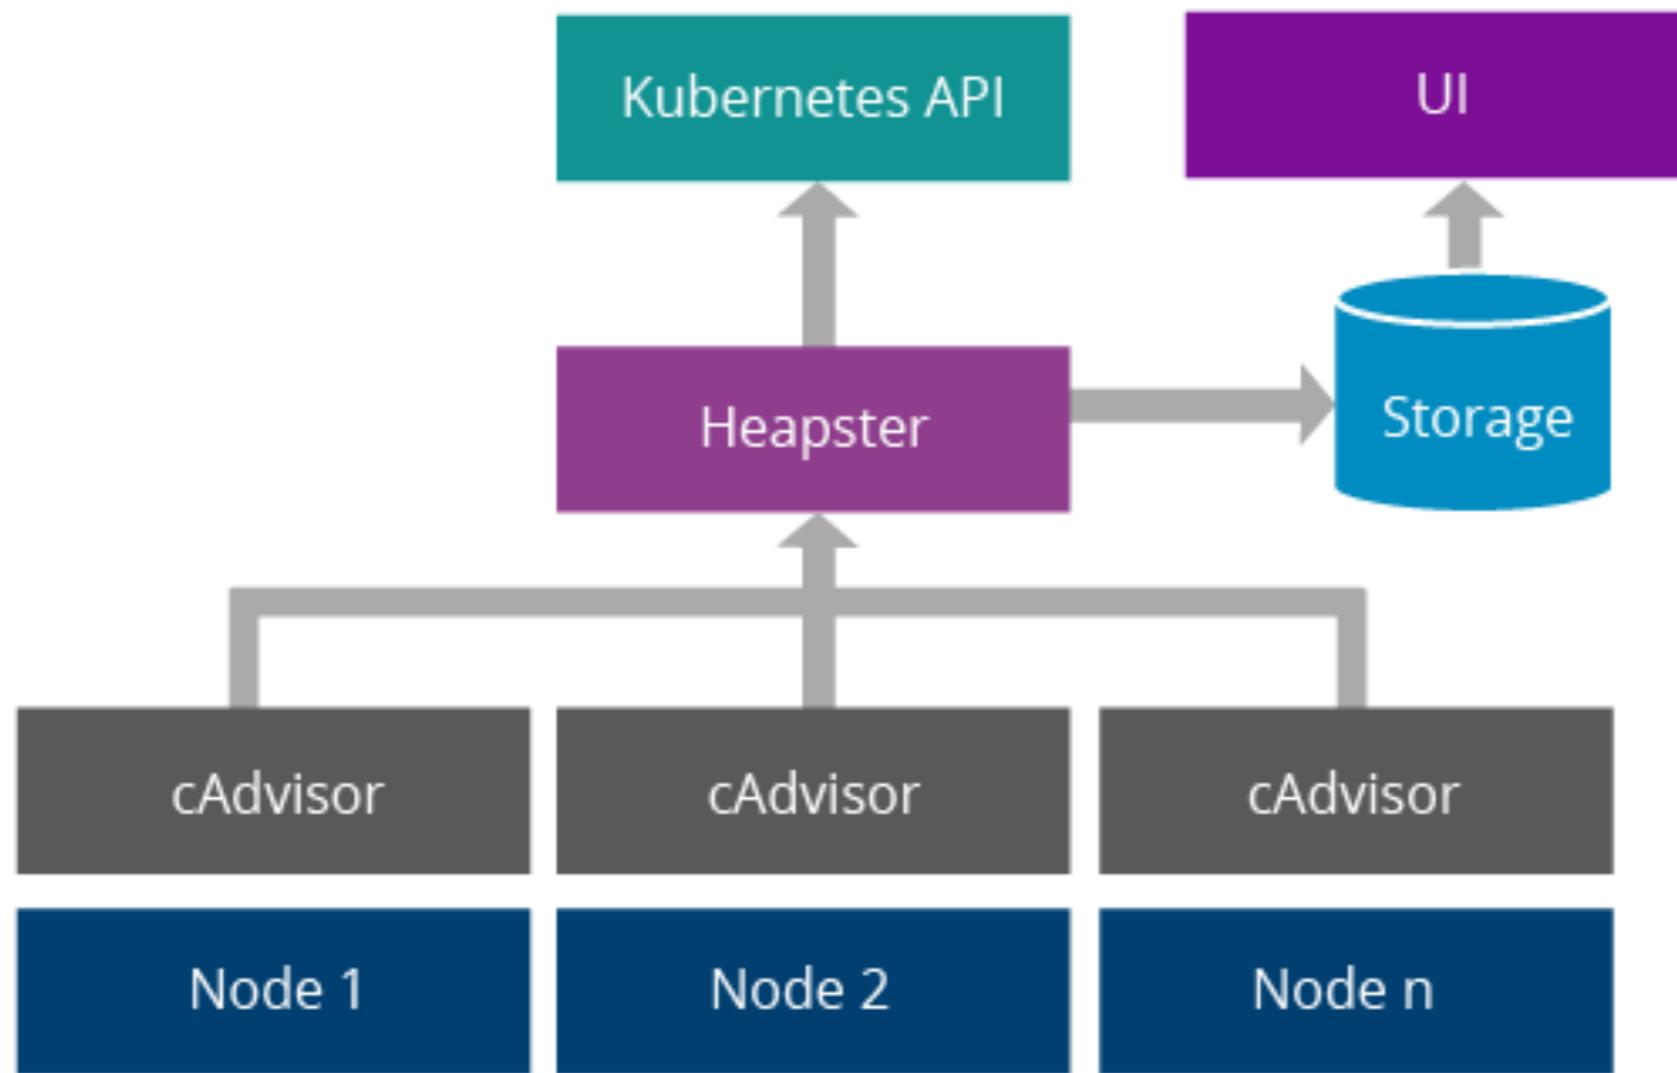
Scaling-up and down for your app need !!



# Resource usage monitoring



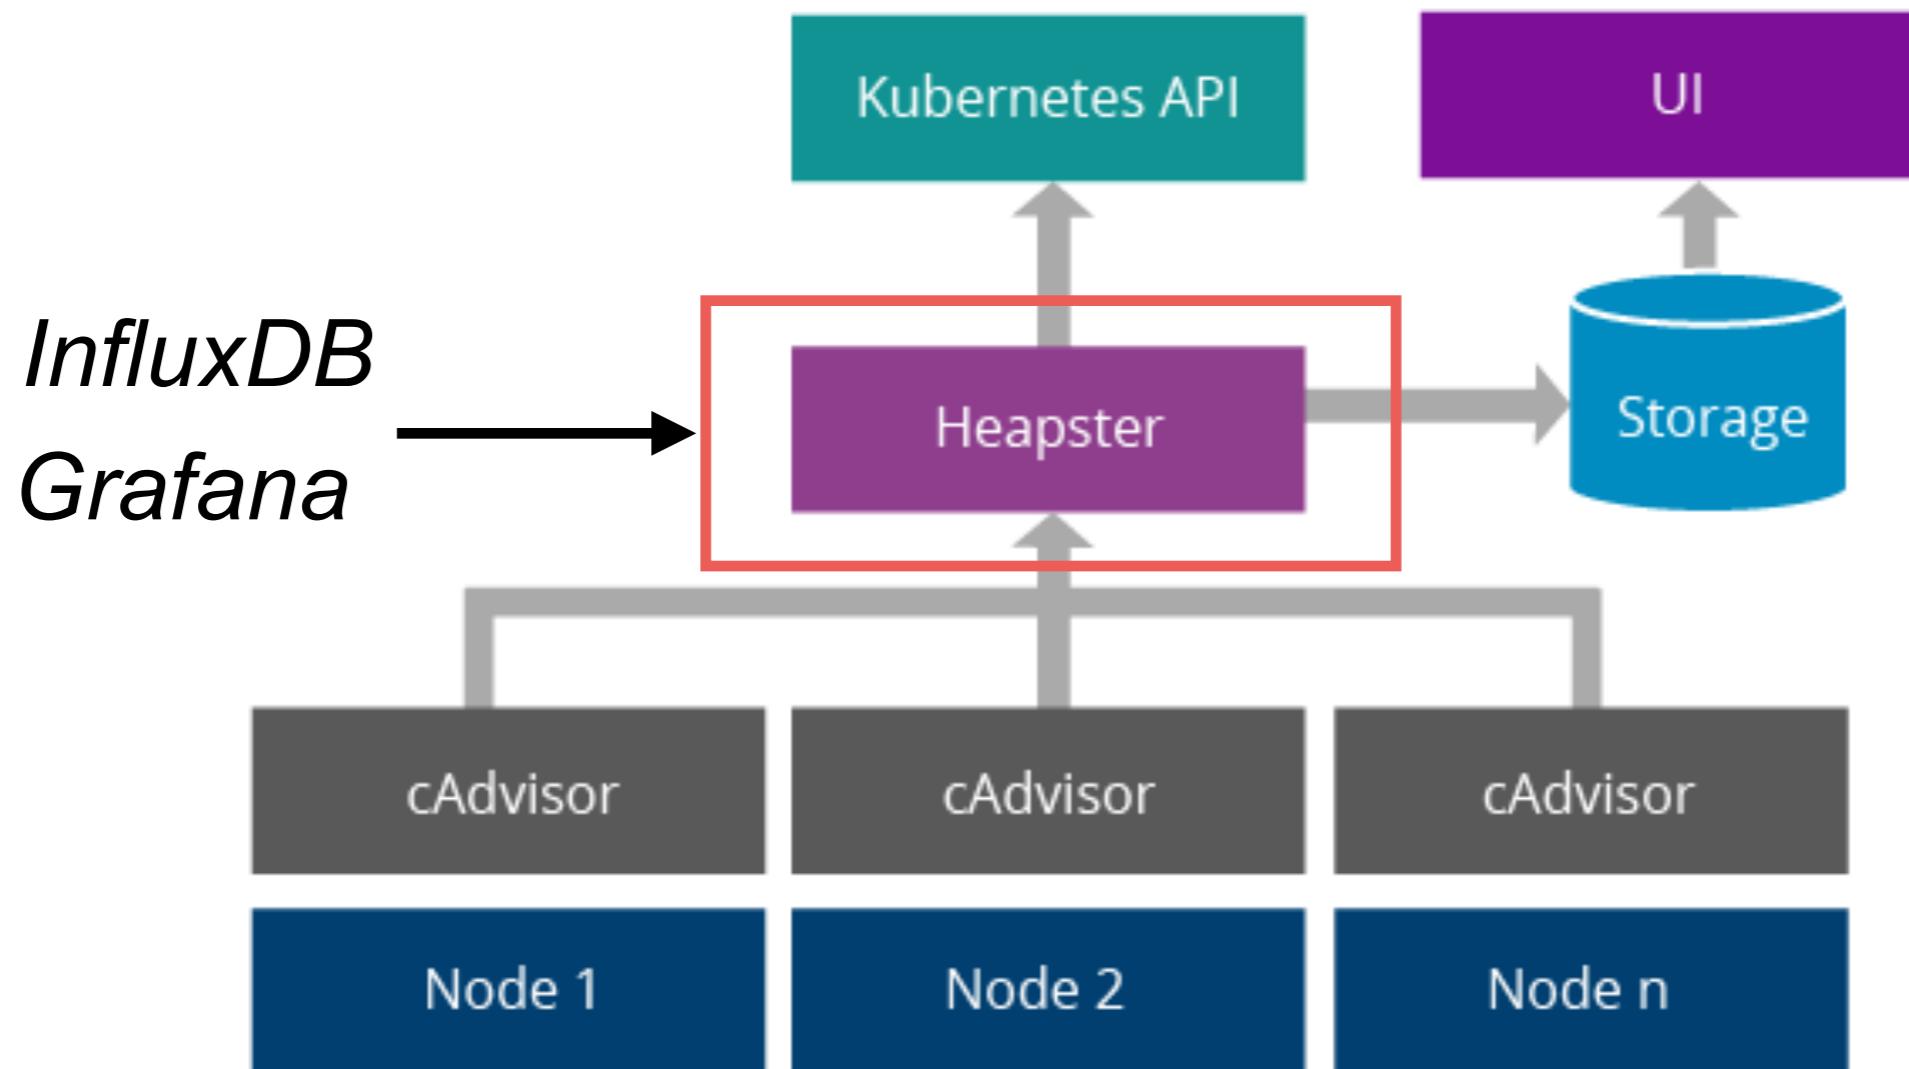
# Resource usage monitoring



<https://github.com/kubernetes/heapster>



# Resource usage monitoring



<https://github.com/kubernetes/heapster>



# Heapster deprecation timeline

Kubernetes	Action	Policy and Support
v 1.11	Initial deprecation	No new feature + fix bugs
v 1.12	Setup removal	The optional to install heapster in Kubernetes
v 1.13	Removal	No new fix bug, move to retired organization

<https://github.com/kubernetes/heapster/blob/master/docs/deprecation.md>



# Try Kubernetes Metric Server

Starting from Kubernetes 1.8

Resources usage metrics => CPU and memory

Use via Metric API and Metric server

<https://github.com/kubernetes-incubator/metrics-server>

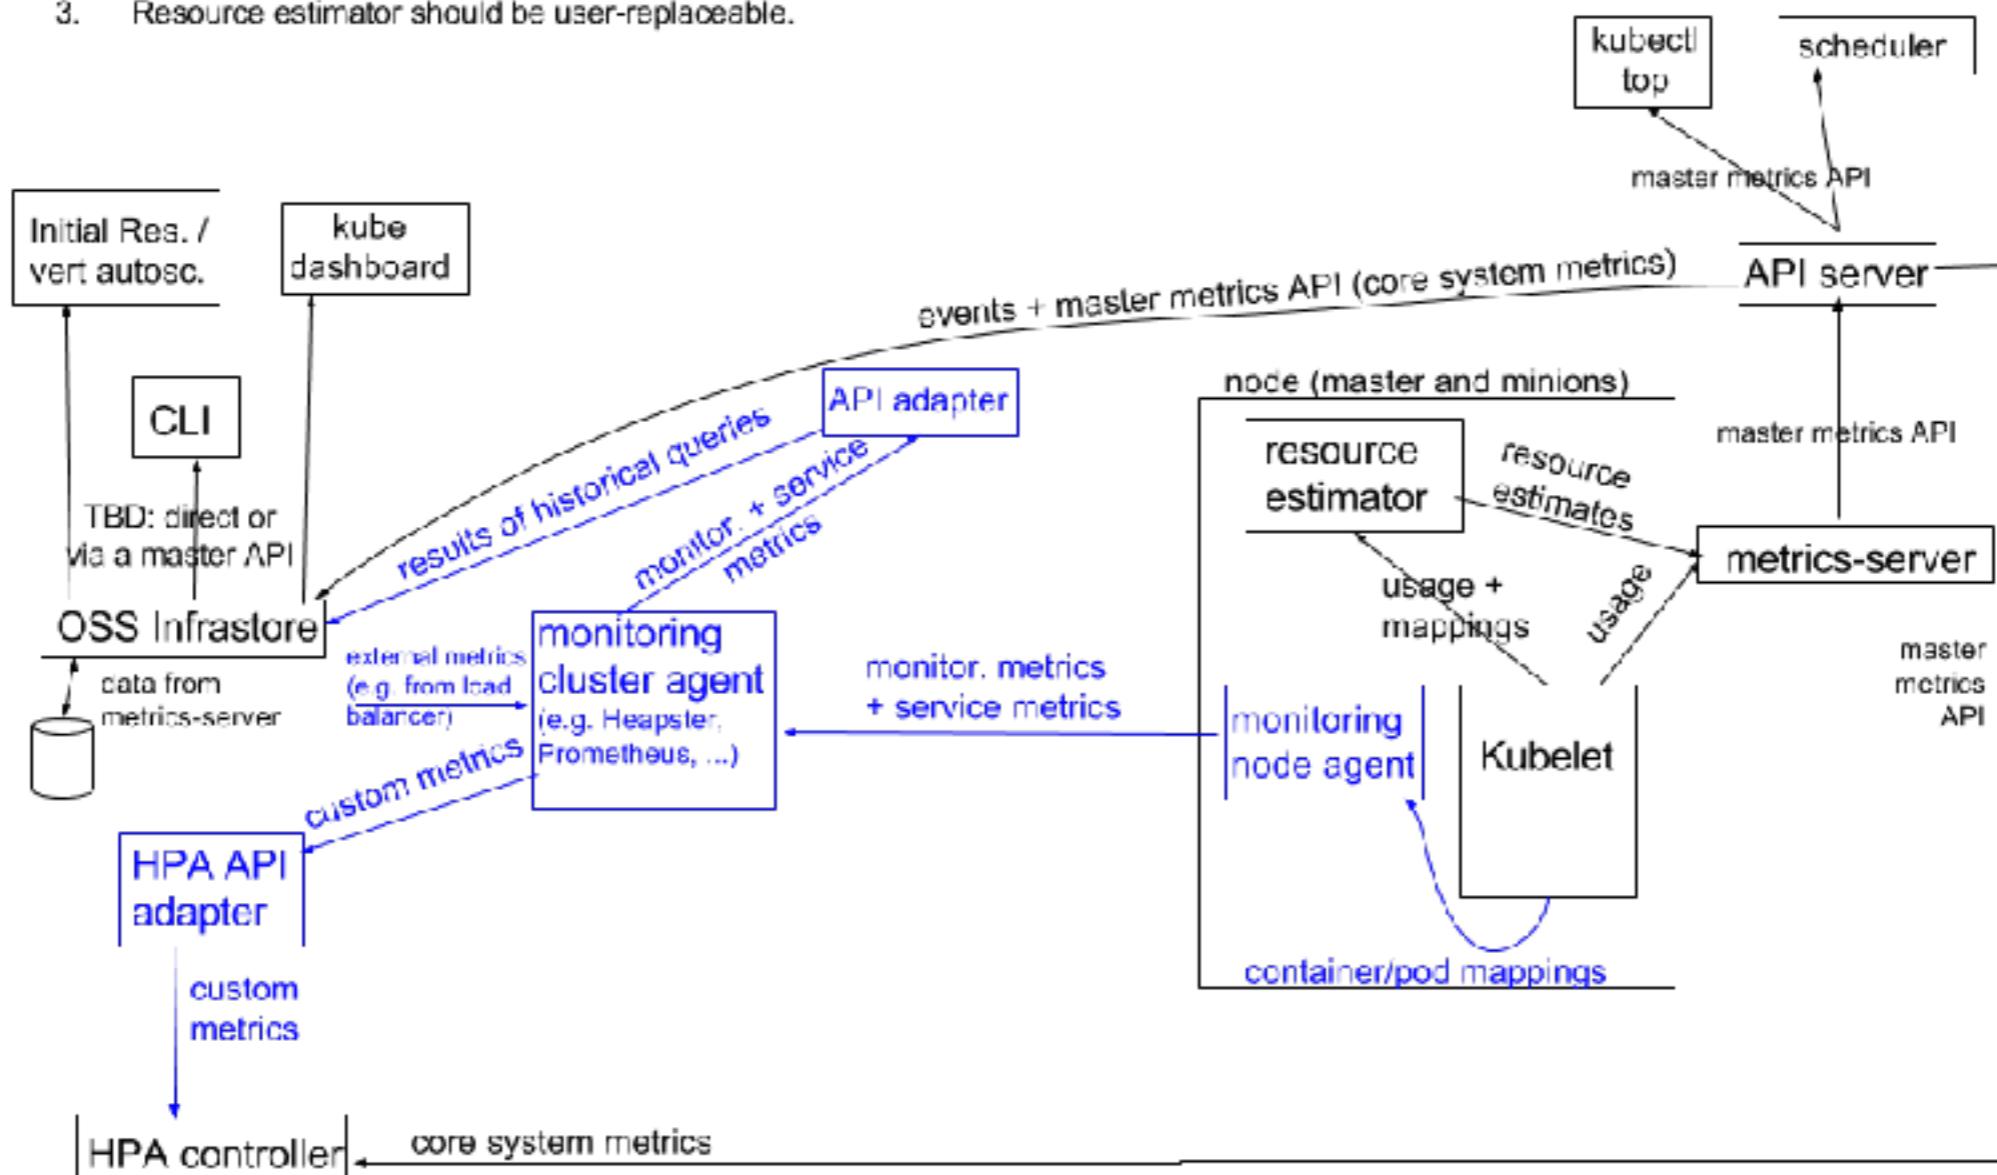


# Architecture diagram

Monitoring architecture proposal: OSS  
(arrows show direction of metrics flow)

## Notes

1. Arrows show direction of metrics flow.
2. Monitoring pipeline is in blue. It is user-supplied and optional.
3. Resource estimator should be user-replaceable.



[https://github.com/kubernetes/community/blob/master/contributors/design-proposals/instrumentation/monitoring\\_architecture.md](https://github.com/kubernetes/community/blob/master/contributors/design-proposals/instrumentation/monitoring_architecture.md)

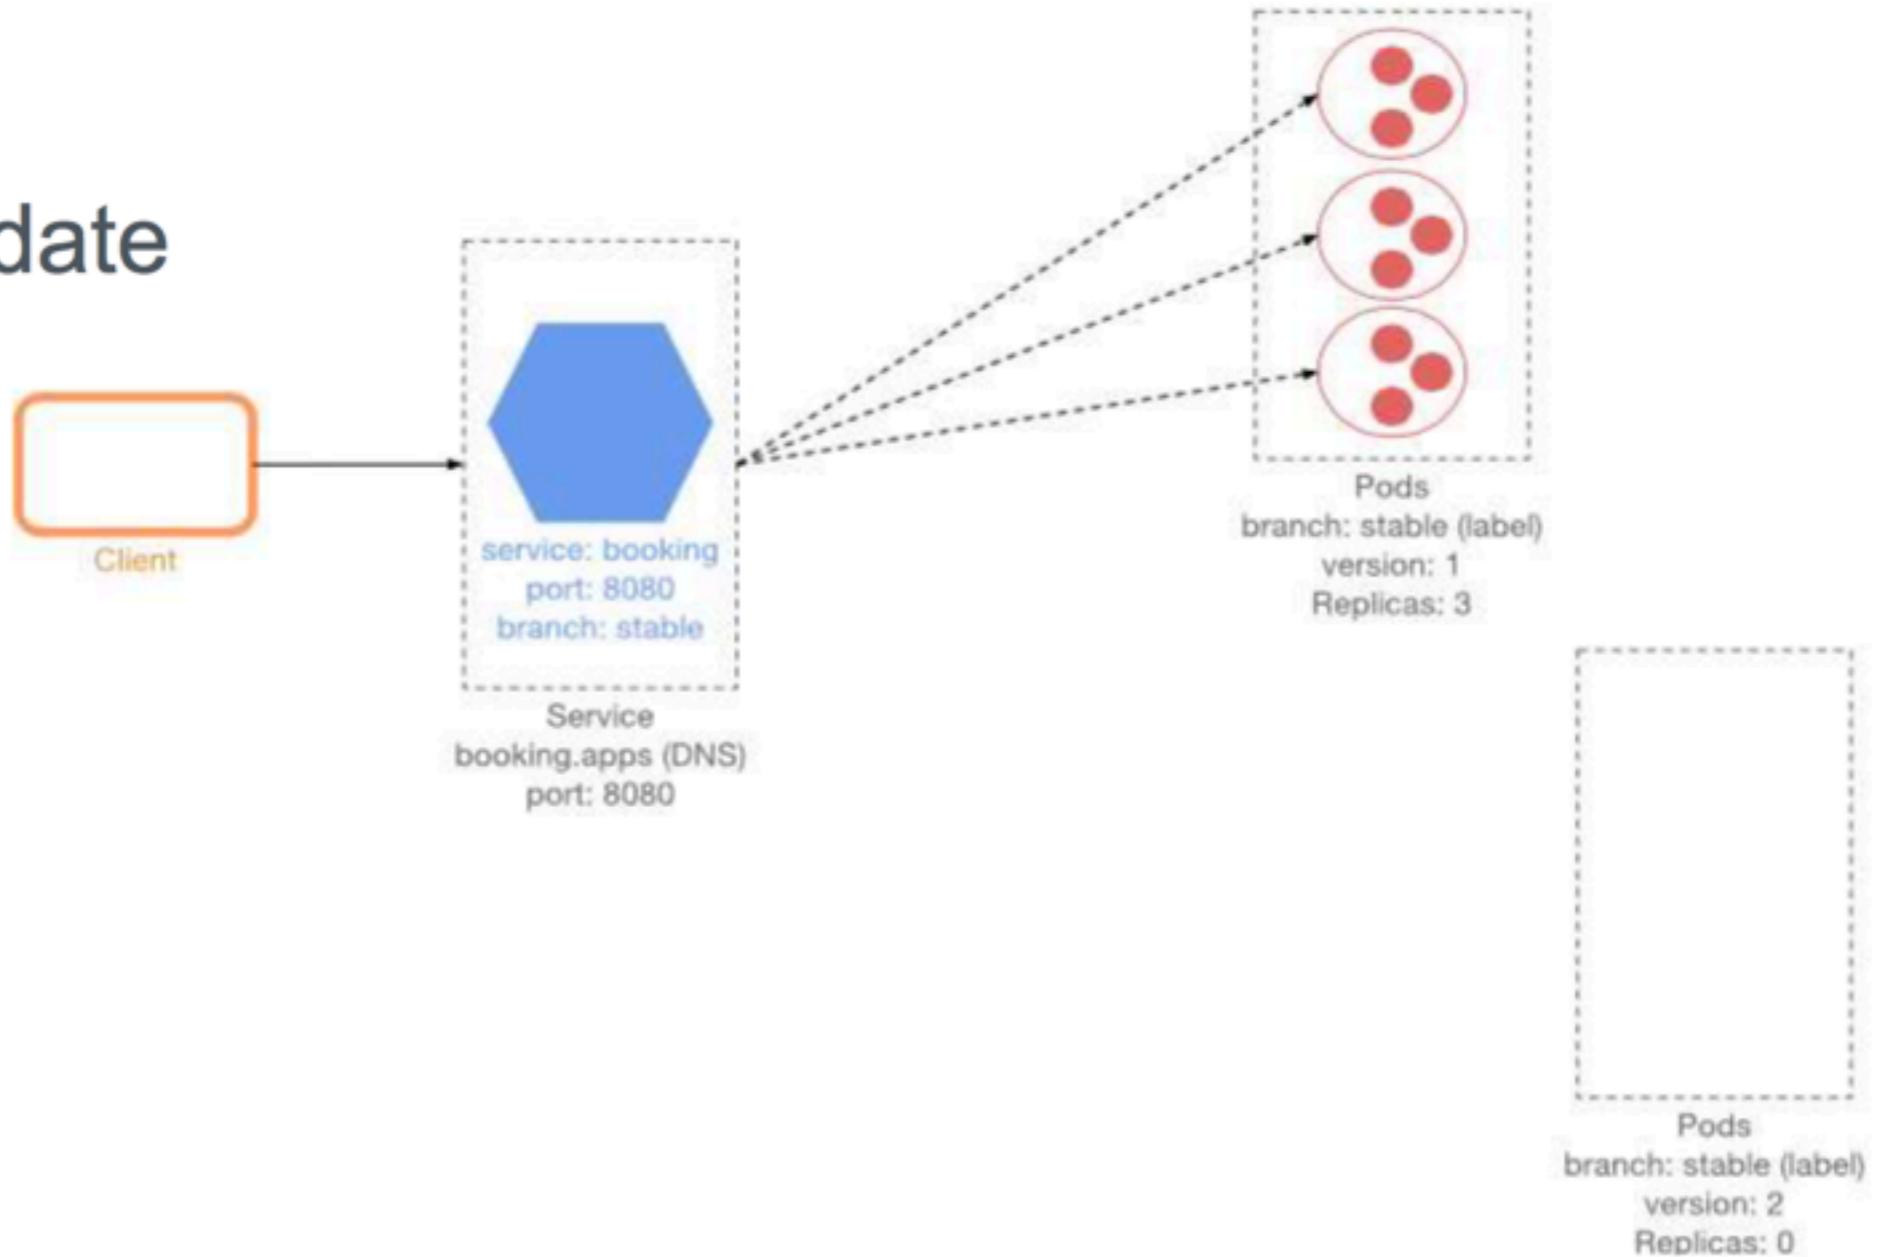


# Automated rollouts & rollbacks



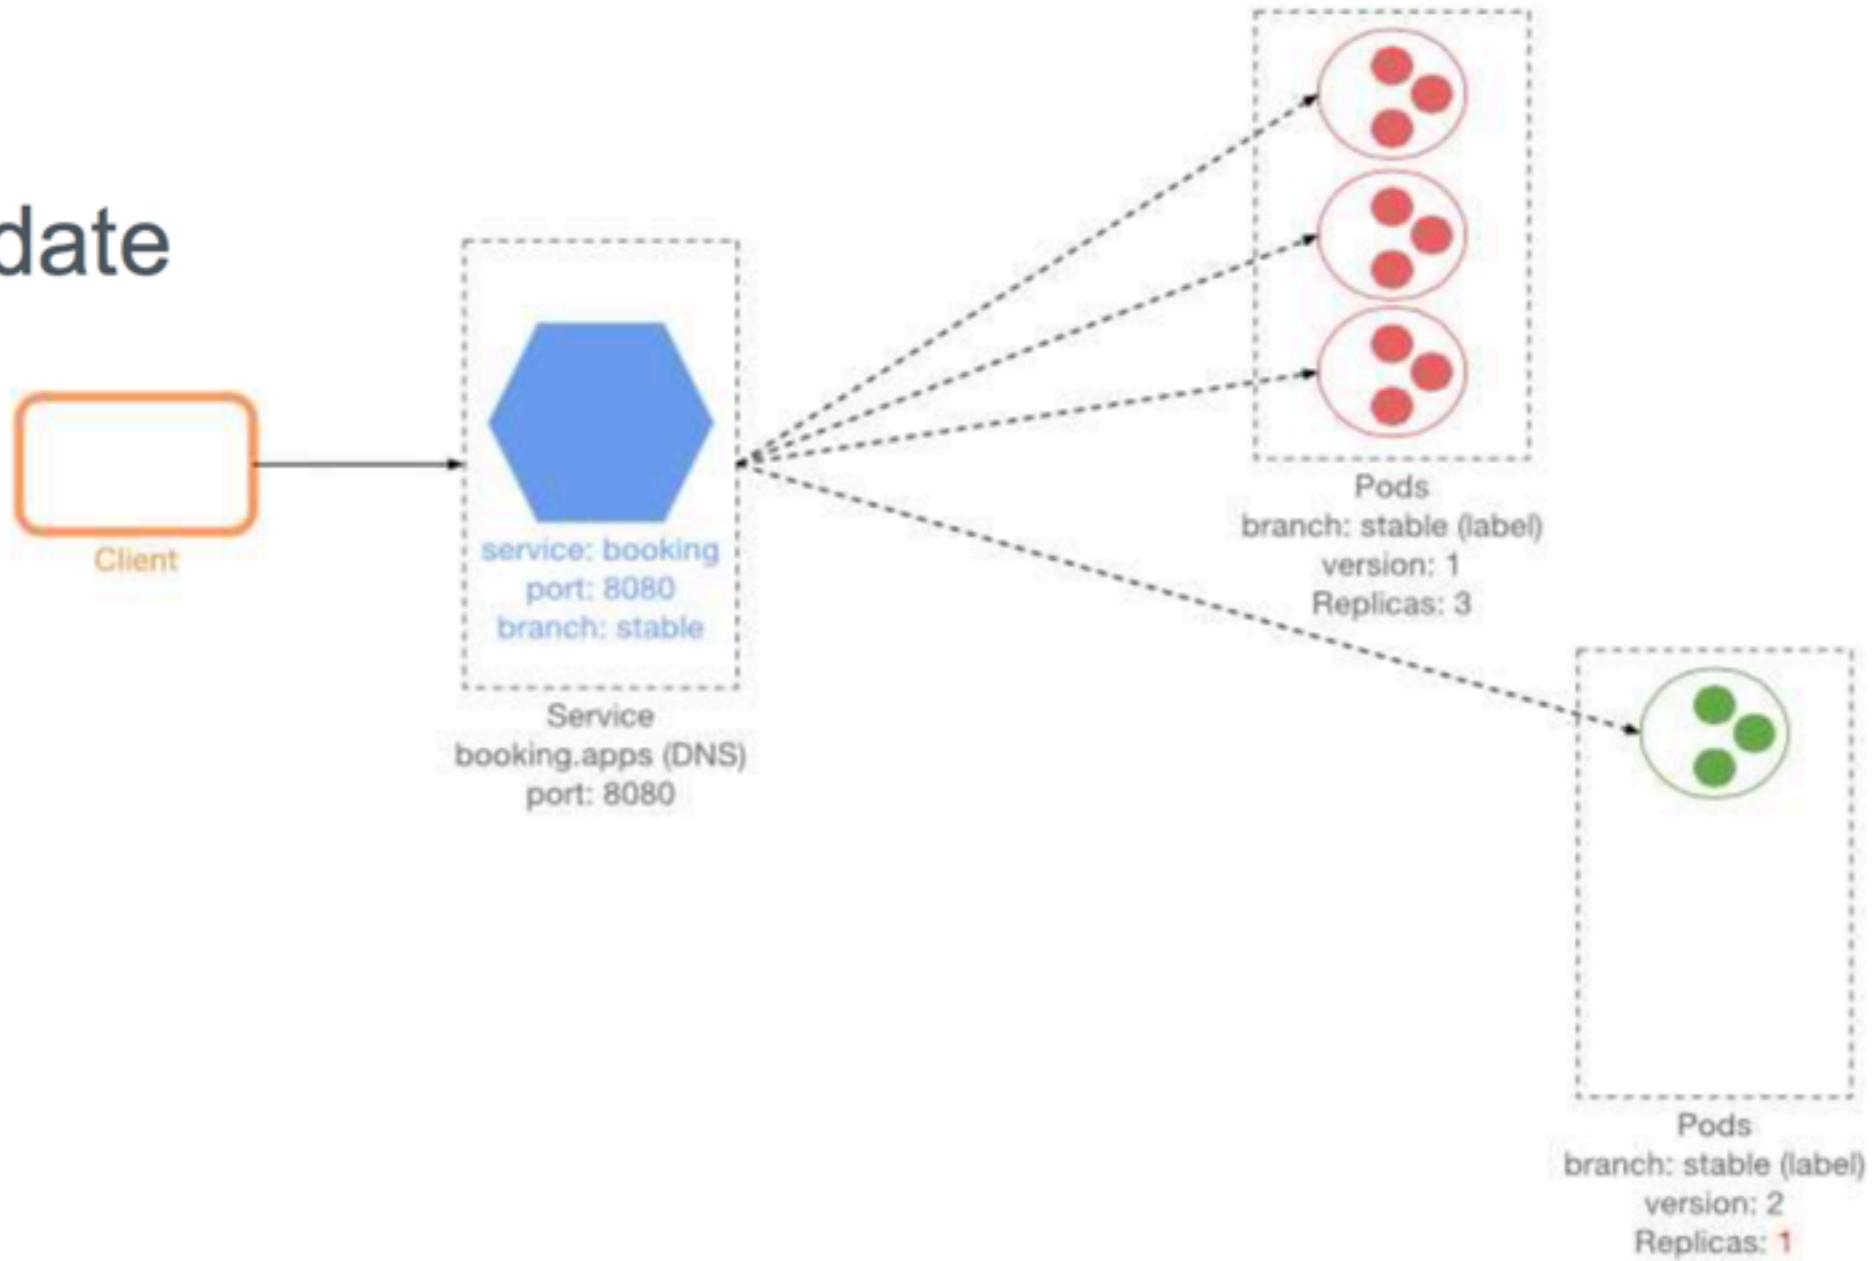
# Rolling update (1)

## Rolling Update



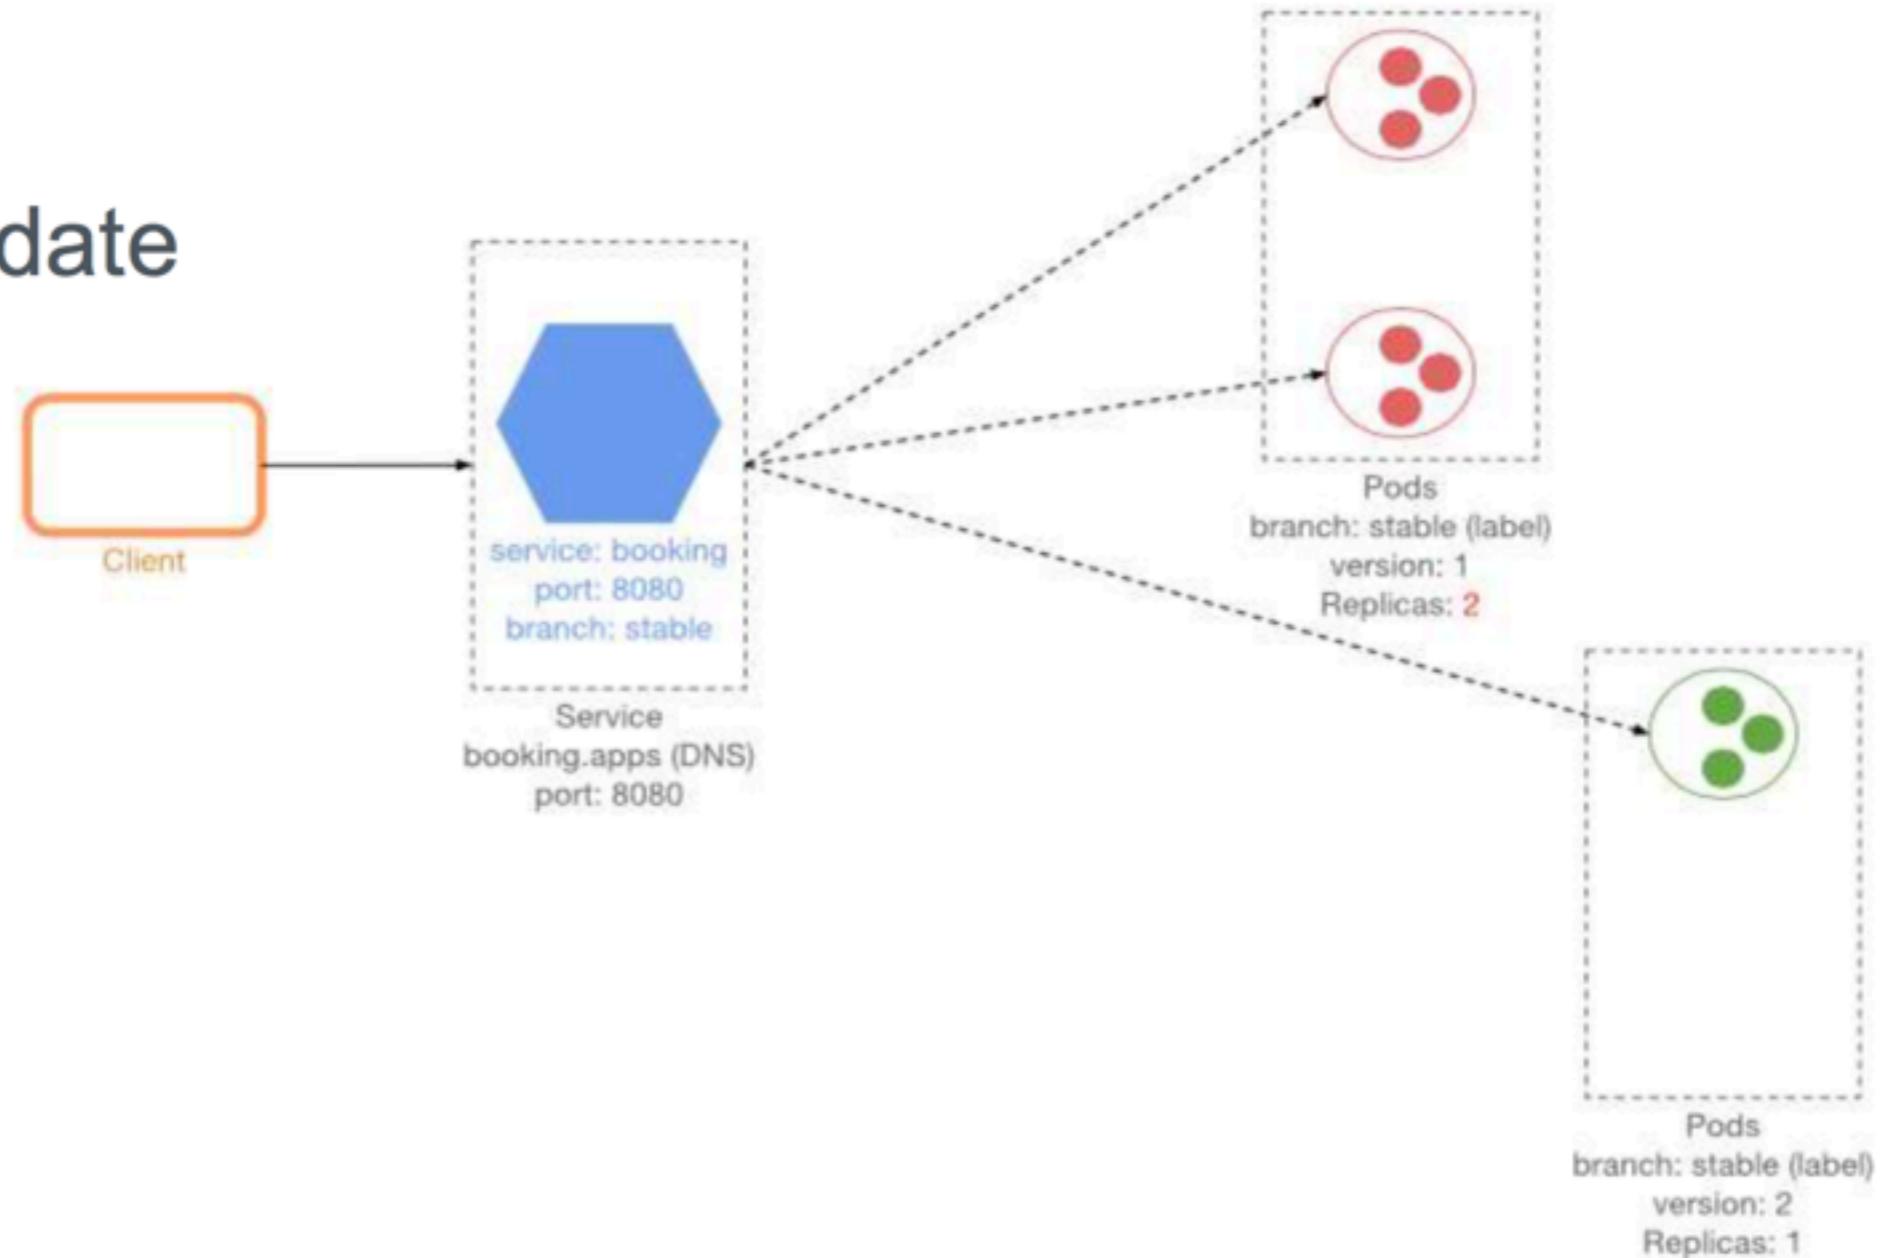
# Rolling update (2)

## Rolling Update



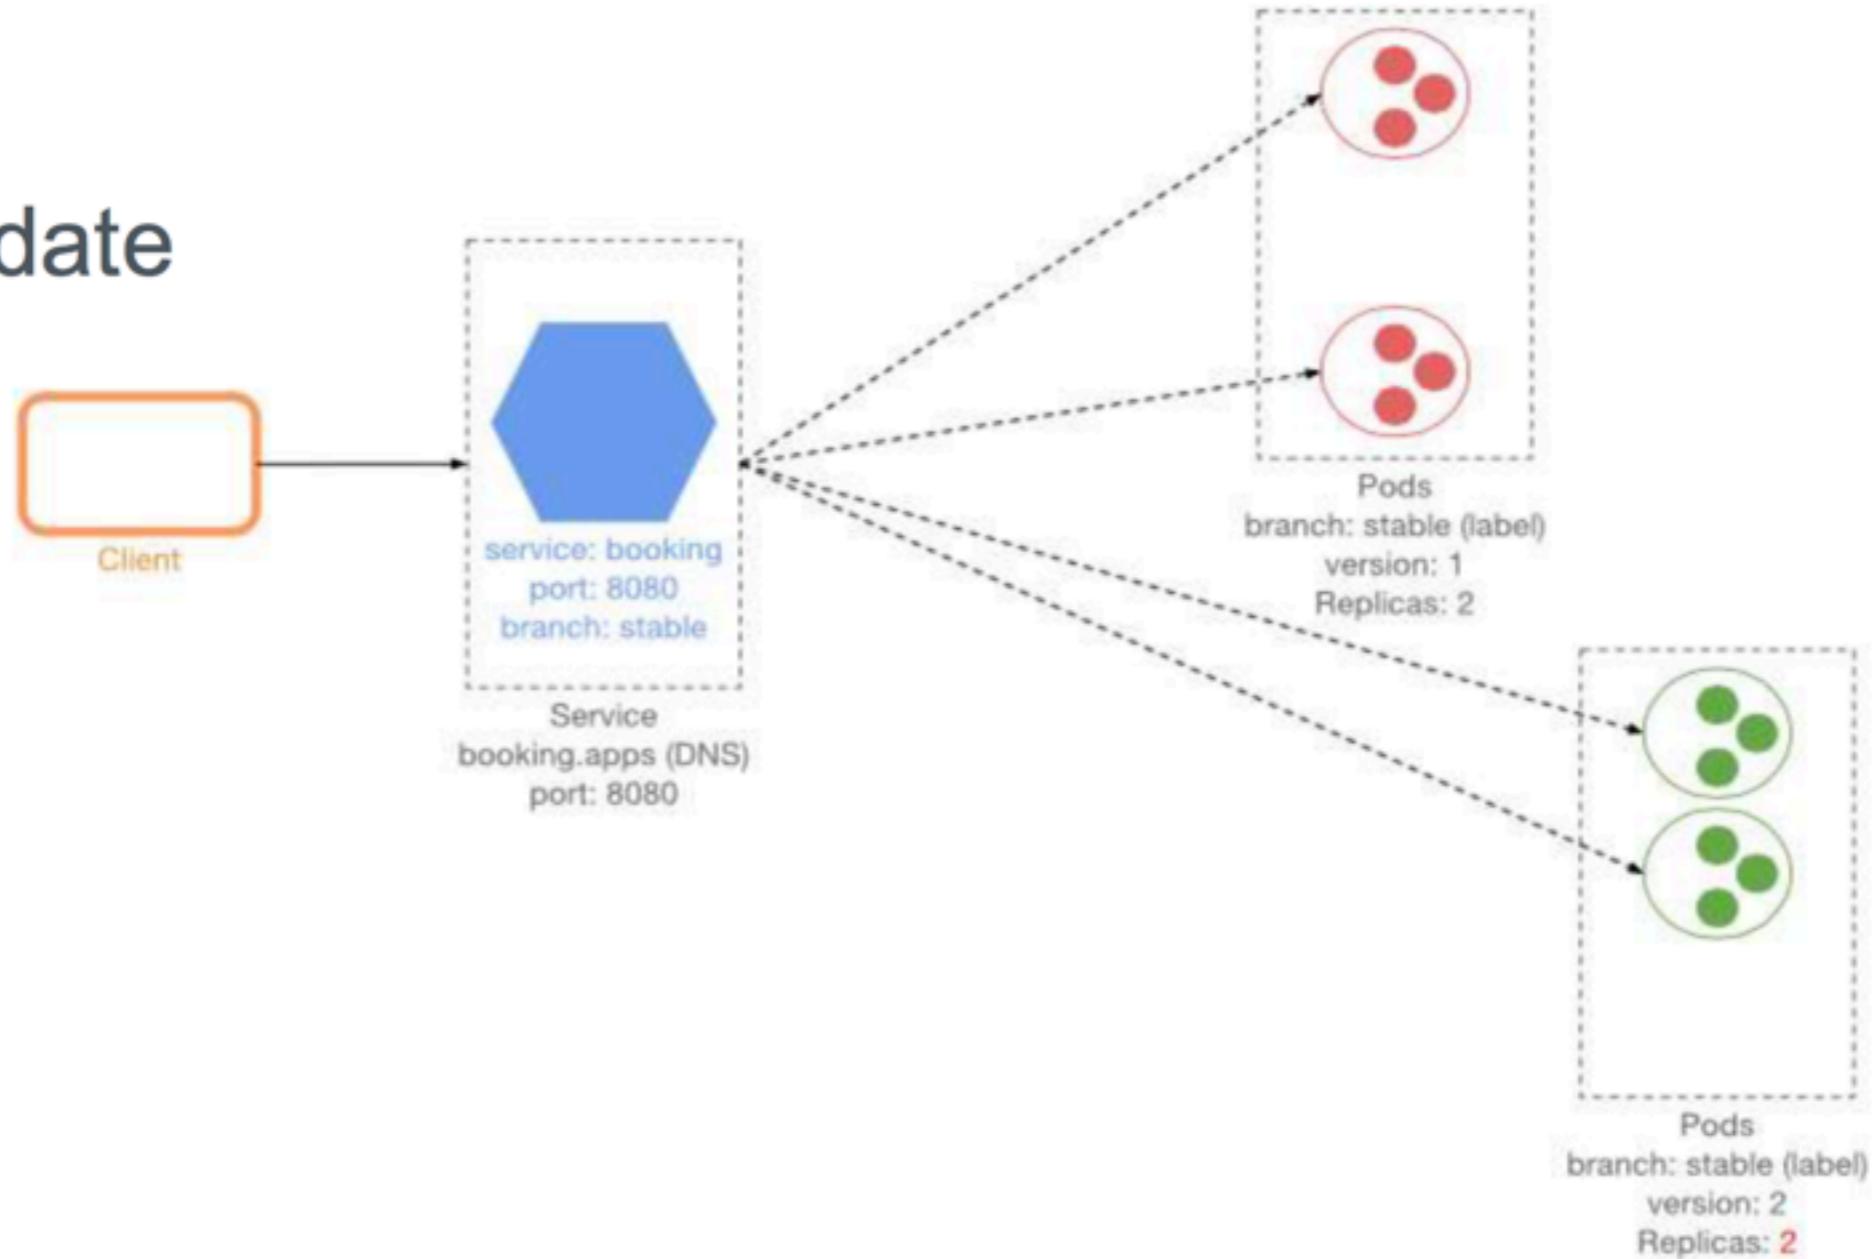
# Rolling update (3)

## Rolling Update



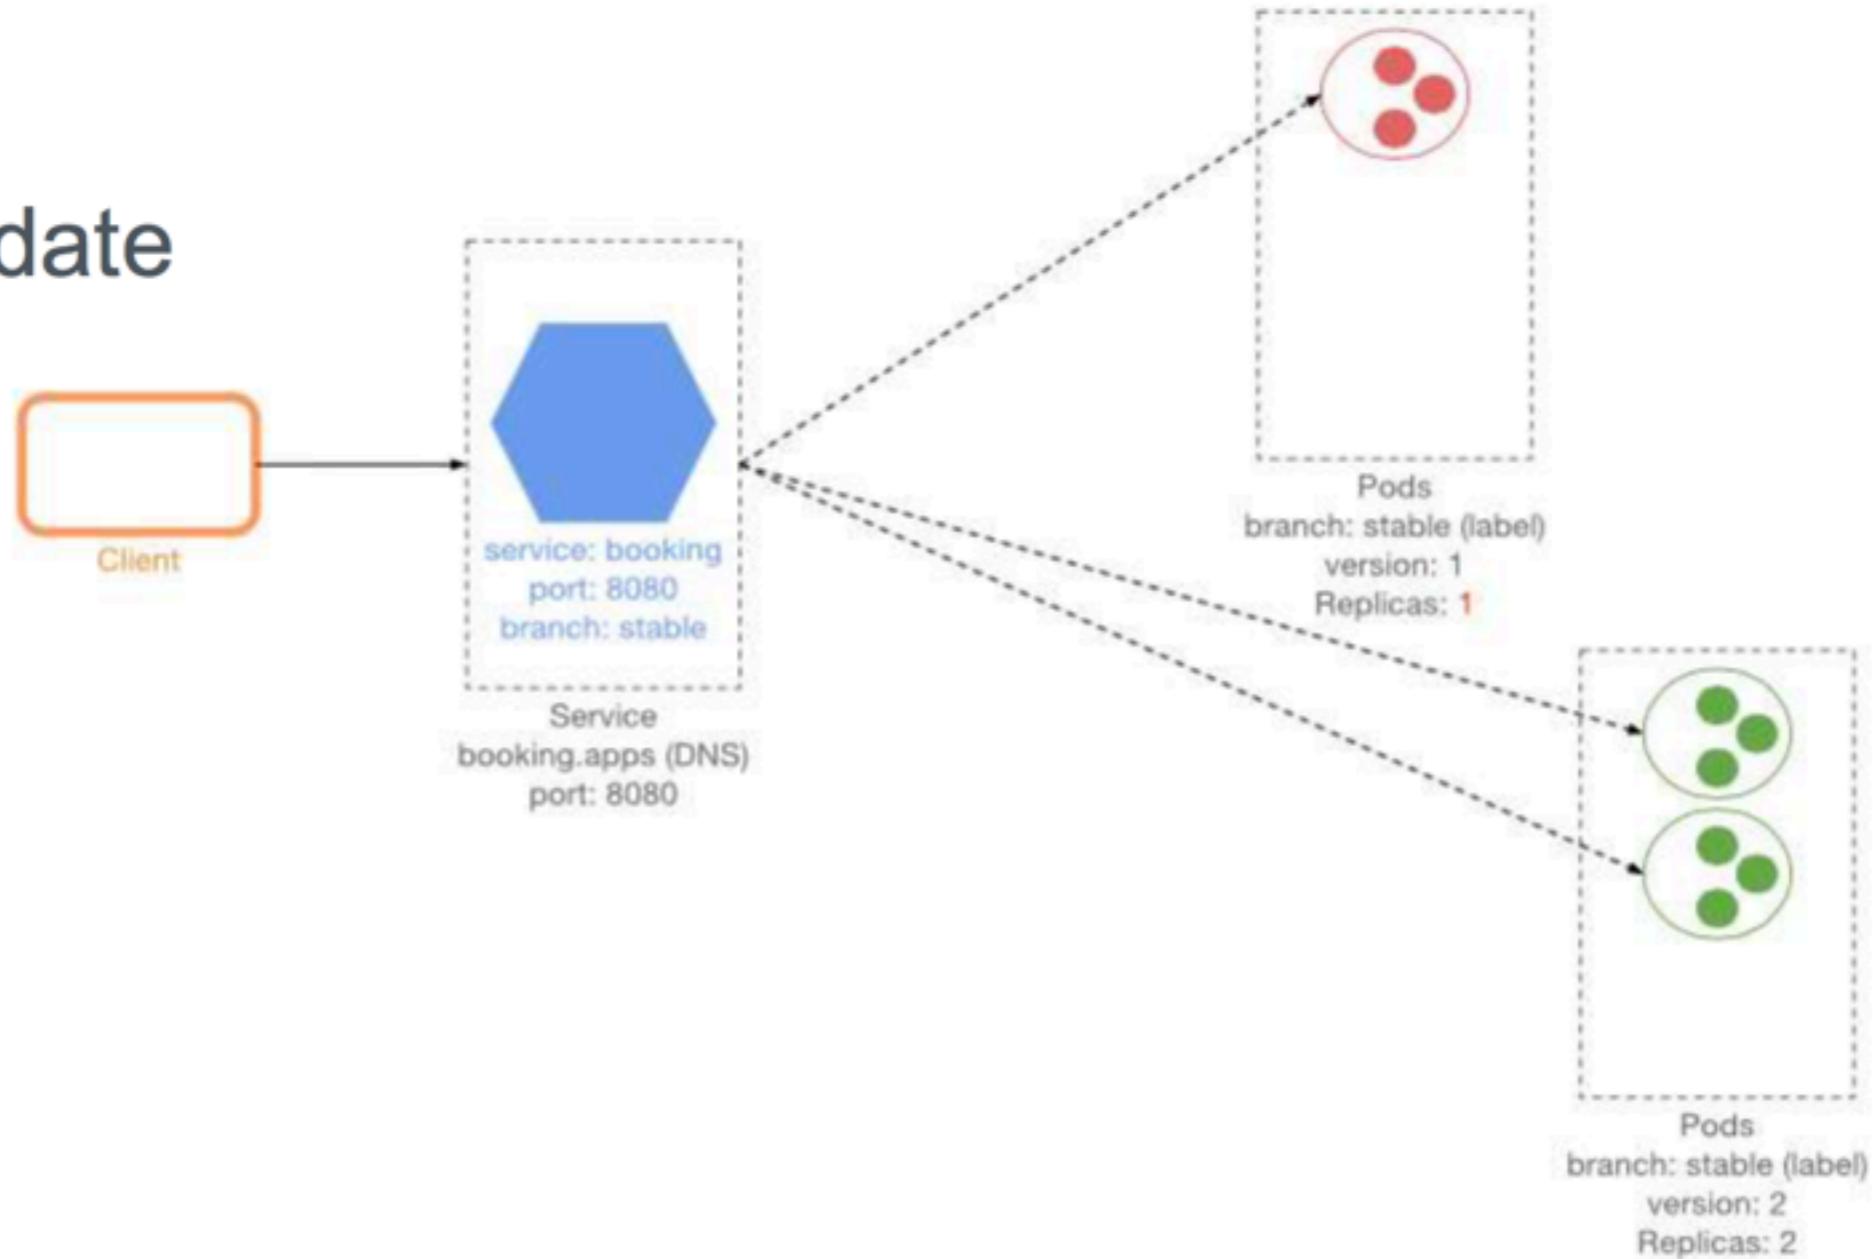
# Rolling update (4)

## Rolling Update



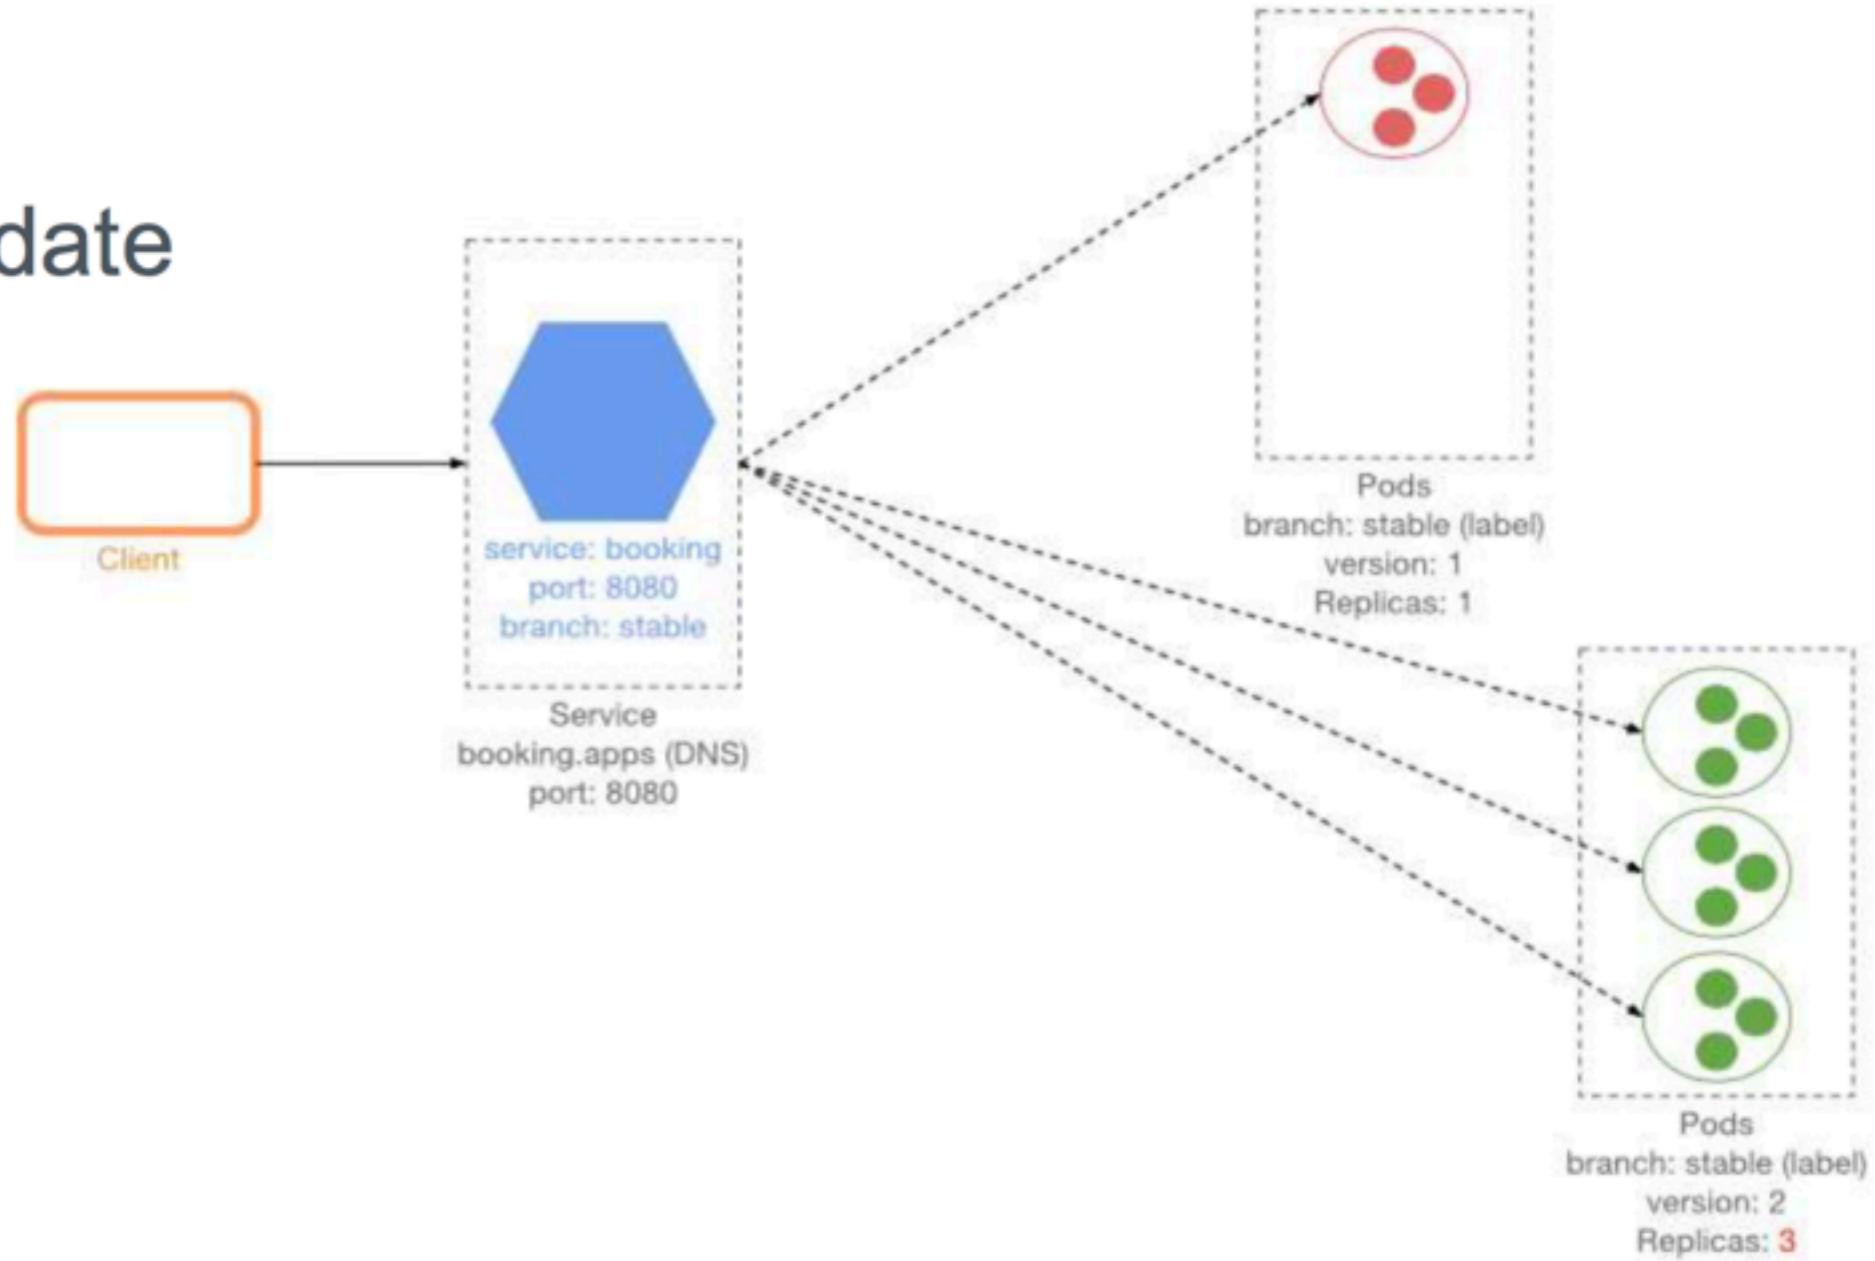
# Rolling update (5)

## Rolling Update



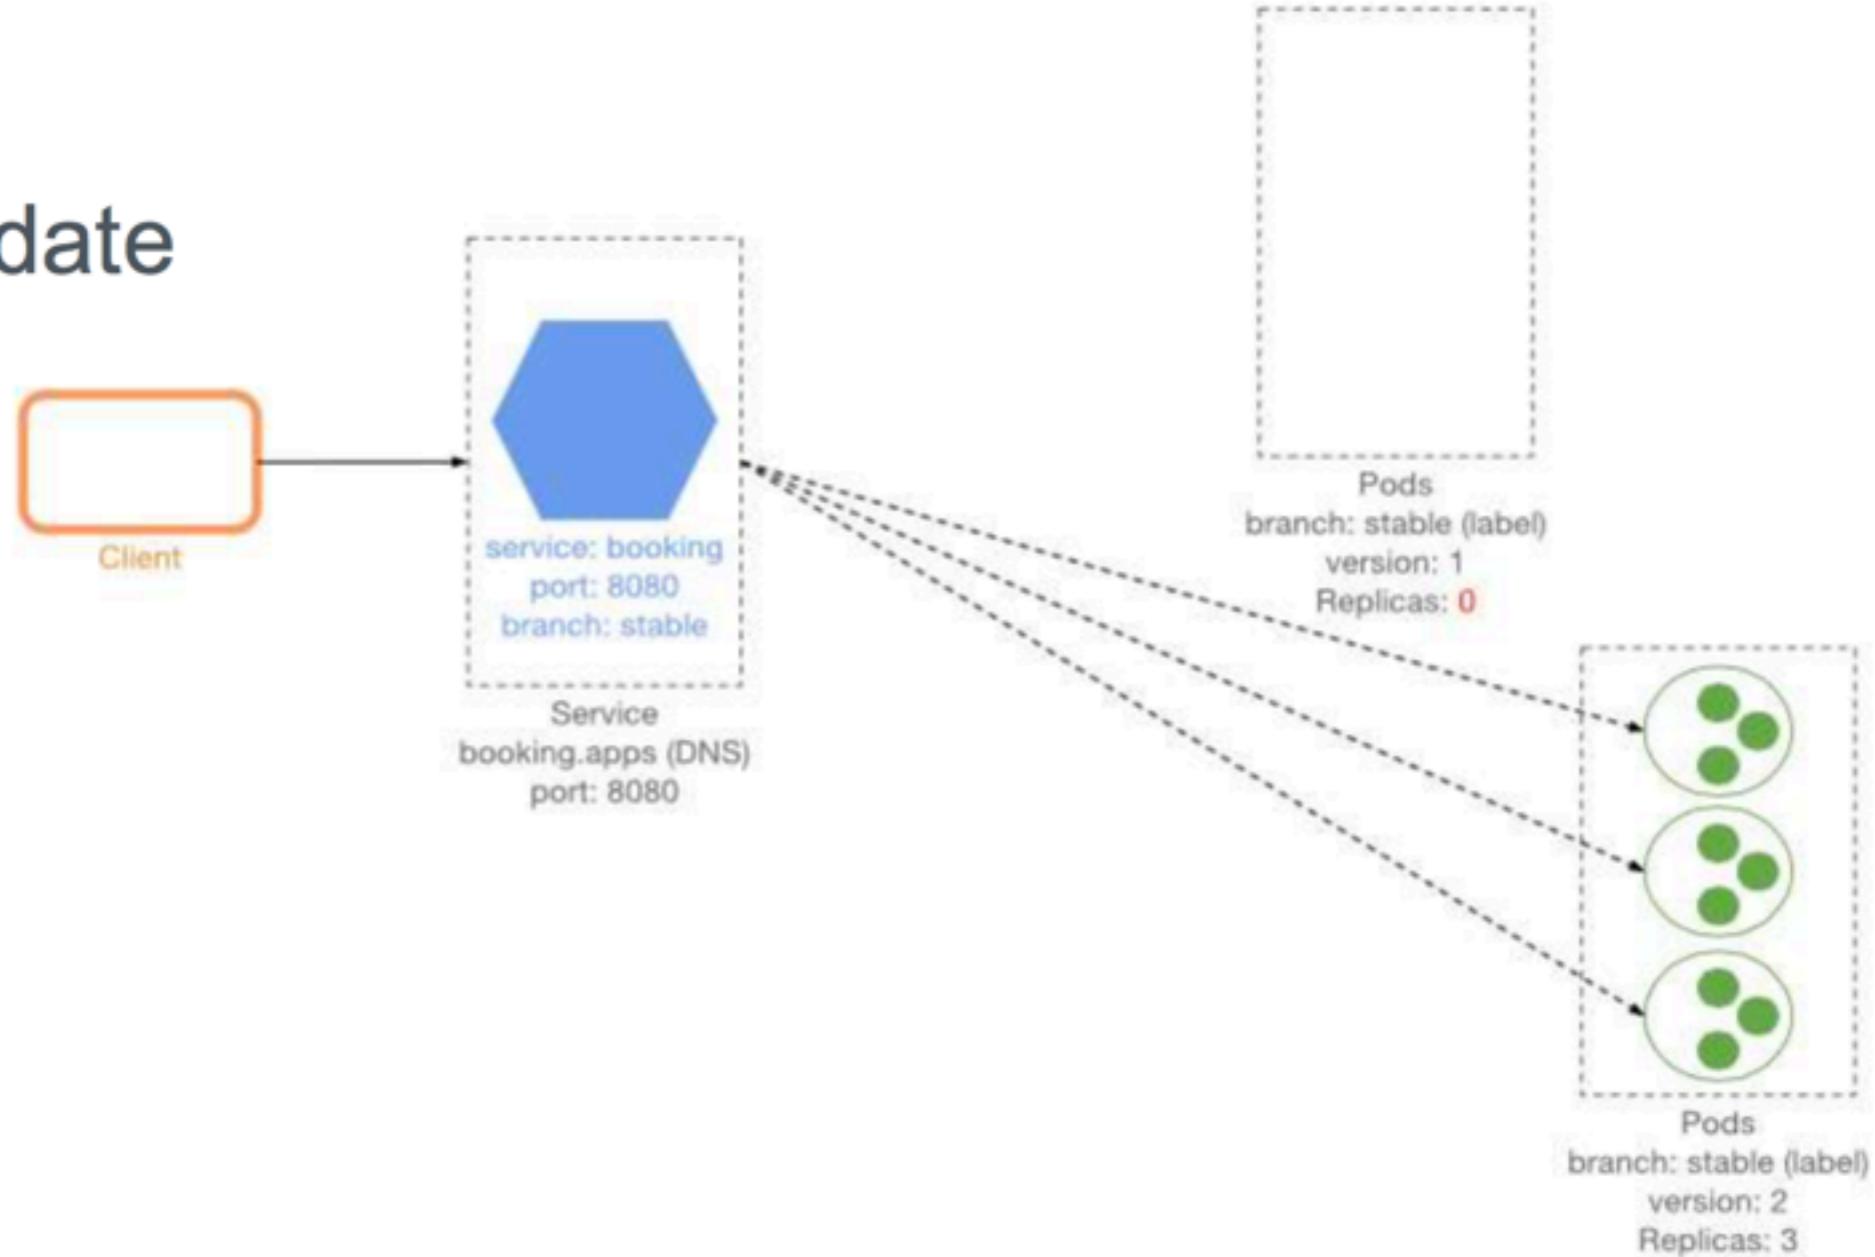
# Rolling update (6)

## Rolling Update



# Rolling update (7)

## Rolling Update



# Storage orchestration



# Storage orchestration

Support many type of storages

1. Local storage
2. Network storage (NFS, GlusterFS, Ceph)
3. Cloud storage (GCE, AWS, Azure )



# Self-healing



# **Provide well known ports for Kubernetes services**



# **Service Discovery and Load balance**

**Service** is the connector/proxy for client to connect with Pod

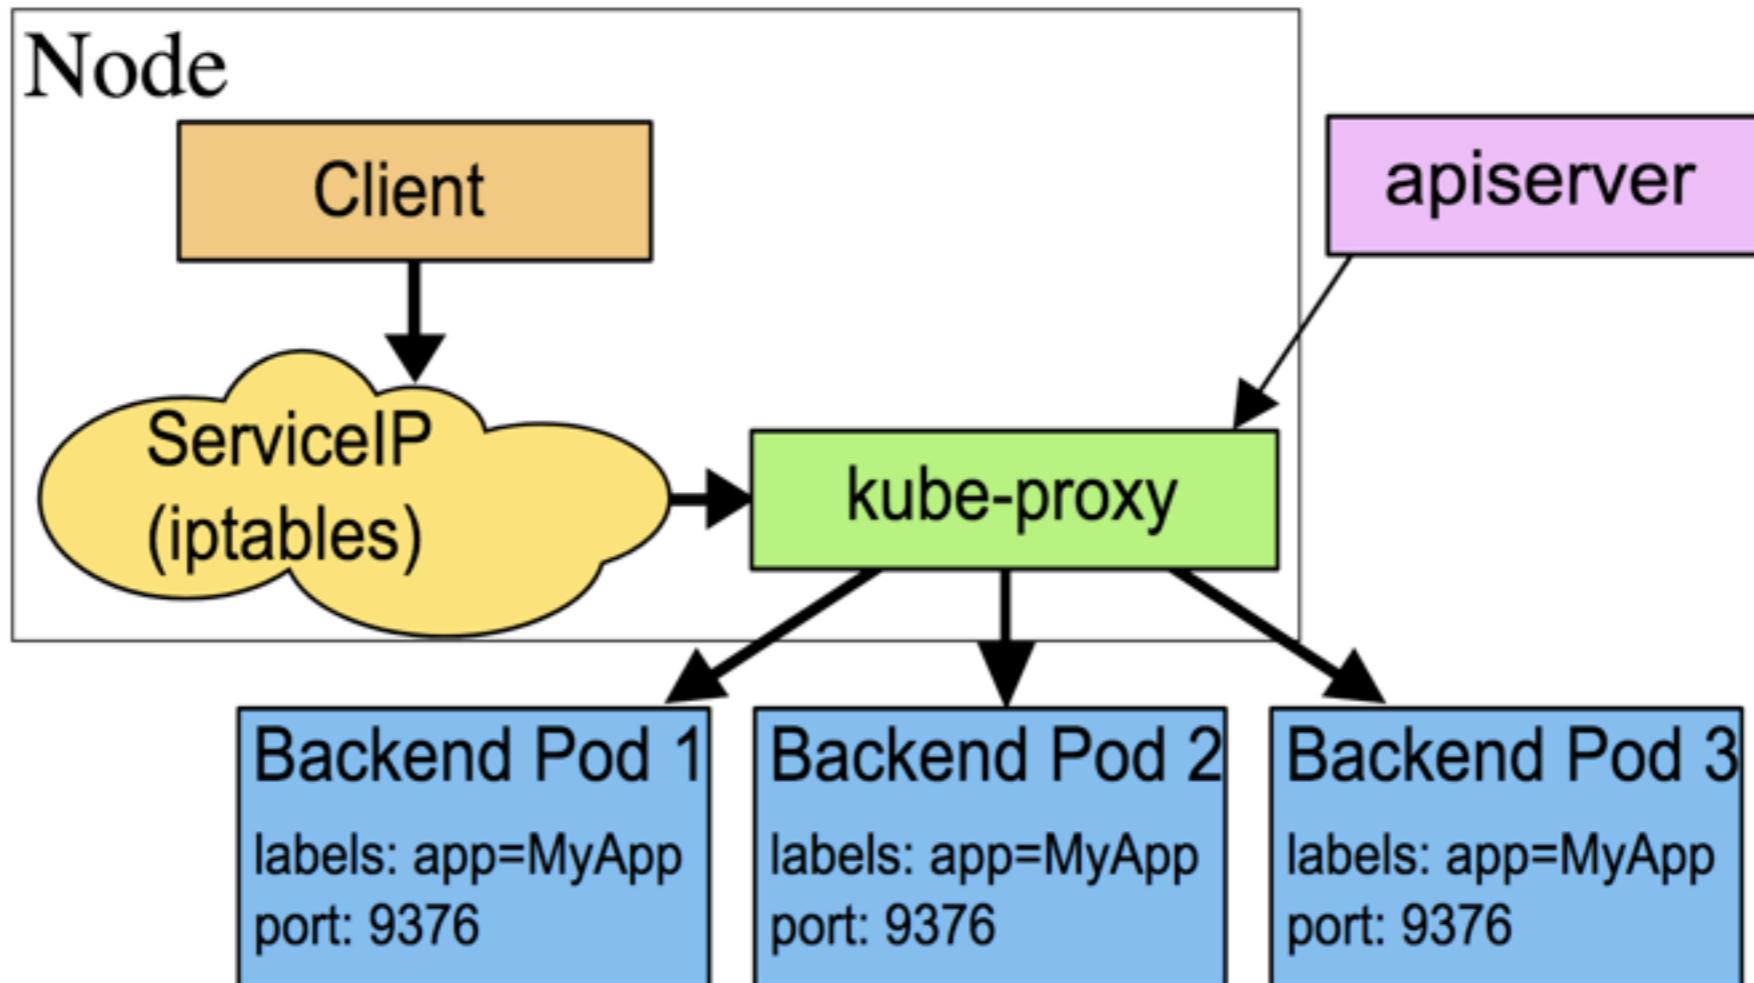
**Discovery** use for service to look Pod

**Support load balance** between Pods (replica)

<https://kubernetes.io/docs/concepts/services-networking/service/>



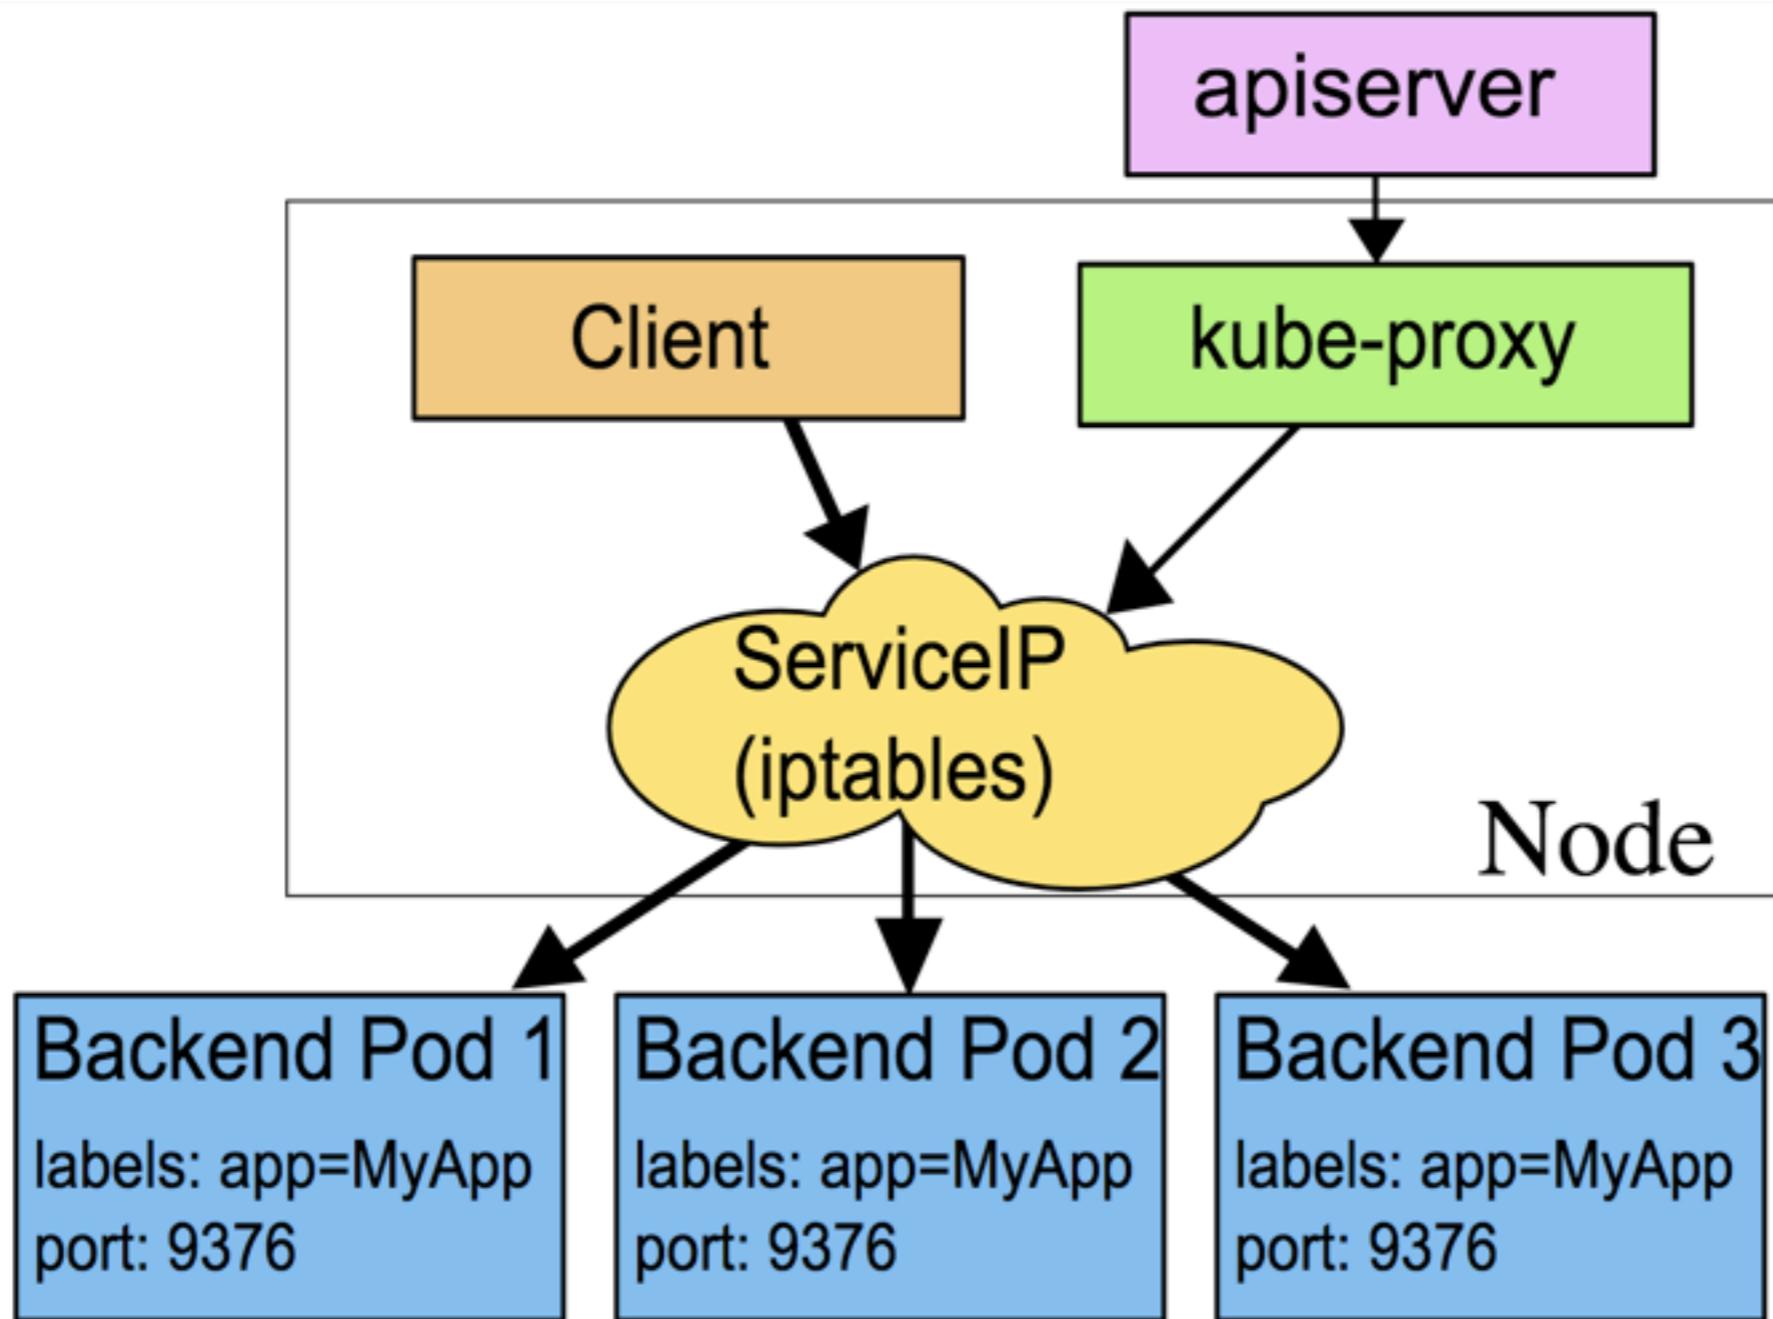
# Proxy-mode: userspace



Note that in the above diagram, `clusterIP` is shown as `ServiceIP`.



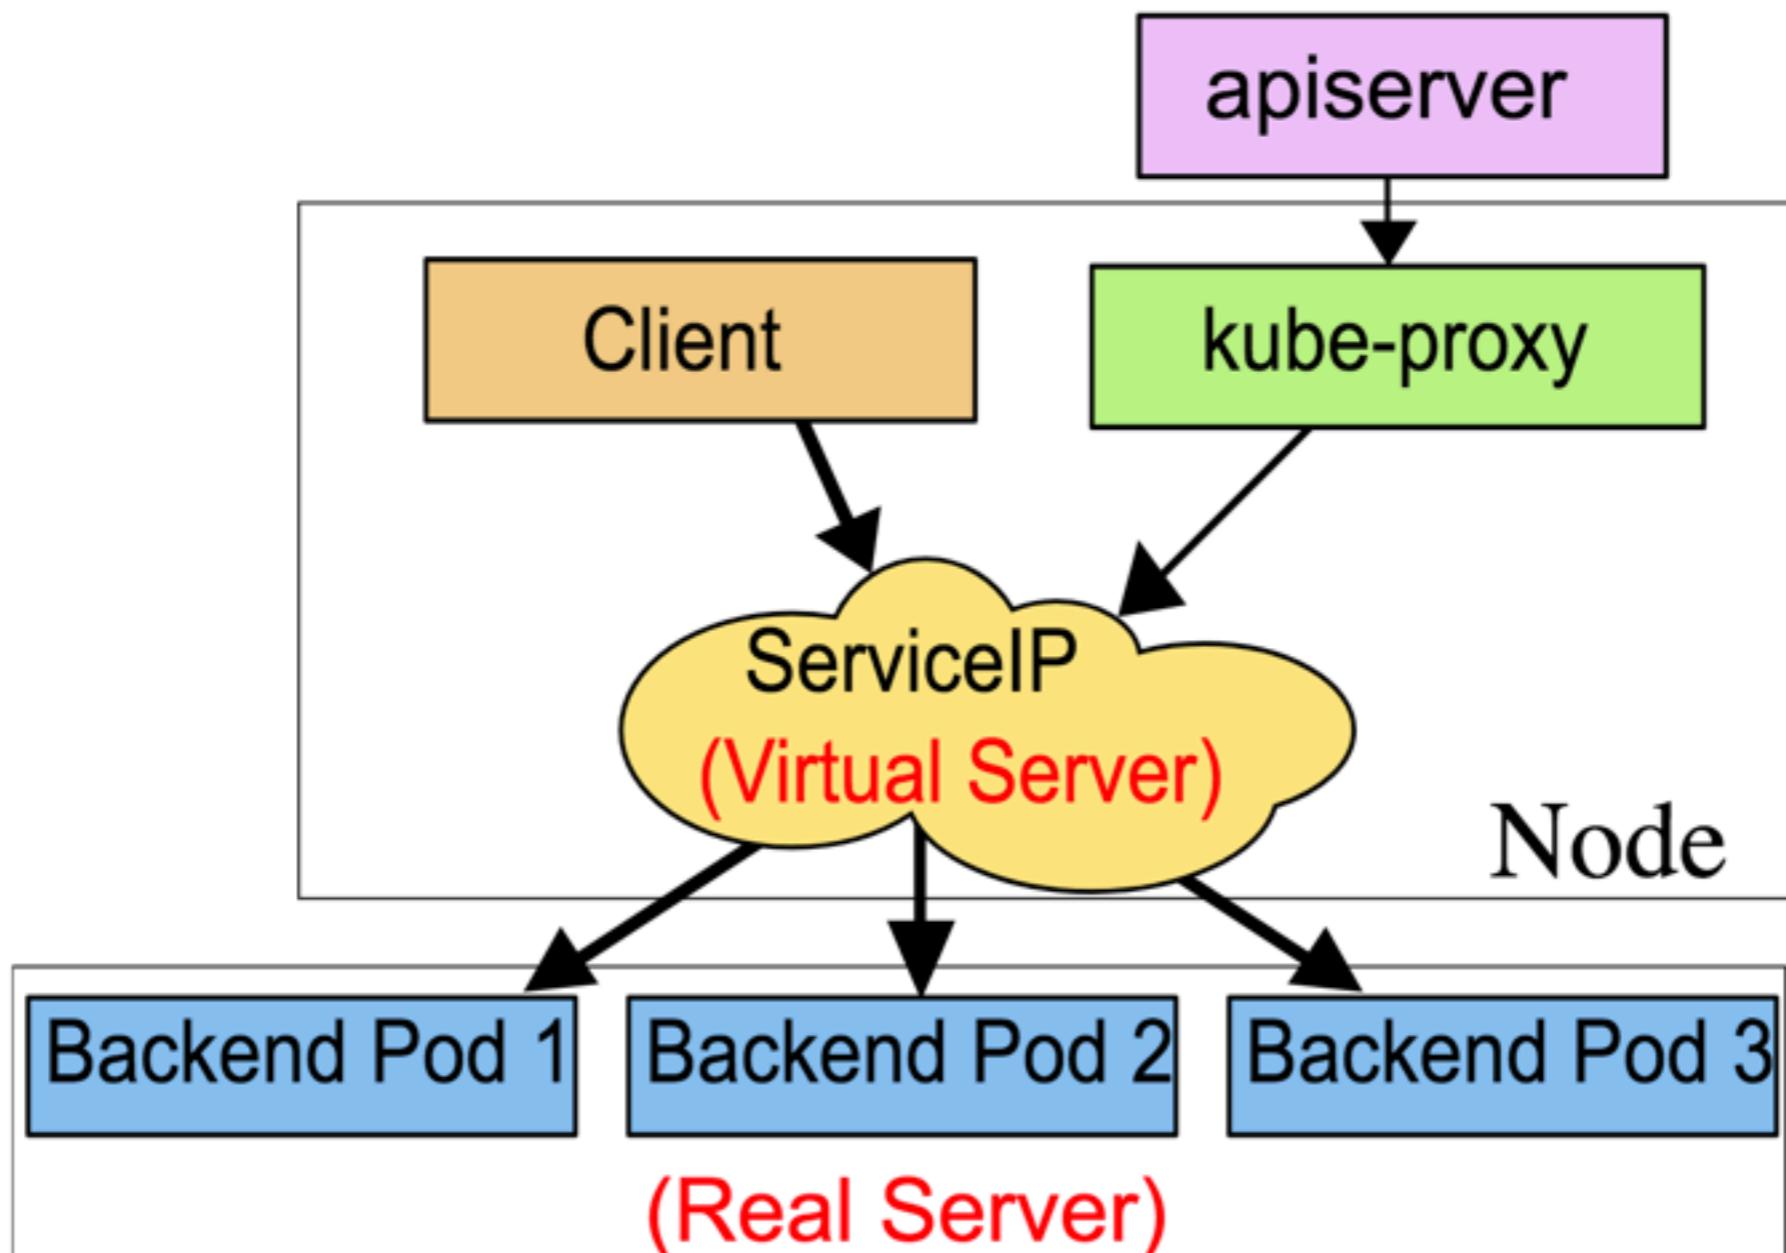
# Proxy-mode: iptables



Note that in the above diagram, `clusterIP` is shown as `ServiceIP`.



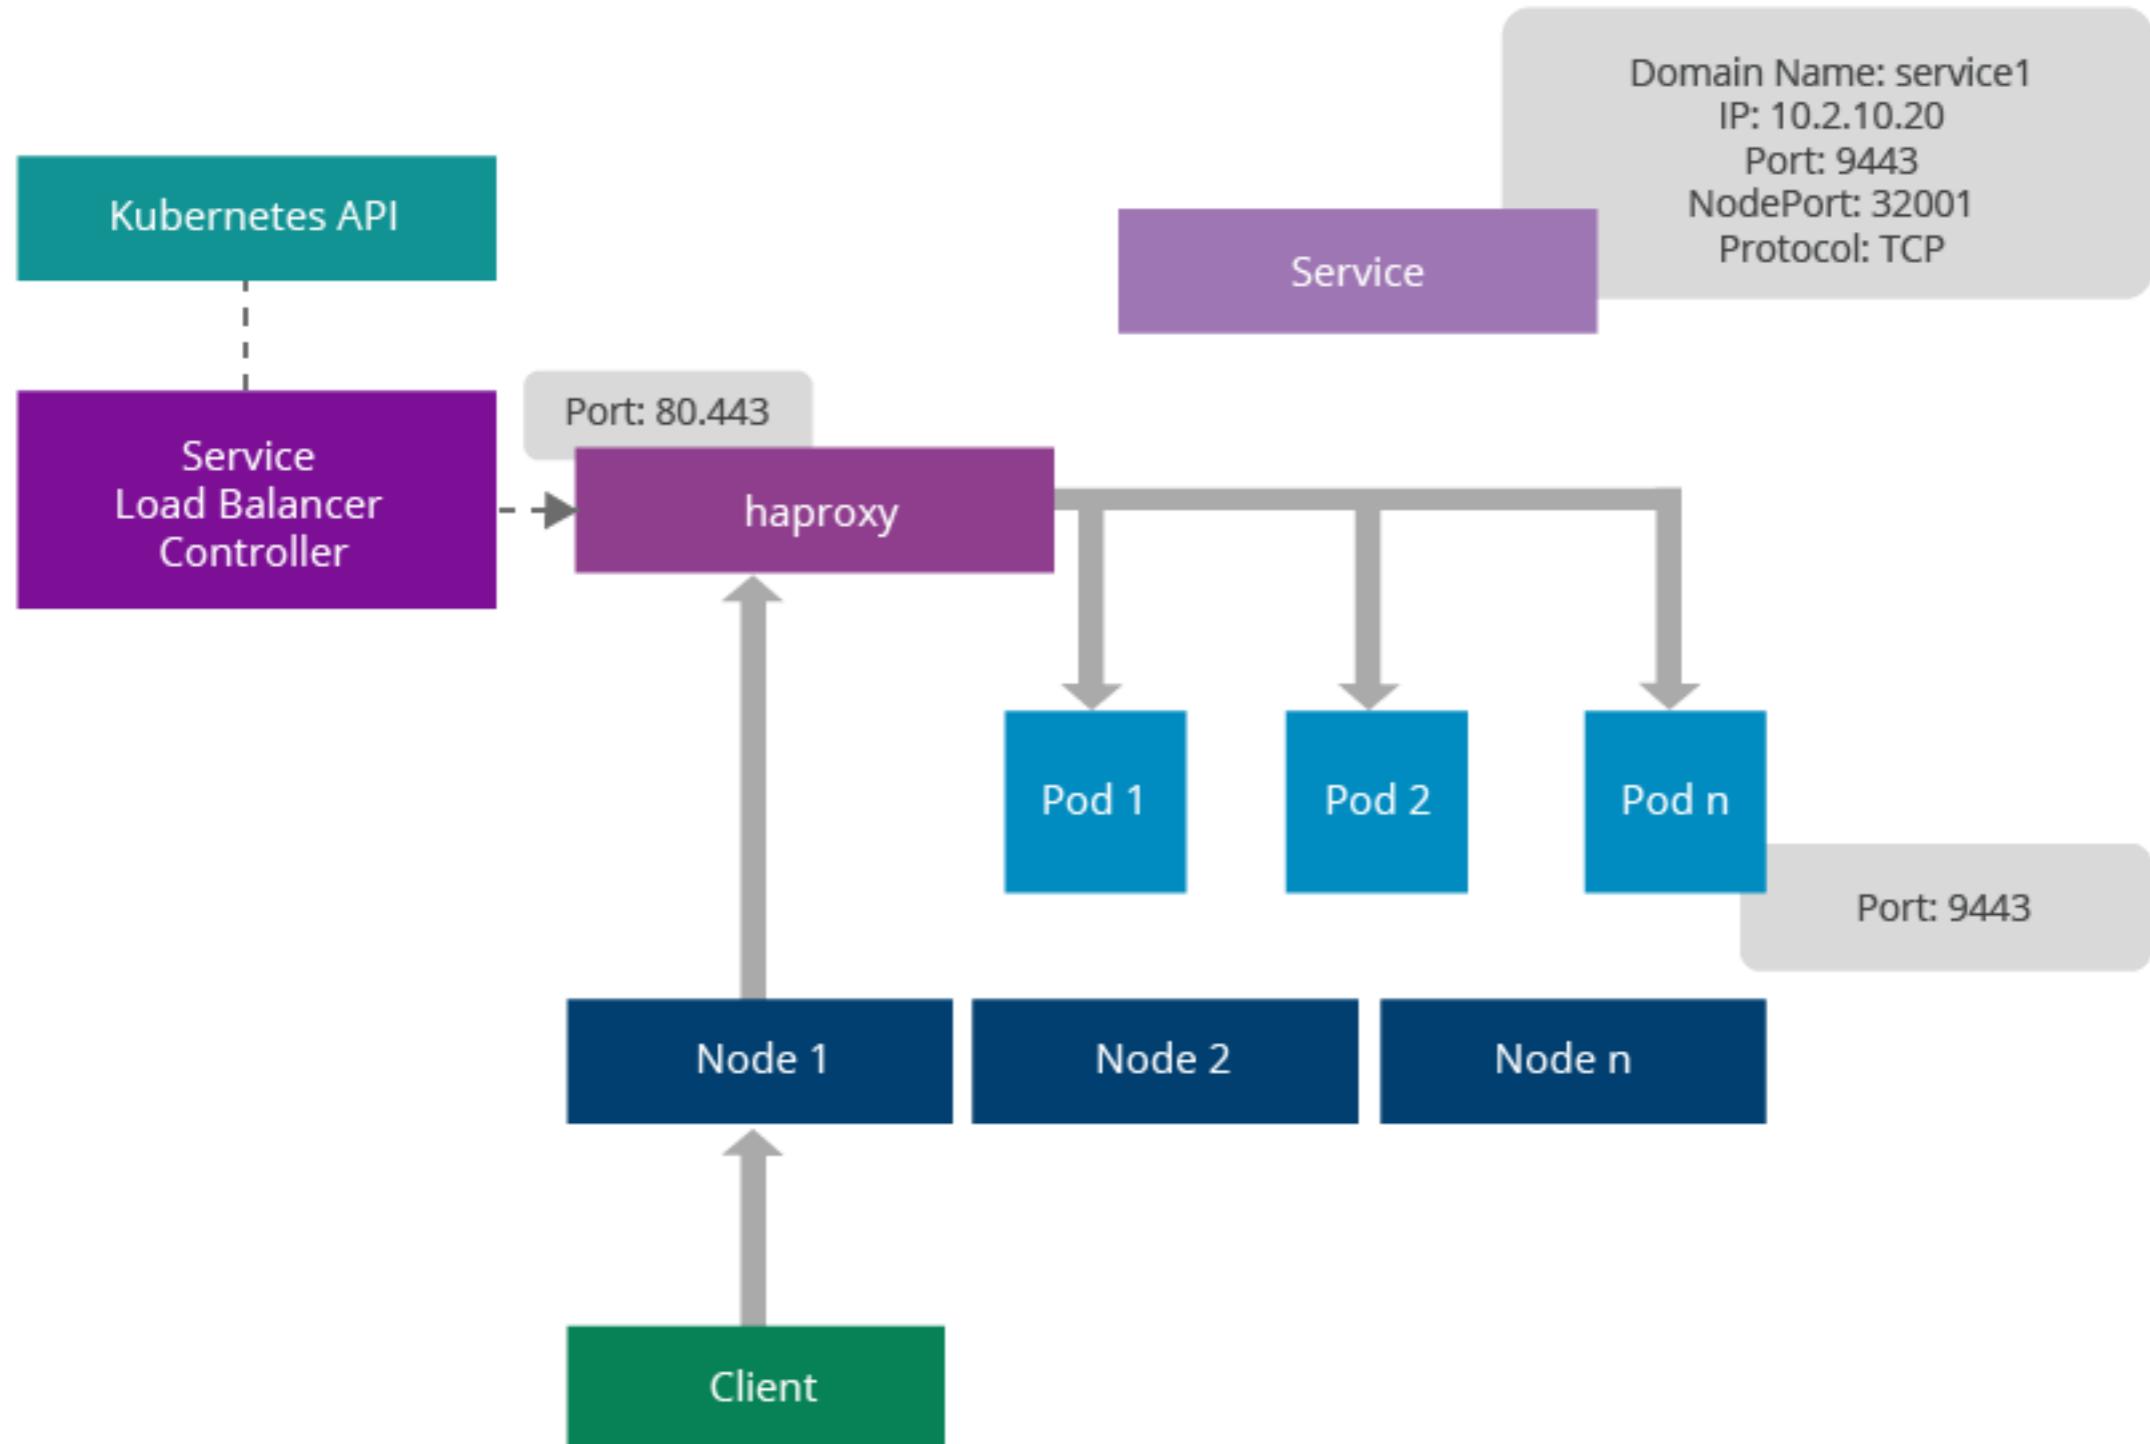
# Proxy-mode: ipvs (beta 1.9)



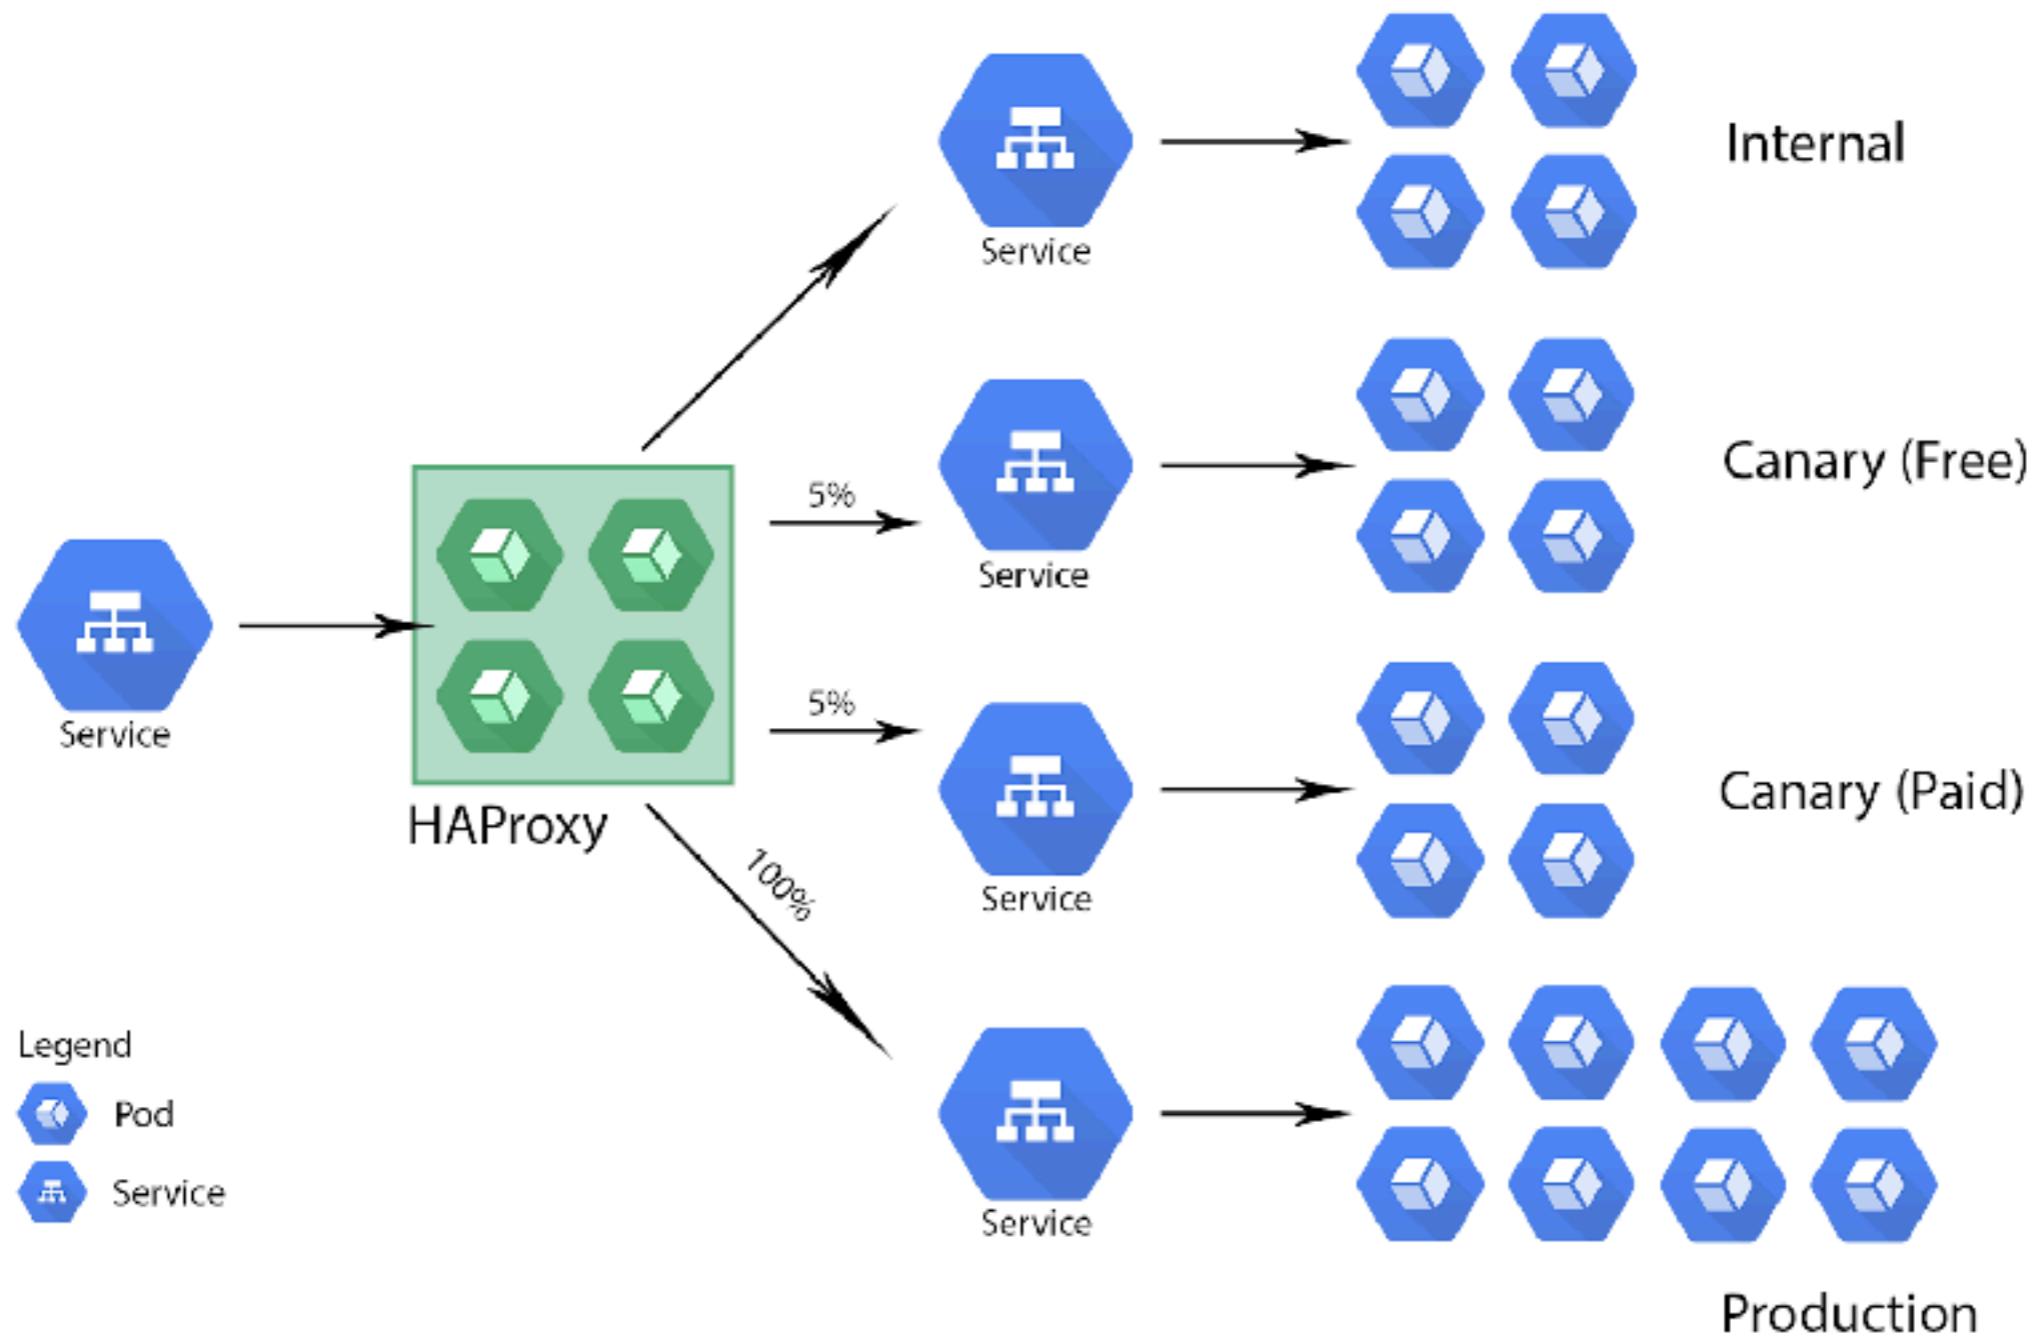
# **Sticky session management using Service Load Balancers**



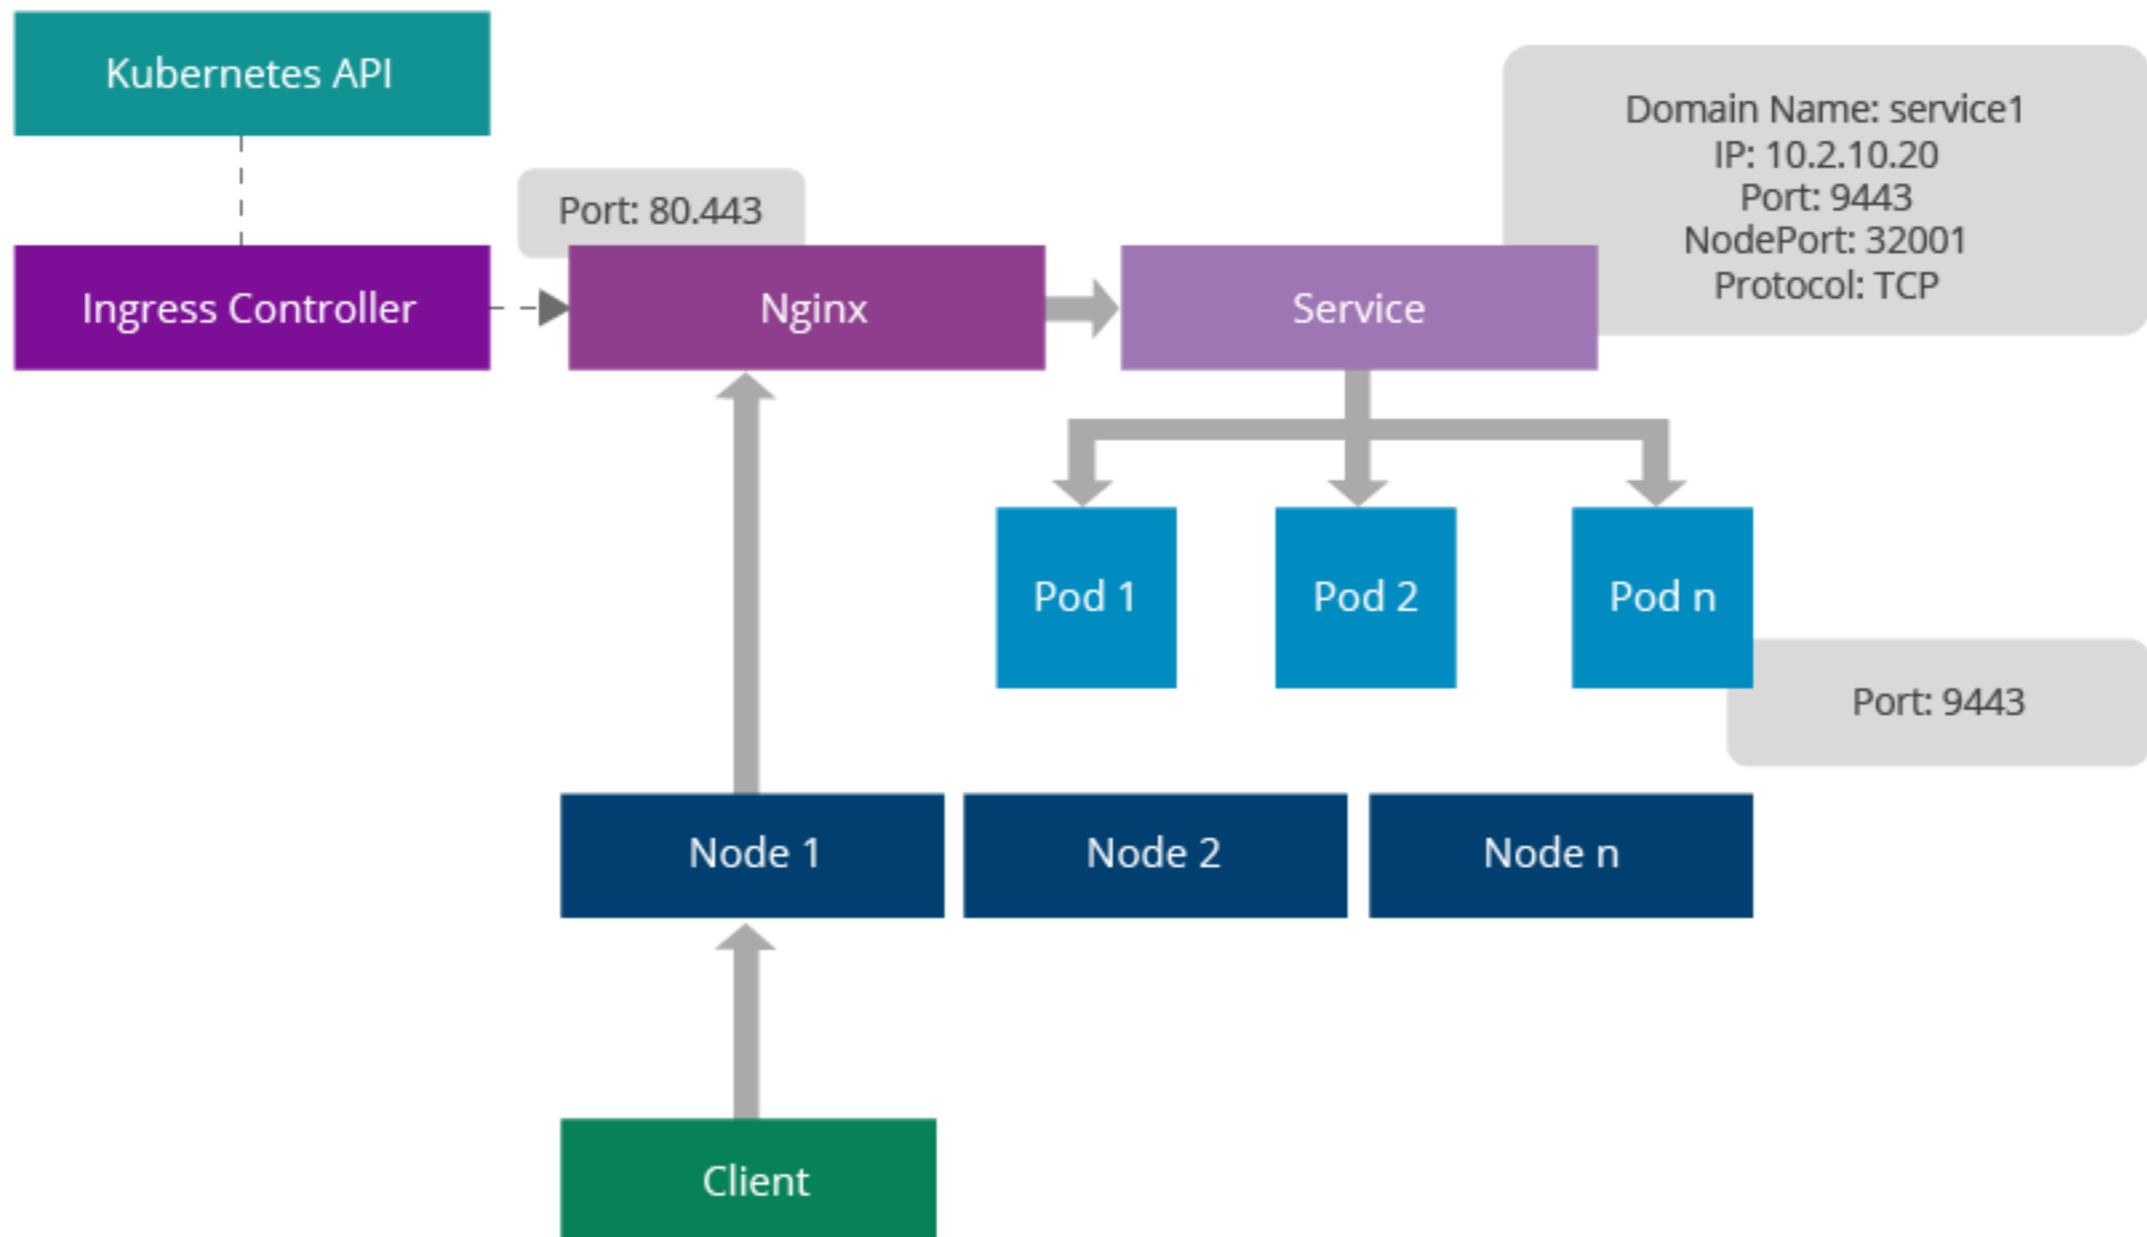
# Service Load Balancer



# Service Load Balancer



# Ingress controller



# **Secret and configuration management**



# Secret and configuration mgt

From **12factor** app

Store config in environments  
(dev/qa/staging/prod)

<https://12factor.net/>



# Secret and configuration mgt

Keep confidential data to running app in  
encryption format

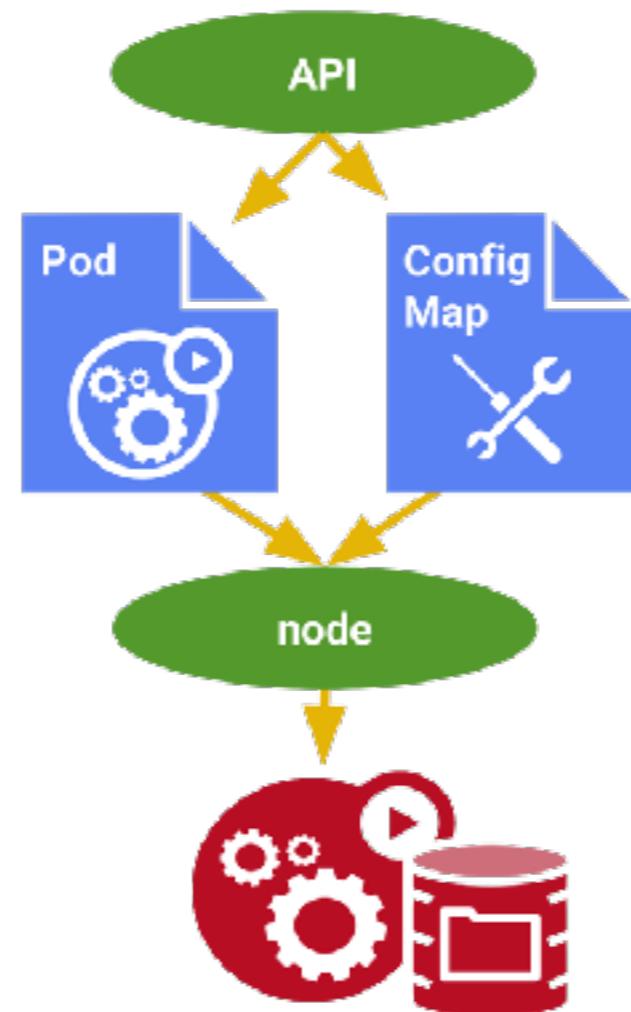
\$kubectl get secret

NAME	TYPE	DATA	AGE
default-token-26swm	kubernetes.io/service-account-token	3	2h



# Secret and configuration mgt

Use **ConfigMap** to define all configuration that reference from Pod

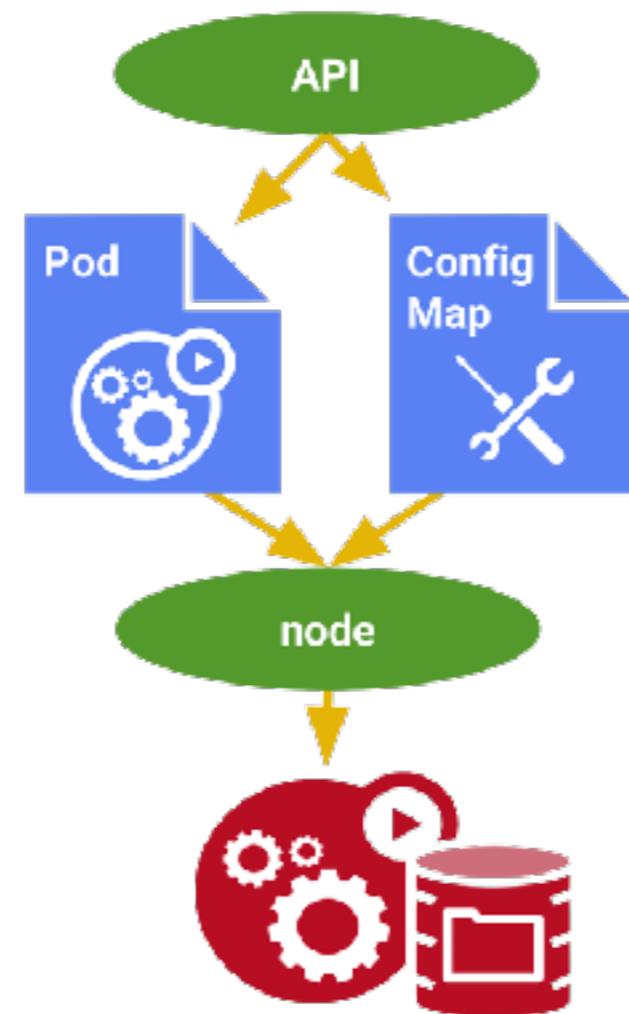


<https://kubernetes.io/docs/tasks/configure-pod-container/configure-pod-configmap/>



# Secret and configuration mgt

Come from **12 factor app** :: config come from the environment



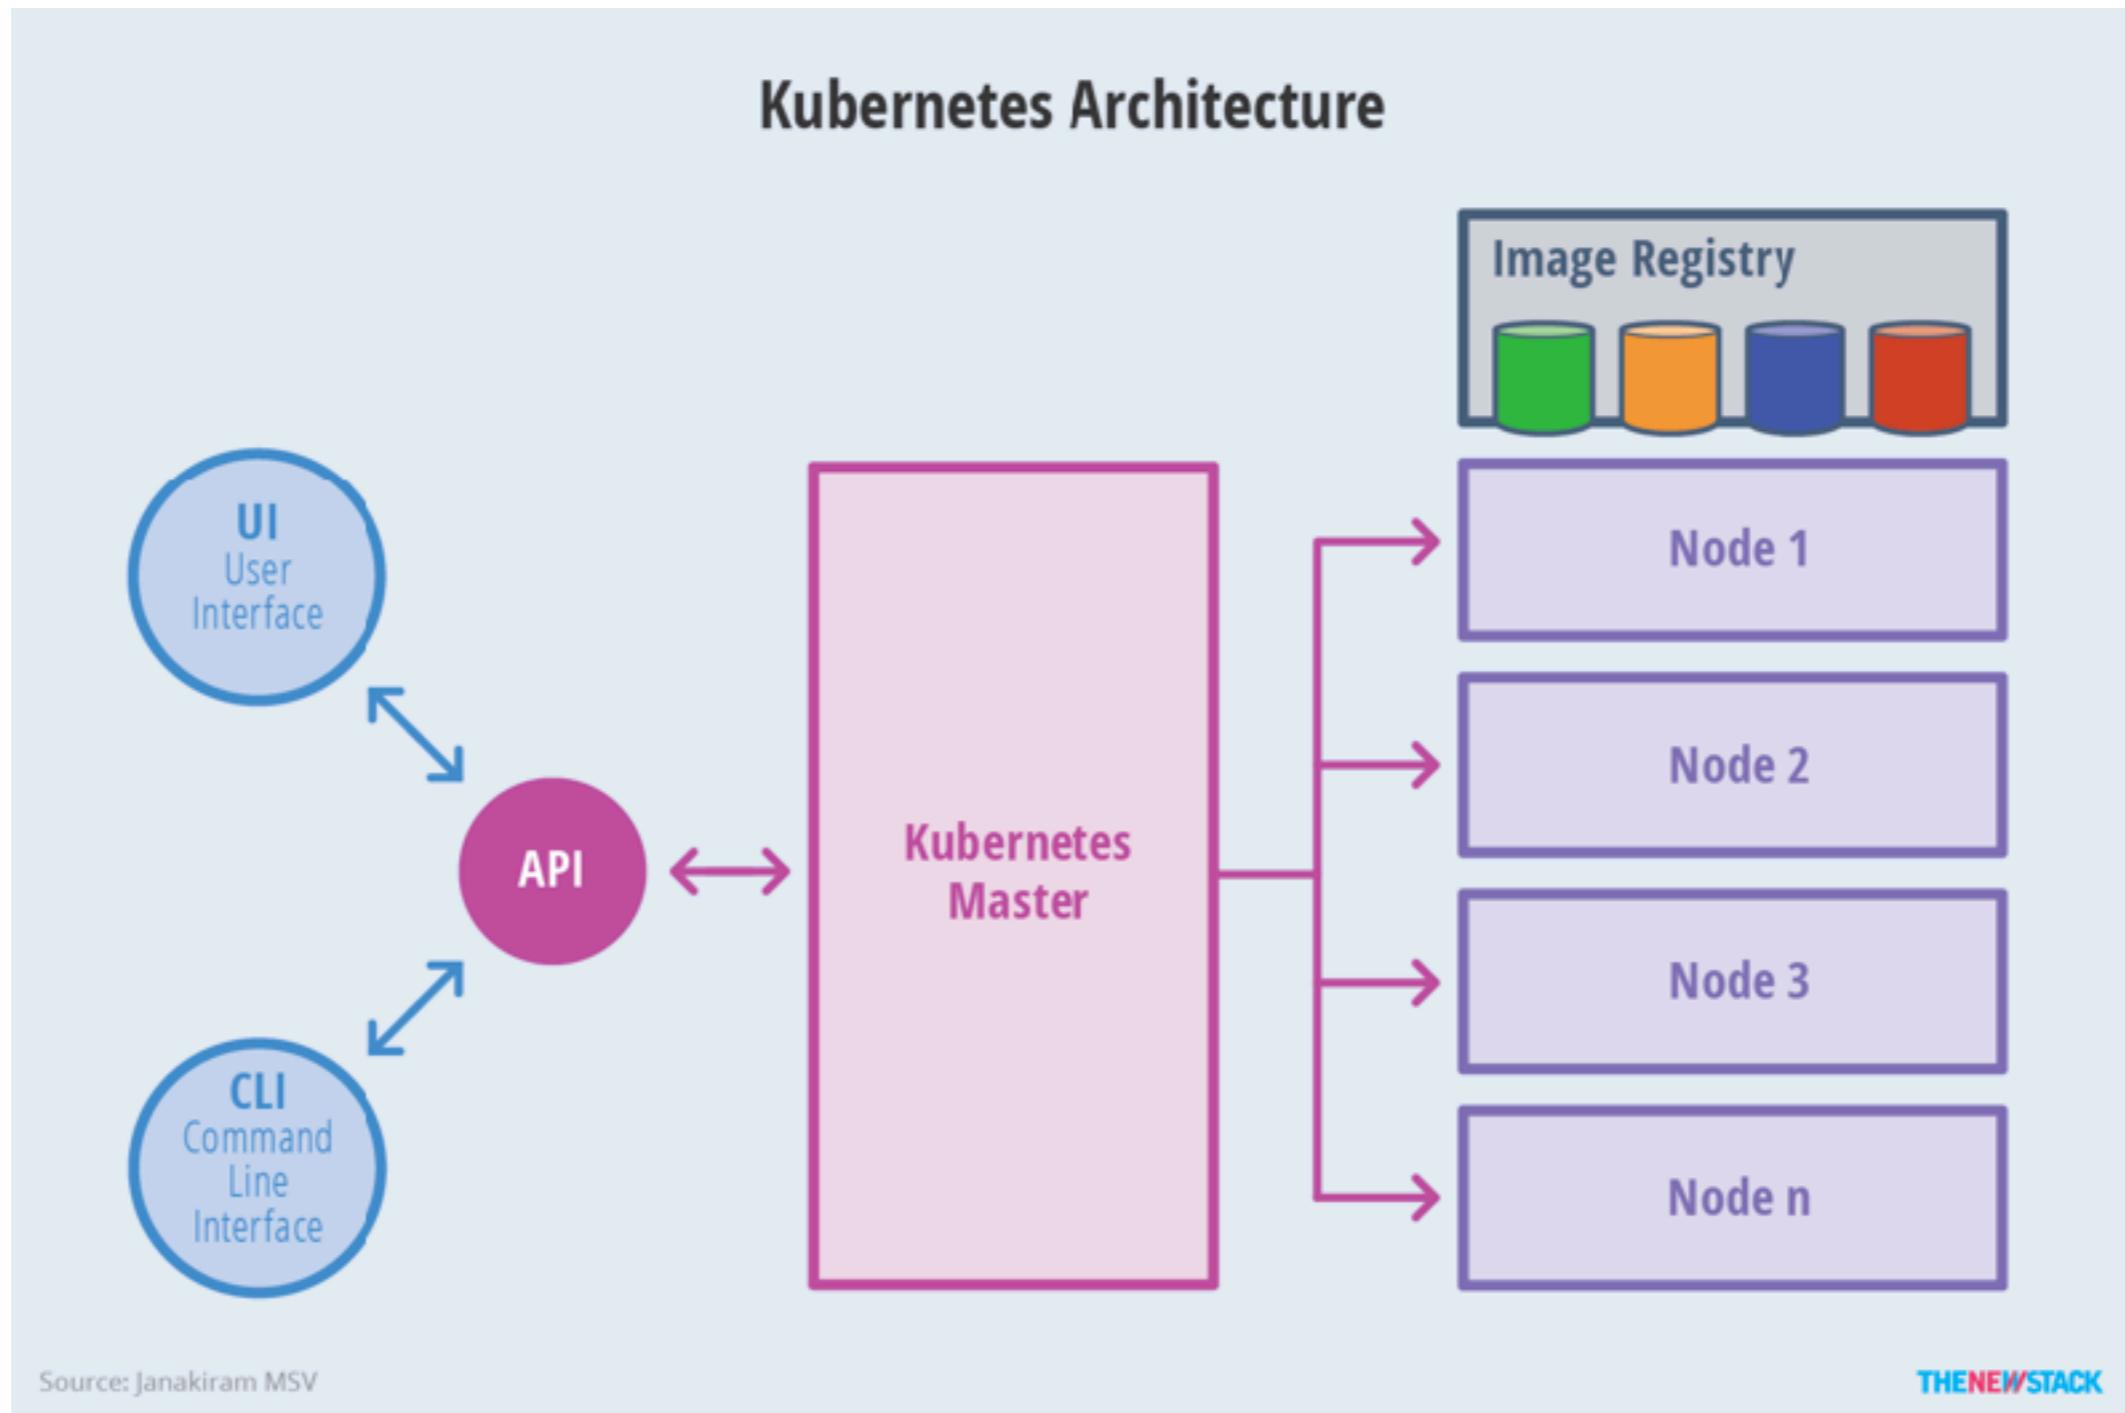
<https://12factor.net/>



# System architecture



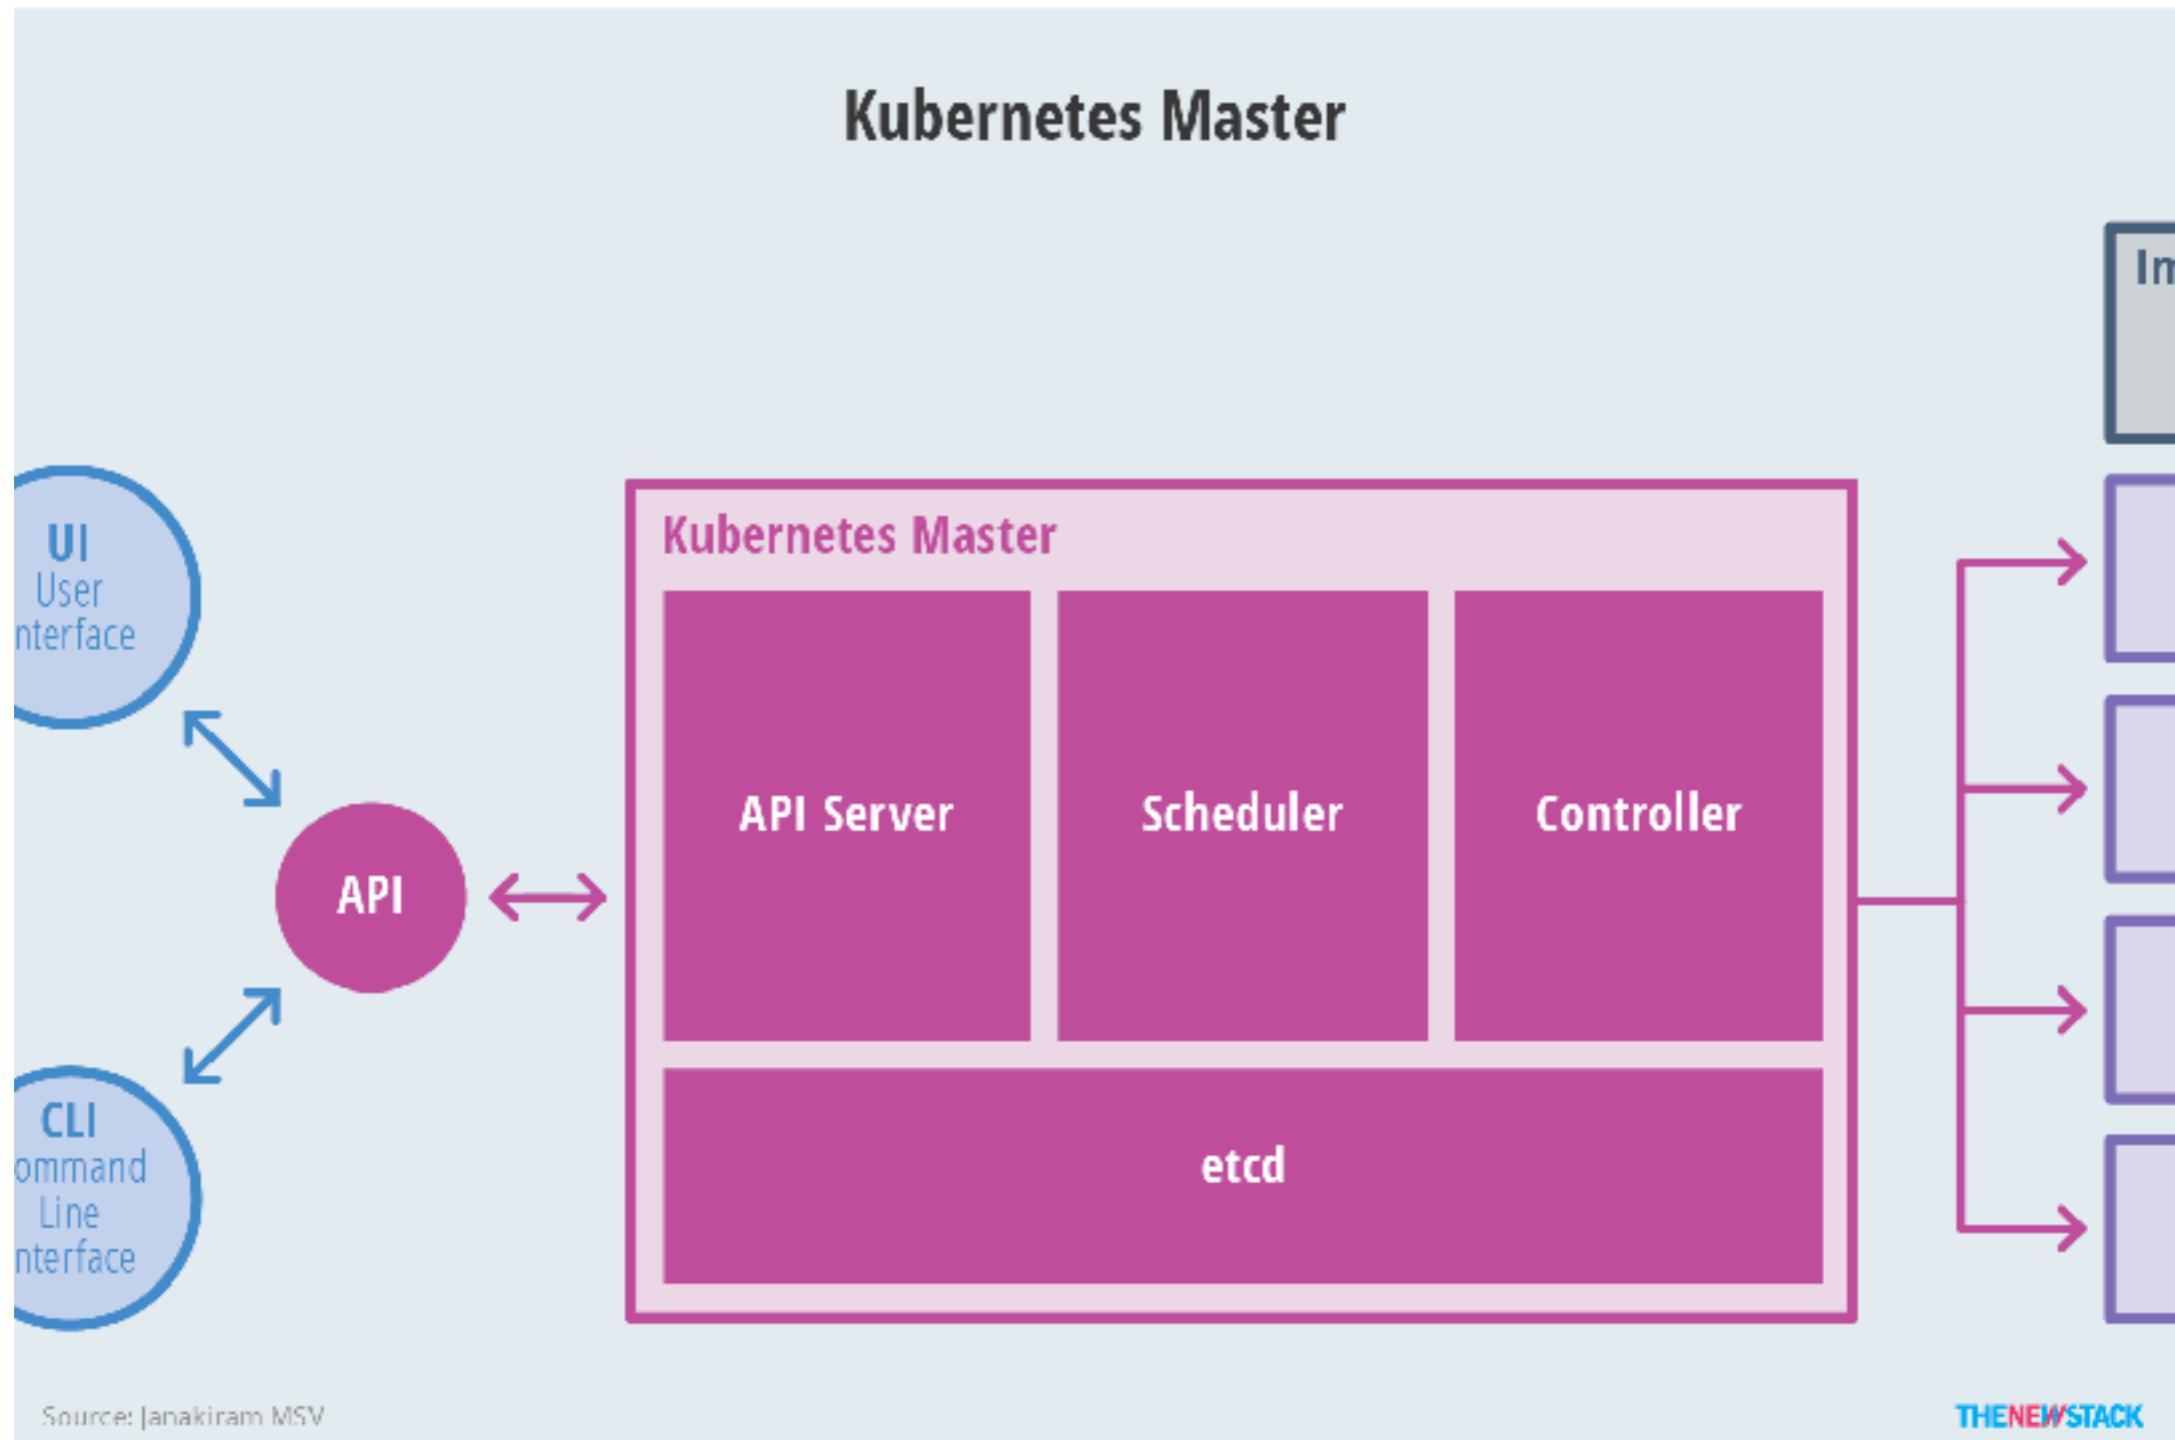
# System architecture



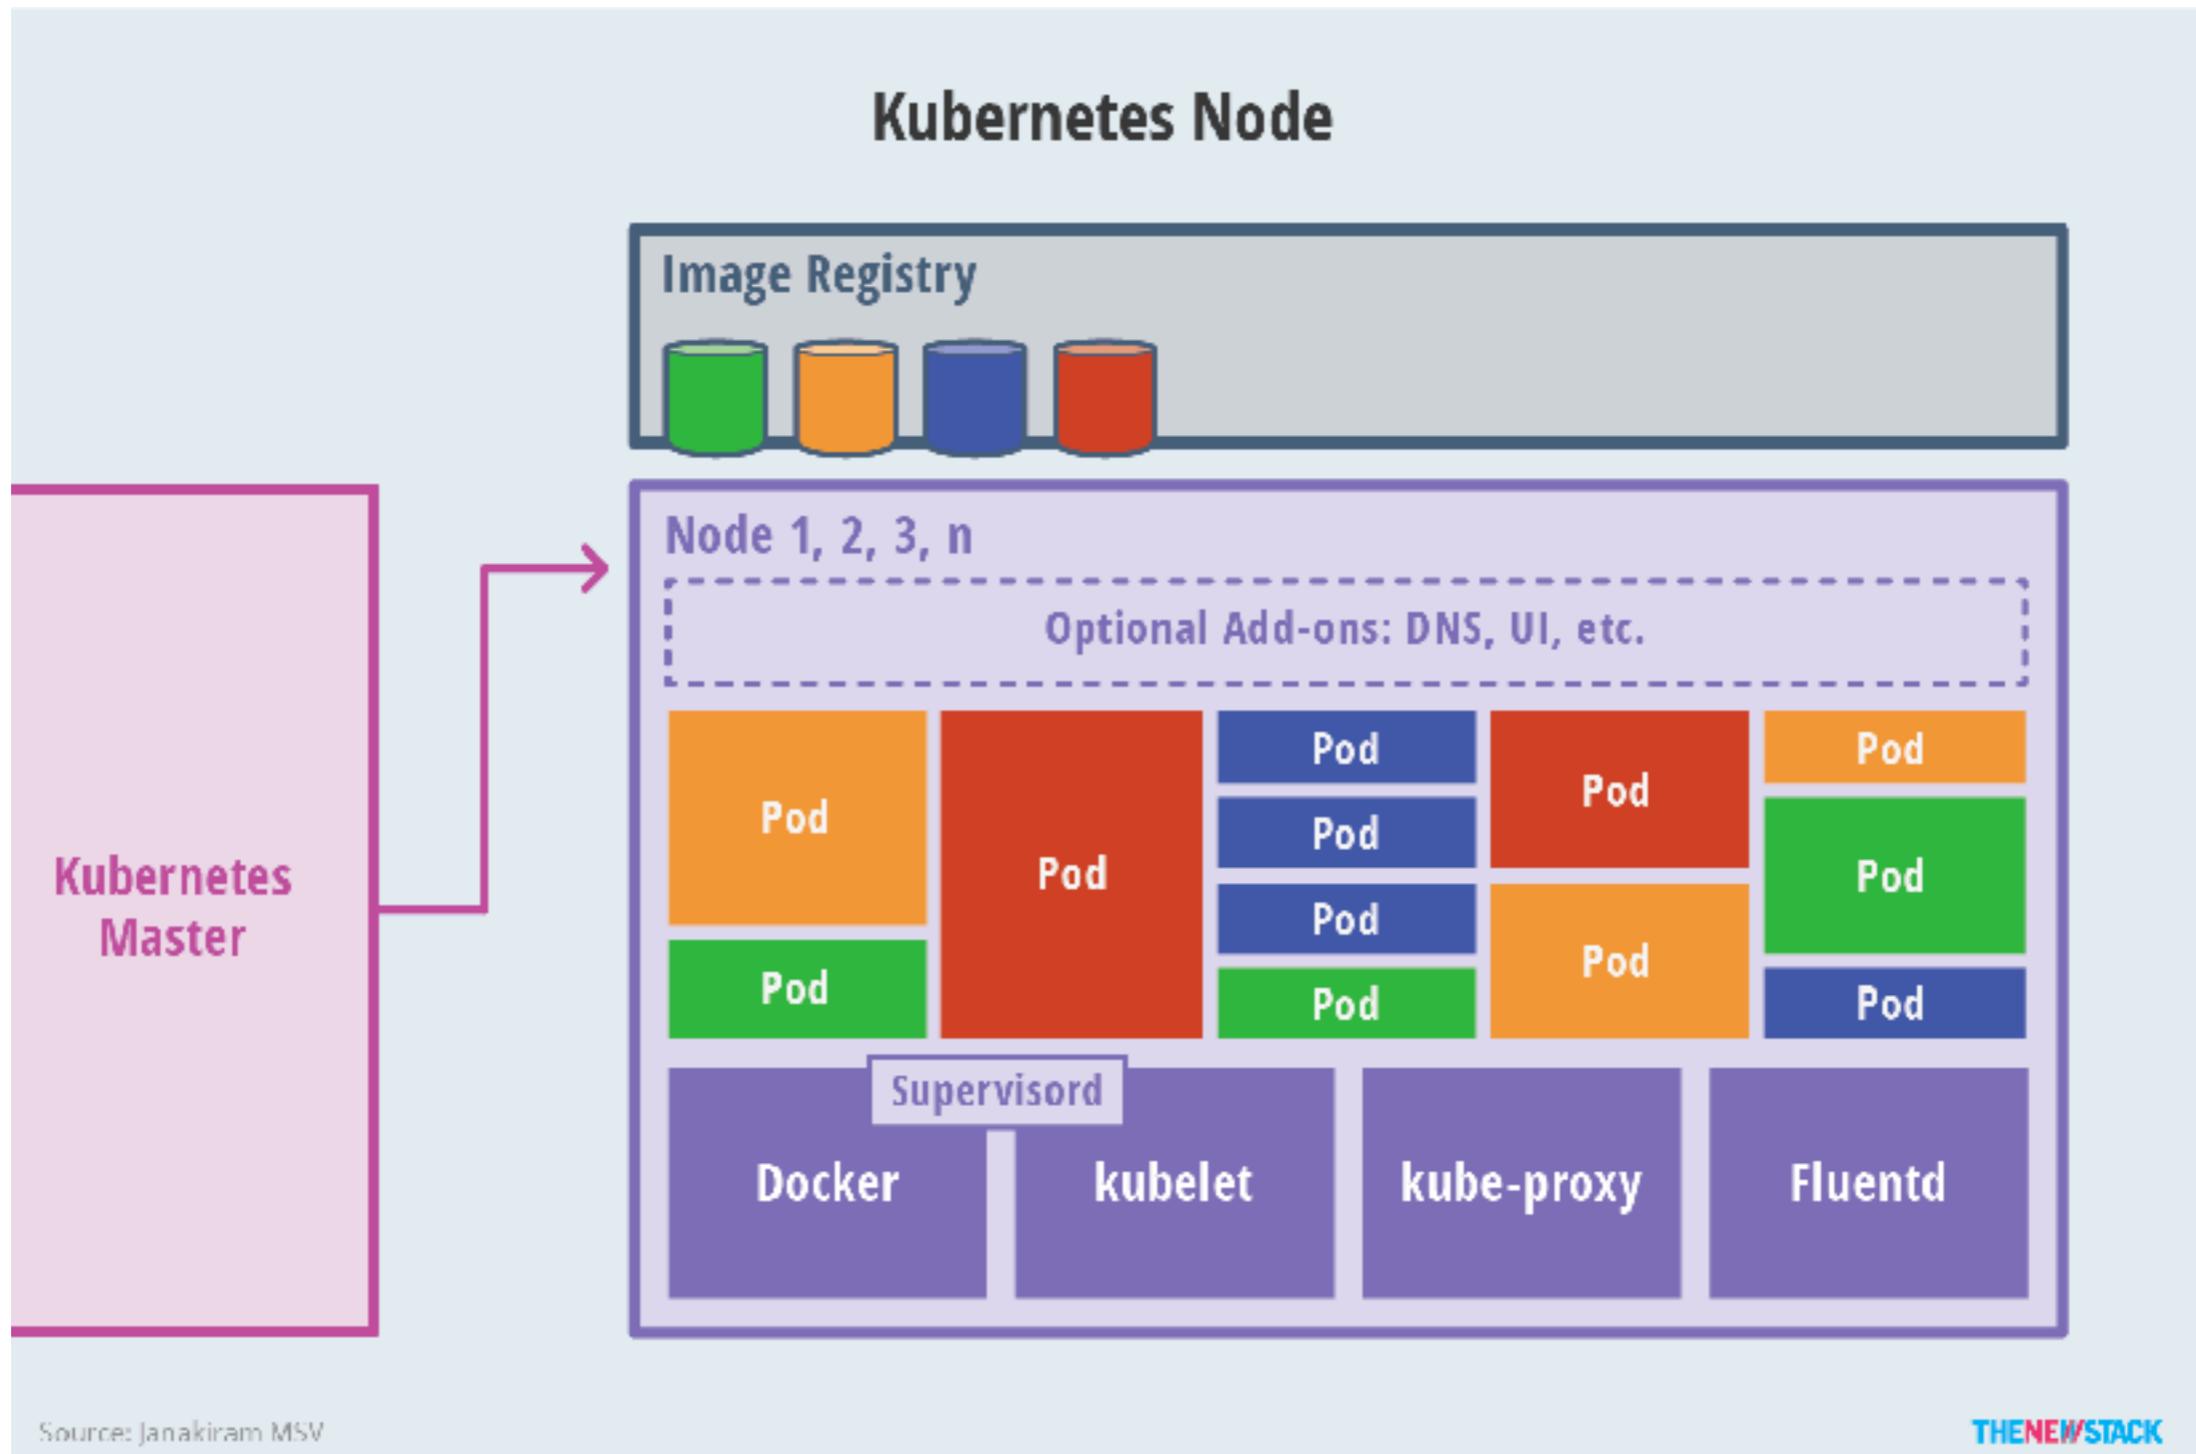
<https://thenewstack.io/kubernetes-an-overview/>



# Master



# Node



<https://thenewstack.io/kubernetes-an-overview/>



# Core concepts of Kubernetes



# Core concepts

Pods vs containers

Services

Replication Controllers

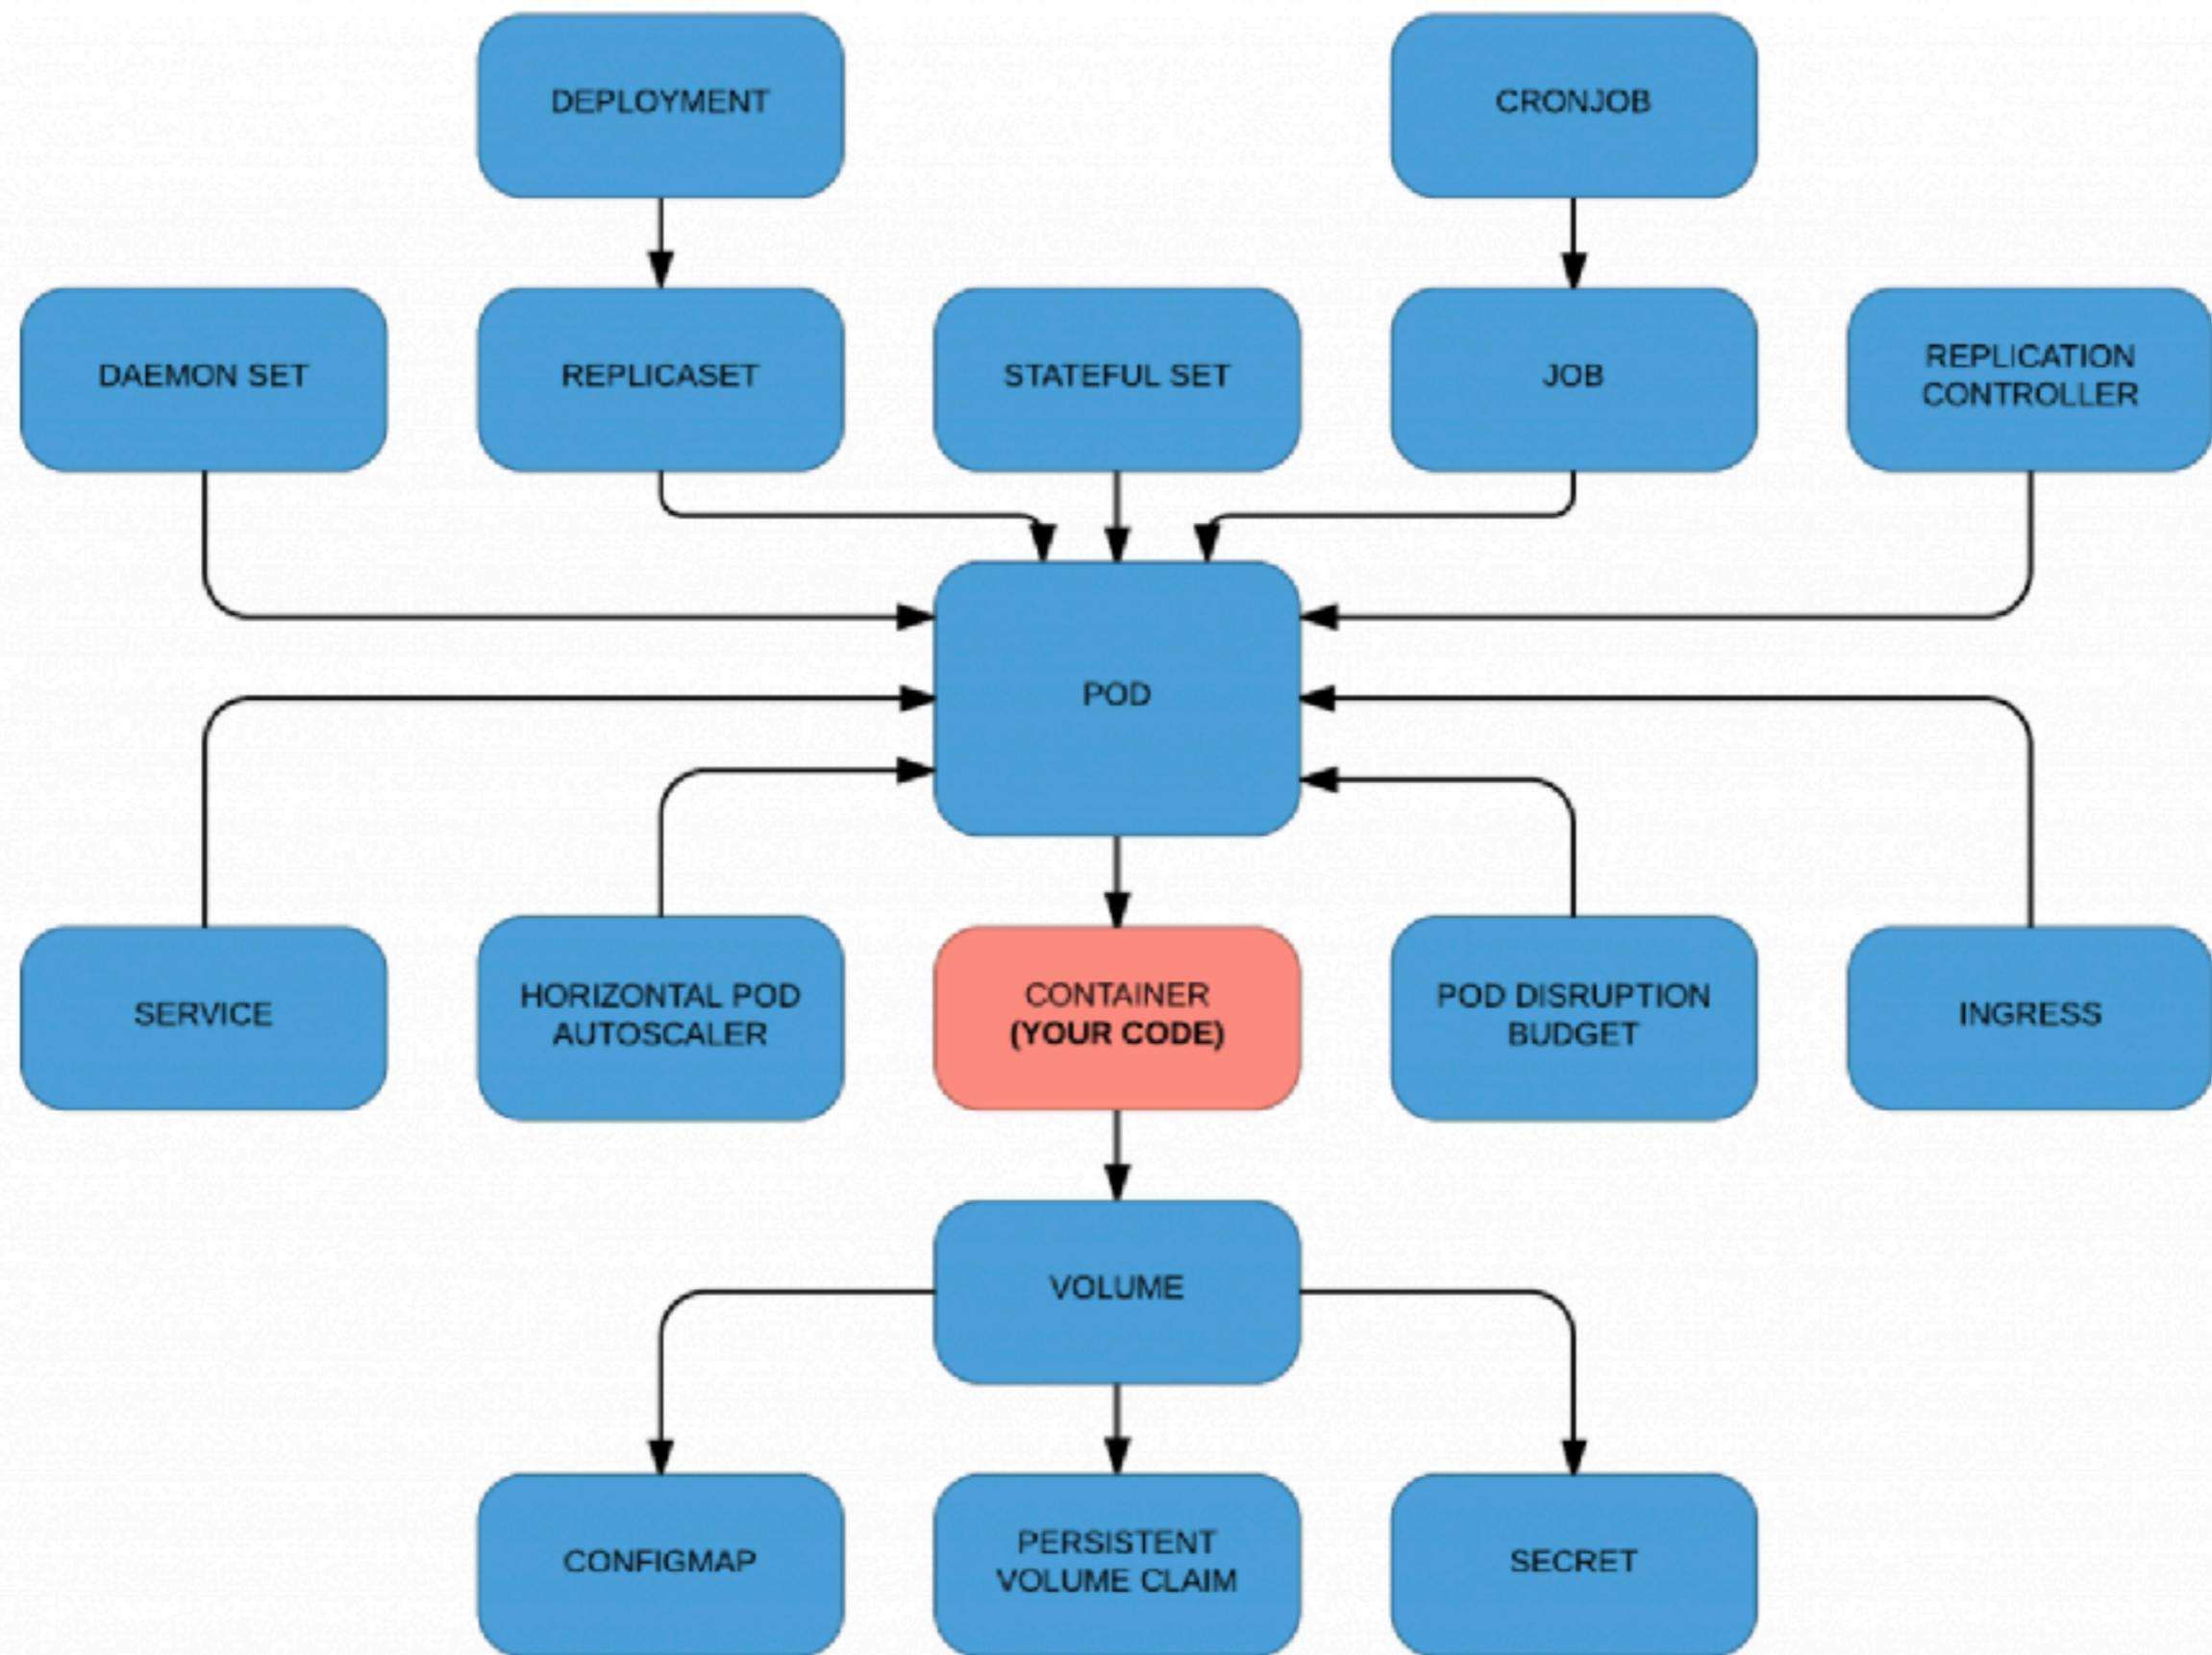
Deployments

ConfigMap and Secret

Volumes

StatefulSets



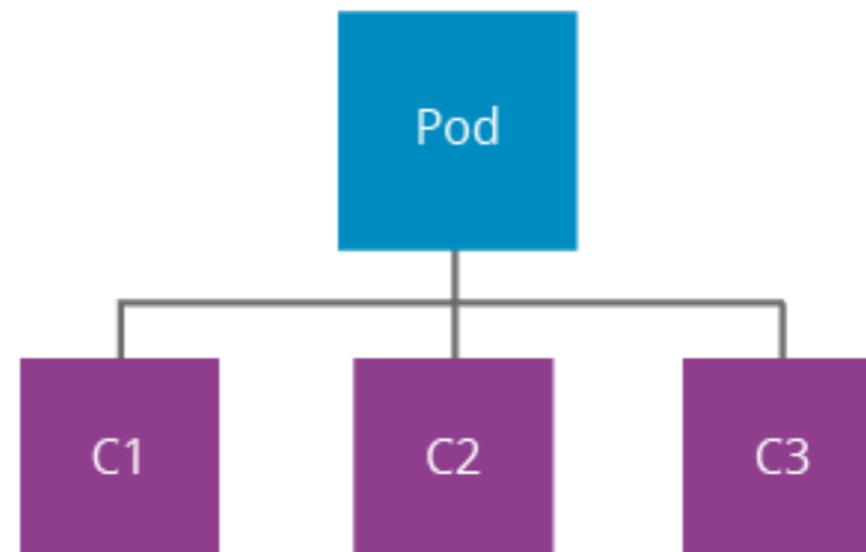


# Pods vs containers



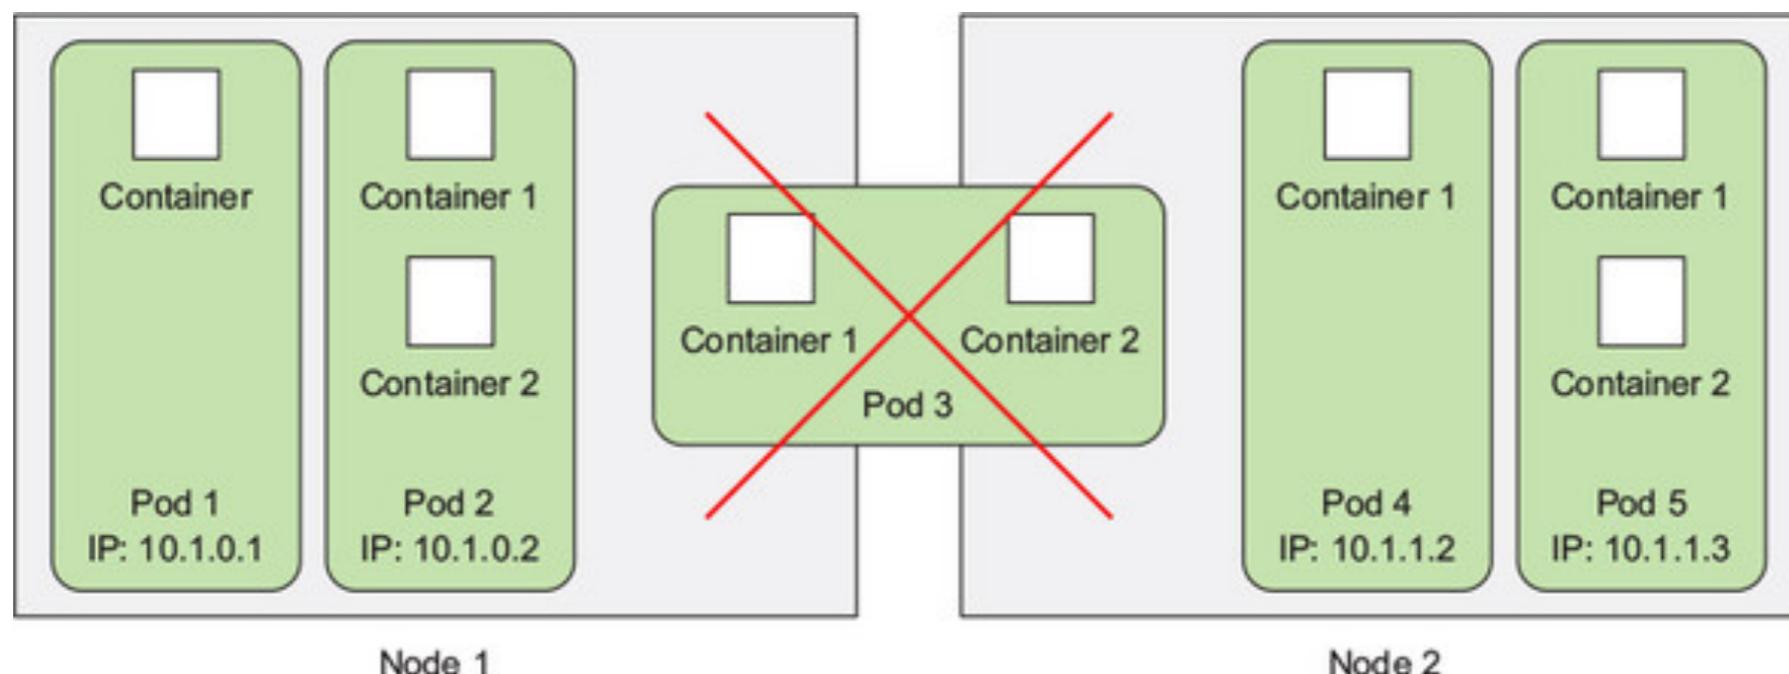
# Pods

Small group of co-located containers  
Optionally shared volume between containers  
Basic deployment unit in Kubernetes



# Pods

1 pods = 1 container  
1 pods = N containers



# All containers in same Pods

Share process ID

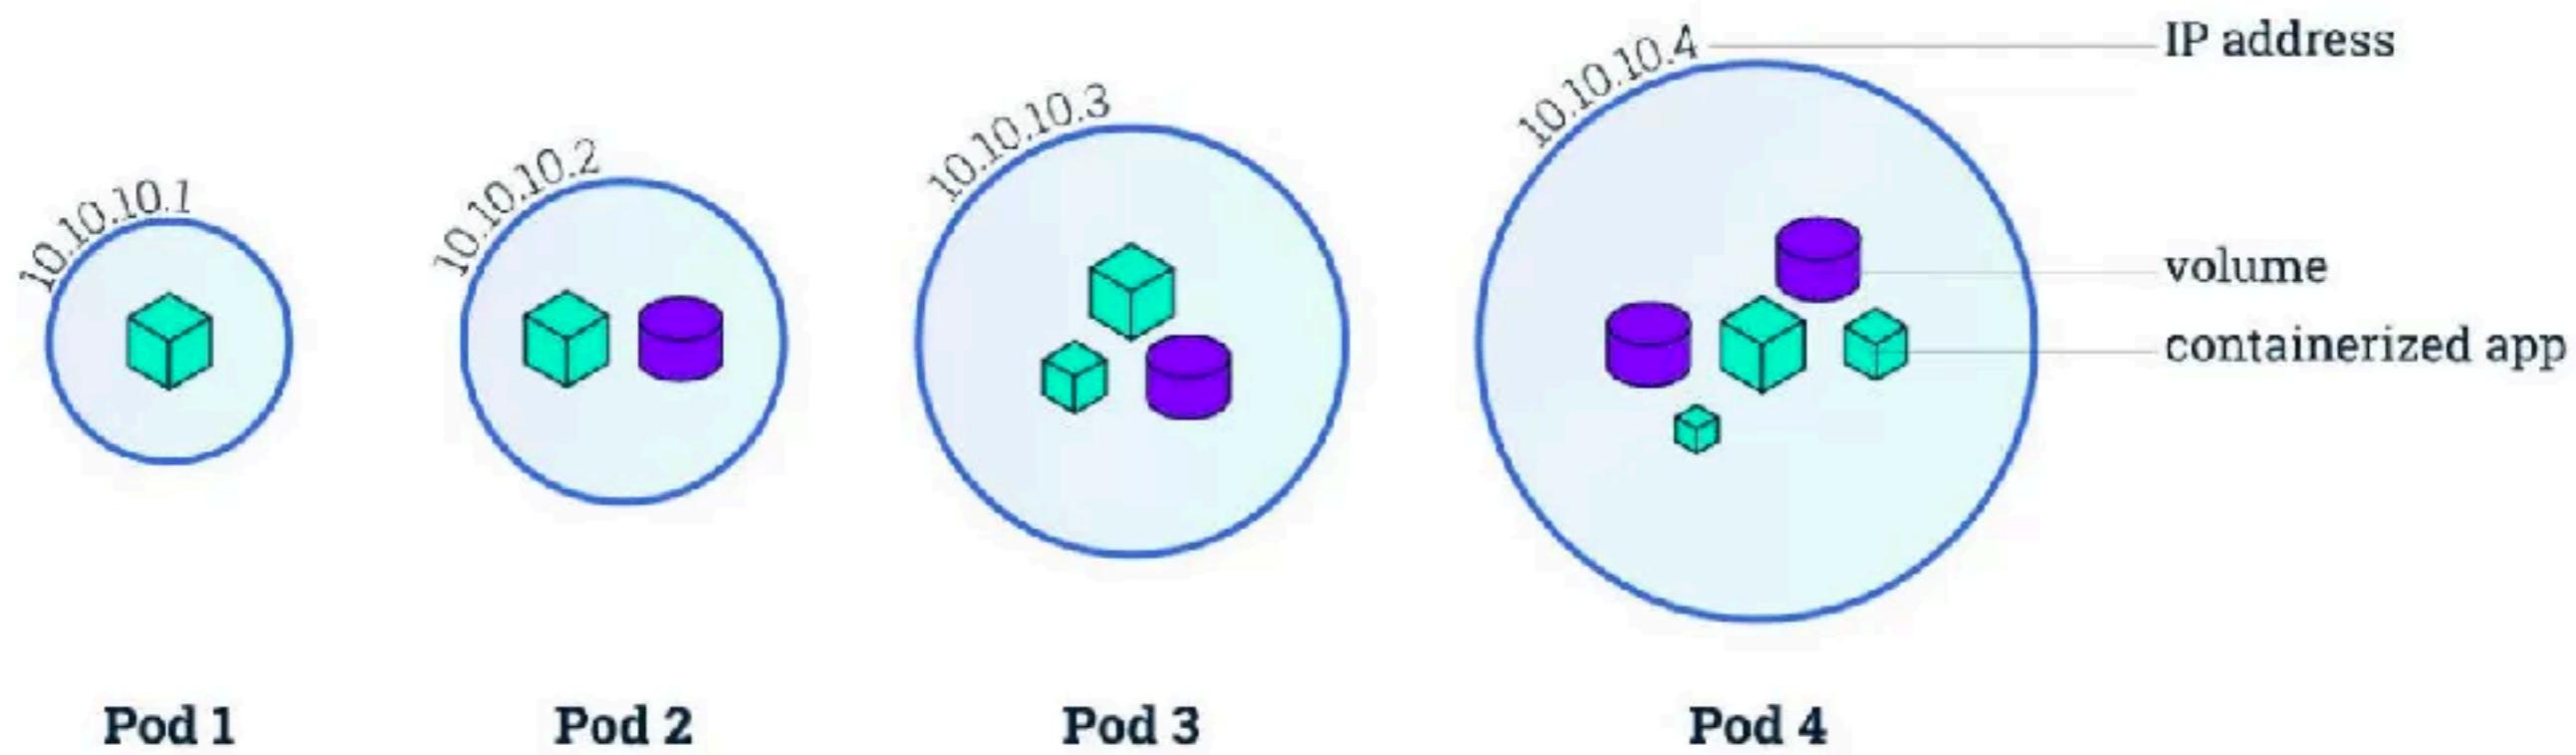
Share network interface

Share Hostname/IP/Port

Share Unix Time Sharing (UTS)

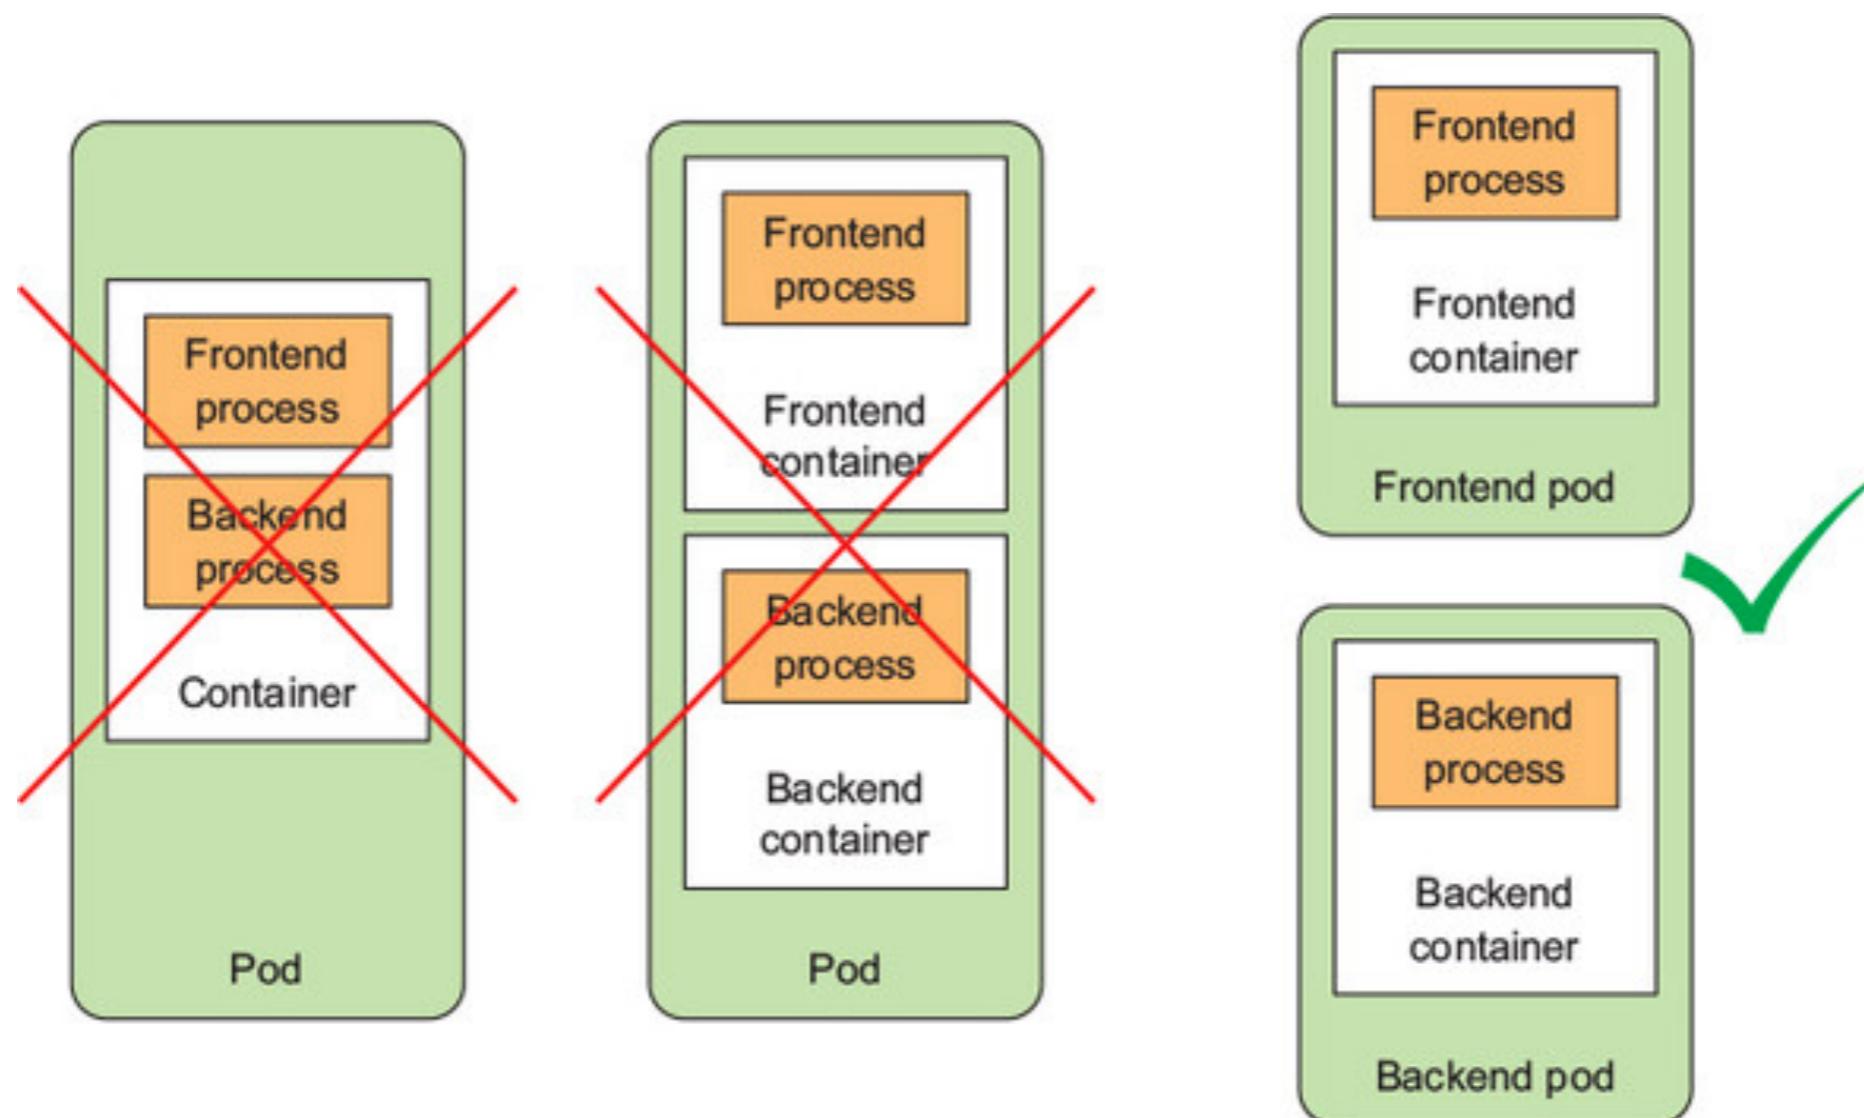


# Pods



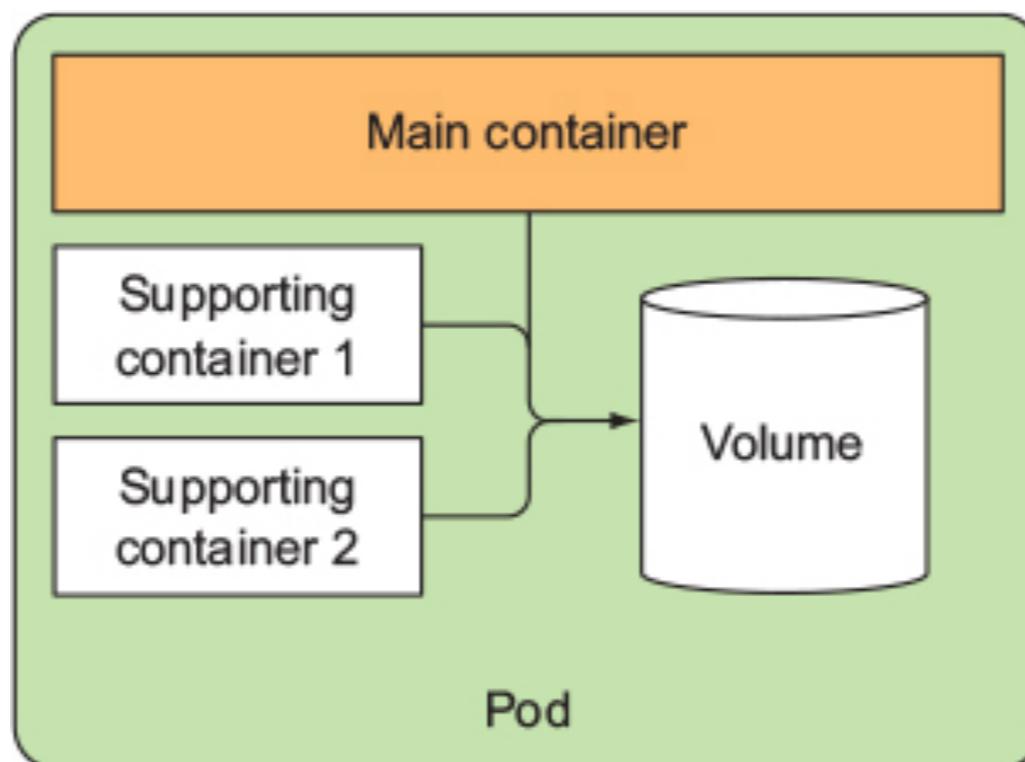
# Organize container across Pods

Split multi-tiers app into multiple pods



# Organize container across Pods

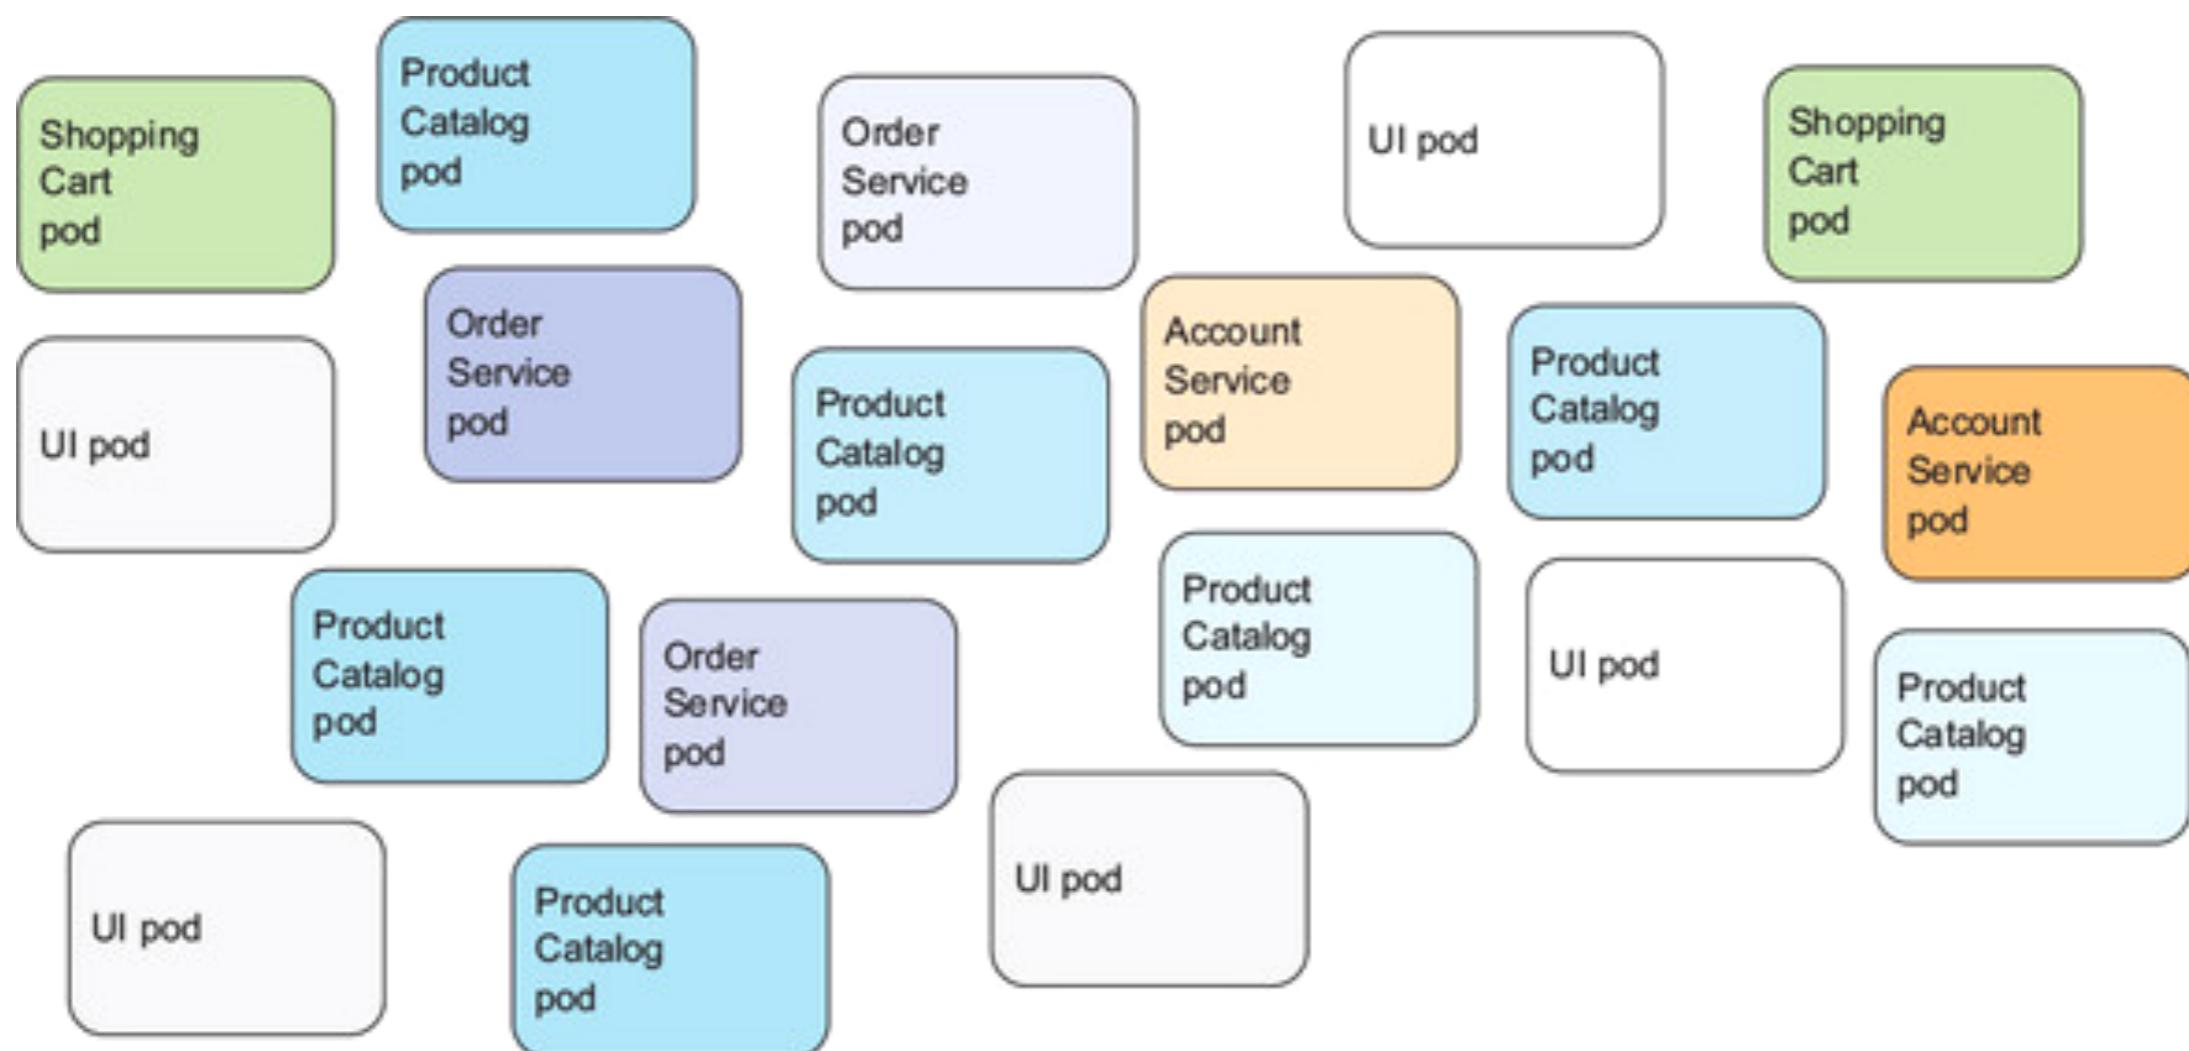
When to use multiple containers in a pods



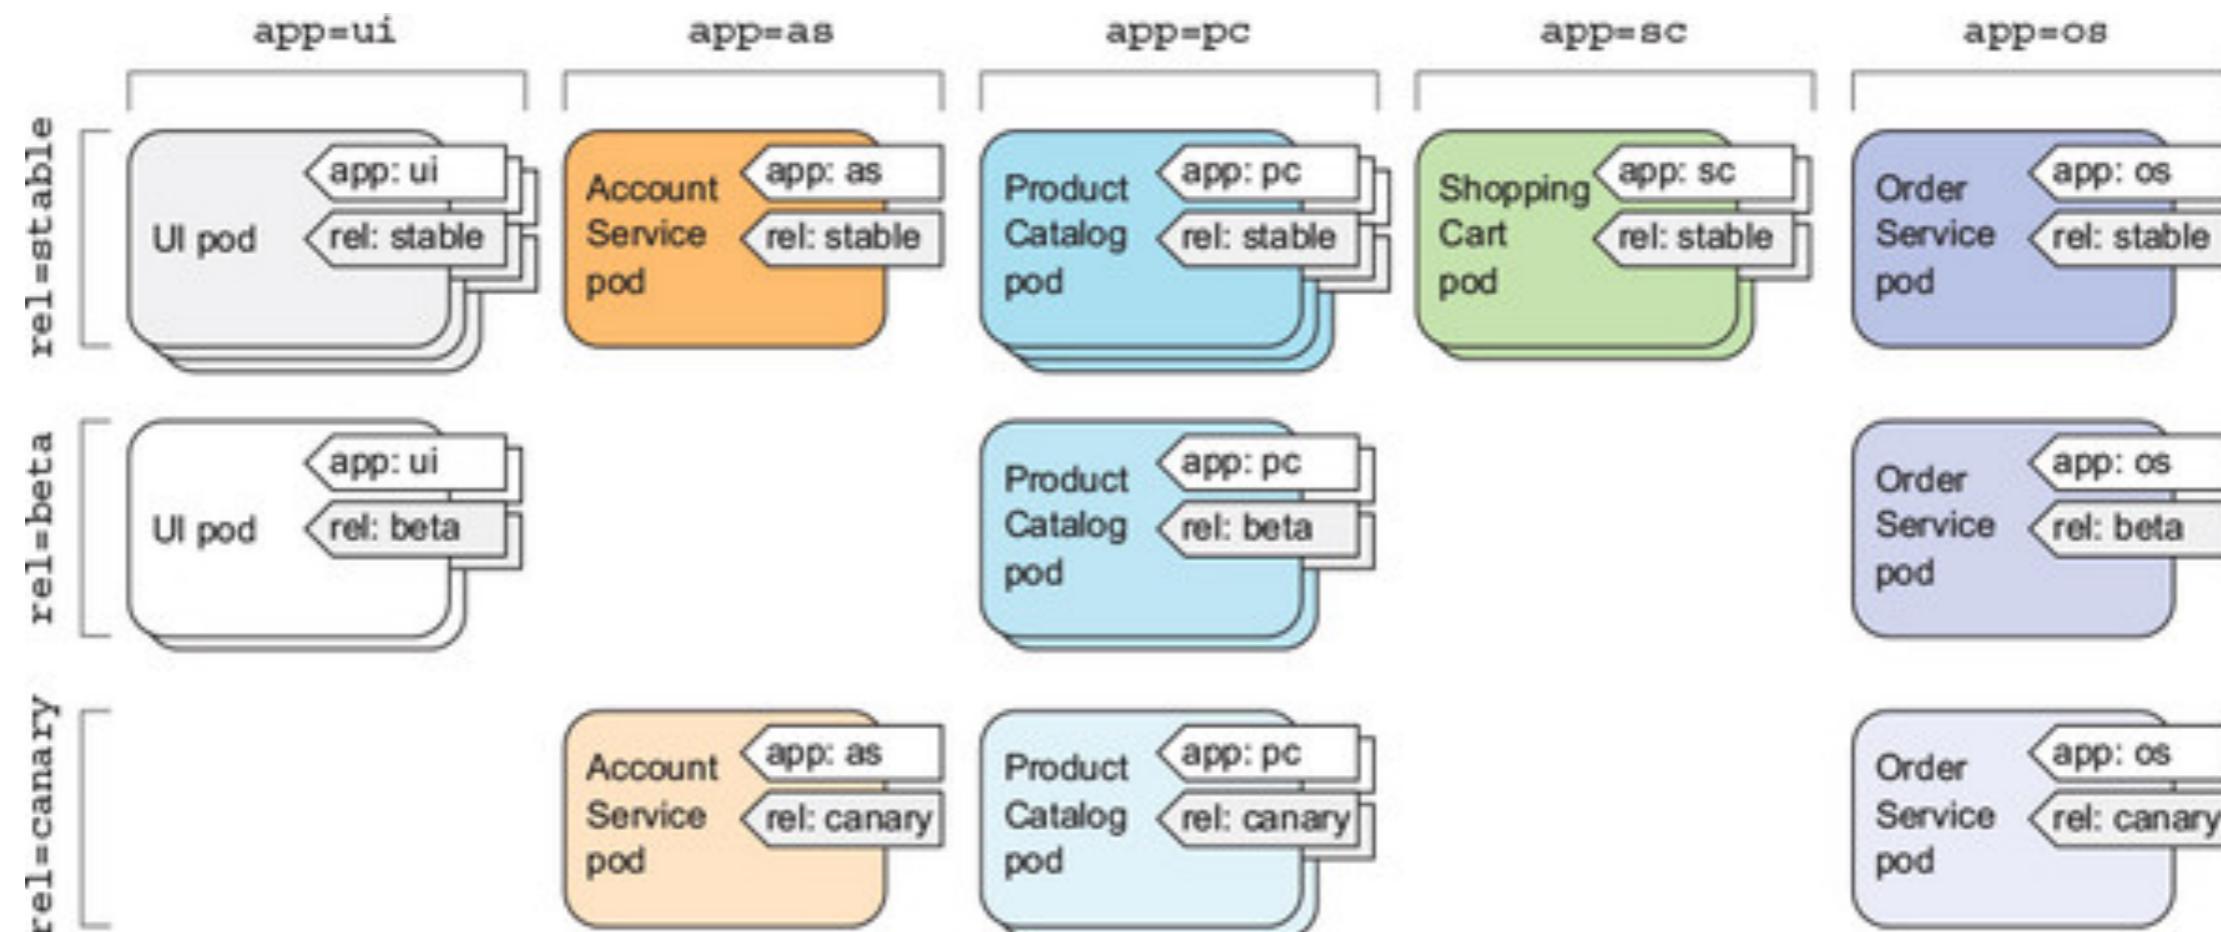
# Organize pods with Labels



# Organize pods with Labels

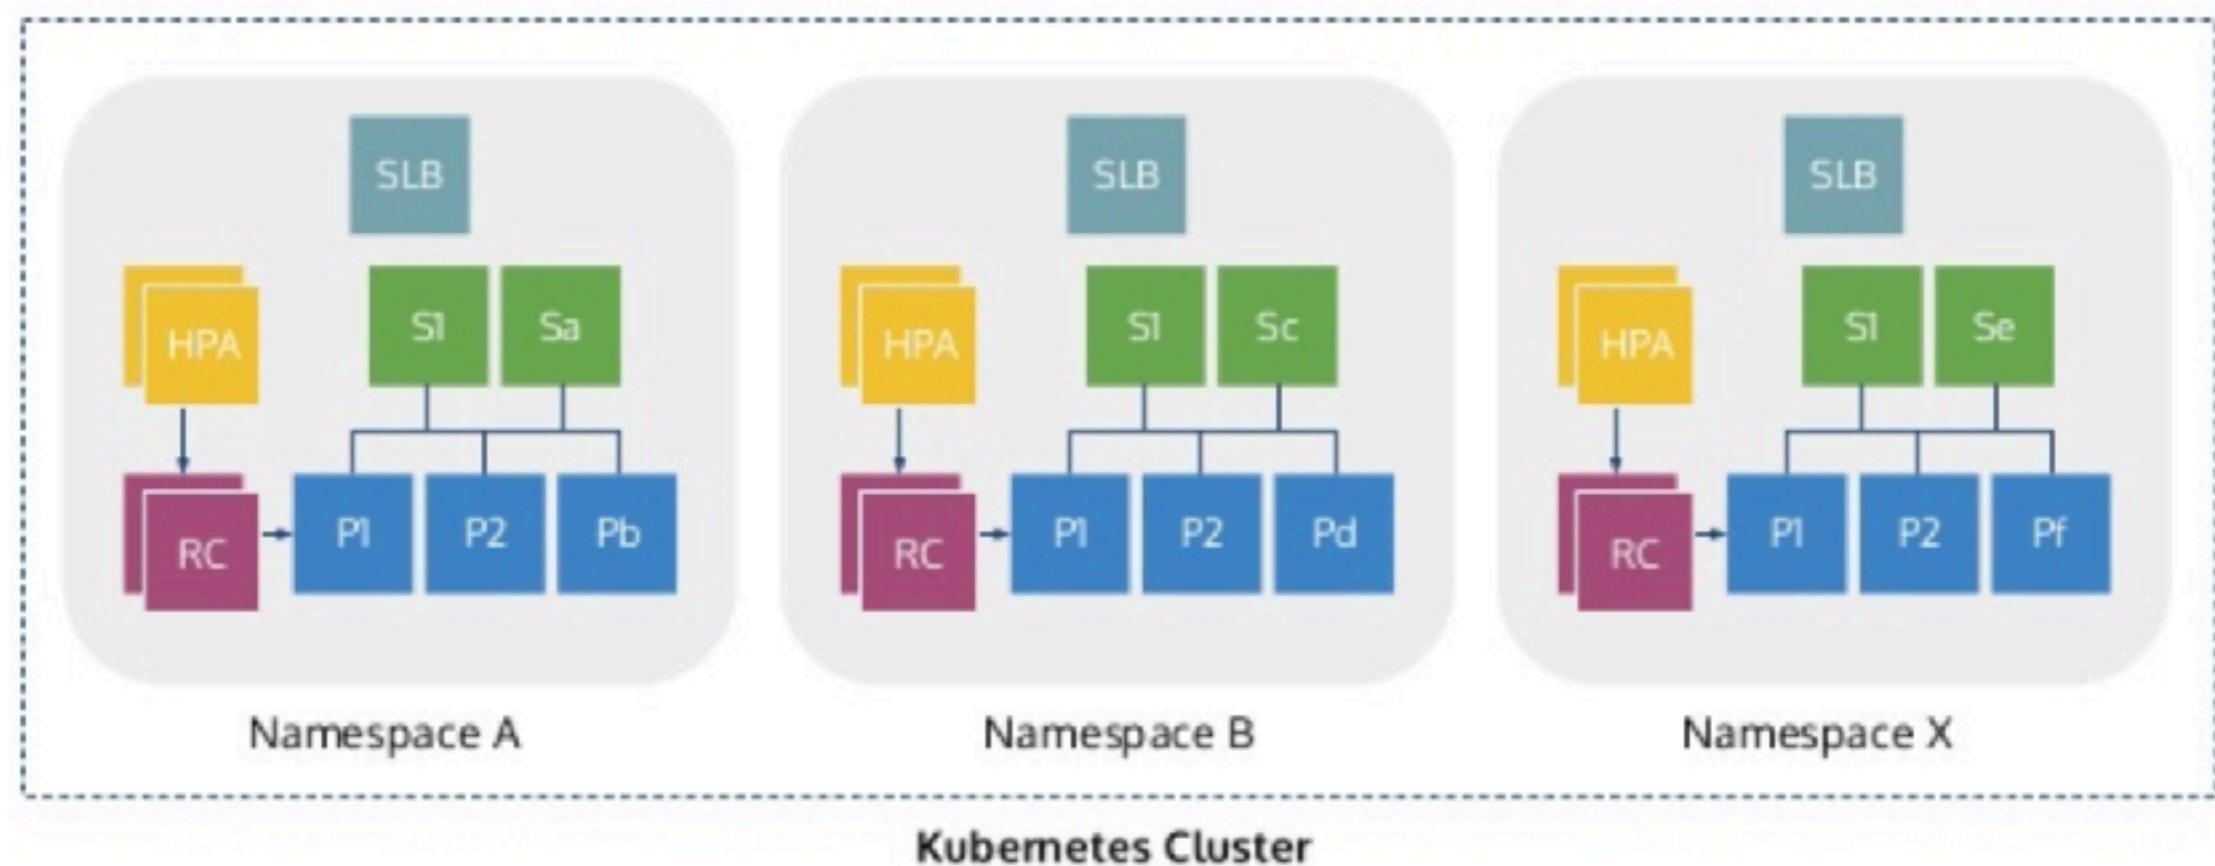


# Organize pods with Labels



# Pods namespaces

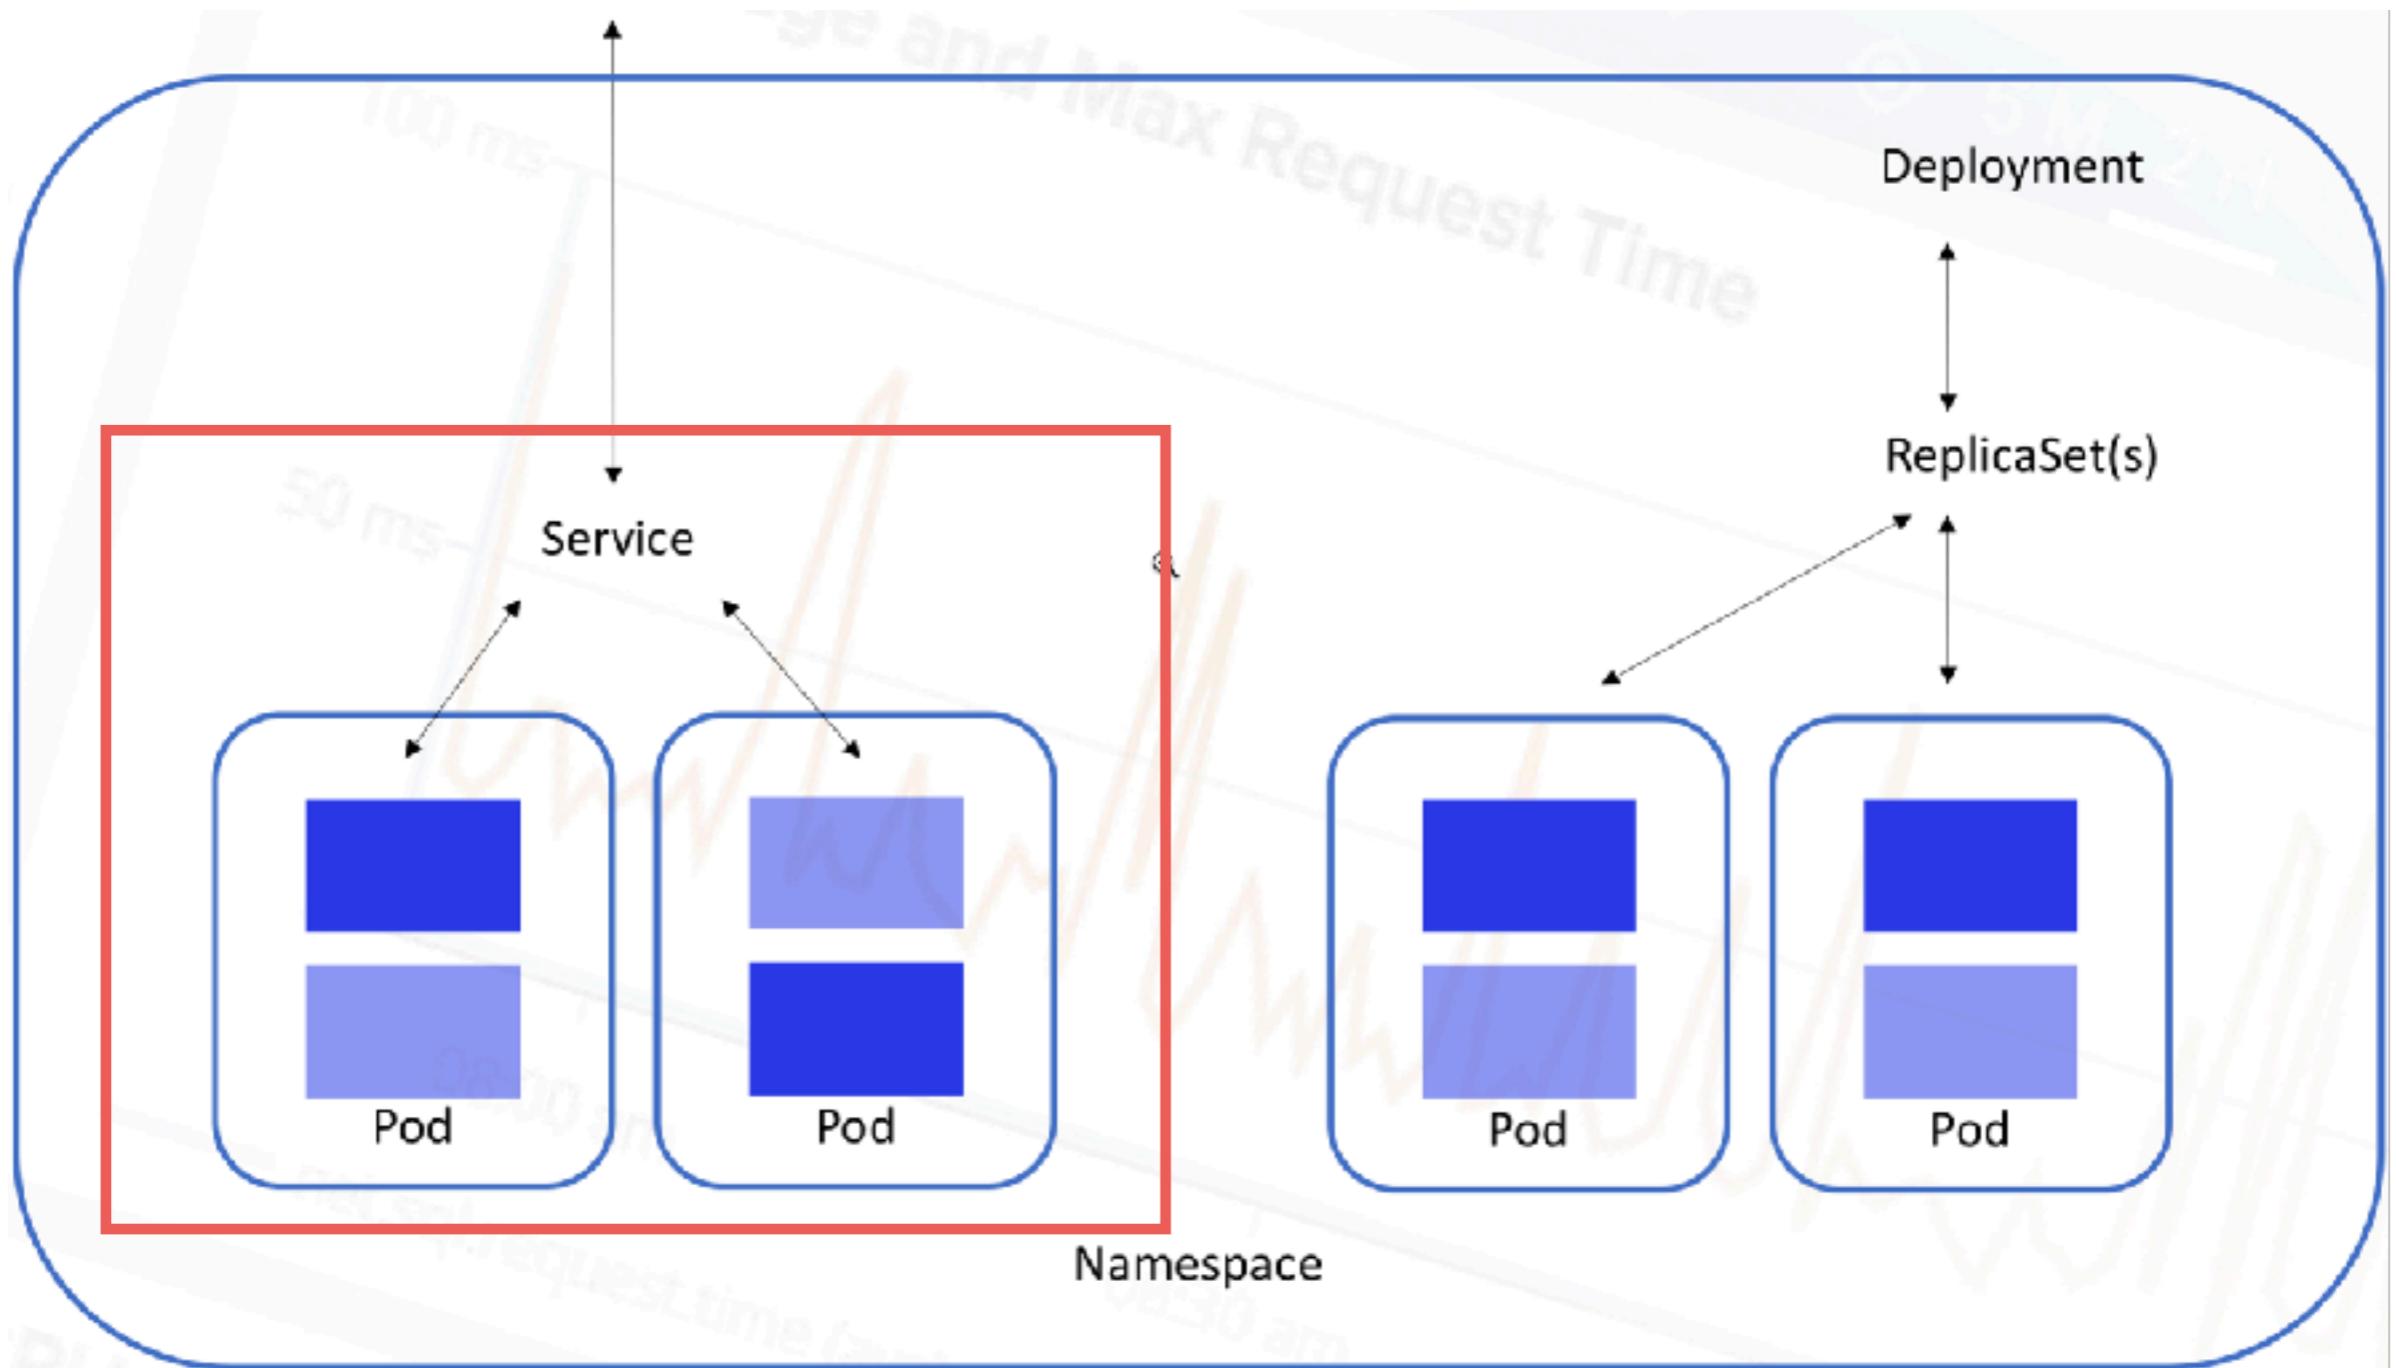
Allow different teams to use the same cluster



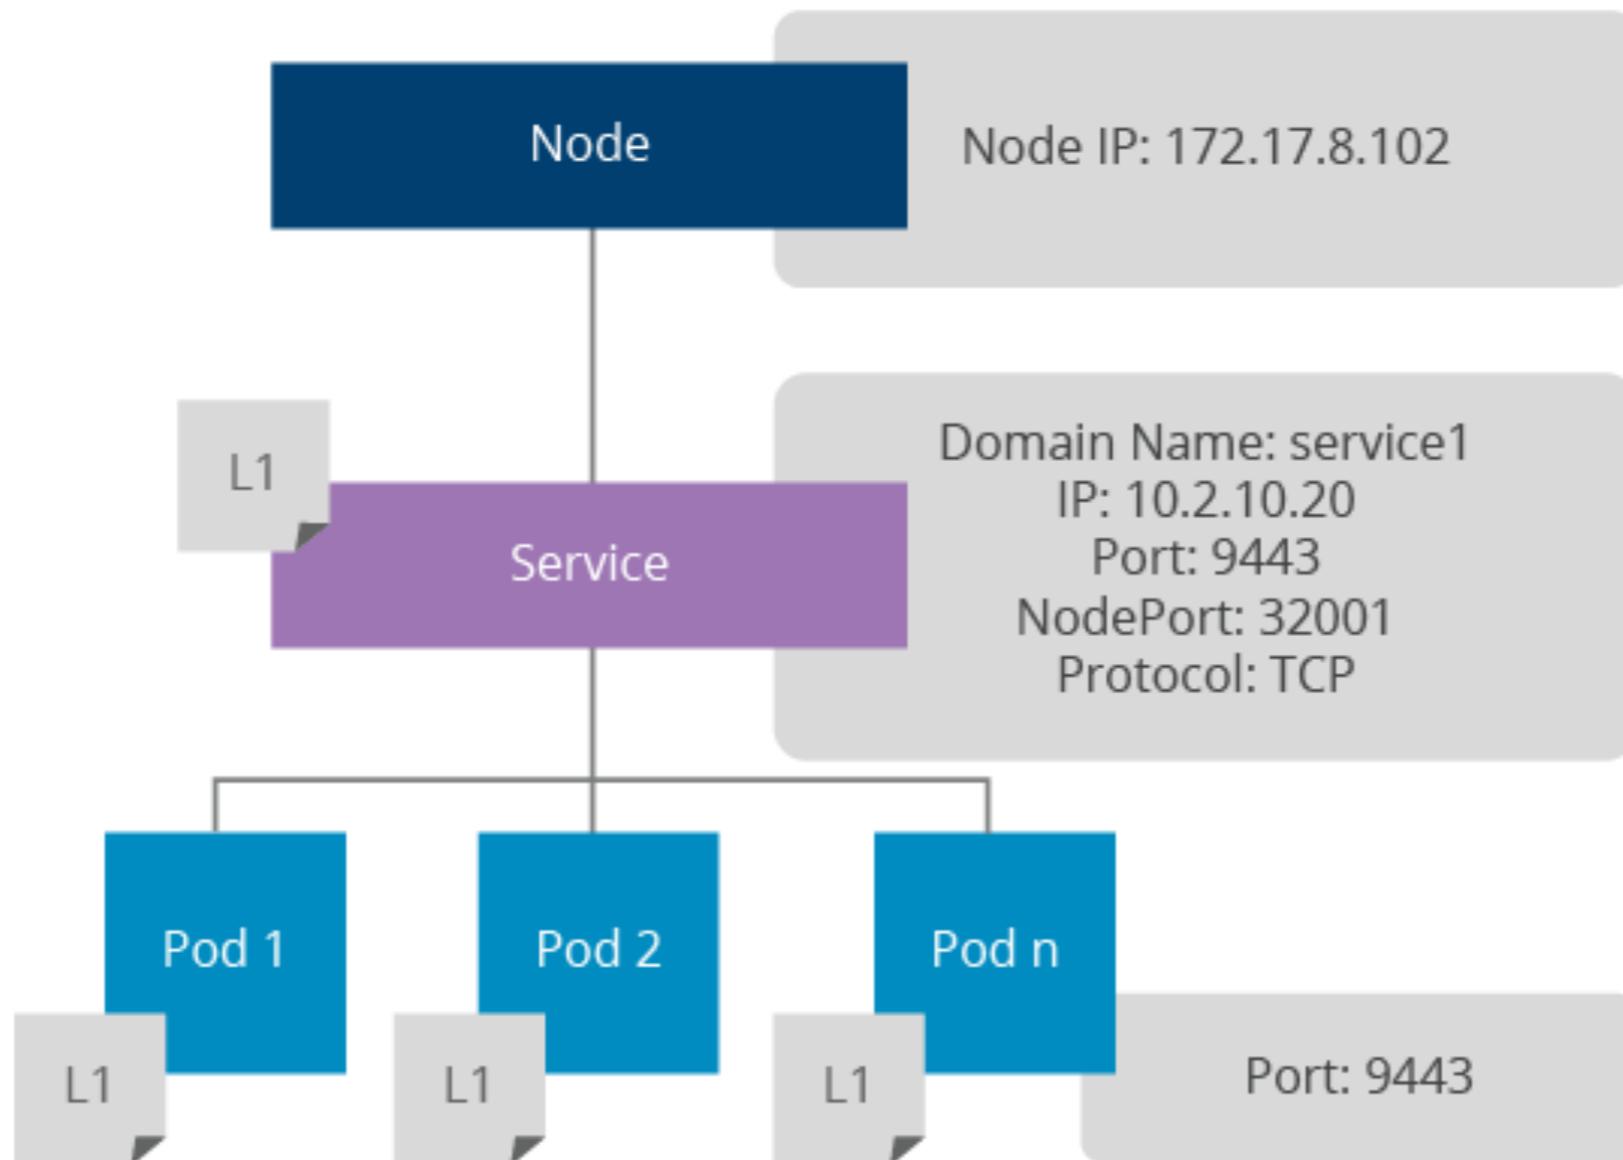
# **Services discovery and Load balancing**



# Services



# Services

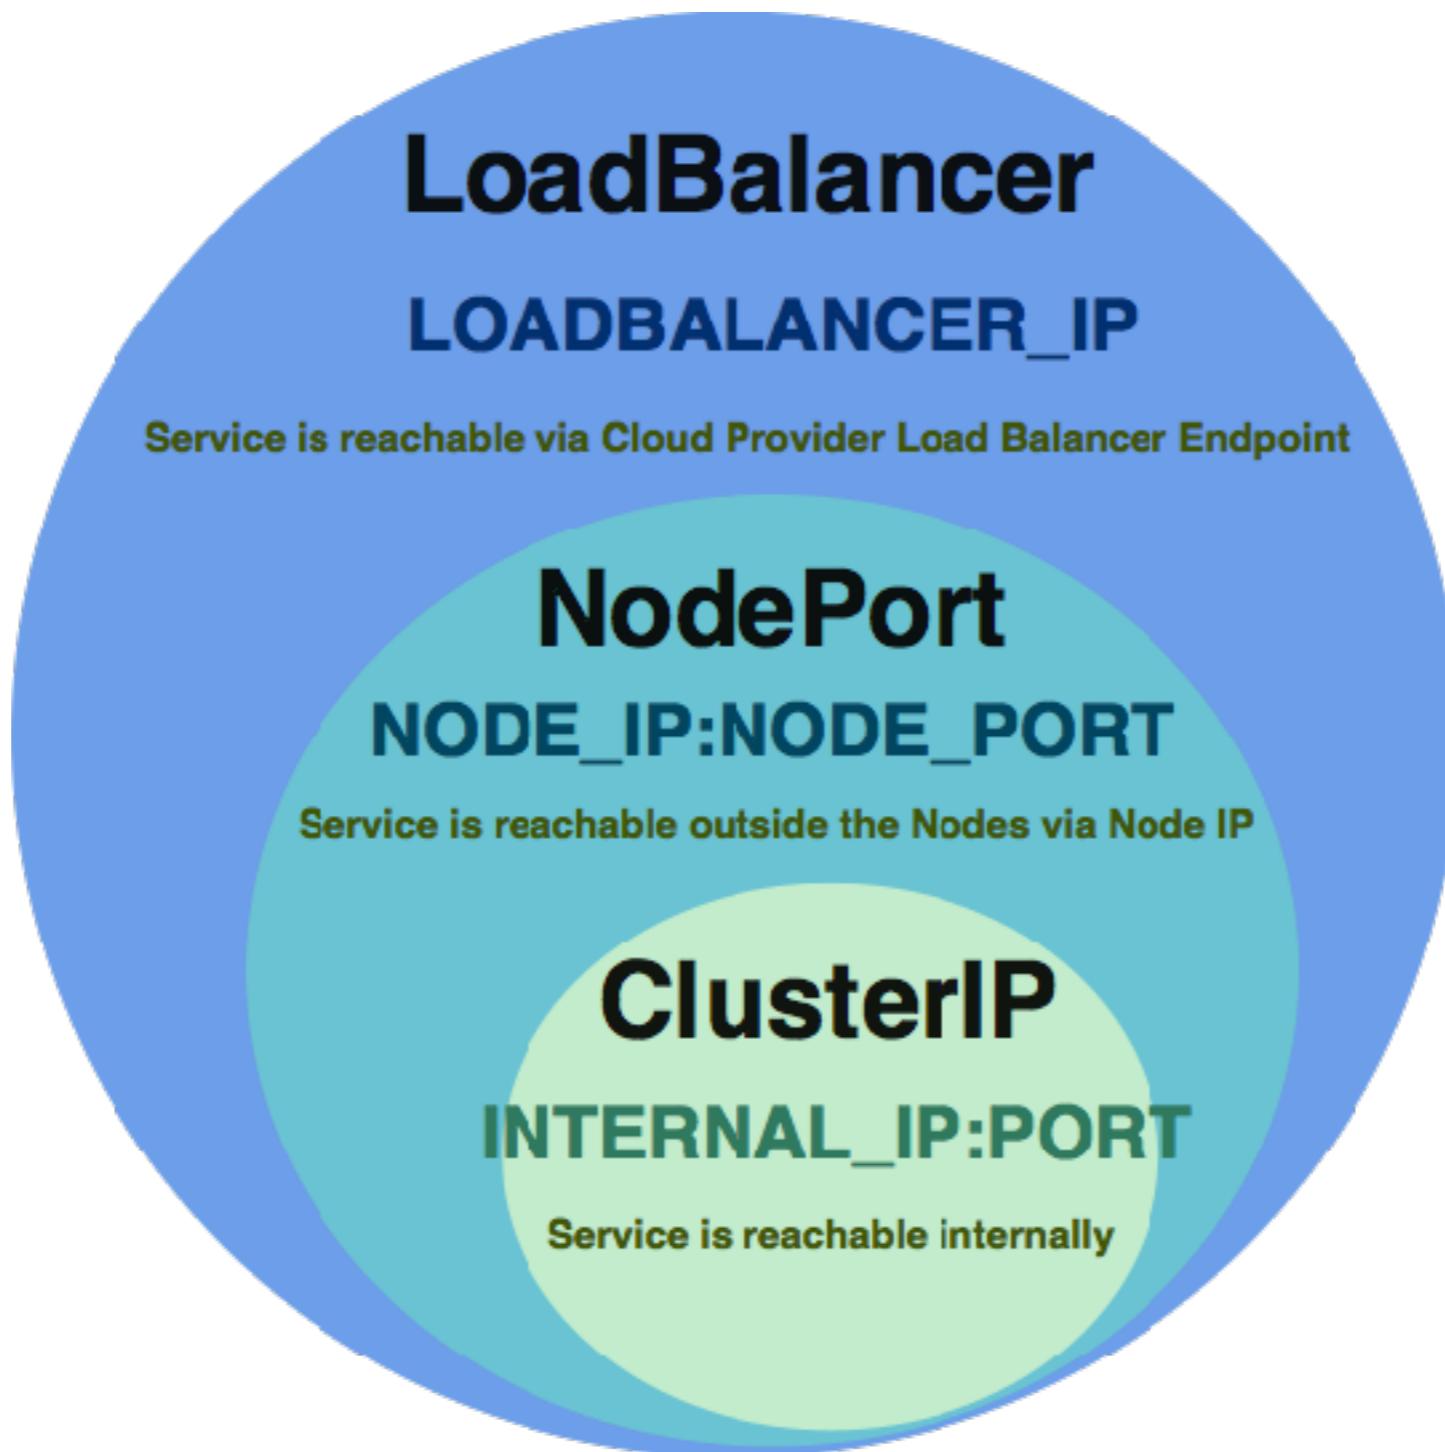


# Services

- Independent from Pods
- Abstraction layer of Pods
- Provide load balance
- Expose access Pods/Load balance
- Find Pods by label selector



# 3 types of services



# Kubernetes Object Management



# Kubernetes Object Management

1. Imperative command
2. Imperative object configuration
3. Declarative object configuration

<https://kubernetes.io/docs/concepts/overview/object-management-kubectl/overview/>



# 1. Imperative command

```
$kubectl run nginx --image nginx
```

```
$kubectl create deployment nginx --image nginx
```



## 2. Imperative object configuration

```
$kubectl create -f nginx.yaml
```

```
$kubectl delete -f nginx.yaml
```



# 3. Declarative object configuration

```
$kubectl apply -f configs/  
$kubectl apply -R -f configs/
```



# Kubernetes Object Management

Management technique	Operates on	Recommended environment	Supported writers	Learning curve
Imperative commands	Live objects	Development projects	1+	Lowest
Imperative object configuration	Individual files	Production projects	1	Moderate
Declarative object configuration	Directories of files	Production projects	1+	Highest

**Warning:** A Kubernetes object should be managed **using only one technique**. Mixing and matching techniques for the same object results in undefined behavior.



# Workshop :: Pods & Services

## 1 container per Pods

file /02-pod-service/single-pod/  
workshop\_instruction.txt



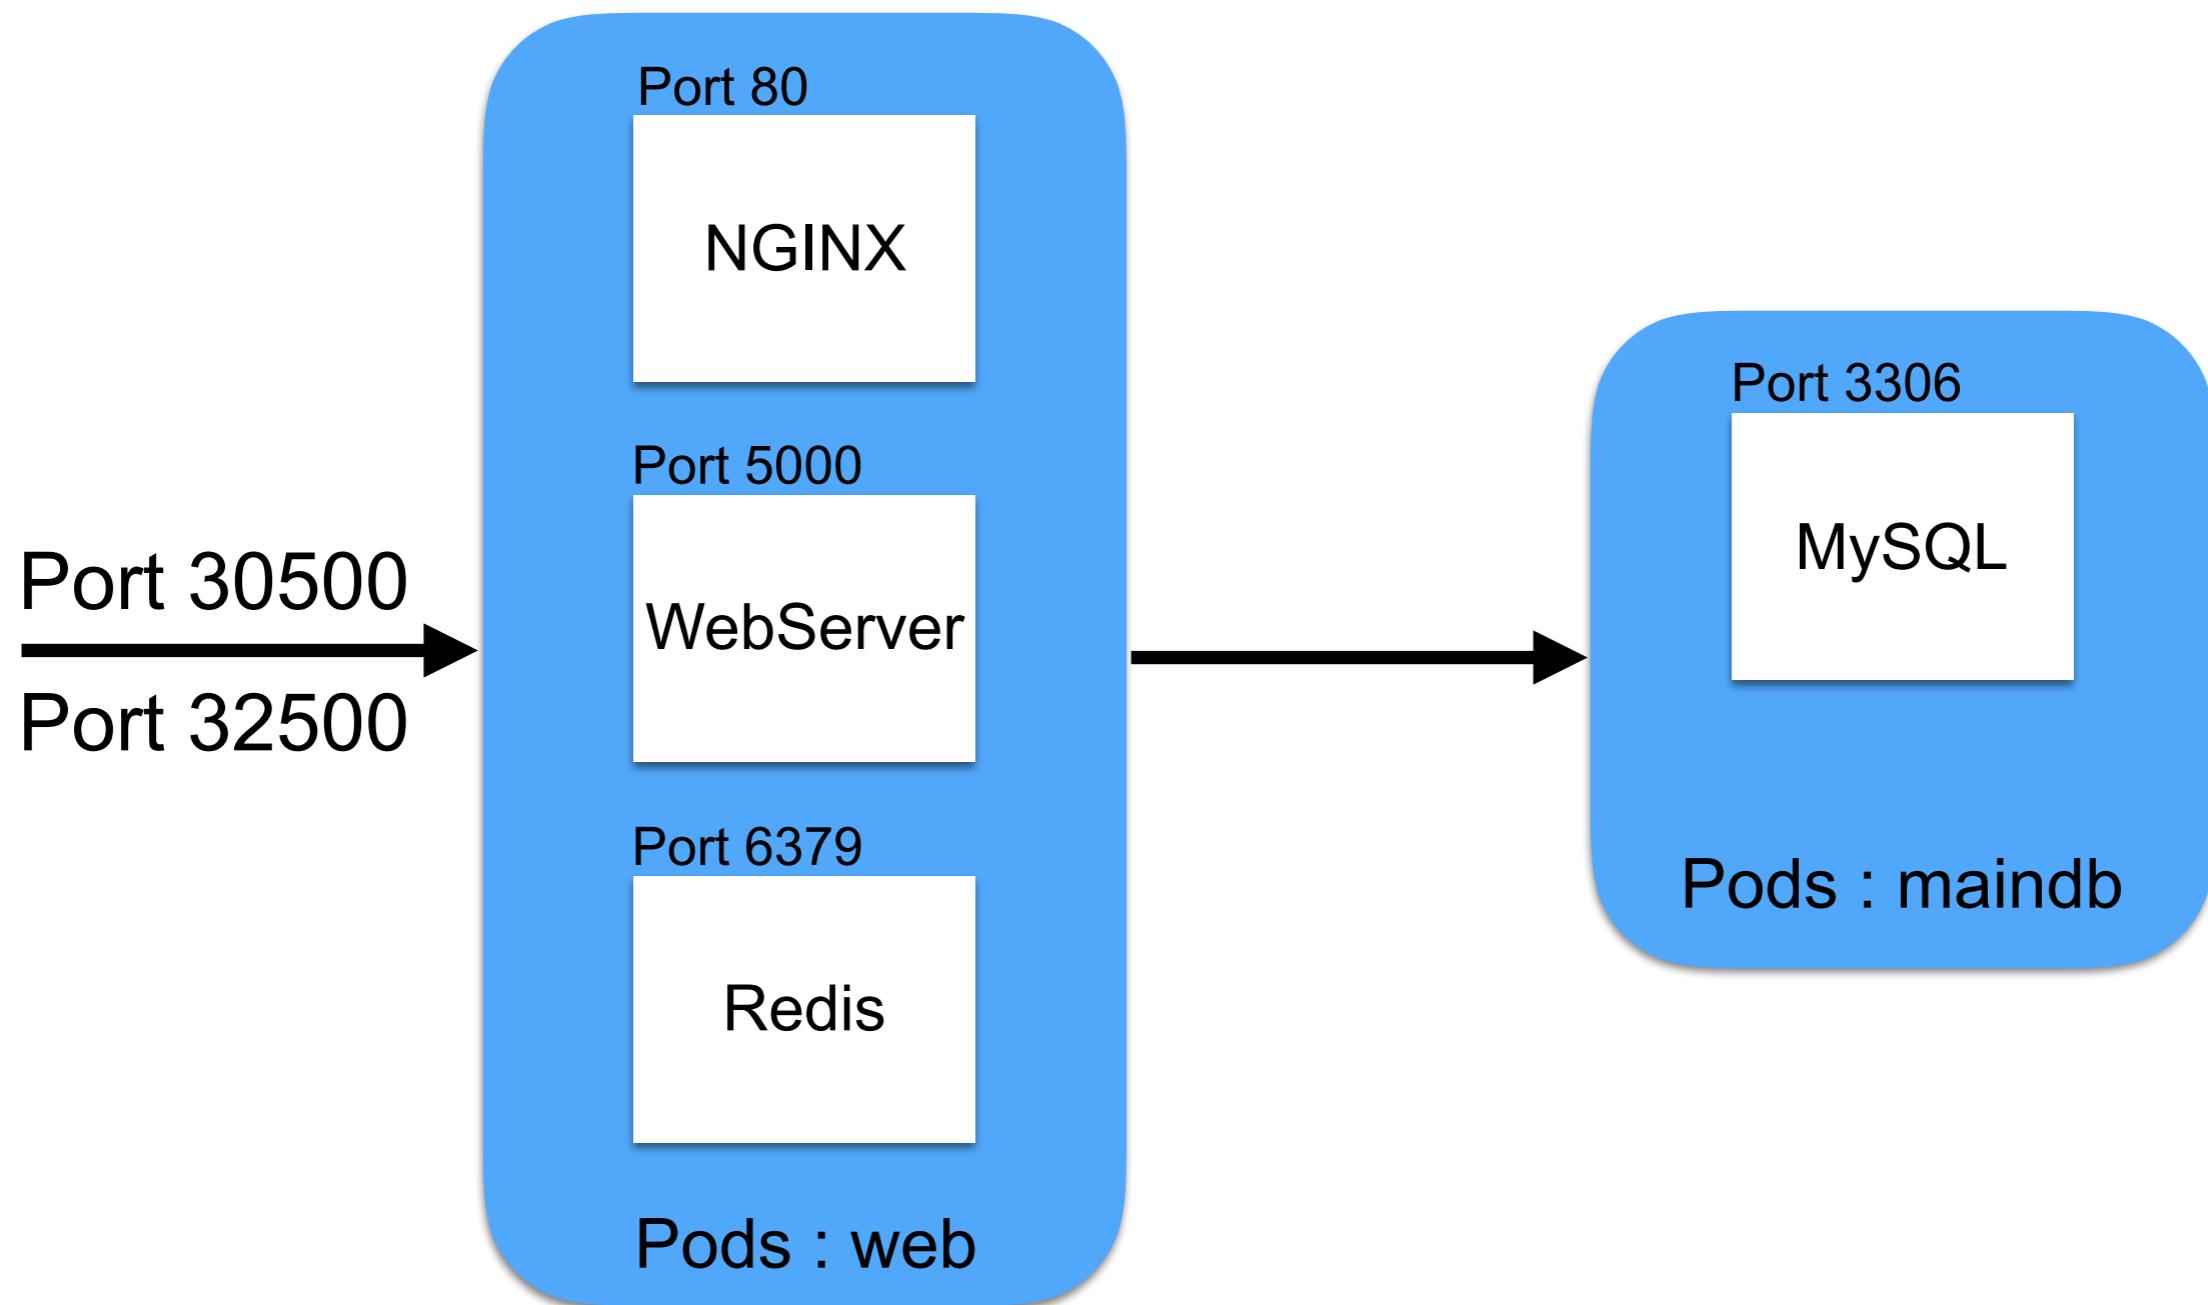
# **Workshop :: Pods & Services**

## **N container per Pods**

**file /02-pod-service/multiple-pod/  
workshop\_instruction.txt**



# Workshop



# **Workshop :: Expose port of Pods to outside cluster**

**file /02-pod-service/expose/expose-pod-service.txt**



# Expose Pods to outside cluster

HostNetwork

HostPort

NodePort of service

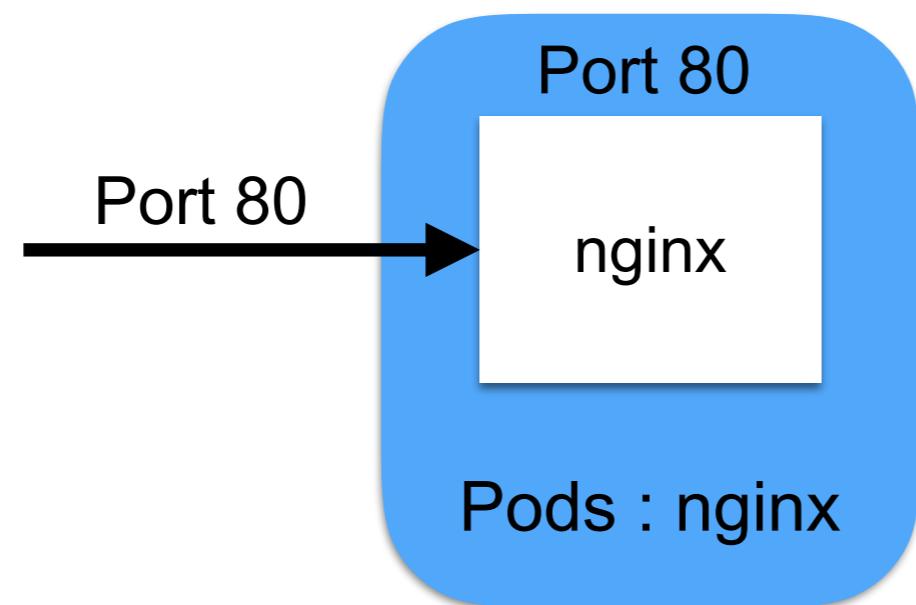
LoadBalance of service

Ingress controller



# HostNetwork

Applications running in such a pod can directly see the network interfaces of the host machine where the pod was started.



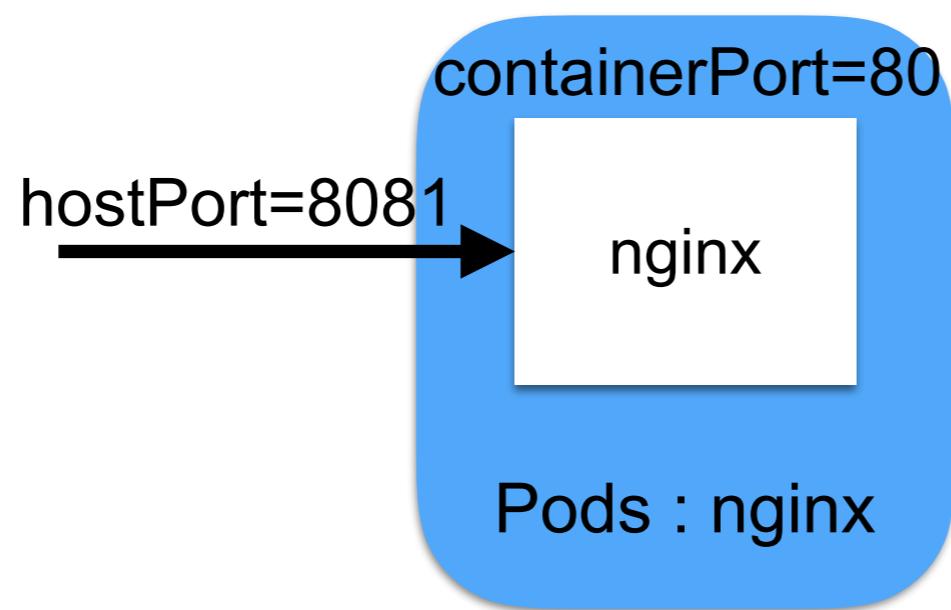
# HostNetwork

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  hostNetwork: true
  containers:
    - name: nginx
      image: nginx
```



# HostPort

The container port will be exposed to the external network at **<hostIP>:<hostPort>**, where the **hostIP** is the IP address of the Kubernetes node where the container is running and the **hostPort** is the port requested by the user.



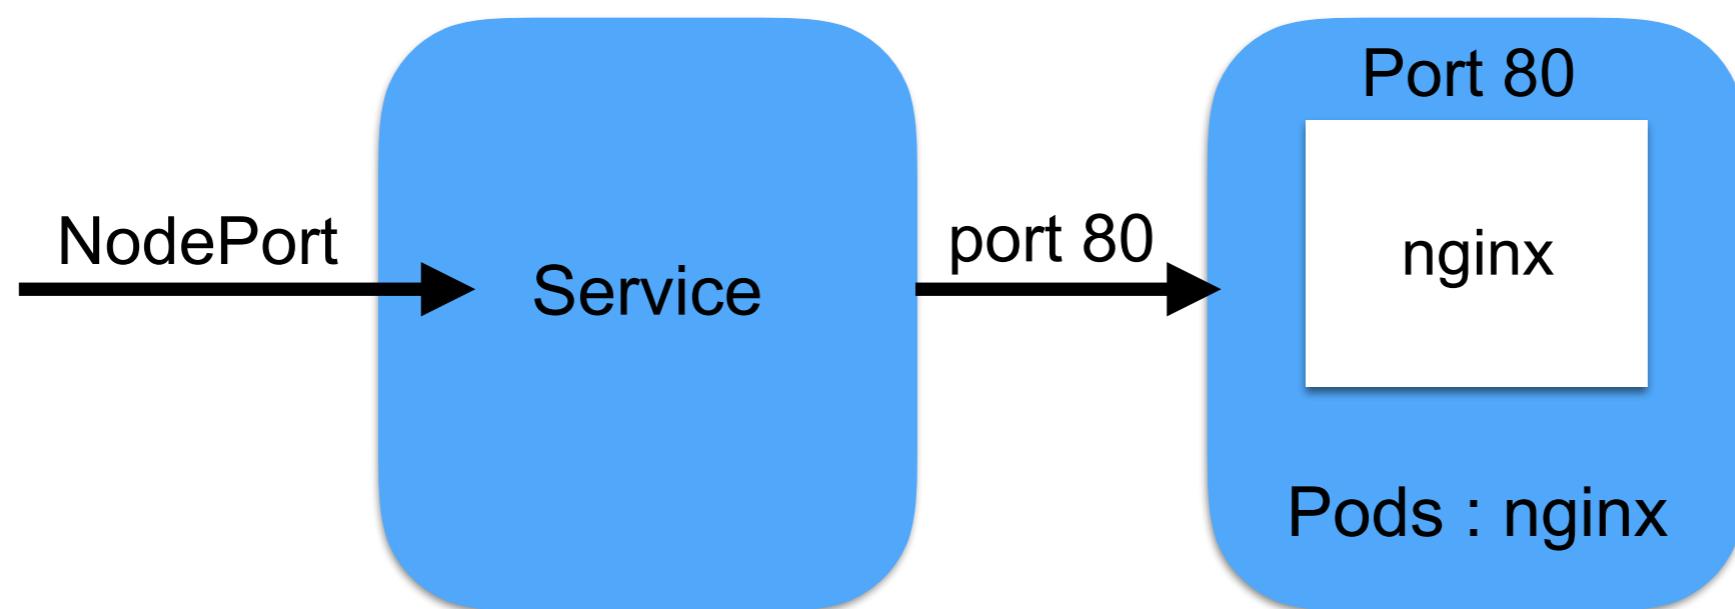
# HostPort

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - name: nginx
      image: nginx
      ports:
        - containerPort: 80
          hostPort: 8081
```



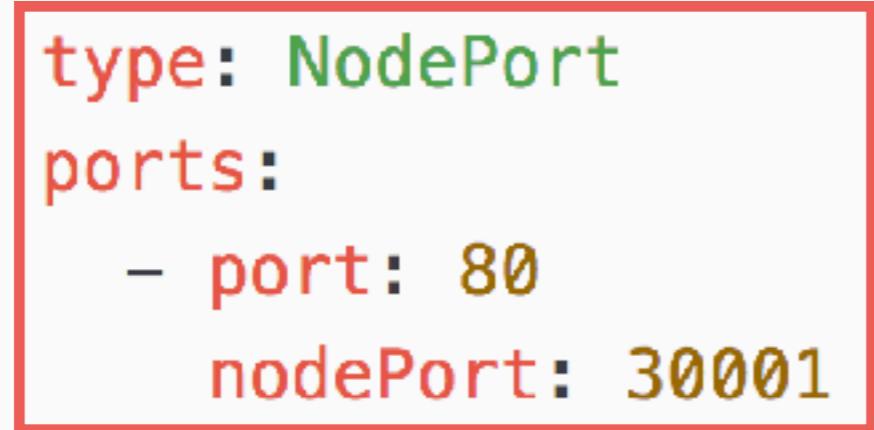
# NodePort

When creating a NodePort service, the user can specify a port from the range **30000-32767**, and each Kubernetes node will proxy that port to the pods selected by the service.



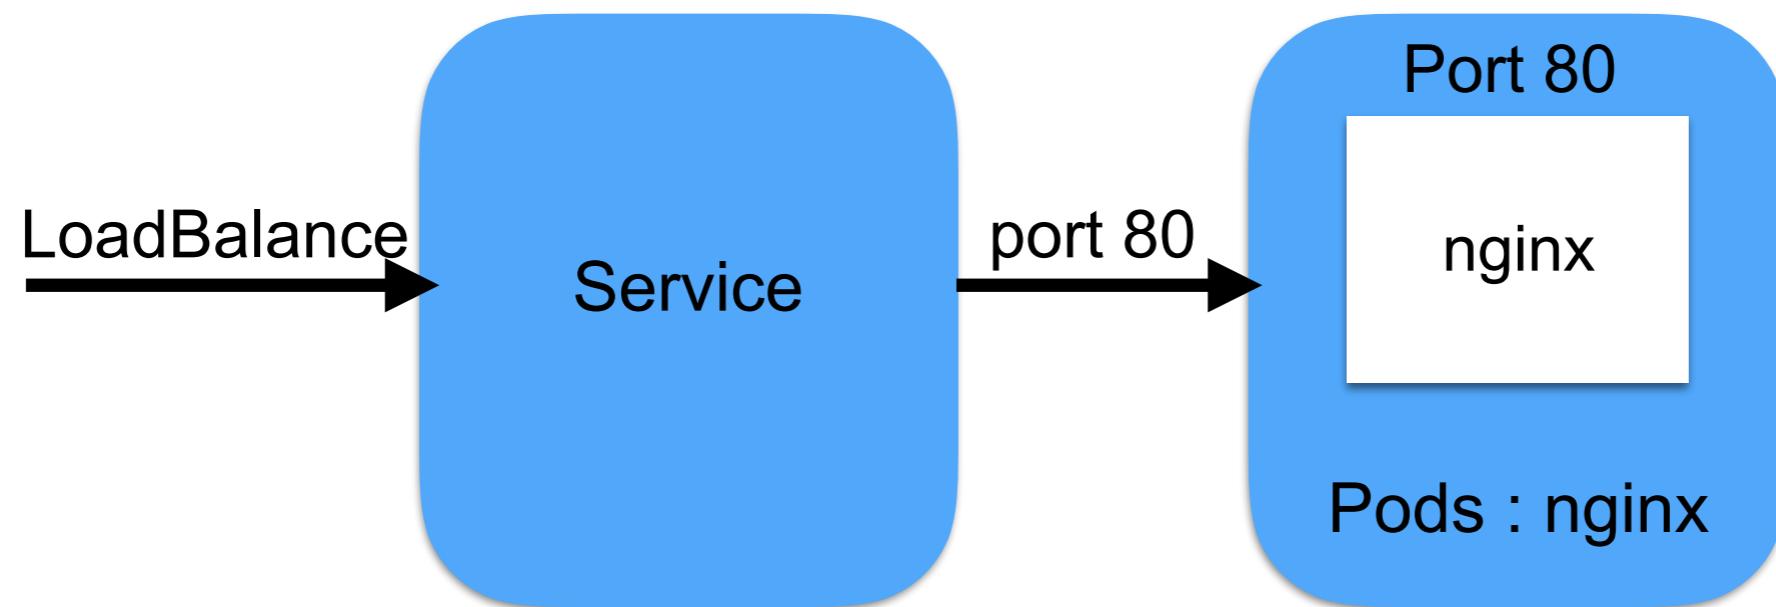
# NodePort

```
kind: Service
apiVersion: v1
metadata:
  name: nginx
spec:
  type: NodePort
  ports:
    - port: 80
      nodePort: 30001
  selector:
    name: nginx
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    name: nginx
spec:
  containers:
    - name: nginx
      image: nginx
      ports:
        - containerPort: 80
```



# LoadBalancer

a cloud provider has to be enabled in the configuration of the Kubernetes cluster. As of version 1.6, Kubernetes can provision load balancers on AWS, Azure, CloudStack, GCE and OpenStack.



# LoadBalancer

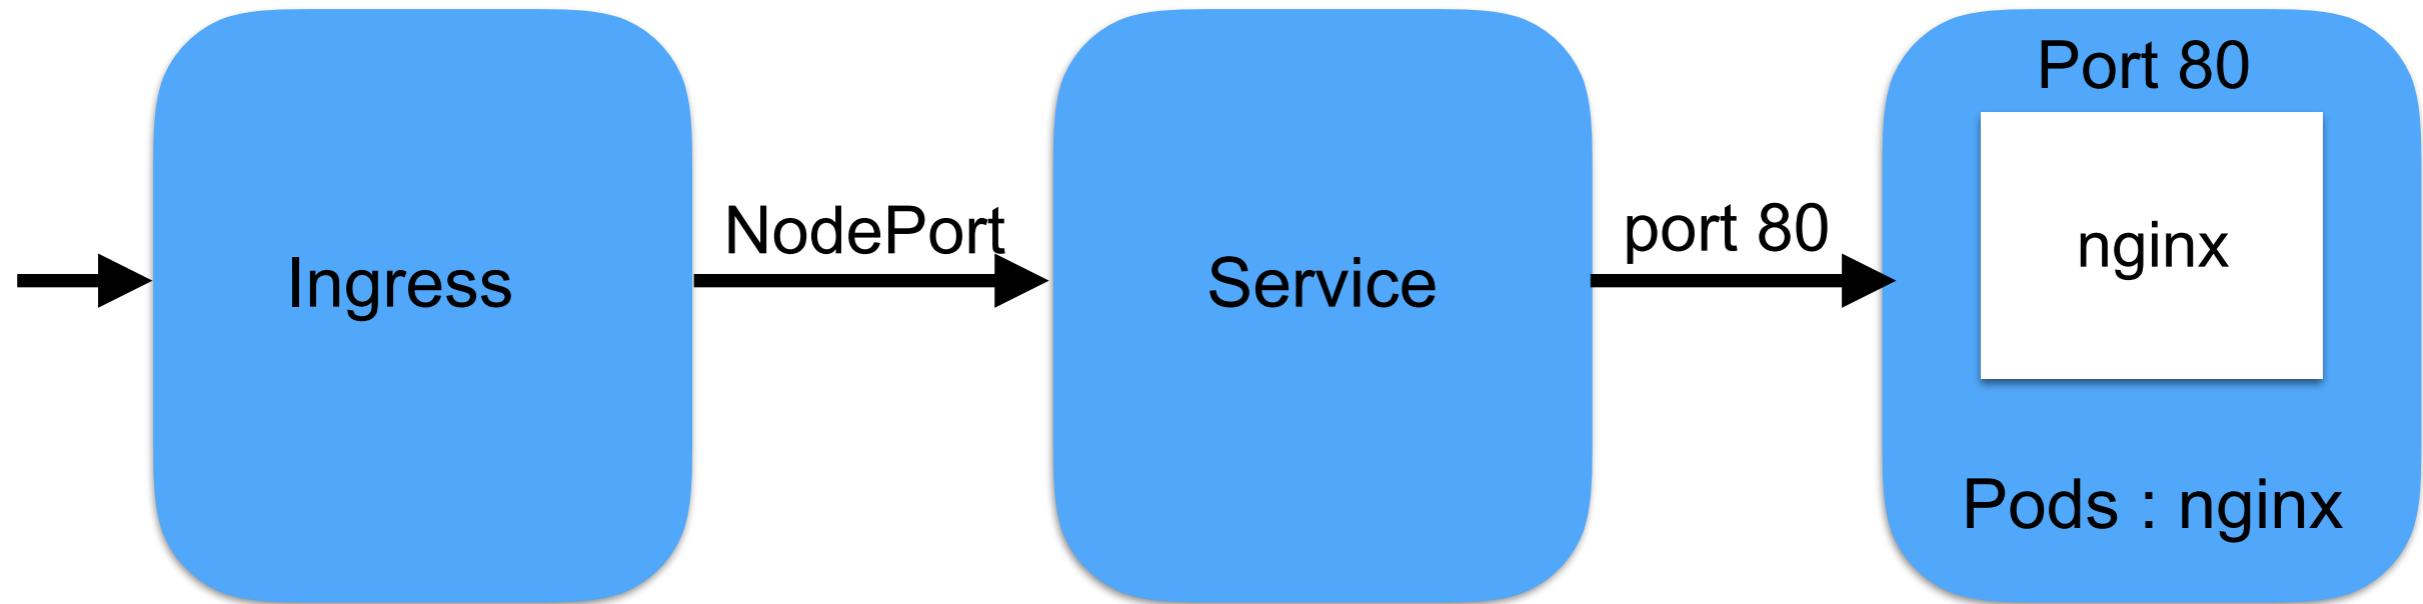
```
kind: Service
apiVersion: v1
metadata:
  name: nginx
spec:
  type: LoadBalancer
  ports:
    - port: 80
selector:
  name: nginx
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    name: nginx
spec:
  containers:
    - name: nginx
      image: nginx
      ports:
        - containerPort: 80
```



# Ingress controller

The Ingress controller is deployed as a Docker container on top of Kubernetes. Its Docker image contains a load balancer like nginx or HAProxy and a controller daemon.



<https://kubernetes.io/docs/concepts/services-networking/ingress/>



# Ingress controller

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: nginx
spec:
  rules:
    - host: nginx.kube.example.com
      http:
        paths:
          - web:
              serviceName: nginx
              servicePort: 80
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    name: nginx
spec:
  containers:
    - name: nginx
      image: nginx
      ports:
        - containerPort: 80
```



# Replication Controller (RC)



# Replication Controller

Create and maintain Pods

Keep copy of Pods by design

Maintain on cluster level

Auto-healing if Pods crash with any reason

Ensure Pods is up and run with desired number



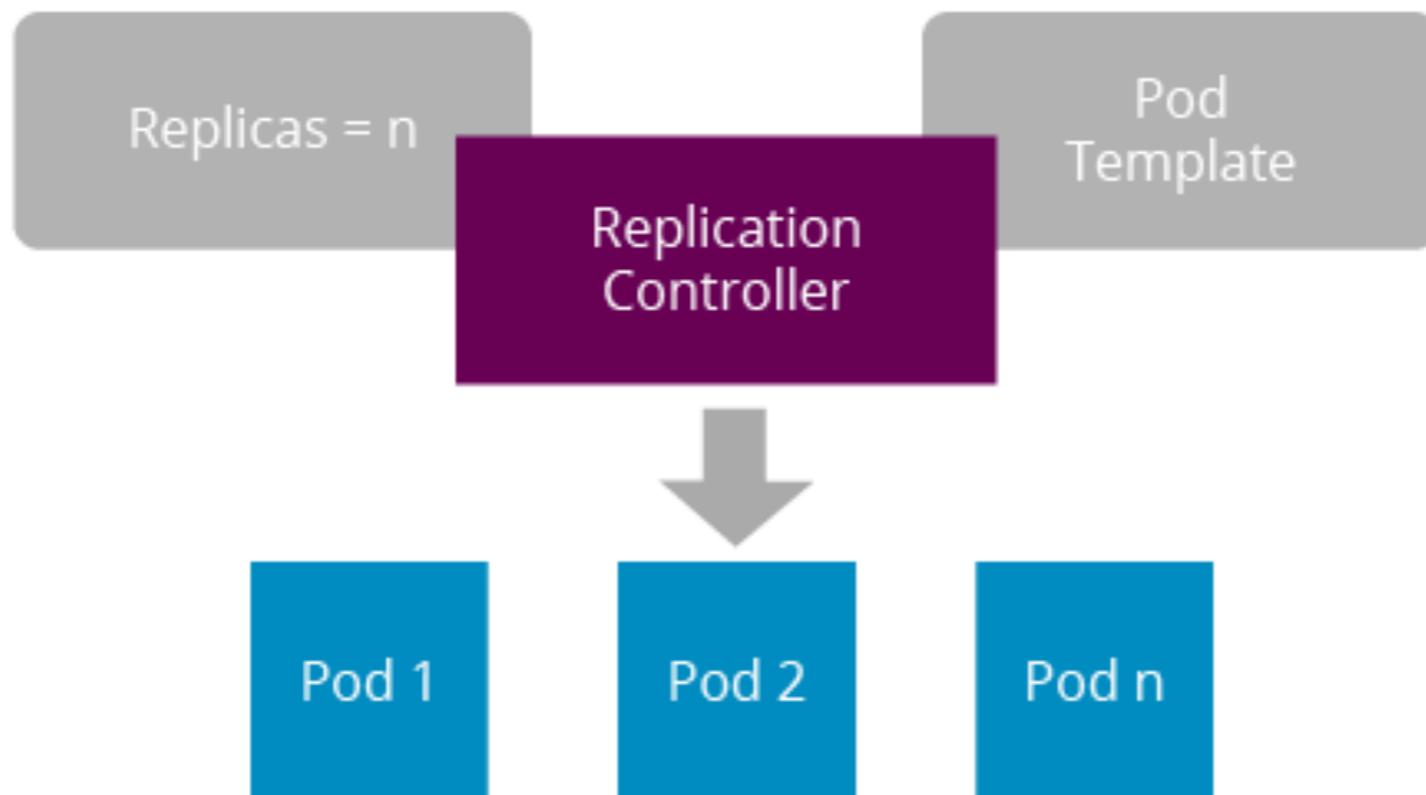
# Replication Controller

Running based on label type  
**Equality-based requirement**

```
selector:  
  name: web  
  version: "1.0"  
  module: WebServer  
  environment: development
```



# Replication Controller



# Scale up replicas of RC

```
$kubectl scale <option>  
--replicas=<number>  
<type or name>
```

```
bash-3.2$ kubectl get rc  
NAME      DESIRED   CURRENT   READY      AGE  
hello     5          5          5          15m  
bash-3.2$ kubectl get pods  
NAME        READY   STATUS    RESTARTS   AGE  
hello-jptv7  1/1     Running   0          43s  
hello-kdbsc  1/1     Running   0          43s  
hello-rn545  1/1     Running   0          15m  
hello-w6t5q  1/1     Running   0          43s  
hello-x6x8k  1/1     Running   0          43s
```



# Detail of RC

## \$kubectl describe rc <name>

### Containers:

```
hello:  
  Image:      somkiat/hello:latest  
  Port:       8080/TCP  
  Host Port:  0/TCP  
  Environment: <none>  
  Mounts:     <none>  
  Volumes:    <none>
```

### Events:

Type	Reason	Age	From	Message
Normal	SuccessfulCreate	16m	replication-controller	Created pod: hello-rn545
Normal	SuccessfulCreate	16m	replication-controller	Created pod: hello-pfnwj
Normal	SuccessfulCreate	16m	replication-controller	Created pod: hello-jfqfl
Normal	SuccessfulCreate	13m	replication-controller	Created pod: hello-6pfqm
Normal	SuccessfulCreate	11m	replication-controller	Created pod: hello-l4lqc
Normal	SuccessfulCreate	11m	replication-controller	Created pod: hello-pfgr2
Normal	SuccessfulDelete	9m	replication-controller	Deleted pod: hello-l4lqc
Normal	SuccessfulDelete	9m	replication-controller	Deleted pod: hello-pfnwj
Normal	SuccessfulDelete	9m	replication-controller	Deleted pod: hello-6pfqm
Normal	SuccessfulDelete	9m	replication-controller	Deleted pod: hello-pfgr2
Normal	SuccessfulCreate	2m	replication-controller	Created pod: hello-kdbsc
Normal	SuccessfulCreate	2m	replication-controller	Created pod: hello-jptv7
Normal	SuccessfulCreate	2m	replication-controller	Created pod: hello-x6x8k

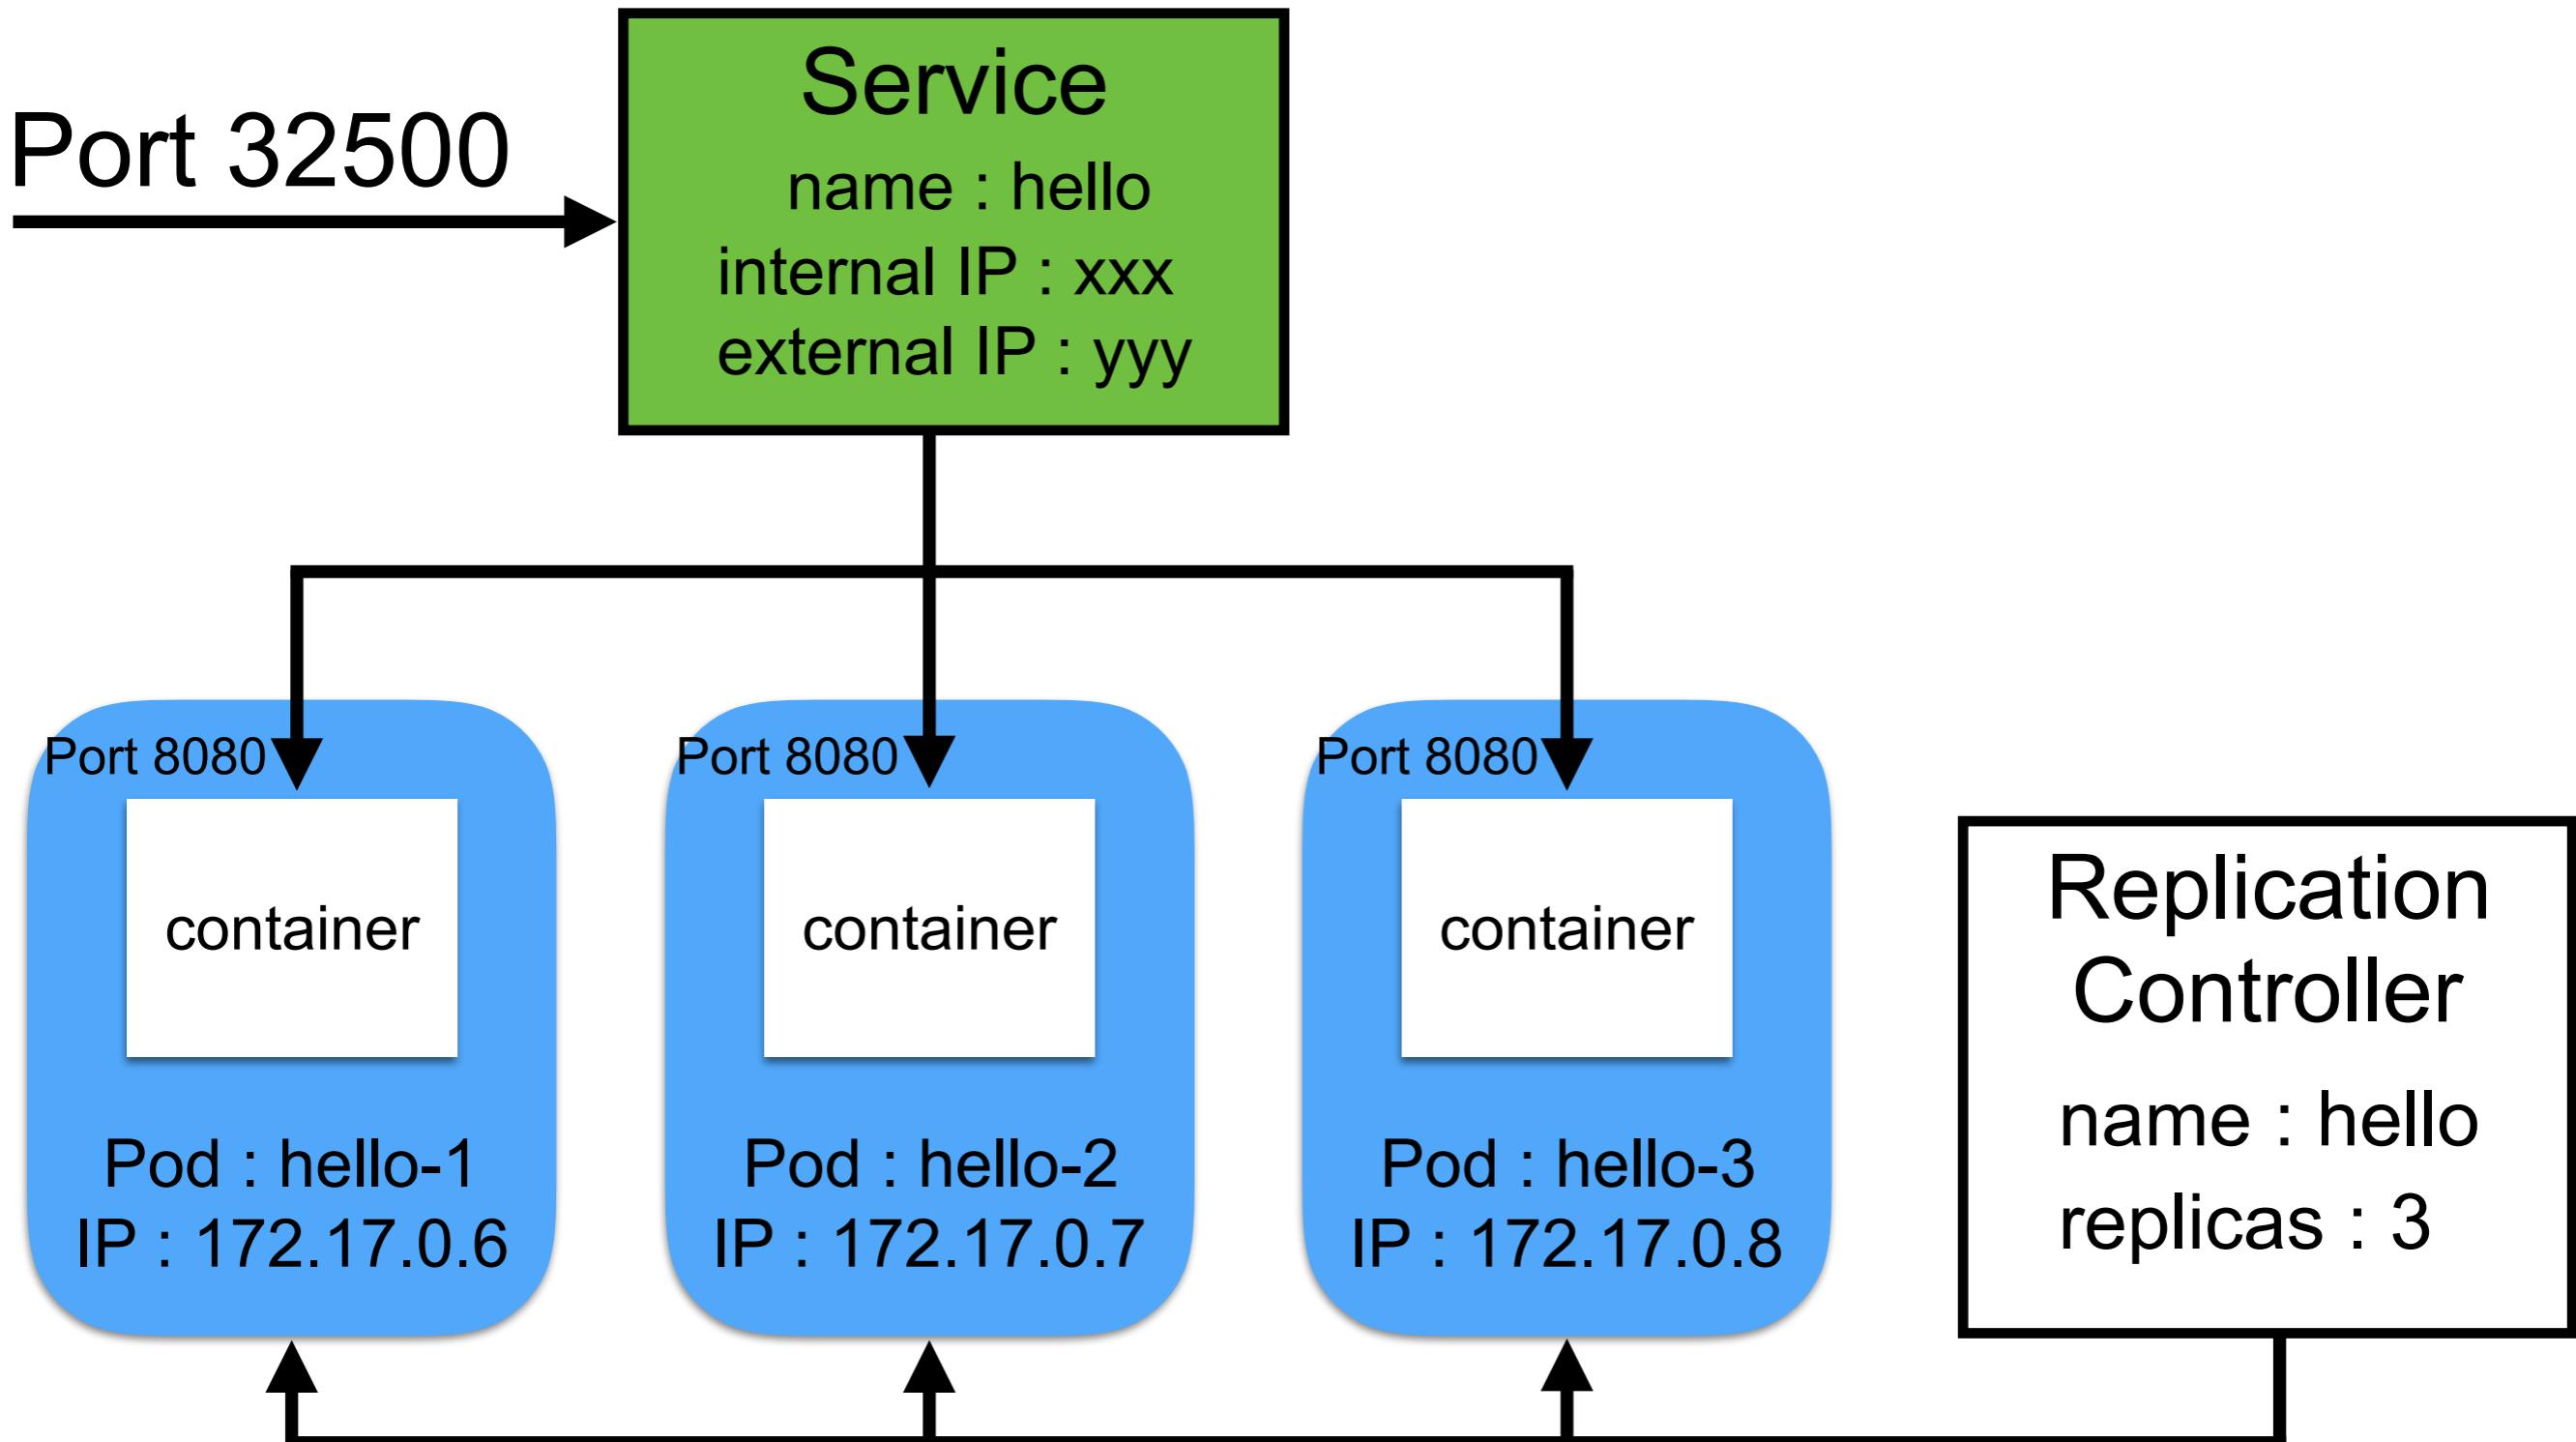


# Workshop Replication Controller

file /03-replication-controller



# Scaling the application

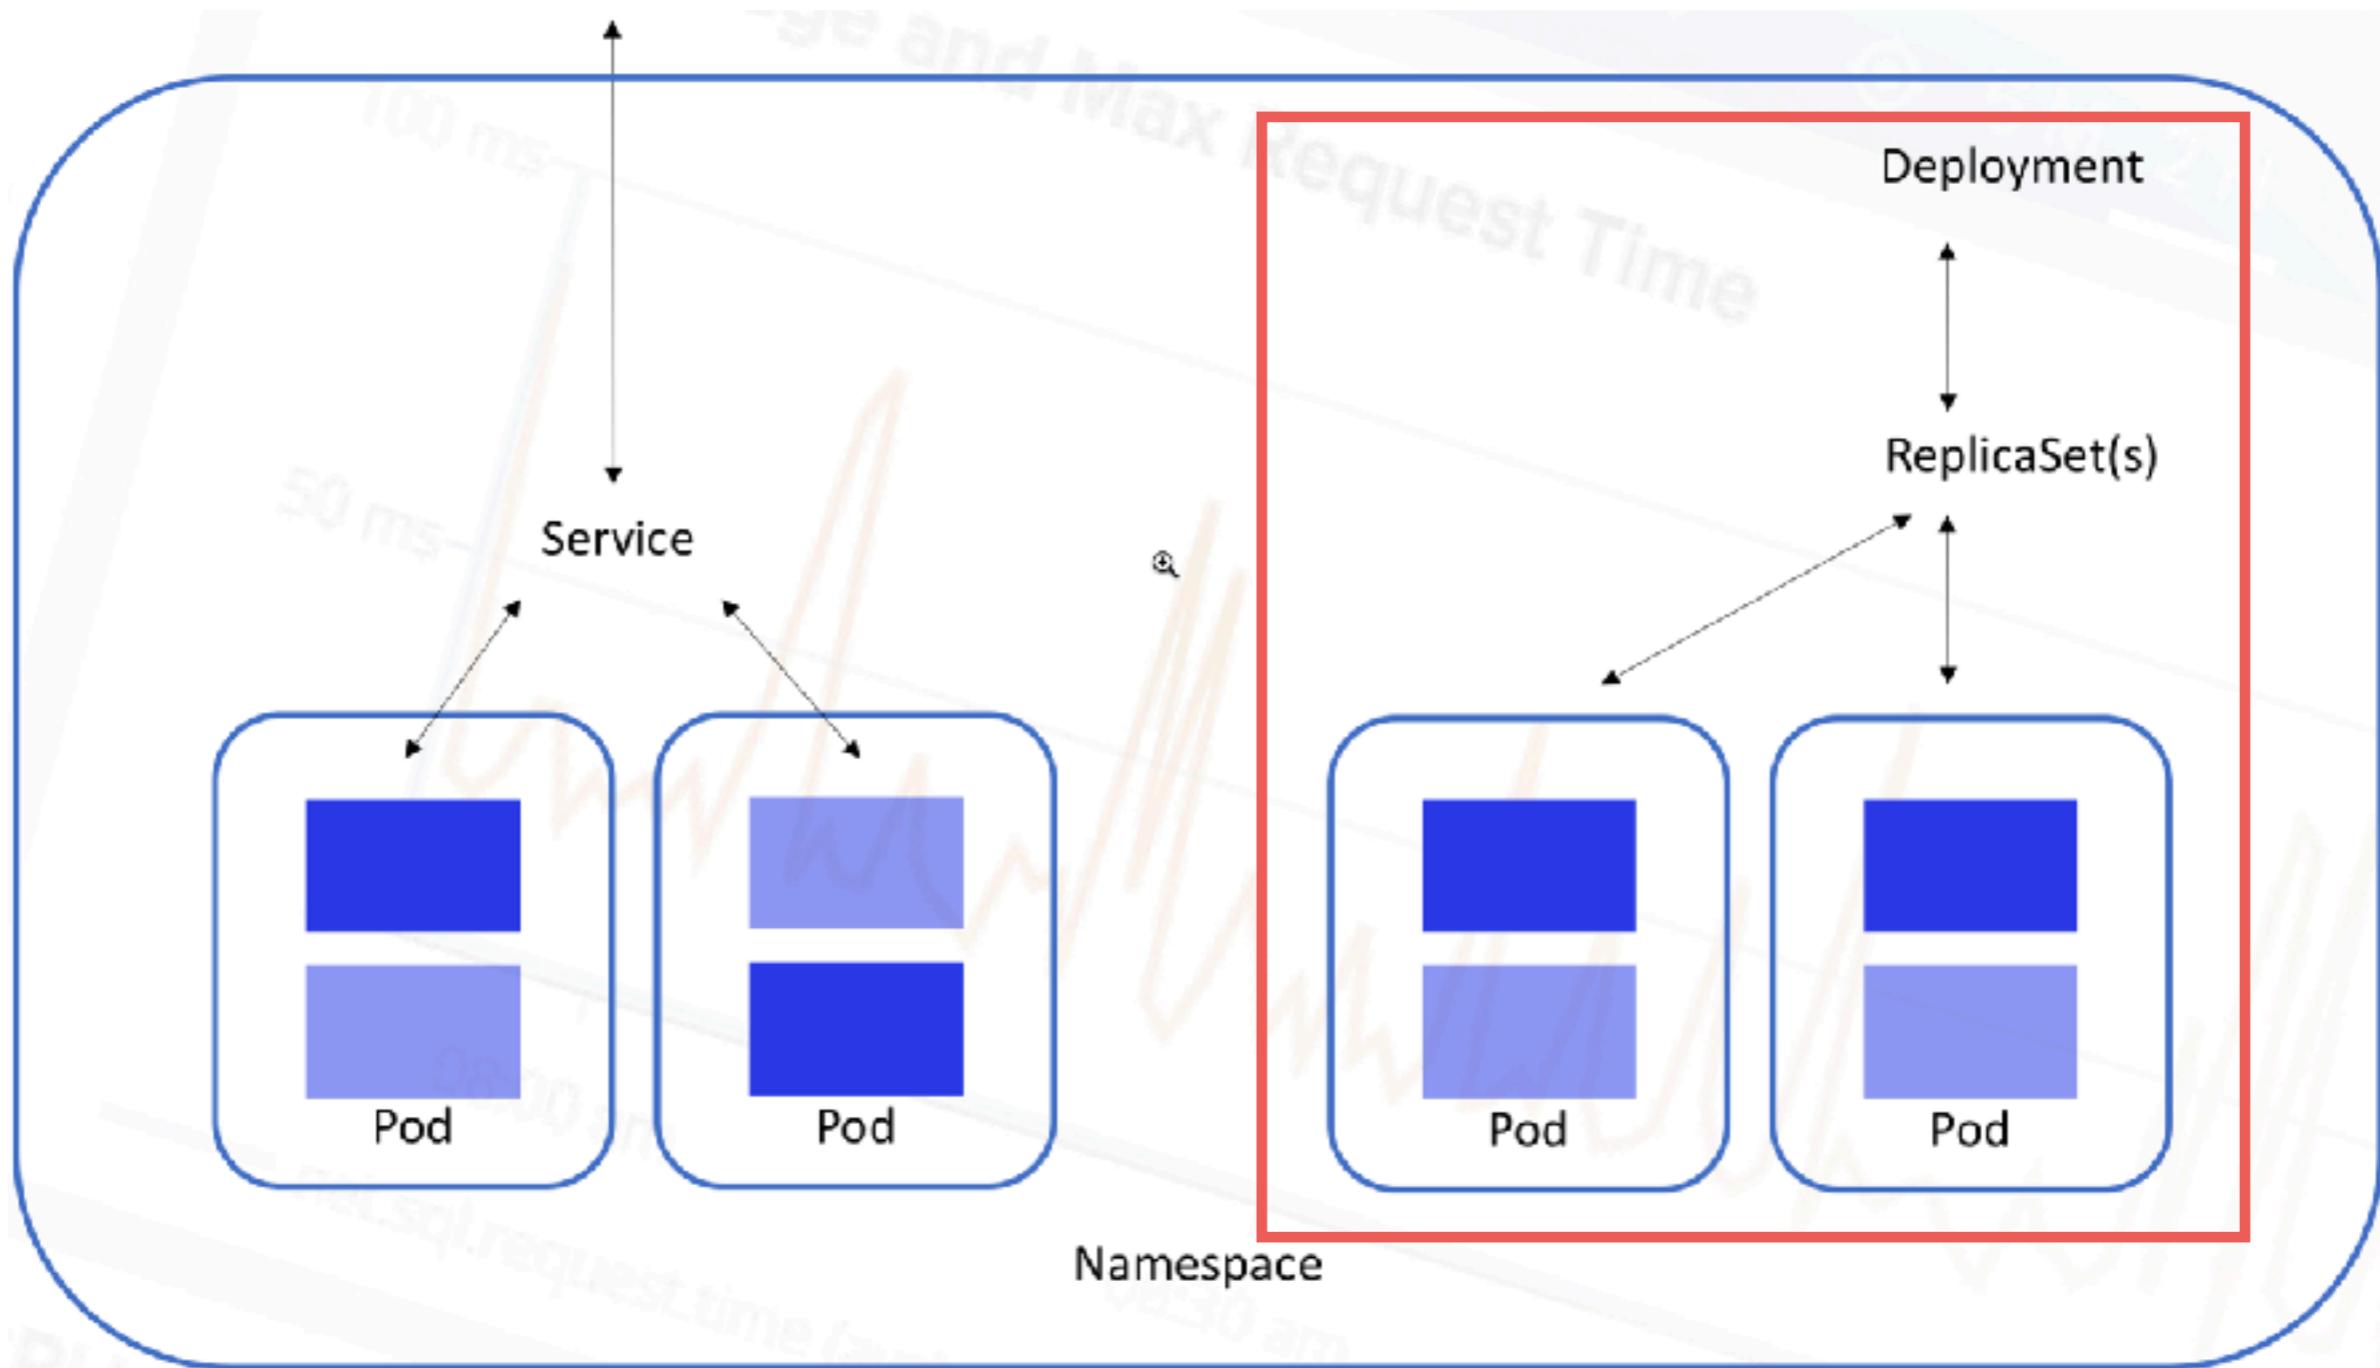


# Deployment and ReplicaSet (RS)

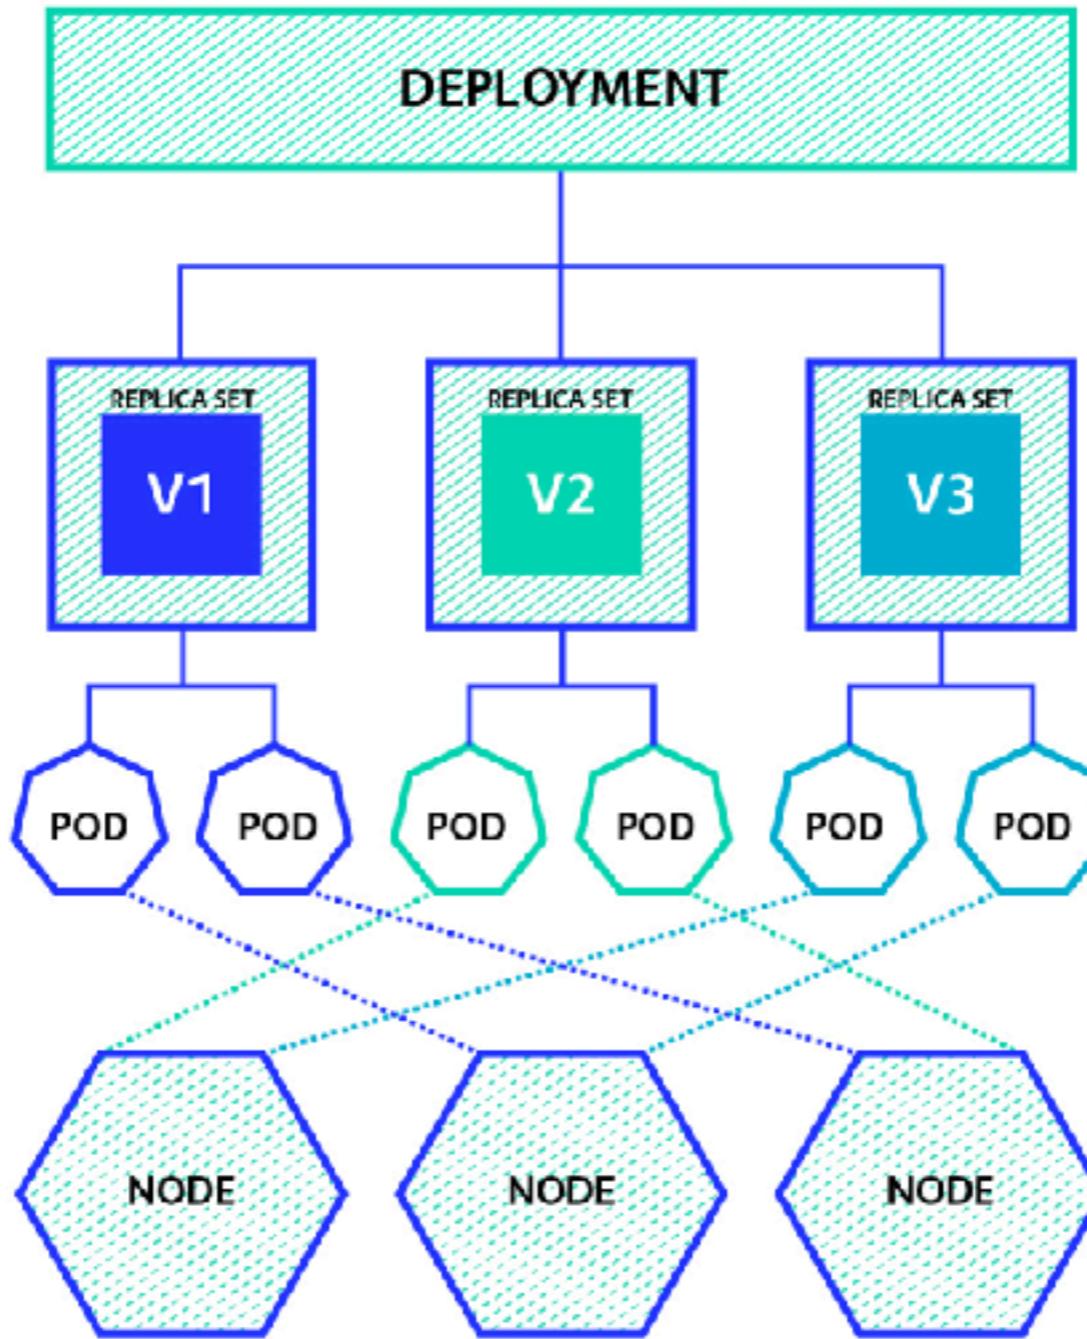
file /04-deployment-and-replica-set



# Deployment and ReplicaSet



# Deployment and ReplicaSet



<https://thenewstack.io/kubernetes-deployments-work/>



# Deployment and ReplicaSet

Next-generation of RC

Provide function to maintain versioning of Pods

Update new version (Rollout)

Revert to old version (Rollback)

Scale a deployment

Pause/Resume process



# New version from Deployment

Create new RS

Start to scale as desired

Scale down existing RS to 0

Delete existing RS



# Deployment/RS vs RC

RS more dynamic than RC

RS support label with methods:

- Equality-based requirement
- Set-based requirement

<https://kubernetes.io/docs/concepts/overview/working-with-objects/labels/>



# Label selector

## Equality-based requirement

```
spec:  
  replicas: 3  
  selector:  
    app: nginx  
    environment: production  
    tier: frontend
```

## Set-based requirement

```
spec:  
  replicas: 3  
  selector:  
    matchLabels:  
      tier: frontend  
    matchExpressions:  
      - {key: tier, operator: In, values: [frontend]}
```



Set-based requirement can be  
mixed with equality-based  
requirement



# Rollout strategy

**Set online**

**Edit online**

**Modify YAML file and apply**



# Set online

```
$kubectl set image deployment/hello  
hello=somkiat/hello:v2
```

```
Waiting for rollout to finish: 1 out of 3 new replicas have been updated...  
Waiting for rollout to finish: 1 out of 3 new replicas have been updated...  
Waiting for rollout to finish: 1 out of 3 new replicas have been updated...  
Waiting for rollout to finish: 1 out of 3 new replicas have been updated...  
Waiting for rollout to finish: 1 out of 3 new replicas have been updated...  
Waiting for rollout to finish: 1 out of 3 new replicas have been updated...  
Waiting for rollout to finish: 2 out of 3 new replicas have been updated...  
Waiting for rollout to finish: 2 out of 3 new replicas have been updated...  
Waiting for rollout to finish: 2 old replicas are pending termination...  
Waiting for rollout to finish: 1 old replicas are pending termination...  
Waiting for rollout to finish: 1 old replicas are pending termination...  
deployment "hello" successfully rolled out
```



# Edit online

## \$kubectl edit deployment hello

```
# Please edit the object below. Lines beginning with a '#' will be ignored,
# and an empty file will abort the edit. If an error occurs while saving this file it
# will be reopened with the relevant failures.
#
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  annotations:
    deployment.kubernetes.io/revision: "6"
    kubectl.kubernetes.io/last-applied-configuration: |
        {"apiVersion":"apps/v1","kind":"Deployment","metadata":{},"annotations":{},"
        "namespace":"default","spec":{"replicas":3,"revisionHistoryLimit":1,"selector":{},"template":{"metadata":{"labels":{"app":"web"}},"
        "spec":{"containers":[{"image":"nginx:1.13","ports":[{"containerPort":8080,"protocol":"TCP"}]}]}}}
  creationTimestamp: 2018-04-01T17:17:51Z
  generation: 10
  labels:
    app: web
  name: hello
  namespace: default
  resourceVersion: "33405"
  selfLink: /apis/extensions/v1beta1/namespaces/default/deployments/hello
  uid: 9b773859-35d0-11e8-9d36-0800275c6c60
```



# Modify YAML file and apply

```
$kubectl apply -f hello_deployment.yml
```



# Show history of rollout

```
$kubectl rollout history deployment/hello
```

```
deployments "hello"
REVISION  CHANGE-CAUSE
2          <none>
3          <none>
4          <none>
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: hello
  labels:
    app: web
spec:
  replicas: 3
  revisionHistoryLimit: 1
  selector:
    matchLabels:
      app: web
```

spec.revisionHistory = 2 (default)  
spec.revisionHistory = 0 (clean all)



# Try to rollback to revision

```
$kubectl rollout undo deployment/hello  
--to-revision=2
```



# Try to scale a deployment

\$kubectl scale deployment hello --replicas=5

```
bash-3.2$ kubectl scale deployment hello --replicas=3
```

```
deployment.extensions "hello" scaled
```

```
bash-3.2$ kubectl get deployment -o wide
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE	CONTAINERS	IMAGES	SELECTOR
hello	3	3	3	3	21m	hello	somkiat/hello	app=web

```
bash-3.2$ kubectl scale deployment hello --replicas=5
```

```
deployment.extensions "hello" scaled
```

```
bash-3.2$ kubectl get deployment -o wide
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE	CONTAINERS	IMAGES	SELECTOR
hello	5	5	5	5	21m	hello	somkiat/hello	app=web

```
bash-3.2$ kubectl scale deployment hello --replicas=1
```

```
deployment.extensions "hello" scaled
```

```
bash-3.2$ kubectl get deployment -o wide
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE	CONTAINERS	IMAGES	SELECTOR
hello	1	1	1	1	21m	hello	somkiat/hello	app=web



# Pause and resume rollout

```
$kubectl rollout pause deployment/hello
```

```
$kubectl rollout resume deployment/hello
```



# **Workshop**

# **Deployment/ReplicaSet**

**file /04-deployment-and-replica-set**



# **Horizontal Pods Autoscale (HPA)**

File /working-with-java/03-hpa



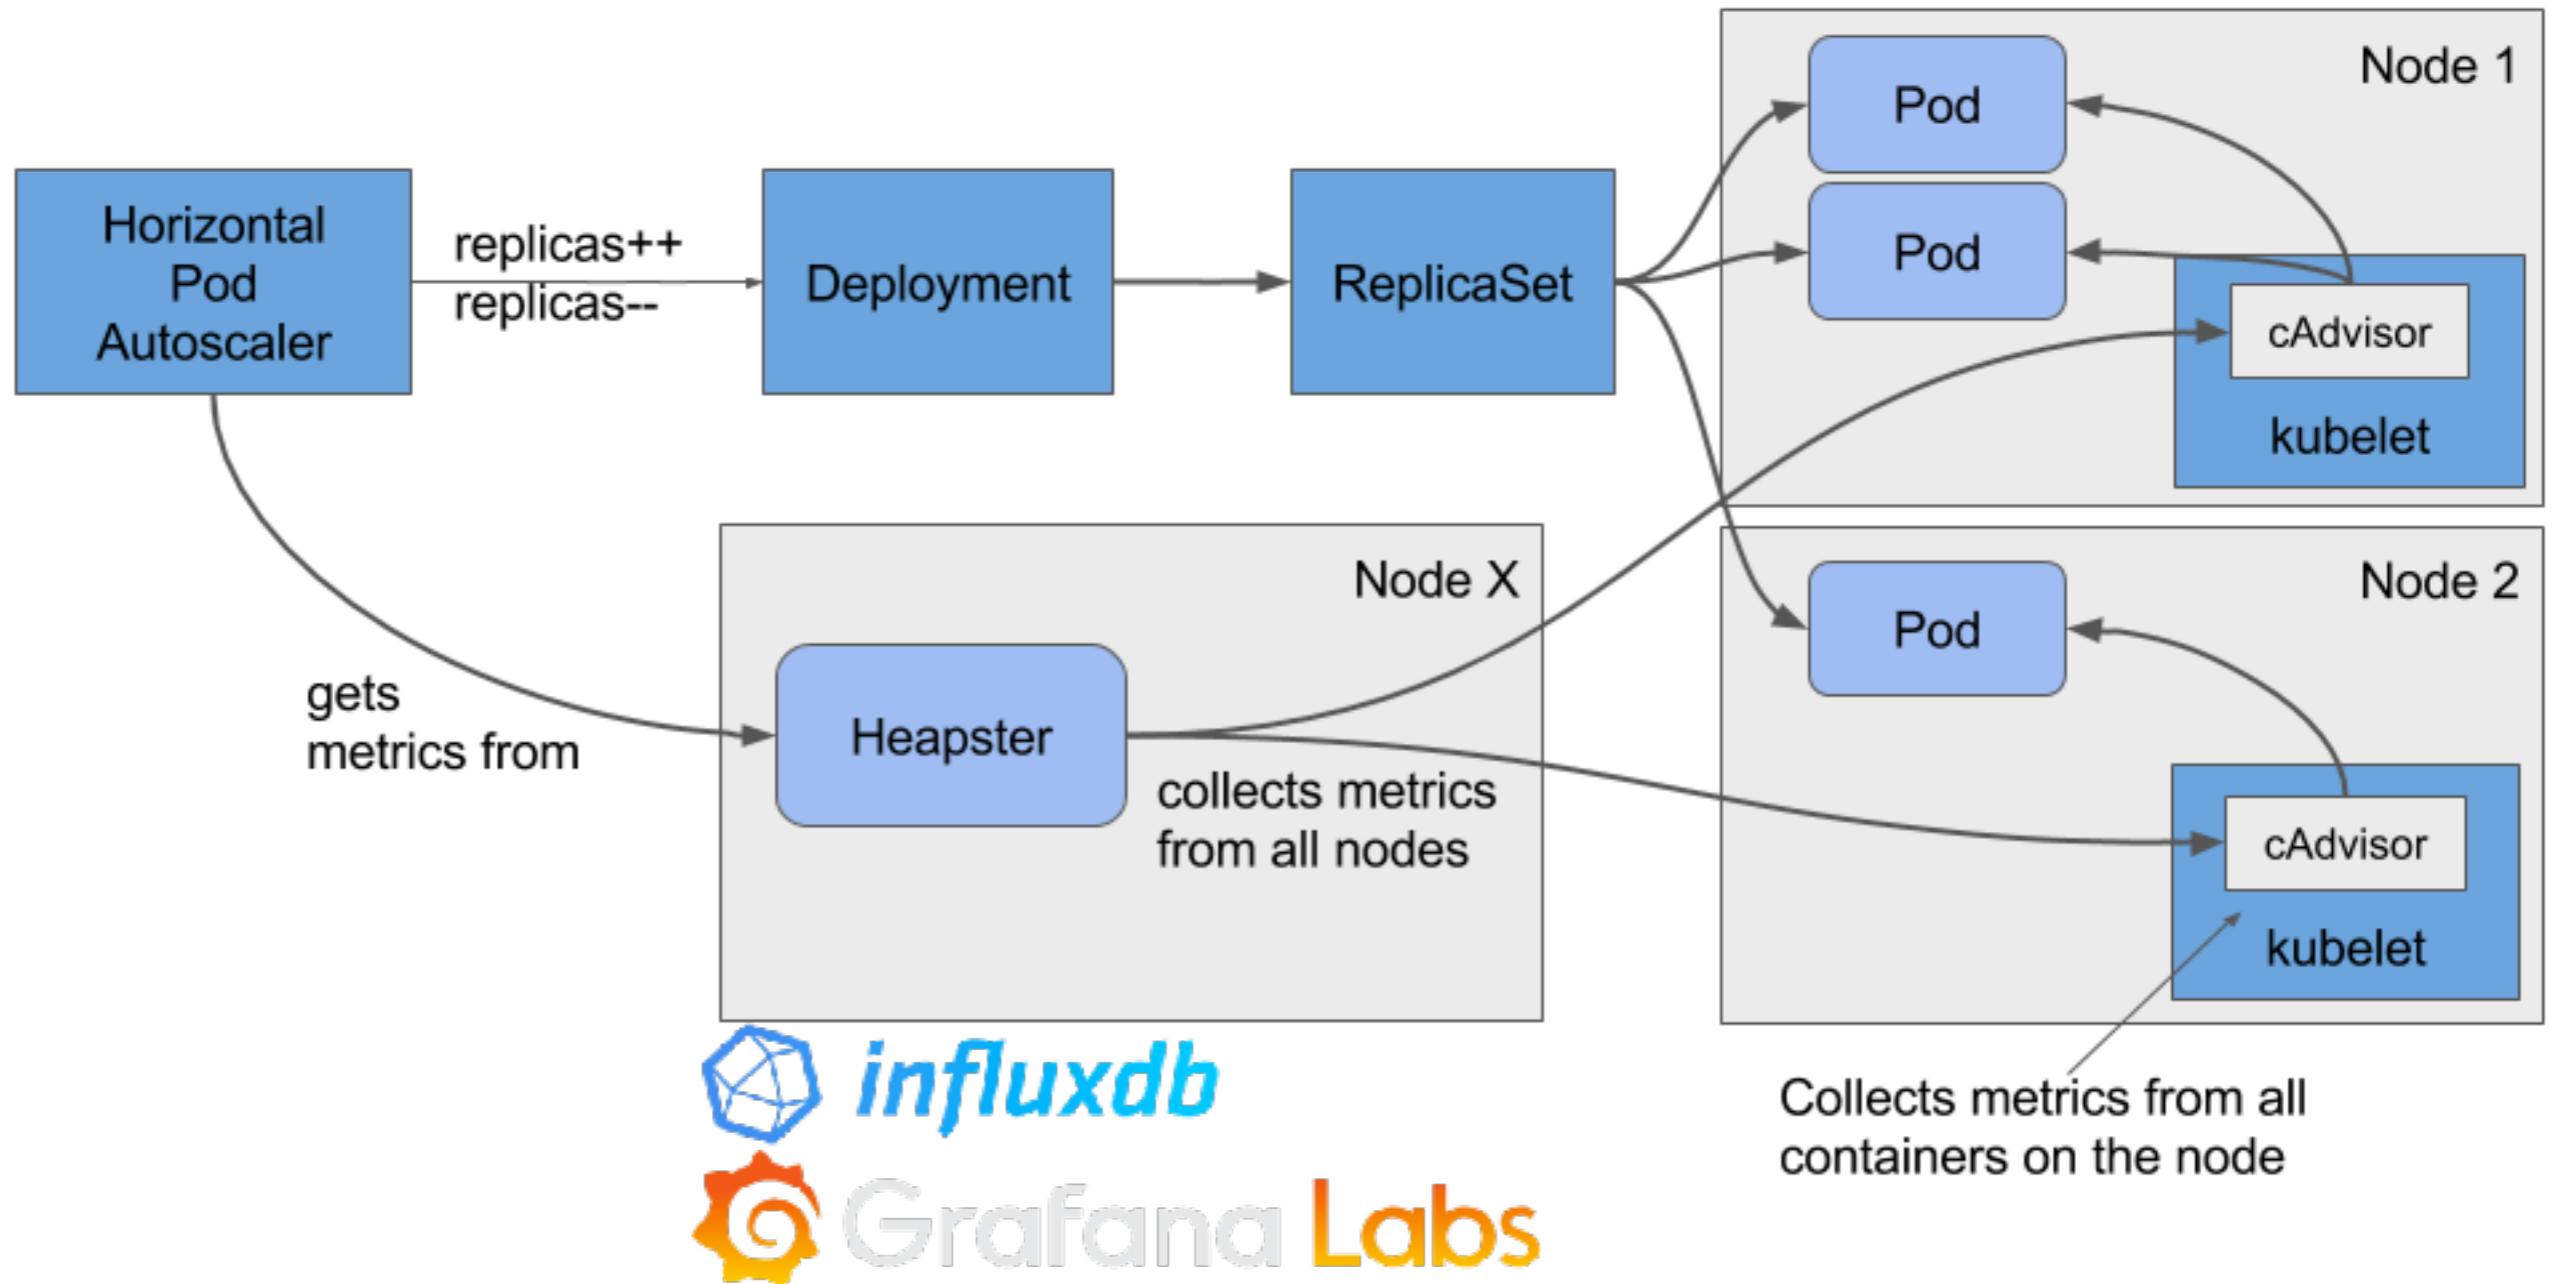
# Horizontal Pod Autoscaler

Monitor workload on Pods (based on CPU)  
and automatic scaling-up application

Scaling-up and down for your app need !!



# Components in Autoscaling



<https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale-walkthrough/>



# 1. Enable hipster in minikube

\$minikube addons enable heapster

\$minikube addons start heapster

\$minikube addons open heapster



# See all pods from heapster

\$kubectl get pods --all-namespaces

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
default	mongo-controller-xdh28	1/1	Running	0	9h
default	spring-boot-service-deployment-74f5df876f-gxn24	1/1	Running	0	4h
kube-system	default-http-backend-kf92s	1/1	Running	0	4h
kube-system	heapster-2ssh2	1/1	Running	0	8h
kube-system	influxdb-grafana-vk6j6	2/2	Running	0	8h
kube-system	kube-addon-manager-minikube	1/1	Running	2	1d
kube-system	kube-dns-54cccfbd8-xkpl6	3/3	Running	6	1d
kube-system	kubernetes-dashboard-77d8b98585-n49sg	1/1	Running	4	1d
kube-system	metrics-server-bb9ffc6b8-5mgnv	1/1	Running	0	8h
kube-system	nginx-ingress-controller-nrk4d	1/1	Running	0	4h
kube-system	storage-provisioner	1/1	Running	2	1d



# See all services from heapster

\$kubectl get service --all-namespaces

NAMESPACE	NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
default	kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	1d
default	mongo-service	ClusterIP	10.99.242.178	<none>	27017/TCP	9h
default	spring-boot-service	NodePort	10.98.105.78	<none>	80:32132/TCP	4h
kube-system	default-http-backend	NodePort	10.102.221.170	<none>	80:30001/TCP	1h
kube-system	heapster	ClusterIP	10.100.221.108	<none>	80/TCP	8h
kube-system	Kube-dns	ClusterIP	10.96.0.10	<none>	53/UDP, 53/TCP	1d
kube-system	kubernetes-dashboard	NodePort	10.99.243.227	<none>	80:30000/TCP	1d
kube-system	metrics-server	ClusterIP	10.107.38.180	<none>	443/TCP	8h
kube-system	monitoring-grafana	NodePort	10.104.125.177	<none>	80:30002/TCP	8h
kube-system	monitoring-influxdb	ClusterIP	10.99.175.232	<none>	8083/TCP, 8086/TCP	8h



# See result in Grafana

<http://192.168.99.100:30002>

The screenshot shows the Grafana Home Dashboard at the URL <http://192.168.99.100:30002/?orgId=1>. The dashboard features a top navigation bar with icons for search, refresh, and settings. Below the navigation is a progress bar titled "Getting Started with Grafana" consisting of five steps: "Install Grafana" (checkmark), "Create your first data source" (checkmark), "Create your first dashboard" (checkmark), "Add Users" (highlighted in green), and "Install apps & plugins" (sun icon). To the left, sections for "Starred dashboards" (Pods) and "Recently viewed dashboards" (Cluster) are shown. On the right, sections for "Installed Apps", "Installed Panels", and "Installed Datasources" all indicate "None installed. Browse Grafana.net".



## 2. Create metric server

Starting from Kubernetes 1.8, resource usage metrics, such as container CPU and memory usage, are available in Kubernetes through the Metrics API. These metrics can be either accessed directly by user, for example by using `kubectl top` command, or used by a controller in the cluster, e.g. Horizontal Pod Autoscaler, to make decisions.

<https://github.com/kubernetes-incubator/metrics-server>



# 3. Deploy Springboot service

```
$kubectl create -f boot-deployment.yaml  
$kubectl create -f boot-service.yaml
```



# 4. Create HPA in command line

```
$kubectl autoscale  
deployment/spring-boot-service-deployment  
--min=1 --max=5 --cpu-percent=5
```

min = minimum of replica

max = maximum of replica

cpu-percent = average of % of CPU usage



# 4. Create HPA in command line

\$kubectl get hpa

REFERENCE	TARGETS	MINPODS	MAXPODS
Deployment/spring-boot-service-deployment	0%/5%	1	5



## 4. Create HPA in command line

**Scale-up can only happen if there was no rescaling  
within the last 3 minutes**

**Scale-down will wait for 5 minutes from the last  
rescaling**

<https://github.com/kubernetes/community/blob/master/contributors/design-proposals/autoscaling/horizontal-pod-autoscaler.md#autoscaling-algorithm>



# 5. Load testing with command line

`http://spring-boot-service.default.svc.cluster.local:  
8080/user`



# 6. See result (1)

\$kubectl get pods

NAME	READY	STATUS	RESTARTS	AGE
load-test-6f86656986-9qvdl	0/1	Error	0	4m
mongo-controller-xdh28	1/1	Running	0	1h
spring-boot-service-deployment-86c7b47b7b-9l9sn	1/1	Running	0	15m
spring-boot-service-deployment-86c7b47b7b-9srkg	1/1	Running	0	1m
spring-boot-service-deployment-86c7b47b7b-kbtrl	1/1	Running	0	1m
spring-boot-service-deployment-86c7b47b7b-pqbsm	1/1	Running	0	1m



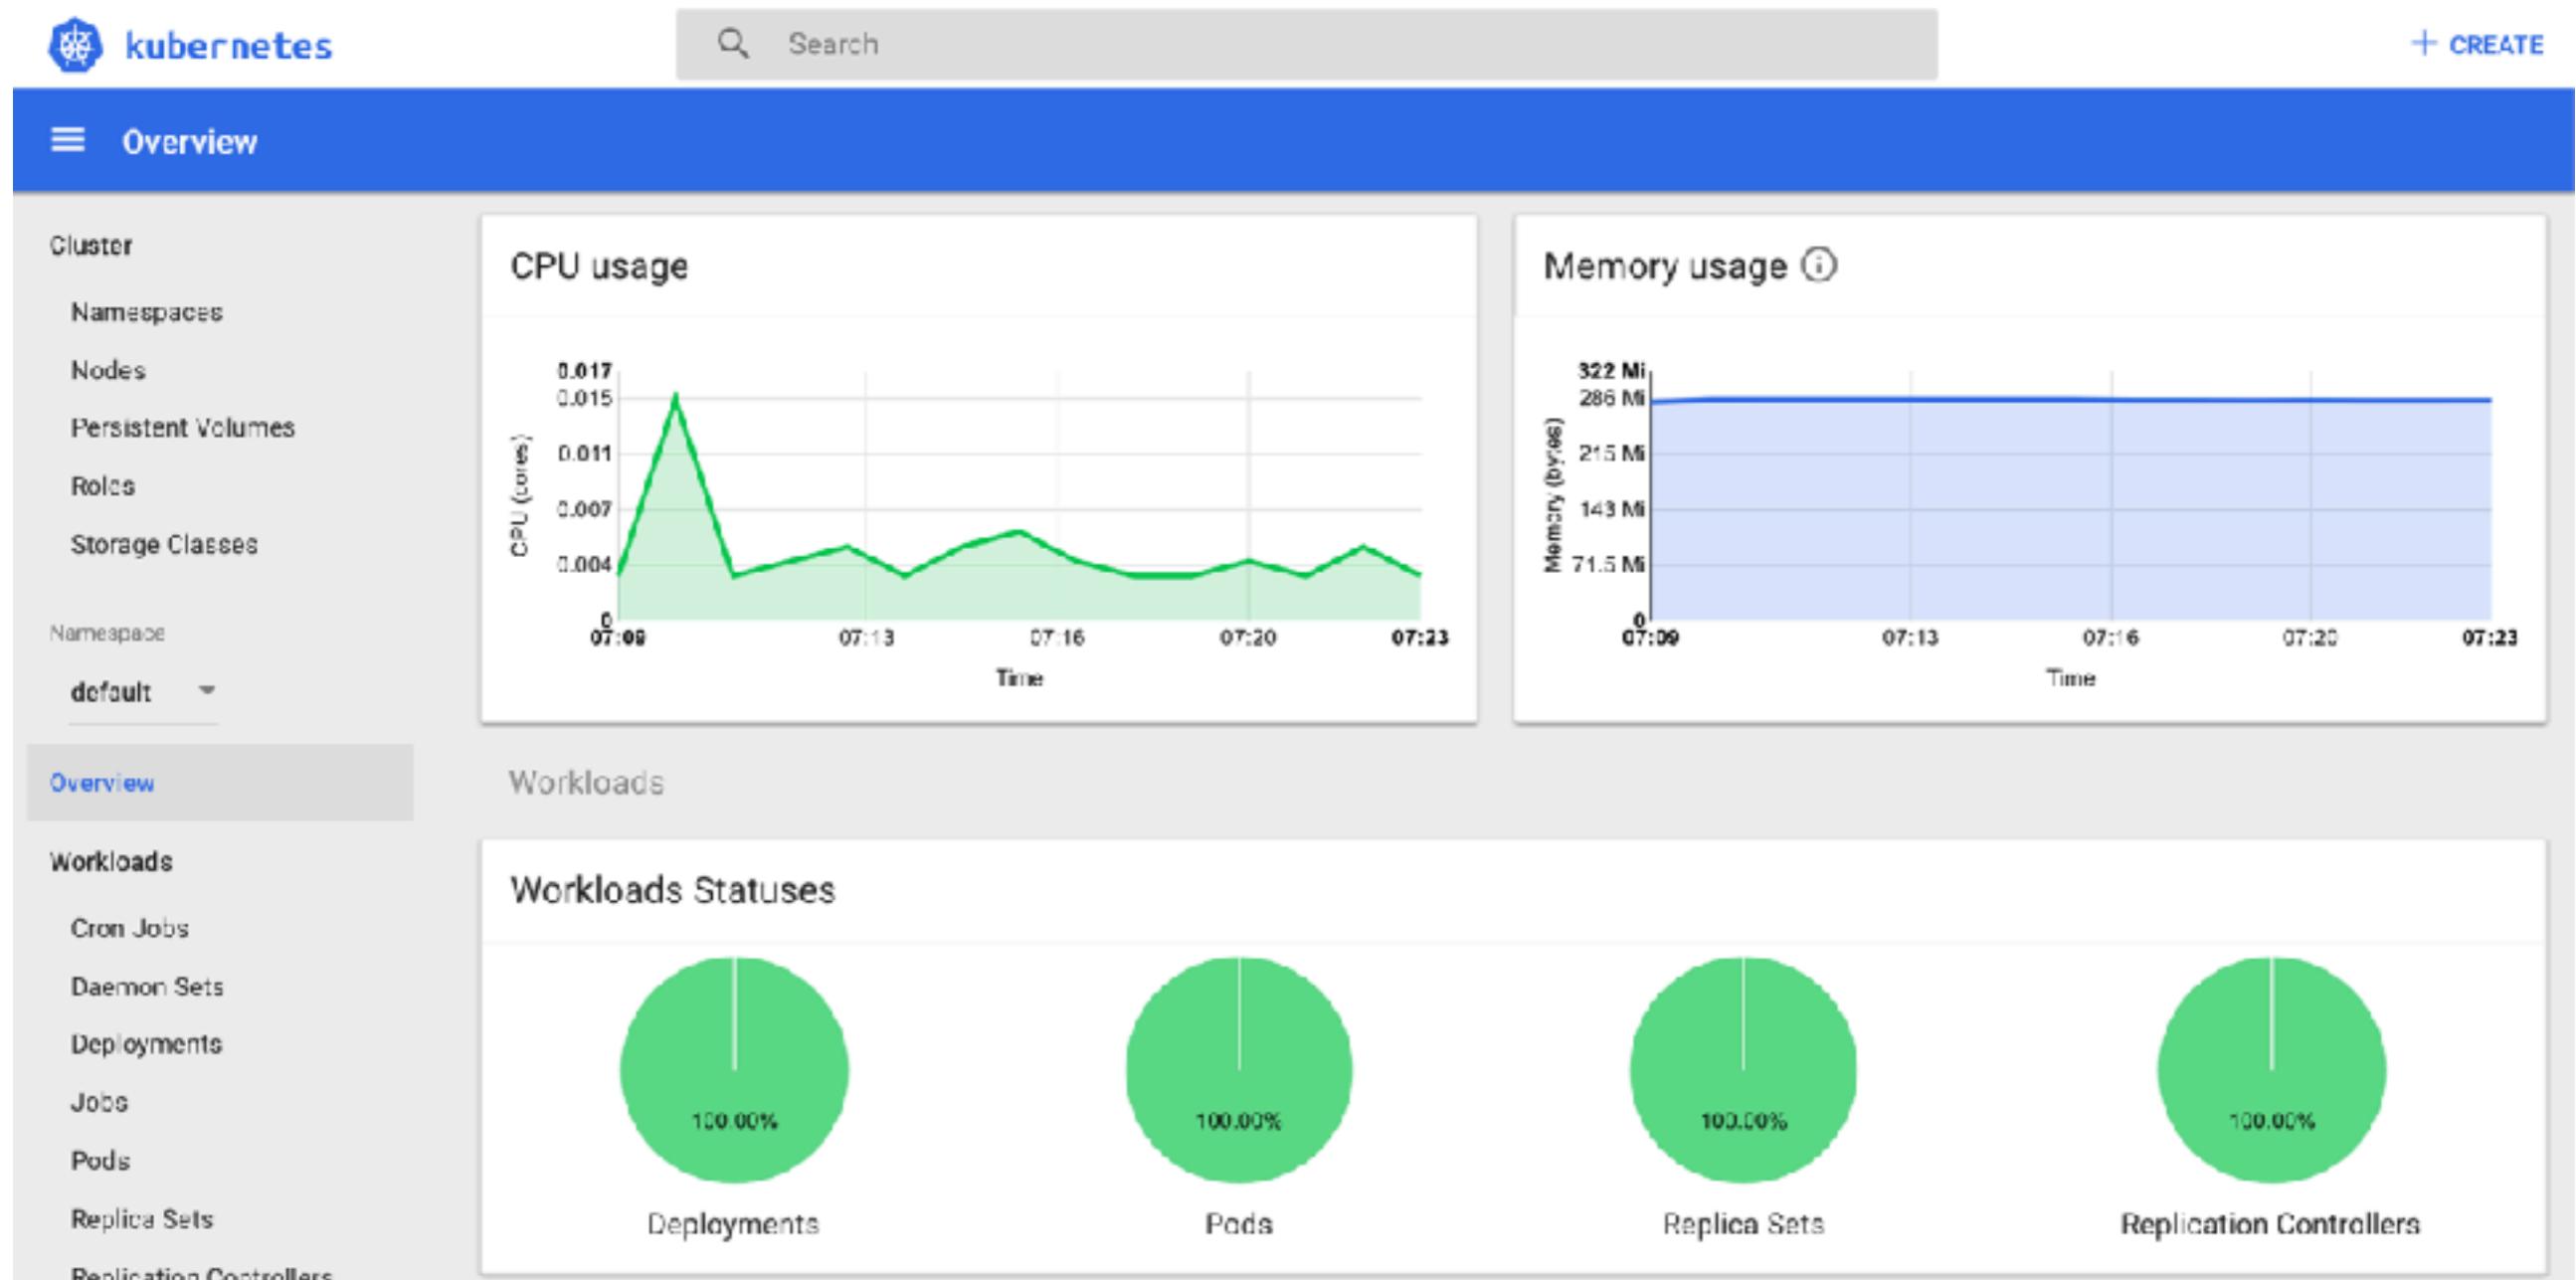
# 6. See result (2)

## \$kubectl describe hpa

```
Metrics:                               ( current / target )
  resource cpu on pods  (as a percentage of request): 134% (269m) / 5%
Min replicas:                           1
Max replicas:                           5
Conditions:
  Type      Status  Reason            Message
  ----      ----   ----
  AbleToScale False   BackoffBoth    the time since the previous scale is still within both the downscale and upscale to
  rbidden windows
  ScalingActive True    ValidMetricFound the HPA was able to successfully calculate a replica count from cpu resource utilization (percentage of request)
  ScalingLimited True   TooManyReplicas the desired replica count is more than the maximum replica count
Events:
  Type      Reason          Age      From            Message
  ----      ----           ----   ----
  Normal   SuccessfulRescale 1m      horizontal-pod-autoscaler  New size: 4; reason: cpu resource utilization (percentage of request) above target
  Warning  FailedGetResourceMetric 19s      horizontal-pod-autoscaler  unable to get metrics for resource cpu: unable to fetch metrics from API: the server could not find the requested resource (get pods.metrics.k8s.io)
  Warning  FailedUpdateStatus     19s      horizontal-pod-autoscaler  Operation cannot be fulfilled on horizontalpodautoscalers.autoscaling "spring-boot-service-deployment": the object has been modified; please apply your changes to the latest version and try again
  Warning  FailedComputeMetricsReplicas 19s      horizontal-pod-autoscaler  failed to get cpu utilization: unable to get metrics for resource cpu: unable to fetch metrics from API: the server could not find the requested resource (get pods.metrics.k8s.io)
```



# 7. See result in dashboard (1)



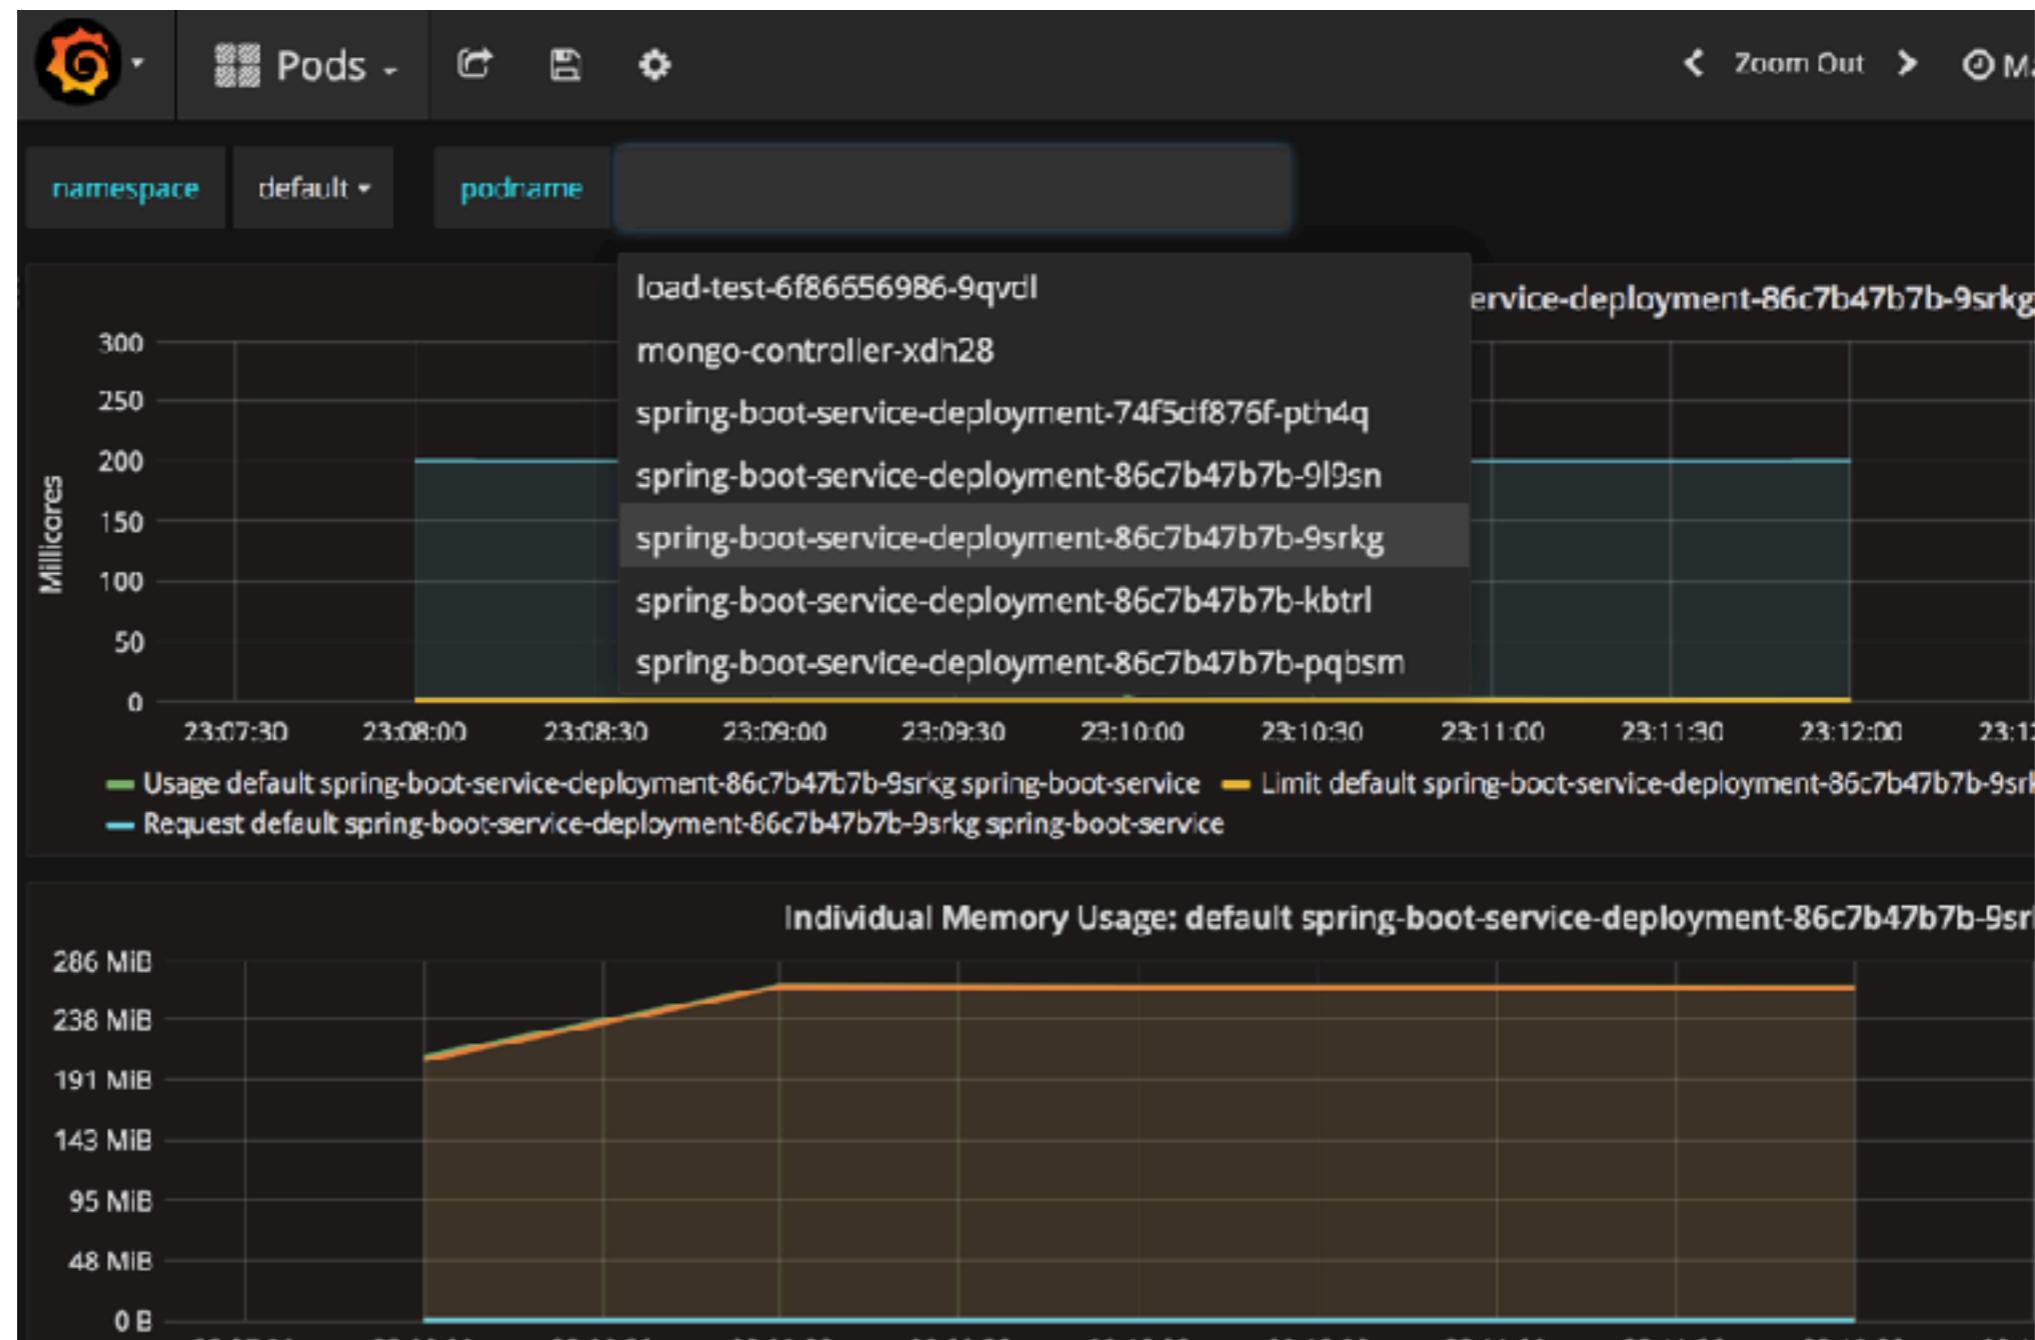
# 7. See result in dashboard (2)

Pods							
	Name	Node	Status	Restarts	Age	CPU (cores)	Memory (bytes)
✓	spring-boot-service-de	minikube	Running	0	2 minutes	 0.249	265.926 Mi
✓	spring-boot-service-de	minikube	Running	0	2 minutes	 0.03	254.156 Mi
✓	spring-boot-service-de	minikube	Running	0	2 minutes	 0.001	94.840 Mi
✓	load-test-6f86656986-	minikube	Running	1	5 minutes	 0	688 Ki
✓	spring-boot-service-de	minikube	Running	0	16 minutes	 0.018	33.211 Mi
✓	mongo-controller-xdh2	minikube	Running	0	an hour	 0.004	24.453 Mi



# 8. See result

## In Grafana dashboard



# Stateful Application Deployment



# Stateless vs Stateful

## Stateless

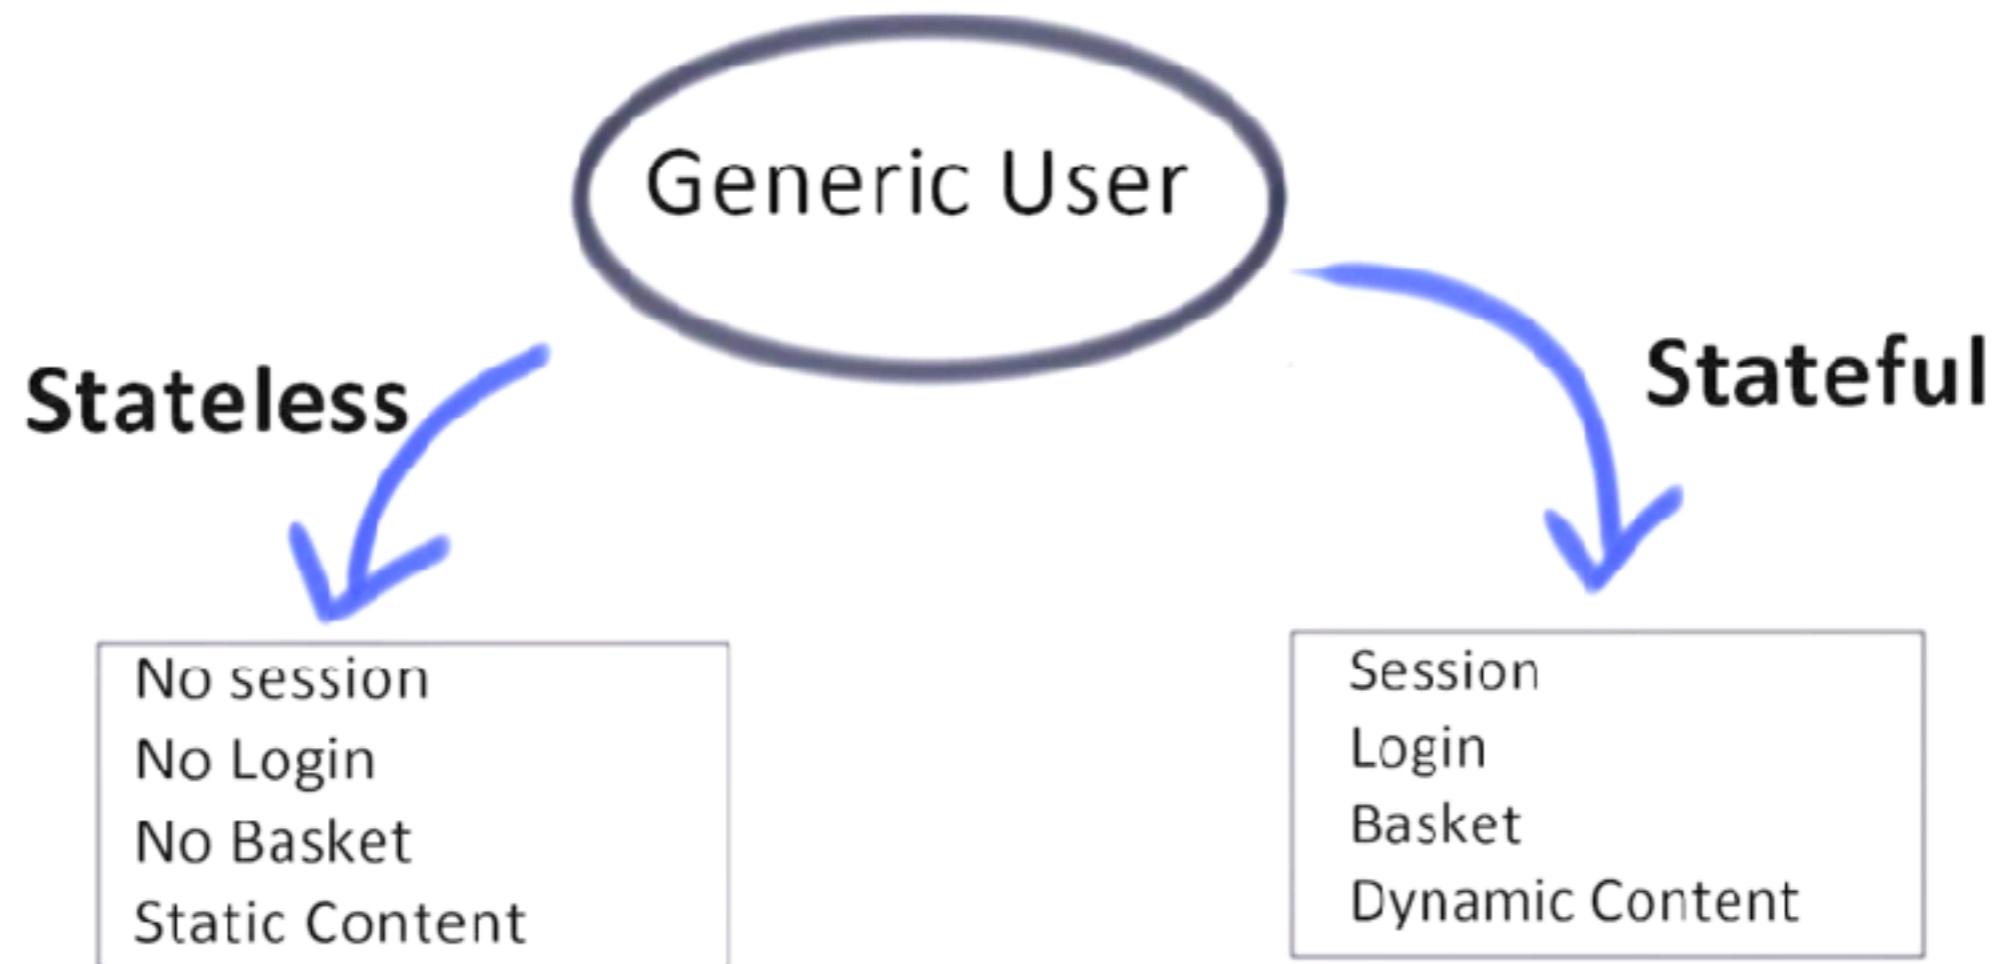
No information about a transaction/request is maintained after it is processed.

## Stateful

State information is kept even after a transaction/request has been processed.



# Stateless vs Stateful



# Container world

Container designed for **Stateless** application

All Load Balance/dispatch process is **not** aware about **state** of application



# Solution for Stateful

Share centralize storage pool for all node

## **Web and Application server**

- Keep session and files in storage pool
- Every servers are read/write data on the same place



# Solution for Stateful

Share centralize storage pool for all node

## Database server

- Active/Active
- Active/Standby
- Keep data in storage pool

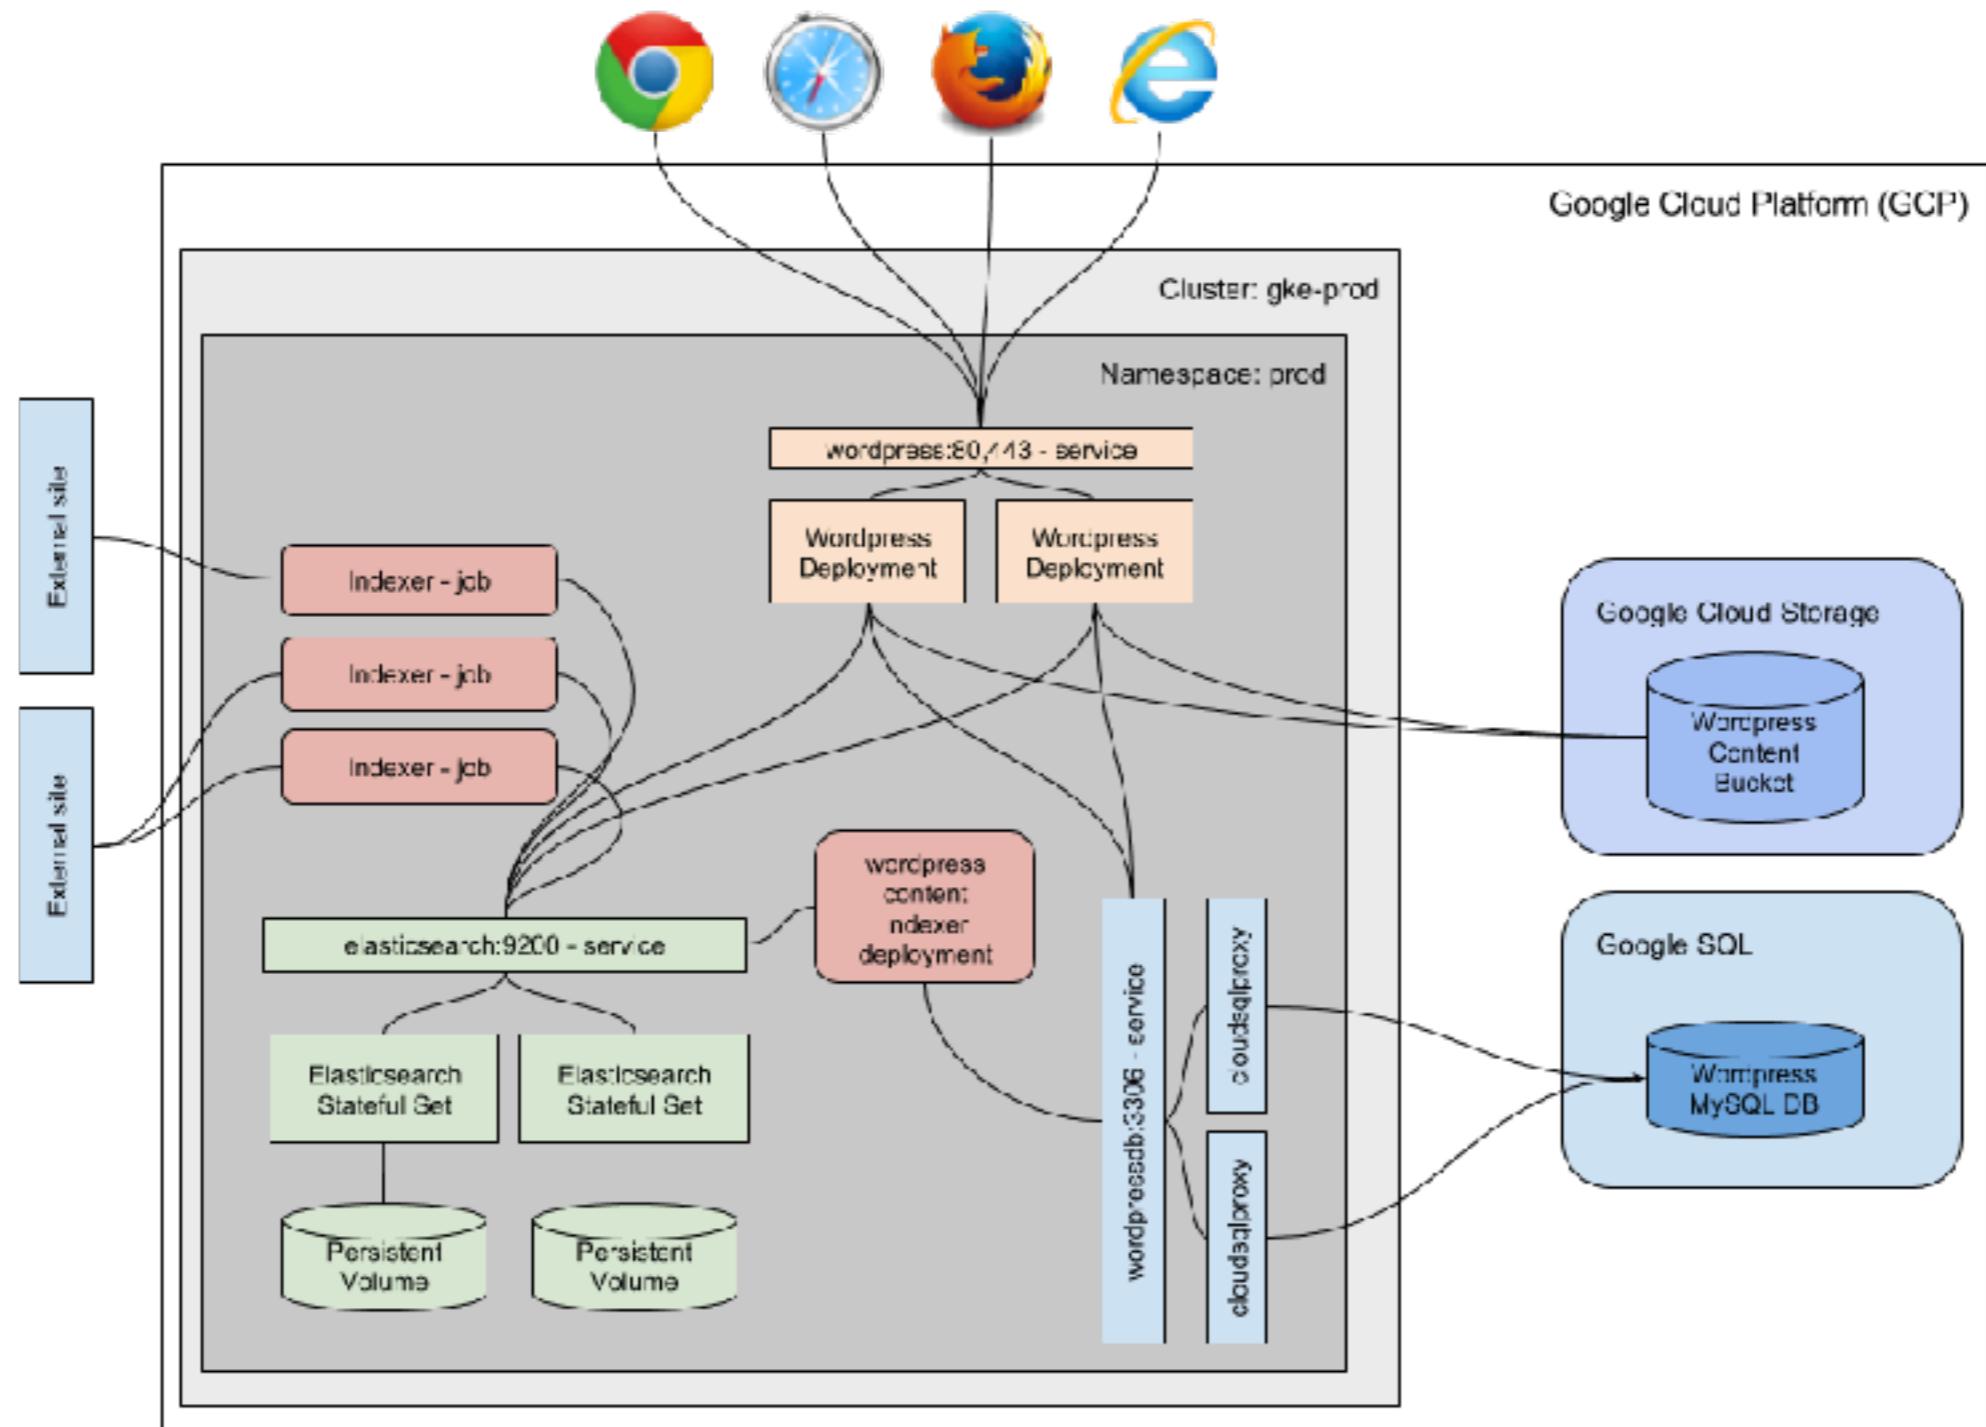


# Stateful with Kubernetes

Volume  
Statefulset



# Stateful with Kubernetes



<http://blogs.avalonconsult.com/blog/enterprise-web/cloud-computing/architecting-for-kubernetes/>

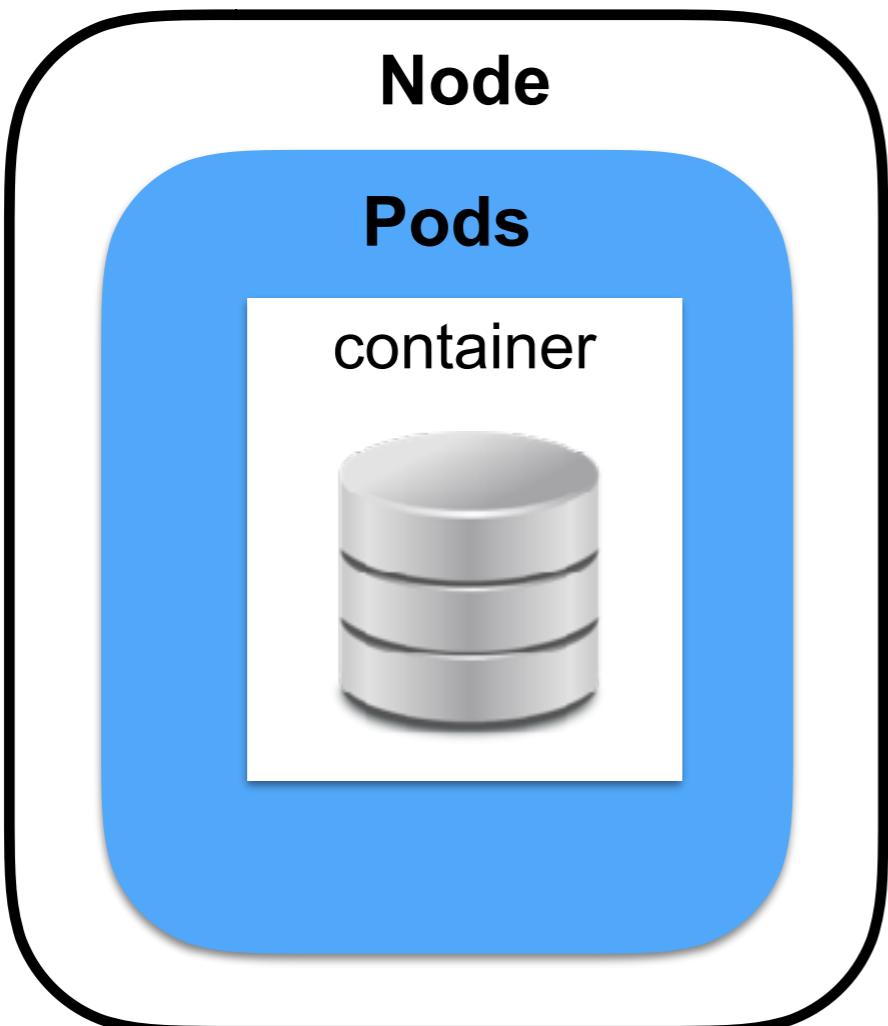


# Volume in Kubernetes

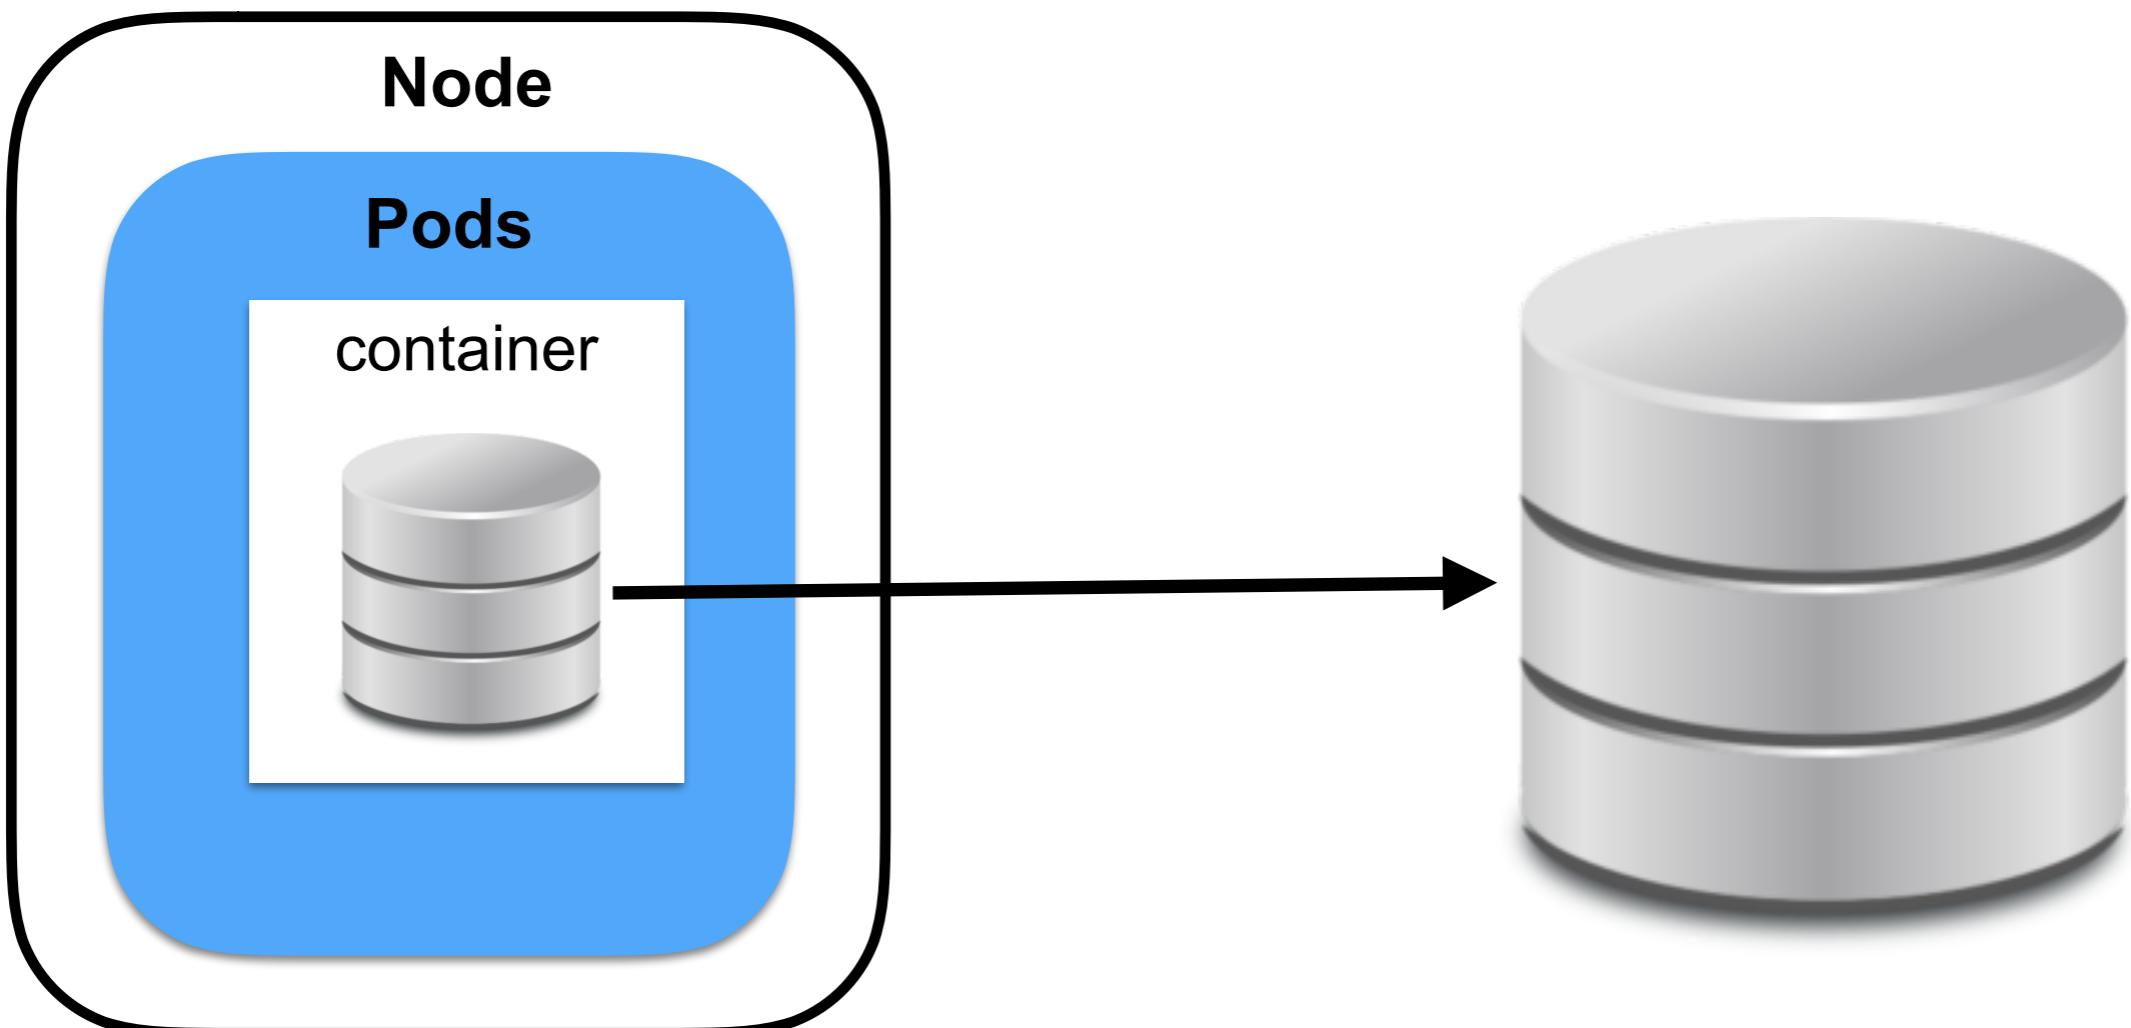
<https://kubernetes.io/docs/concepts/storage/volumes/>



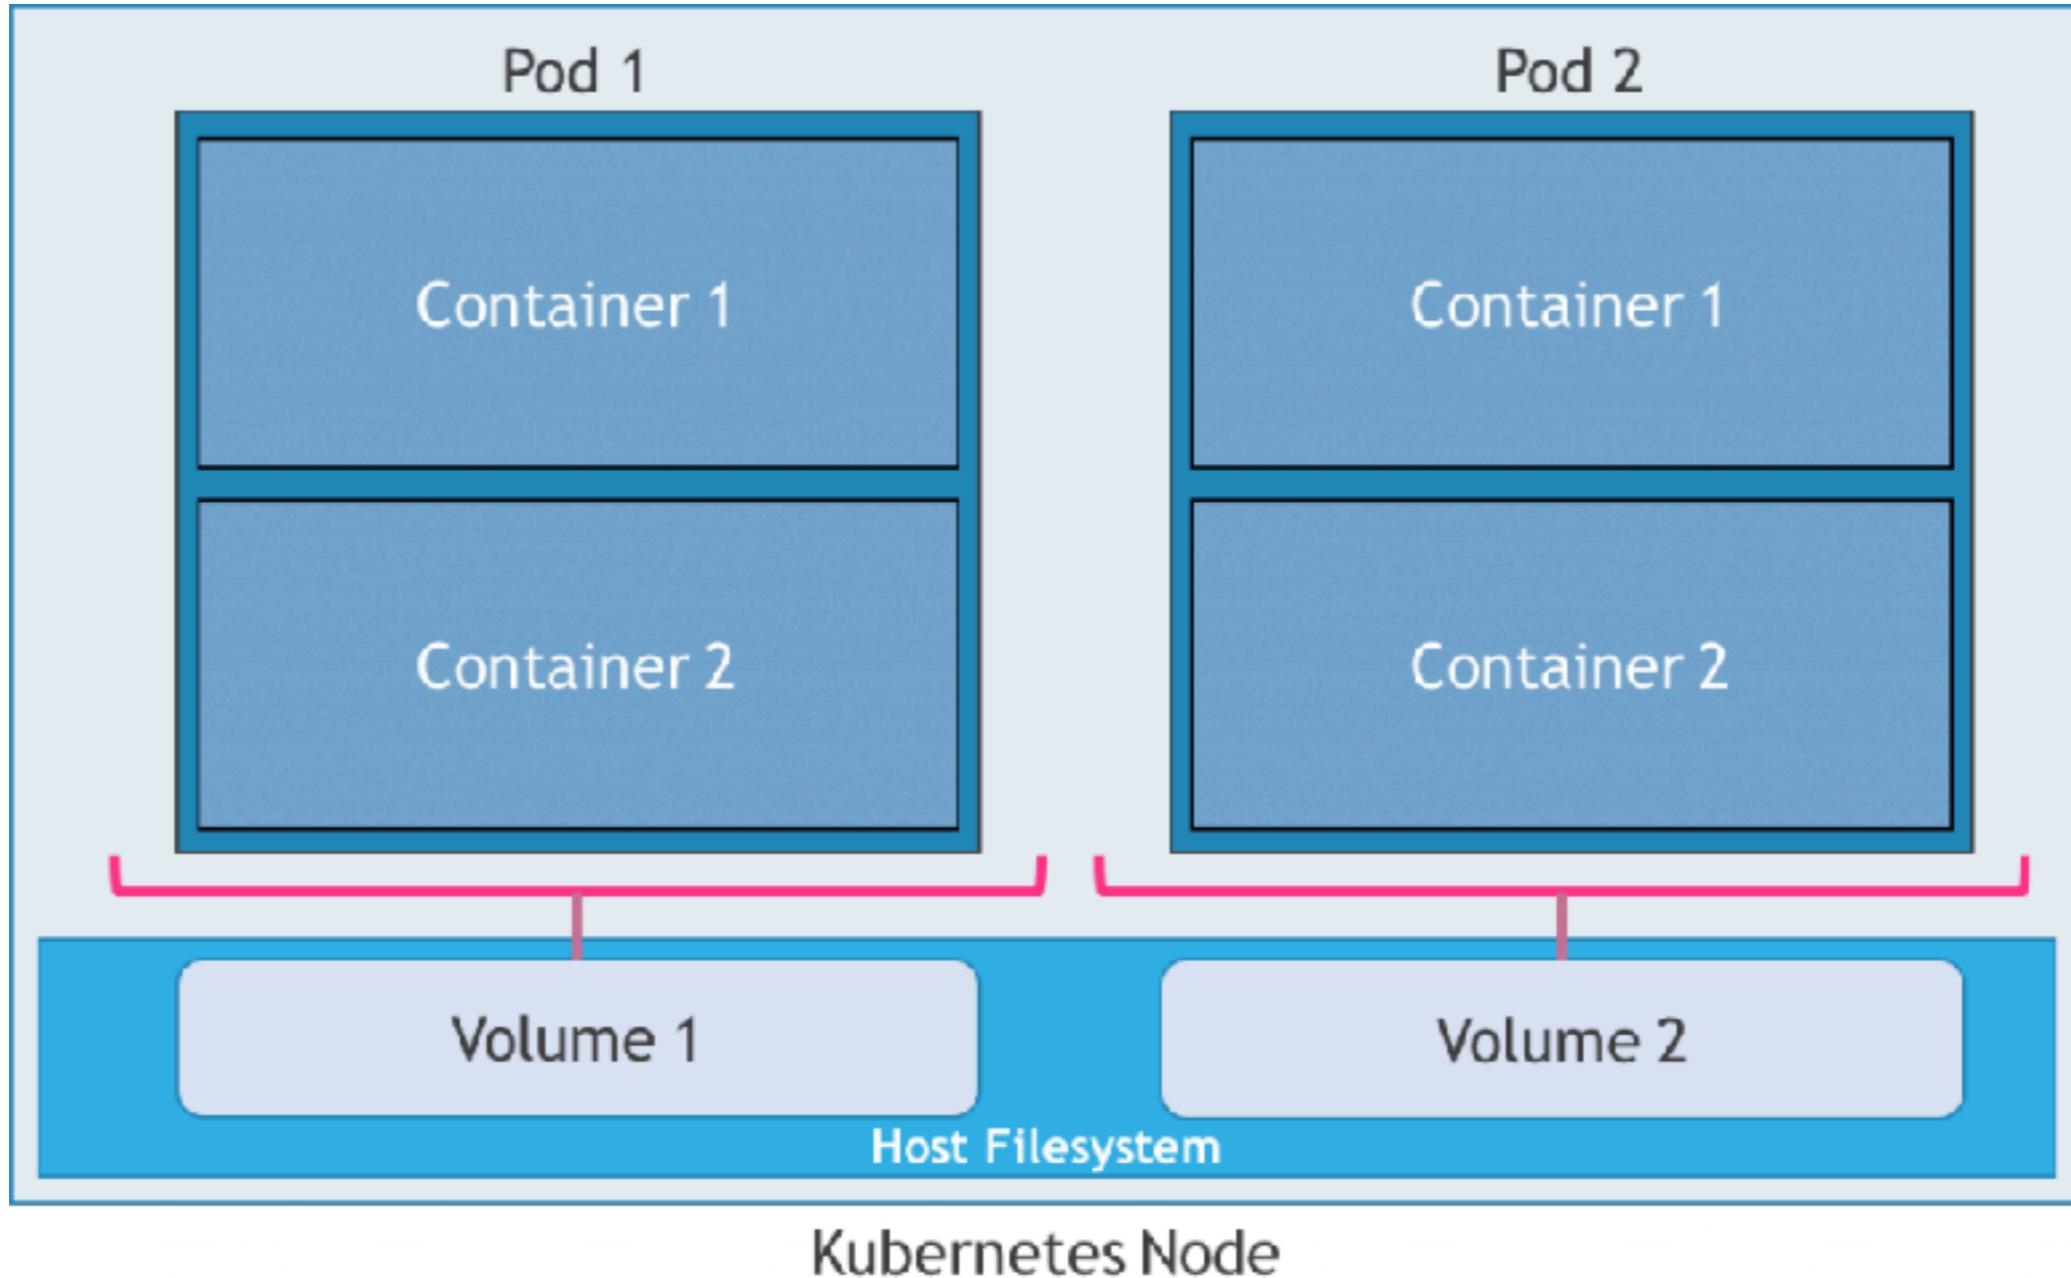
# Volume



# Volume



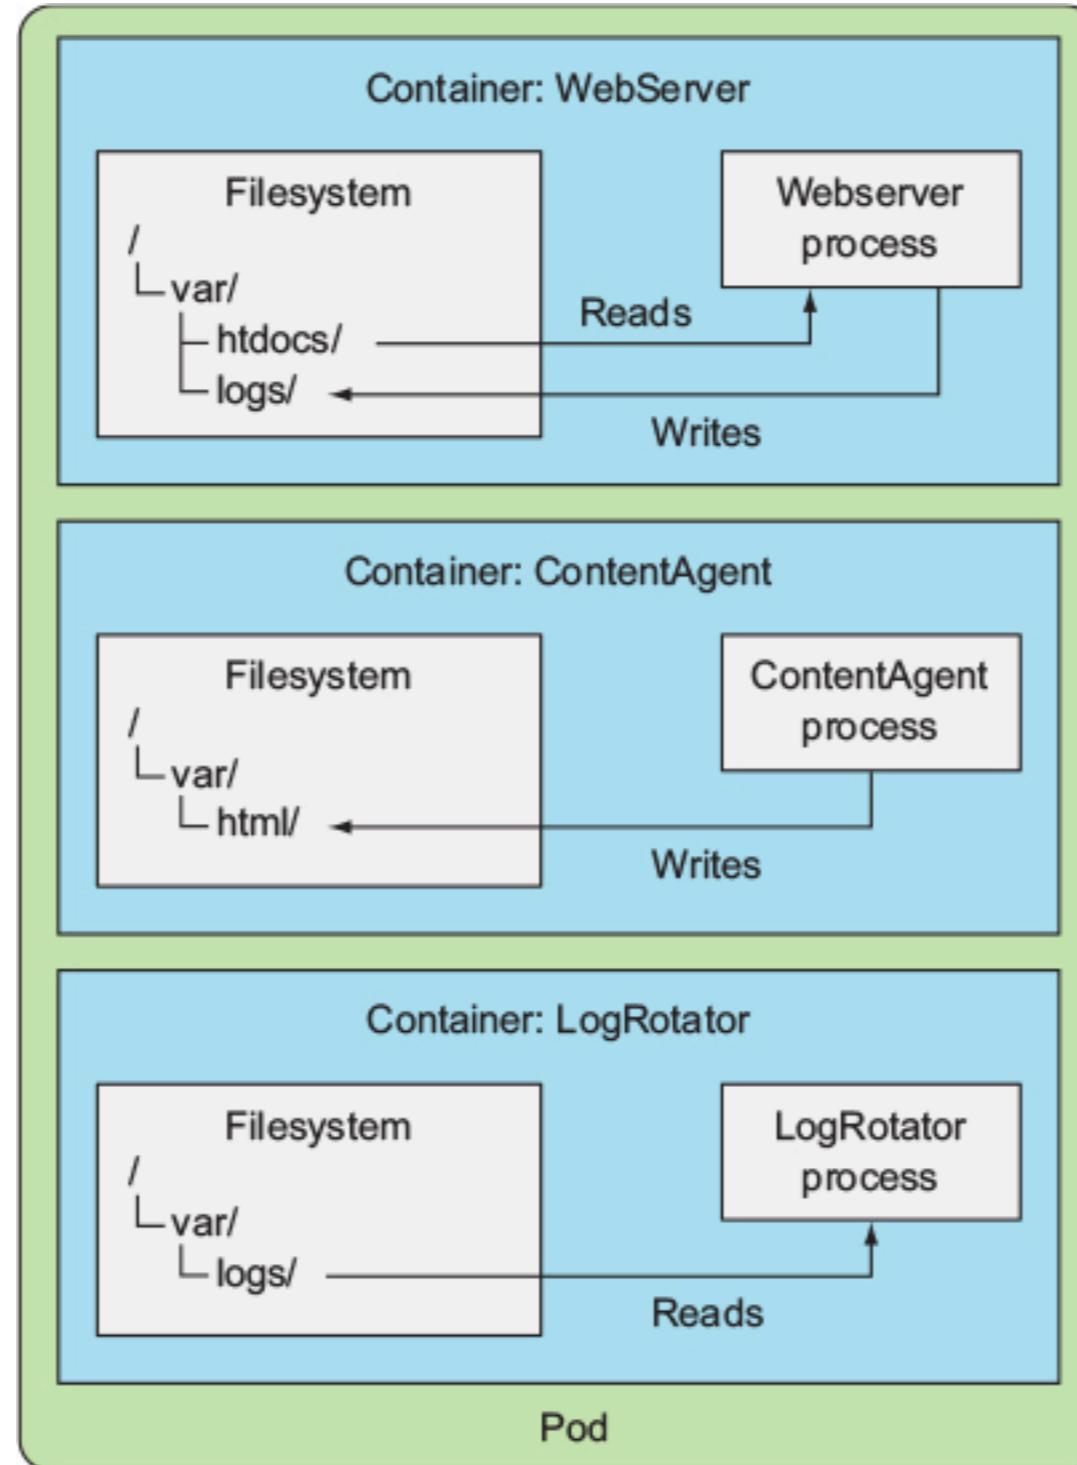
# Containers shared volume



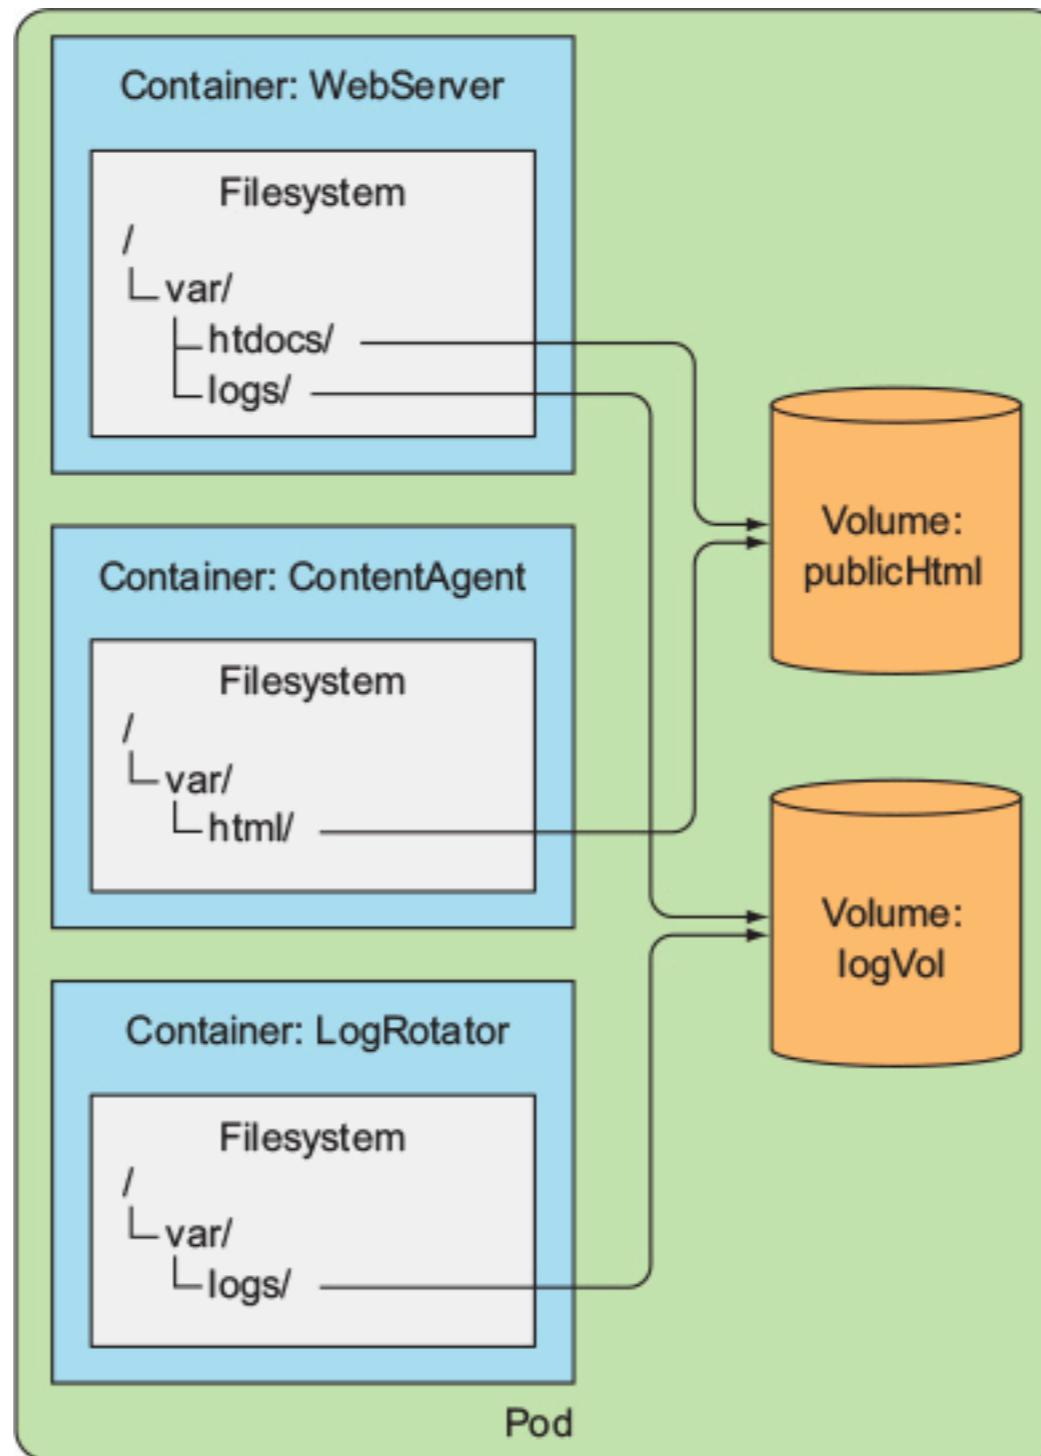
<https://thenewstack.io/strategies-running-stateful-applications-kubernetes-volumes/>



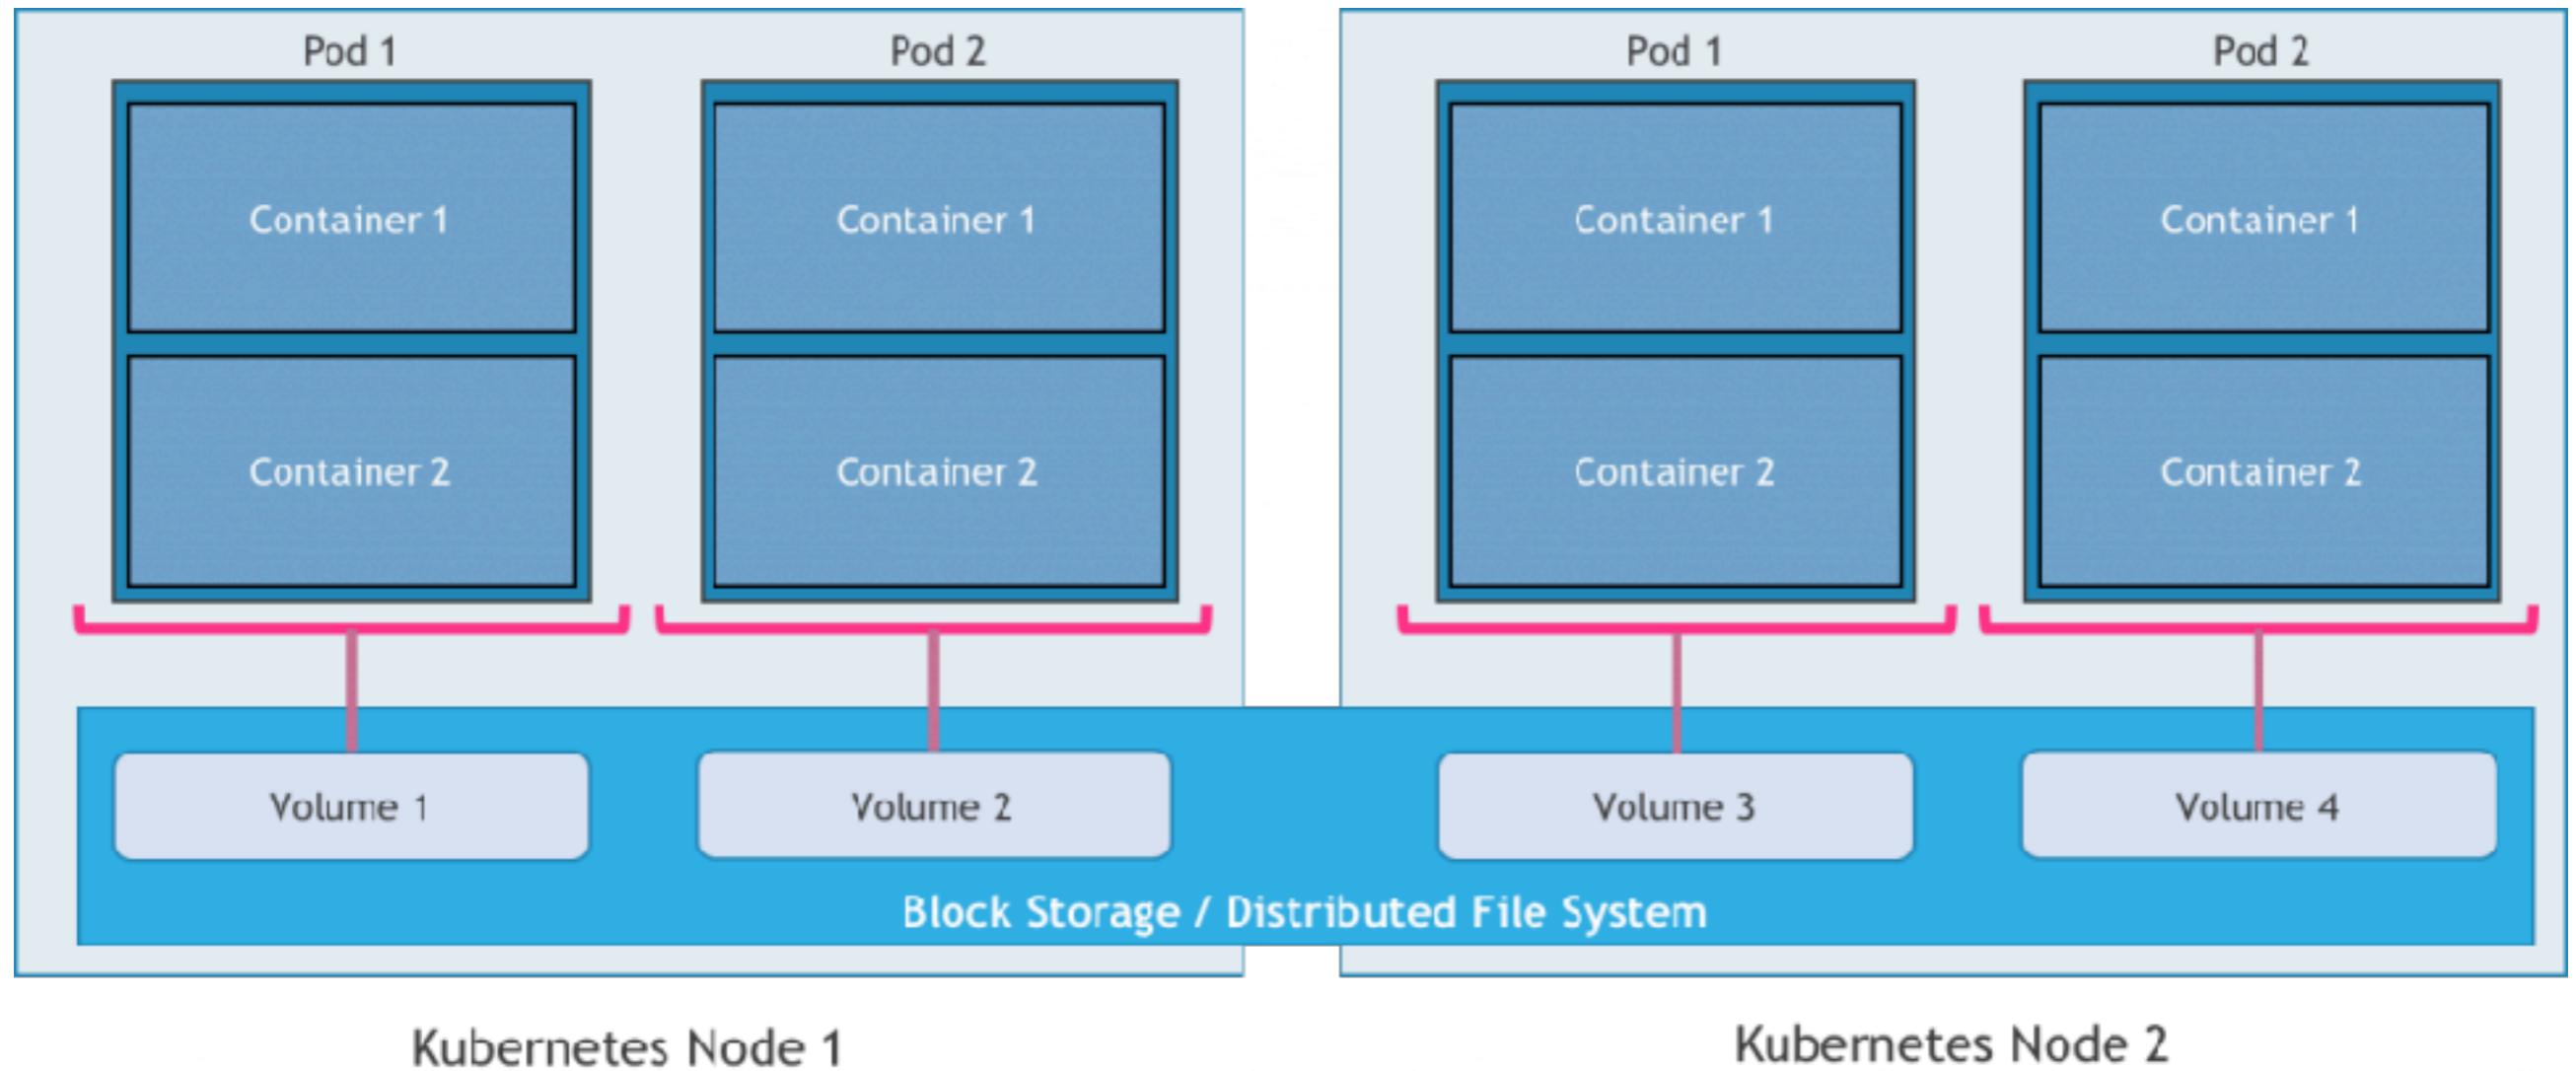
# Without shared storage



# Containers shared volume



# Nodes shared volume



<https://thenewstack.io/strategies-running-stateful-applications-kubernetes-volumes/>



# Volume

Containers in the same Pods share storage  
(EmptyDir)

- Read and write for all containers in Pods
- Container crash not effect to this storage
- When Pods is deleted then Empty will deleted

Data of Pods/container is **ephemeral**  
all data may loss when Pods is restarted



# Volume on Kubernetes

`EmptyDir`

`hostPath`

`gitRepo`

`NFS`

`icePersistentDisk`

`Flocker`

## **Persistent Volume Claim (PVC)**

<https://kubernetes.io/docs/concepts/storage/volumes/>



# Volume on Kubernetes

EmptyDir

HostPath

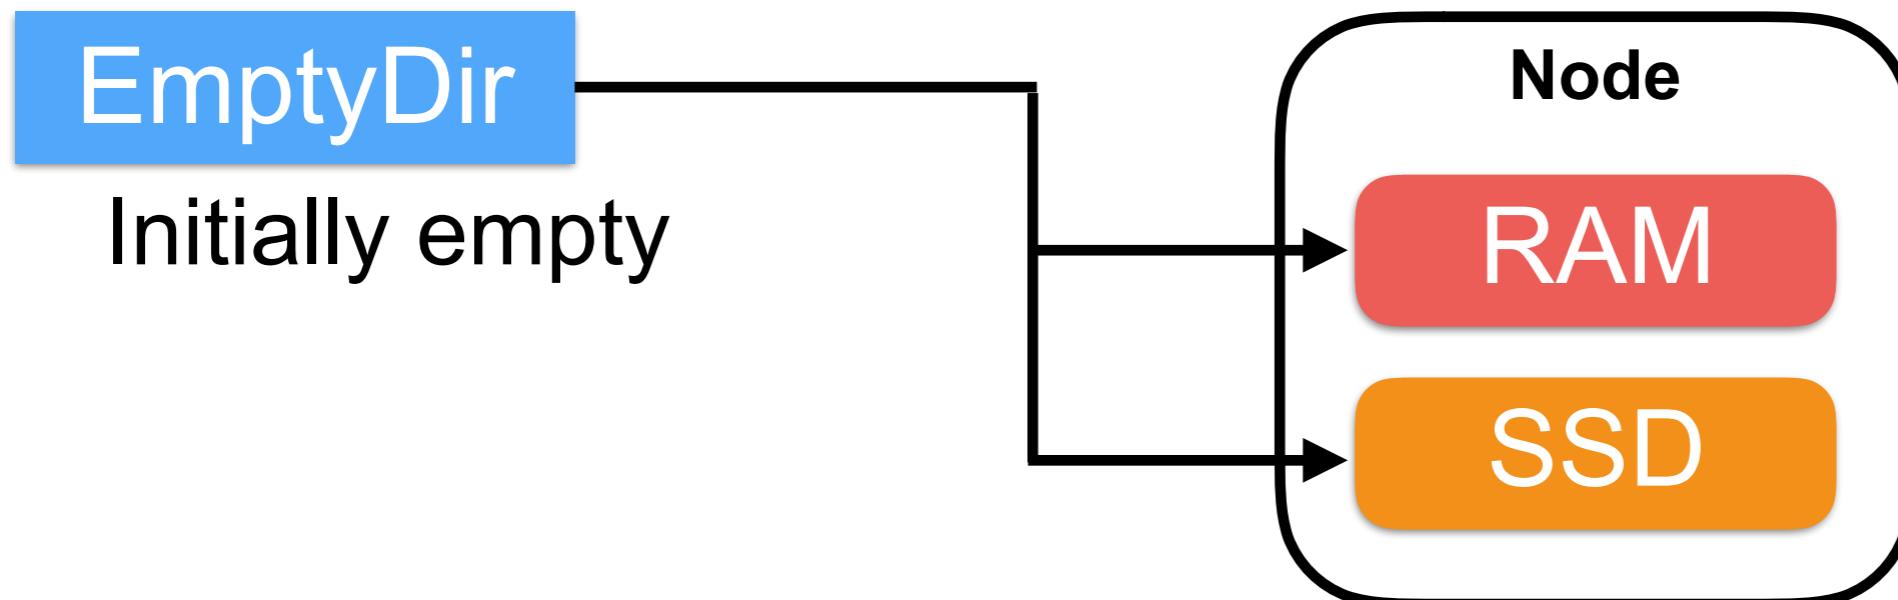
NFS

Cloud Volume

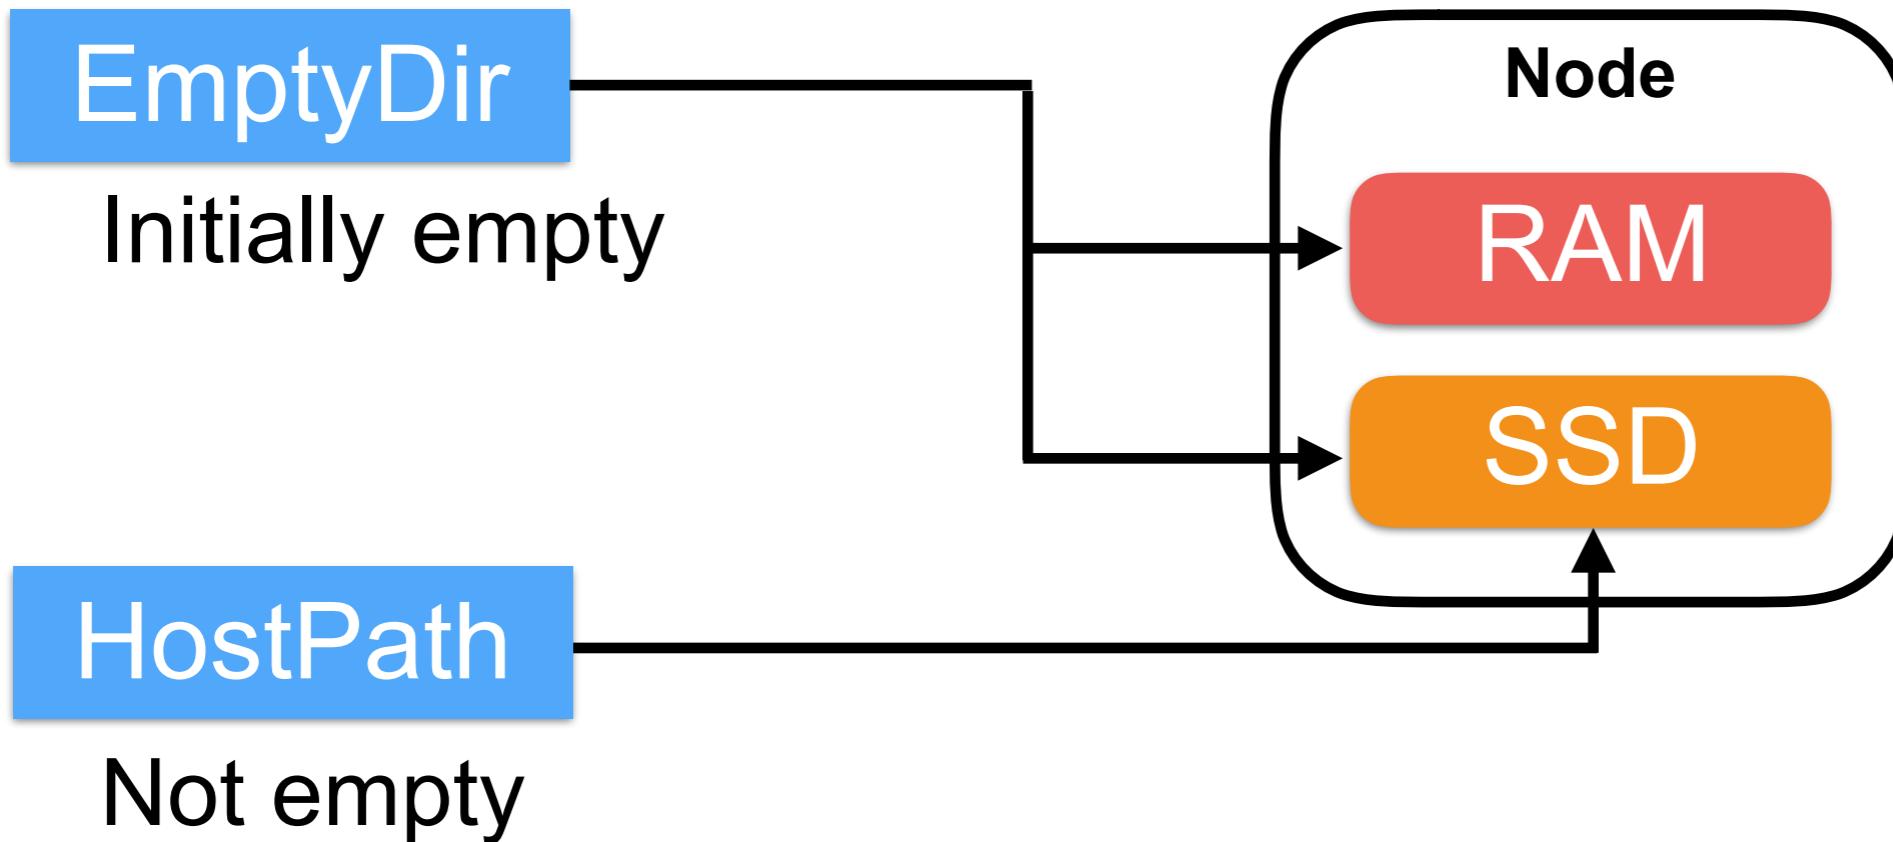
Persistent Volume Claim



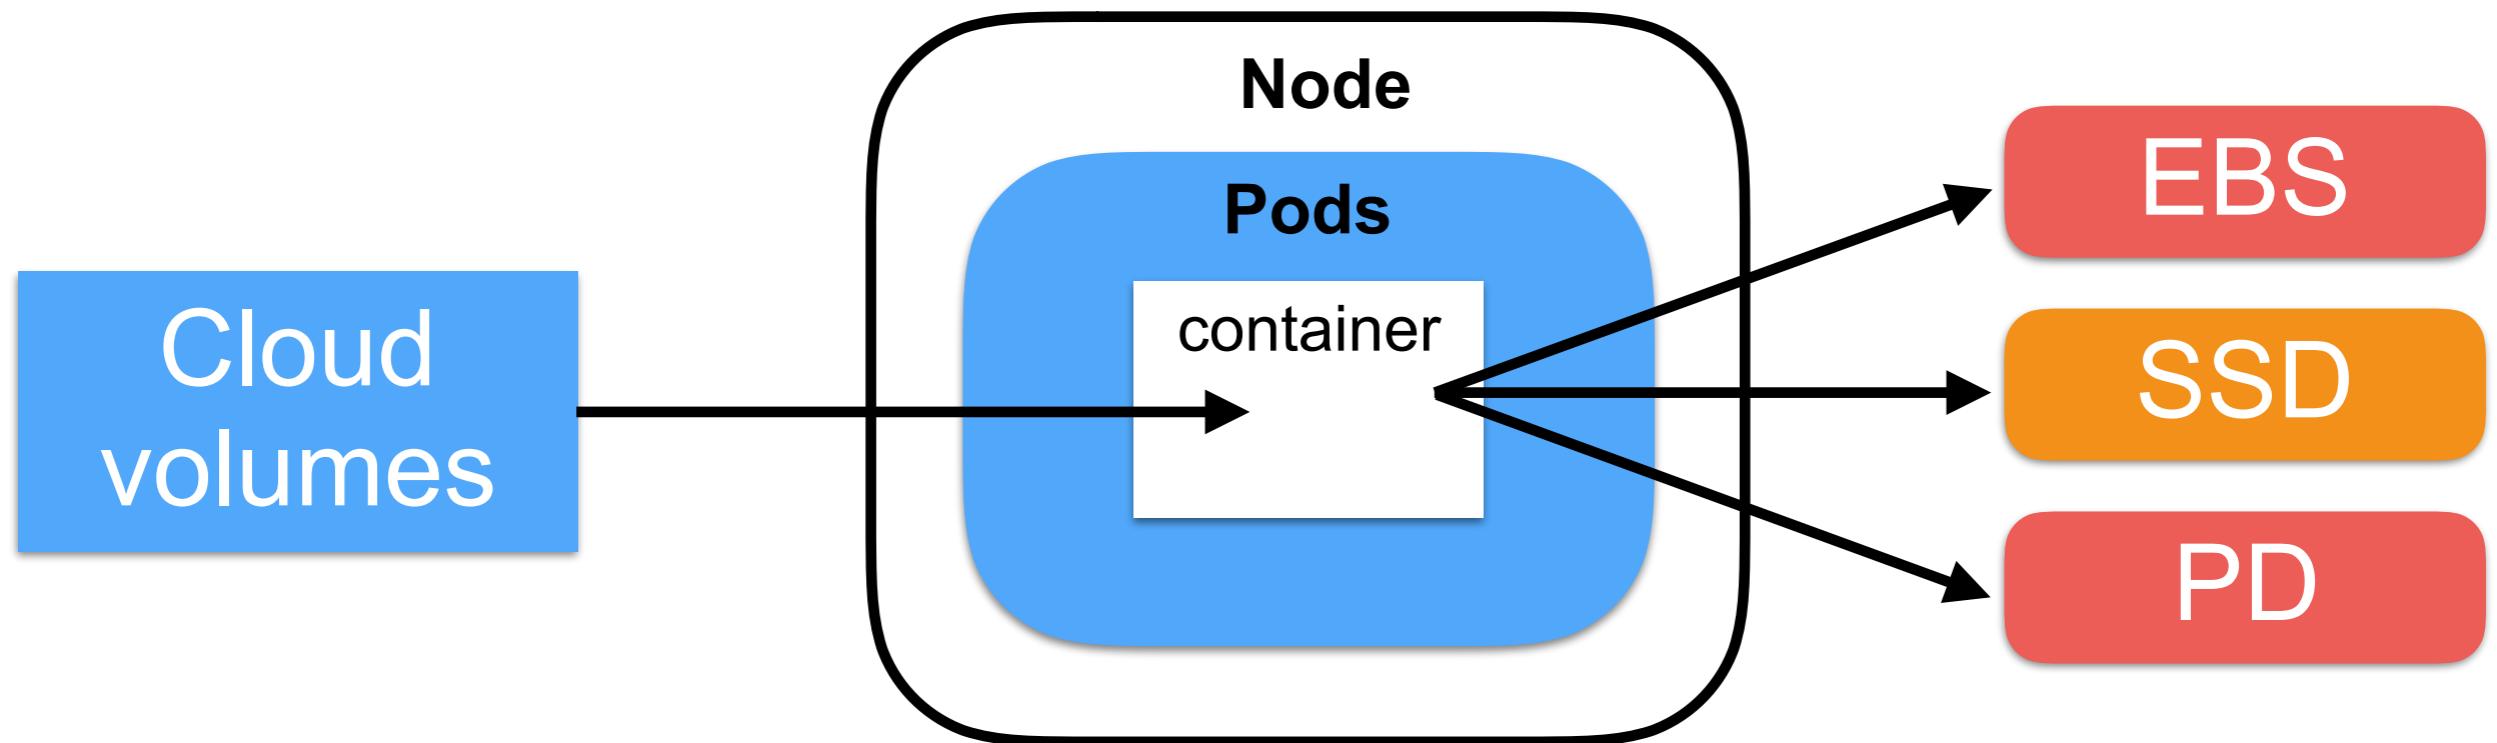
# Volume on Kubernetes



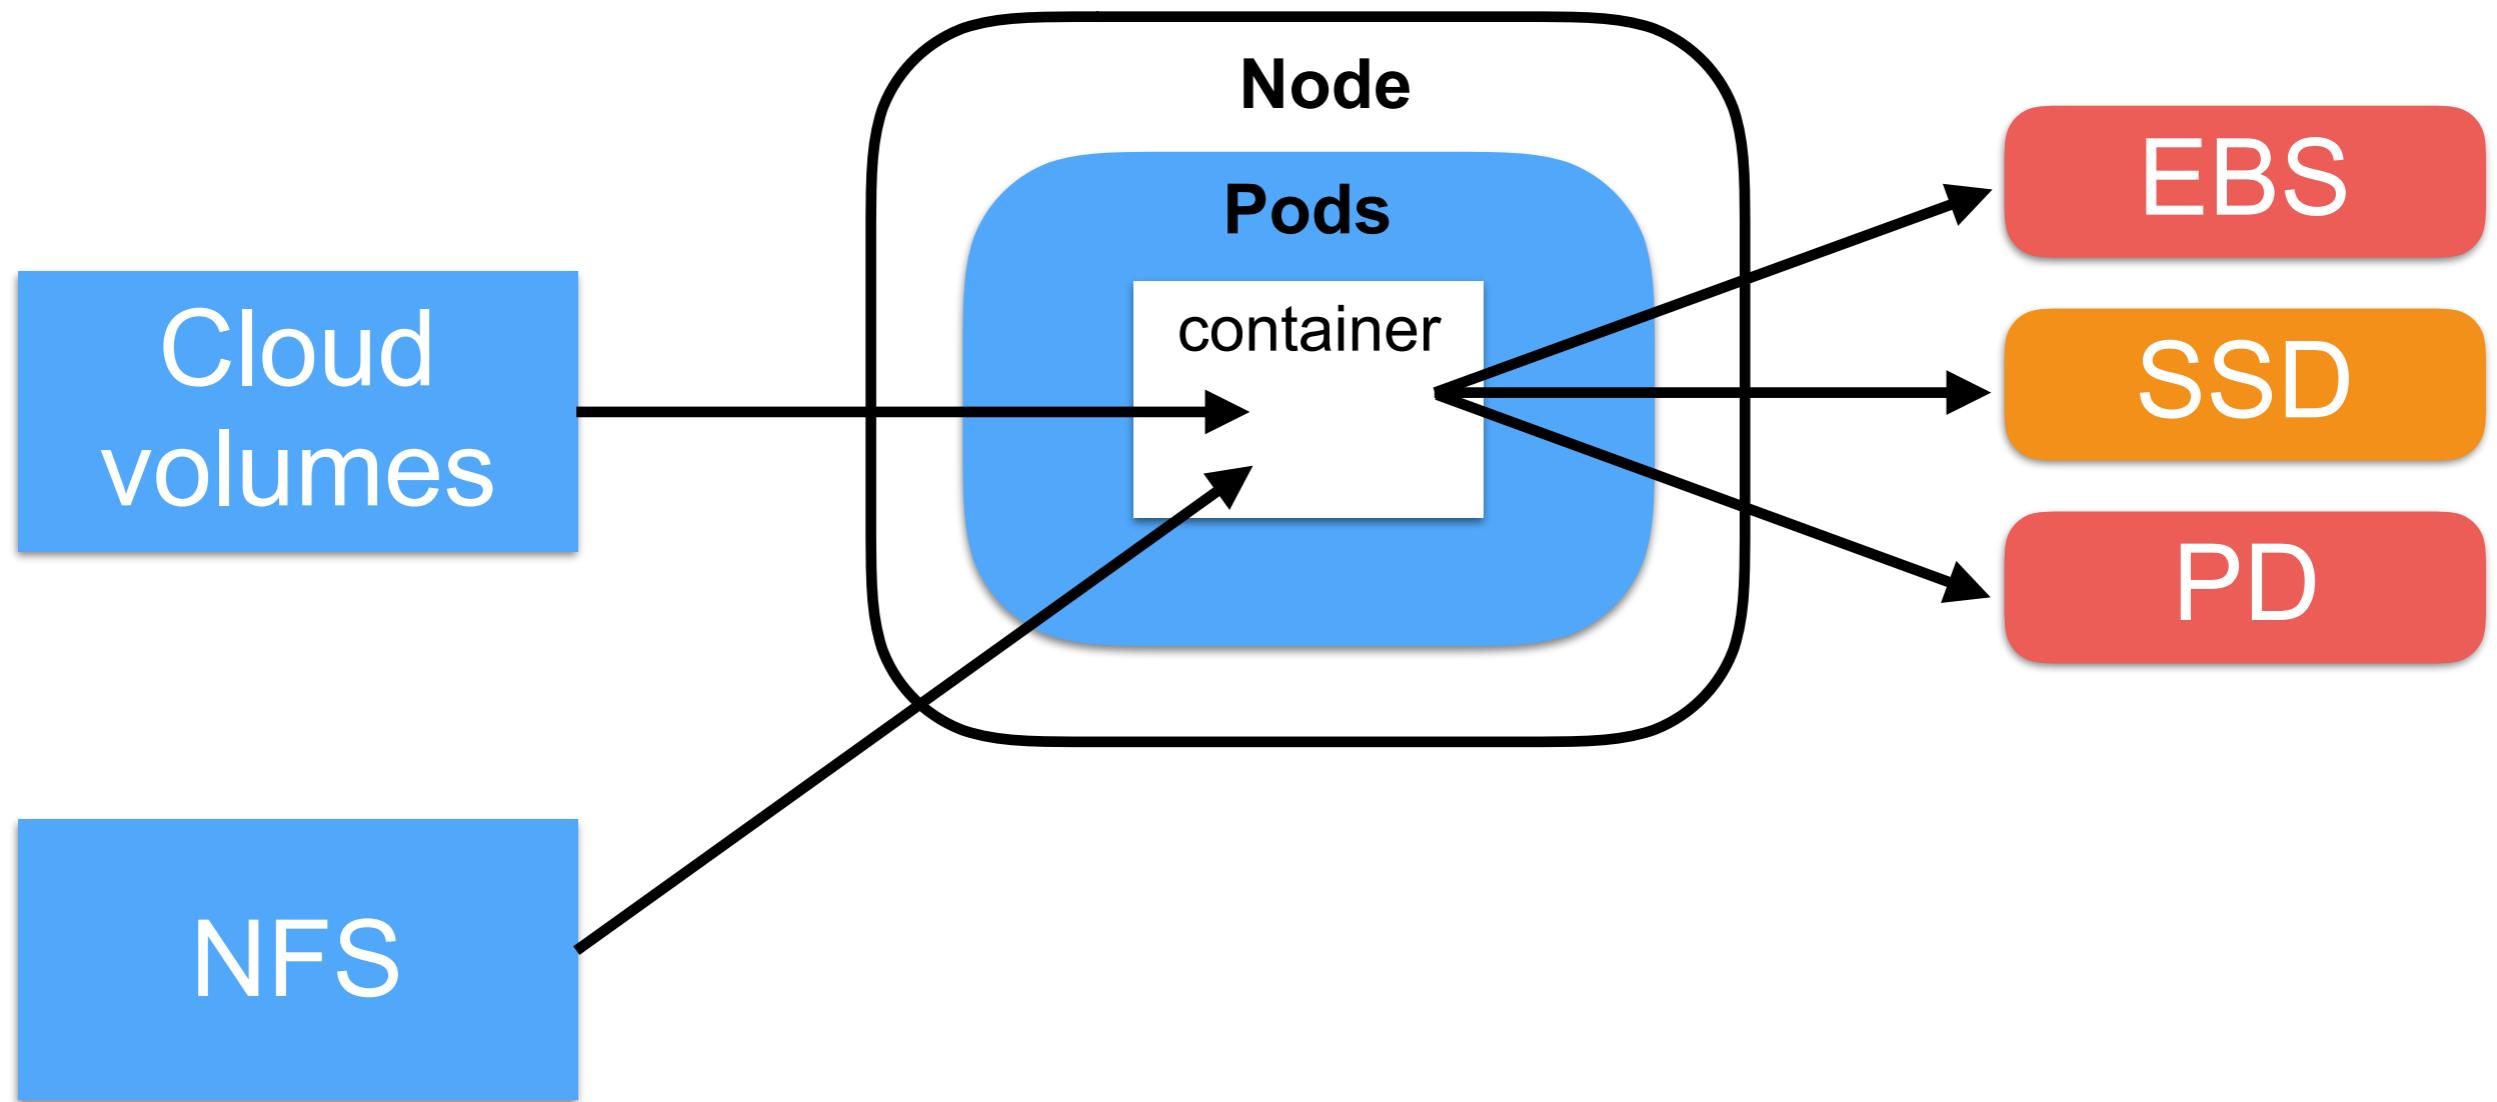
# Volume on Kubernetes



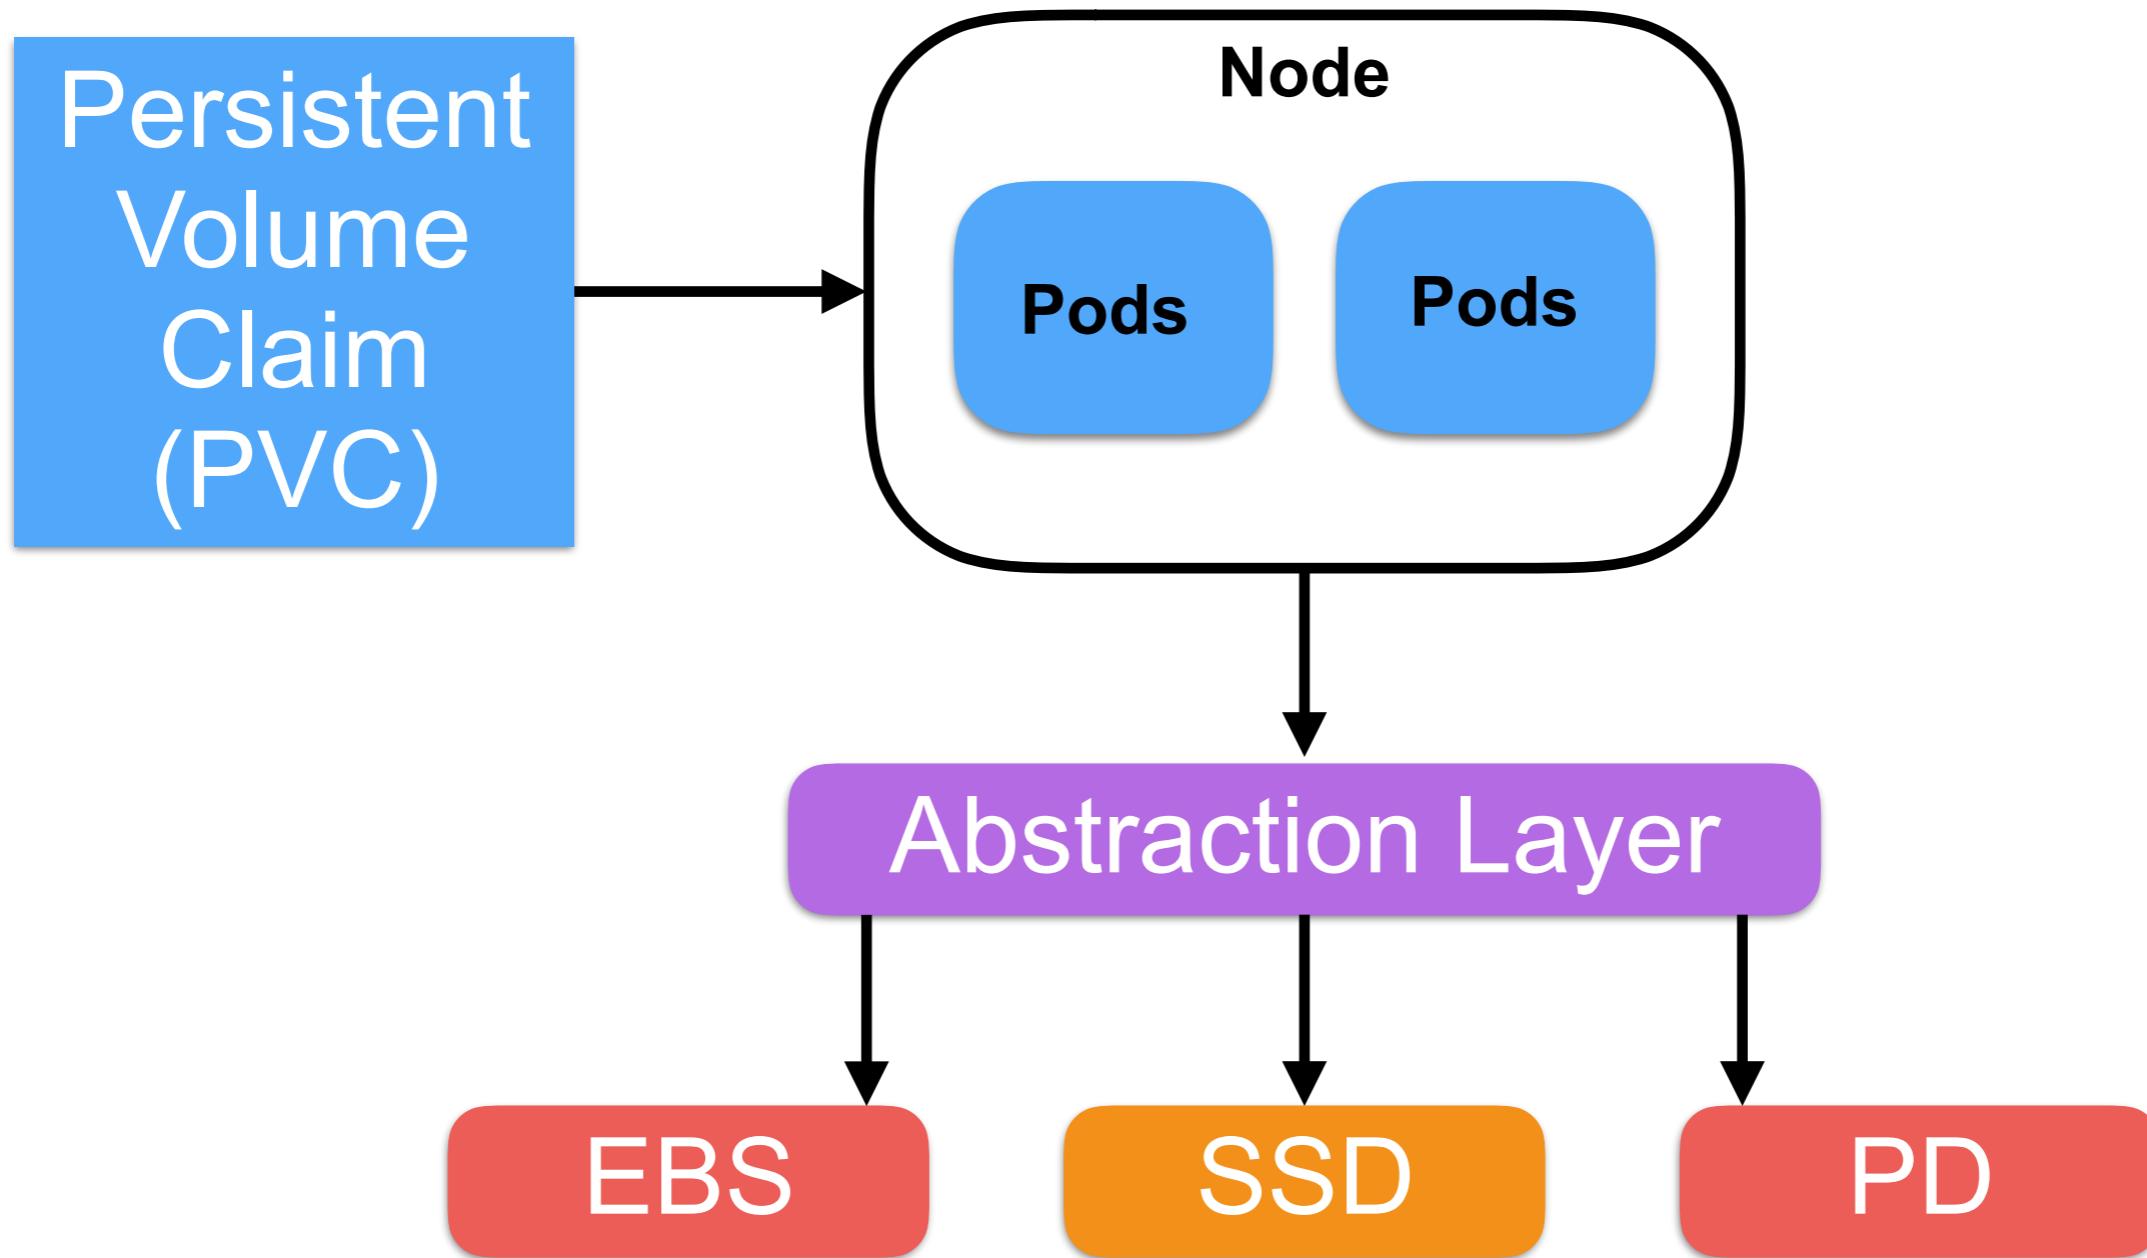
# Volume on Kubernetes



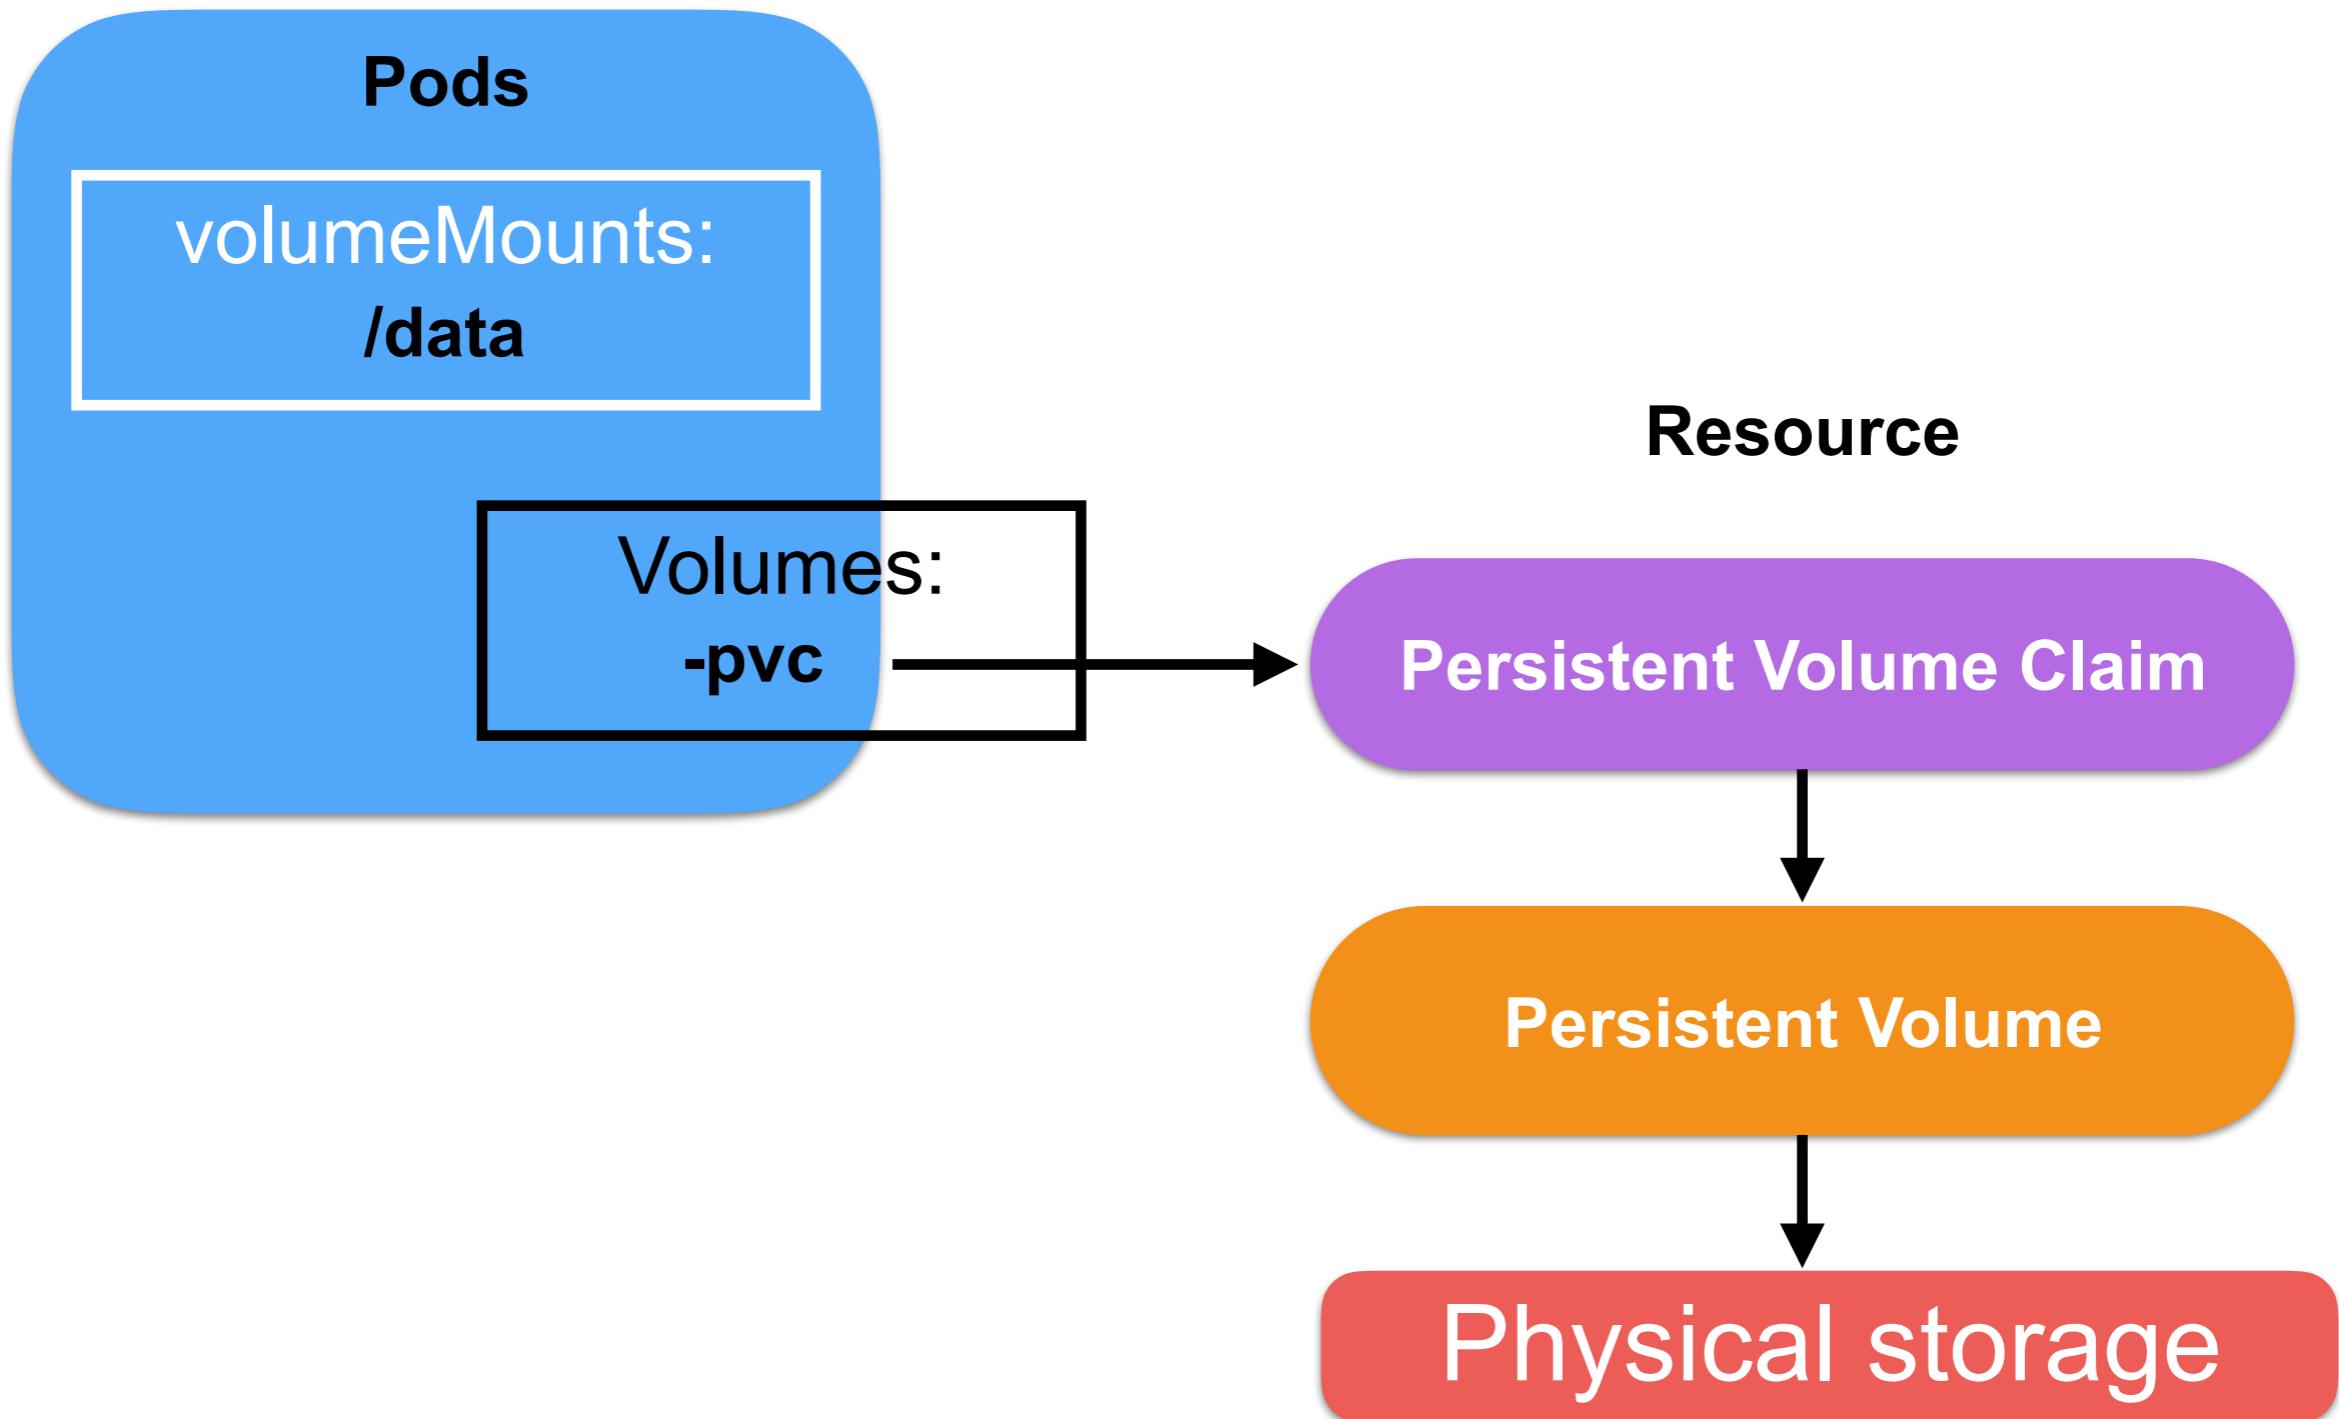
# Volume on Kubernetes



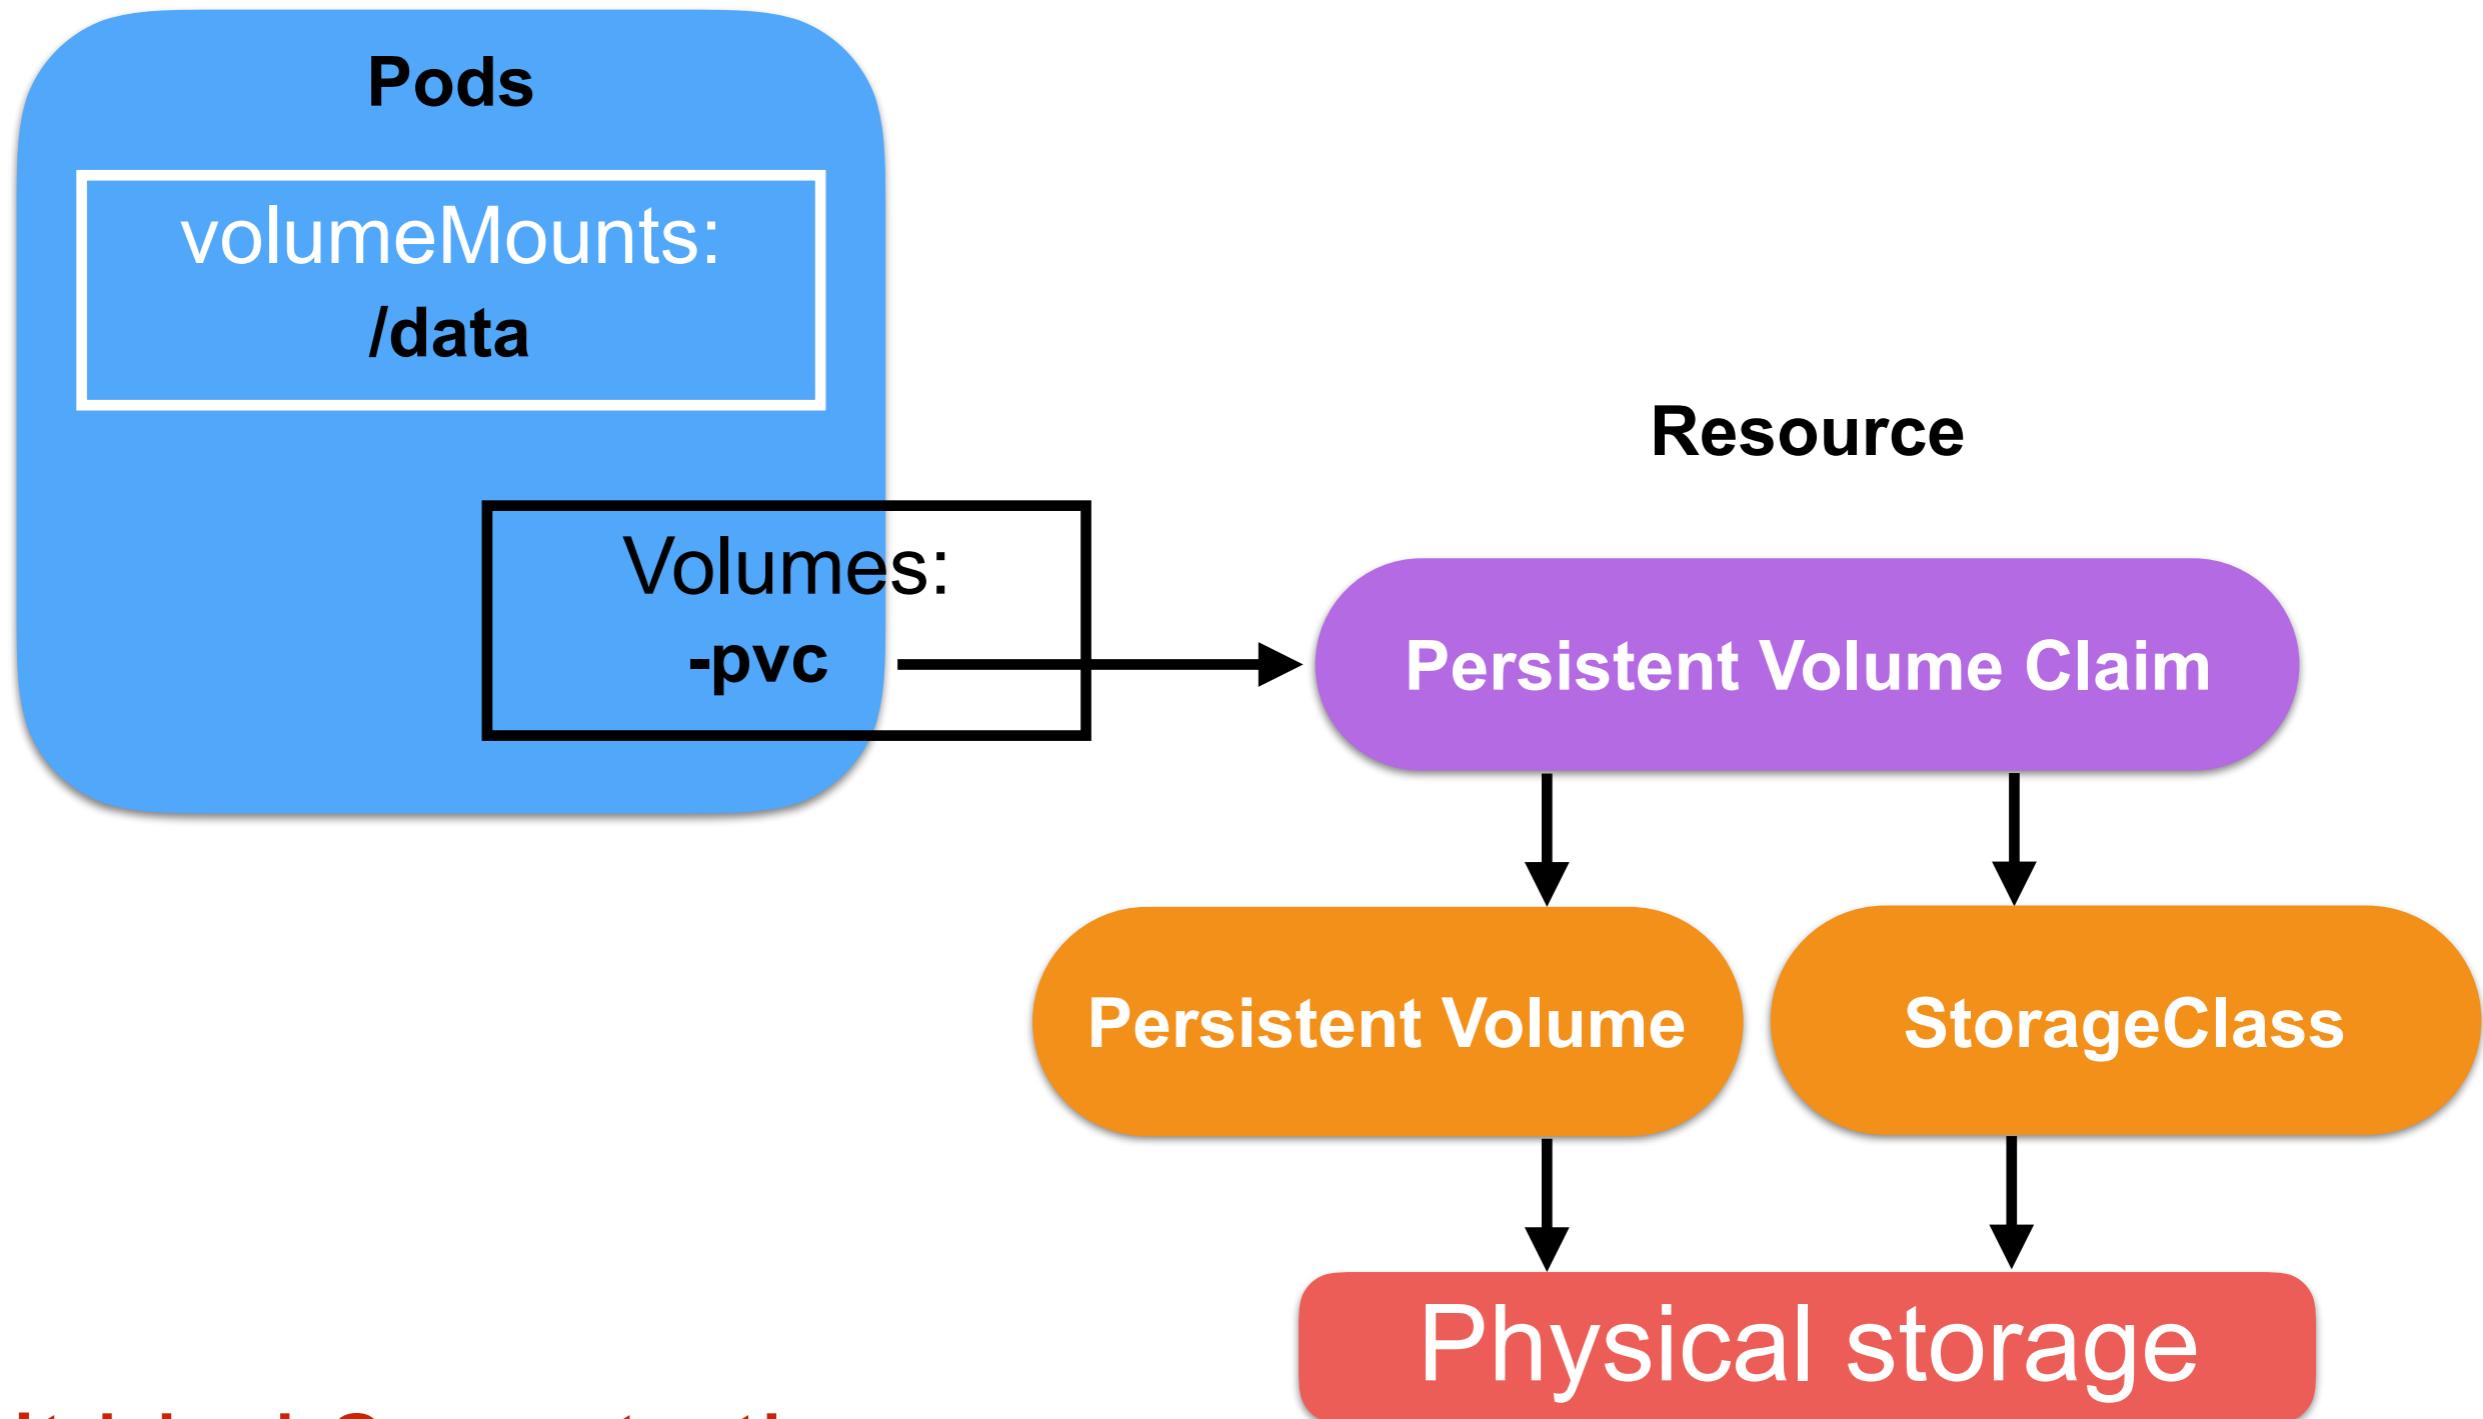
# Volume on Kubernetes



# Volume on Kubernetes



# Volume on Kubernetes



Can't bind 2 pvc to the same pv.



# Persistent Volume

Provide and APIs for user and administrator  
Abstraction layer to hide detail of storage

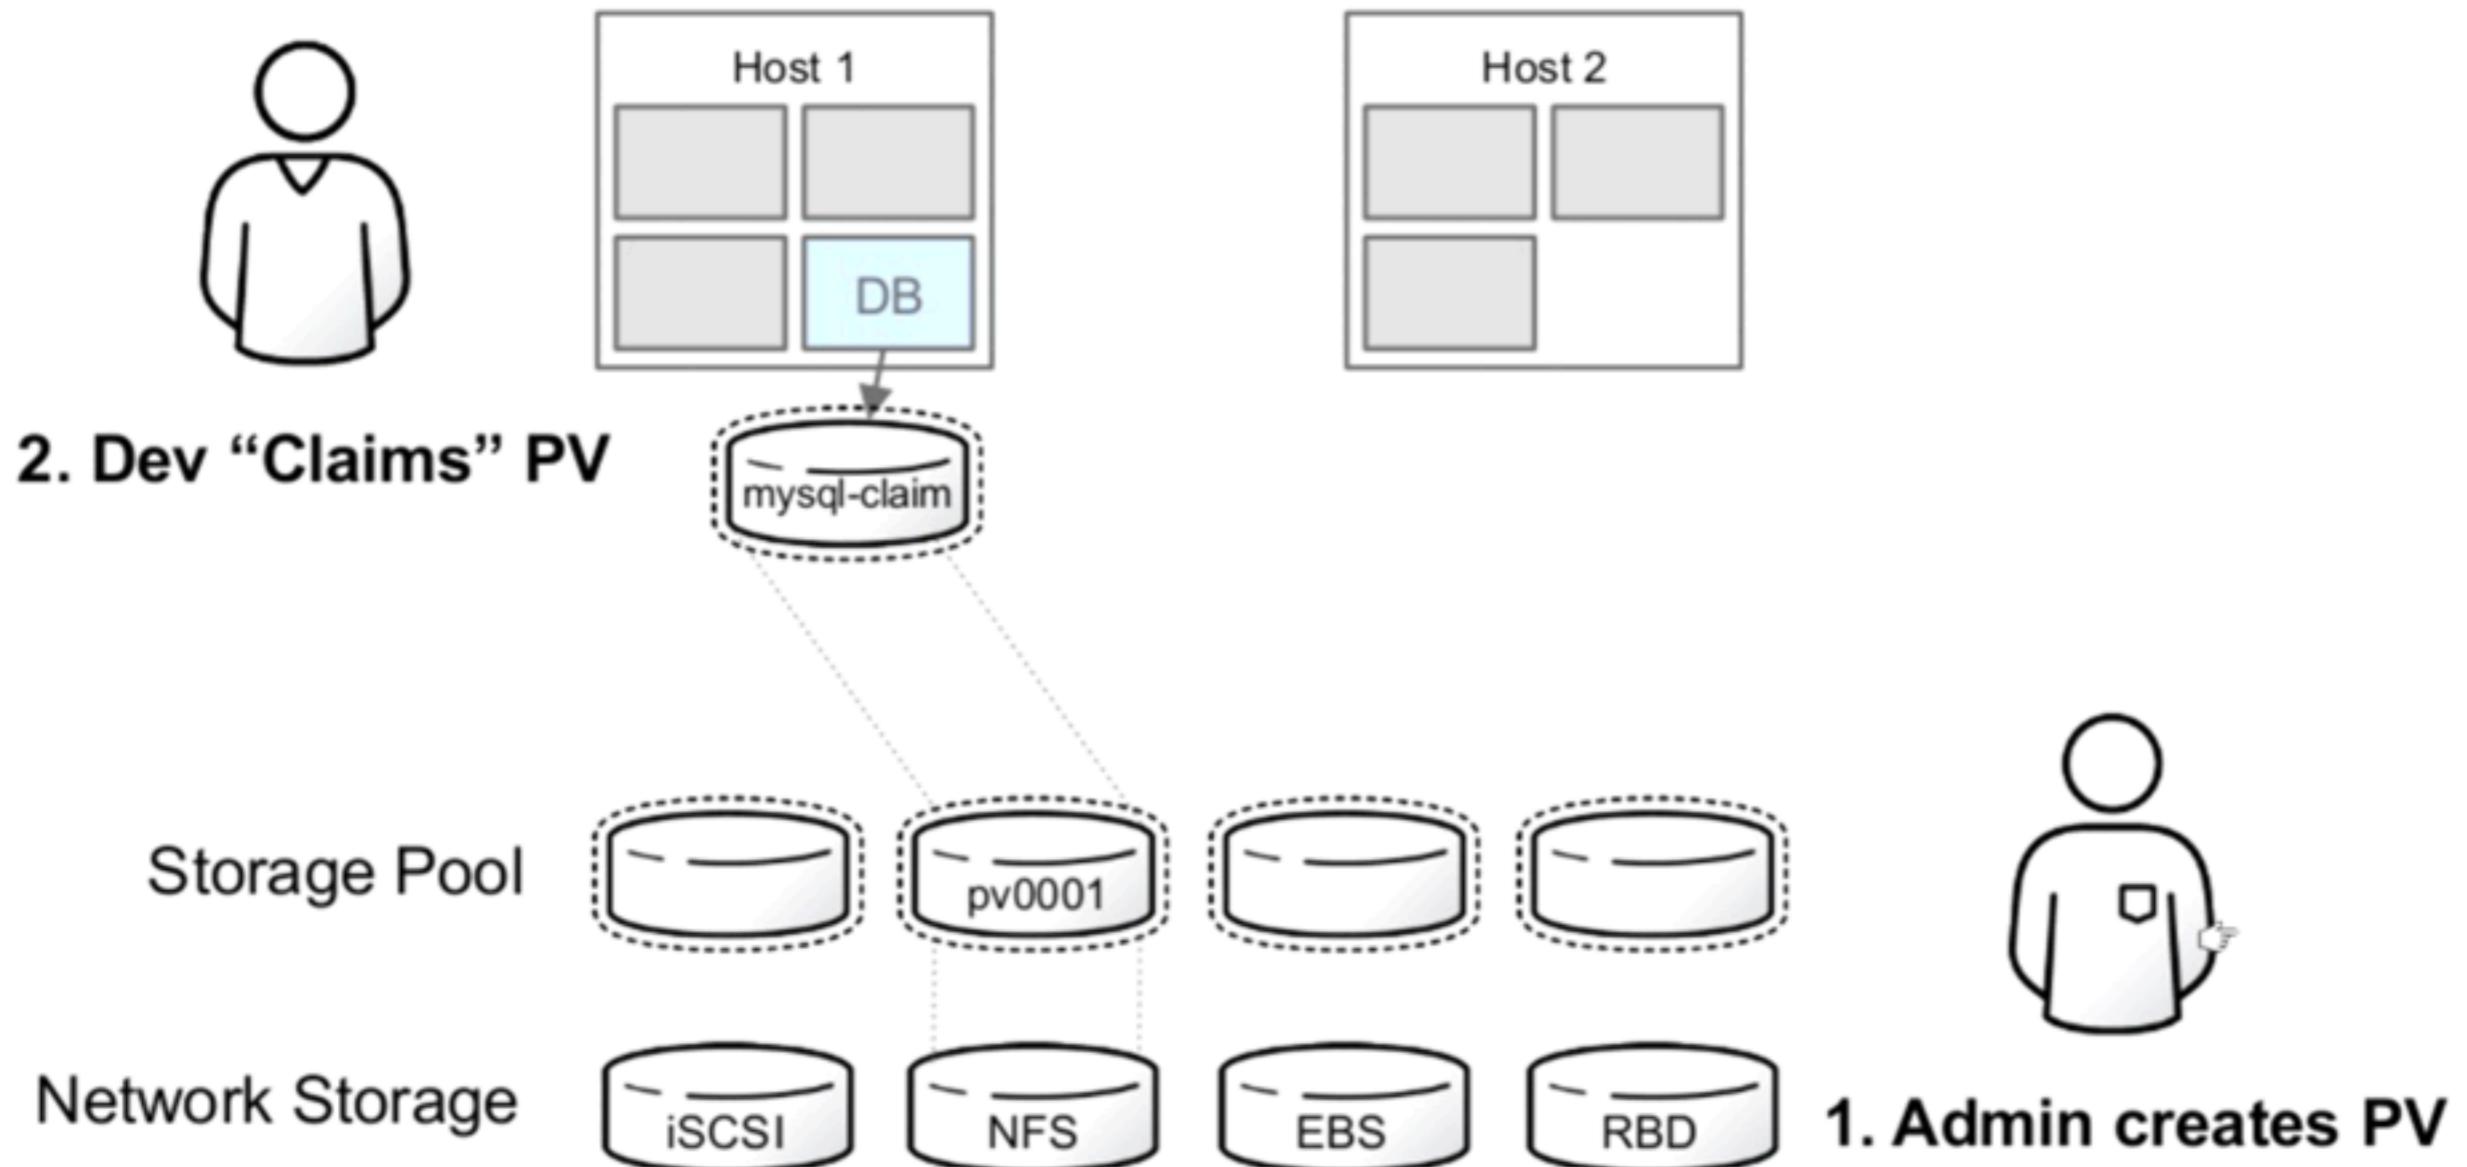
Provide 2 apis for manage your storage

1. PersistentVolume (PV)
2. PersistentVolumeClaim (PVC)

<https://kubernetes.io/docs/concepts/storage/persistent-volumes/>



# Persistent Volume



# PersistentVolume (PV)

PV is a piece of network storage in the cluster  
Piece of storage that independence from Pods  
Long lifecycle independent of any Pods  
Provisioning by administrator



# PersistentVolume (PV)

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv0003
spec:
  capacity:
    storage: 5Gi
  volumeMode: Filesystem
  accessModes:
    - ReadWriteOnce
  persistentVolumeReclaimPolicy: Recycle
  storageClassName: slow
  mountOptions:
    - hard
    - nfsvers=4.1
  nfs:
    path: /tmp
    server: 172.17.0.2
```



# PersistentVolumeClaim (PVC)

Request for resources/storage by user  
Allows specific resource requests eg. size,  
access modes



# PersistentVolumeClaim (PVC)

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: myclaim
spec:
  accessModes:
    - ReadWriteOnce
  volumeMode: Filesystem
  resources:
    requests:
      storage: 8Gi
  storageClassName: slow
  selector:
    matchLabels:
      release: "stable"
    matchExpressions:
      - {key: environment, operator: In, values: [dev]}
```



# Access modes

## **ReadWriteOnce (RWO)**

the volume can be mounted as read-write by a single node

## **ReadOnlyMany (ROX)**

the volume can be mounted read-only by many nodes

## **ReadWriteMany (RWX)**

the volume can be mounted as read-write by many nodes



# Access modes

Volume Plugin	ReadWriteOnce	ReadOnlyMany	ReadWriteMany
AWSElasticBlockStore	✓	-	-
AzureFile	✓	✓	✓
AzureDisk	✓	-	-
CephFS	✓	✓	✓
Cinder	✓	-	-
FC	✓	✓	-
FlexVolume	✓	✓	-
Flocker	✓	-	-
GCEPersistentDisk	✓	✓	-
Glusterfs	✓	✓	✓
HostPath	✓	-	-
iSCSI	✓	✓	-
Quobyte	✓	✓	✓
NFS	✓	✓	✓



# Lifecycle of volume and claim



# Lifecycle of volume

1. Provisioning
2. Binding
3. Using
4. Reclaim (retain, recycle, delete)



# Phase of volume

## 1. Available

Free resource that is not yet bound to a claim

## 2. Bound

Volume is bound to a claim

## 3. Released

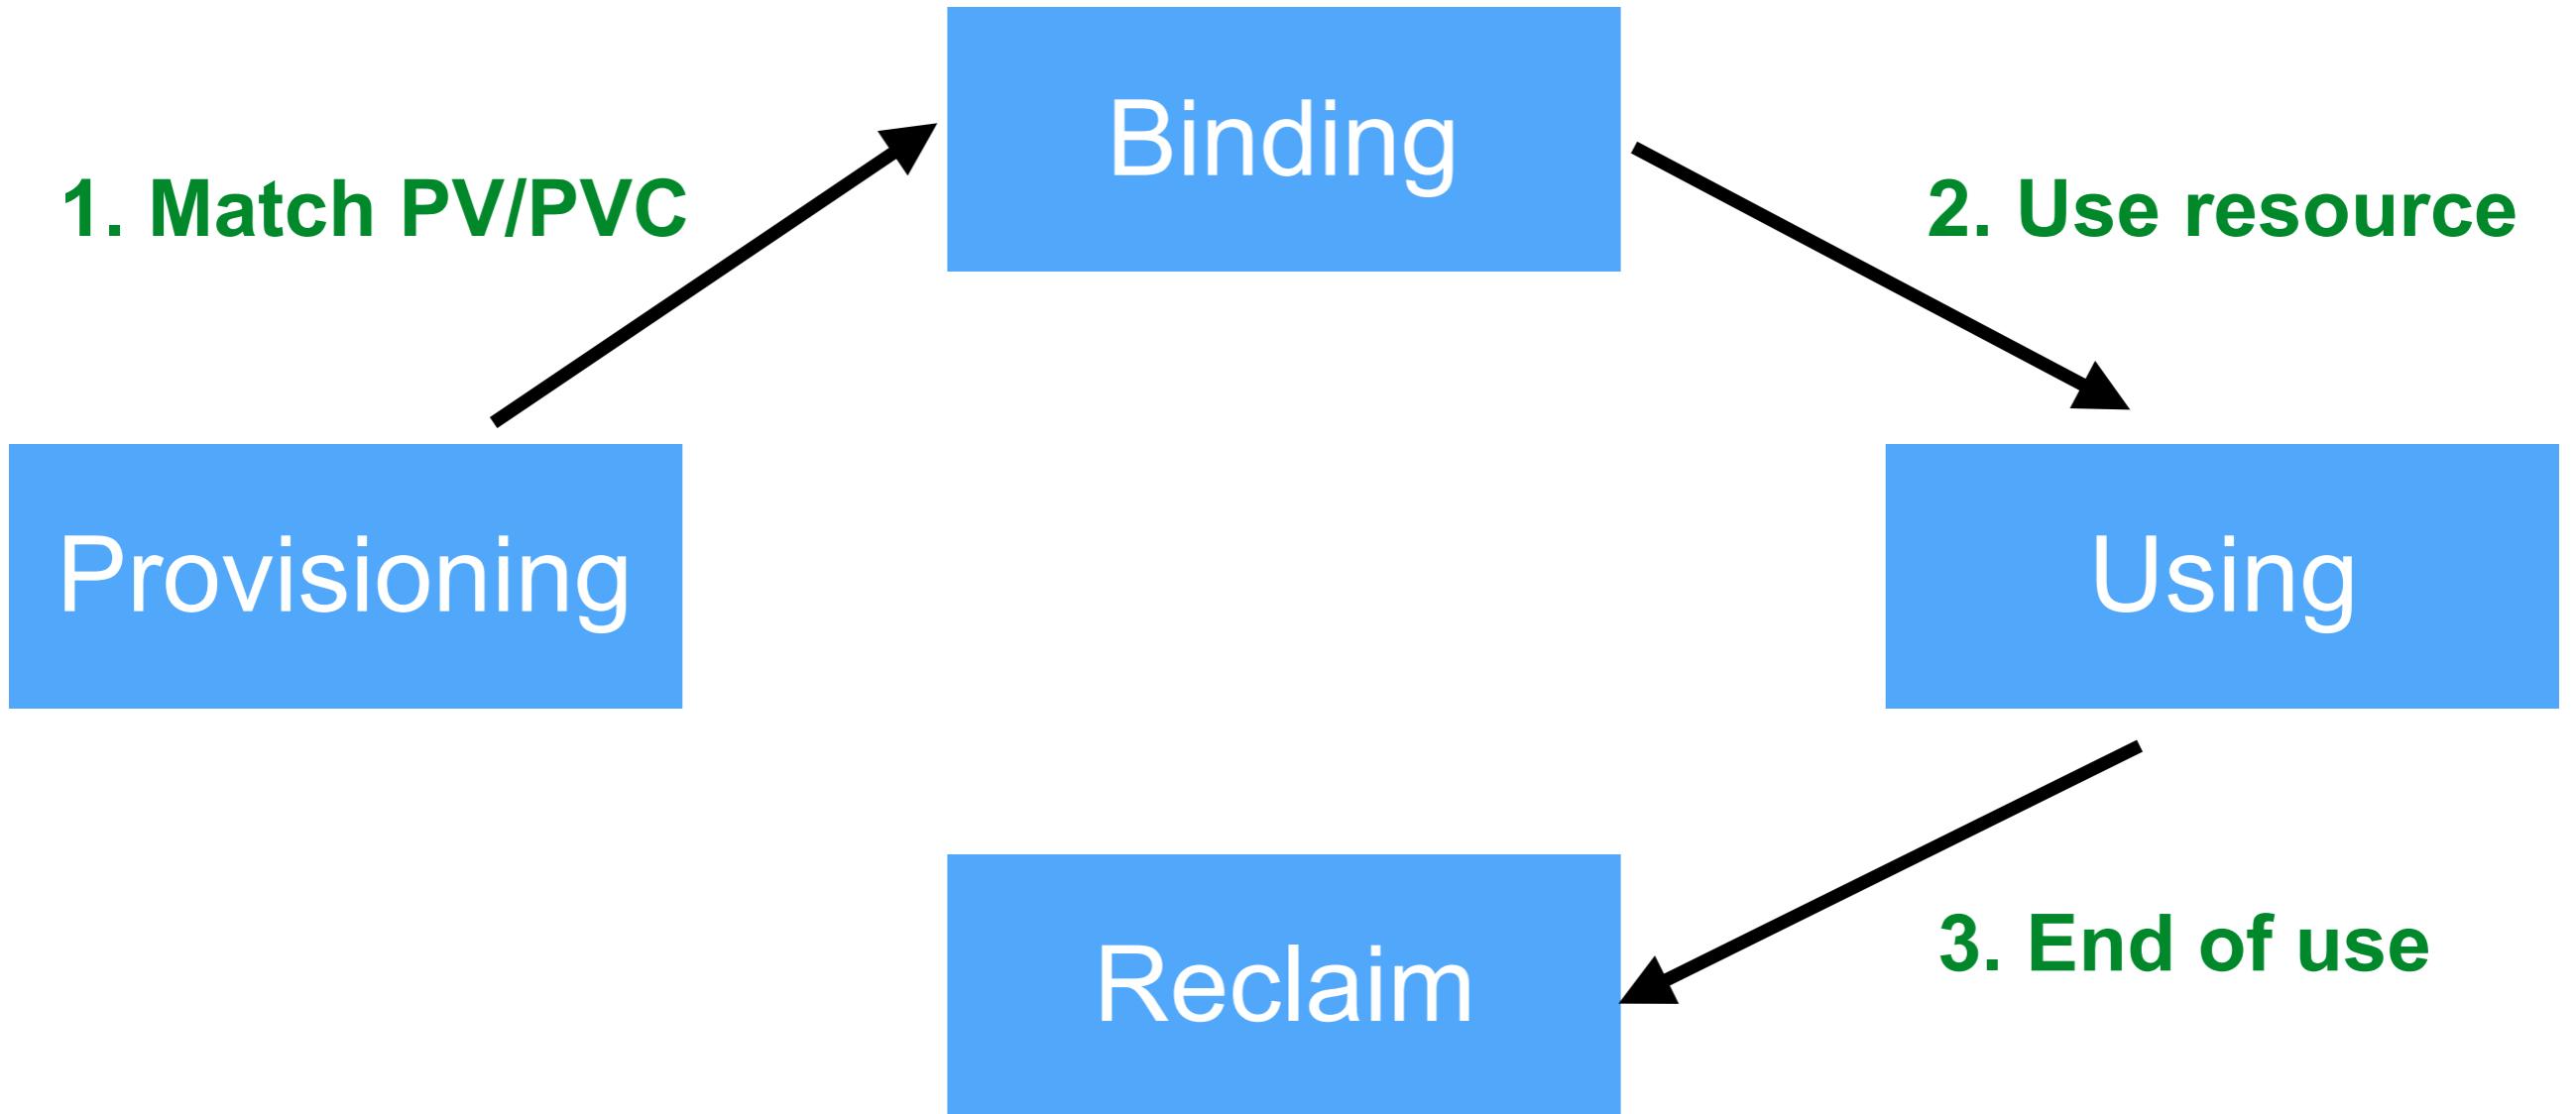
The claim has been deleted, but the resource is not yet reclaimed by the cluster

## 4. Failed

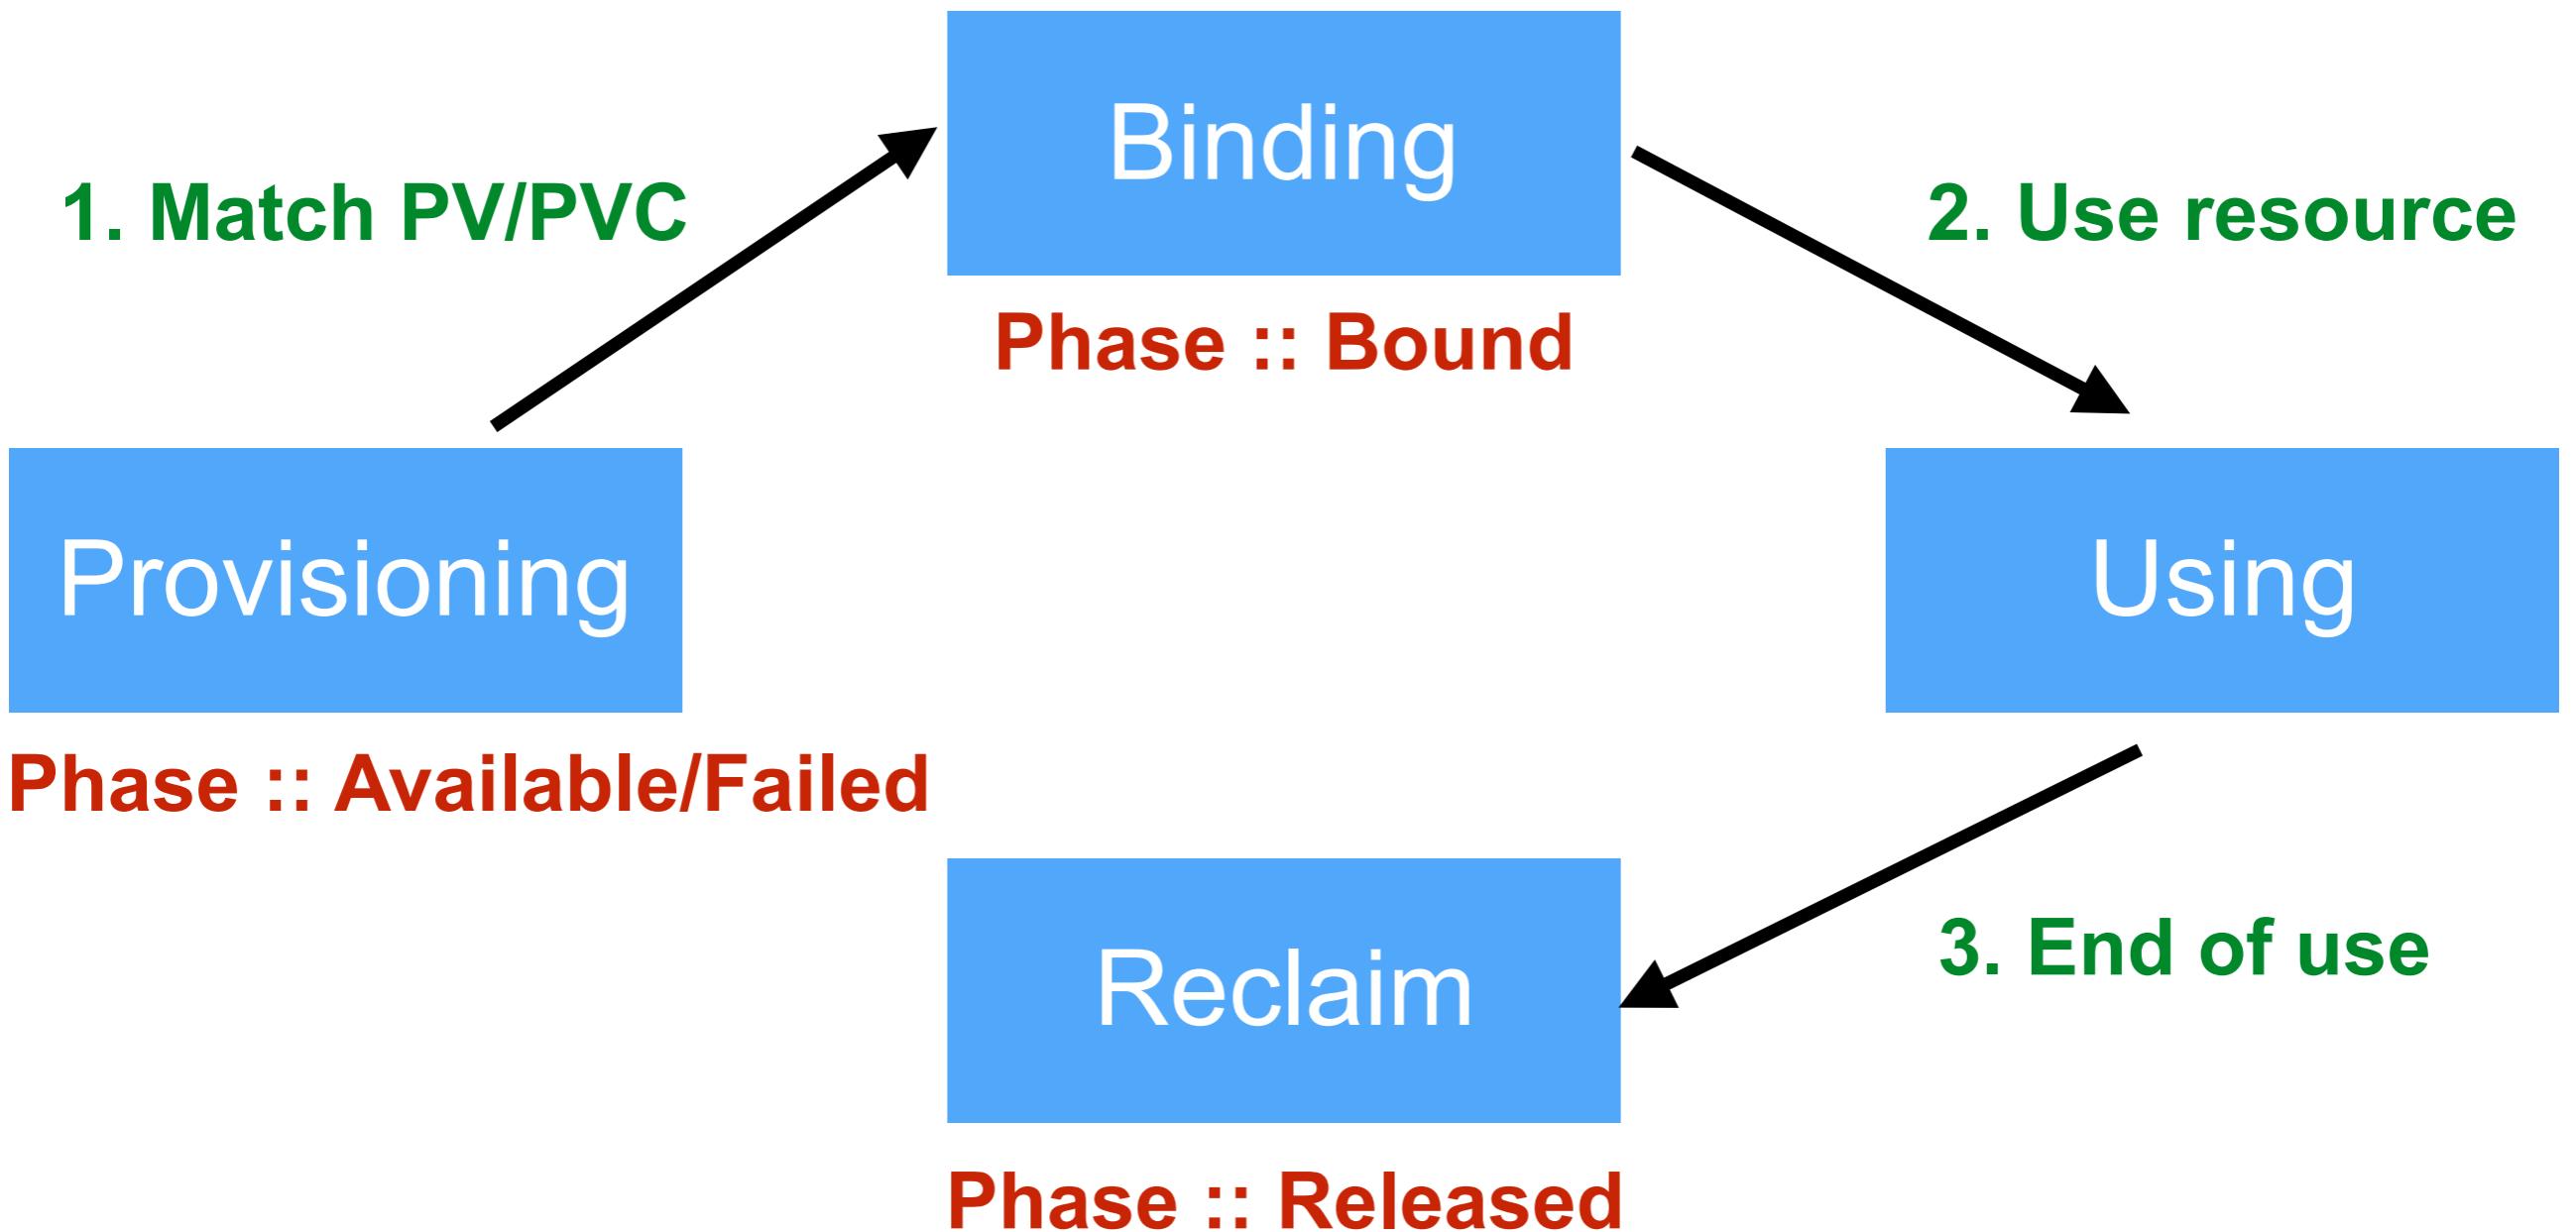
Volume has failed its automatic reclamation



# Lifecycle of volume



# Phase of volume



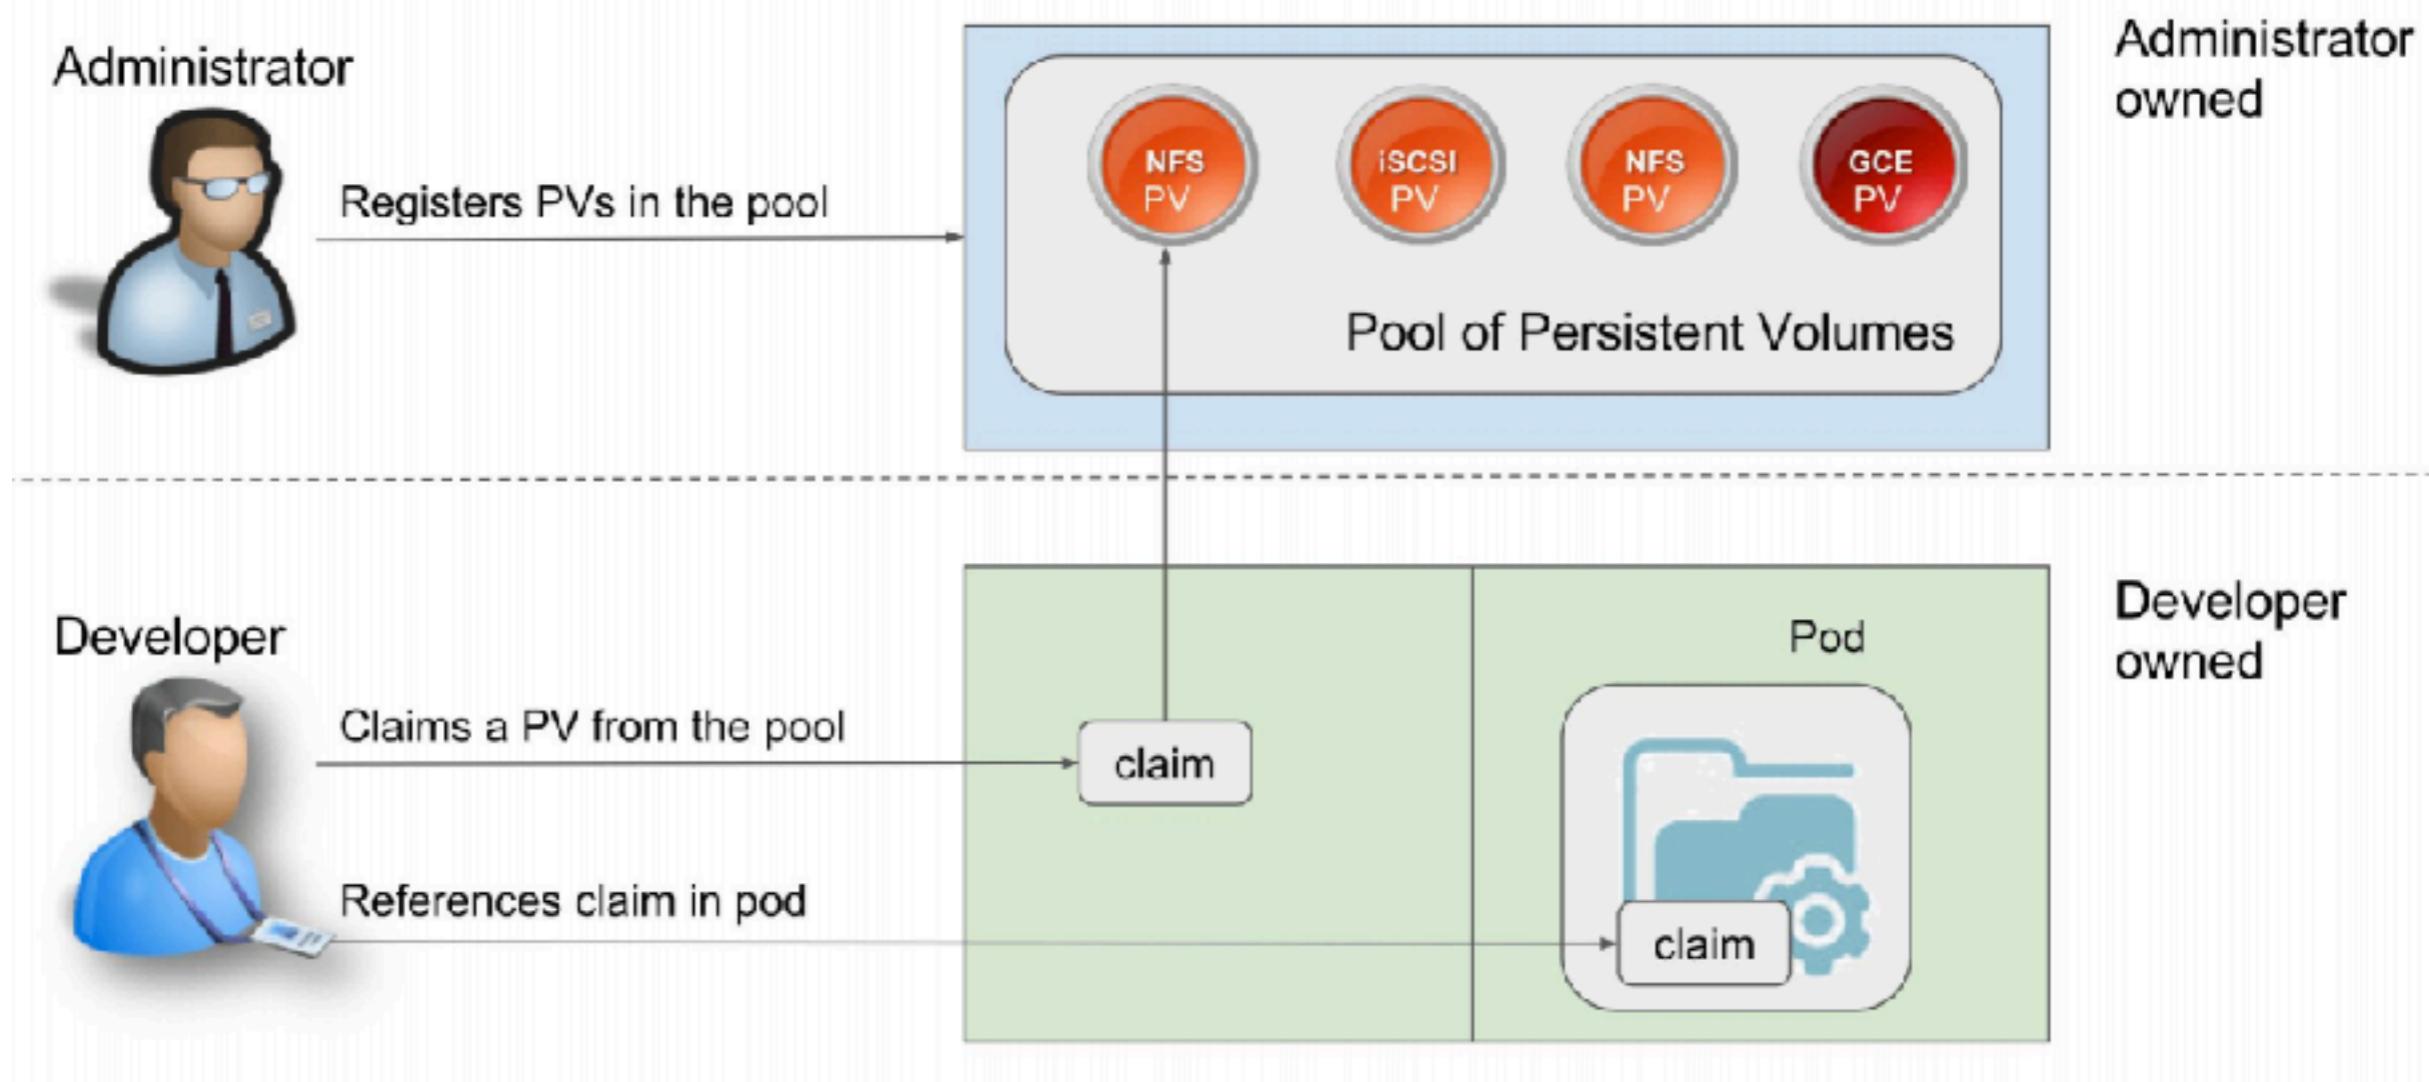
# Provisioning

Static  
Dynamic



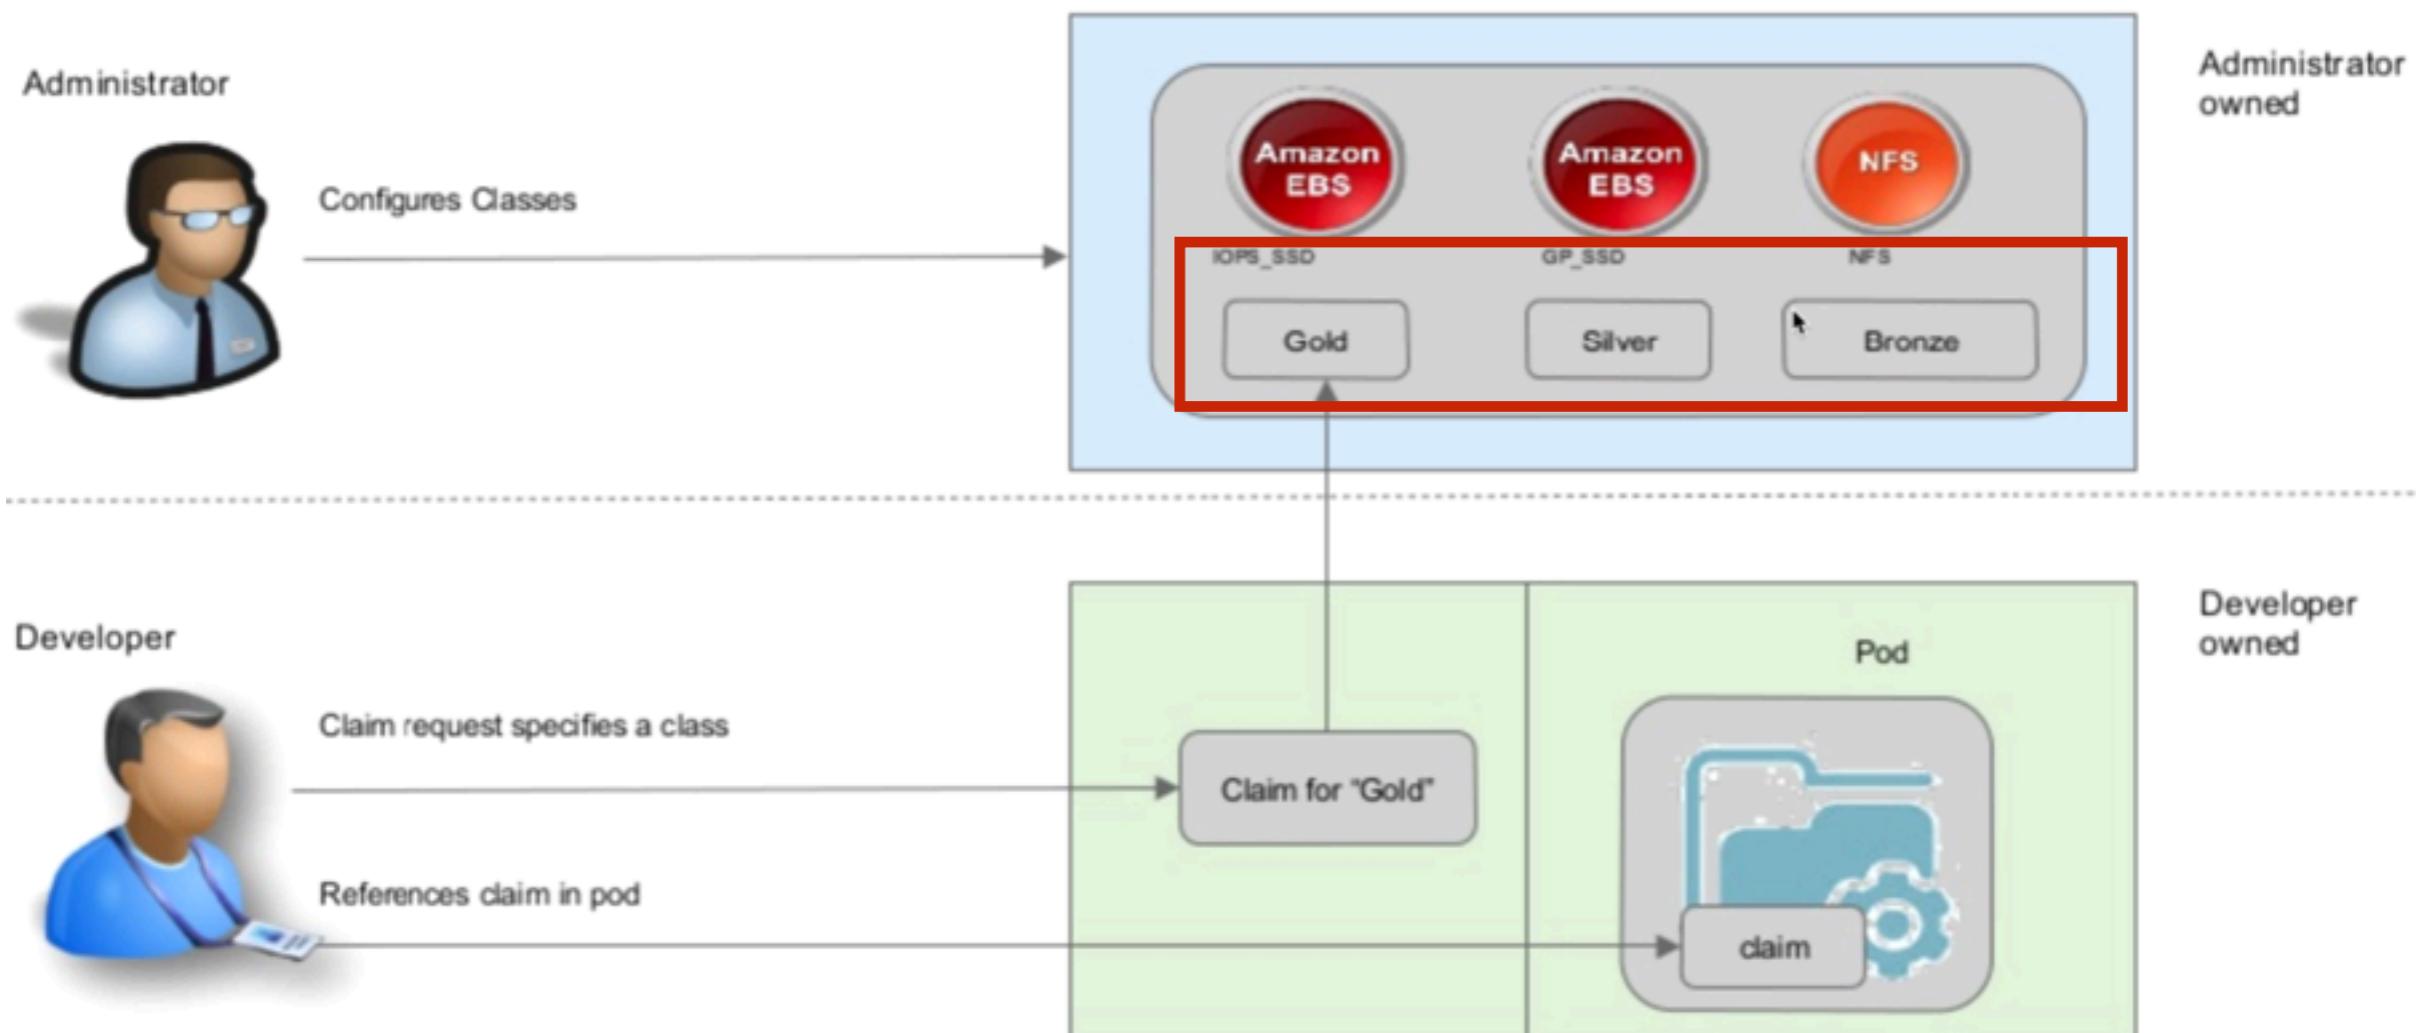
# Static

Administrator create all PV  
User need to know the detail of storage in each PV



# Dynamic

Provisioning based on **StorageClass**  
Create and configure by administrator



# Using PV/PVC from Pods

```
kind: Pod
apiVersion: v1
metadata:
  name: mypod
spec:
  containers:
    - name: myfrontend
      image: dockerfile/nginx
      volumeMounts:
        - mountPath: "/var/www/html"
          name: mypd
  volumes:
    - name: mypd
      persistentVolumeClaim:
        claimName: myclaim
```



# Workshop with Volume

File /working-with-volume/nfs

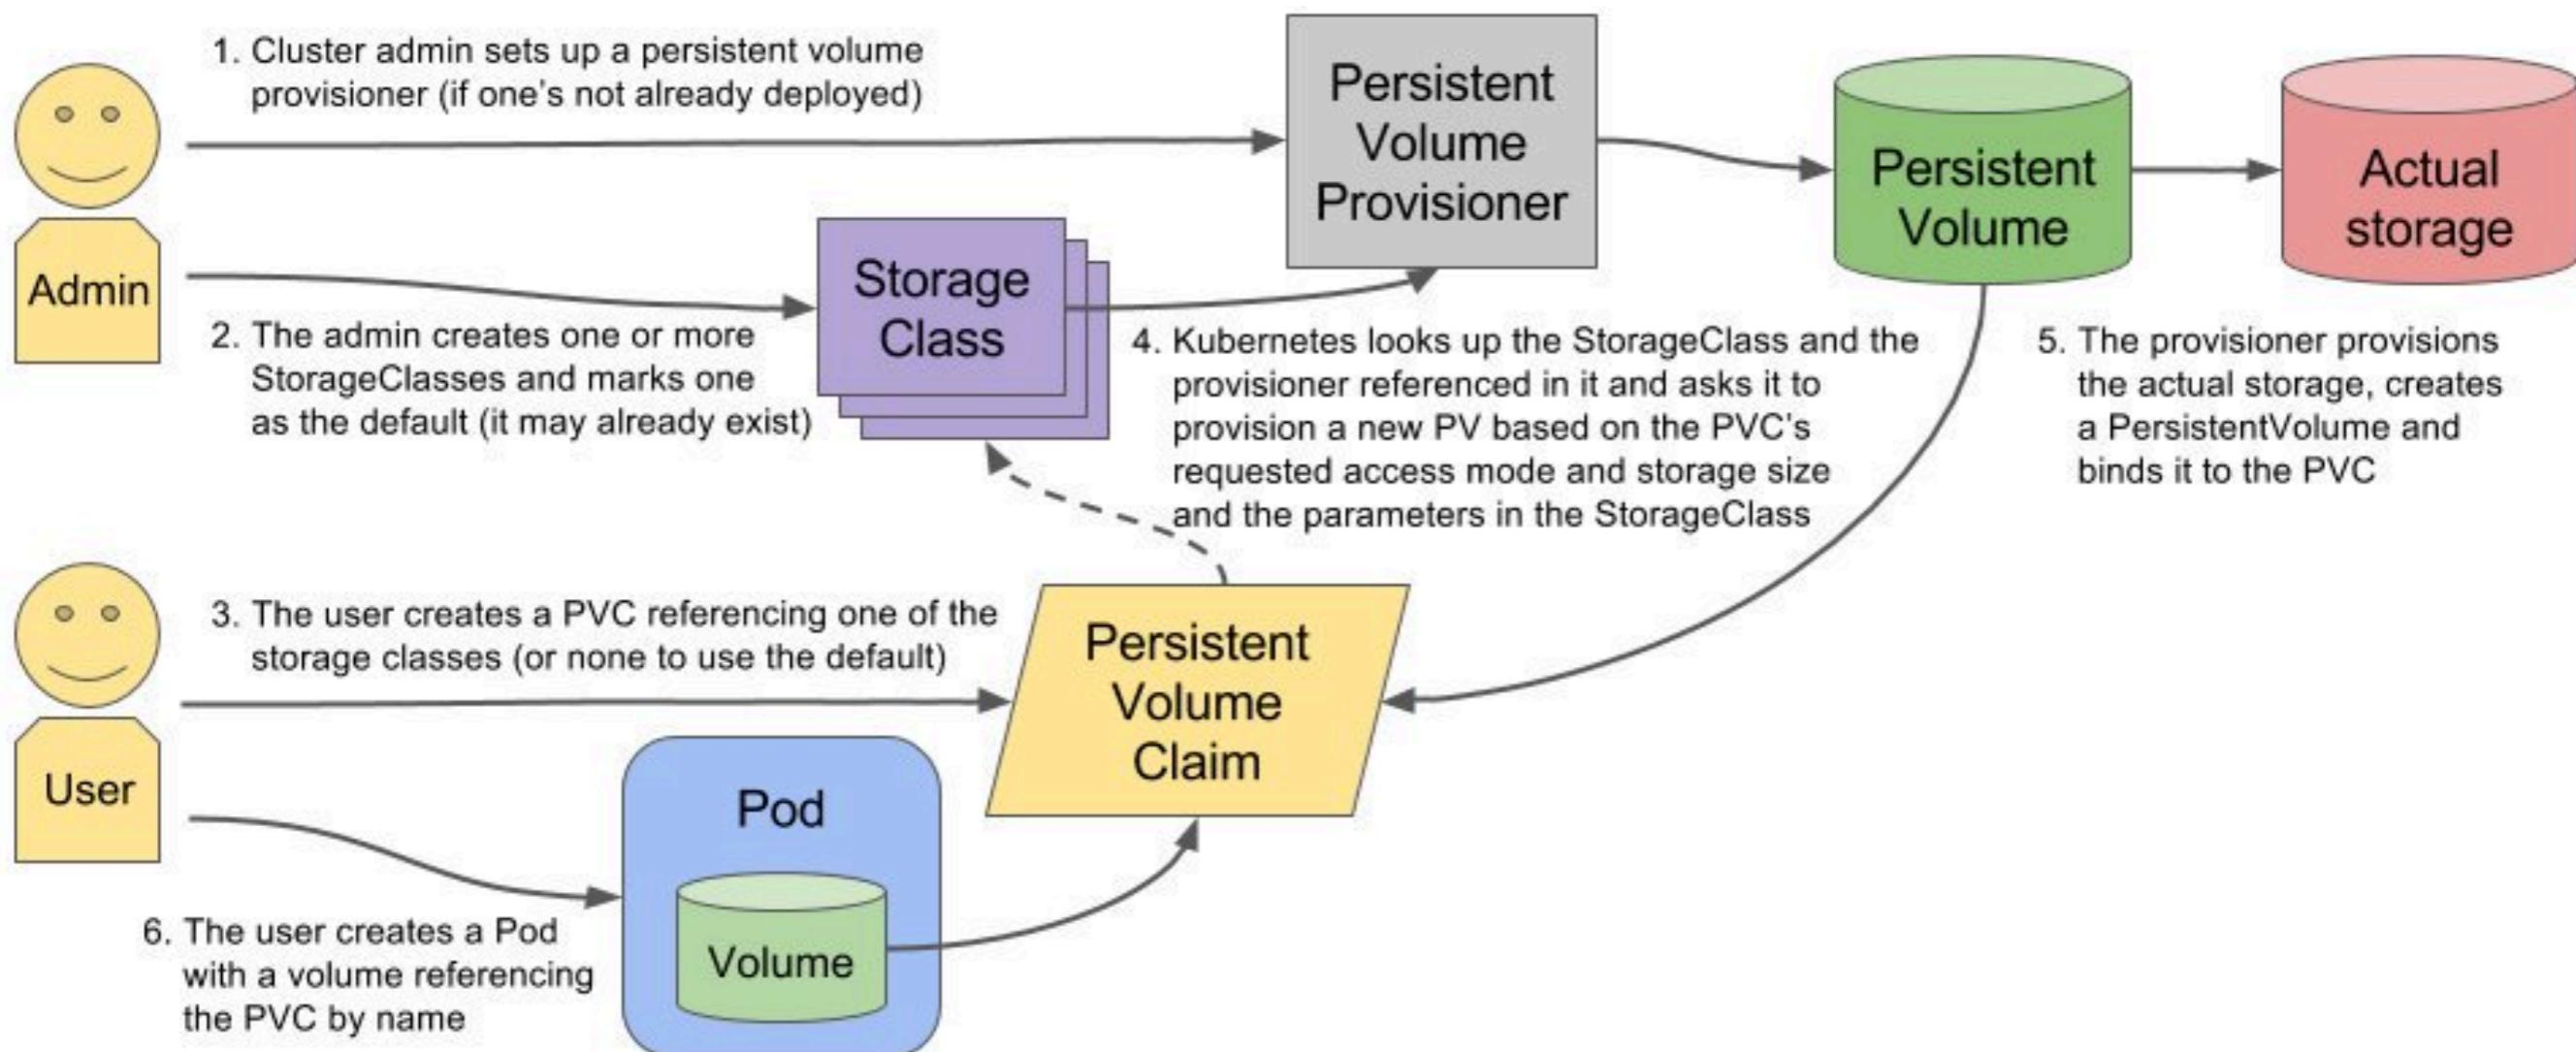


# Step to working with PV/PVC

1. Create PV
2. Create PVC
3. Create Deployment
4. Testing



# Step to working with PV/PVC



# Setup storage engine

NFS (Network File System)  
Cloud storage



# Create PV and PVC

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pvc-nfs
  labels:
    name: pvc-nfs
    environment: development
spec:
  accessModes:
    - ReadWriteMany
  storageClassName: ""
resources:
  requests:
    storage: 500Mi
selector:
  matchLabels:
    name: pv-nfs
    environment: development
```

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv-nfs
  labels:
    name: pv-nfs
    environment: development
spec:
  capacity:
    storage: 1Gi
  storageClassName: ""
  accessModes:
    - ReadWriteMany
  nfs:
    server: 10.148.0.2
    path: "/var/nfs/general"
```



# Use PV with PVC from deployment

```
spec:  
  containers:  
    - name: nginx  
      image: nginx:1.8  
      ports:  
        - containerPort: 80  
      volumeMounts:  
        - name: try-nfs  
          mountPath: "/usr/share/nginx/html"  
  volumes:  
    - name: try-nfs  
      persistentVolumeClaim:  
        claimName: pvc-nfs
```

```
apiVersion: v1  
kind: PersistentVolumeClaim  
metadata:  
  name: pvc-nfs  
  labels:  
    name: pvc-nfs  
    environment: development  
spec:  
  accessModes:  
    - ReadWriteMany  
  storageClassName: ""  
  resources:  
    requests:  
      storage: 500Mi  
  selector:  
    matchLabels:  
      name: pv-nfs  
      environment: development
```



# StatefulSet

<https://kubernetes.io/docs/concepts/workloads/controllers/statefulset/>



# StatefulSet

Bringing the concept of ReplicaSets to **stateful** Pods

Enable running Pods in **cluster** mode

Ideal for deploy **highly** available database workload



# StatefulSet for application

Stable, unique network identifiers

Stable, persistent storage

Ordered, graceful deployment and scaling

Ordered, graceful deletion and termination



# Key concepts

Depend on a **Headless service** for Pods

Each Pods get s DNS name accessible to other Pods

Leverage PV and PVC

Each Pods is suffix with a predictable  
(mysql-01, mysql-02)



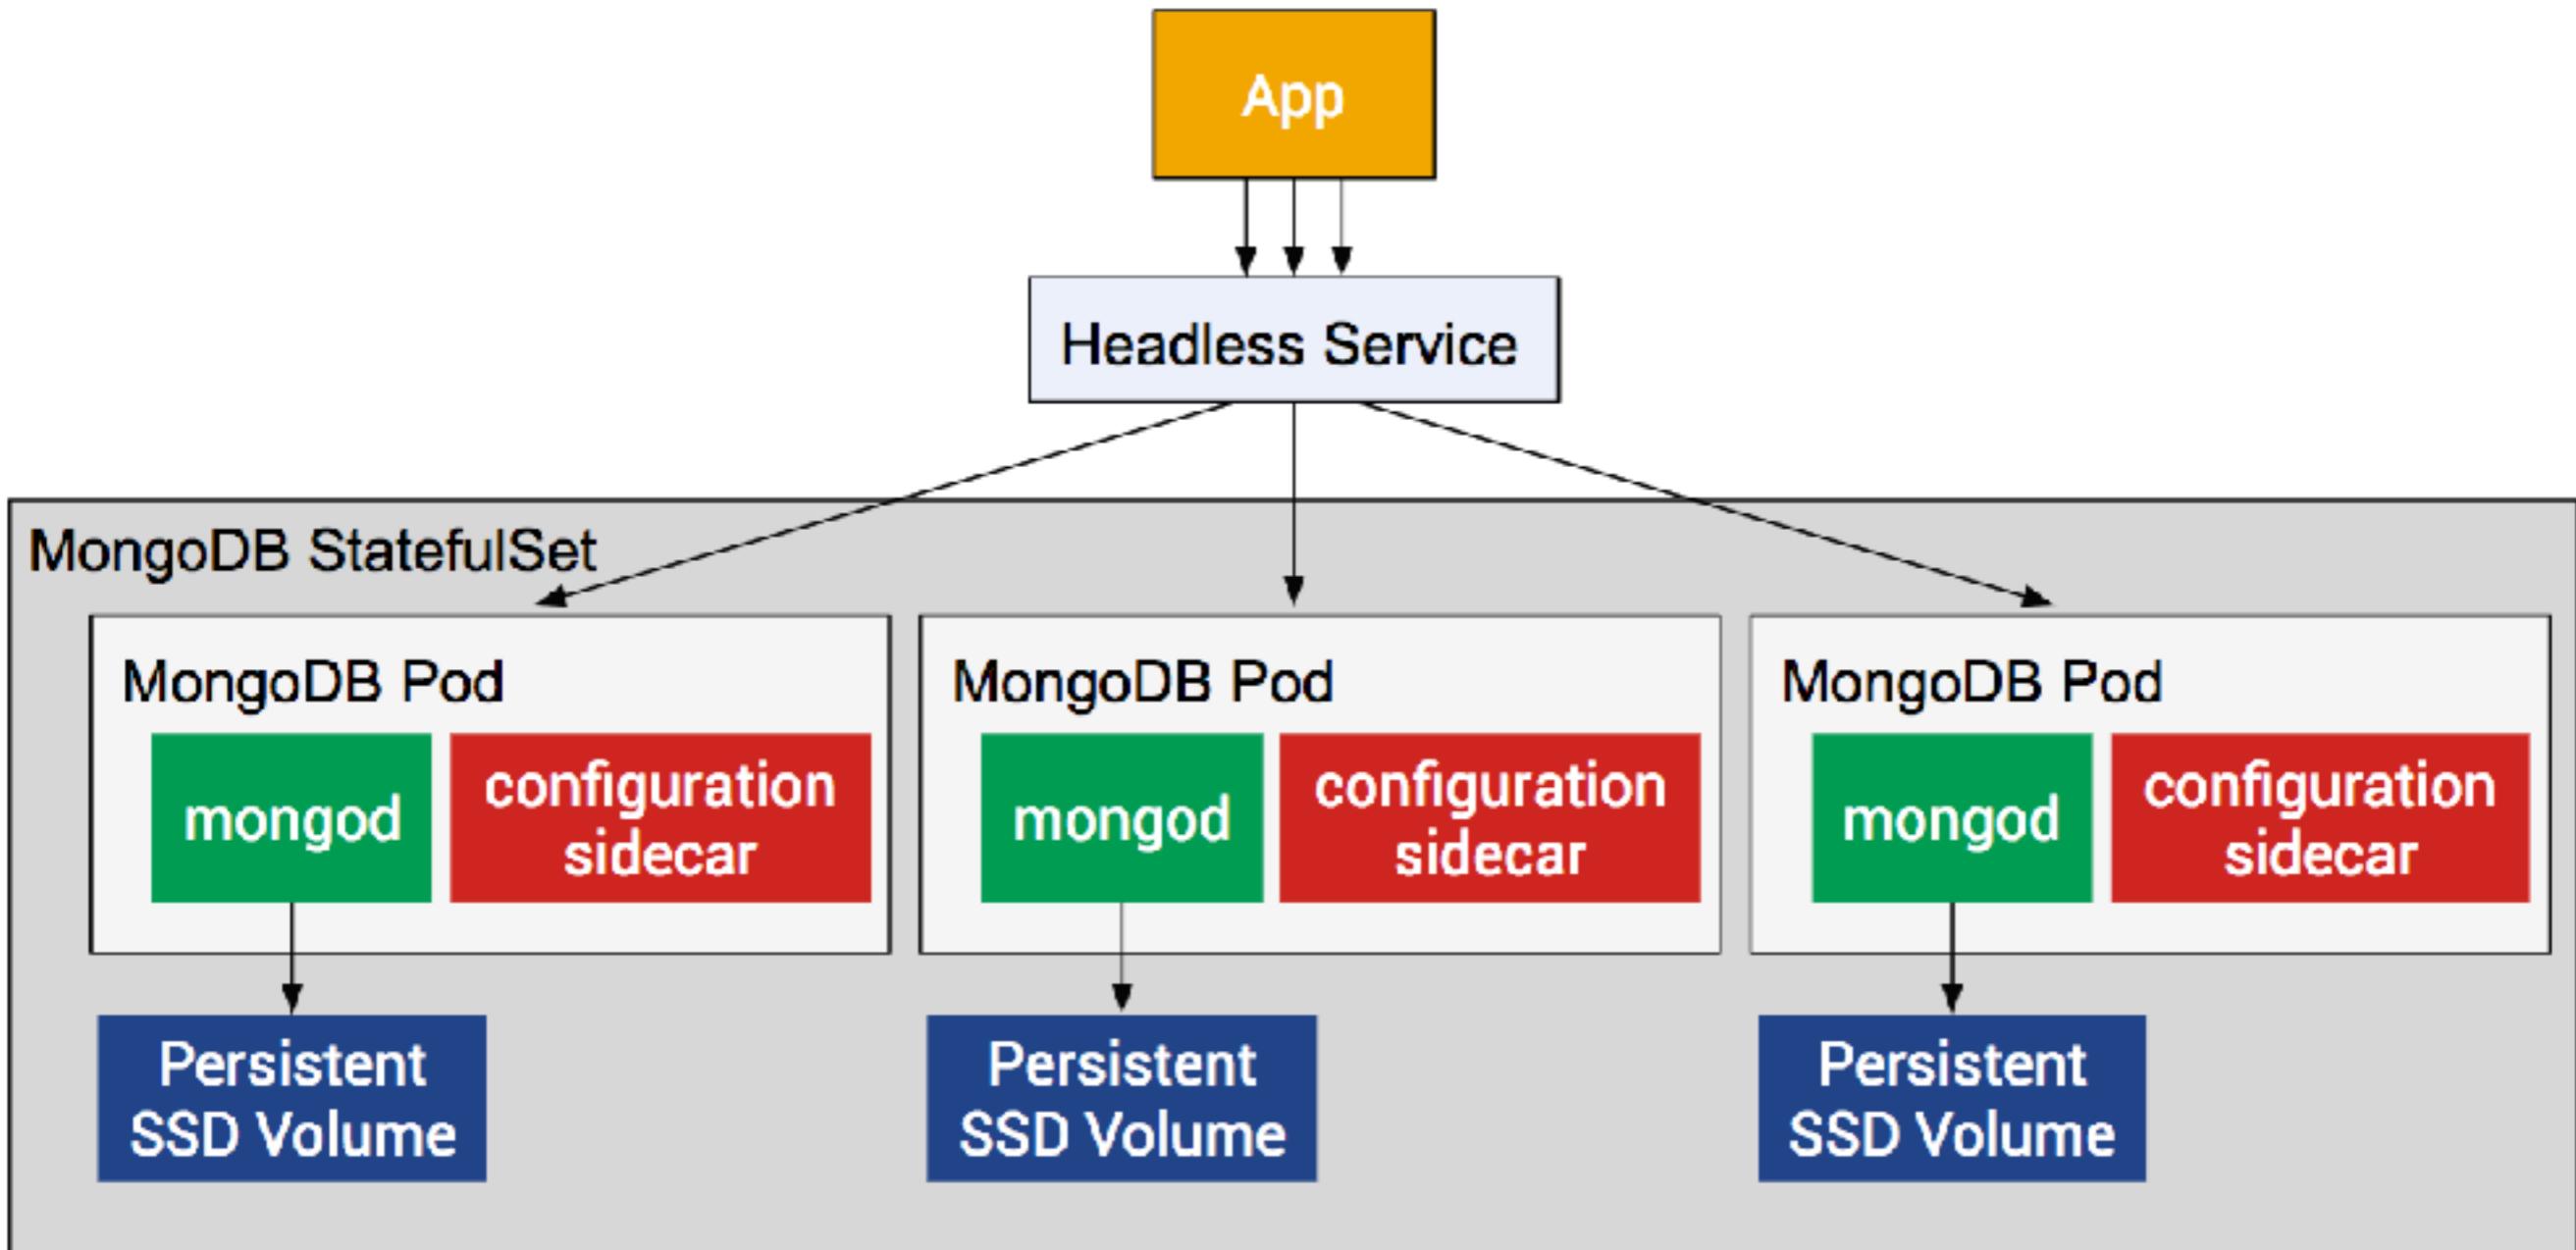
# Key concepts

Pods are created sequentially

Pods are terminated in LiFo (Last in, First out)



# StatefulSet



<https://kubernetes.io/blog/2017/01/running-mongodb-on-kubernetes-with-statefulsets/>



# Key concepts

Pods are created sequentially

Pods are terminated in LiFo (Last in, First out)



# Workshop with StatefulSet

File /working-with-stateful-set



# Application Deployment Strategies



# Strategies to deploy

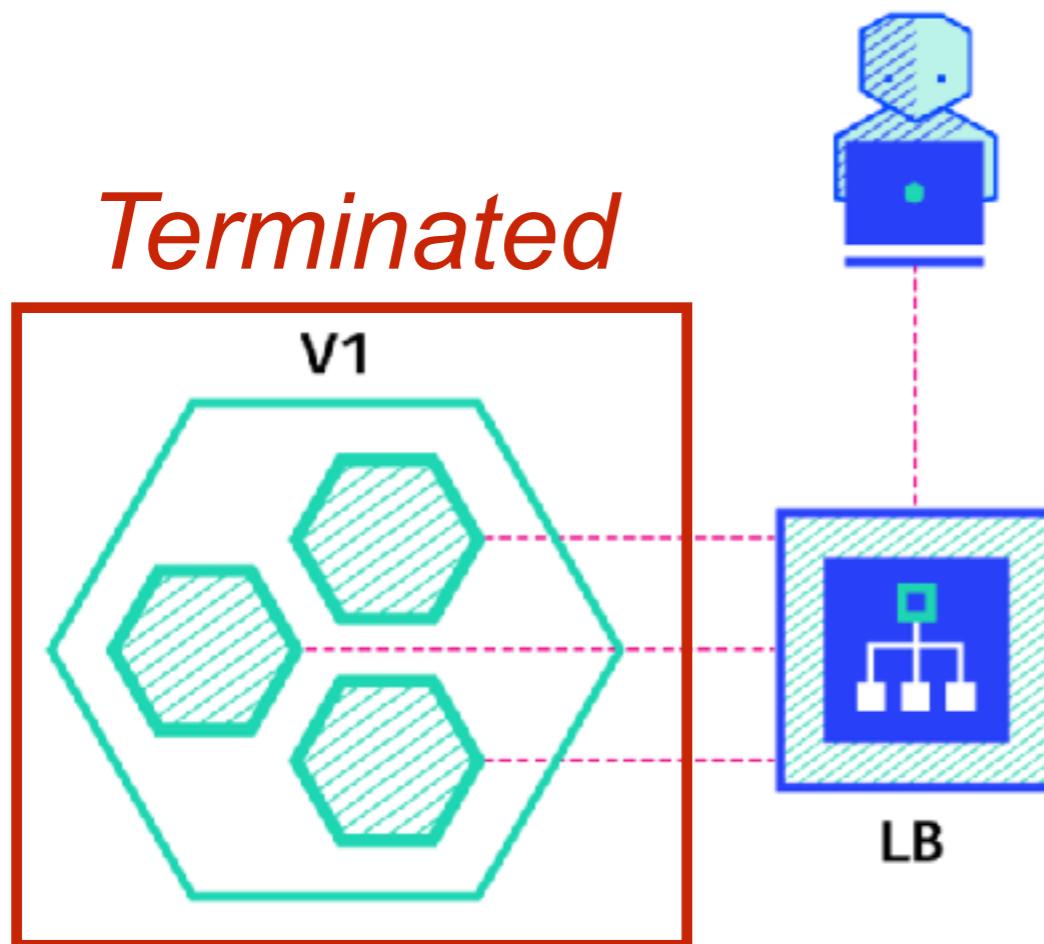
Recreate  
Ramped  
Blue/Green  
Canary  
A/B testing  
Shadow

<https://thenewstack.io/deployment-strategies/>



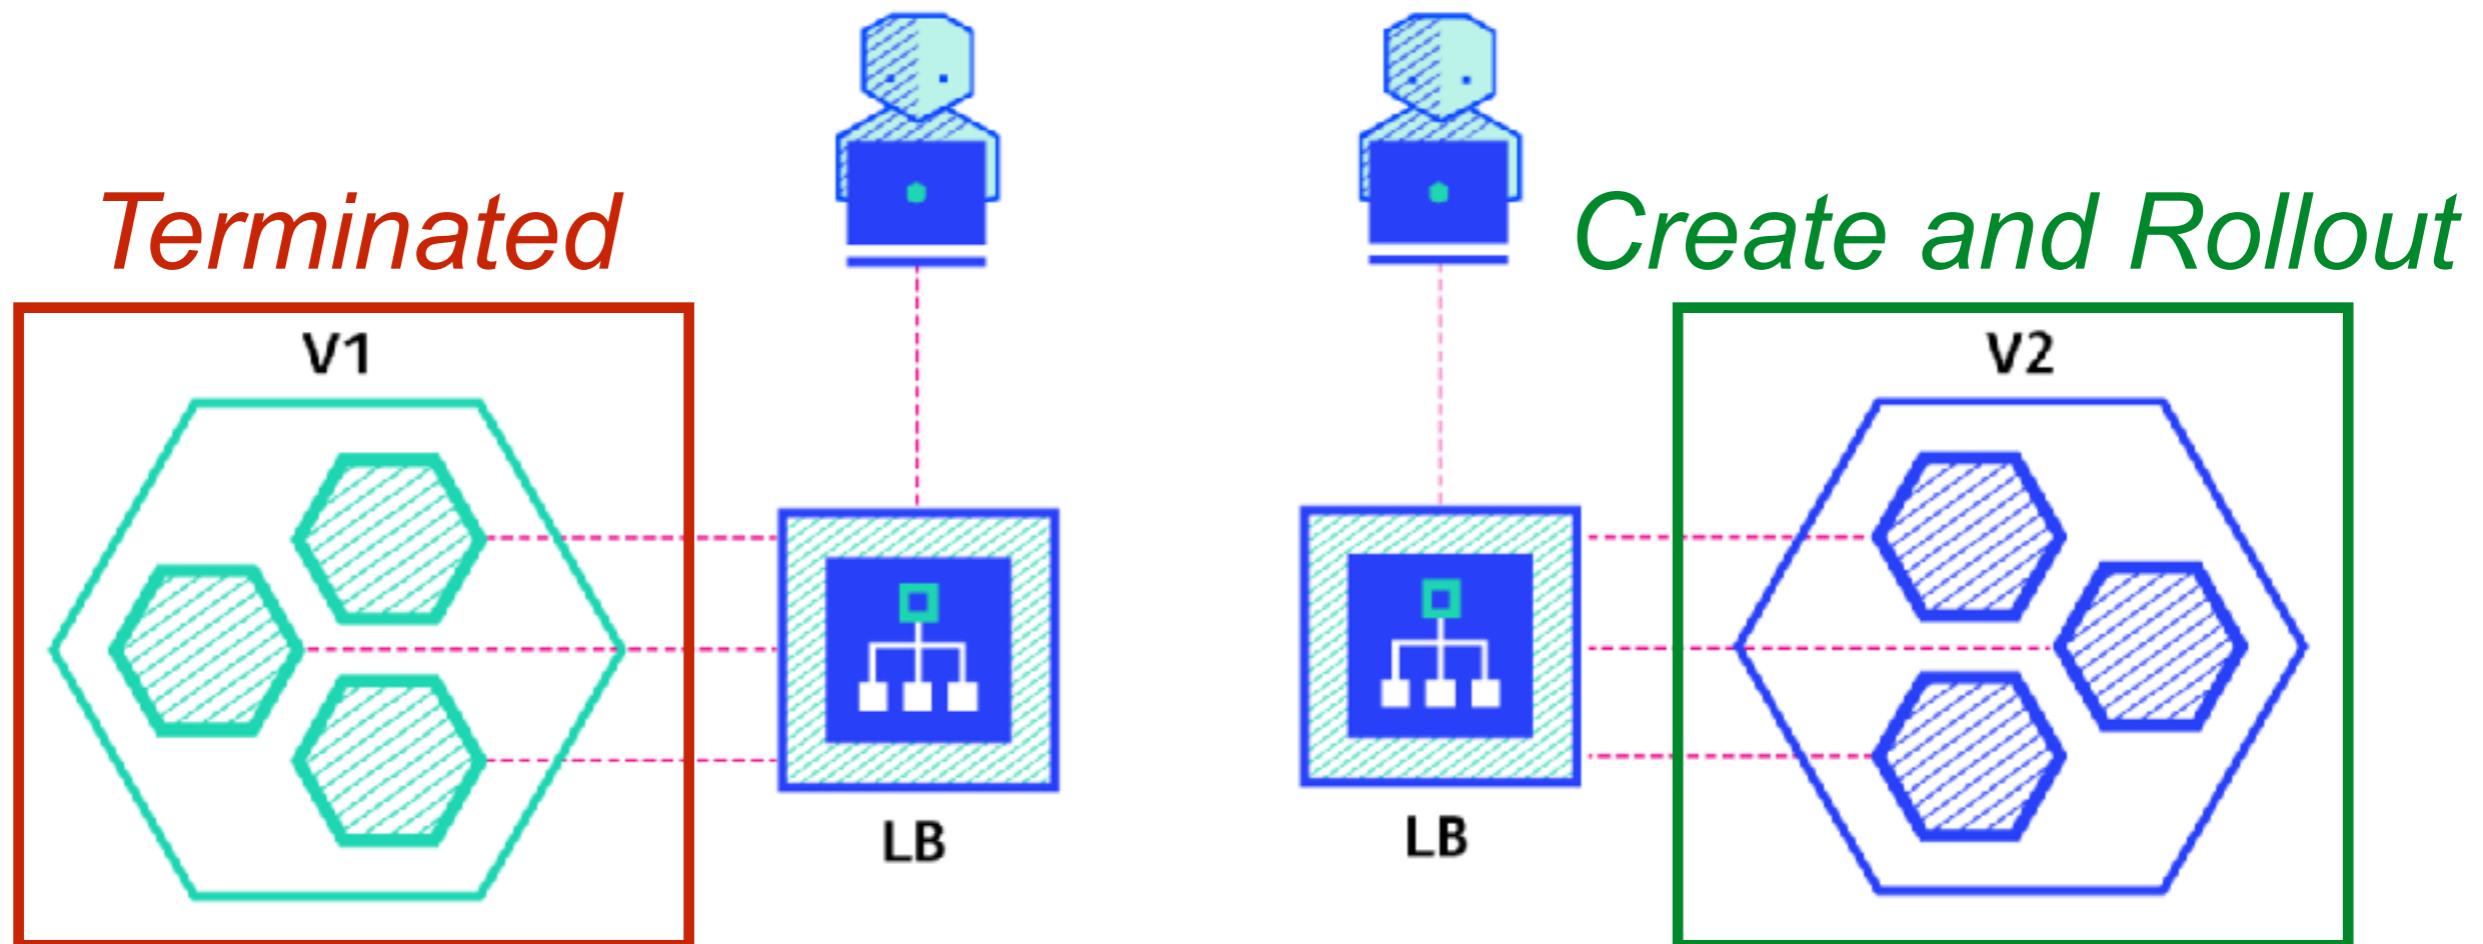
# 1. Recreate

Version A is terminated then version B is rollout



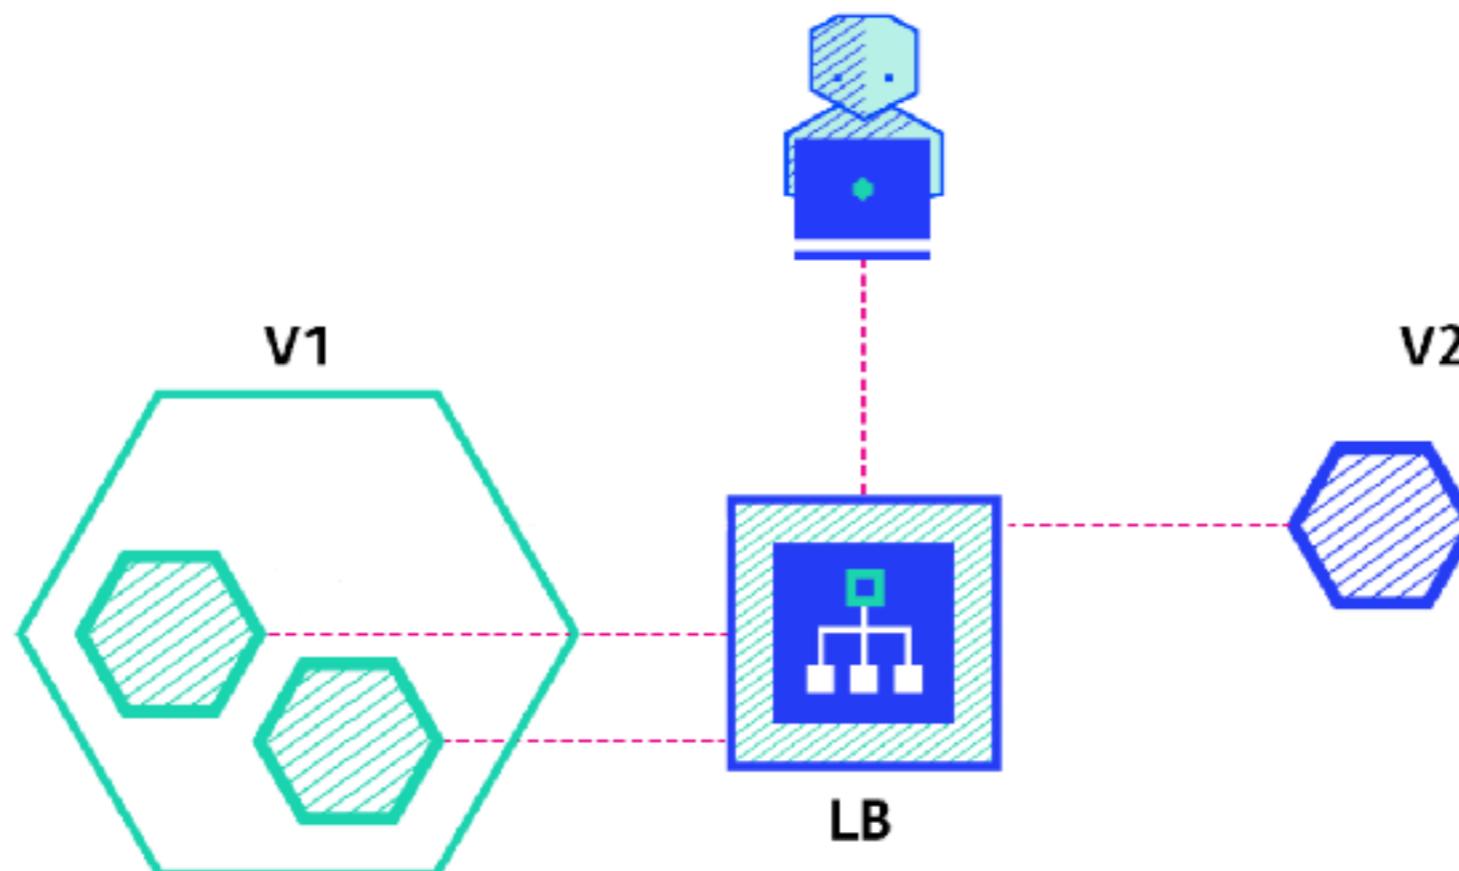
# 1. Recreate

Version A is terminated then version B is rollout



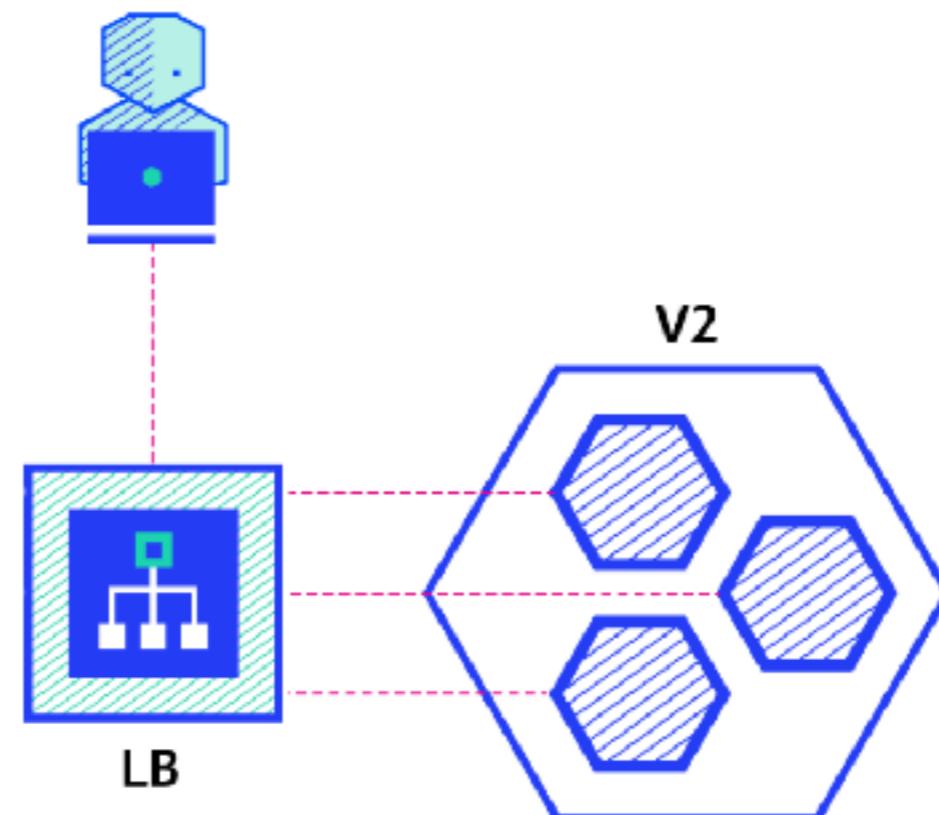
## 2. Ramped

Slow roll out by replace instance one-by-one



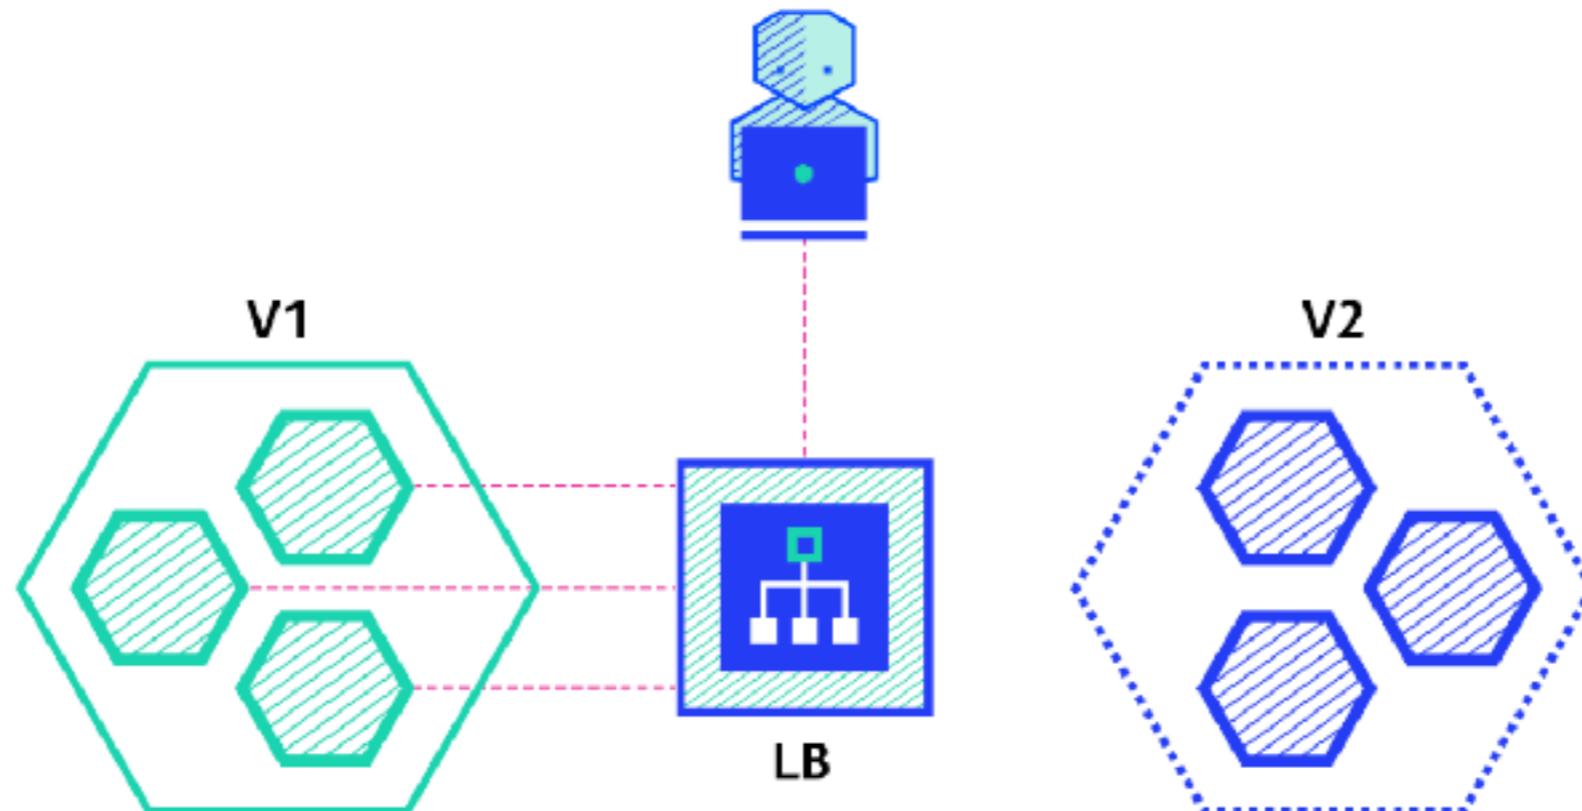
## 2. Ramped

Slow roll out by replace instance one-by-one



# 3. Blue/Green

Current version is called **Blue**  
New version is called **Green**

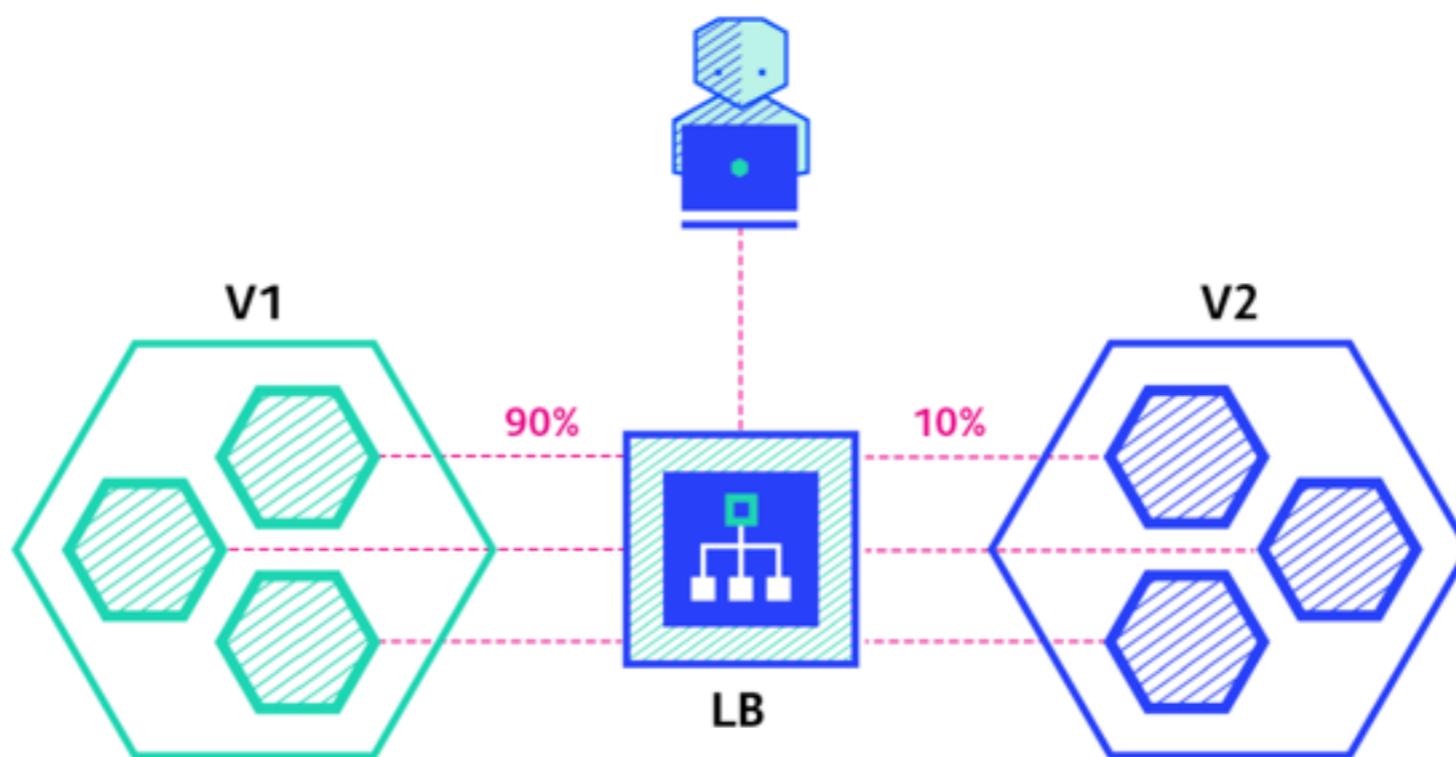


# 4. Canary

Shift production traffic from version A to B

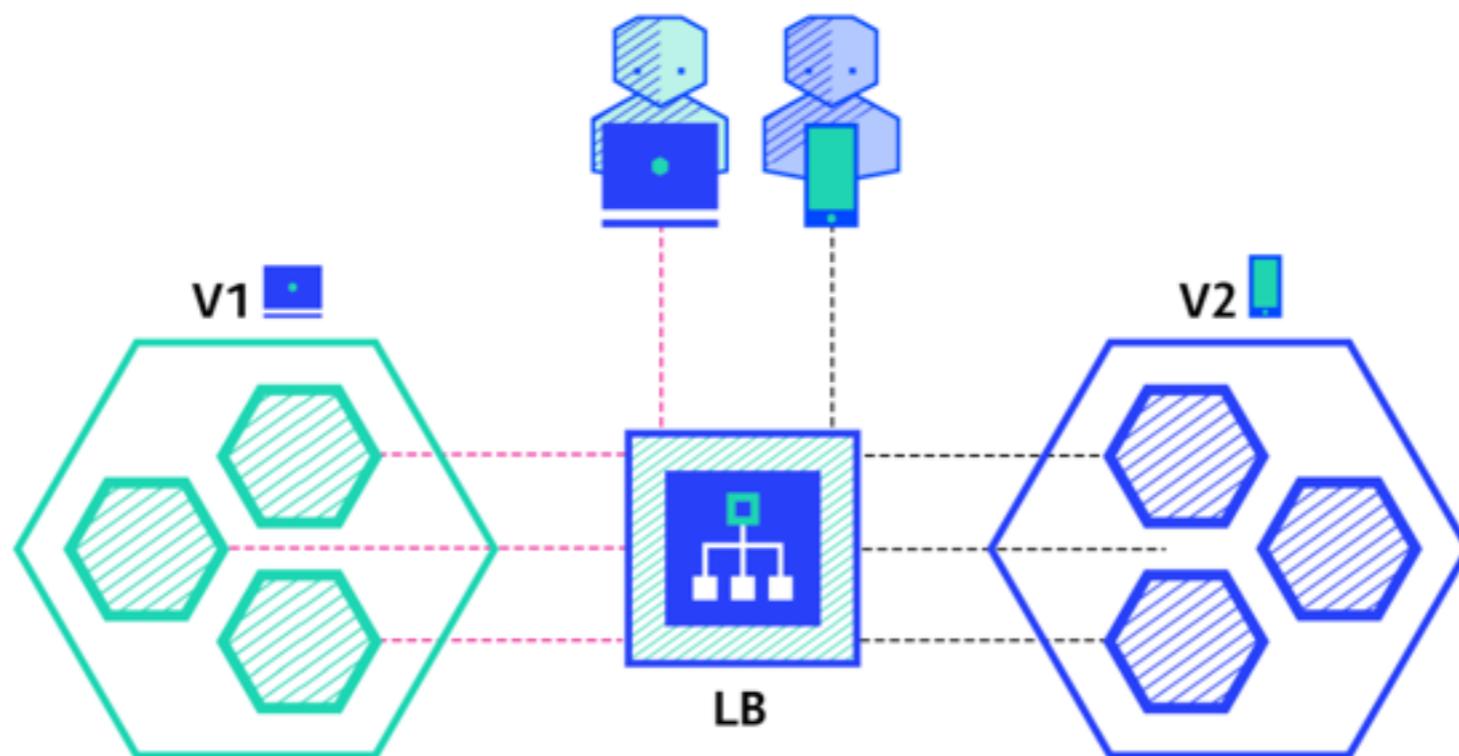
Traffic is split based-on weight

Use when tests are lacking/not reliable and less confident in system



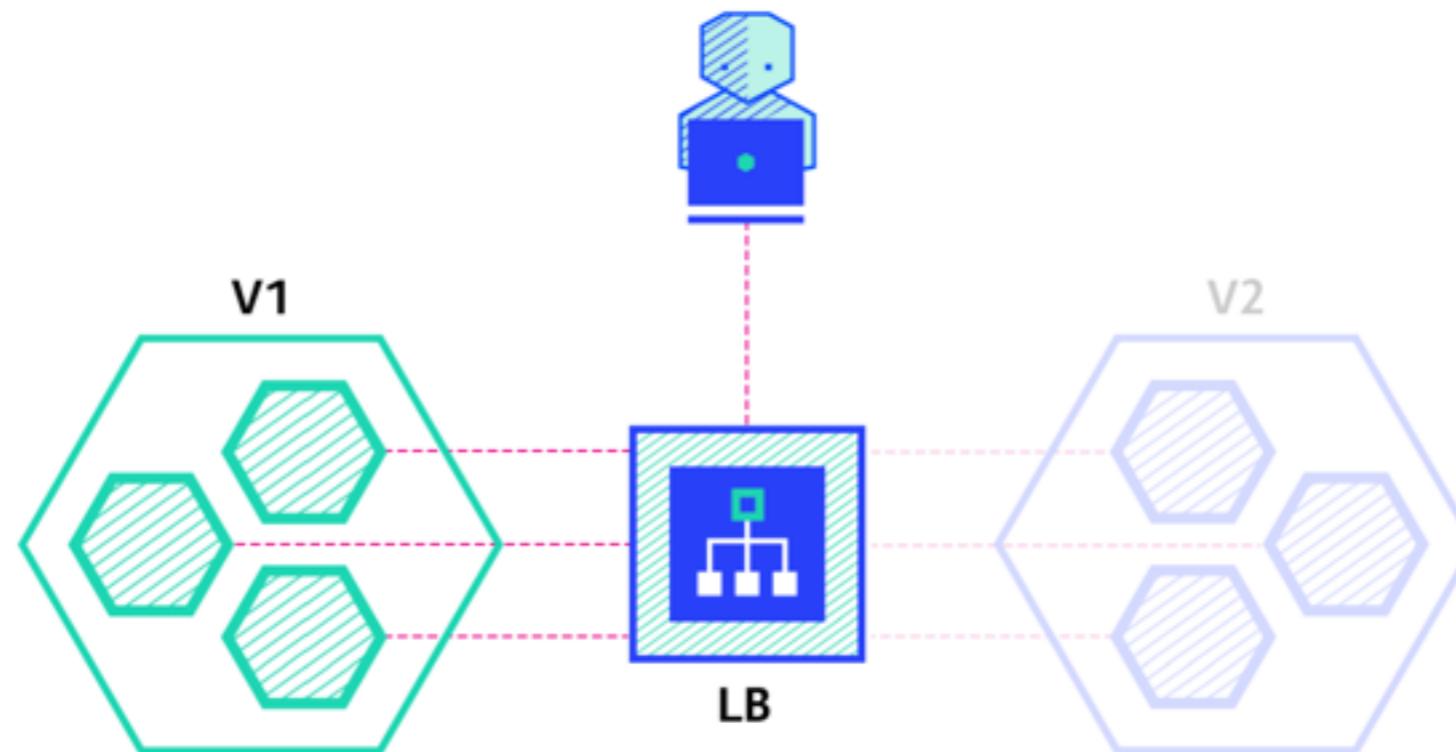
# 5. A/B testing

Routing the subset of users to new services under the specific condition



# 6. Shadow

Release version B alongside version A  
Send request's A to B without production impact



# DEPLOYMENT STRATEGIES

When it comes to production, a ramped or blue/green deployment is usually a good fit, but proper testing of the new platform is necessary.

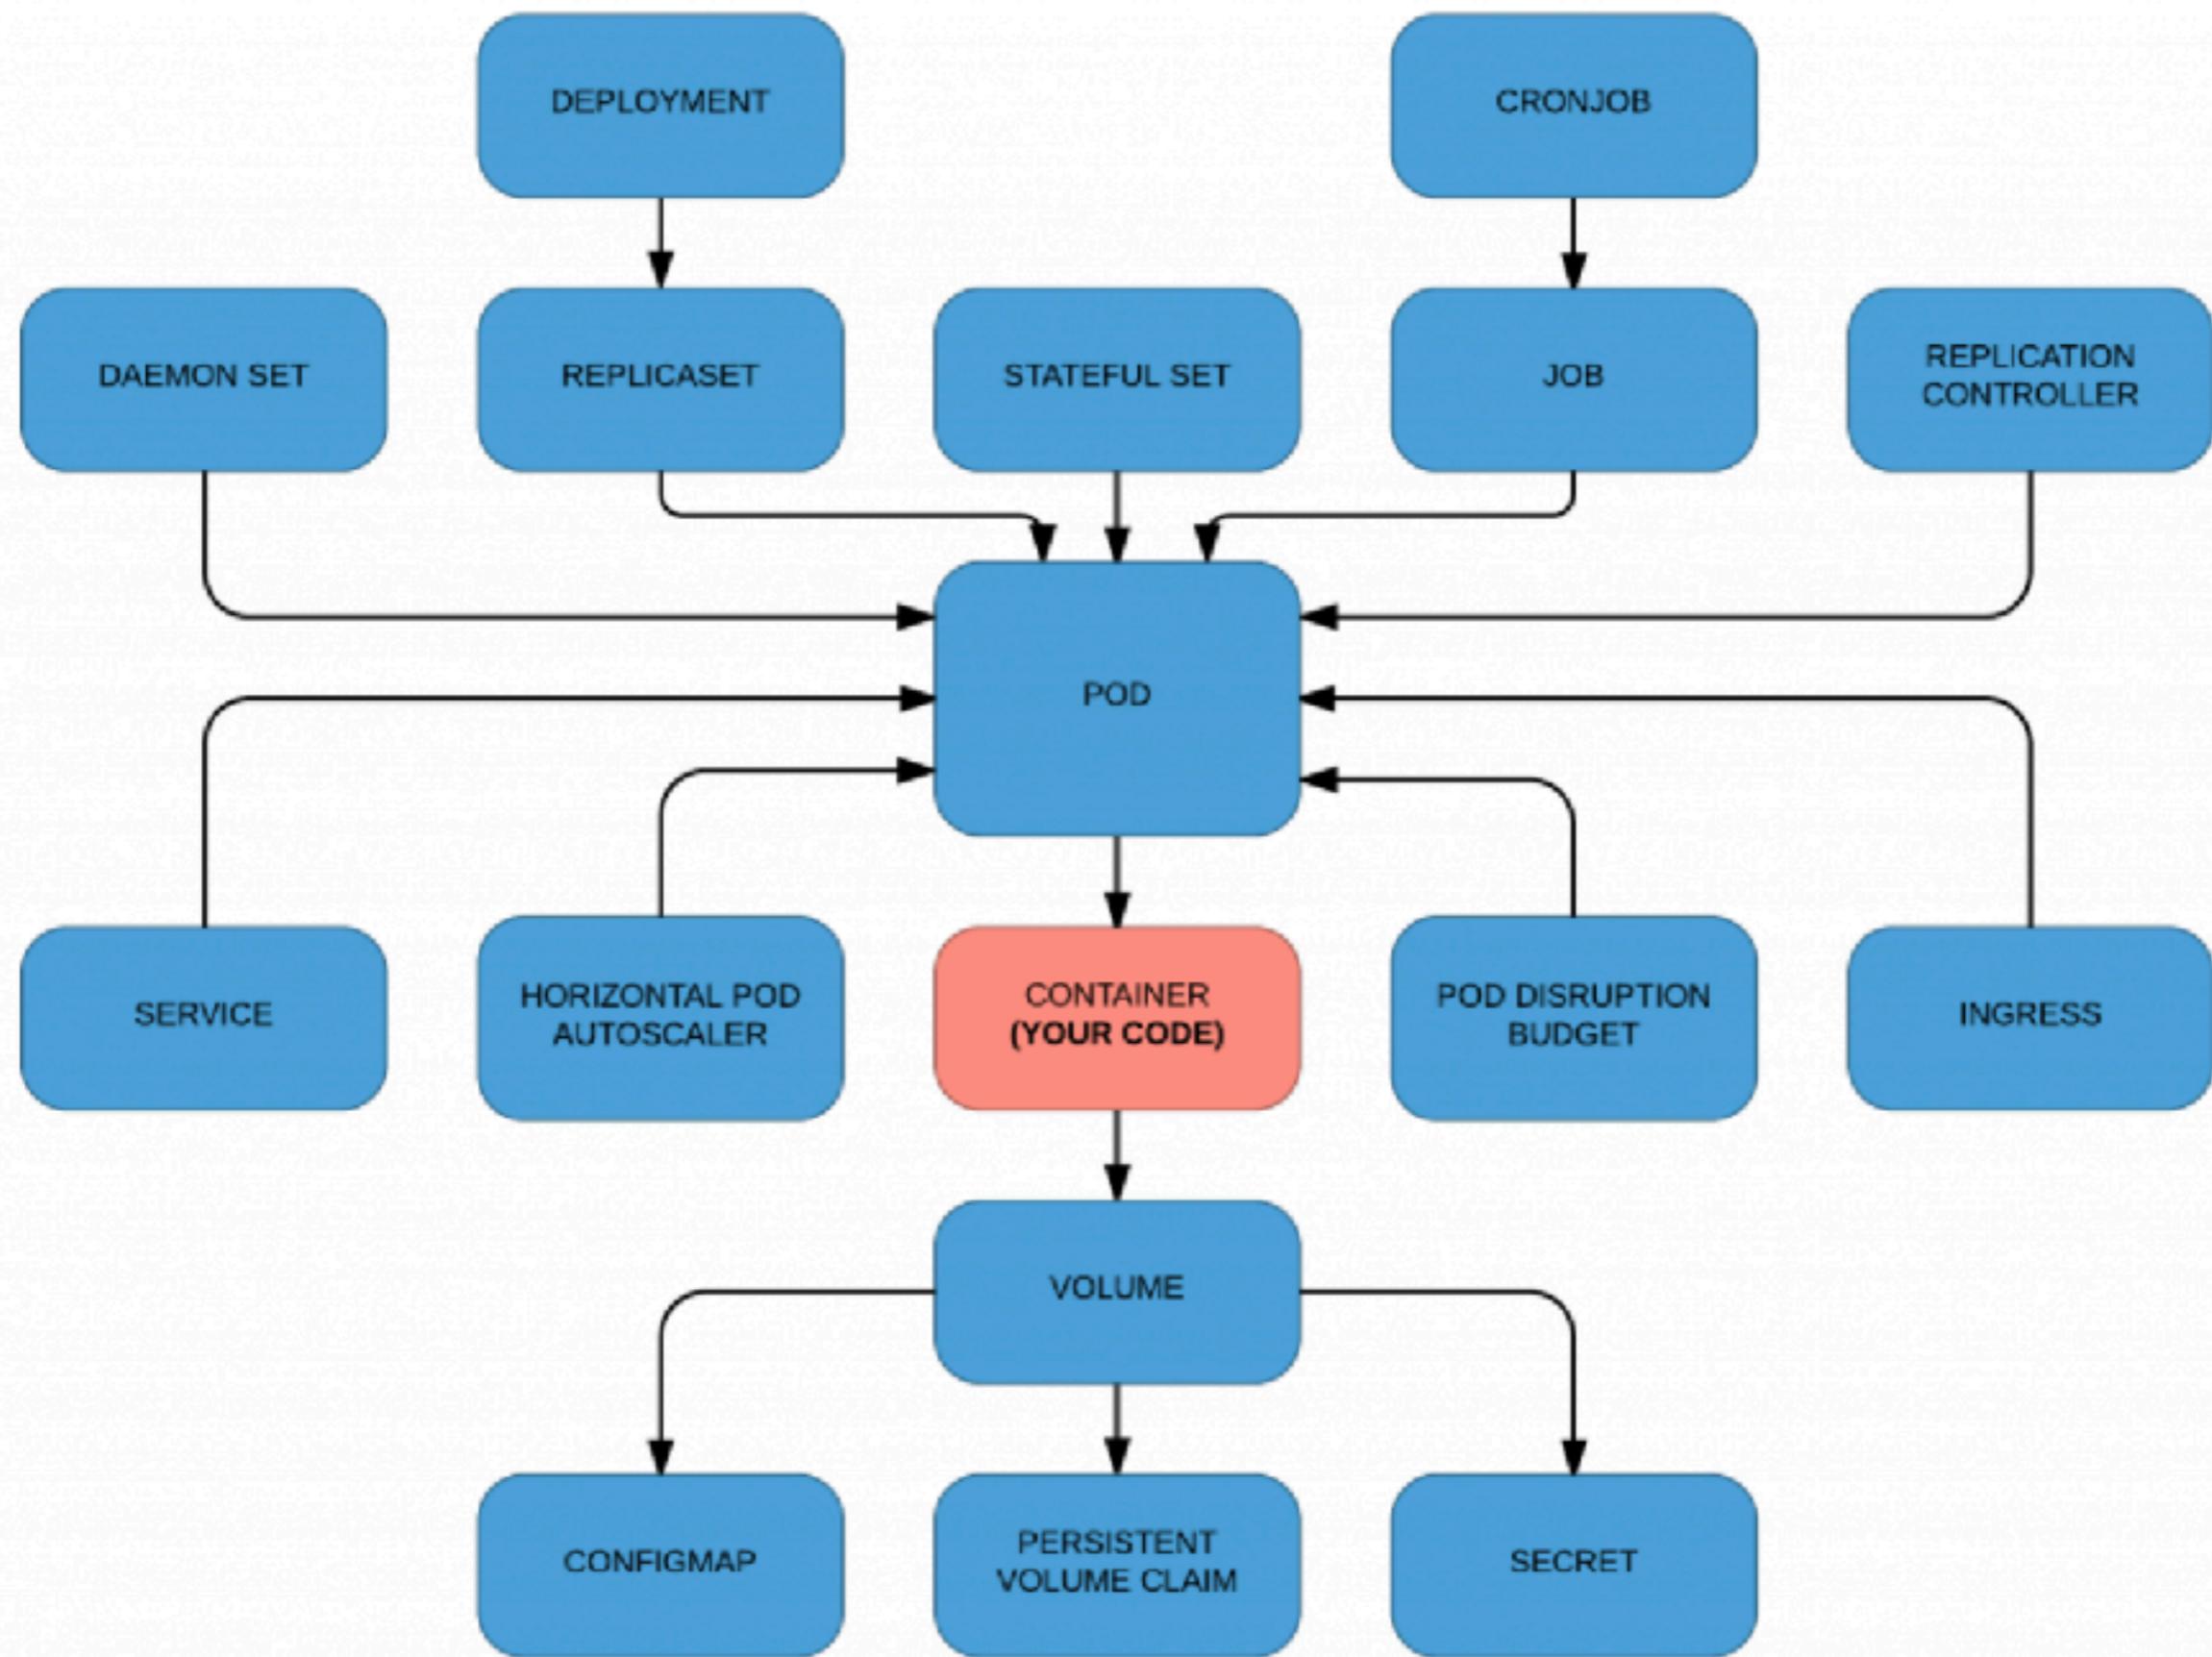
Blue/green and shadow strategies have more impact on the budget as it requires double resource capacity. If the application lacks in tests or if there is little confidence about the impact/stability of the software, then a canary, a/b testing or shadow release can be used.

If your business requires testing of a new feature amongst a specific pool of users that can be filtered depending on some parameters like geolocation, language, operating system or browser features, then you may want to use the a/b testing technique.



Strategy	ZERO DOWNTIME	REAL TRAFFIC TESTING	TARGETED USERS	CLOUD COST	ROLLBACK DURATION	NEGATIVE IMPACT ON USER	COMPLEXITY OF SETUP
<b>RECREATE</b> version A is terminated then version B is rolled out	✗	✗	✗	■ ■ ■	■ ■ ■	■ ■ ■	□ □ □
<b>RAMPED</b> version B is slowly rolled out and replacing version A	✓	✗	✗	■ ■ ■	■ ■ ■	■ □ □	■ □ □
<b>BLUE/GREEN</b> version B is released alongside version A, then the traffic is switched to version B	✓	✗	✗	■ ■ ■	□ □ □	■ ■ □	■ ■ □
<b>CANARY</b> version B is released to a subset of users, then proceed to a full rollout	✓	✓	✗	■ ■ ■	□ □ □	■ □ □	■ ■ □
<b>A/B TESTING</b> version B is released to a subset of users under specific condition	✓	✓	✓	■ ■ ■	□ □ □	■ □ □	■ ■ ■
<b>SHADOW</b> version B receives real world traffic alongside version A and doesn't impact the response	✓	✓	✗	■ ■ ■	□ □ □	□ □ □	■ ■ ■



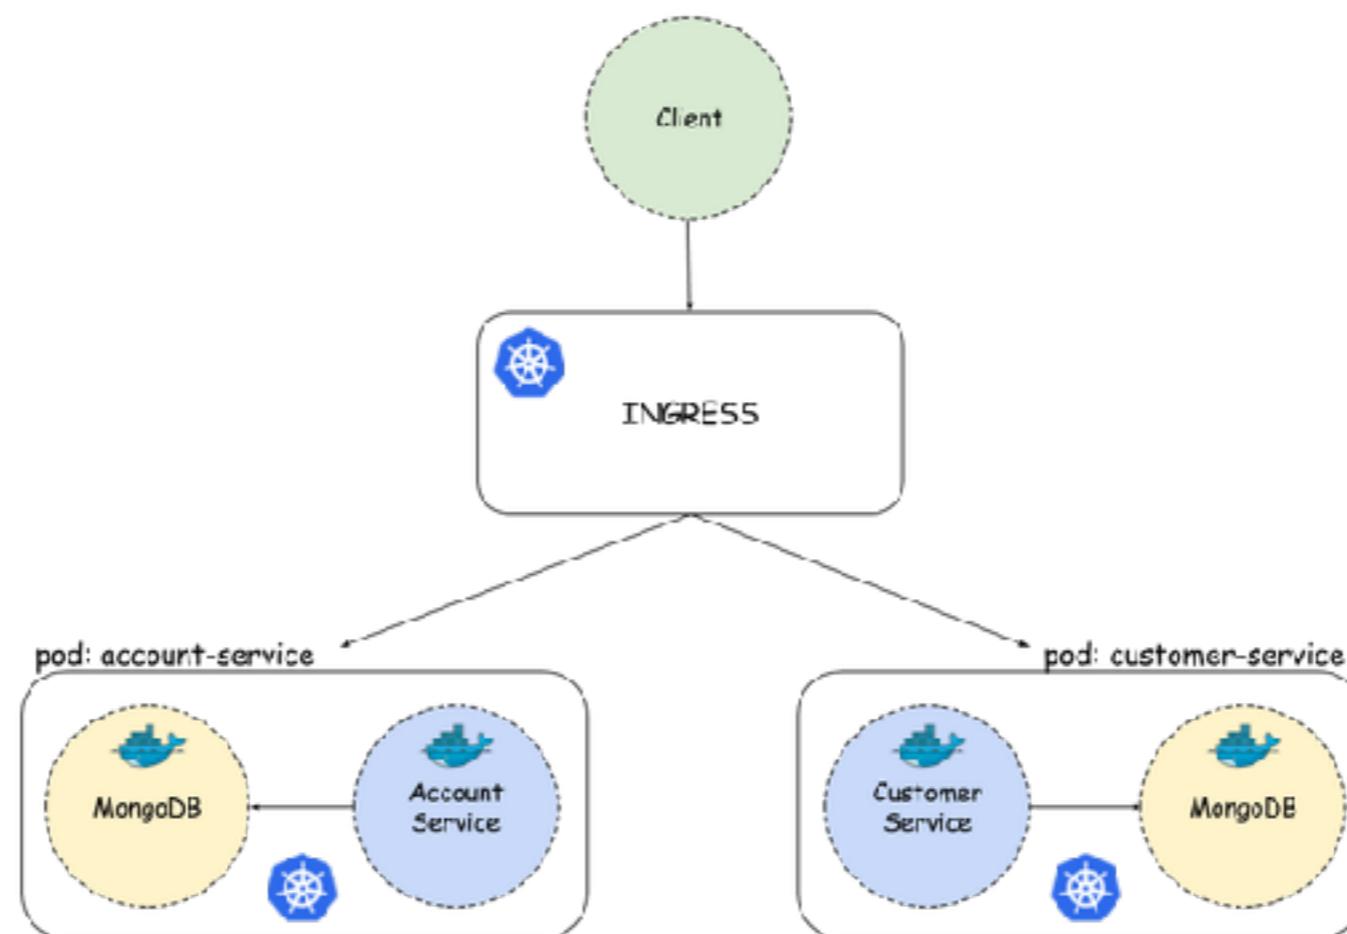


# Ingress Network



# Ingress Network

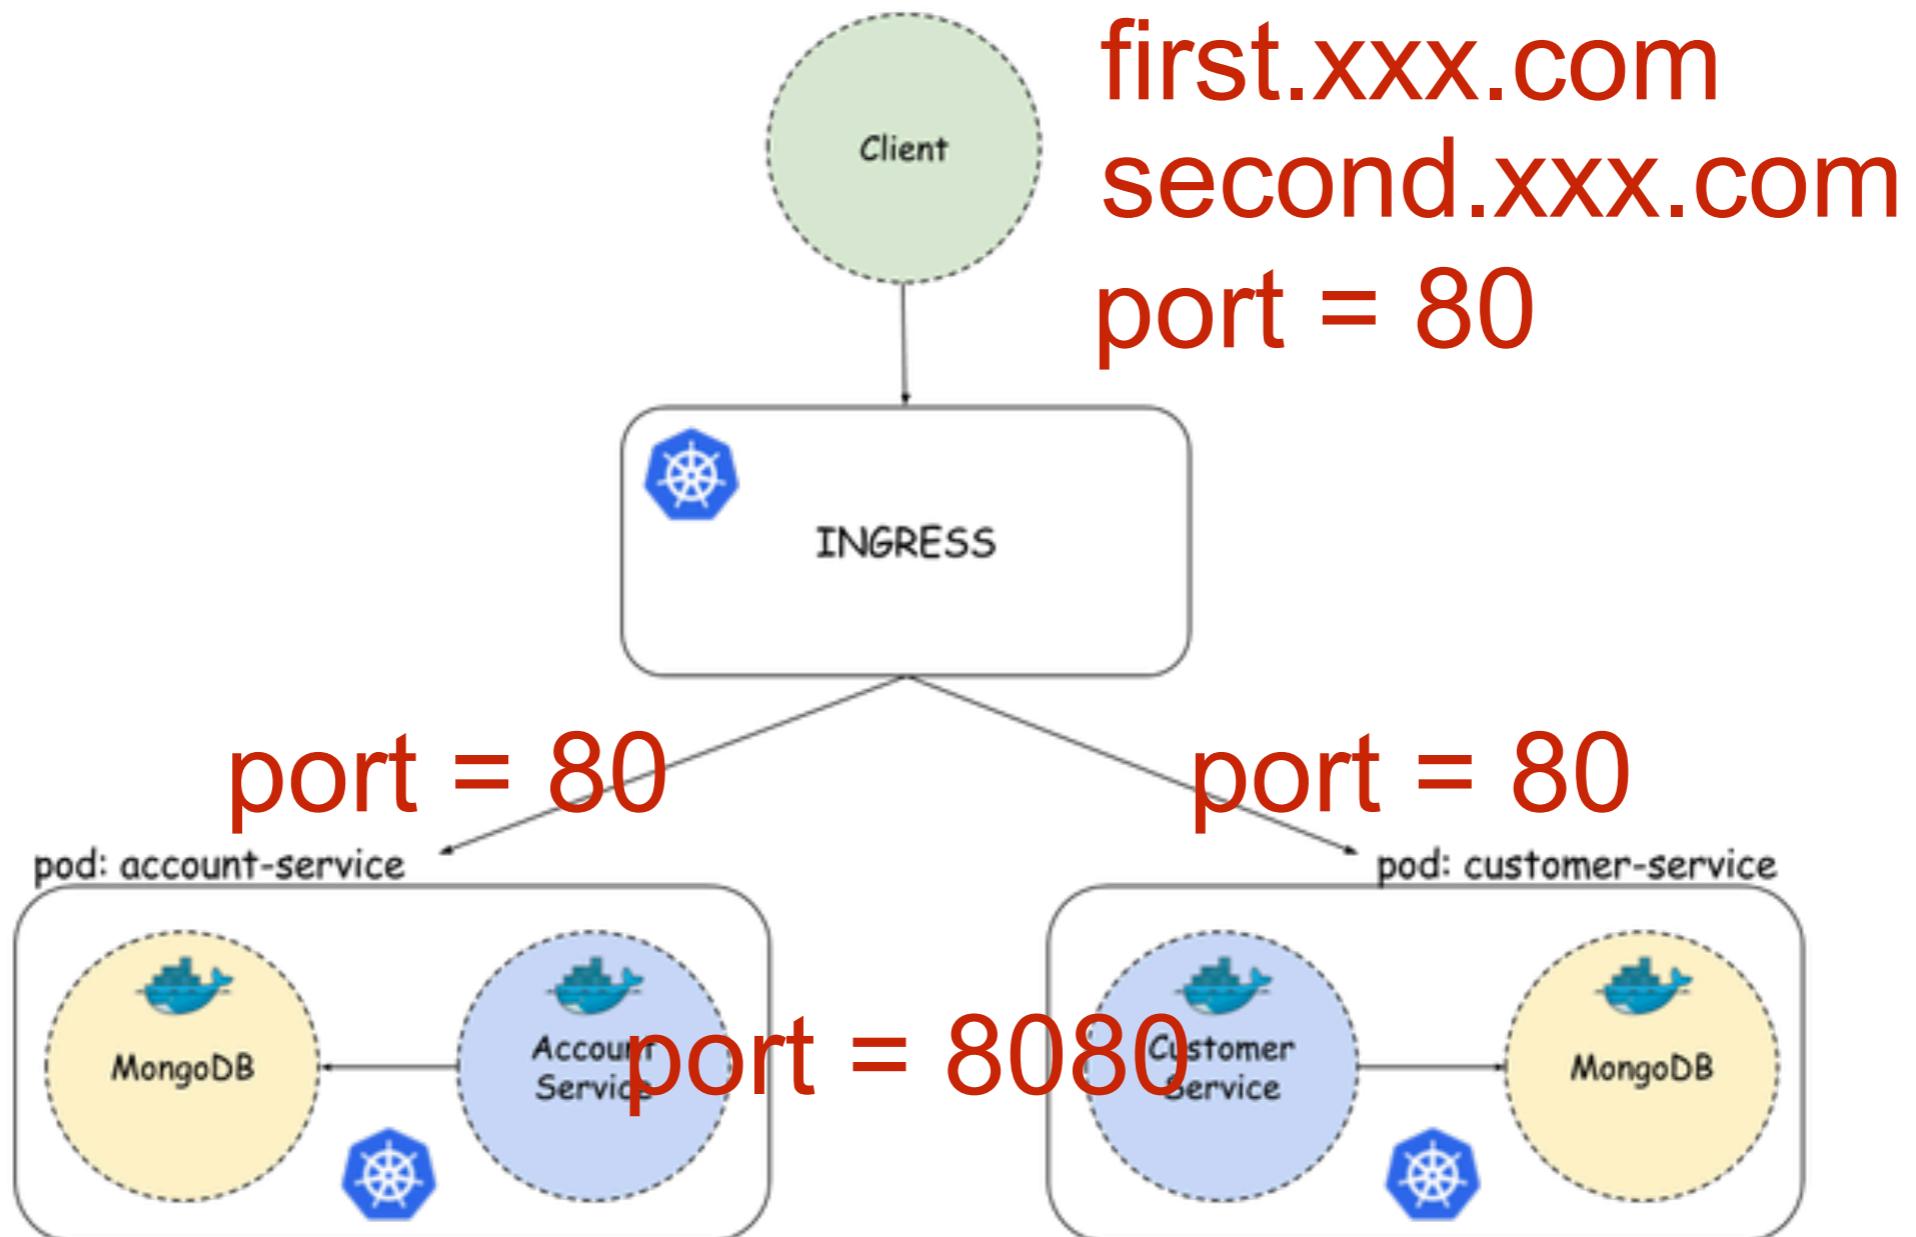
How to handle multiple services in same port ?  
How to limit protocol to access ?



<https://kubernetes.io/docs/concepts/services-networking/ingress/>



# Ingress Network



# Create ingress

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: ingresswebtest
spec:
  rules:
    - host: first.xxx.com
      http:
        paths:
          - backend:
              serviceName: spring-boot-service
              servicePort: 80
    - host: second.xxx.com
      http:
        paths:
          - backend:
              serviceName: spring-boot-service
              servicePort: 80
```



# Service

```
apiVersion: v1
kind: Service
metadata:
  name: spring-boot-service
spec:
  selector:
    app: spring-boot-service
  type: NodePort
  ports:
    - port: 80
      name: http
      targetPort: 8080
      protocol: TCP
```



# Workshop

# Ingress Network

File /working-with-java/05-ingress



# Liveness and Readiness probe



# Job and Cronjob



# Thank you



# Wordpress and MySQL

File /workshop-wordpress

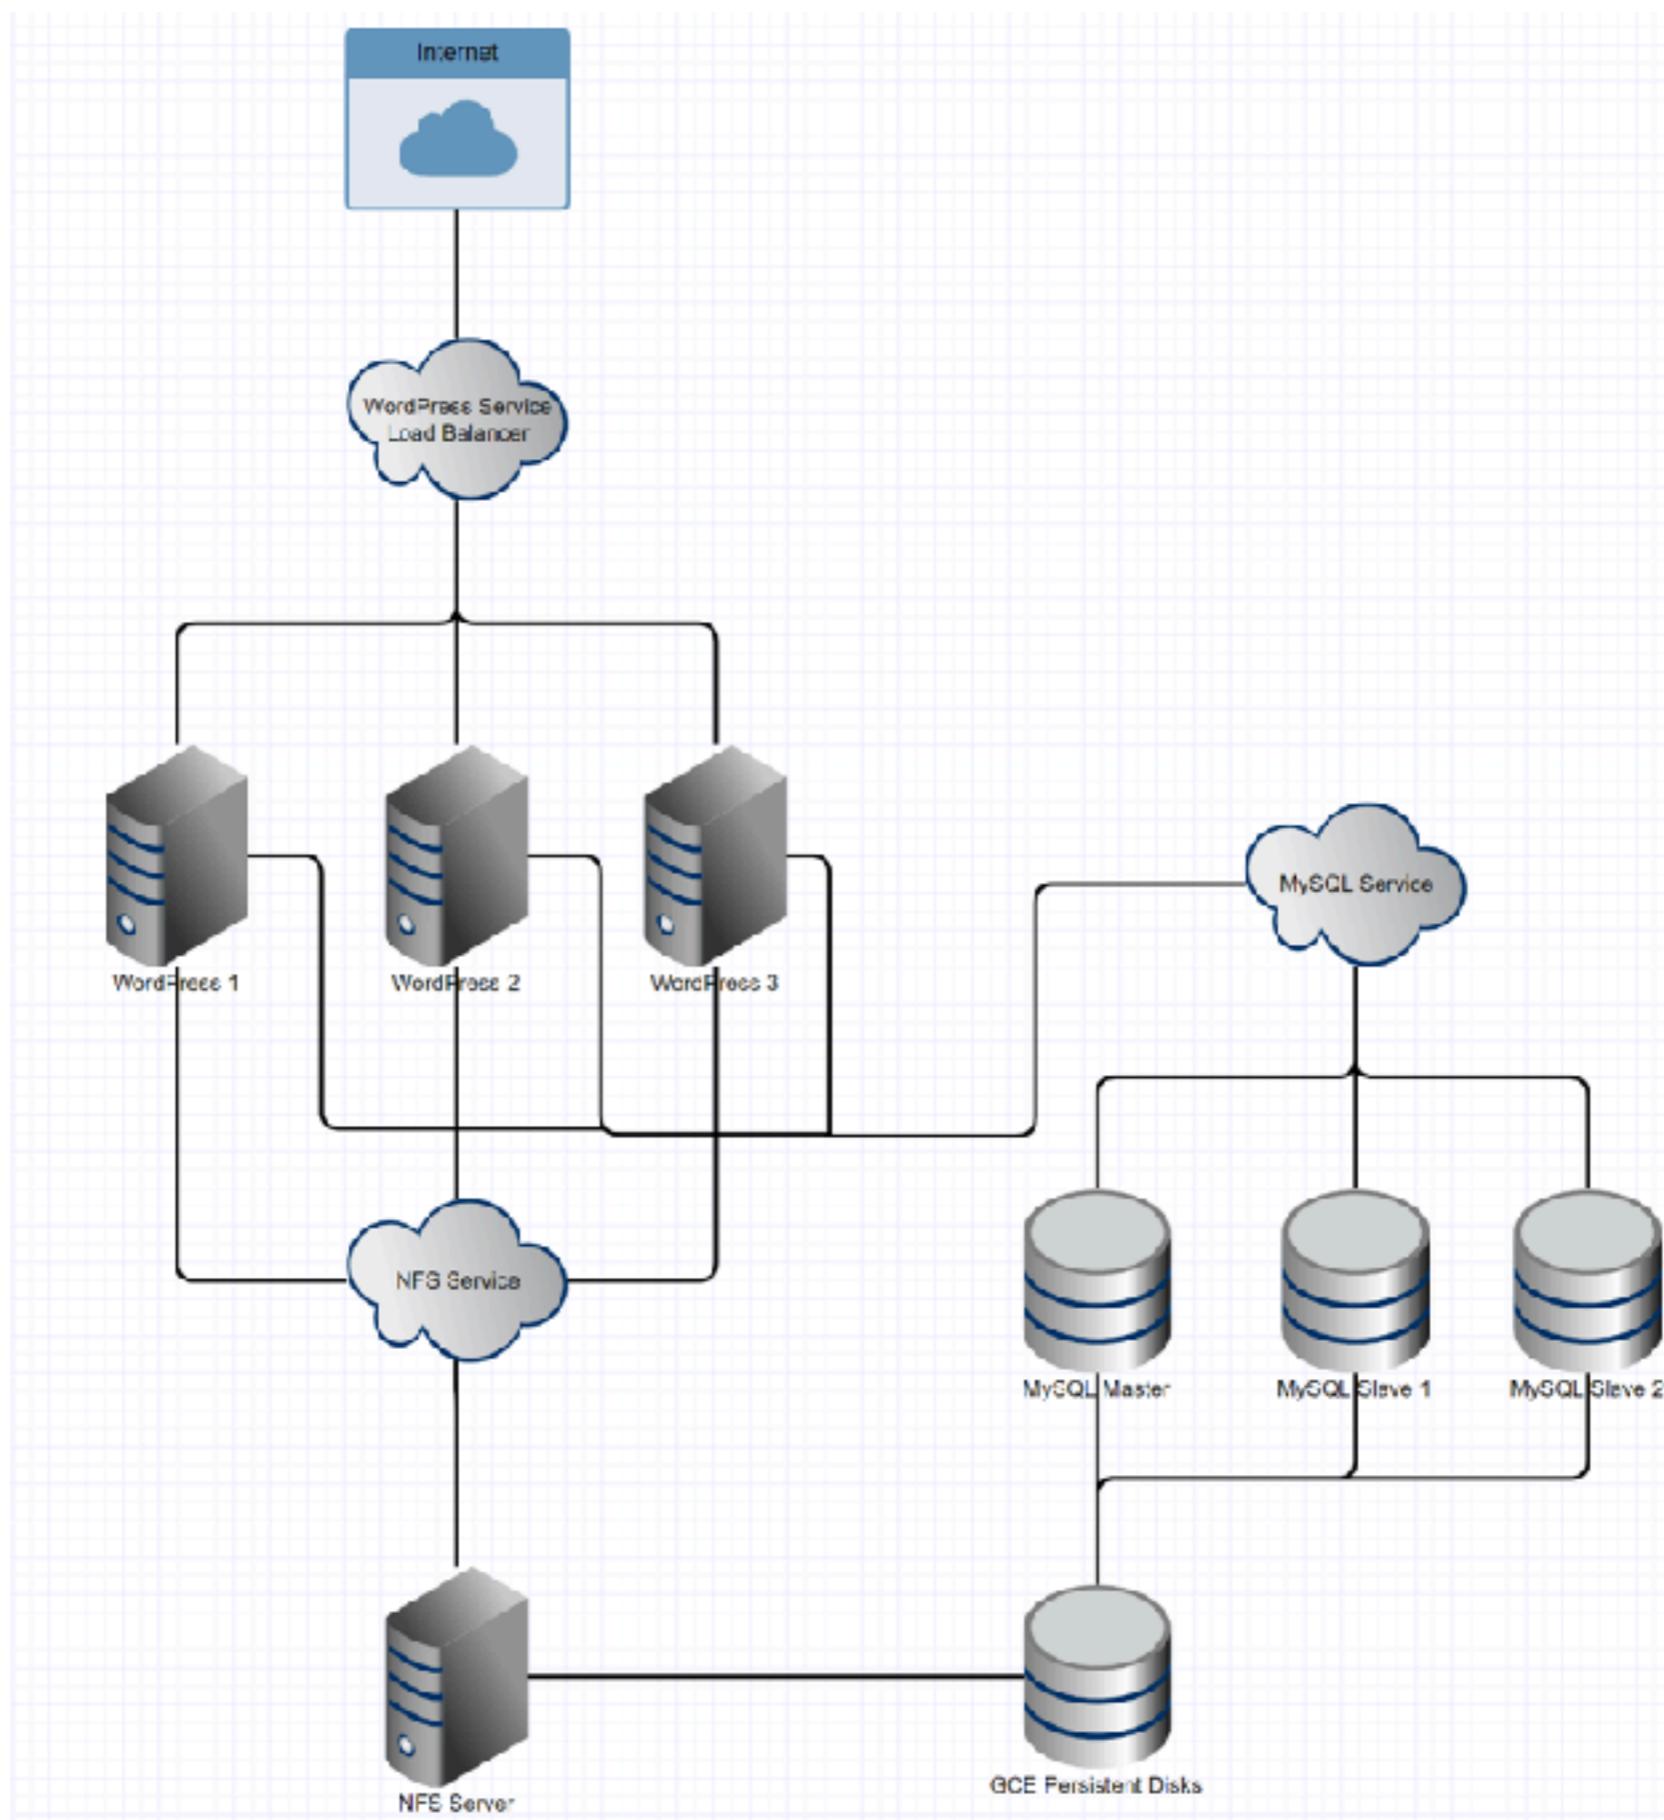


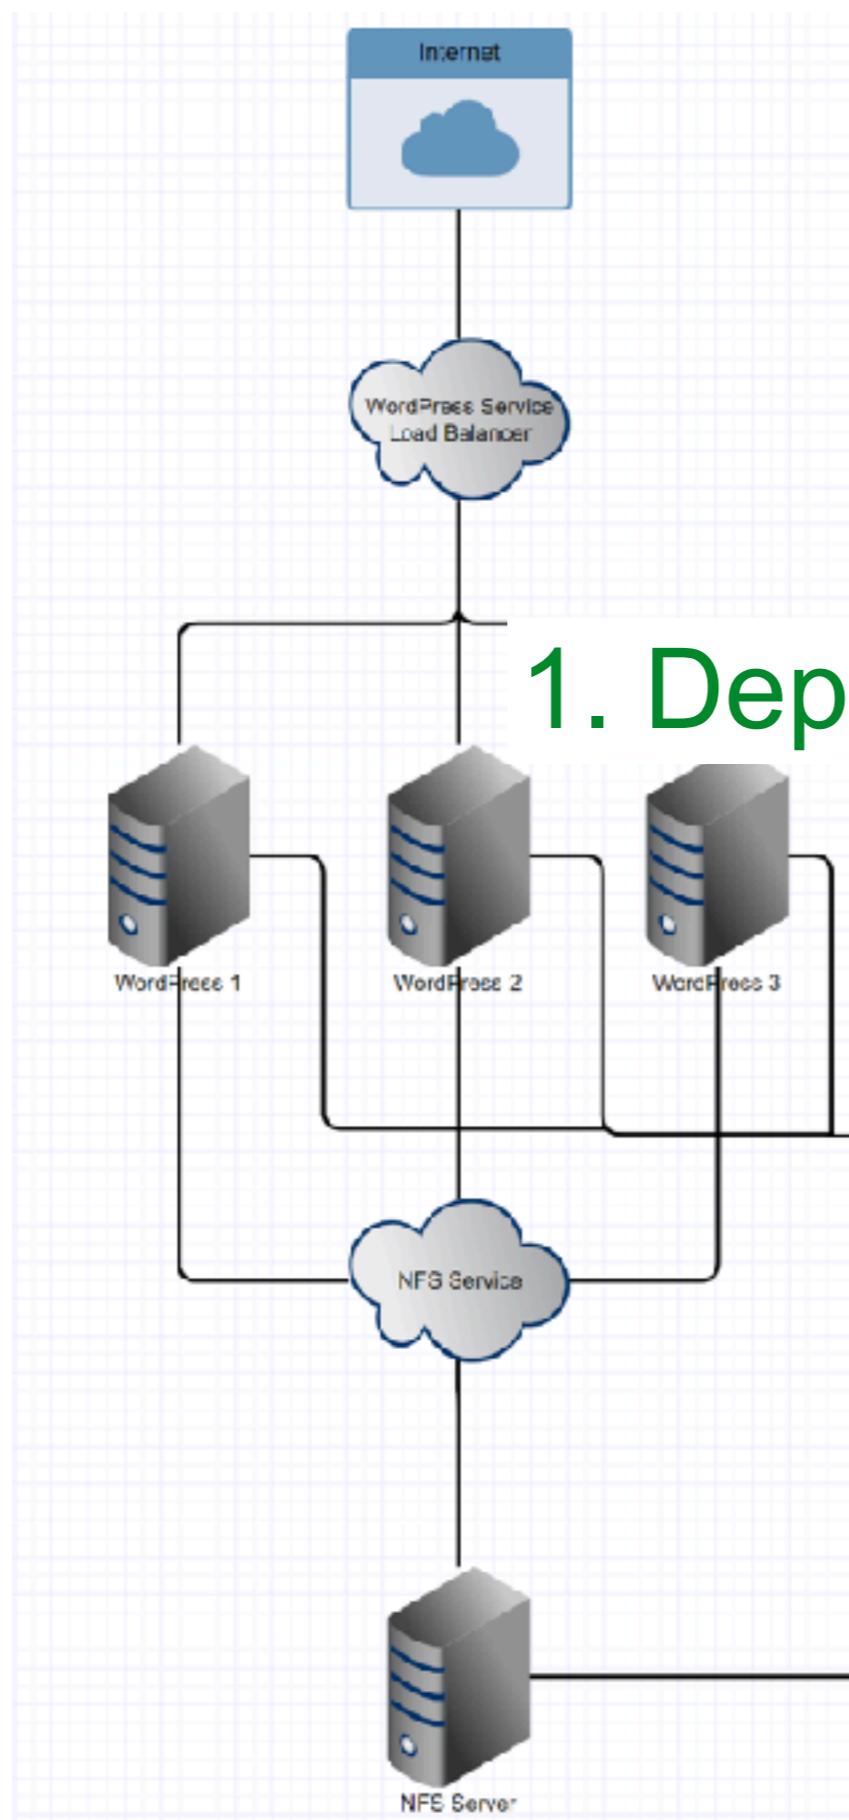
# Architecture

Storage for Wordpress application file with **NFS**  
Database cluster with **MySQL**

Use Kubernetes load balance and **service** networking

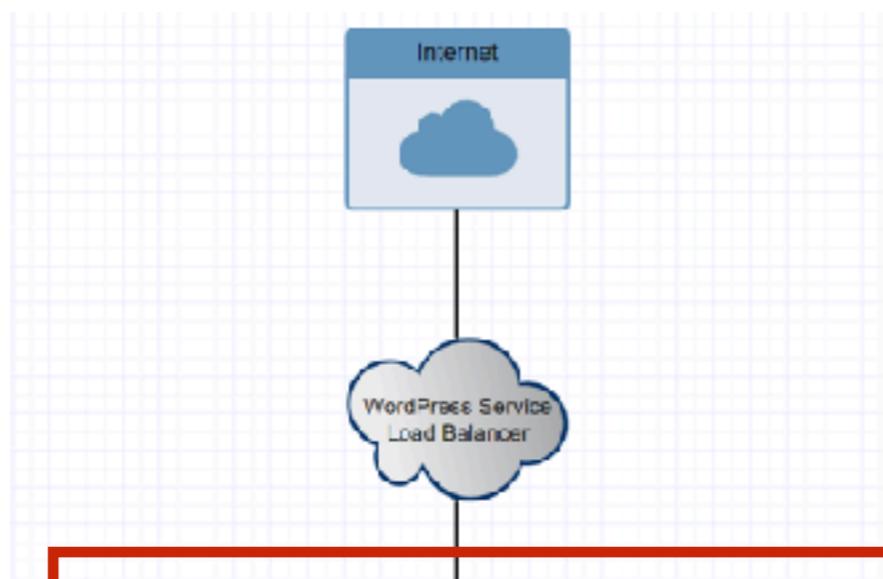




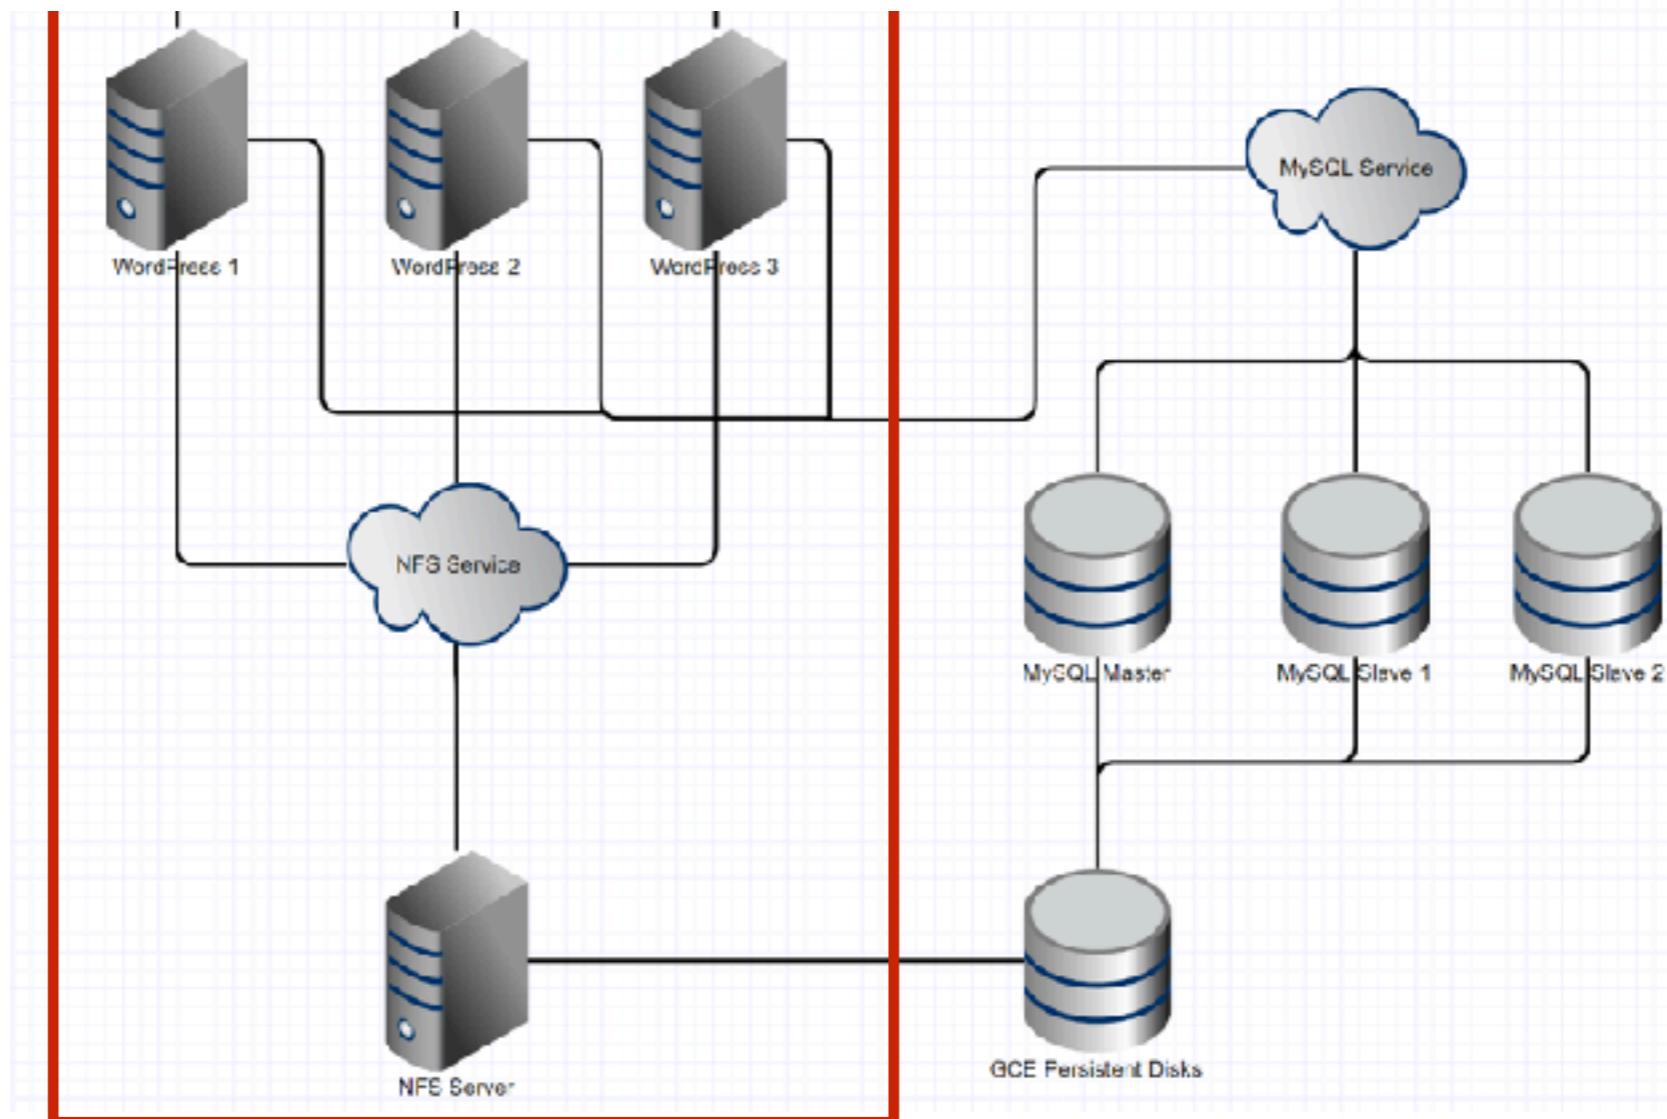


## 1. Deploy MySQL with StatefulSet





## 2. Share wordpress files with NFS



### 3. Expose port with service

