



# Kubernetes

In Practice







Somkiat Puisungnoen

Update Info 1 View Activity Log 10+ ...

Timeline About Friends 3,138 Photos More

When did you work at Opendream? X

... 22 Pending Items

Post Photo/Video Live Video Life Event

What's on your mind?

Public Post

Intro

Software Craftsmanship

Software Practitioner at สยามชั่นนาฎกิจ พ.ศ. 2556

Agile Practitioner and Technical at SPRINT3r

Somkiat Puisungnoen 15 mins · Bangkok · ...

Java and Bigdata



somkiat.cc

Page Messages Notifications 3 Insights Publishing Tools Settings Help ▾

somkiat.cc  
@somkiat.cc

Home Posts Videos Photos

Liked Following Share ... + Add a Button



# Agenda Day 1

1. Cloud Native Application
2. Kubernetes architecture
3. Key-features
4. Pods and Containers
5. Service
6. Replication Controller (RC)
7. Deployment and ReplicaSet (RS)
8. Volume



# Agenda Day 2

1. Resource management
2. Horizontal Pods Autoscaler (HPA)
3. ConfigMap and Secret
4. Log and monitoring
5. Ingress network
6. Batch execution
7. Persistence storage

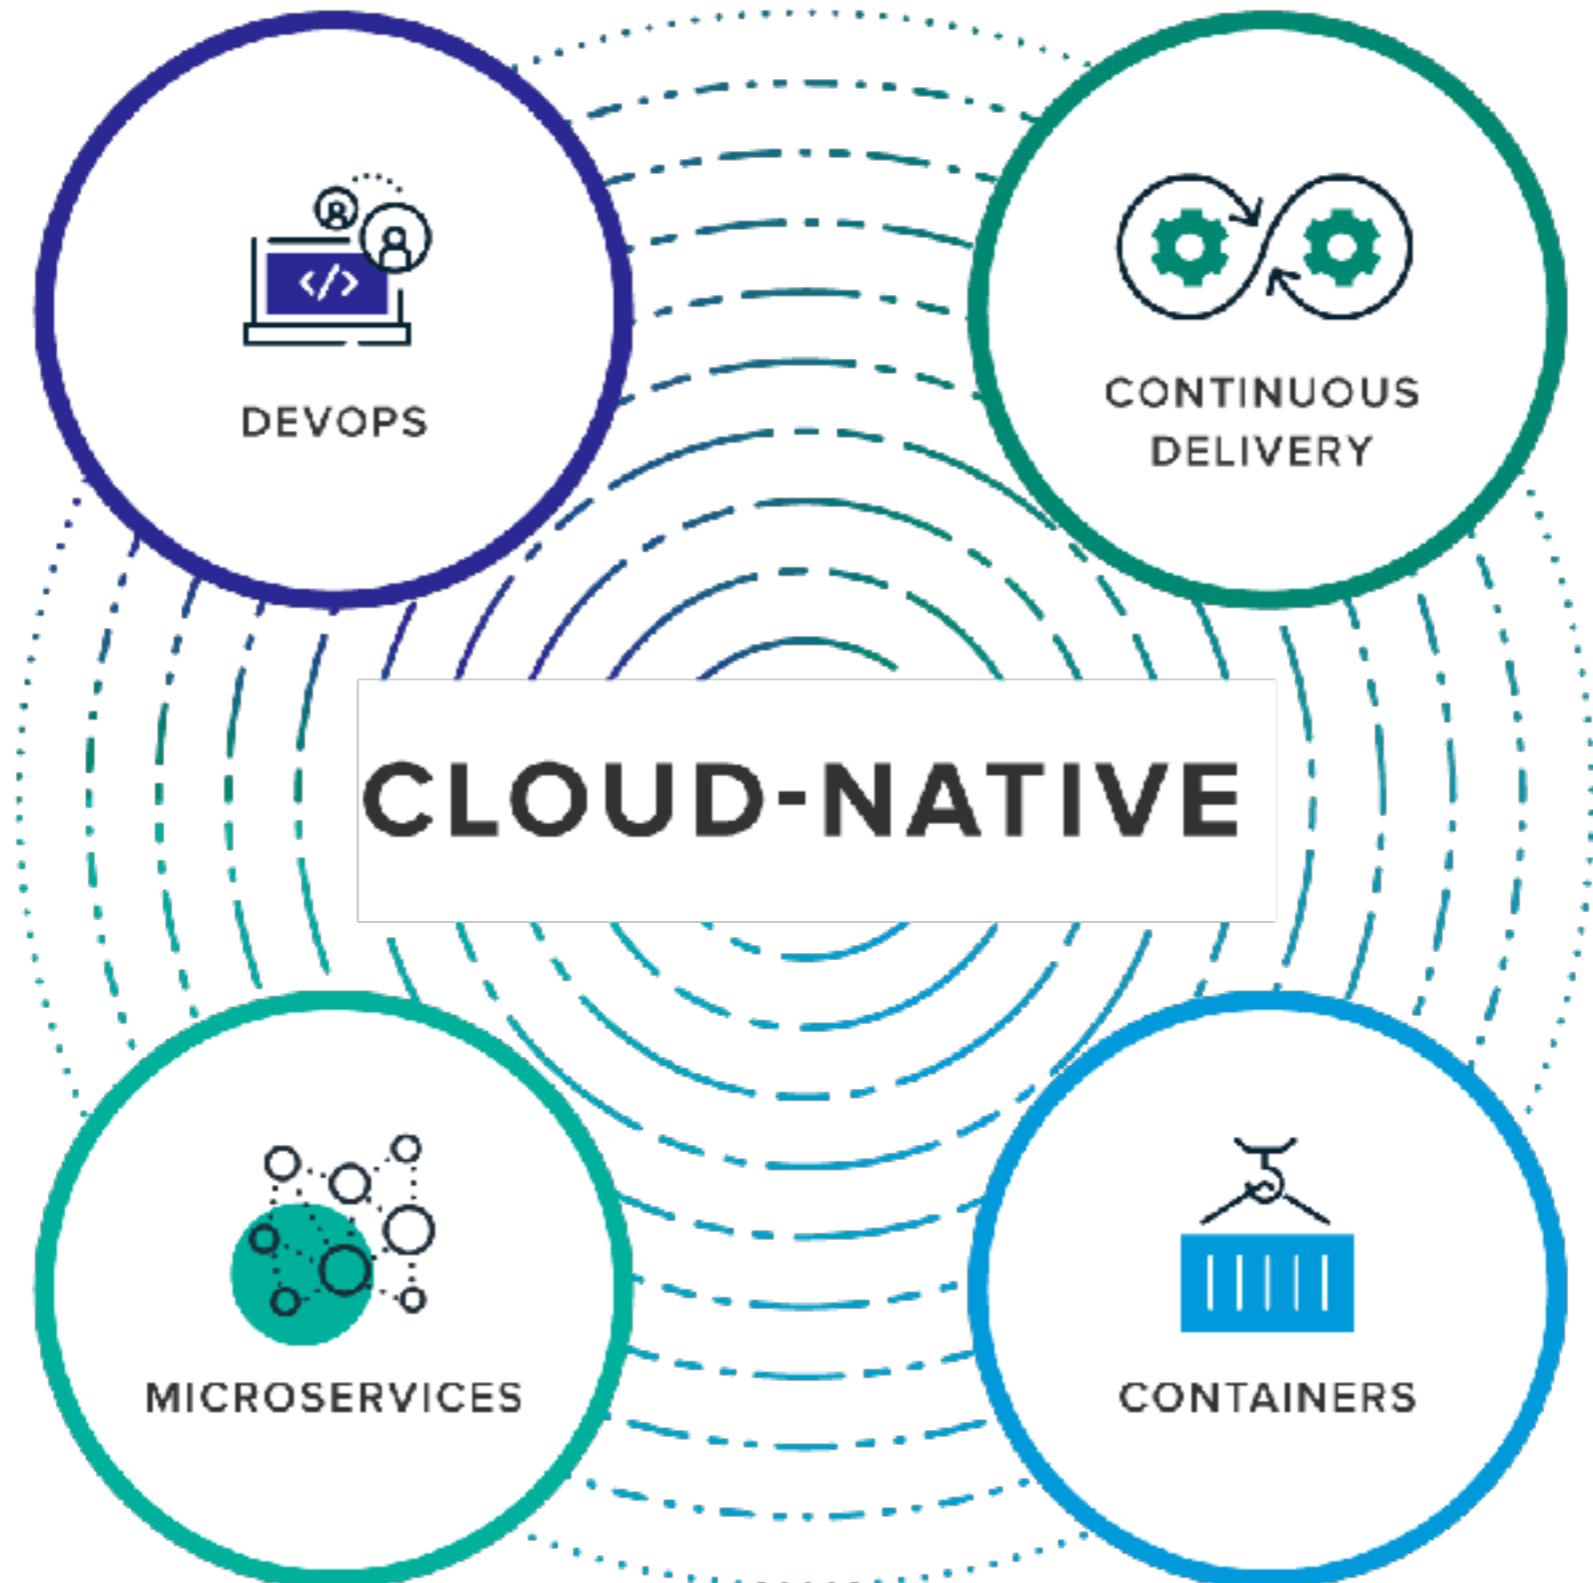


**<https://github.com/up1/course-kubernetes-in-practice>**



# Cloud Native Application





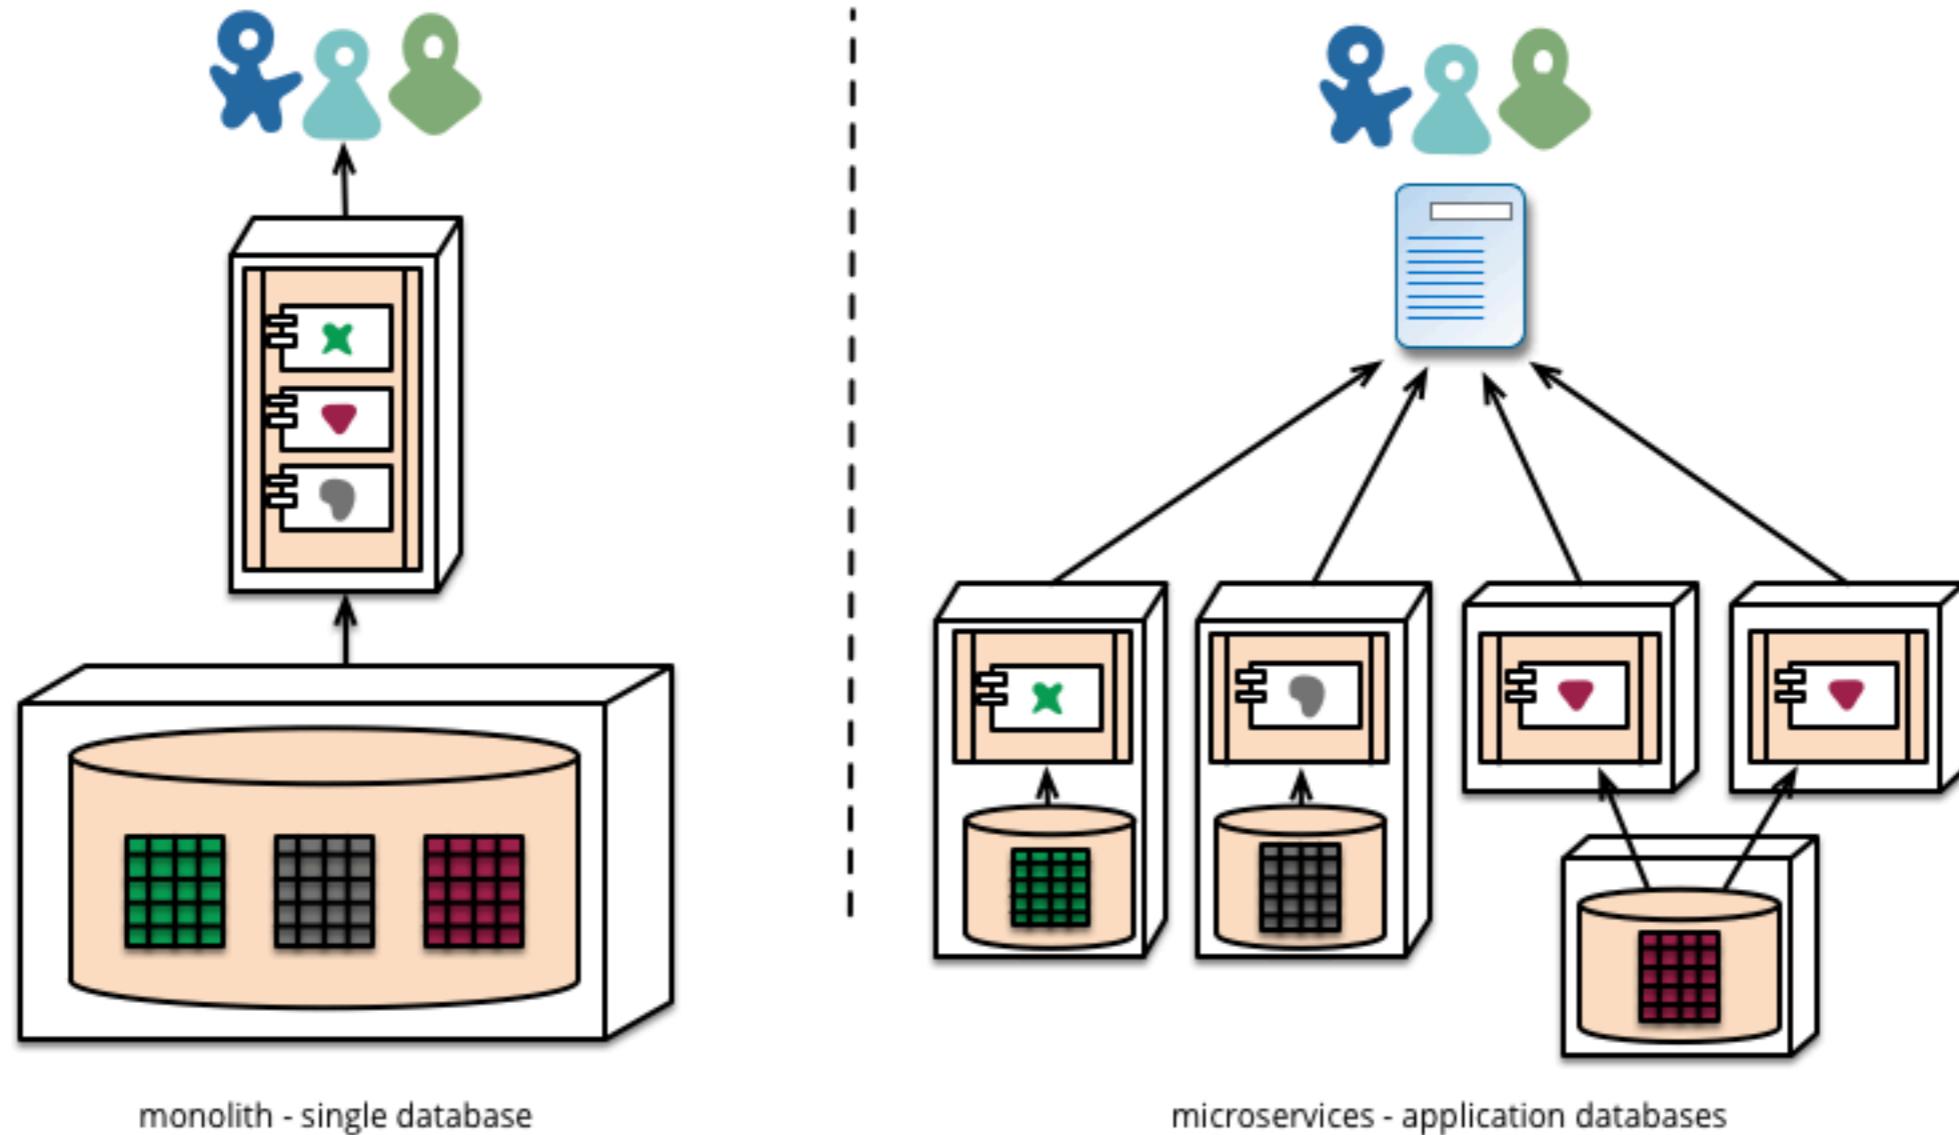
<https://pivotal.io/cloud-native>



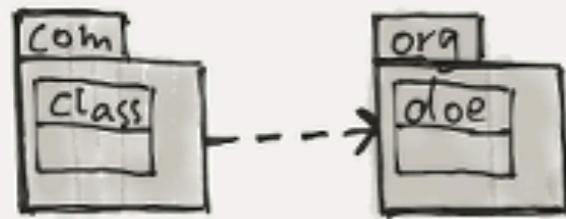
# Evolution of Architecture



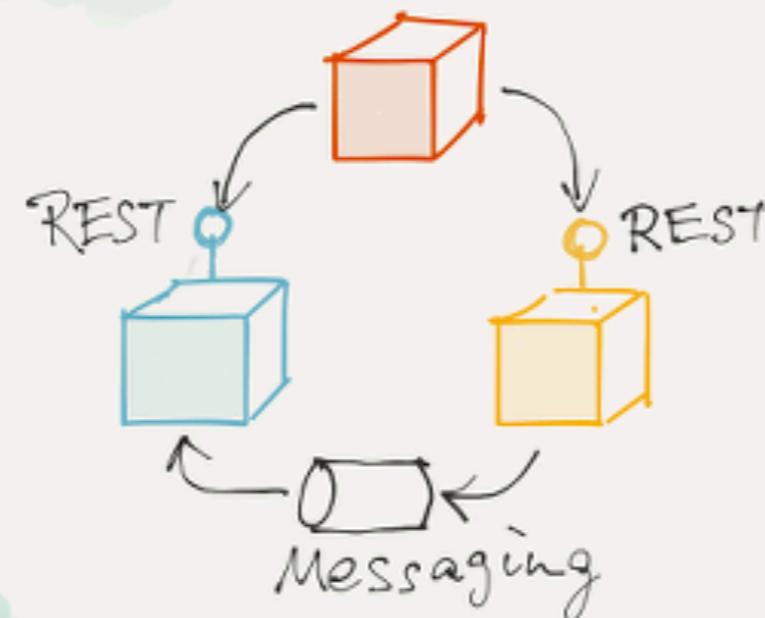
# Microservices



# Architecture



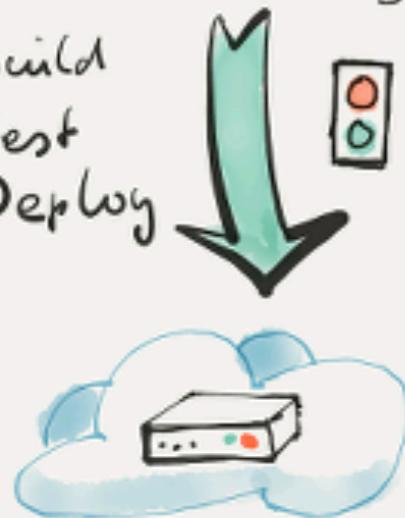
# Microservices



# Deployment

## Continuous Delivery

`{ var i=1; }`  
Build  
Test  
Deploy



# Infrastructure



# People & Teams



Communication  
Collaboration

# Monitoring

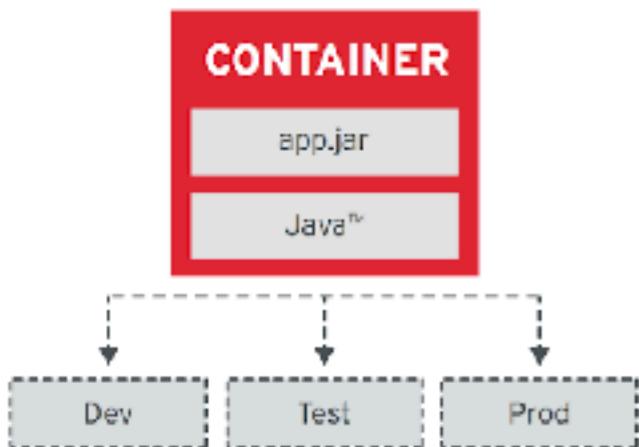


Features & Technology



# Container design principles

Image Immutability Principle



High Observability Principle



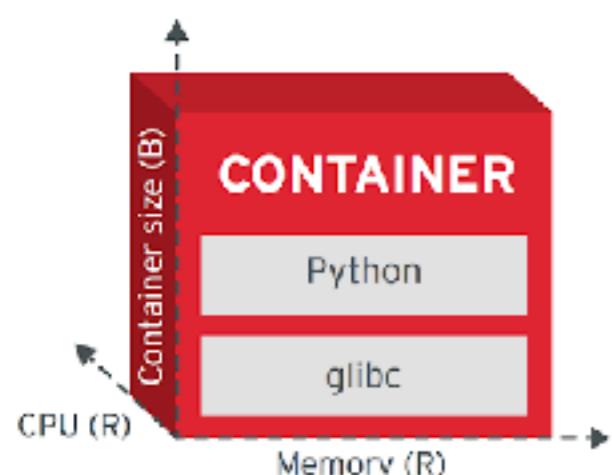
Process Disposability Principle



Lifecycle Conformance Principle



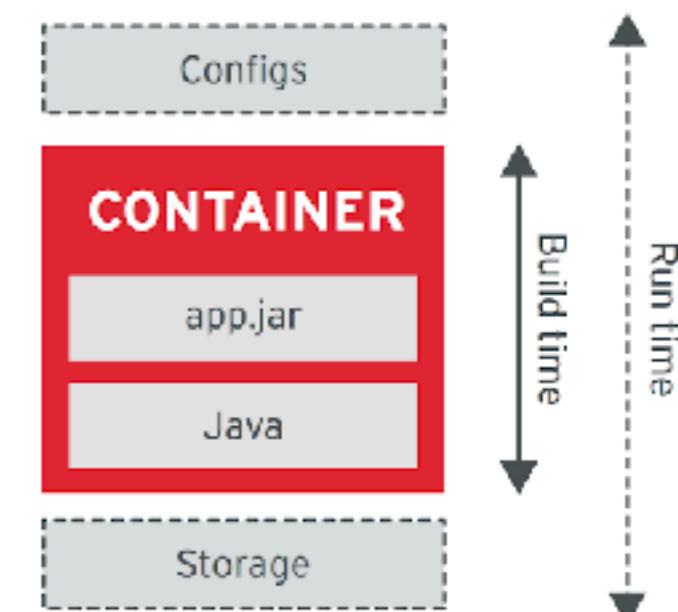
Runtime Confinement Principle



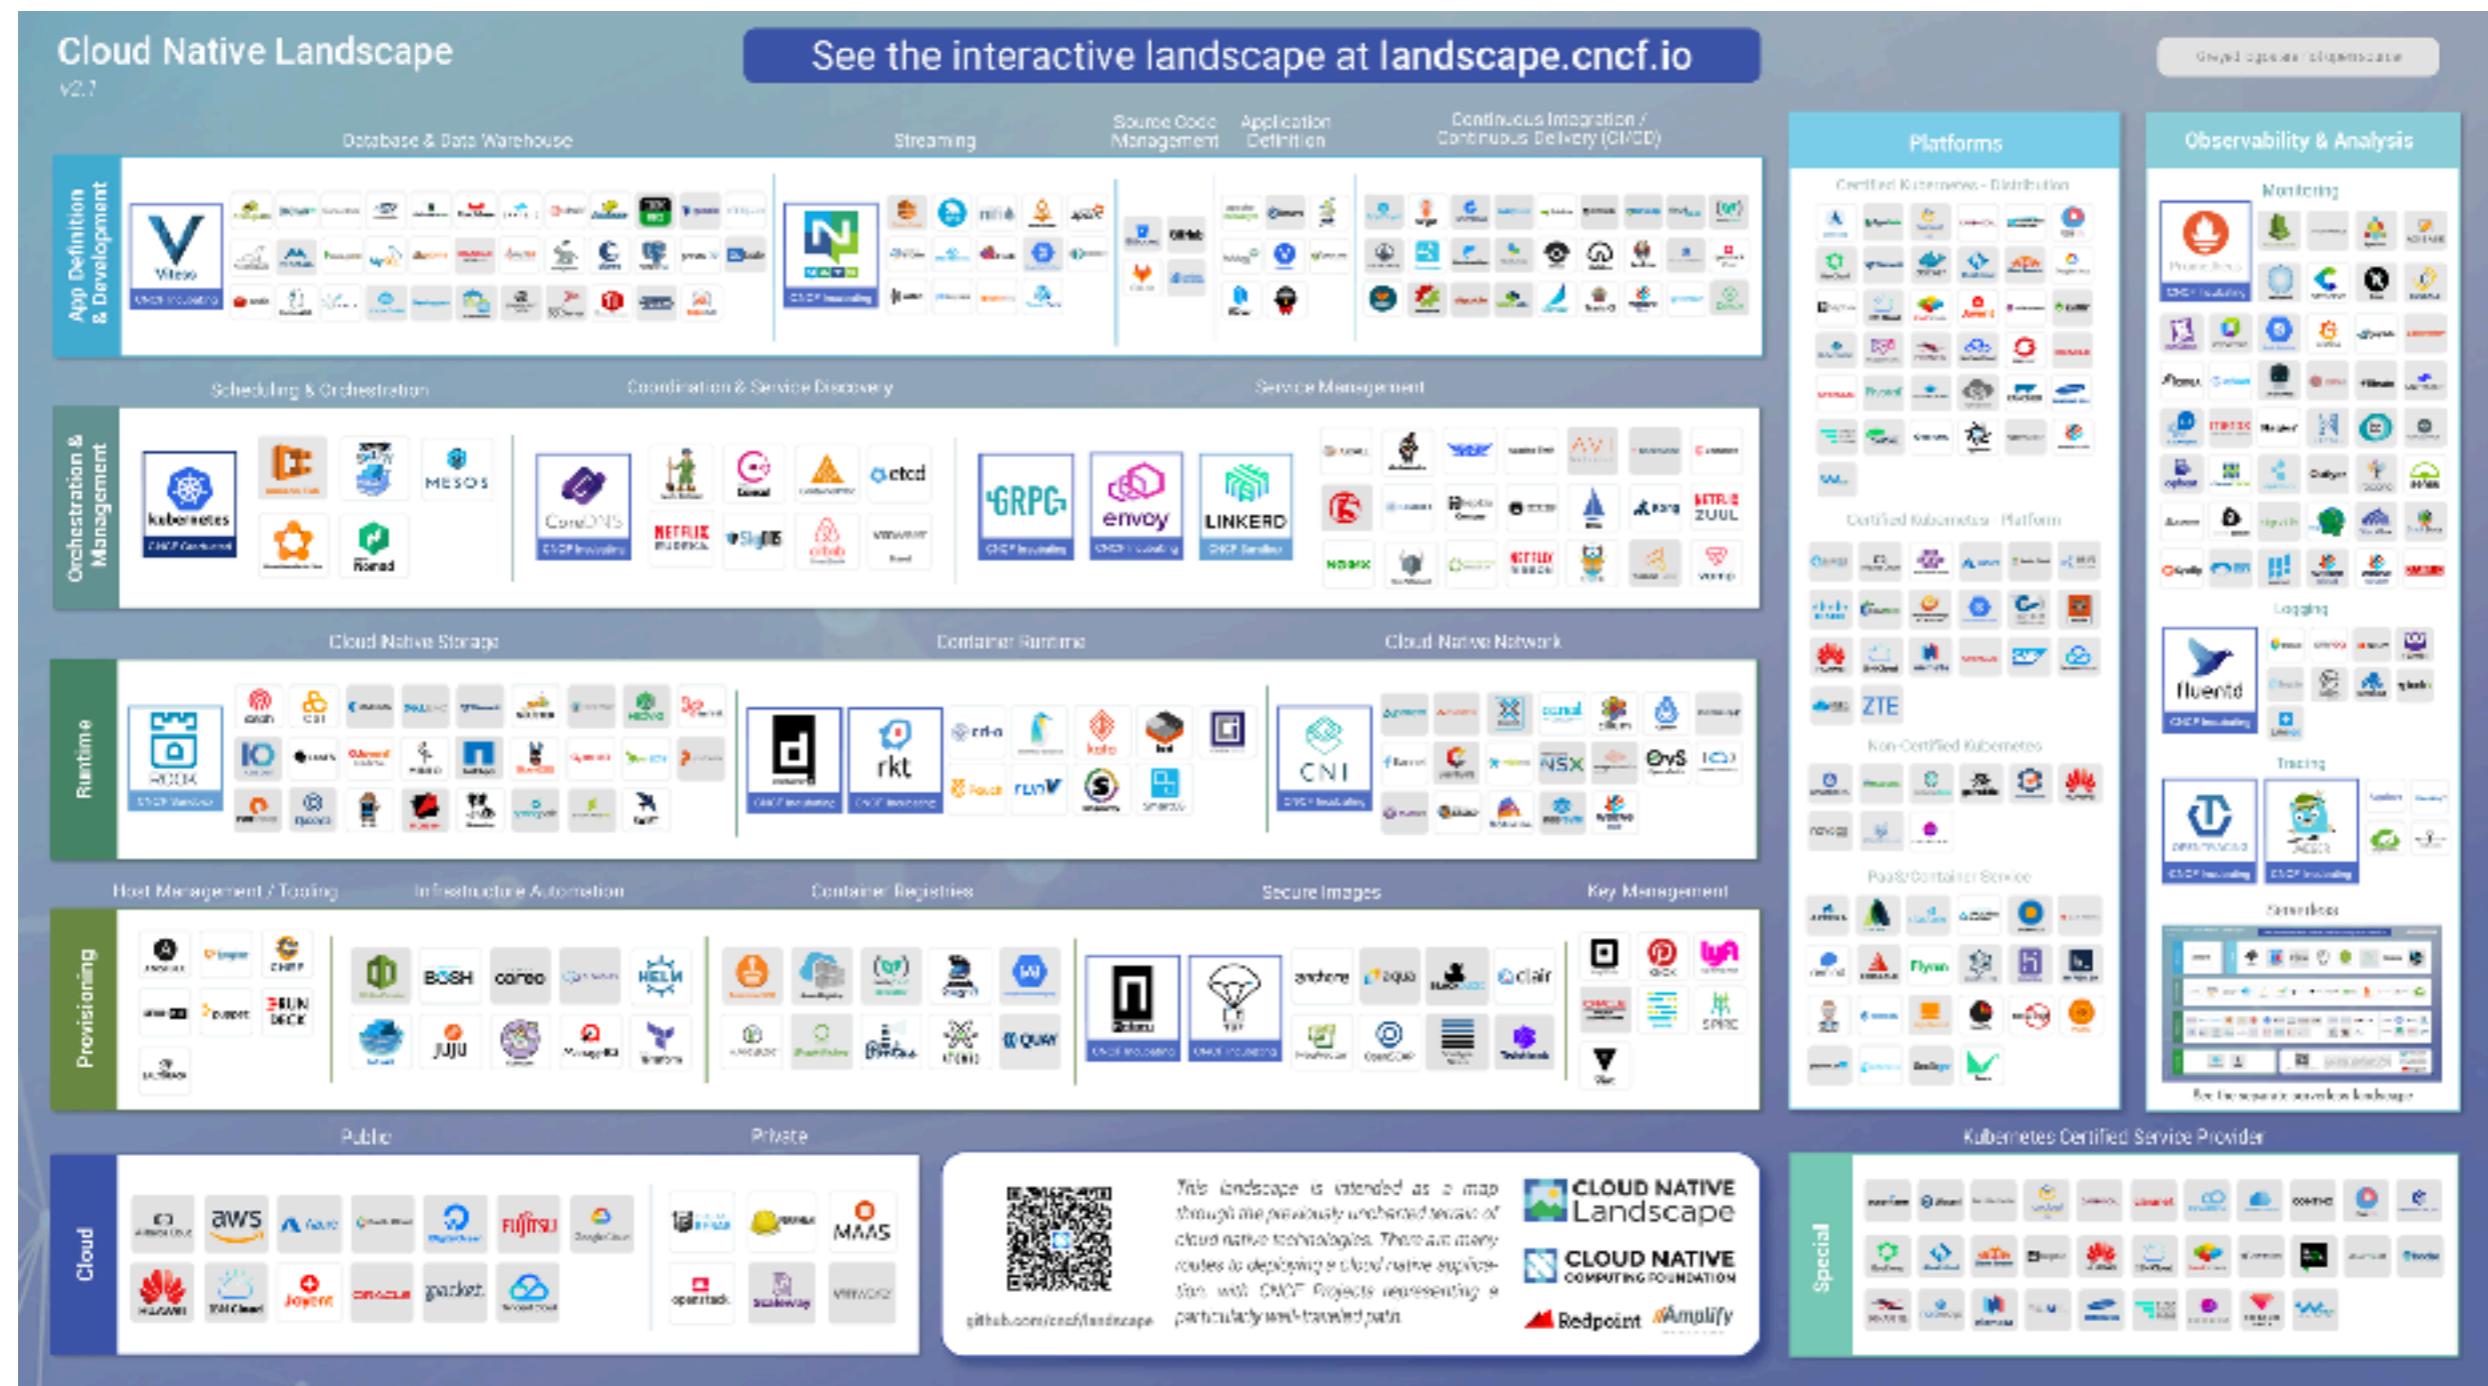
Single Concern Principle



Self-Containment Principle



# Cloud native landscape



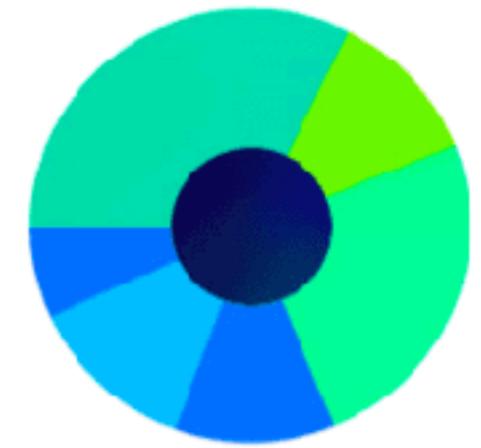
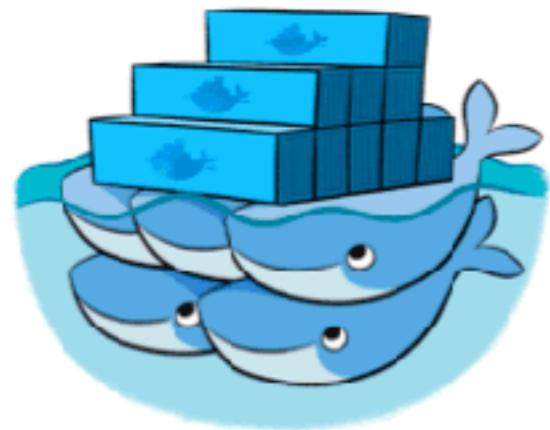
<https://github.com/cncf/landscape>







MESOS



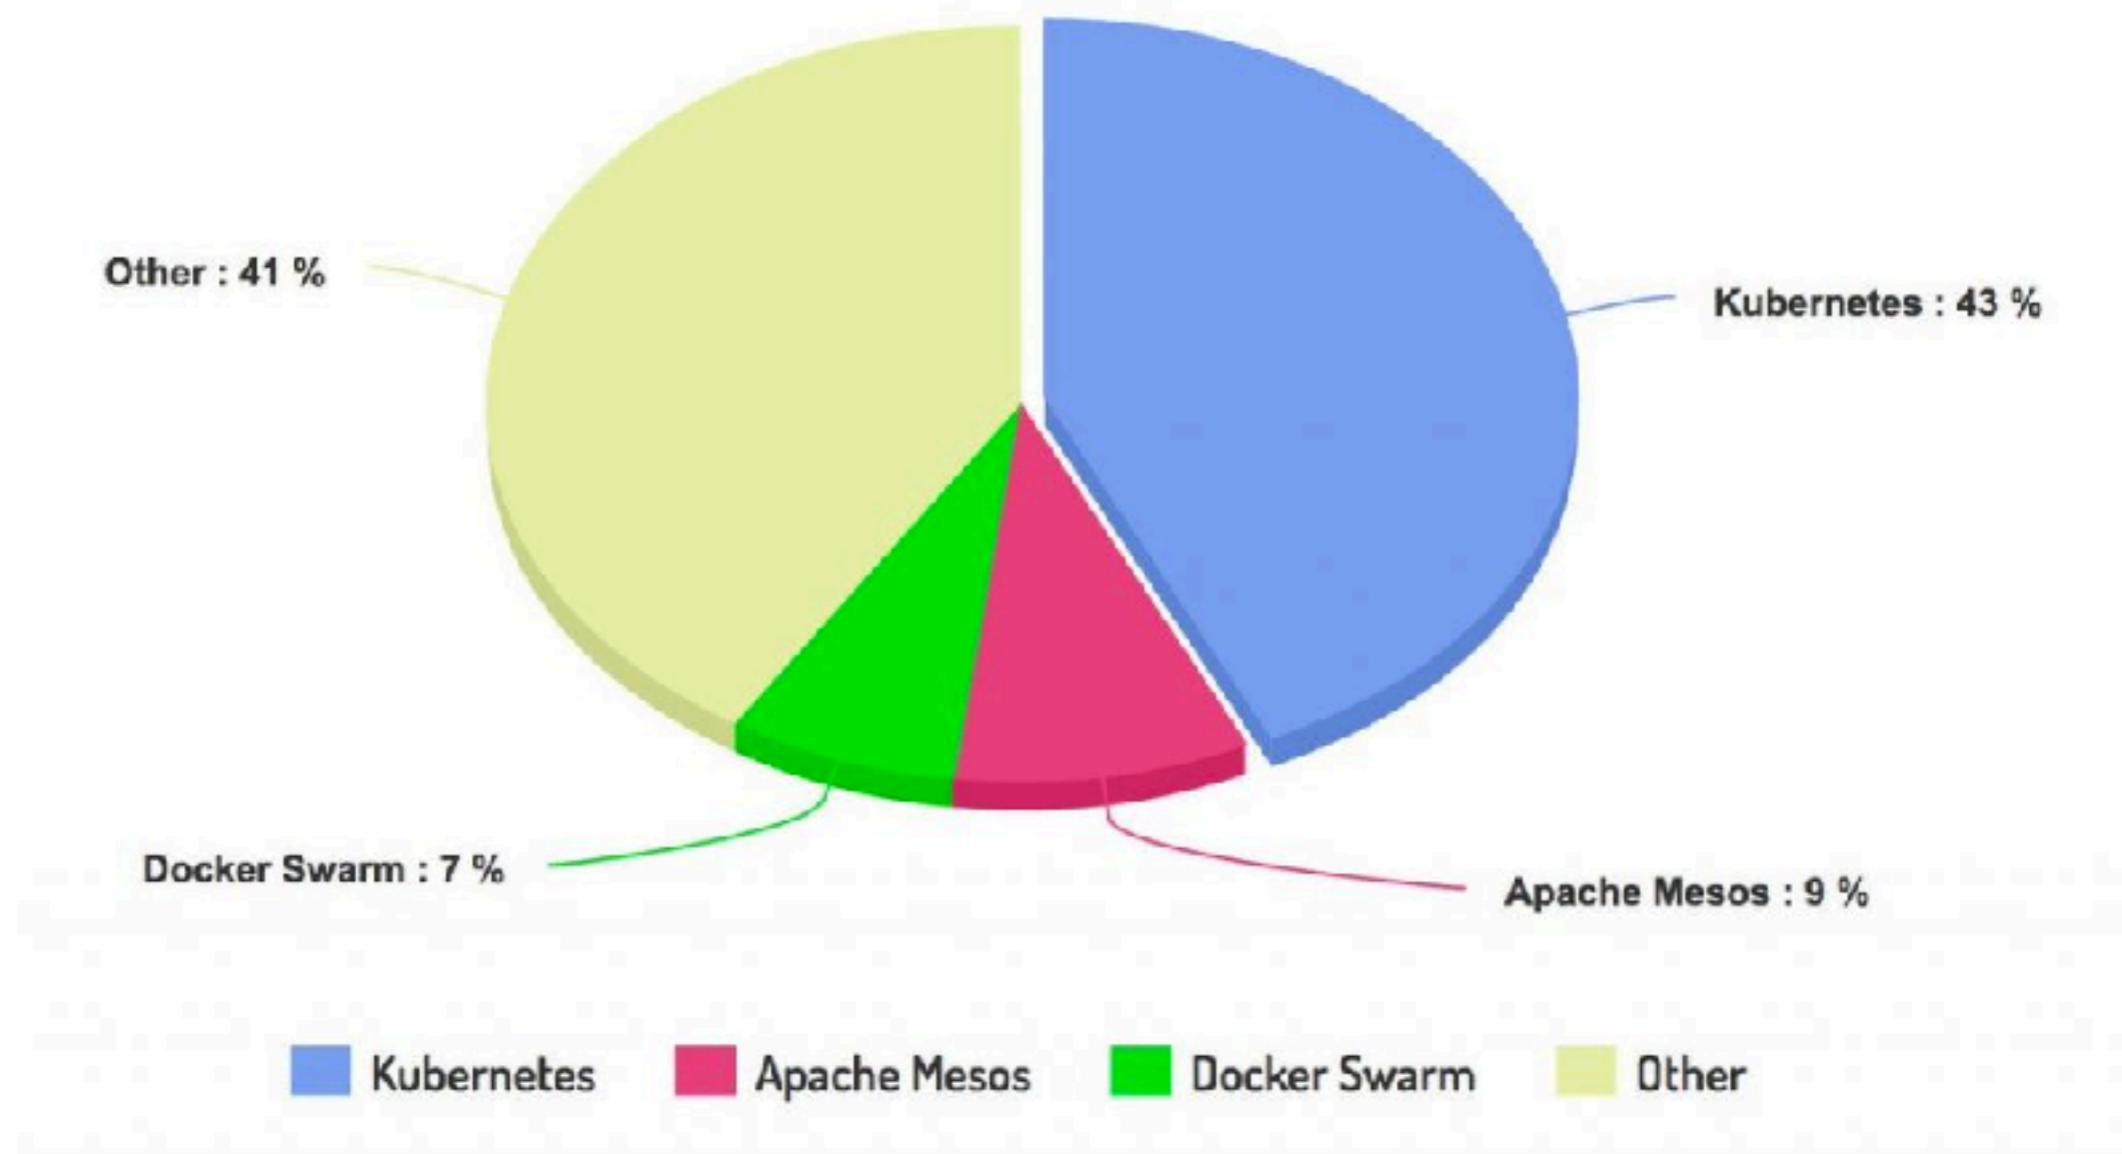
Marathon



kubernetes  
by Google



# Container Orchestrators in Sysdig's 2017 Docker Usage Report



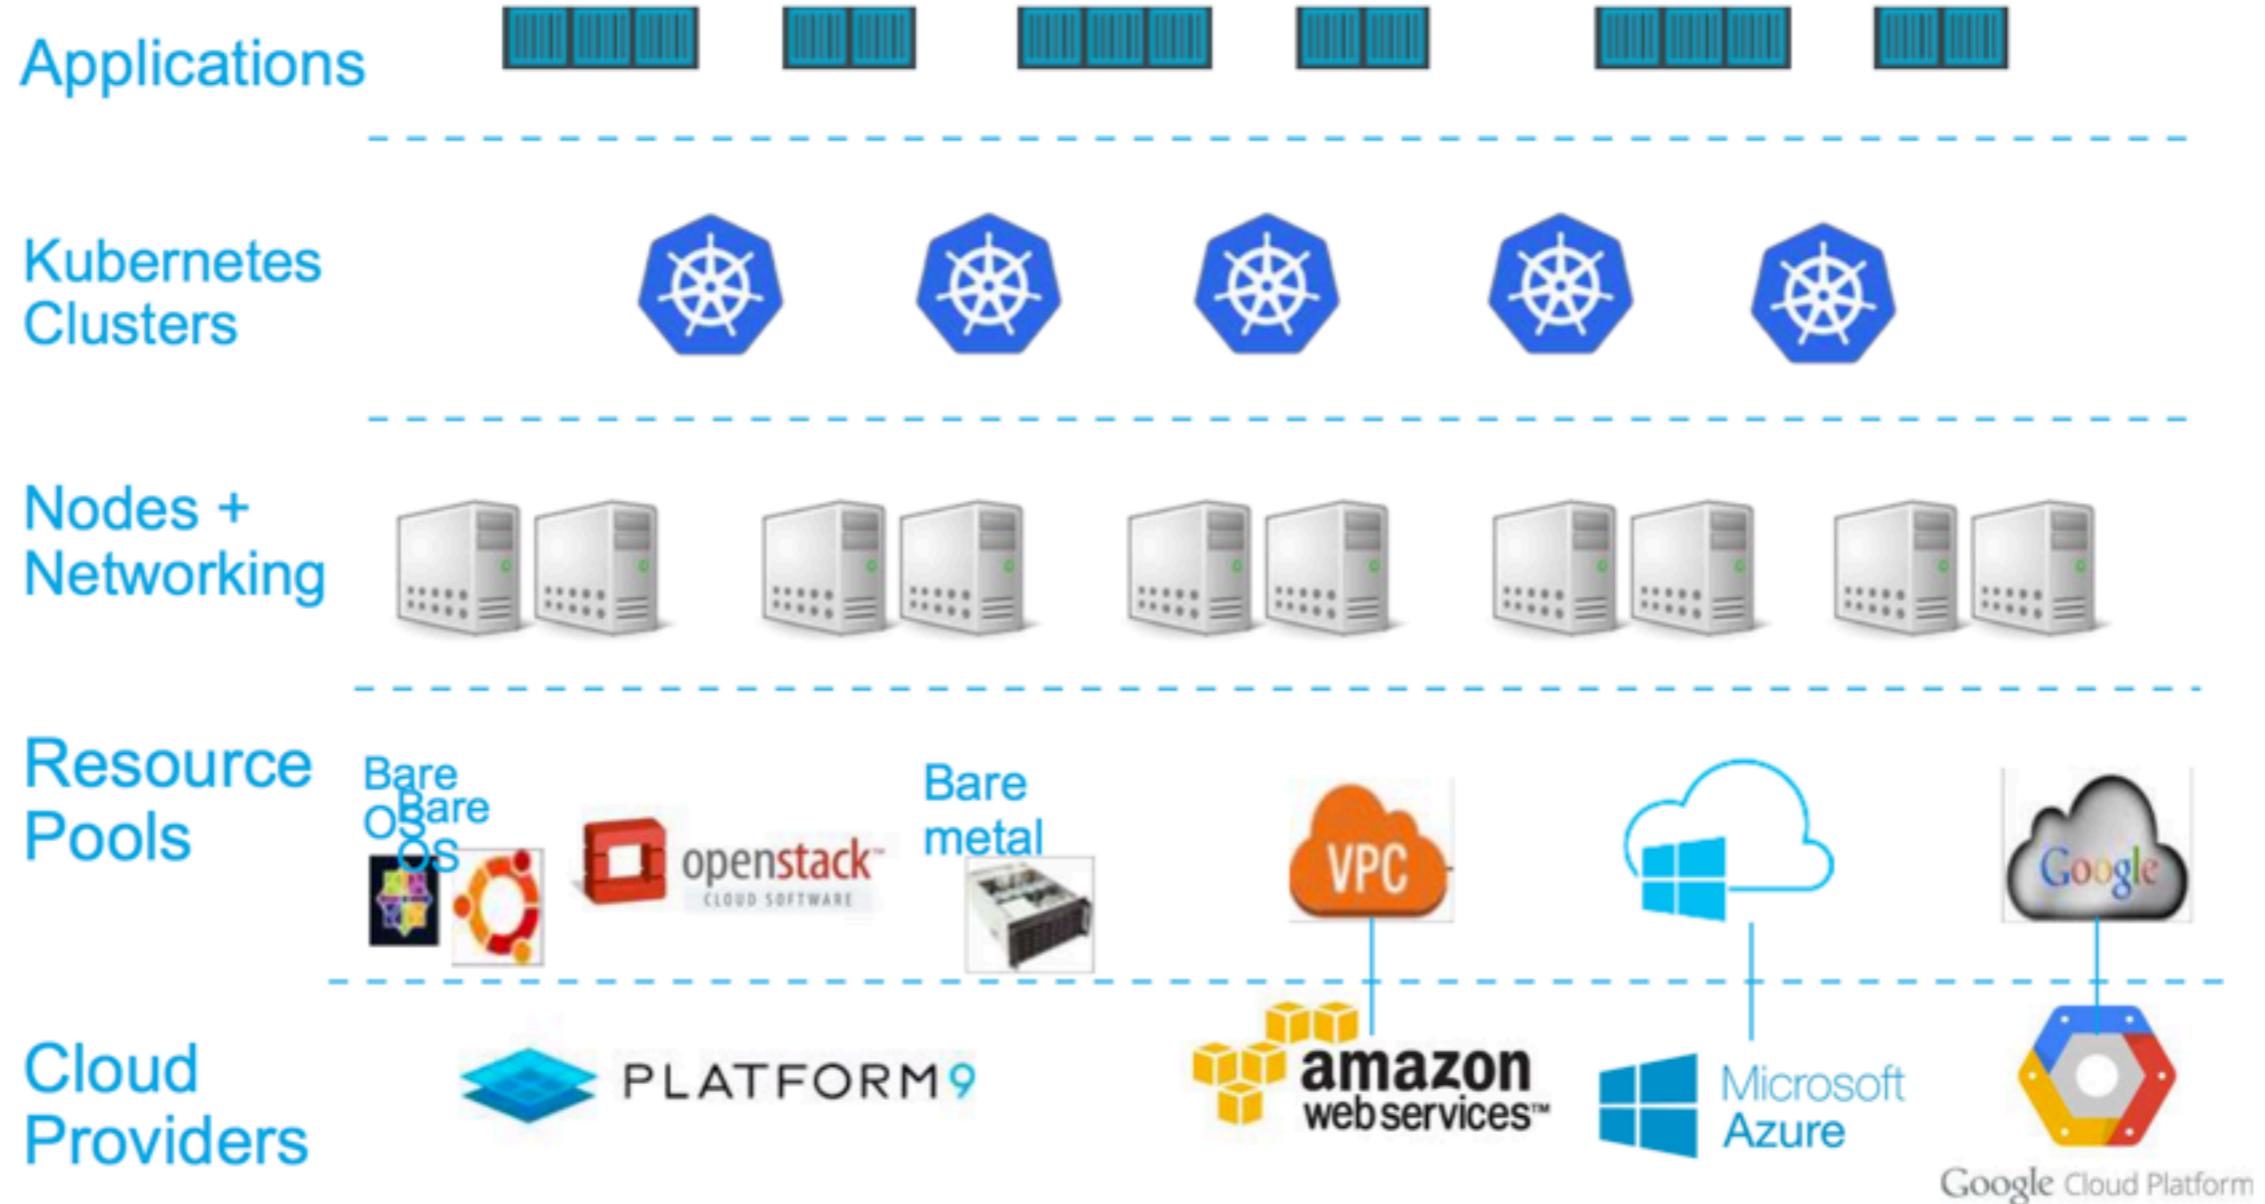
<https://sysdig.com/blog/sysdig-docker-usage-report-2017/>



# Why Kubernetes ?



# Write once, run anywhere



# Write once, run anywhere

Eliminate infrastructure lock-in

Use containers

Provides management for containers



# Modular app design

Monolithic app makes everything worse

- Larger teams slow thing down

- Spaghetti dependencies

Lack of ownership for sharing components

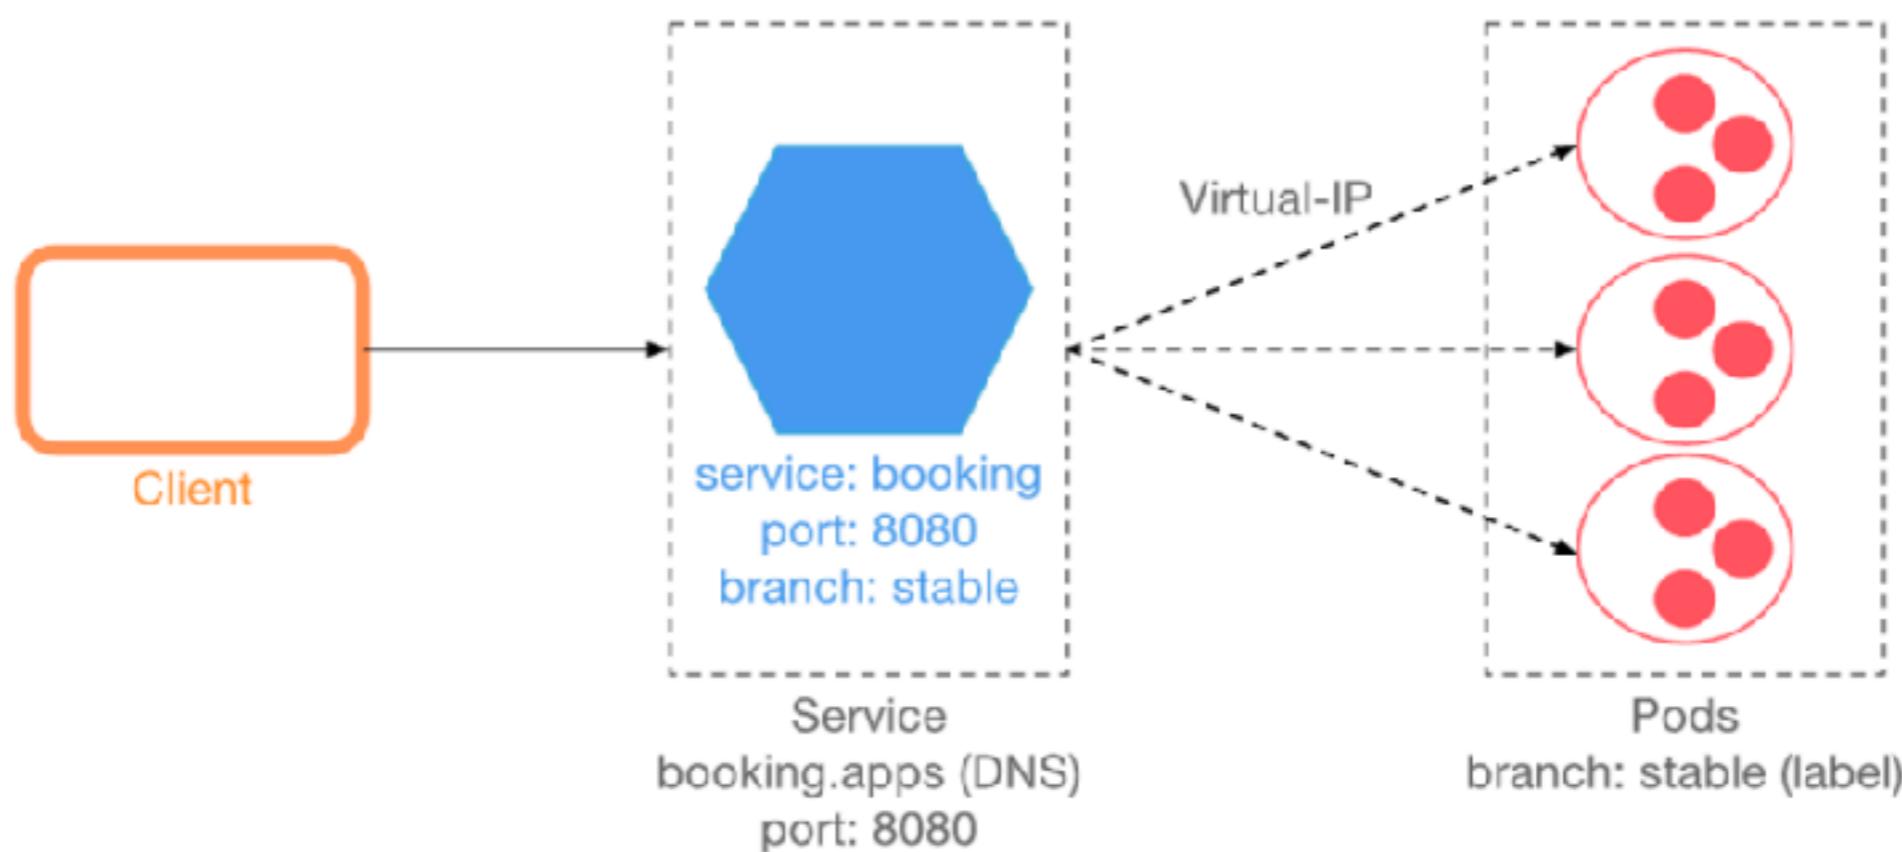
- Complexity to testing

- Slow building



# Modular app design

Container/Image boundary similar to class  
How to use/manage a collection of container ?



# Fault-tolerant by design

Design for failure

Infrastructure provisioning/re-provisioning

Configuration networking and load balance

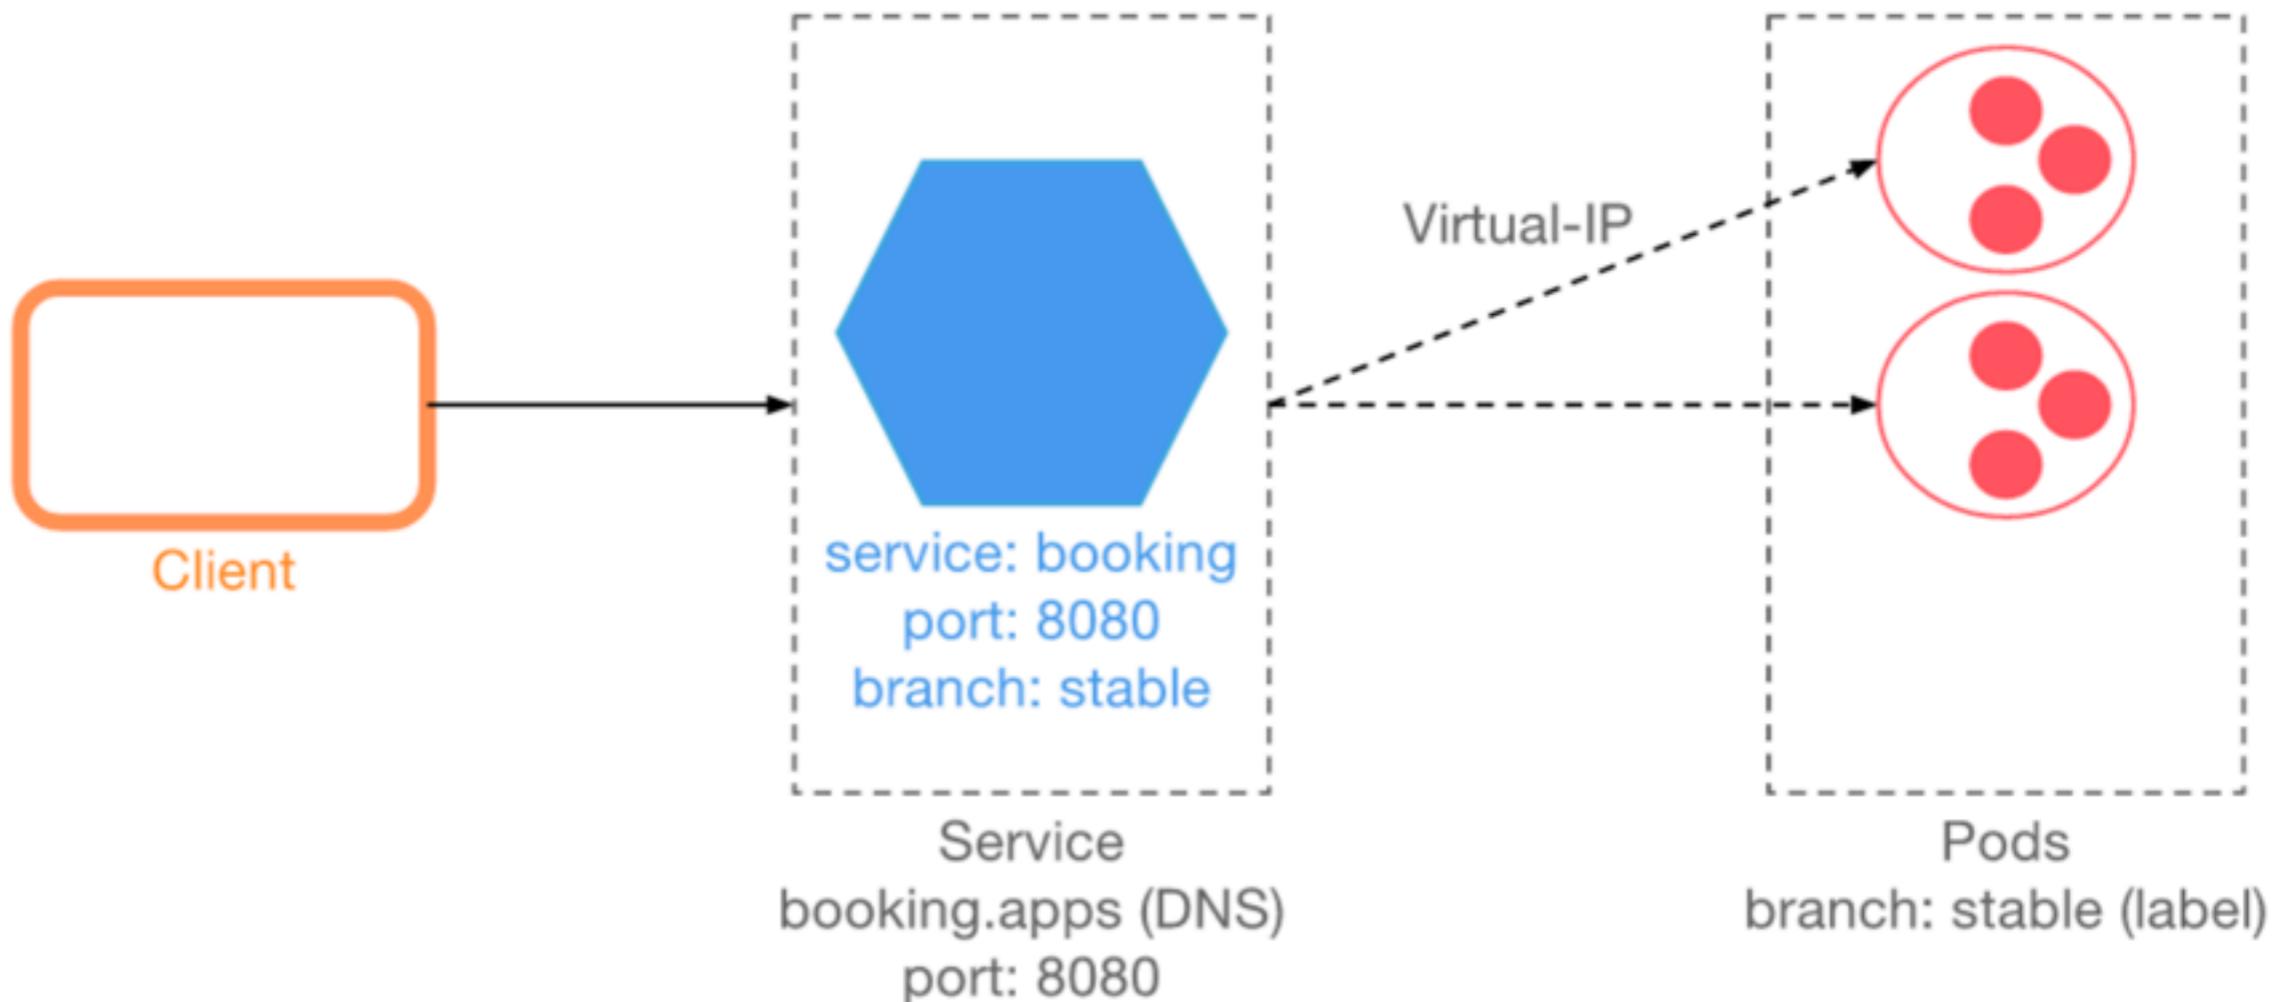
Redundancy (scale-out)

Lifecycle management (Software update)



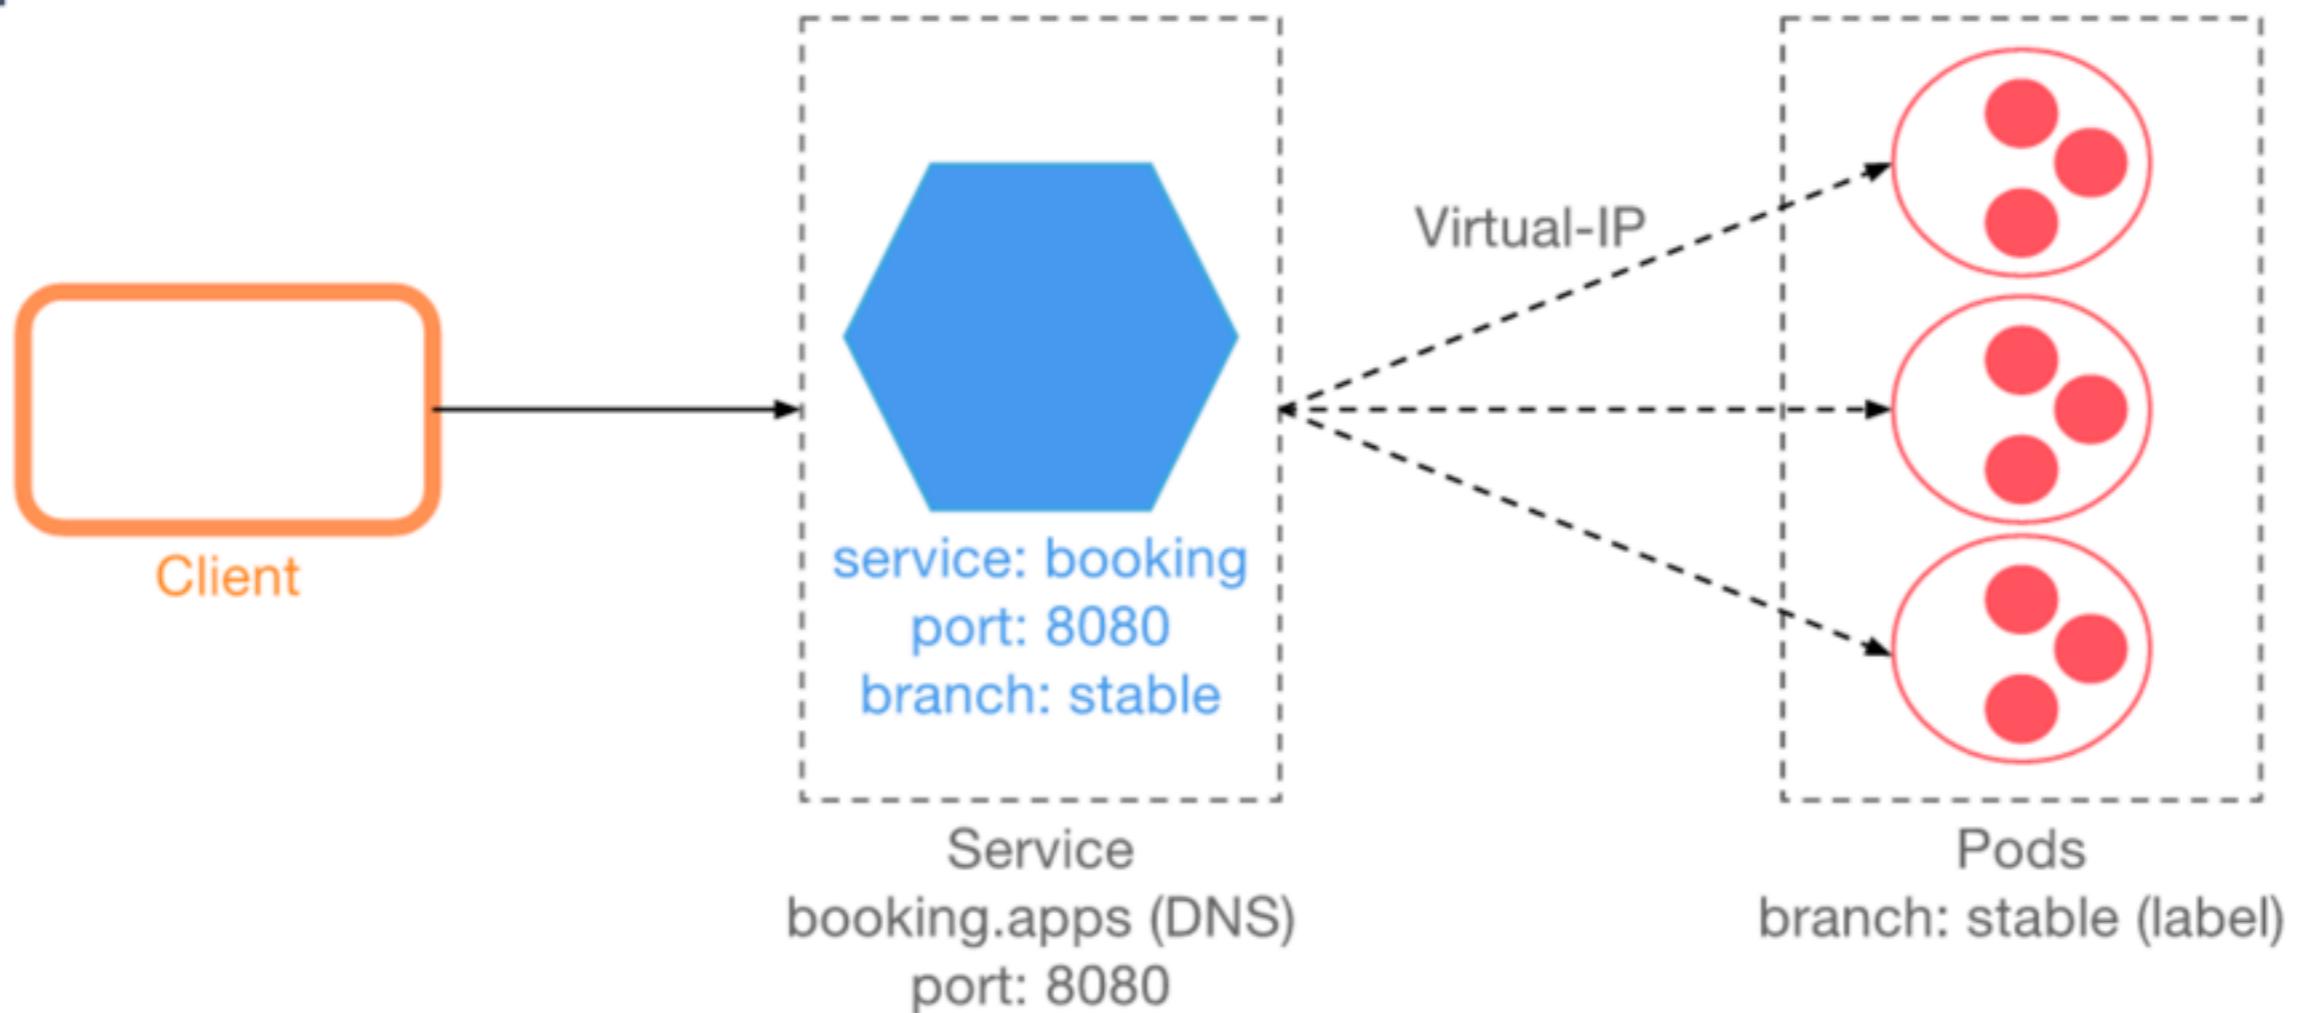
# Fault-tolerant by design

replicas = 2



# Fault-tolerant by design

replicas = 3



# Deployment, not Infrastructure

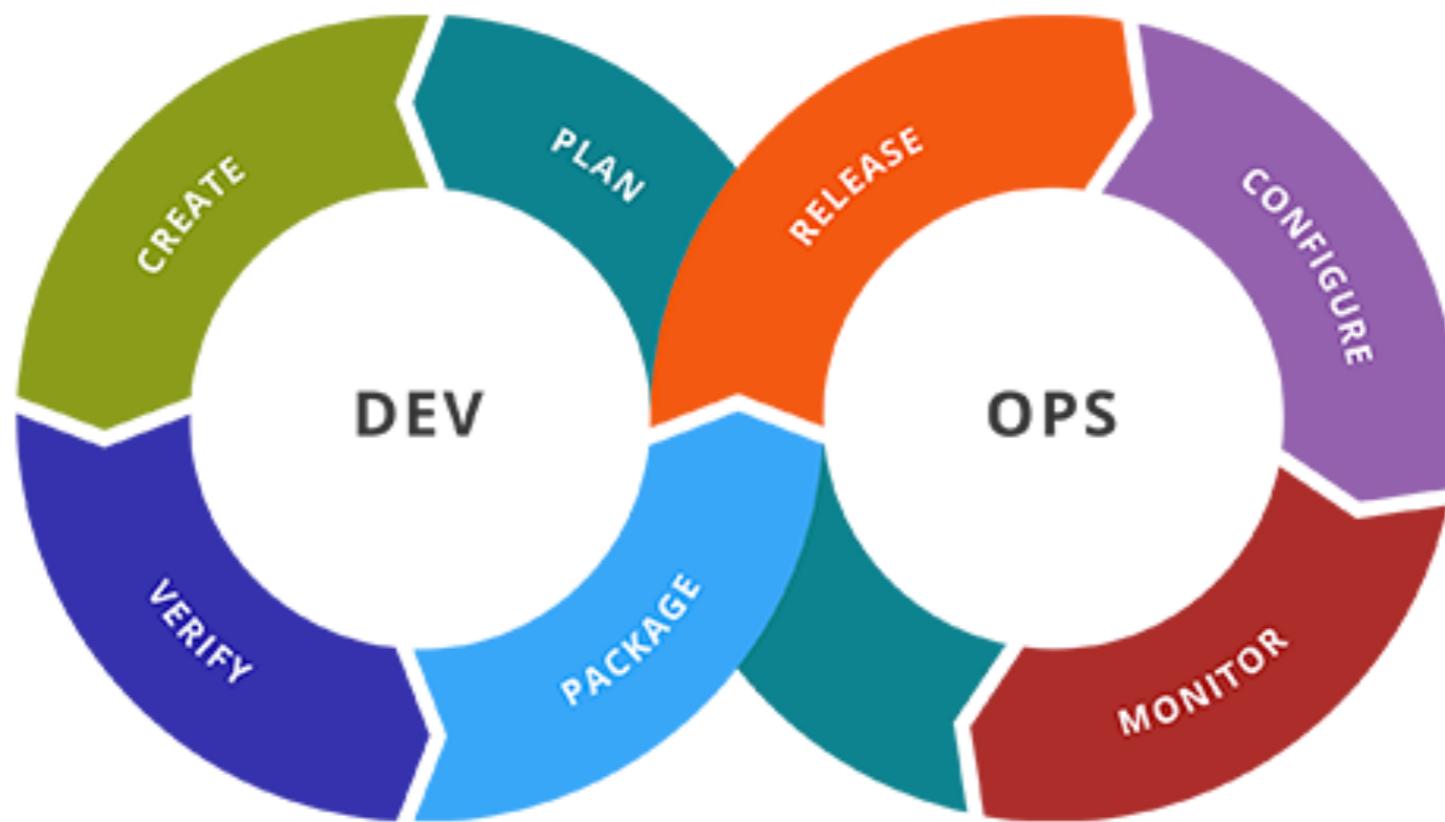
Software deployment is hard  
Infrastructure provisioning/re-provisioning  
Configuration networking and load balance  
Redundancy (scale-out)  
Lifecycle management (Software update)



# Deployment, not Infrastructure

Kubernetes support for deployment  
Controllers are in focus  
Scale-out service  
Rolling update for new version  
Rollback to a previous version  
Pause and resume a deployment  
Horizontal auto-scaling  
Canary deployment





# Let's start !!



# Installation



# Software requirement

minkube

Kubernetes command-line tool



# Starting minikube

```
$minikube start  
$minikube status
```



# Version of Kubernetes

```
$kubectl get nodes -o yaml
```

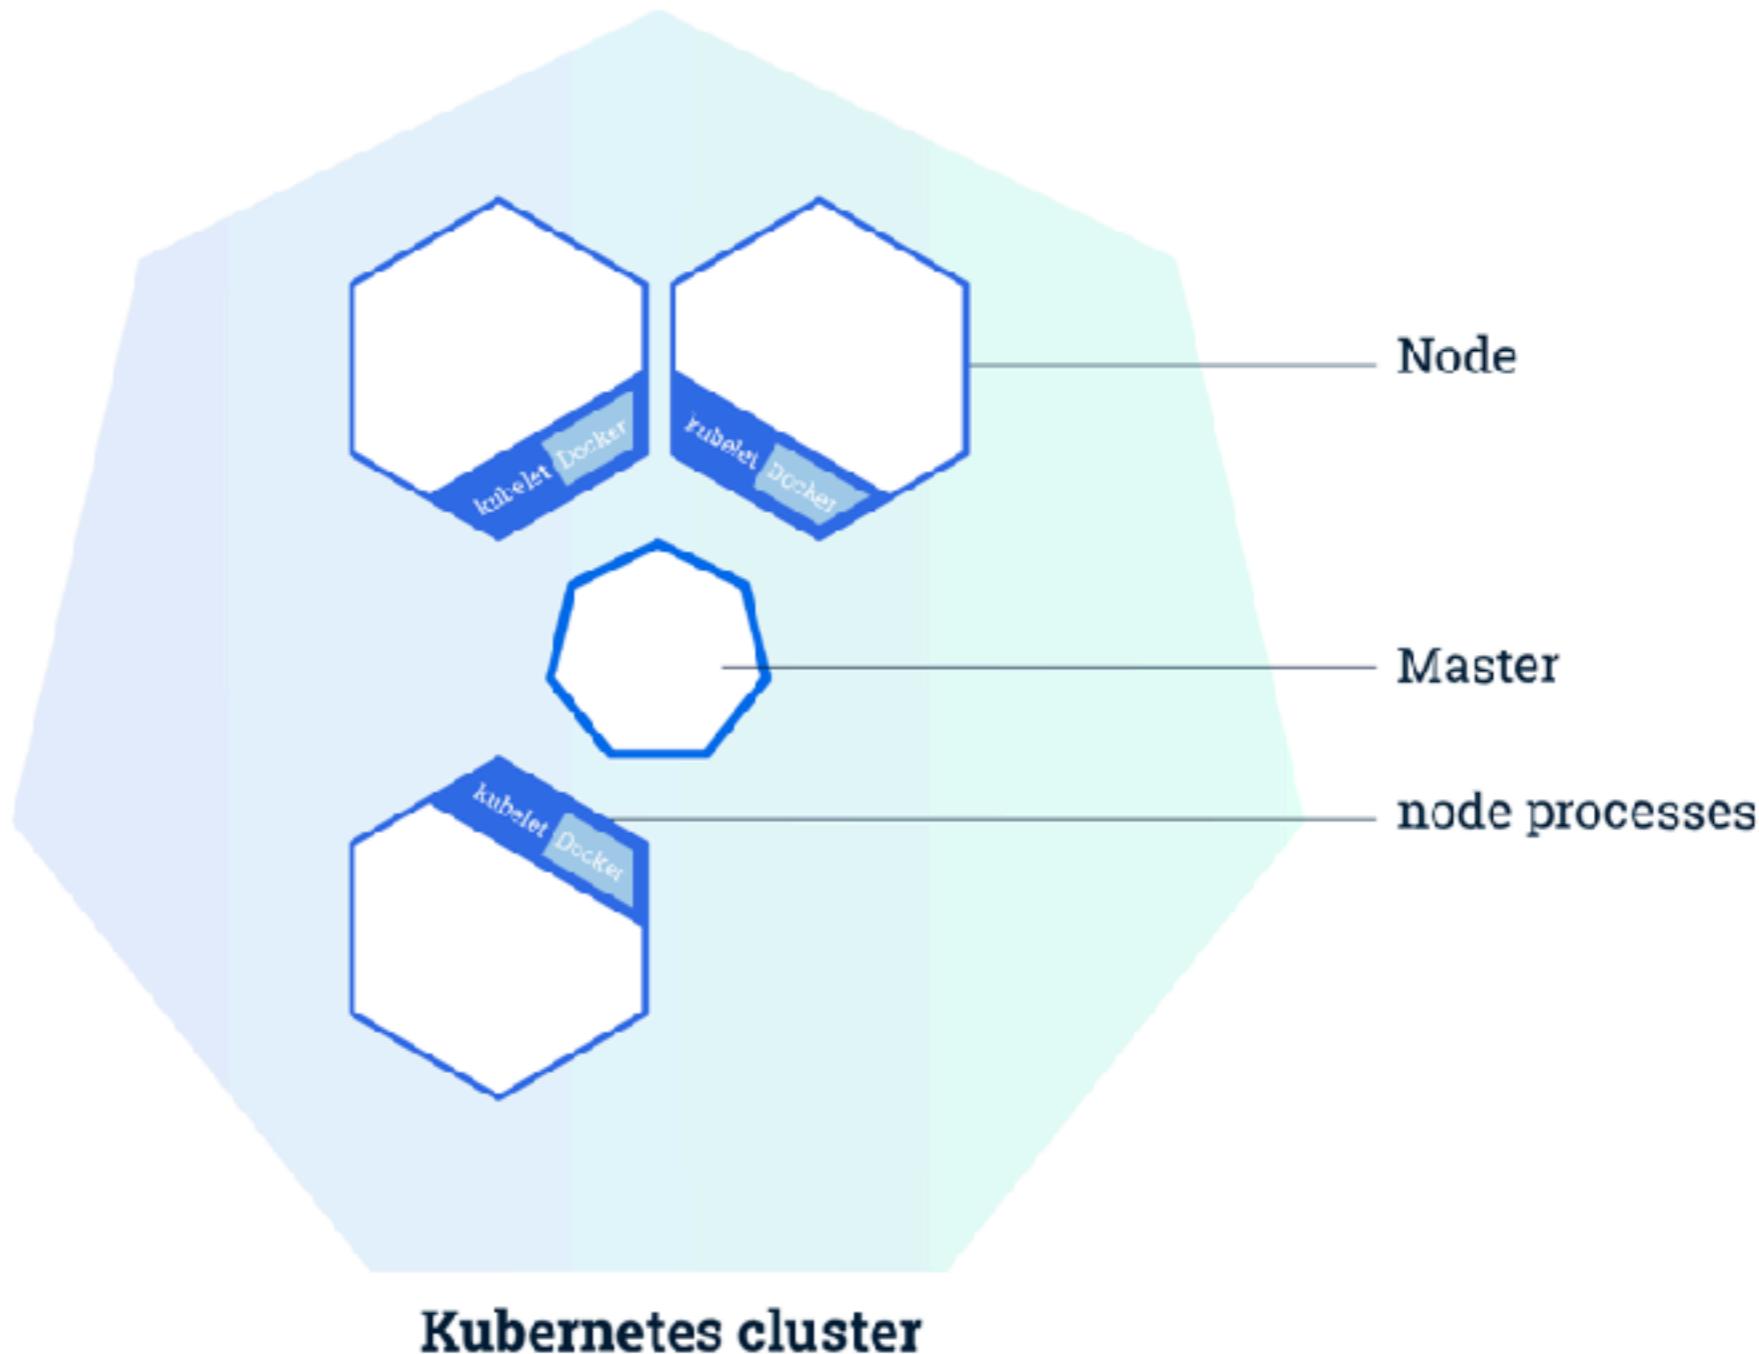
```
nodeInfo:  
  architecture: amd64  
  bootID: 7c769465-27e5-4dd8-a89f-319ba1b8ef57  
  containerRuntimeVersion: docker://17.9.0  
  kernelVersion: 4.9.64  
  kubeProxyVersion: v1.9.4  
  kubeletVersion: v1.9.4  
  machineID: edffd4ca8bf24253a736ea86d2448185  
  operatingSystem: linux  
  osImage: Buildroot 2017.11  
  systemUUID: 98BAE0EF-8C9E-45DF-9AC7-F0E05806B189
```



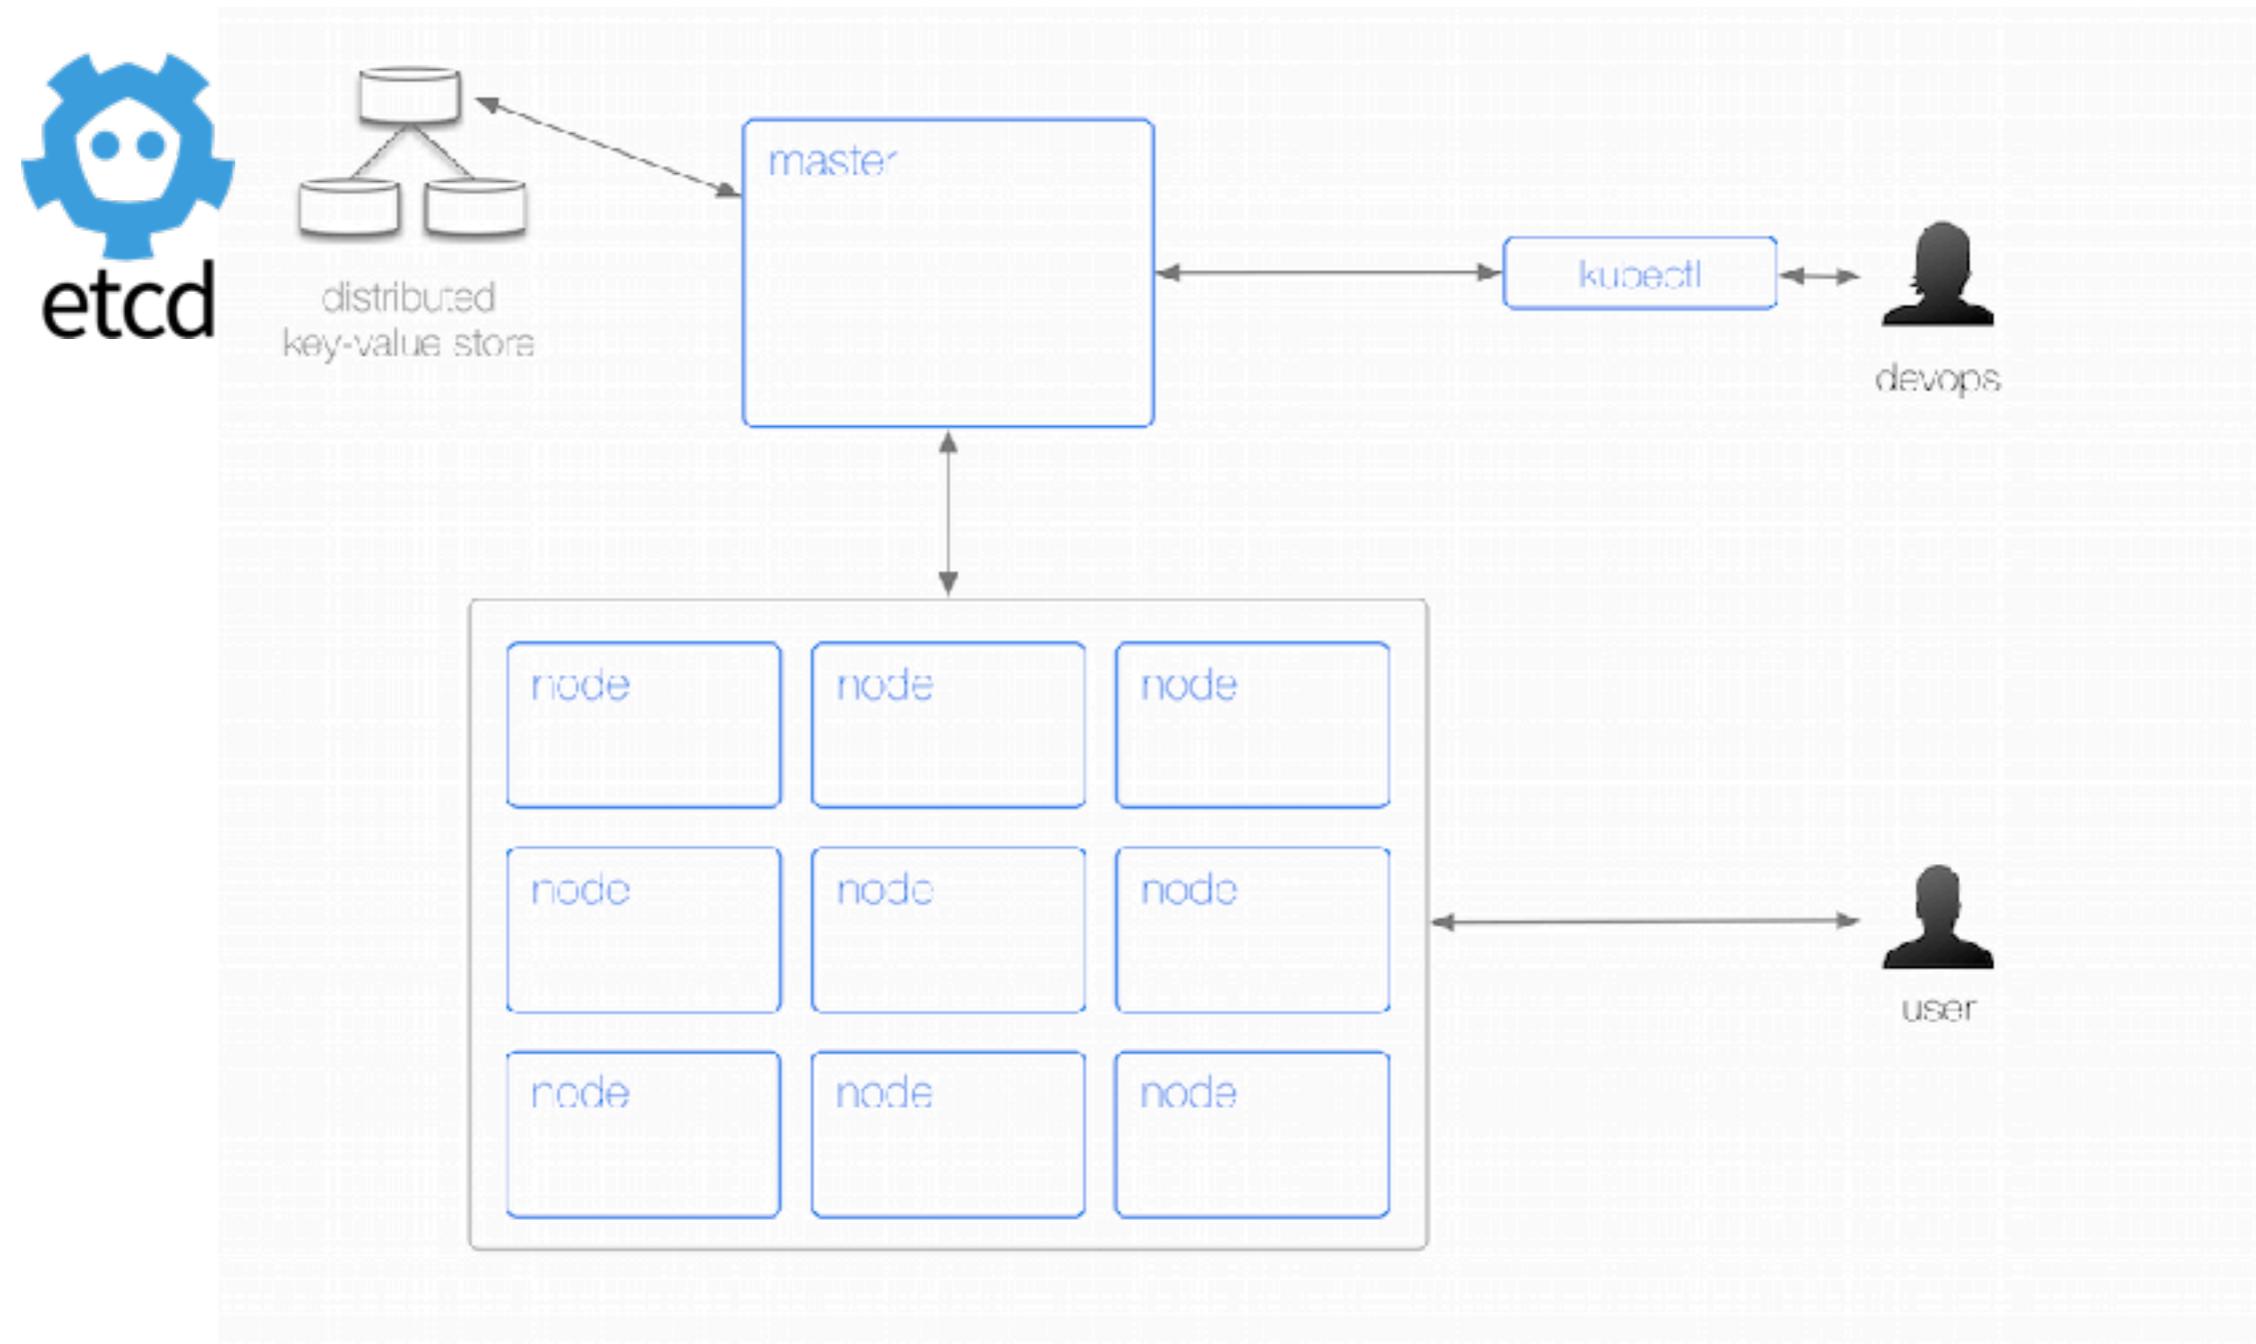
# Basic of Kubernetes



# Kubernetes cluster



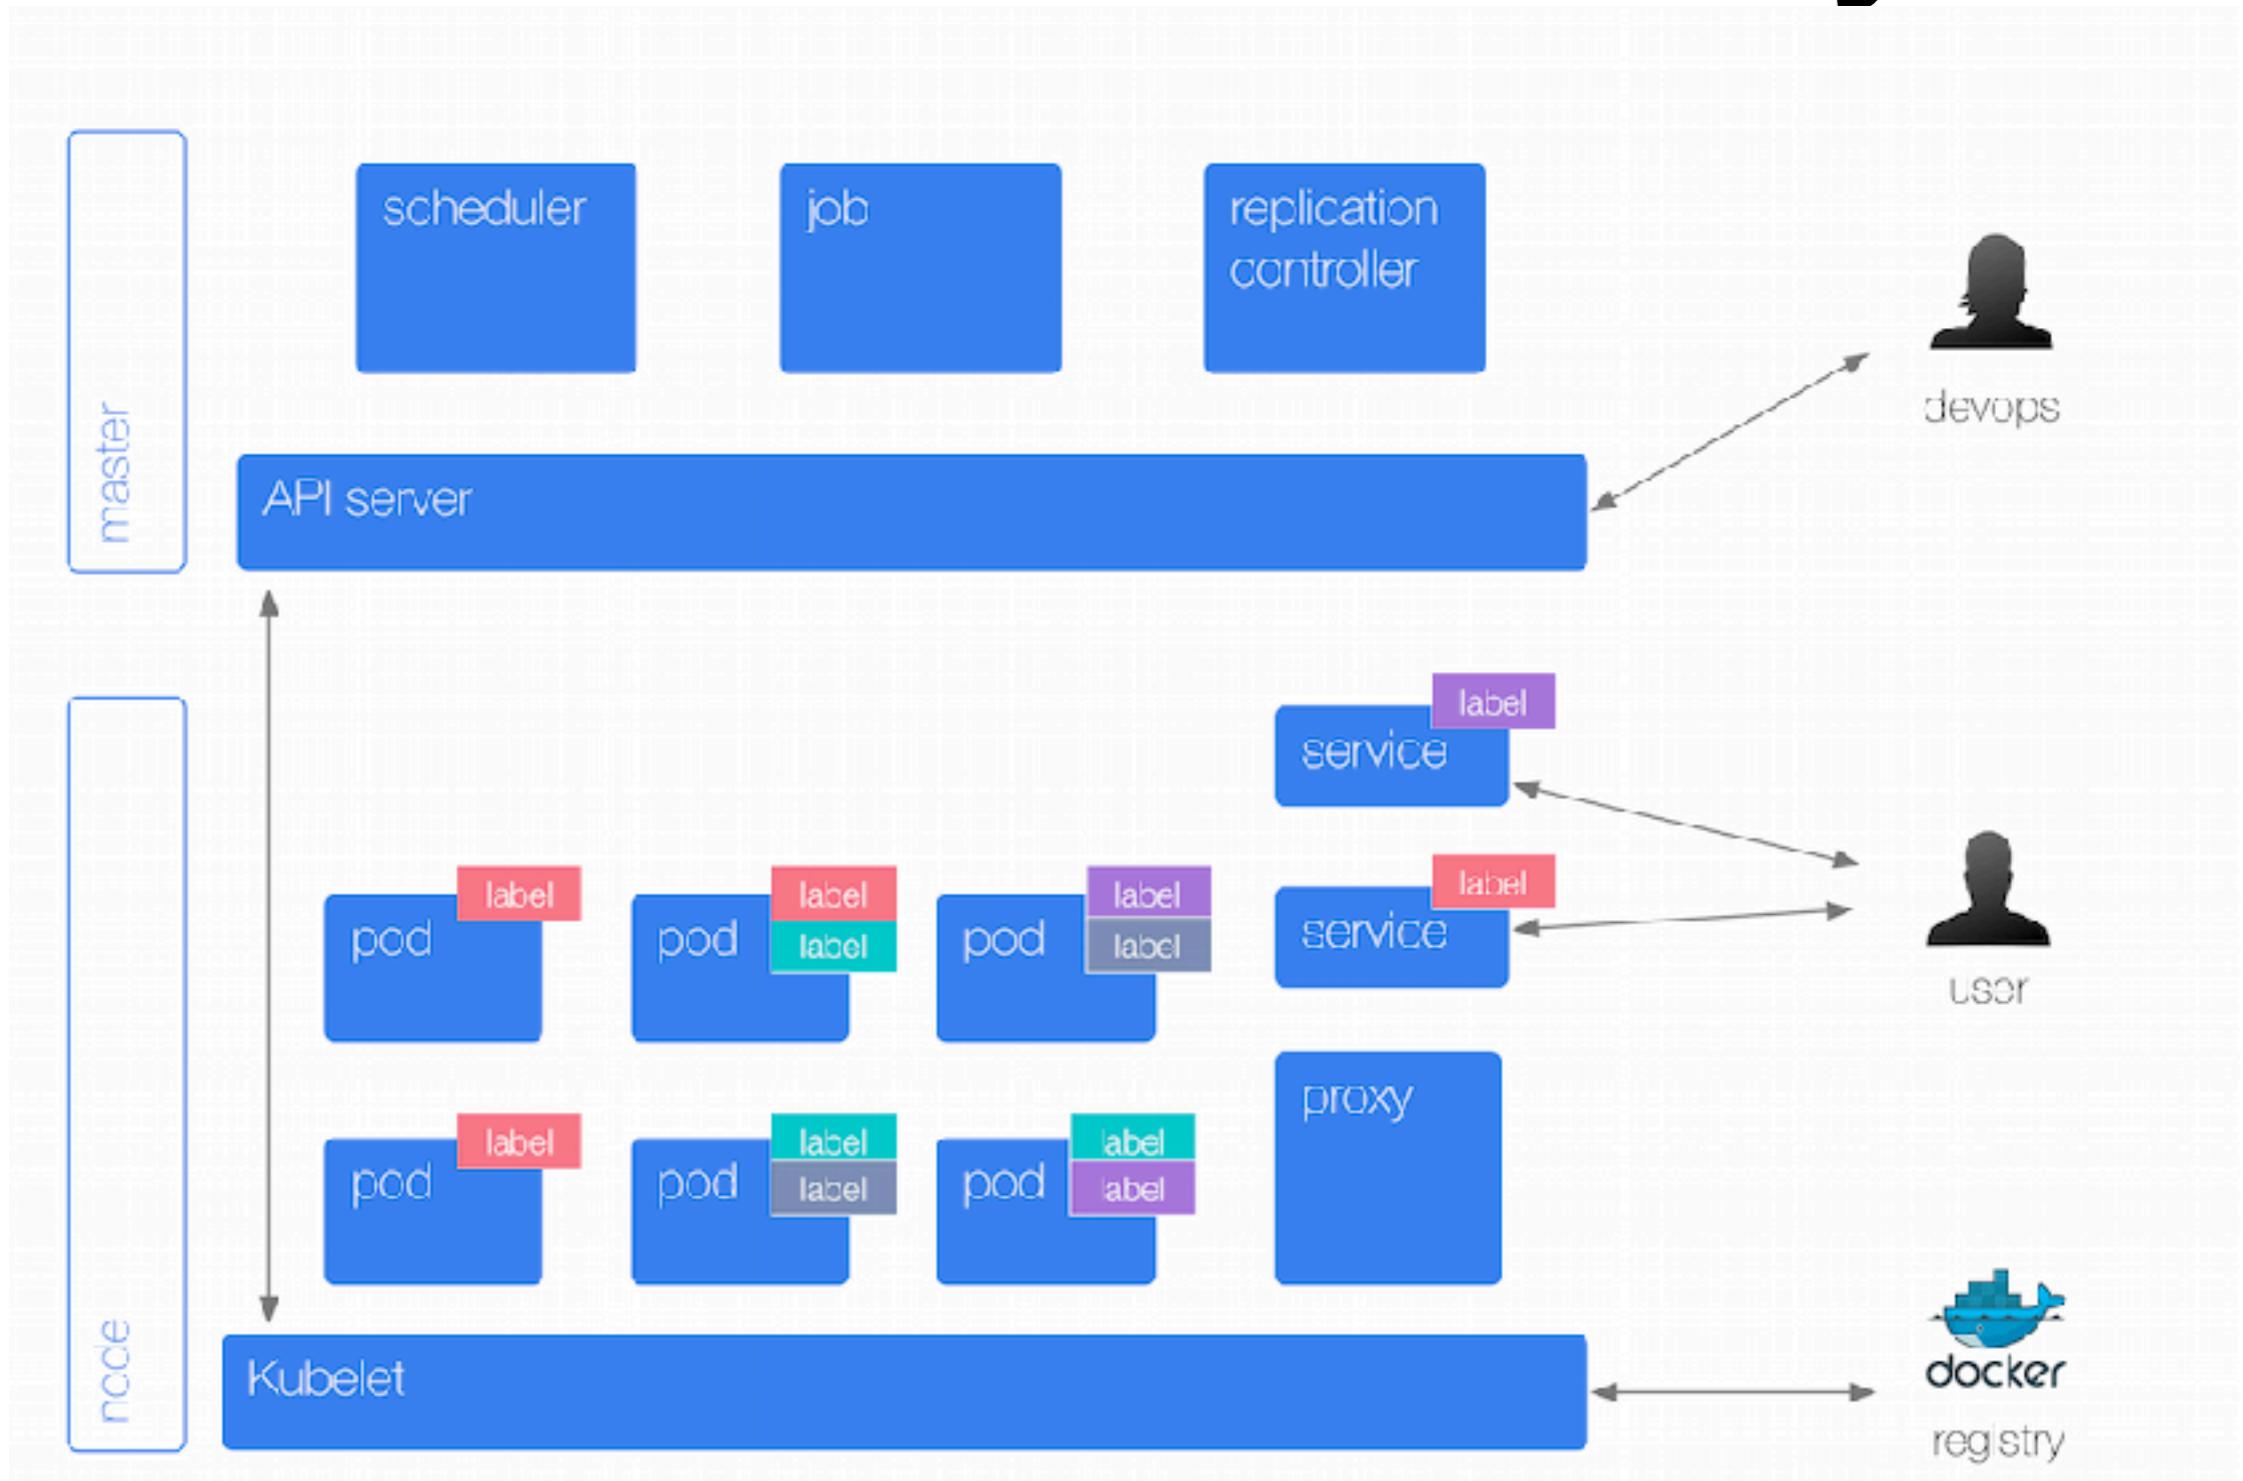
# Kubernetes physical layout



<http://k8s.info/cs.html>



# Kubernetes abstraction layout



<http://k8s.info/cs.html>



# Hello with Kubernetes

file /01-hello/instruction.txt

file /01-hello/instruction2.txt



# Hello Kubernetes

Working with **kubectl** in command line  
Try to run image from Docker hub



# Try to use somkiat/hello

The screenshot shows the Docker Hub interface for the repository `somkiat/hello`. At the top, there's a dark header bar with a search bar, a dashboard icon, and navigation links for Dashboard, Explore, Organizations, Create, and a user profile for `somkiat`. Below the header, the repository name `somkiat/hello` is displayed with a star icon, indicating it's a public repository. A note says "Last pushed: 2 months ago". There are tabs for Repo Info, Tags, Collaborators, Webhooks, and Settings, with the Tags tab currently selected. A table lists three tags: `v3`, `v2`, and `latest`, each with a compressed size of 5 MB and a last updated date of 2 months ago. Each tag row has a delete icon on the right.

Tag Name	Compressed Size	Last Updated
v3	5 MB	2 months ago
v2	5 MB	2 months ago
latest	5 MB	2 months ago

<https://hub.docker.com/r/somkiat/hello>



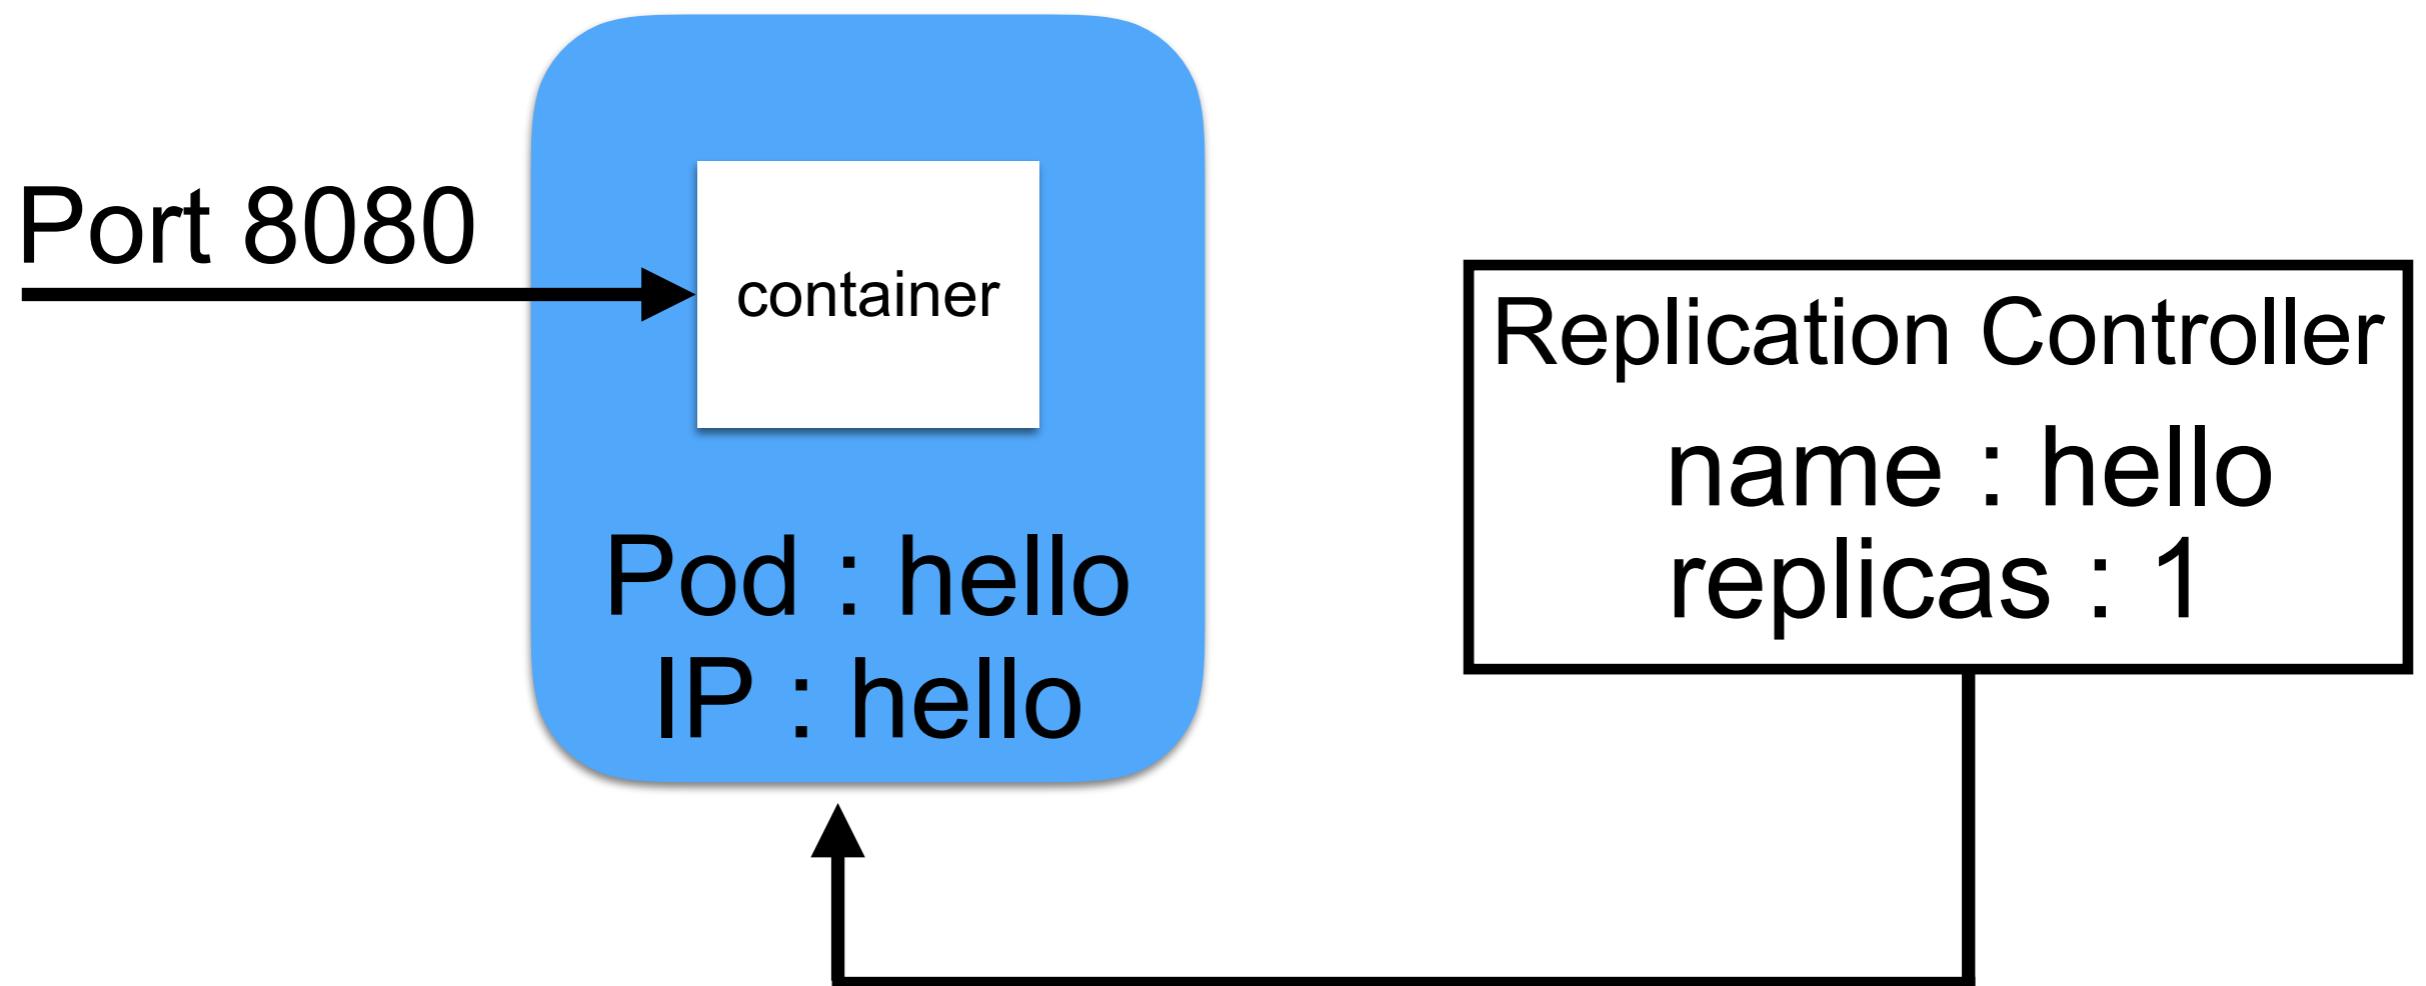
# Create a container

```
$kubectl run hello --image=somkiat/hello  
--port=8080 --generator=run/v1
```

<https://kubernetes.io/docs/reference/kubectl/conventions/#generators>



# Create a container

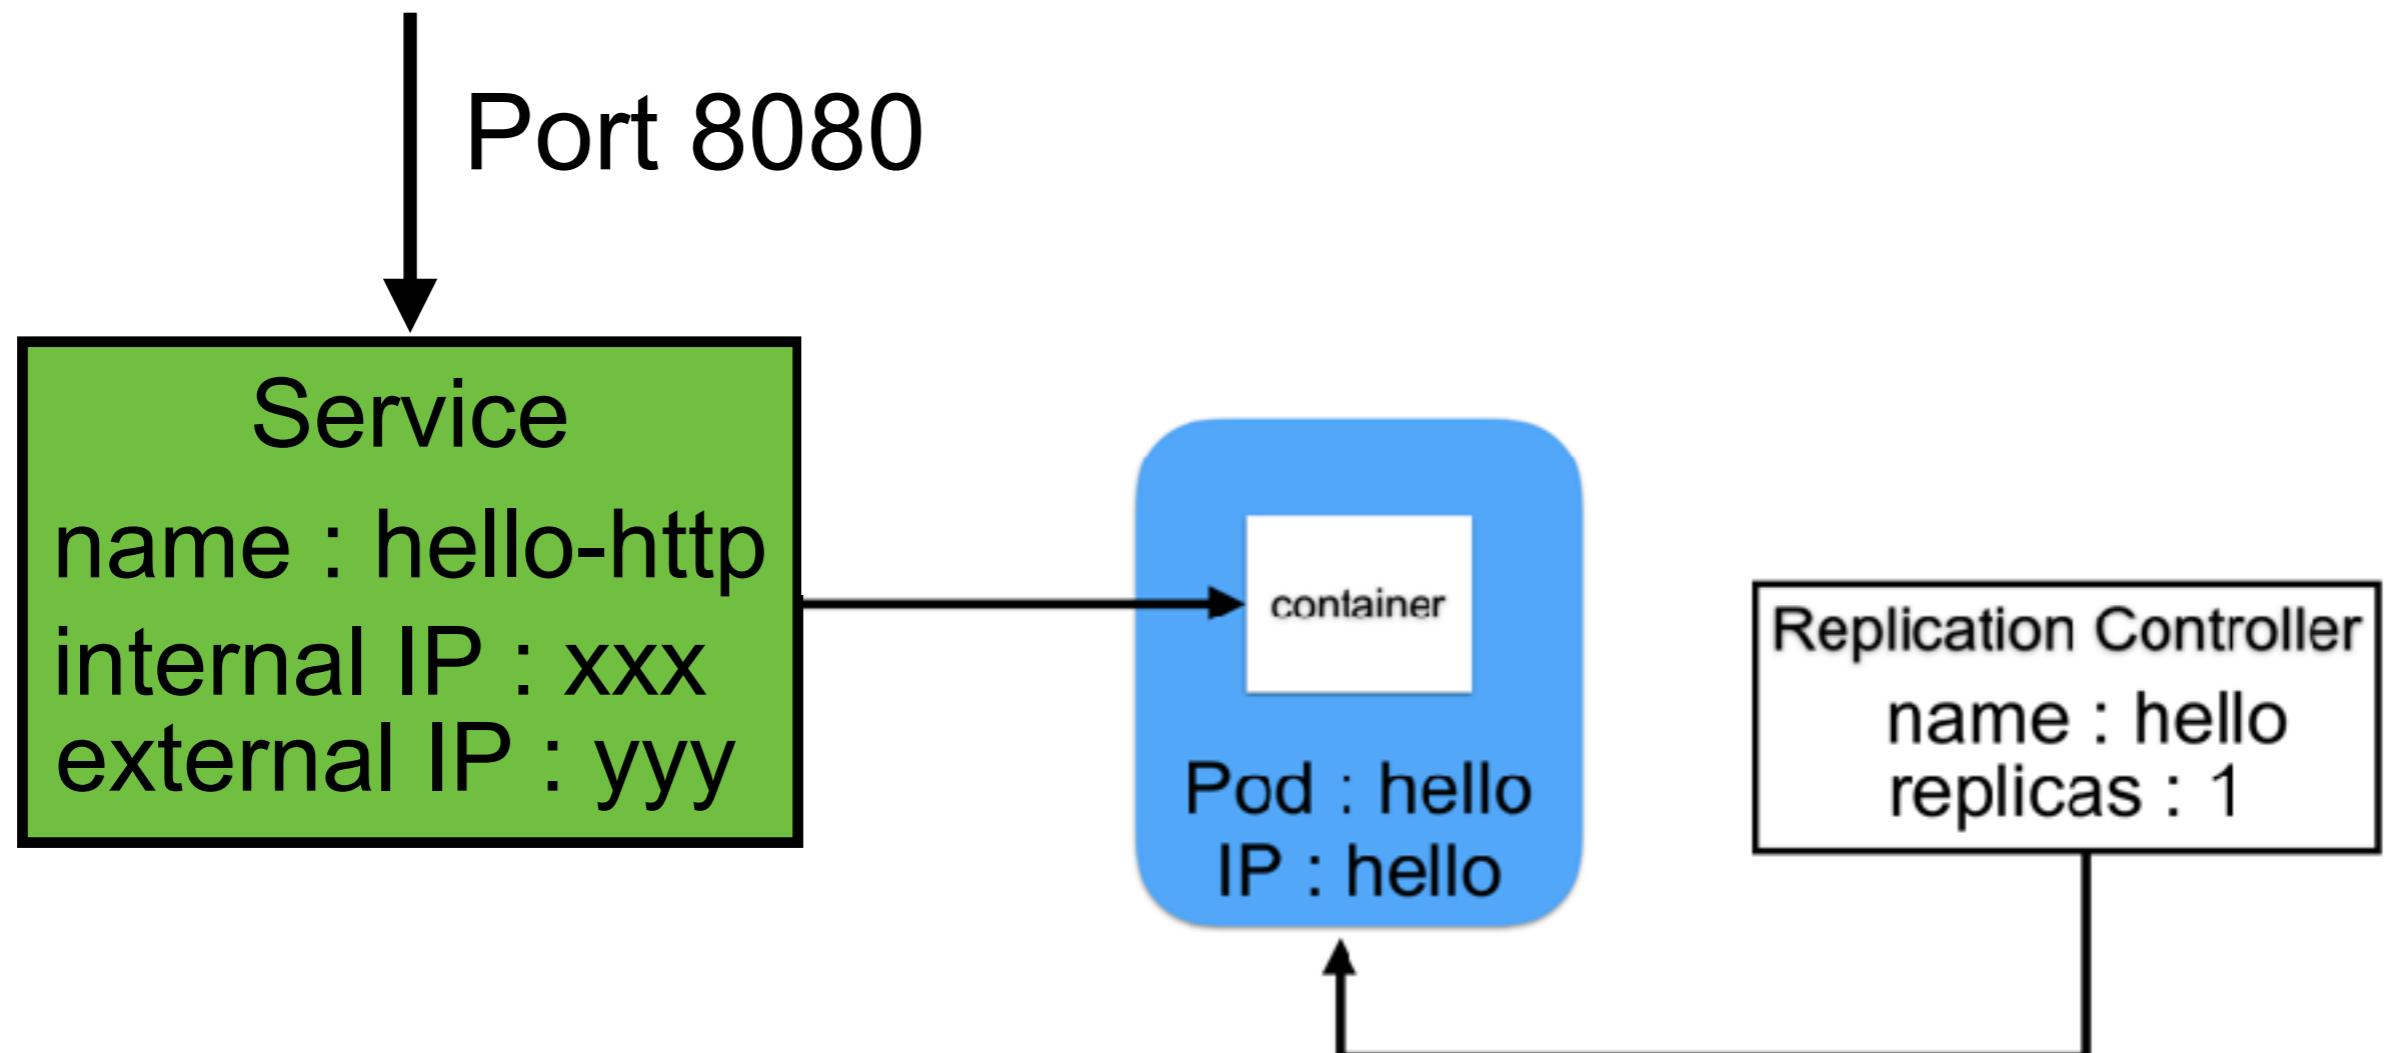


# Expose RC with service

```
$kubectl expose rc hello  
--type=LoadBalancer --name hello-http
```



# Expose RC with service



# Access service with minikube

```
$minikube service hello-http
```

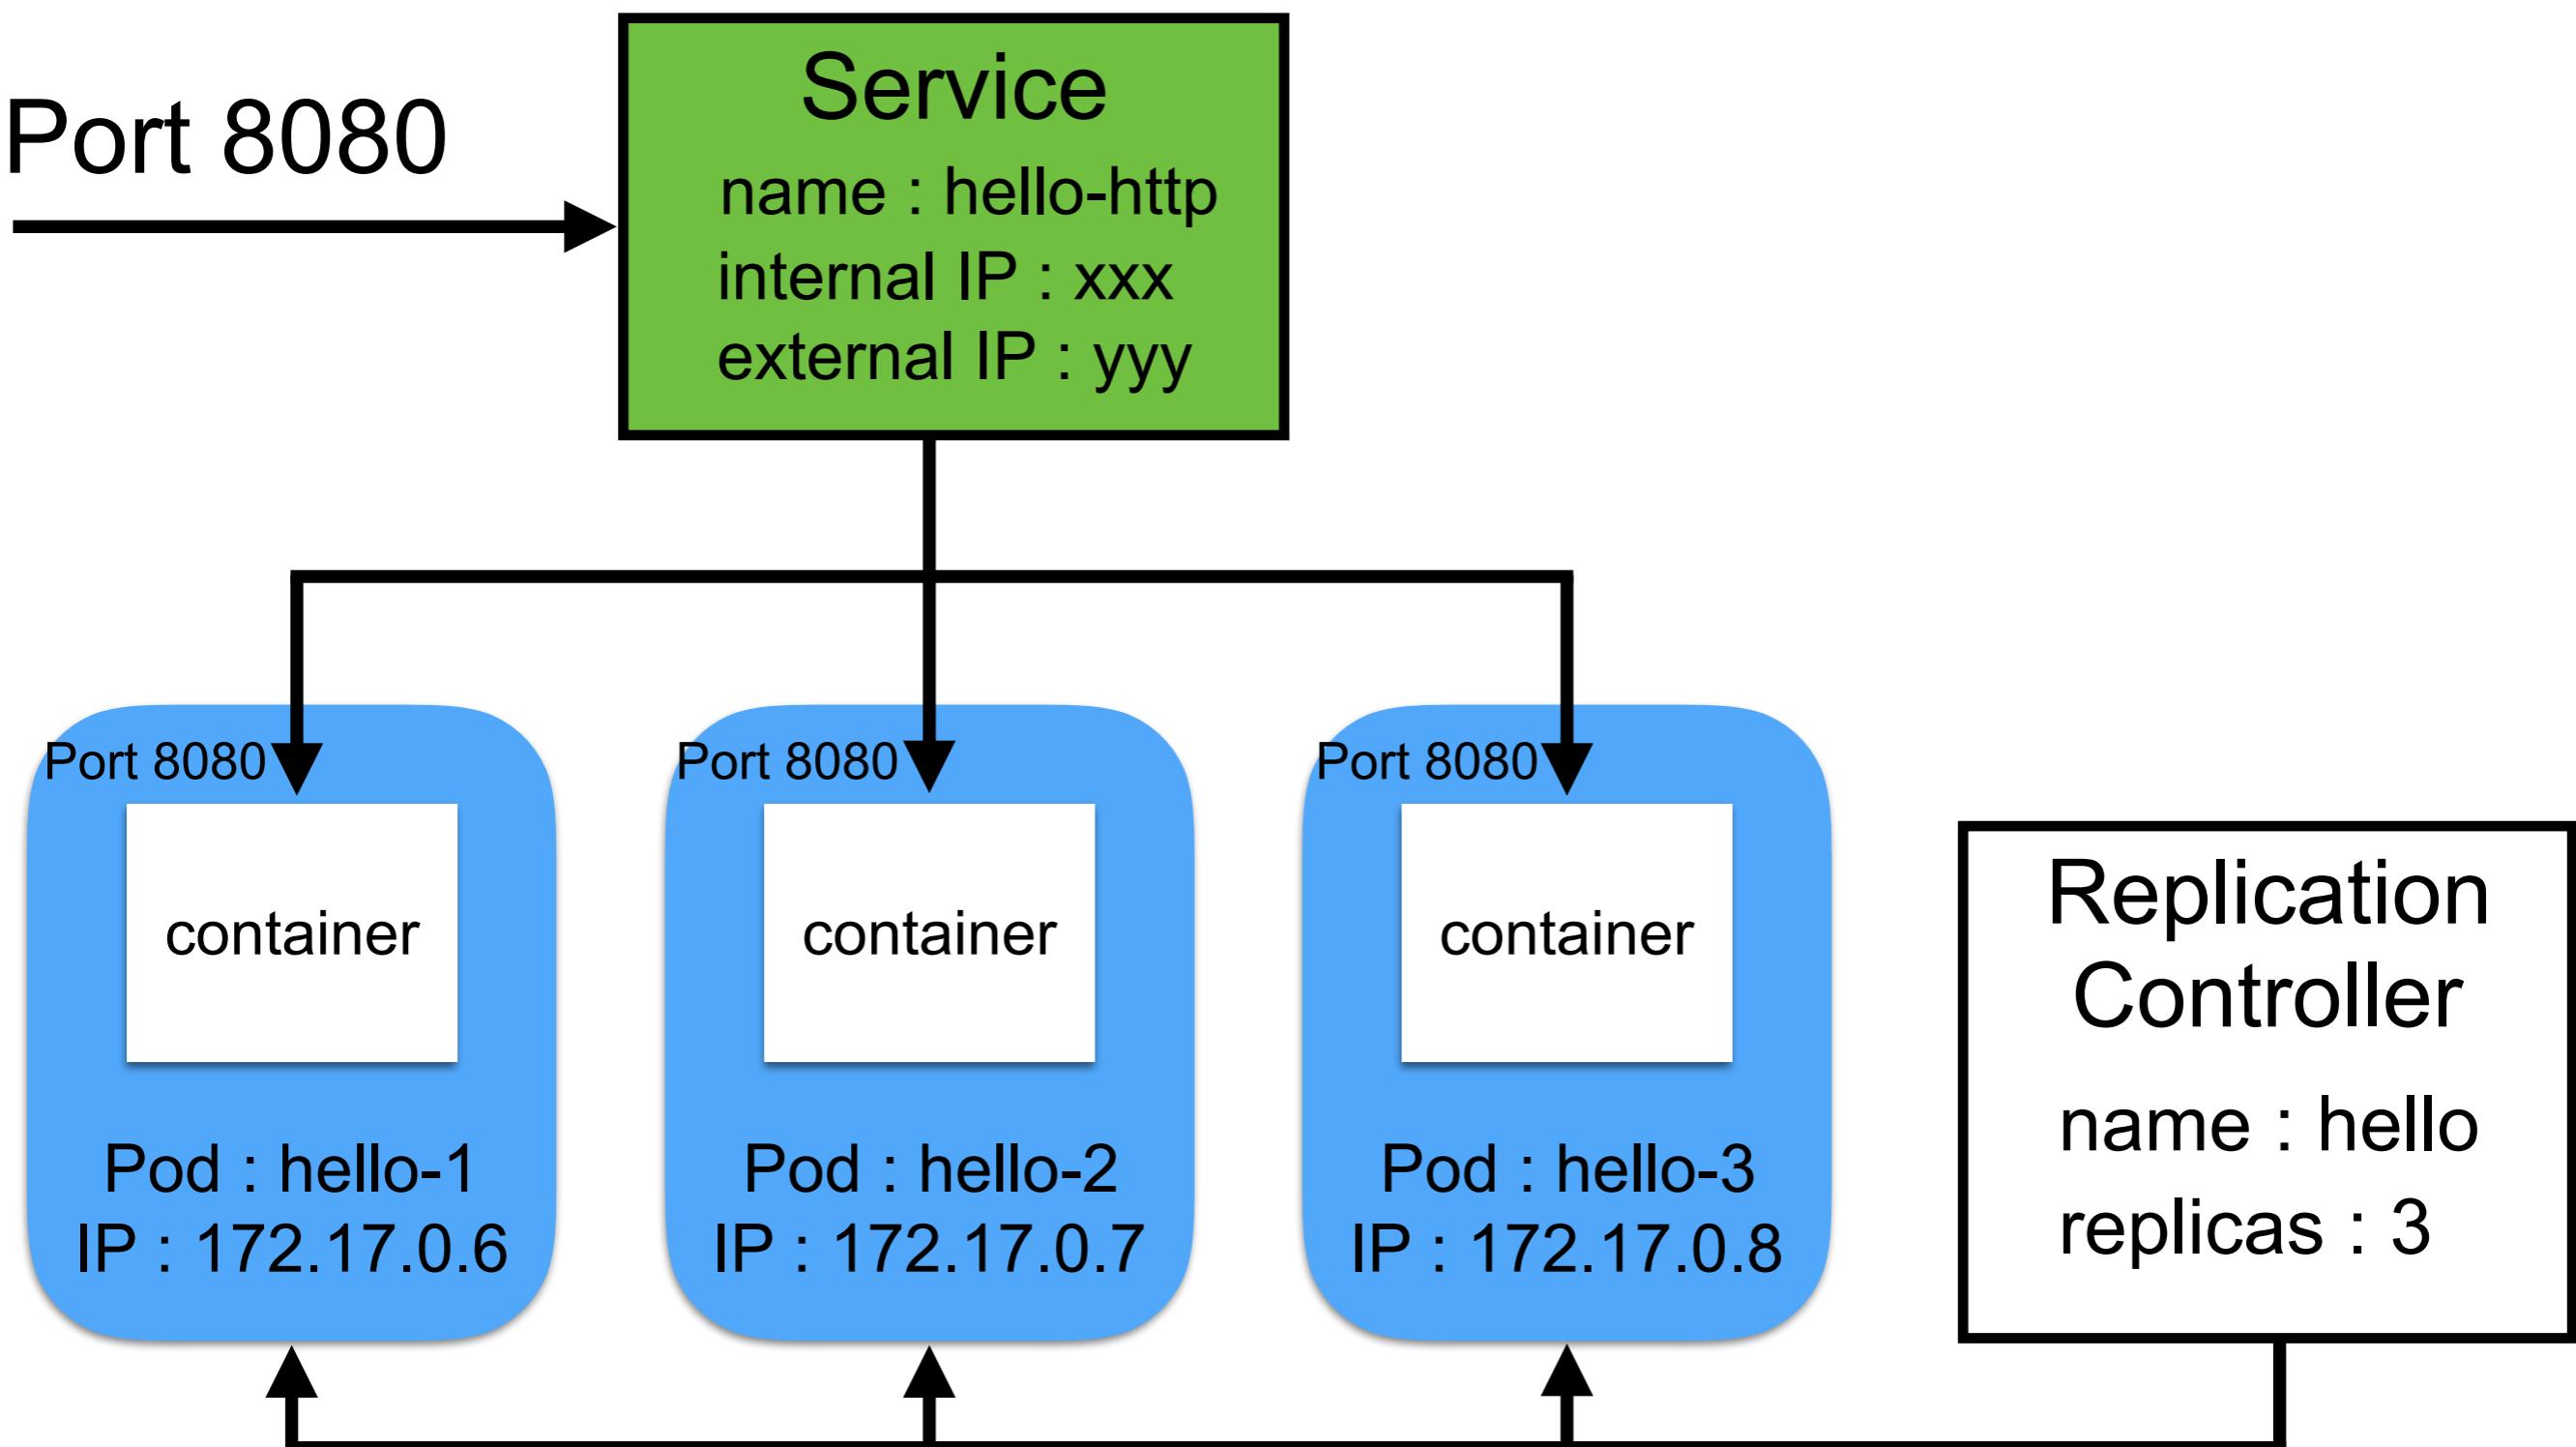


# Scaling the application

```
$ kubectl scale rc hello --replicas=3
```



# Scaling the application



# List of RC

\$kubectl get rc

NAME	DESIRED	CURRENT	READY	AGE
hello	3	3	3	35m



# List of Pods

\$kubectl get pod -o wide

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
hello-lw2p7	1/1	Running	0	29m	172.17.0.8	minikube
hello-pqnpz	1/1	Running	0	36m	172.17.0.6	minikube
hello-wwsnh	1/1	Running	0	29m	172.17.0.7	minikube



# Testing

```
export SERVICE=192.168.99.100:32614
```

```
$curl http://$SERVICE  
Hello, "/" on hello-wwsnh
```

```
$curl http://$SERVICE  
Hello, "/" on hello-pqnzp
```

```
$curl http://$SERVICE  
Hello, "/" on hello-lw2p7
```



# Kubernetes dashboard

## \$minikube dashboard

The screenshot shows the Kubernetes dashboard interface. At the top, there is a navigation bar with a Kubernetes logo, a search bar, and a '+ CREATE' button. Below the navigation bar, a blue header bar displays the text 'Overview'. On the left side, there is a sidebar under the heading 'Cluster' containing links for Namespaces, Nodes, Persistent Volumes, Roles, and Storage Classes. A dropdown menu for 'Namespaces' is open, showing 'default' as the selected option. Below the sidebar, there are two main sections: 'Workloads' and 'Workloads Statuses'. The 'Workloads Statuses' section contains two green circular charts, both labeled '100.00%' and representing 'Pods' and 'Replication Controllers'. In the 'Workloads' section, there is a table titled 'Pods' listing three entries: 'hello-lw2p7', 'hello-wwsnh', and 'hello-pqnpz', all running on the 'minikube' node. The table includes columns for Name, Node, Status, Restarts, and Age.

Name	Node	Status	Restarts	Age
hello-lw2p7	minikube	Running	0	53 minutes
hello-wwsnh	minikube	Running	0	53 minutes
hello-pqnpz	minikube	Running	0	59 minutes



# Try to use jboss/wildfly

The screenshot shows the Docker Hub interface for the jboss/wildfly repository. At the top, there's a dark header bar with a search bar containing 'Search' and icons for Docker Hub. To the right are links for 'Explore', 'Help', 'Sign up' (in a blue button), and 'Sign in'. Below the header, the repository details are shown: 'PUBLIC | AUTOMATED BUILD' followed by the repository name 'jboss/wildfly' with a star icon indicating it's favorited. A note says 'Last pushed: a month ago'. Below this, there are tabs for 'Repo Info', 'Tags', 'Dockerfile', and 'Build Details', with 'Build Details' being the active tab. Under 'Short Description', it says 'WildFly application server image'. Under 'Full Description', it says 'WildFly Docker image'.

<https://hub.docker.com/r/jboss/wildfly/>



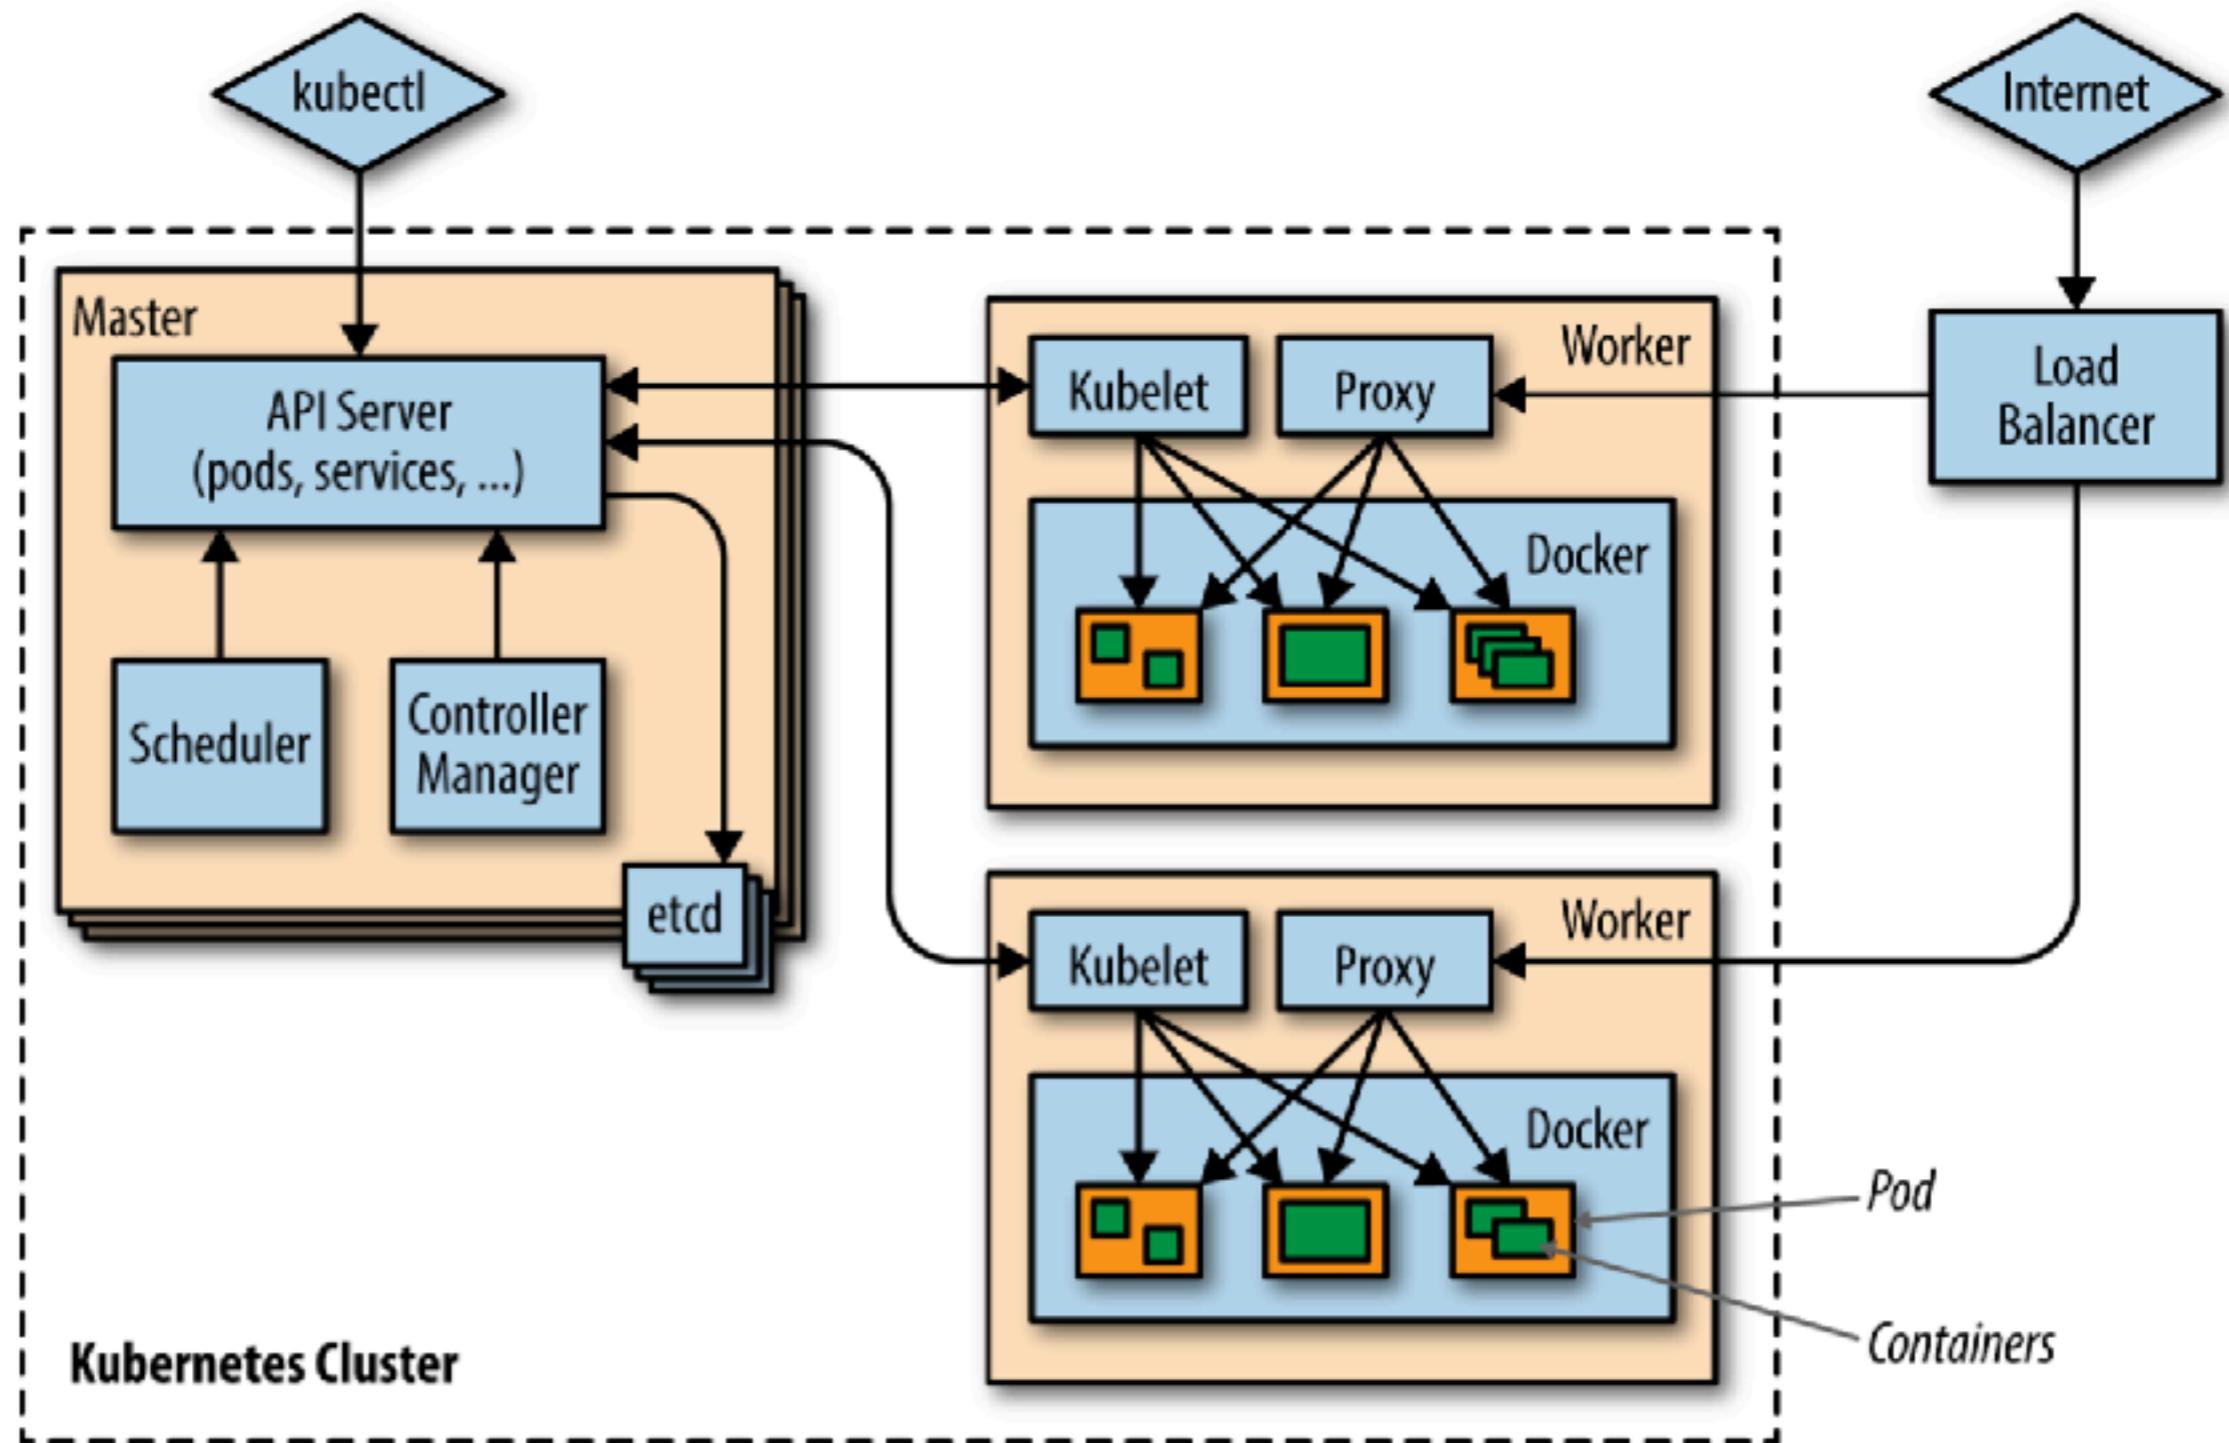
# Welcome to Kubernetes



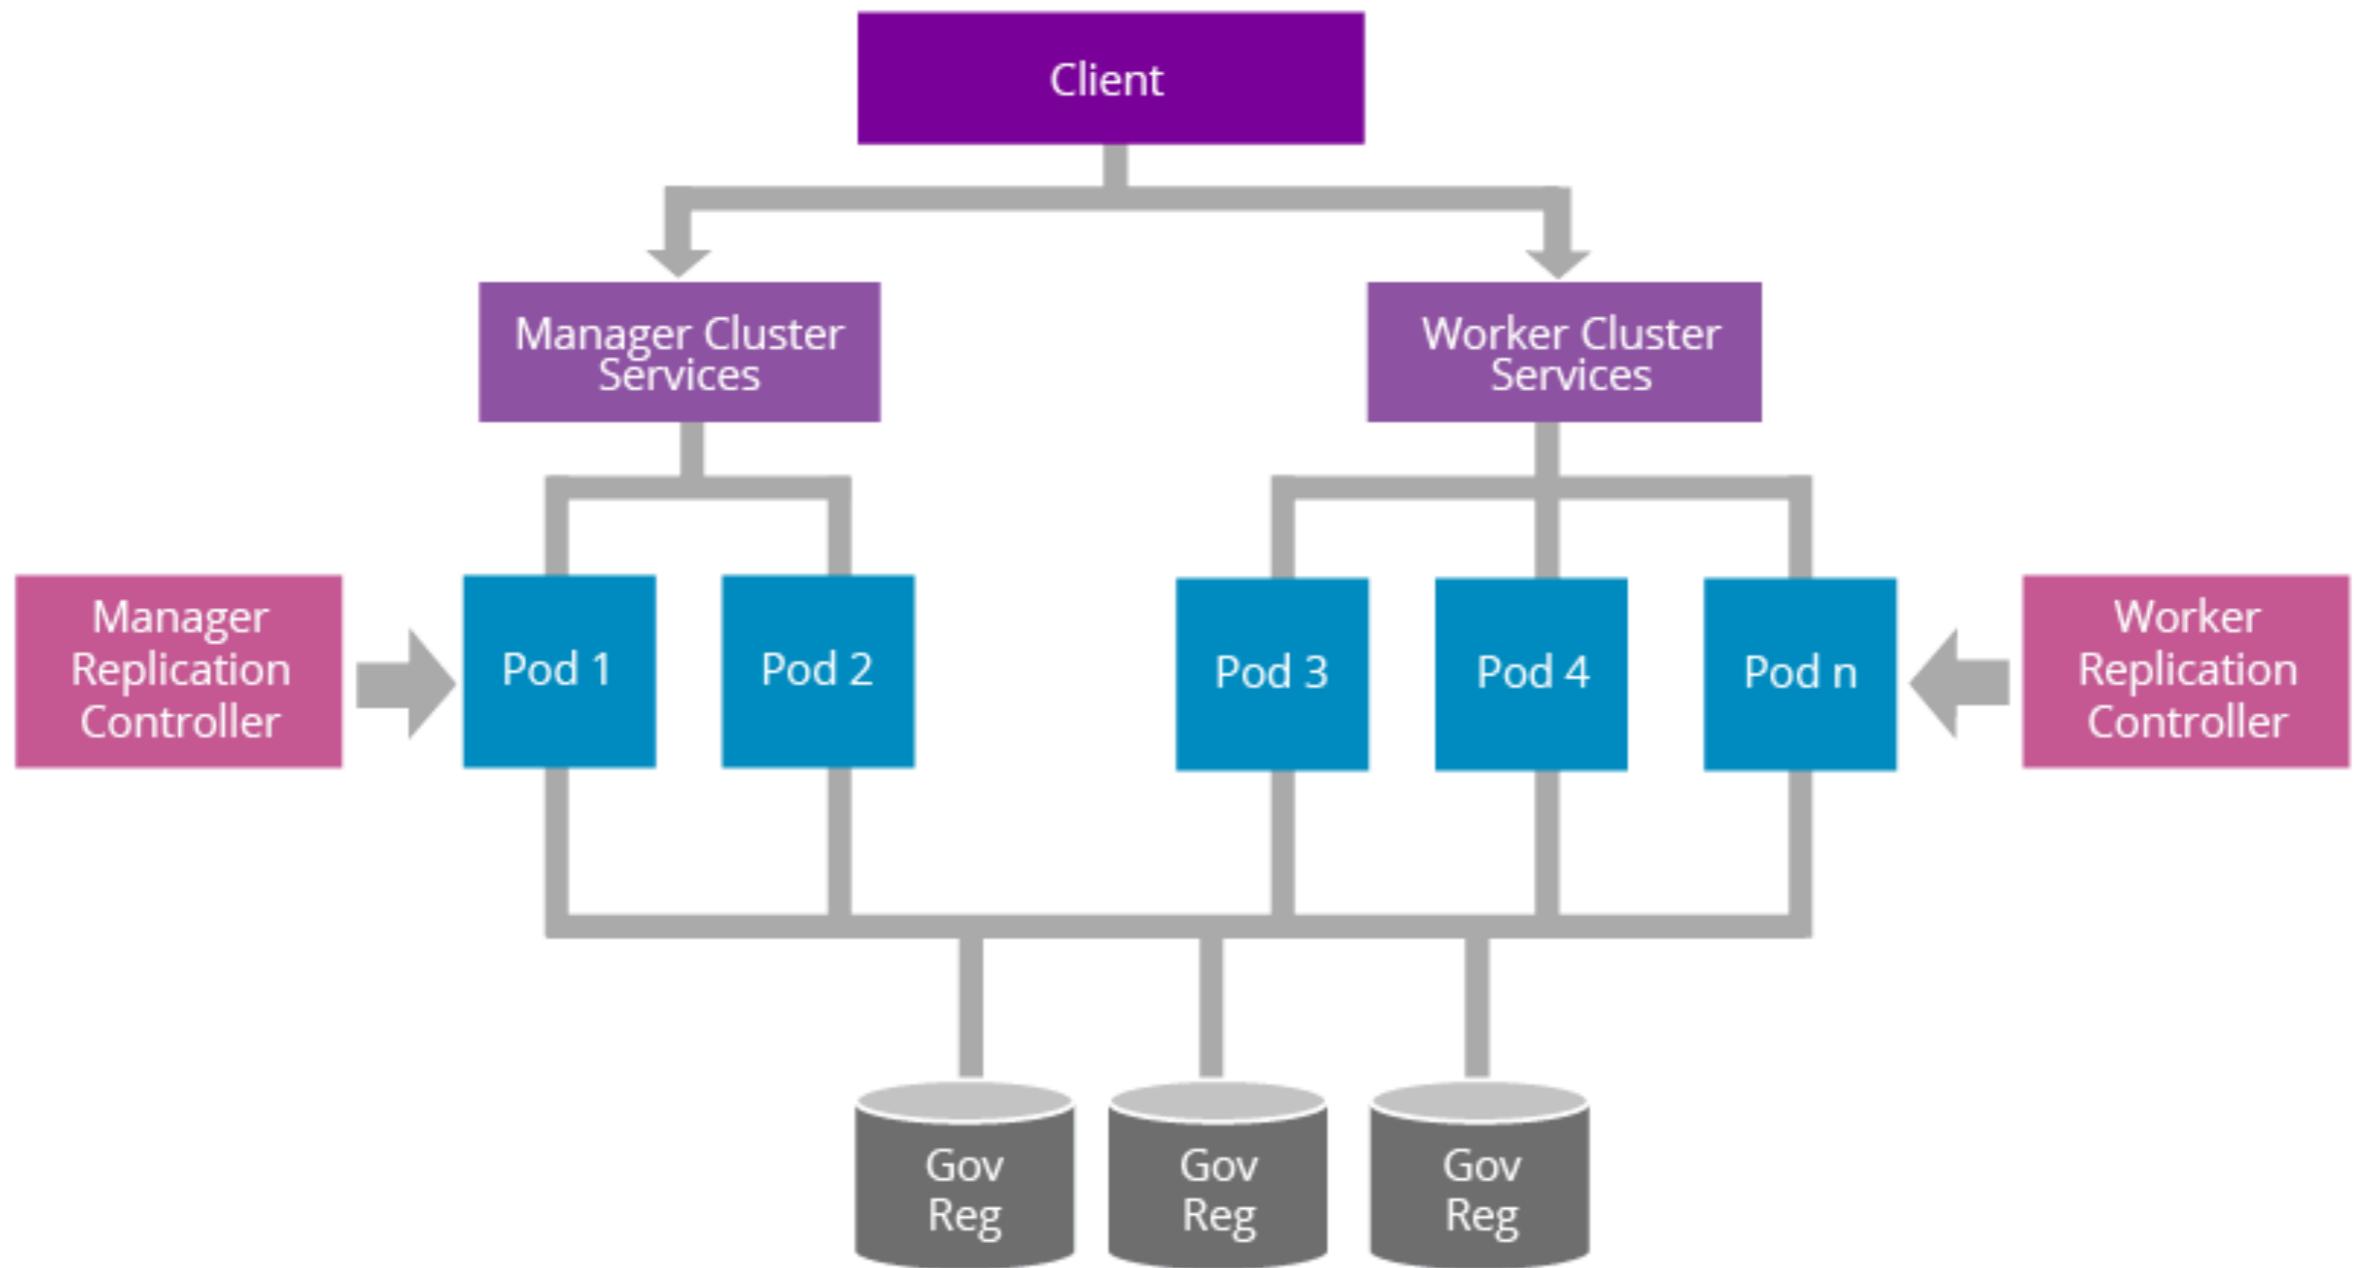
# Kubernetes Architecture



# Kubernetes architecture



# Kubernetes architecture



<https://wso2.com/whitepapers/a-reference-architecture-for-deploying-wso2-middleware-on-kubernetes/>



# Key features

<https://kubernetes.io/>



# Key features

- Automatic binpacking
- Horizontal Pod Autoscaler (HPA)
- Automated rollouts and rollbacks
- Storage orchestration



# Key features

Self-healing

Service discovery and Load Balancing (LB)

Secret and config management

Batch execution



# Automatic binpacking



# Automatic binpacking

Limit of CPU/Memory can define on Pods  
Schedule will select a node that have enough resources



# Automatic binpacking

When reach the Memory limit ?

1. Kill current Pod
2. If Pod have restart flag, try to create in other node



# Automatic binpacking

When reach the CPU limit ?

1. Schedule not kill Pod
2. Schedule waiting it back to normal state



# Automatic binpacking

Try to check resources of node

\$kubectl describe node

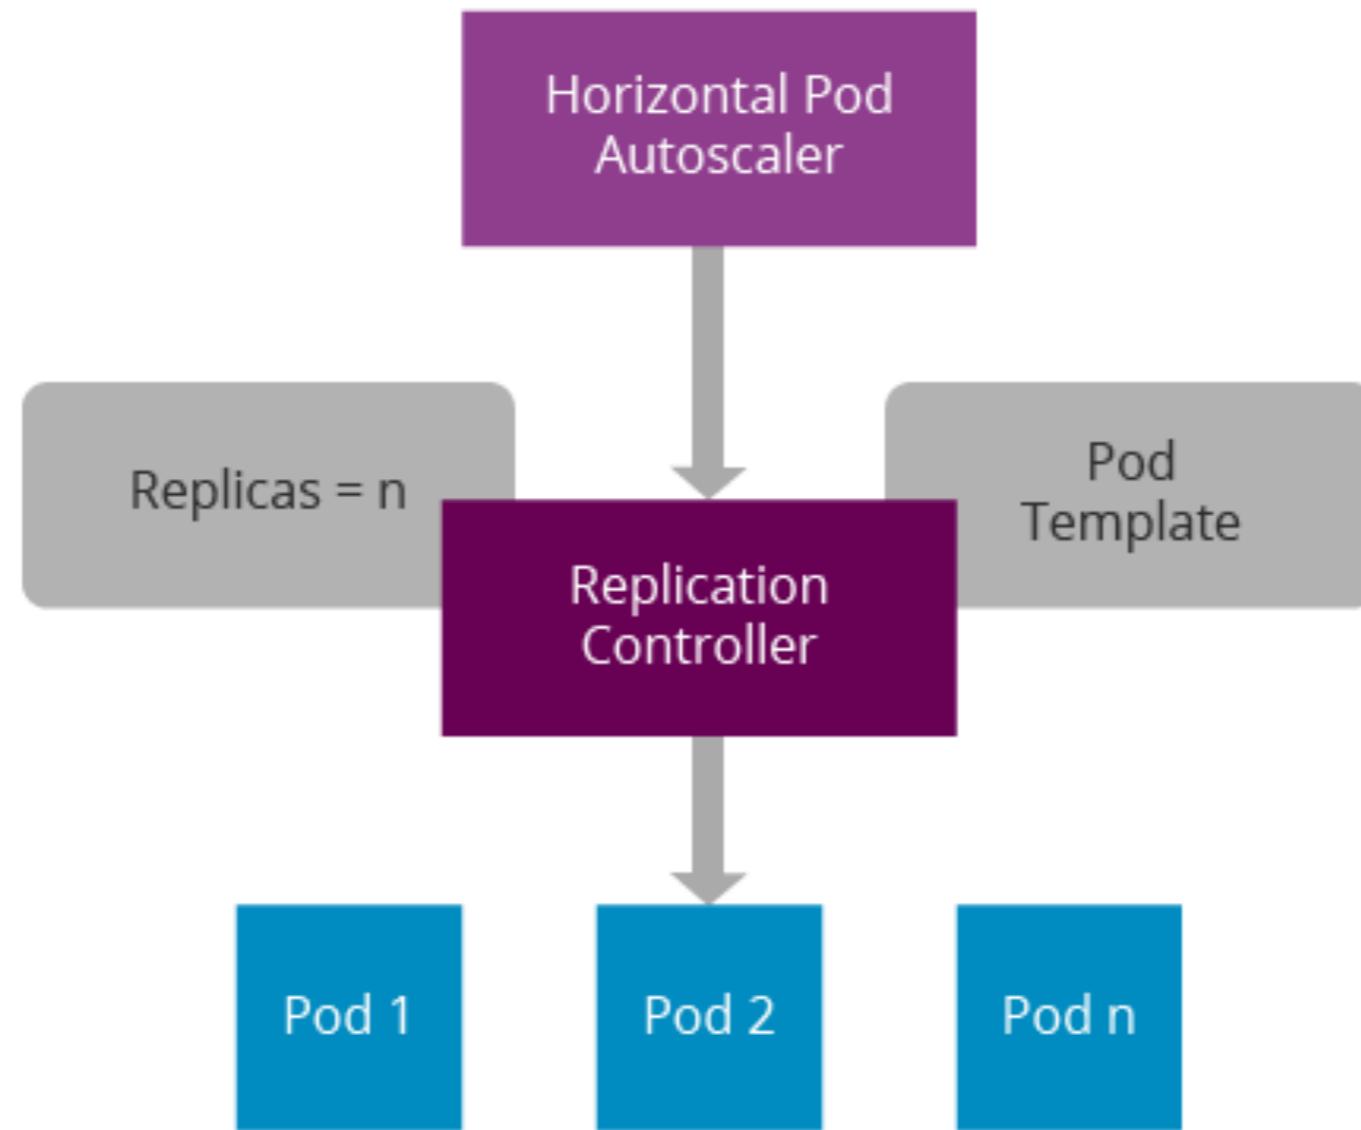
CPU Requests	CPU Limits	Memory Requests	Memory Limits
5m (0%)	0 (0%)	50Mi (2%)	0 (0%)
260m (13%)	0 (0%)	110Mi (5%)	170Mi (8%)
0 (0%)	0 (0%)	0 (0%)	0 (0%)
0 (0%)	0 (0%)	0 (0%)	0 (0%)



# Horizontal Pod Autoscaler



# Horizontal Pod Autoscaler

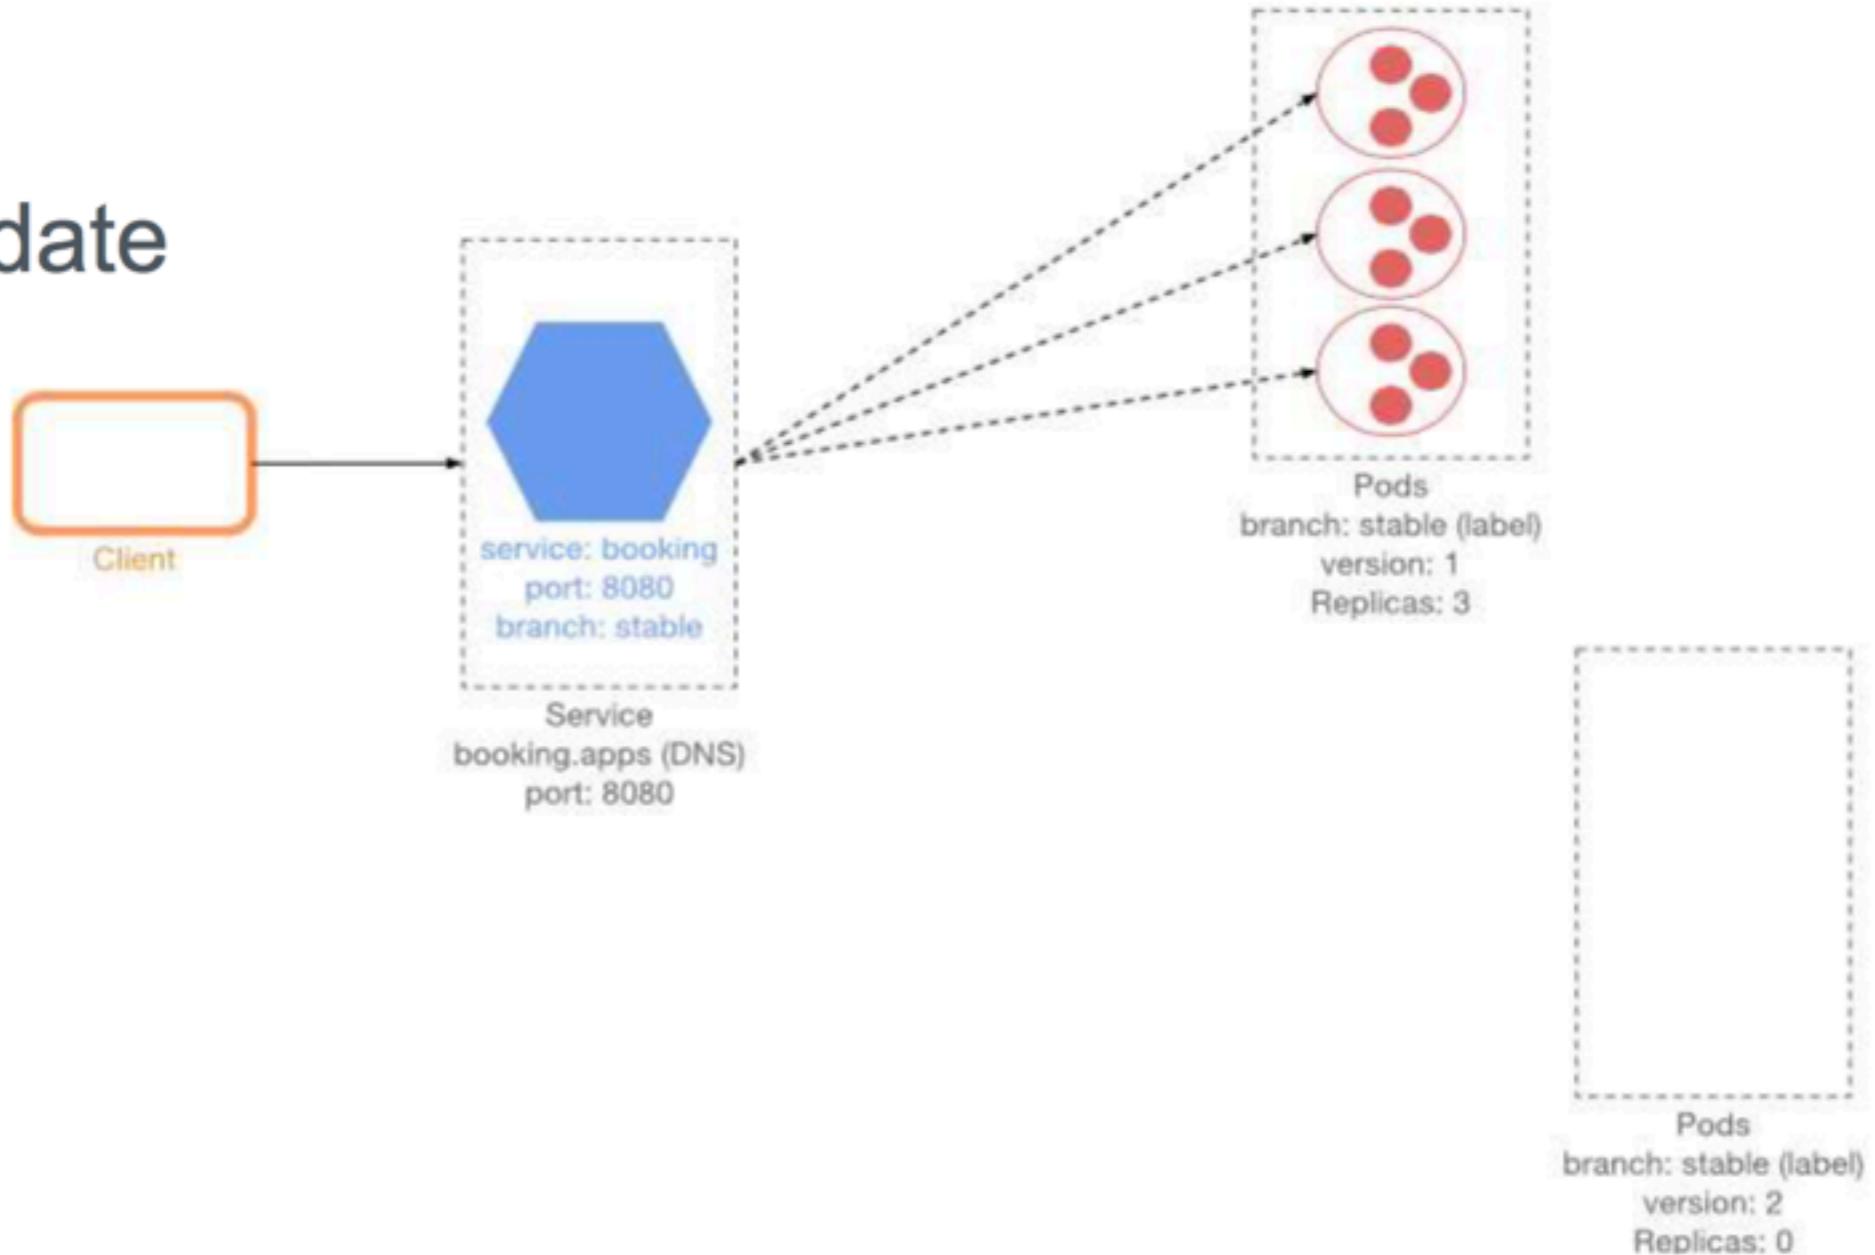


# Automated rollouts & rollbacks



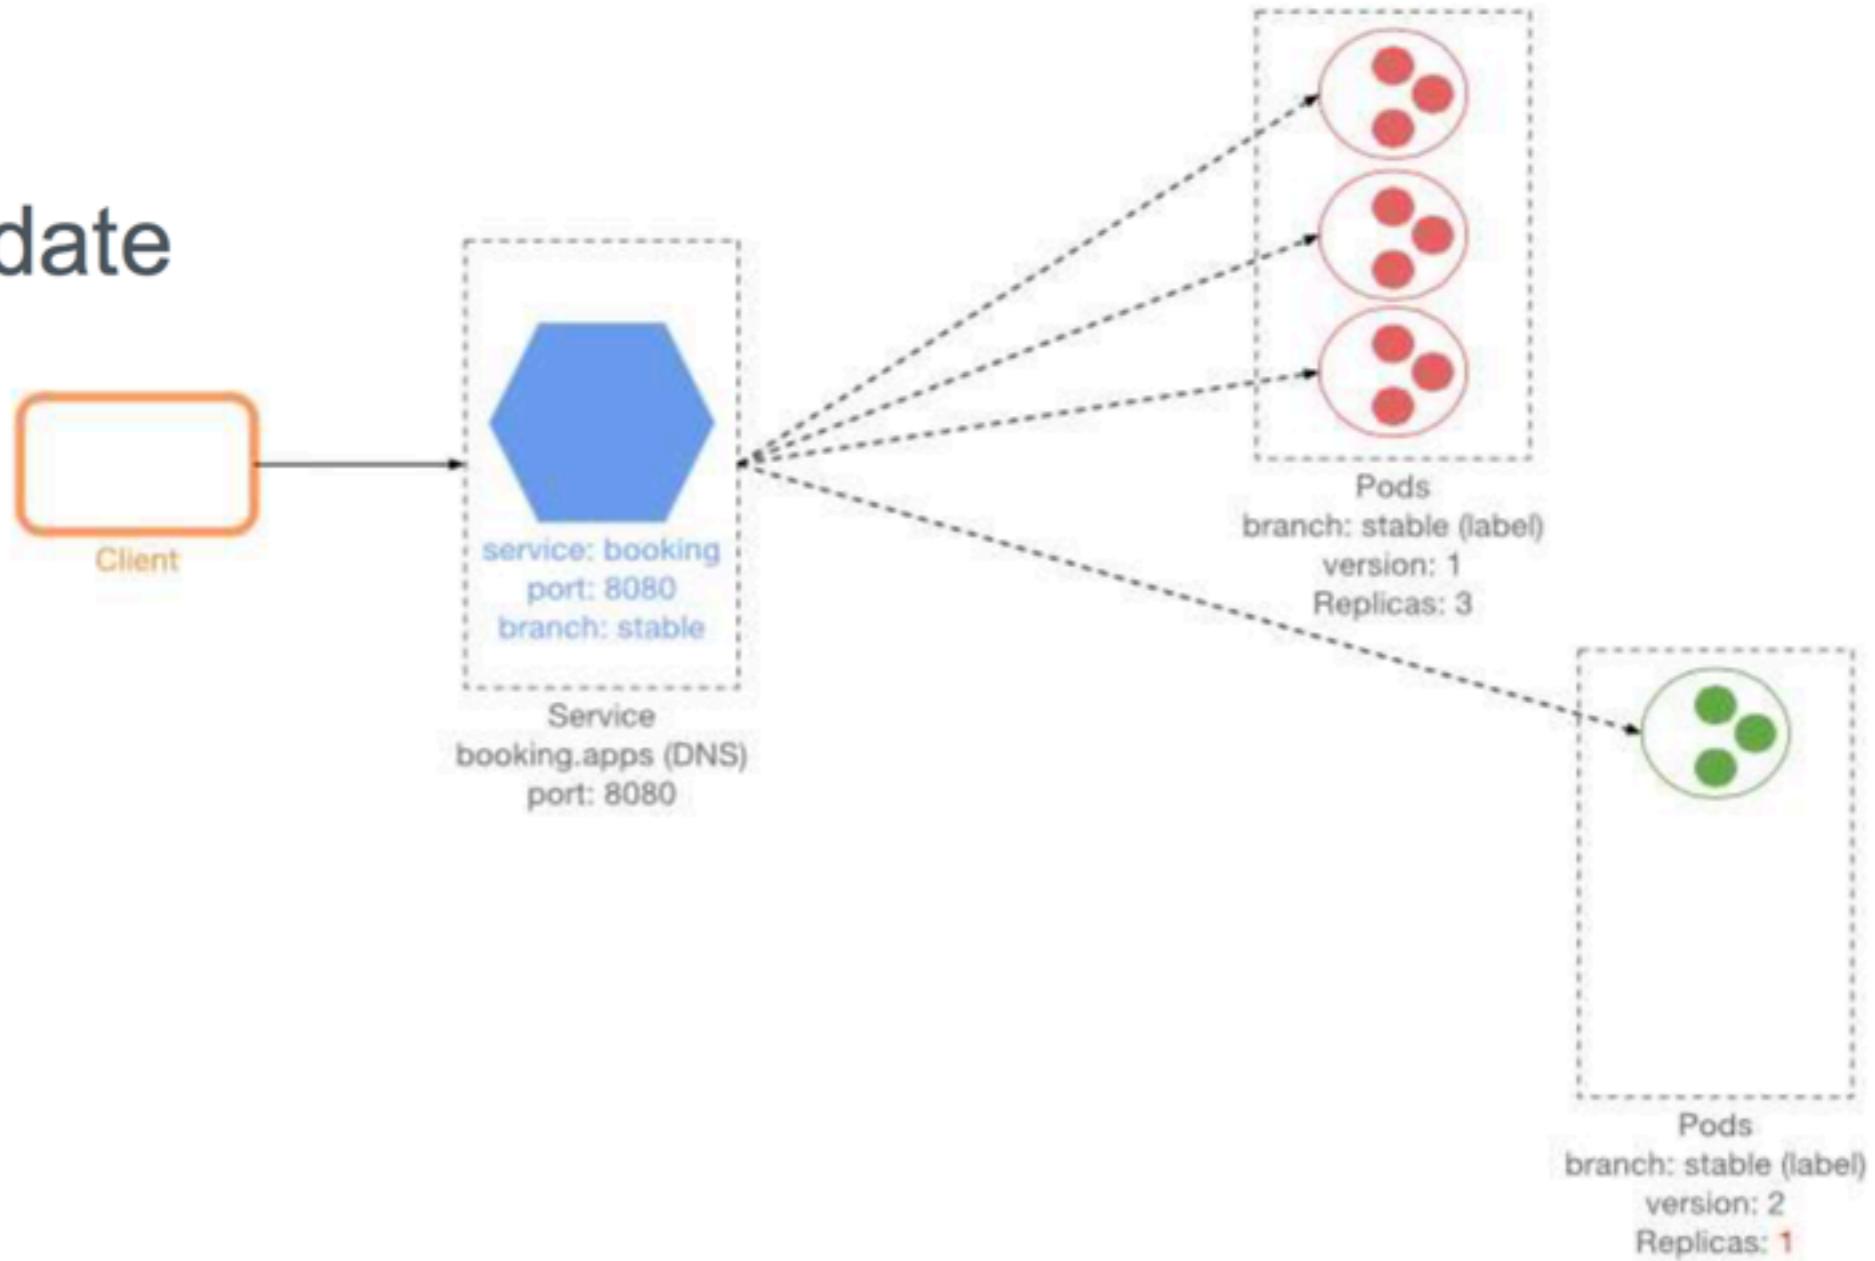
# Rolling update (1)

## Rolling Update



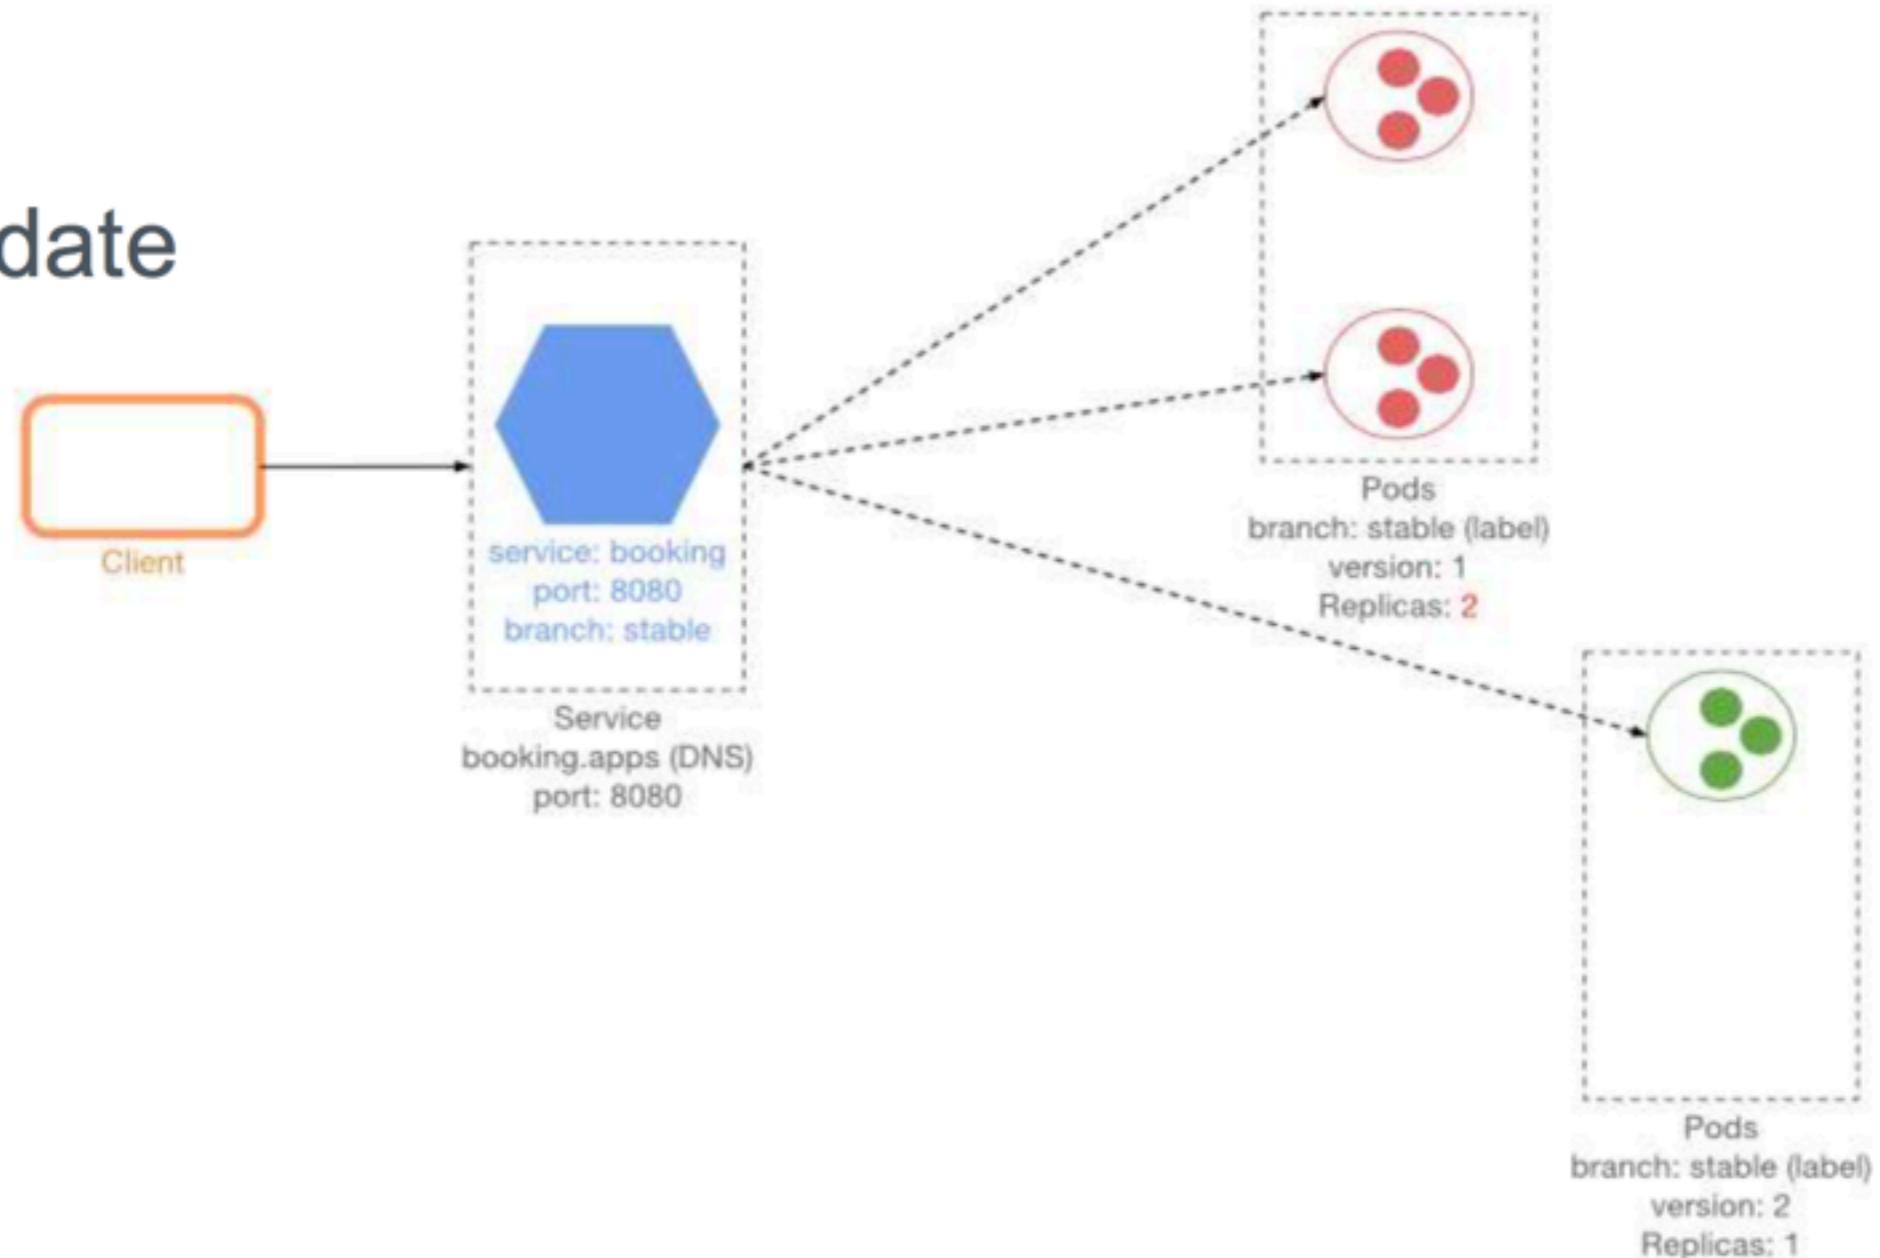
# Rolling update (2)

## Rolling Update



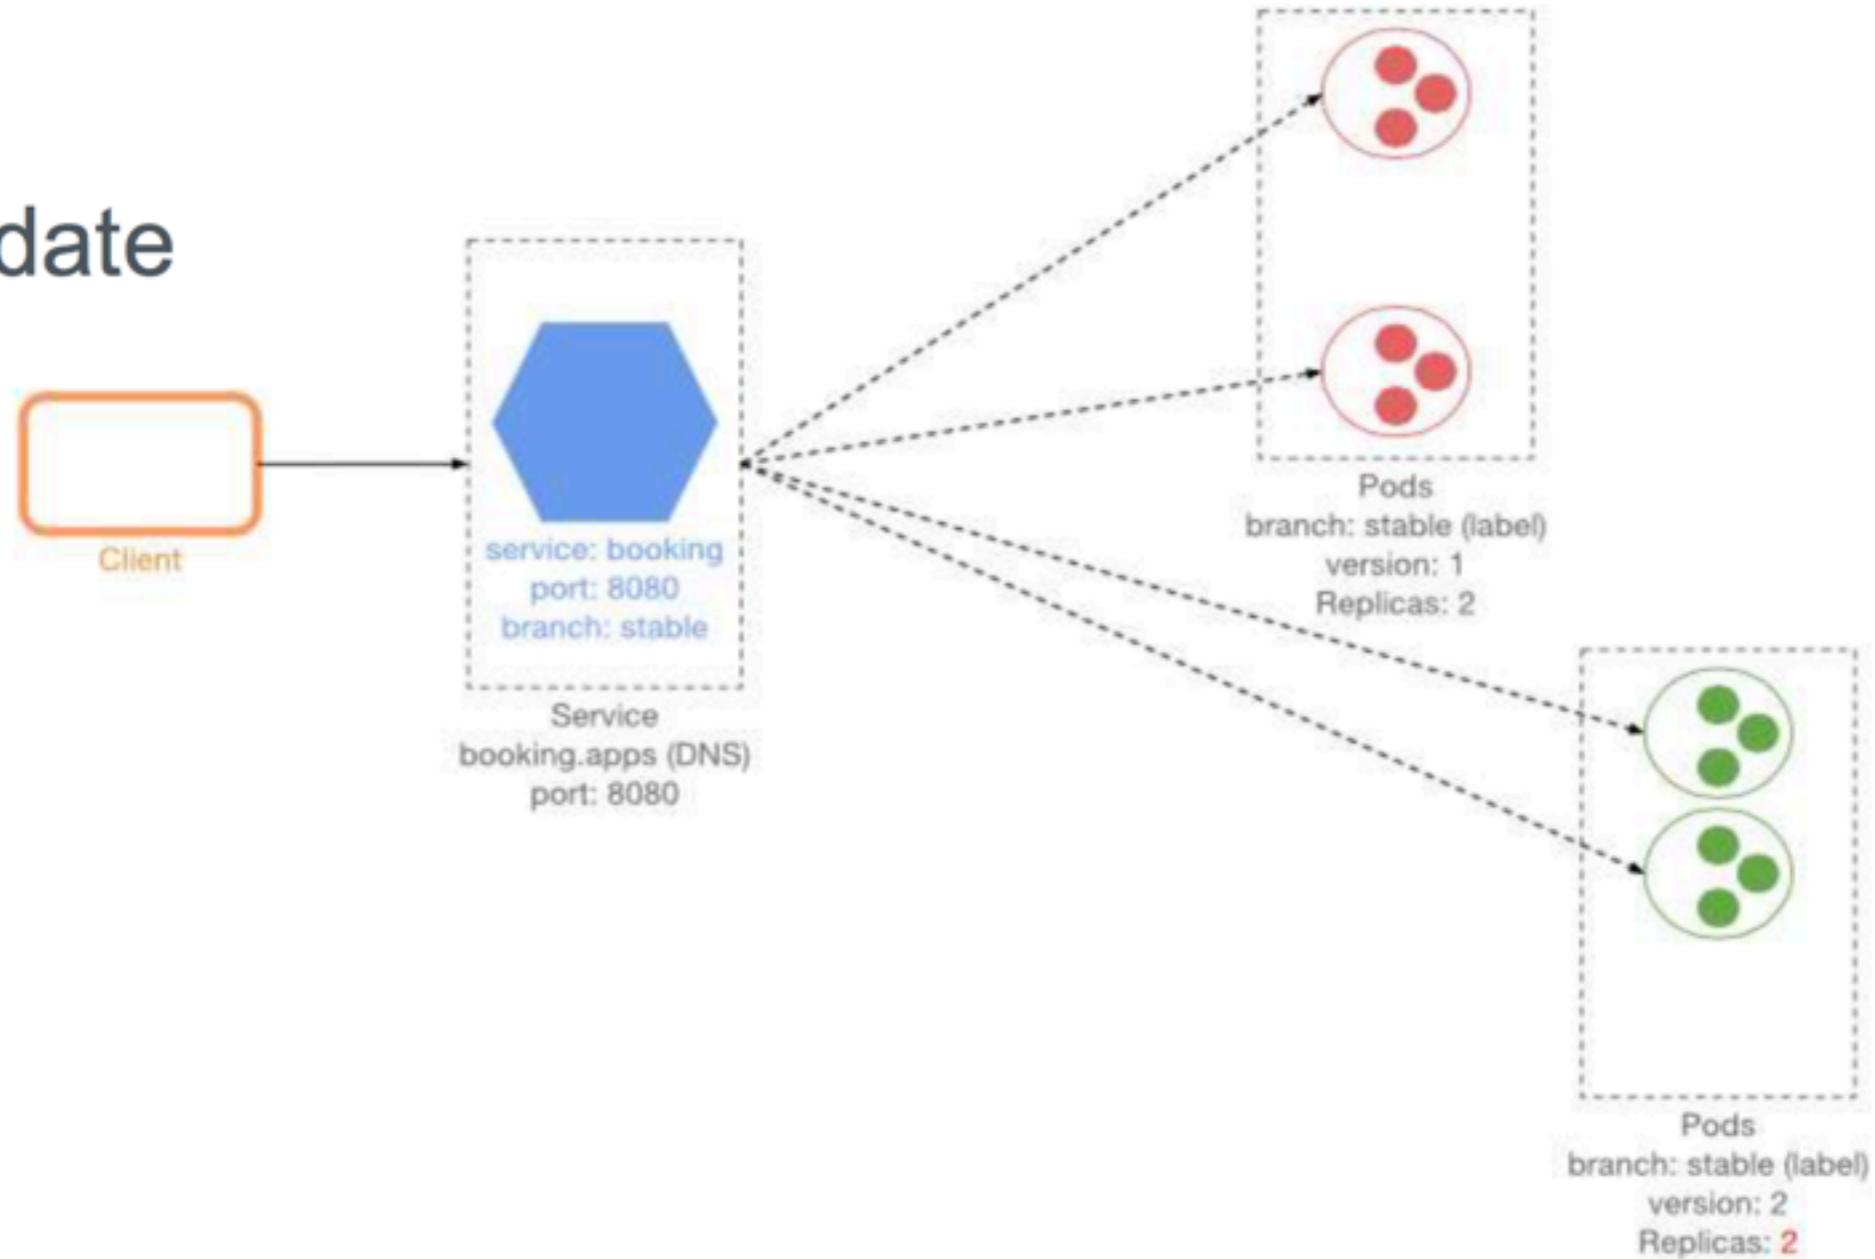
# Rolling update (3)

## Rolling Update



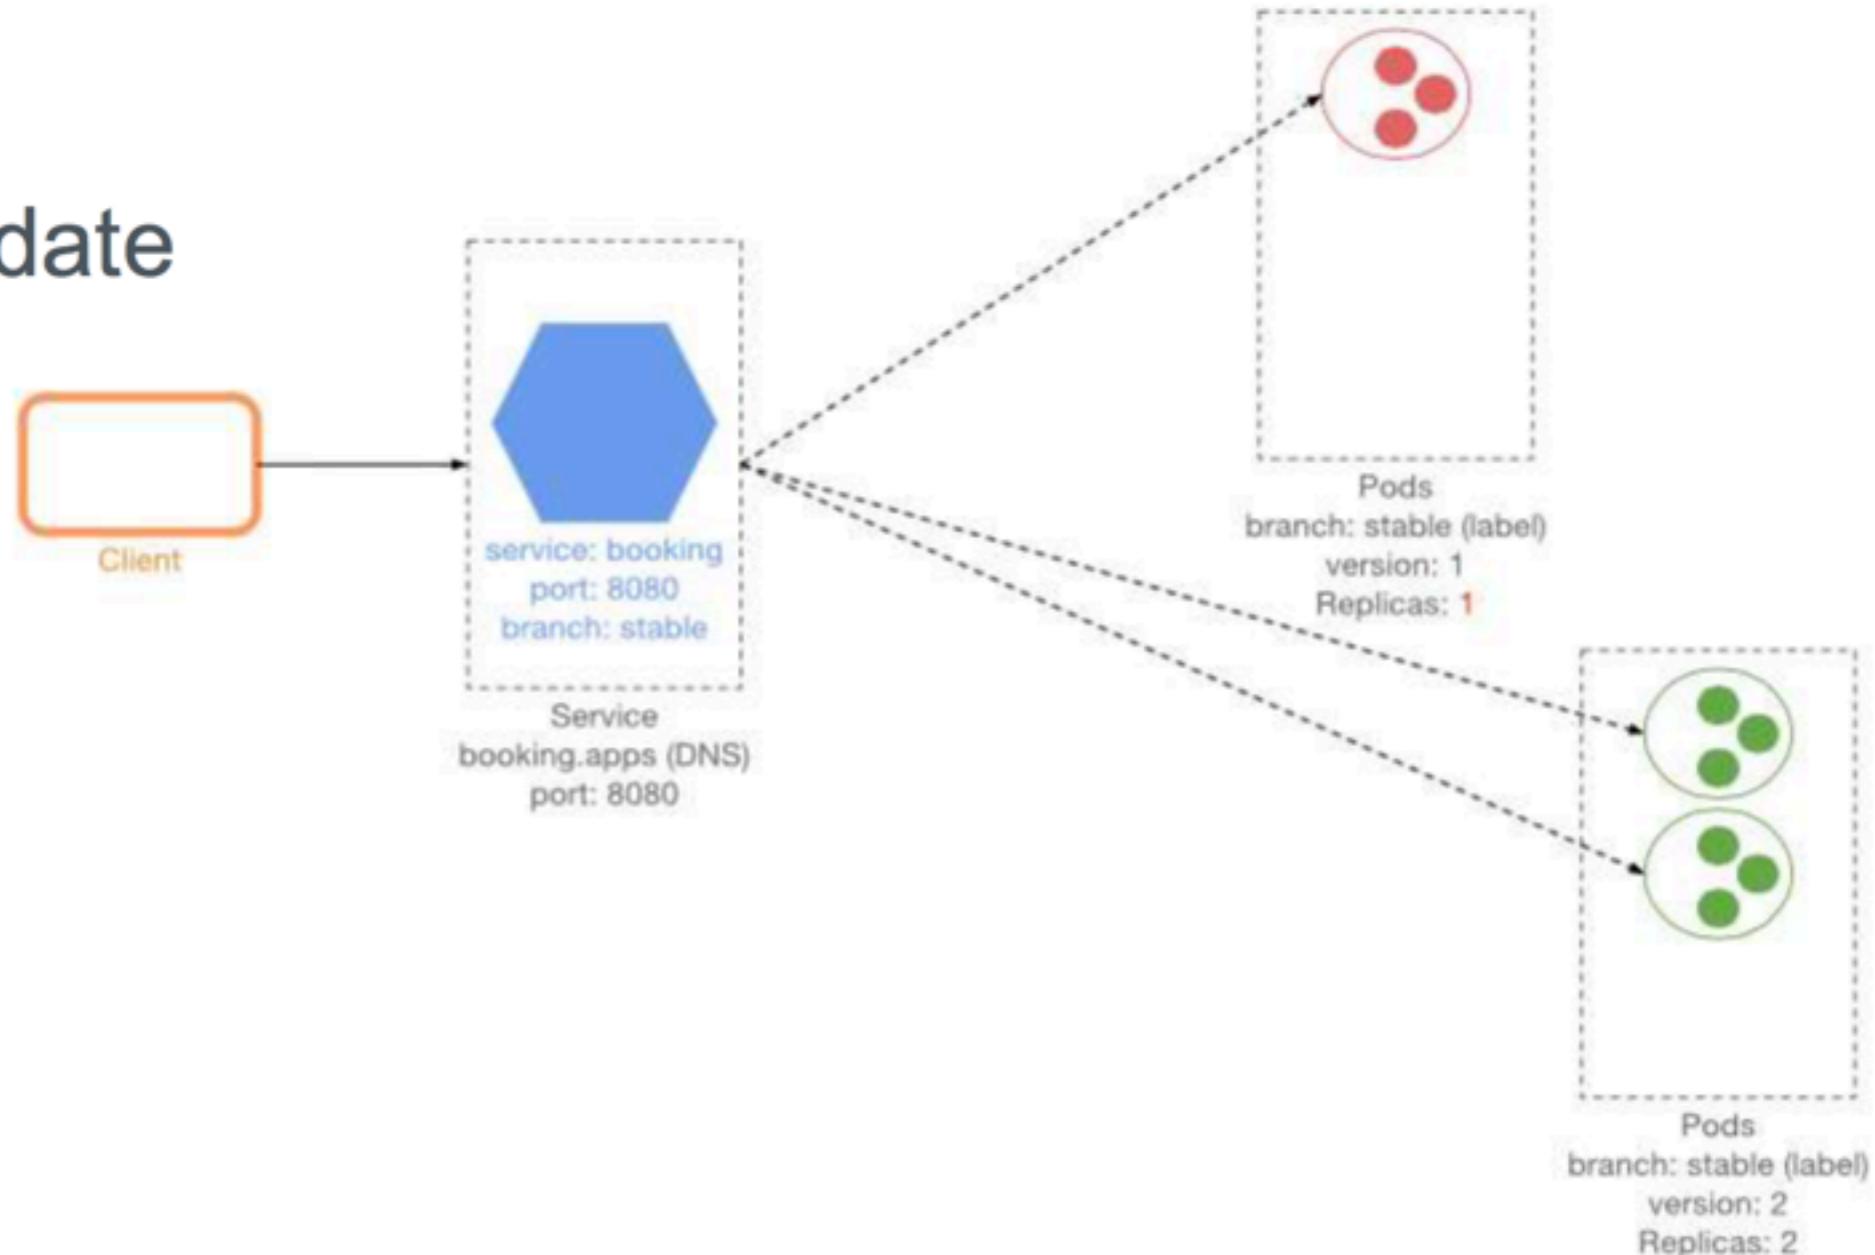
# Rolling update (4)

## Rolling Update



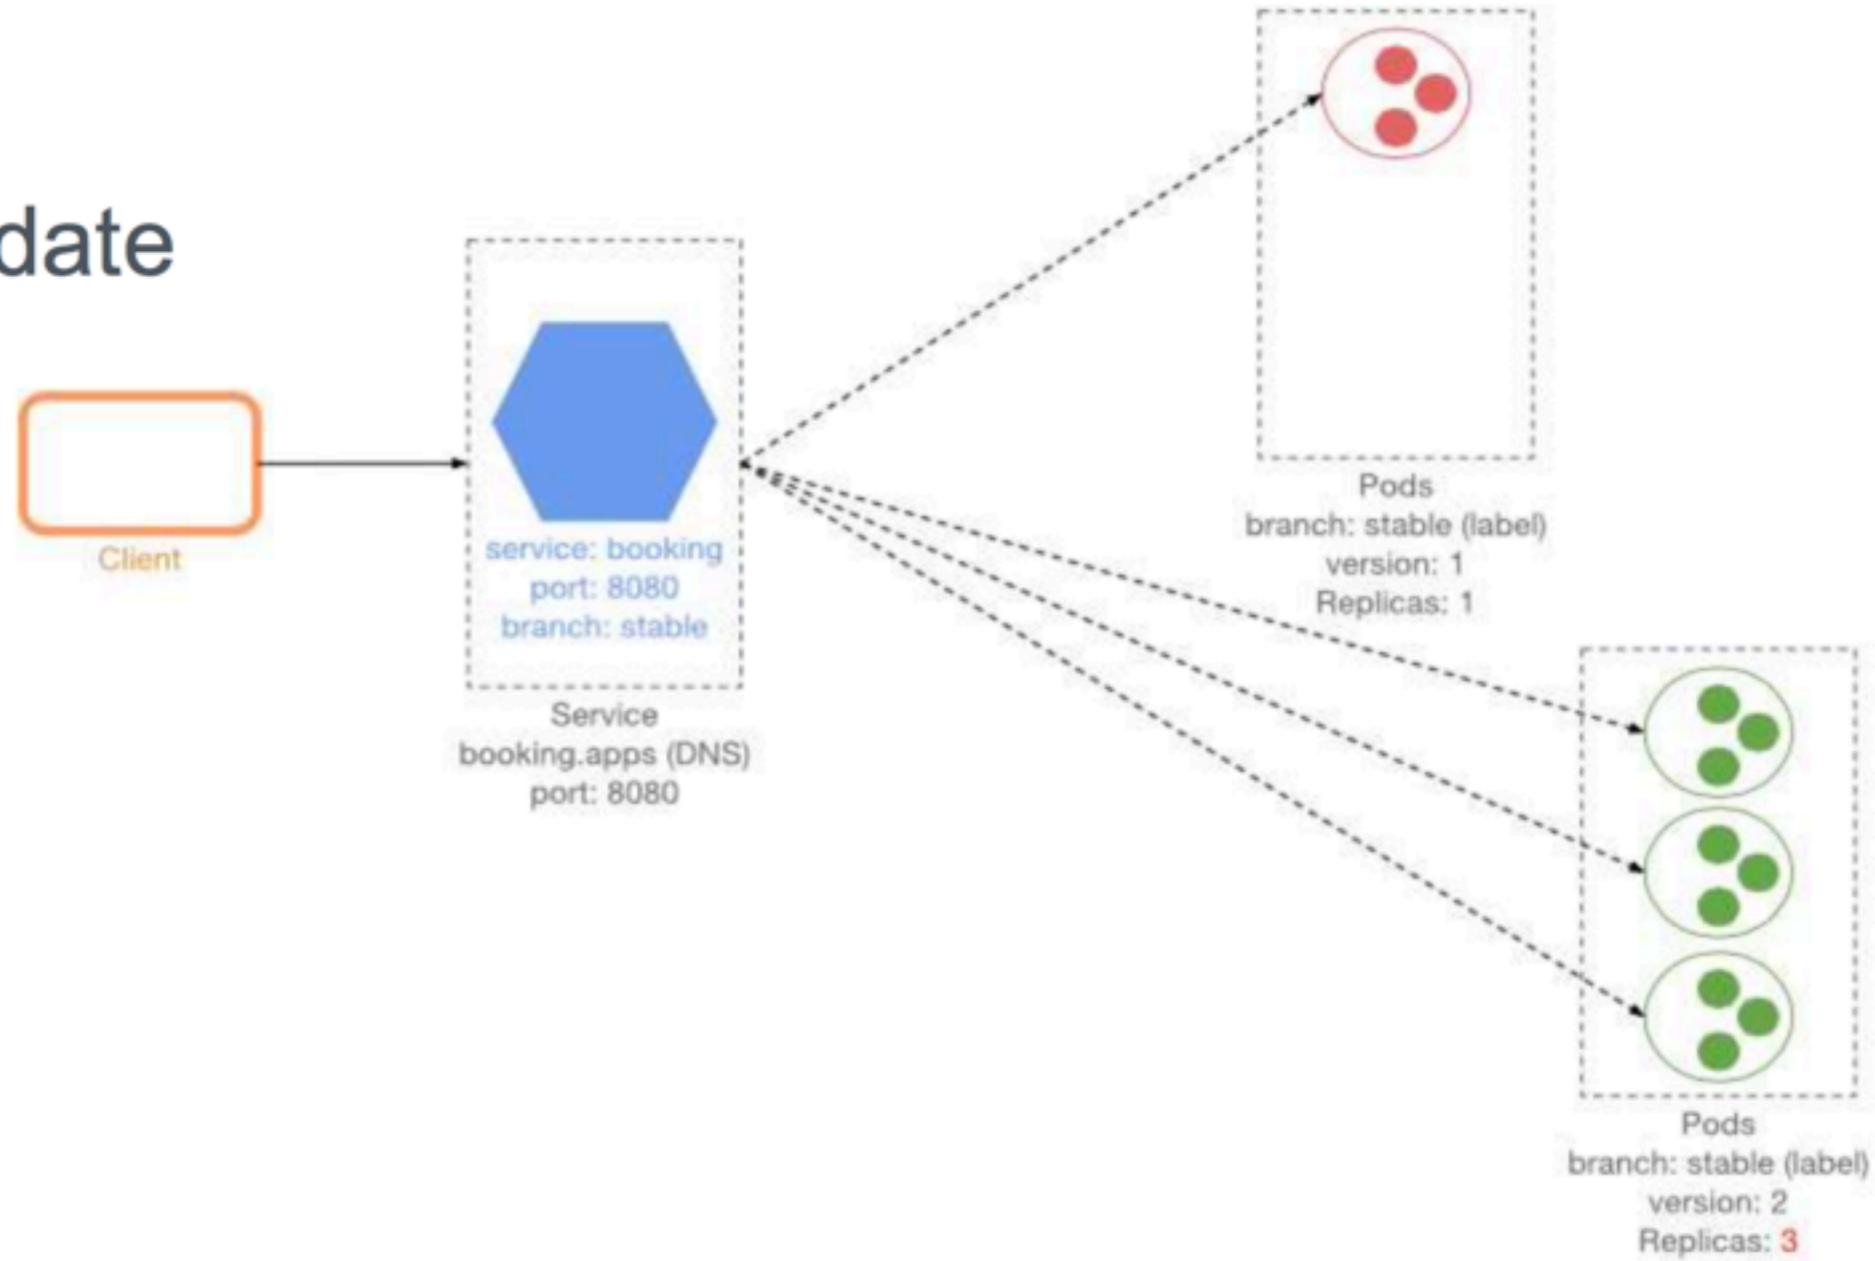
# Rolling update (5)

## Rolling Update



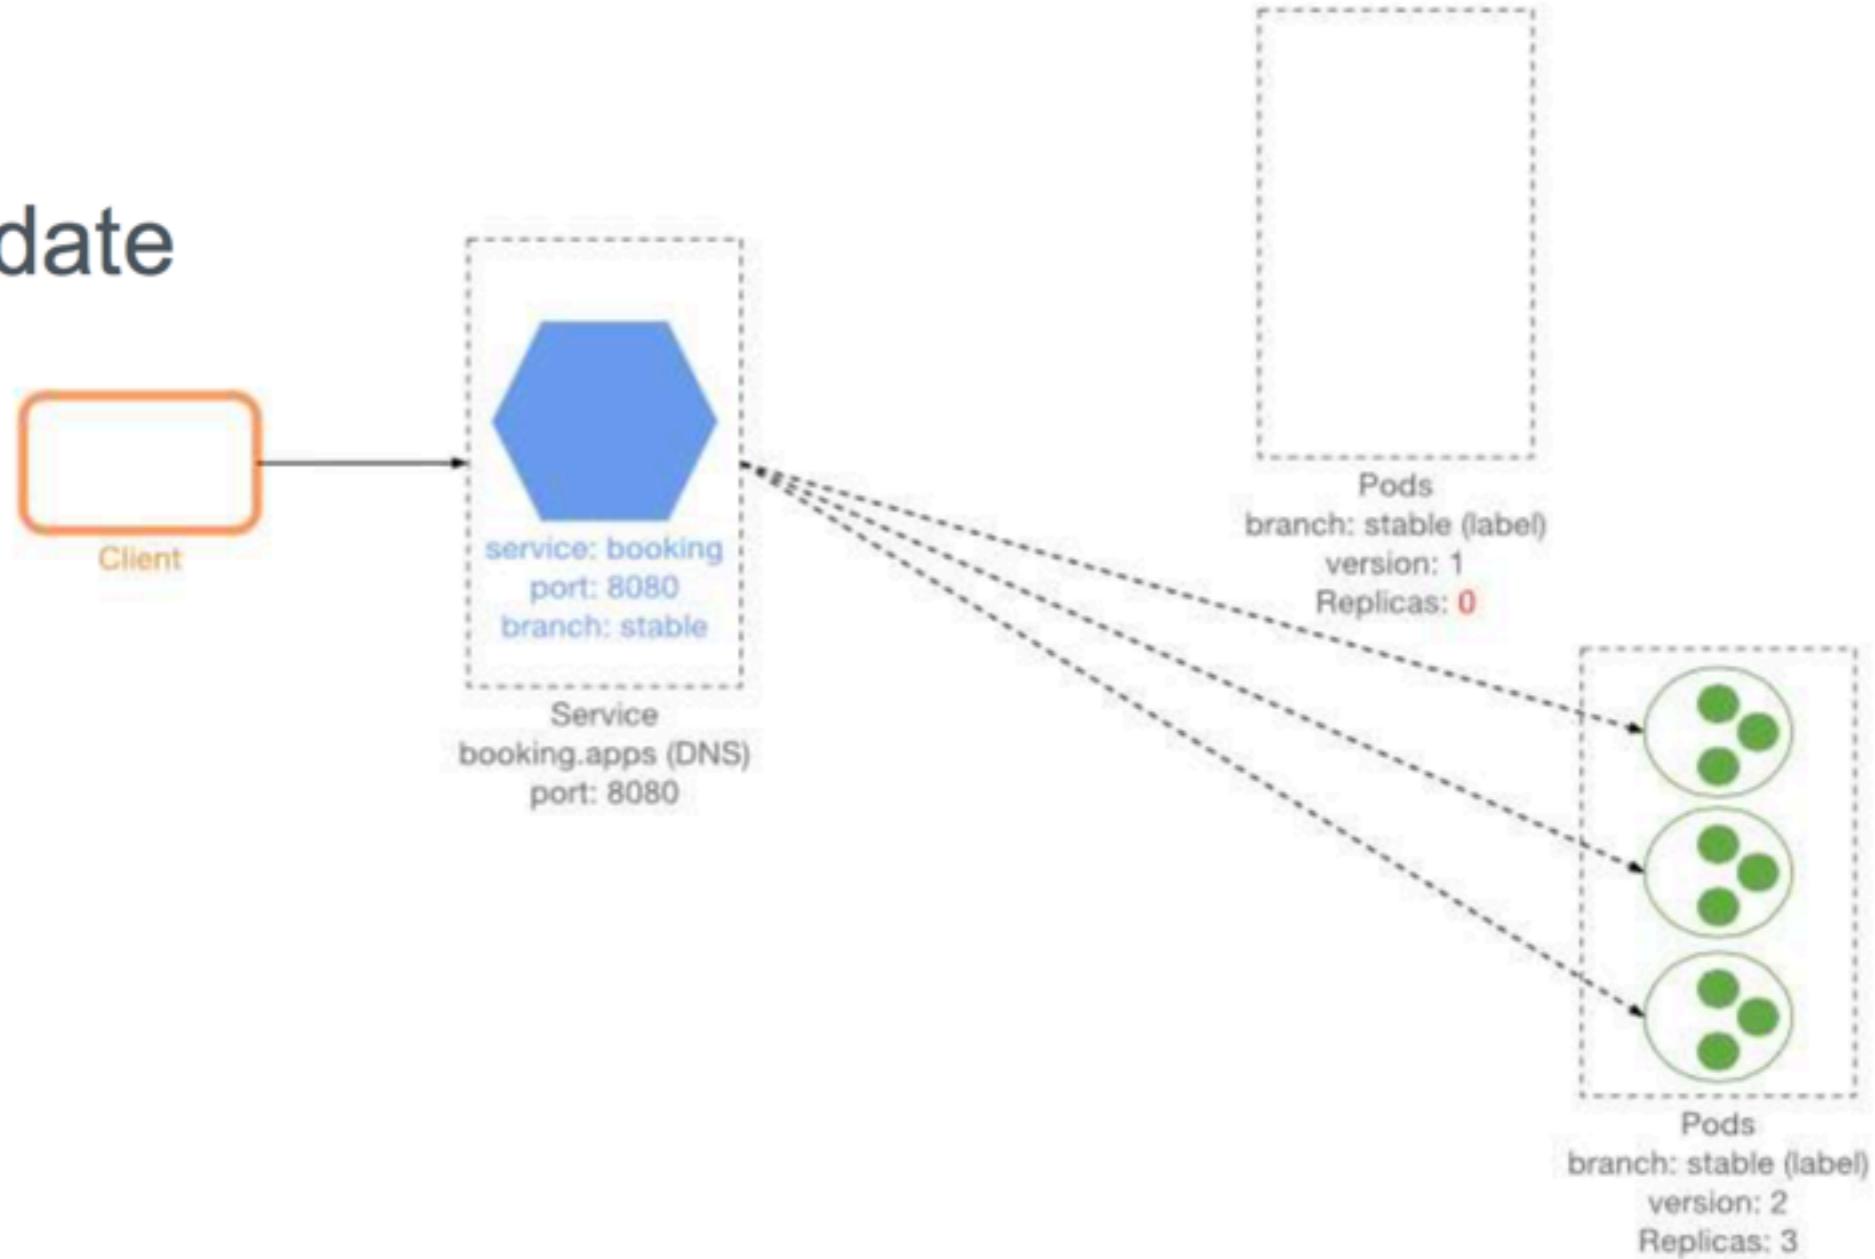
# Rolling update (6)

## Rolling Update



# Rolling update (7)

## Rolling Update



# Storage orchestration



# Storage orchestration

Support many type of storages

1. Local storage
2. Network storage (NFS, GlusterFS, Ceph)
3. Cloud storage (GCE, AWS, Azure )



# Self-healing



# **Provide well known ports for Kubernetes services**



# **Service Discovery and Load balance**

**Service** is the connector/proxy for client to connect with Pod

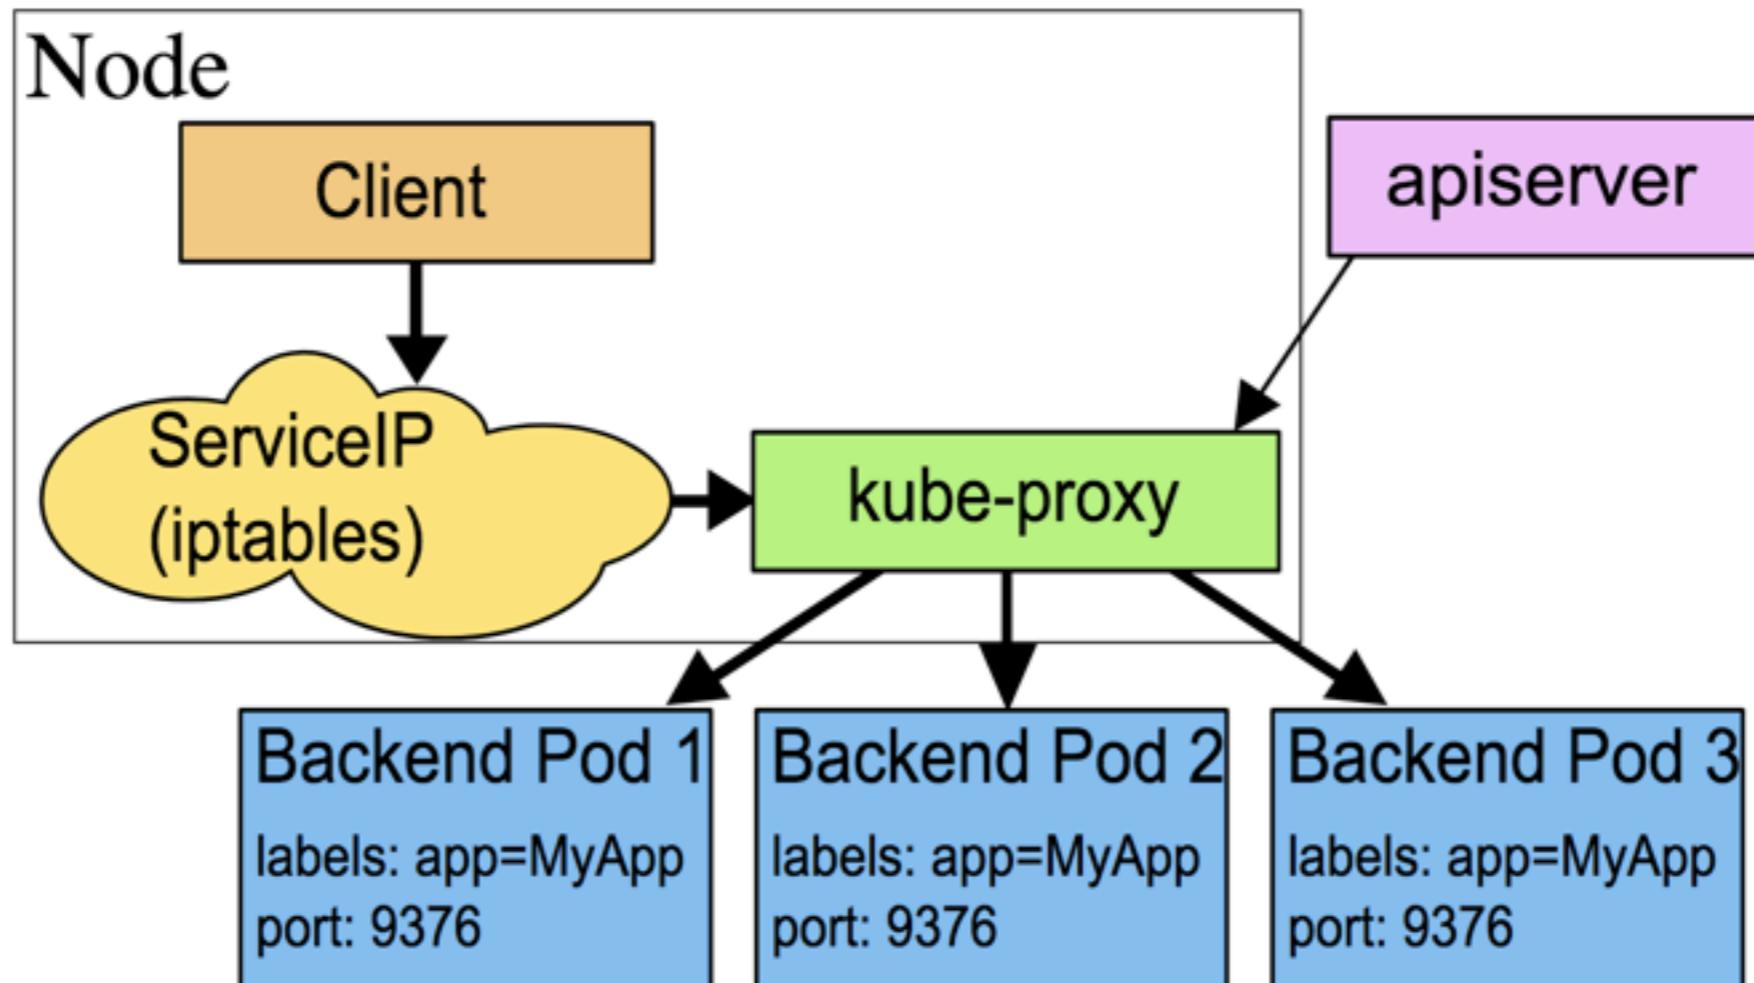
**Discovery** use for service to look Pod

**Support load balance** between Pods (replica)

<https://kubernetes.io/docs/concepts/services-networking/service/>



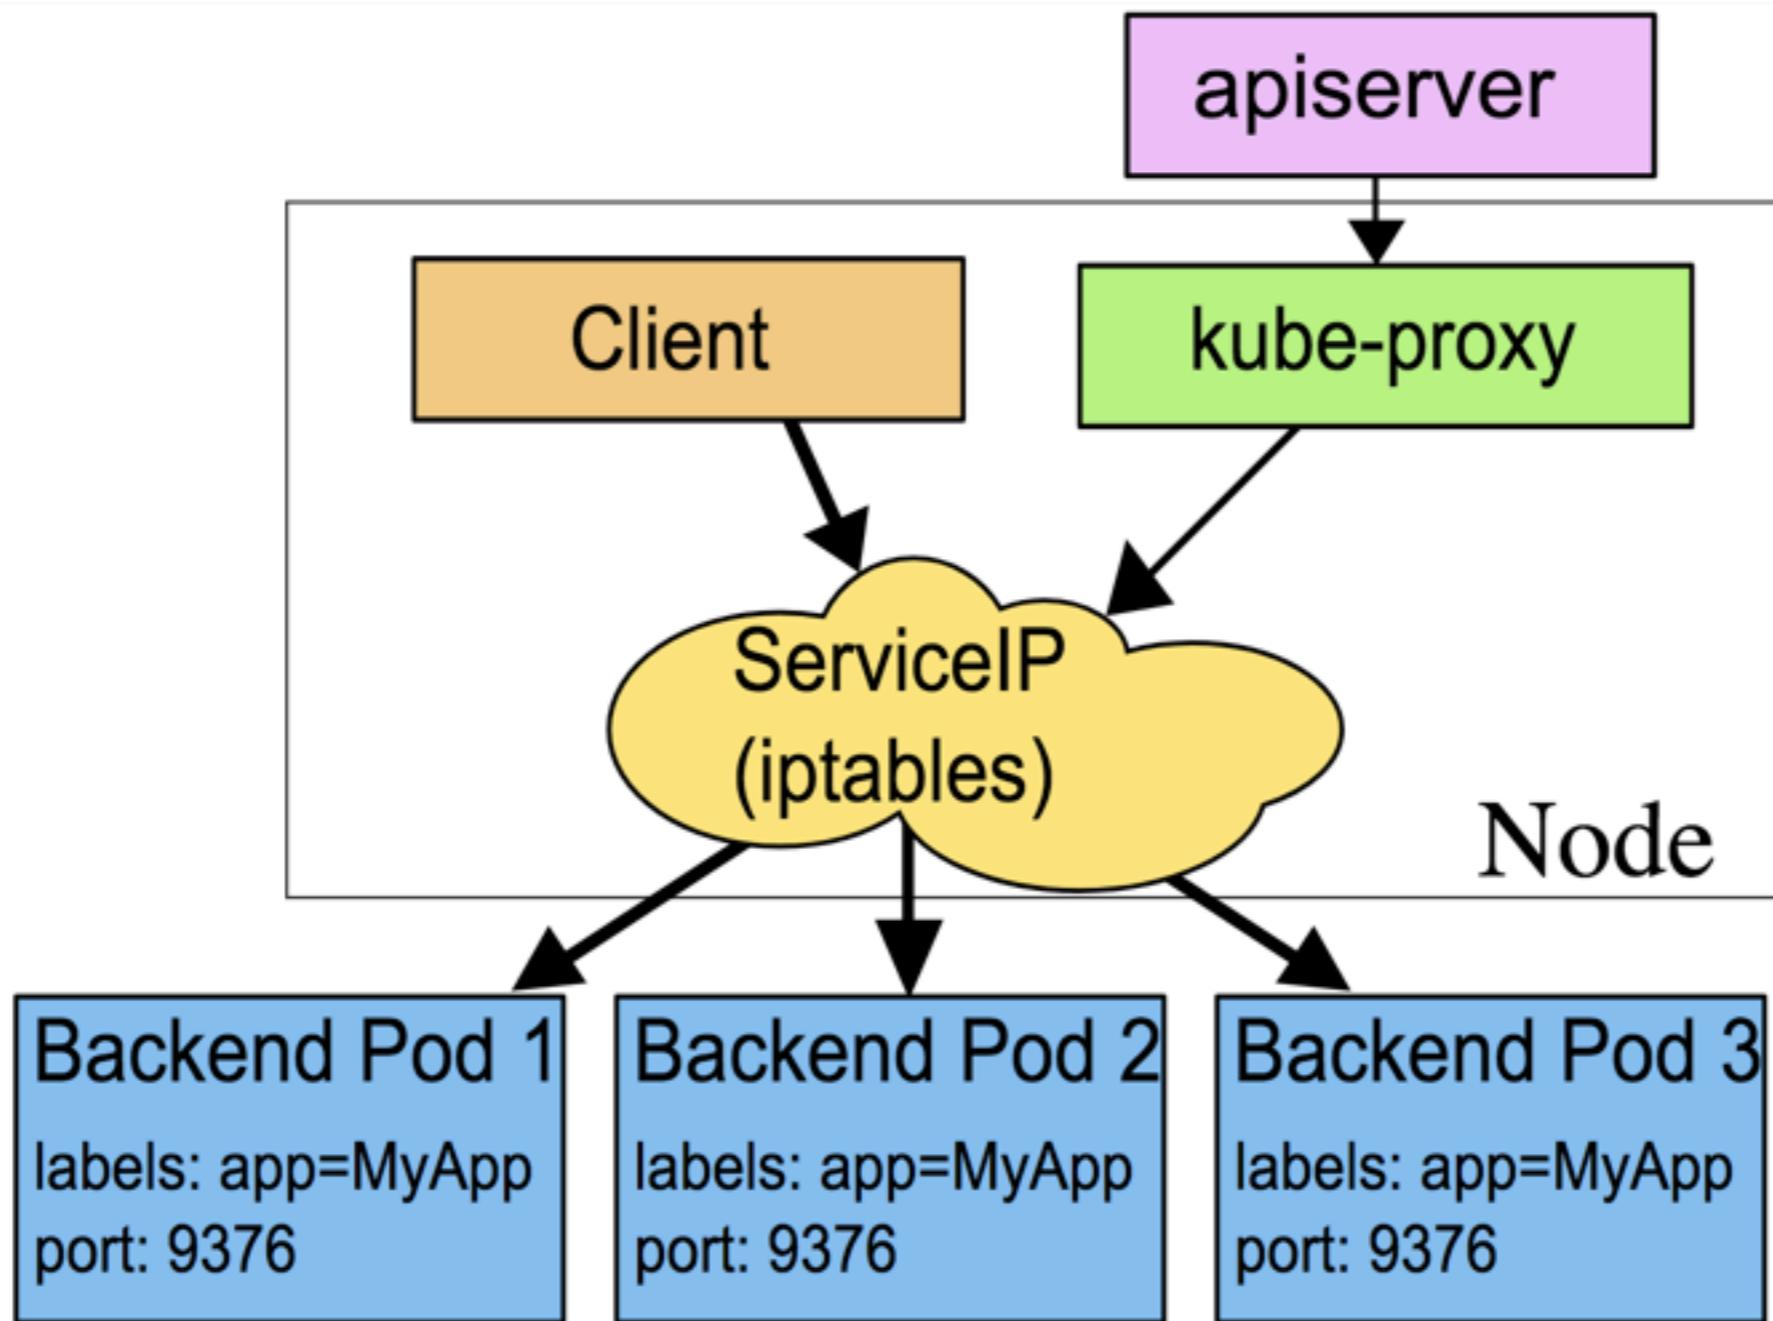
# Proxy-mode: userspace



Note that in the above diagram, `clusterIP` is shown as `ServiceIP`.



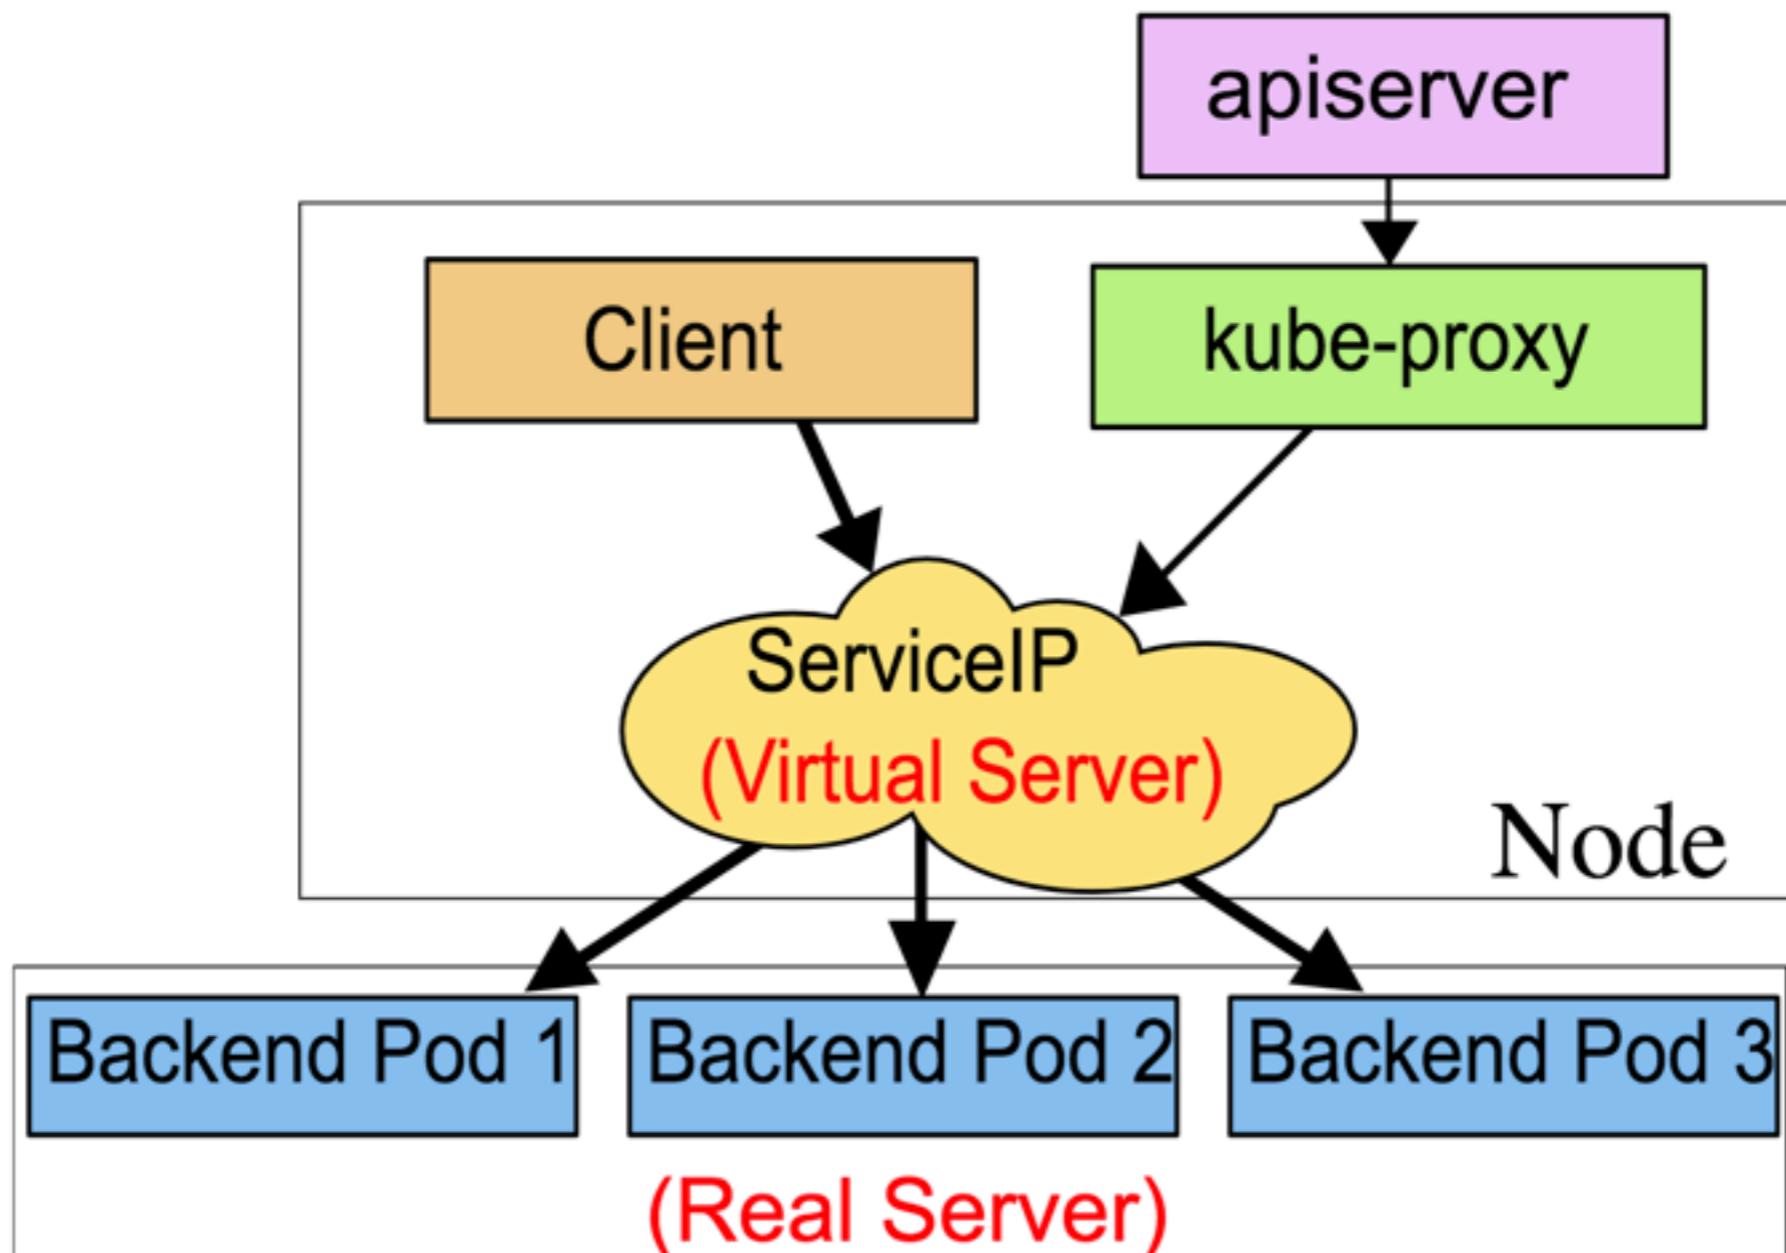
# Proxy-mode: iptables



Note that in the above diagram, `clusterIP` is shown as `ServiceIP`.



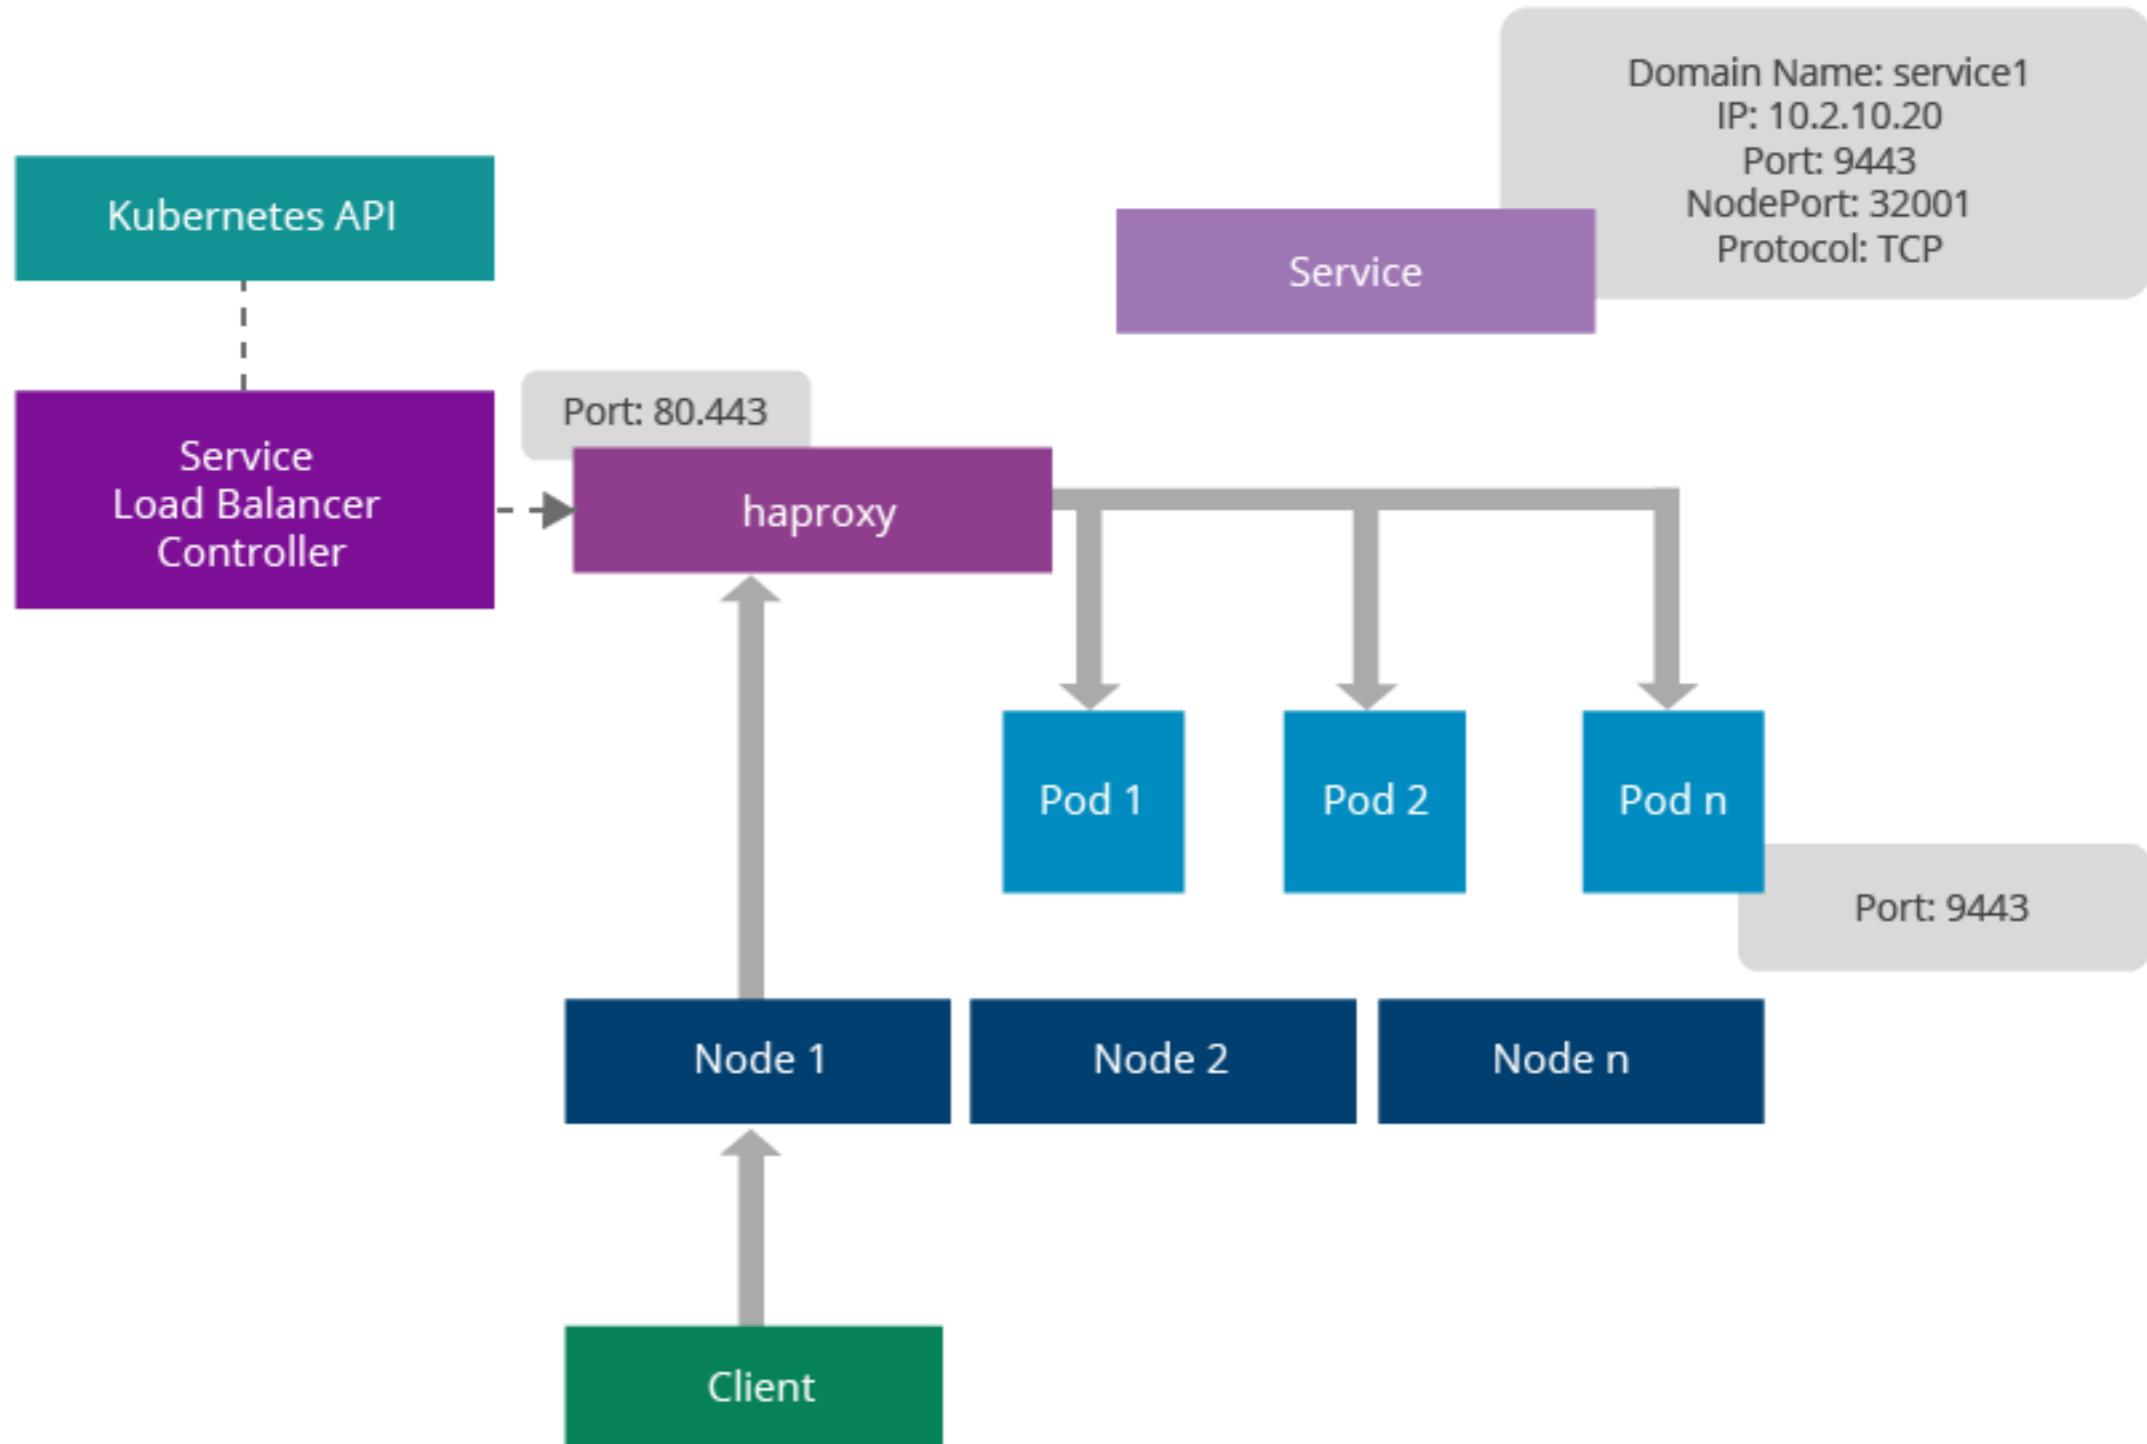
# Proxy-mode: ipvs (beta 1.9)



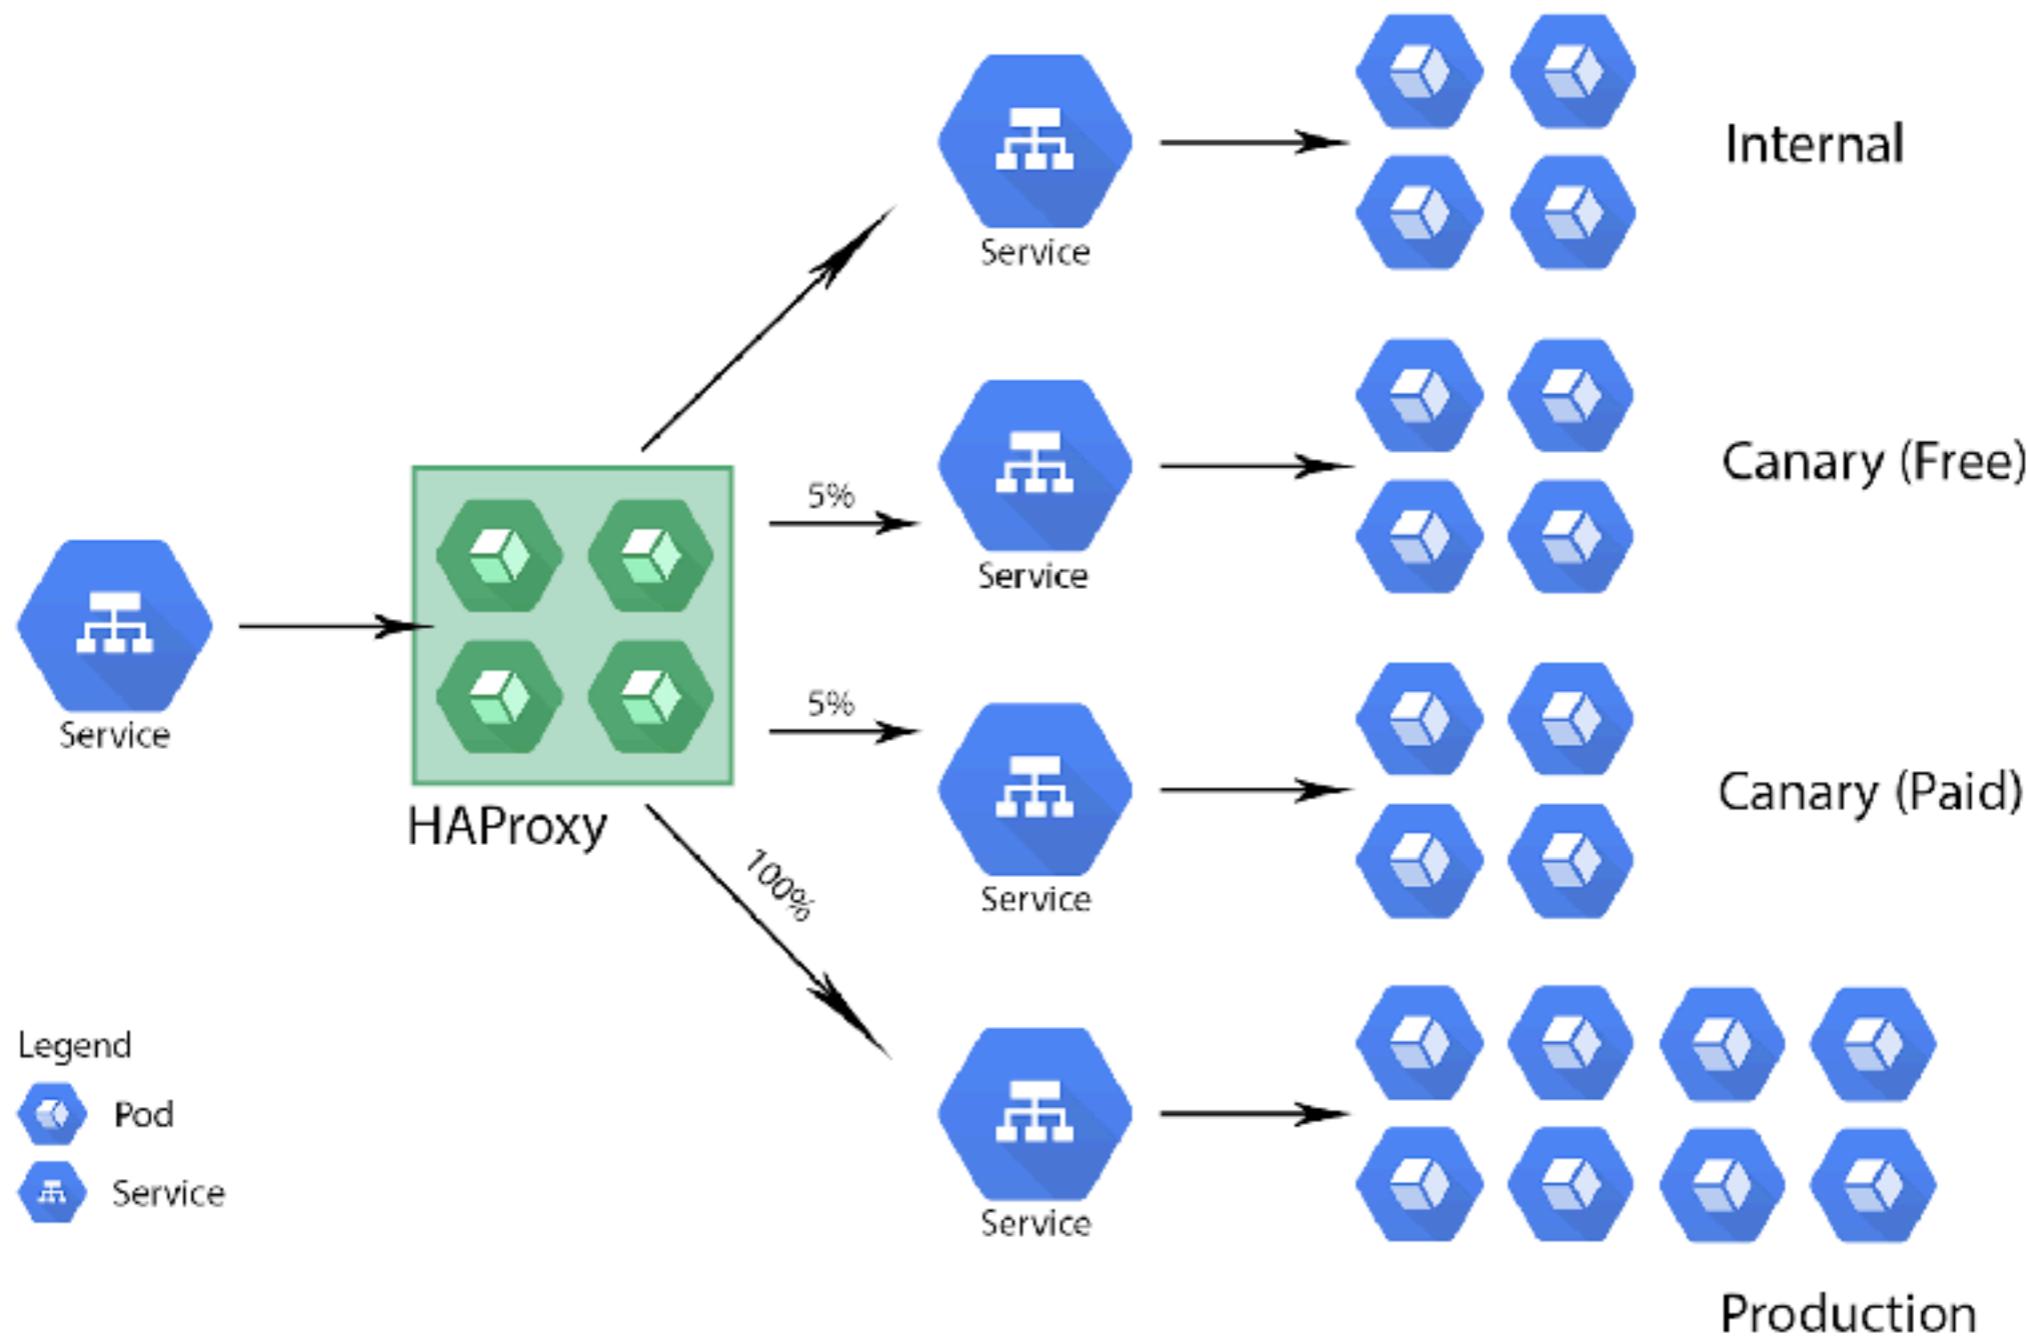
# **Sticky session management using Service Load Balancers**



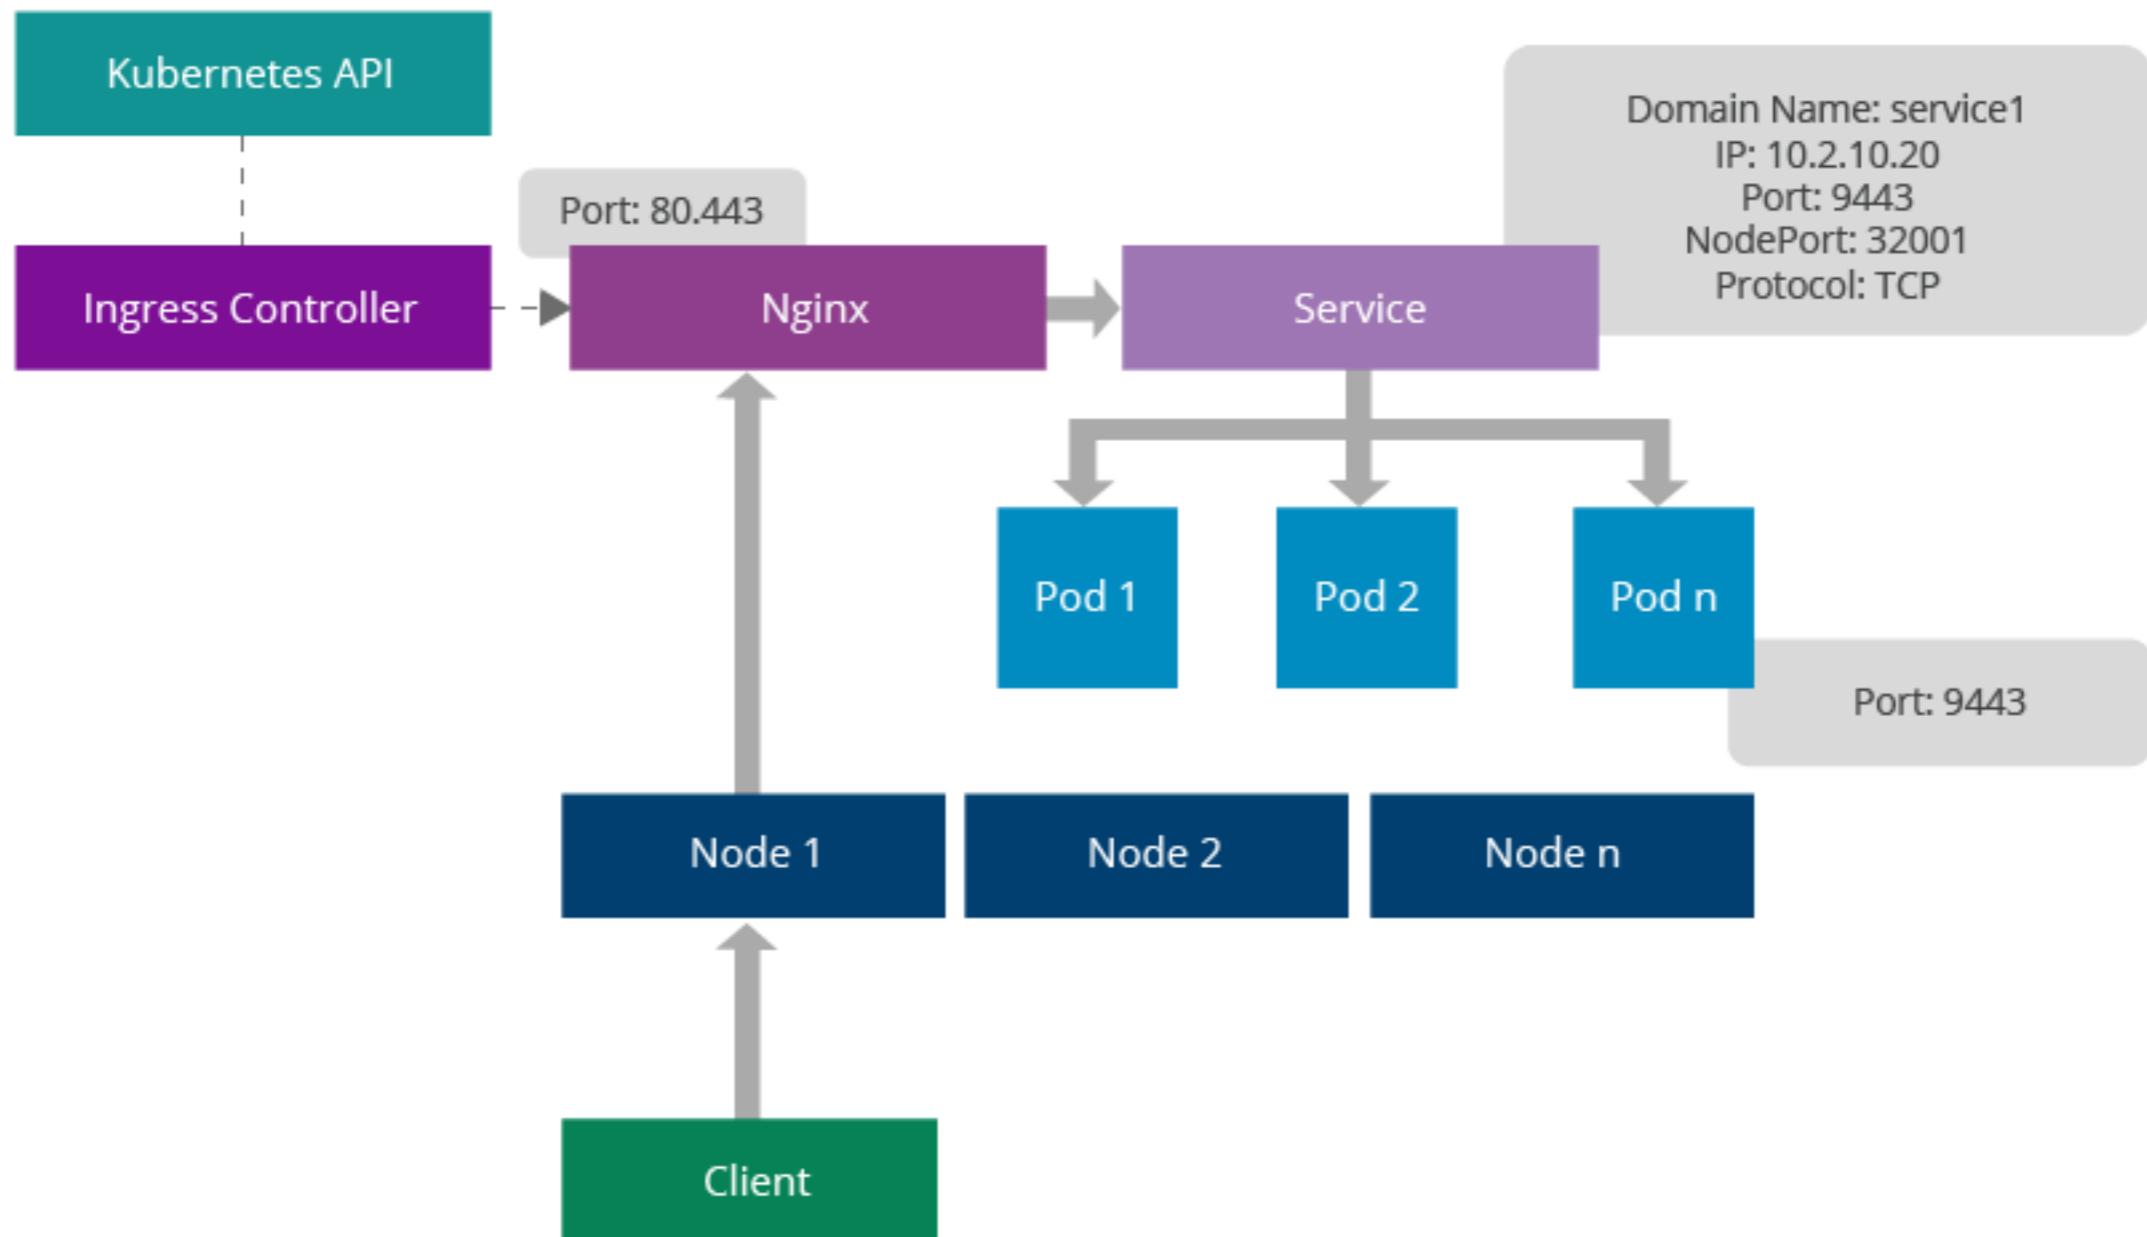
# Service Load Balancer



# Service Load Balancer



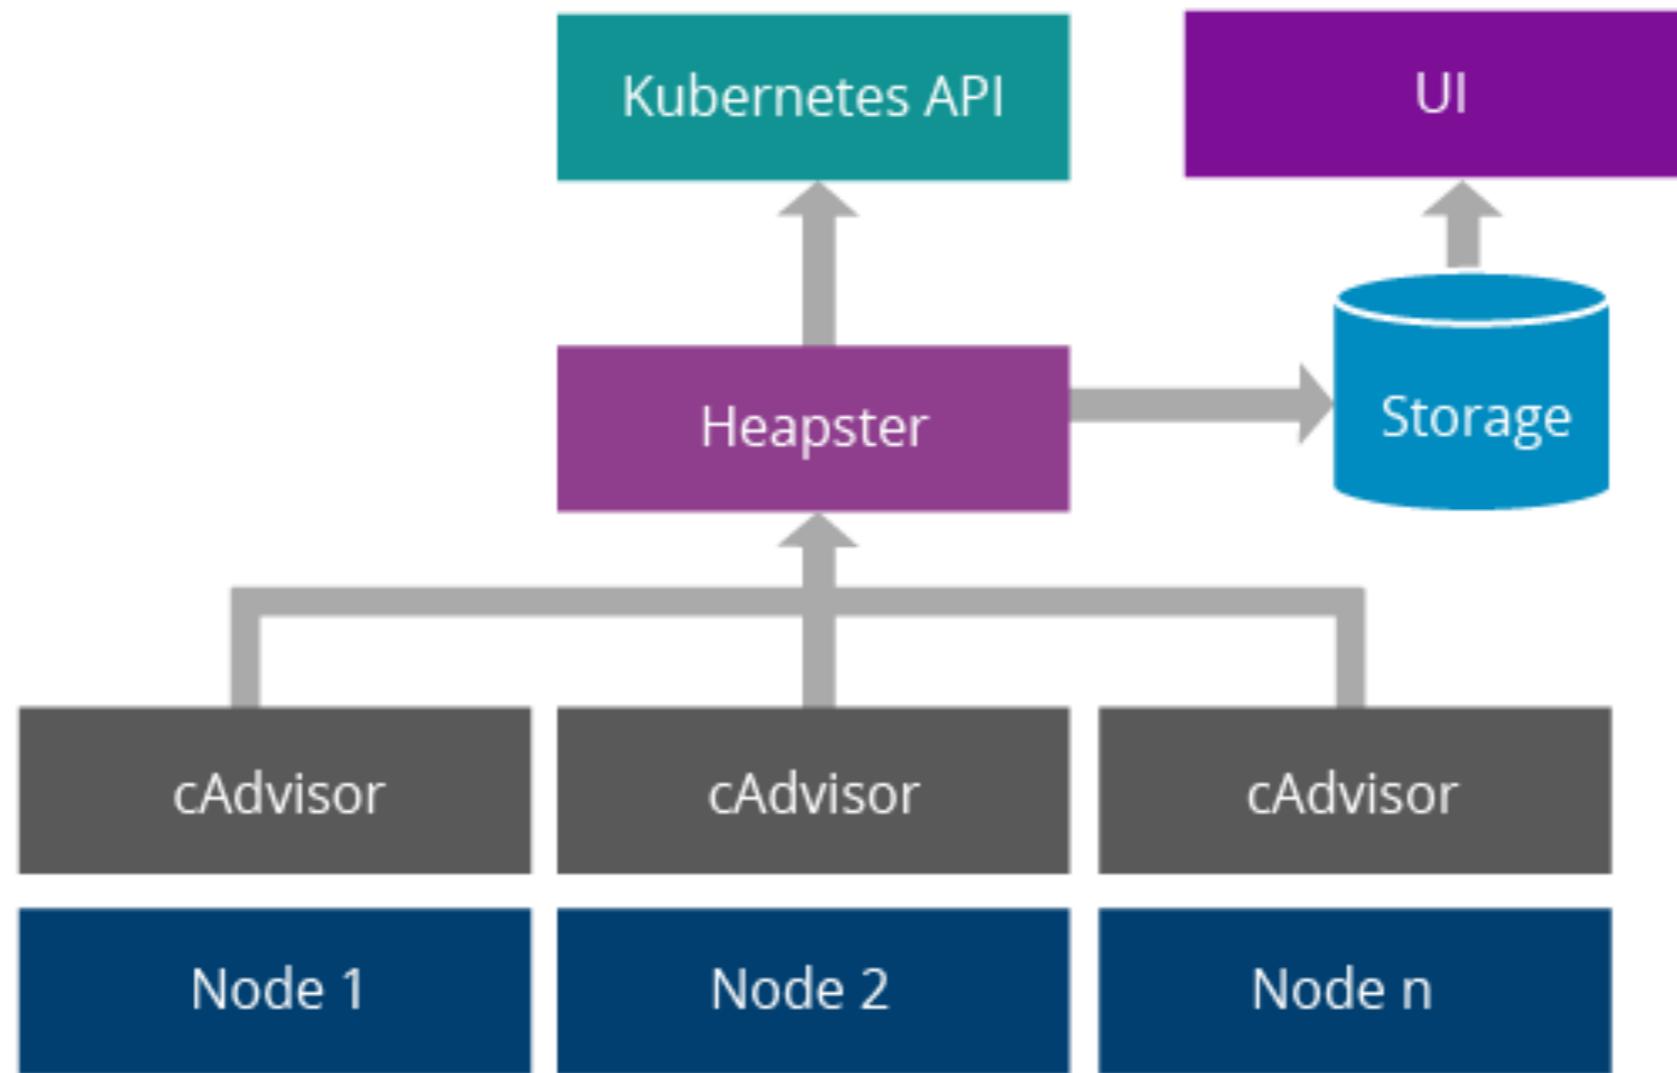
# Ingress controller



# Resource usage monitoring



# Resource usage monitoring



<https://github.com/kubernetes/heapster>



# **Secret and configuration management**



# Secret and configuration mgt

Keep confidential data to running app in  
encryption format

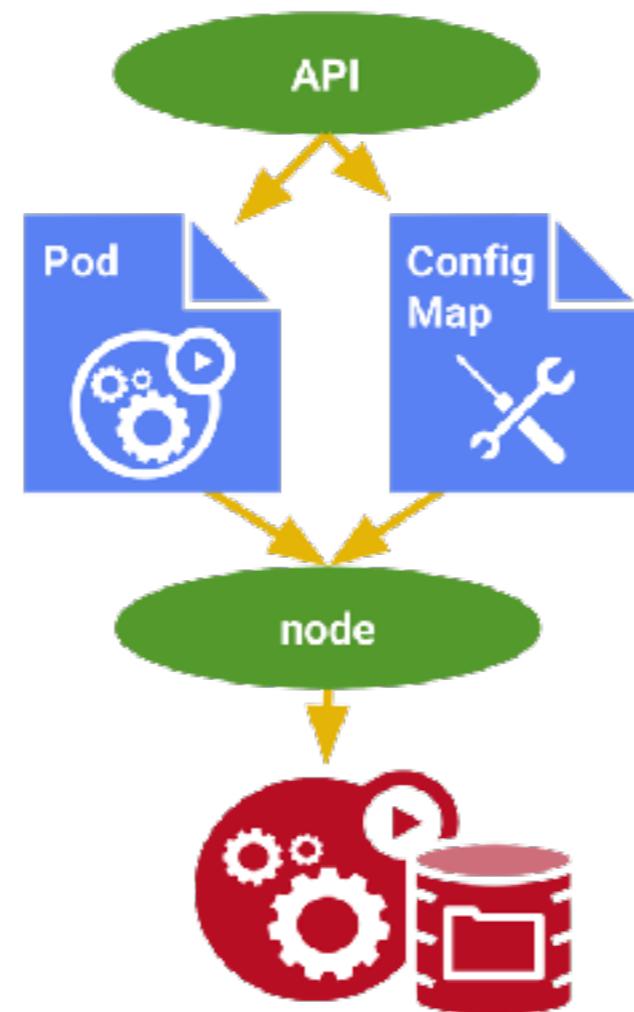
\$kubectl get secret

NAME	TYPE	DATA	AGE
default-token-26swm	kubernetes.io/service-account-token	3	2h



# Secret and configuration mgt

Use **ConfigMap** to define all configuration that reference from Pod

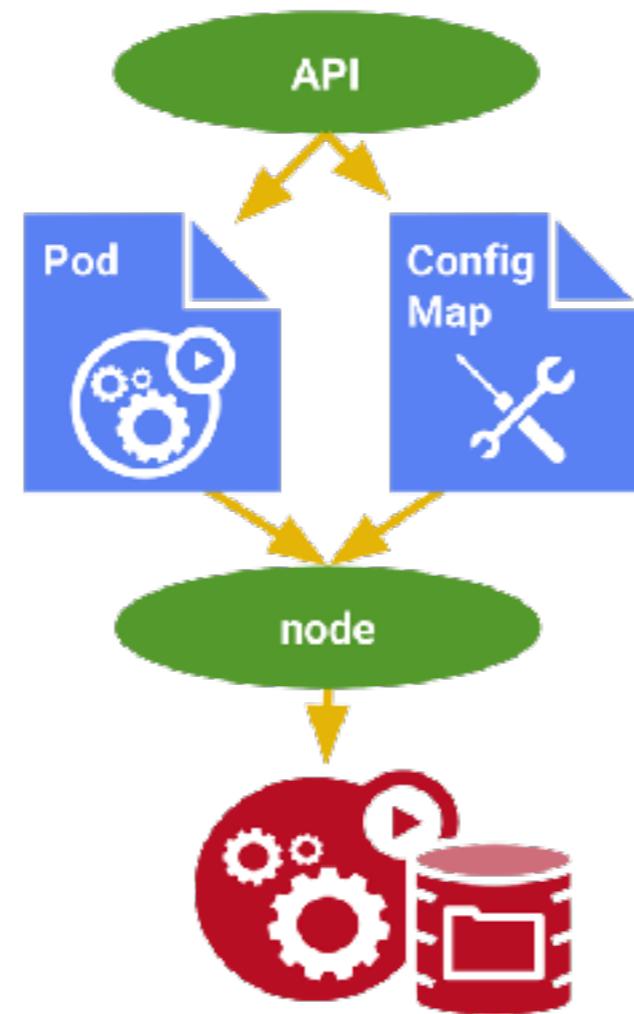


<https://kubernetes.io/docs/tasks/configure-pod-container/configure-pod-configmap/>



# Secret and configuration mgt

Come from **12 factor app** :: config come from the environment



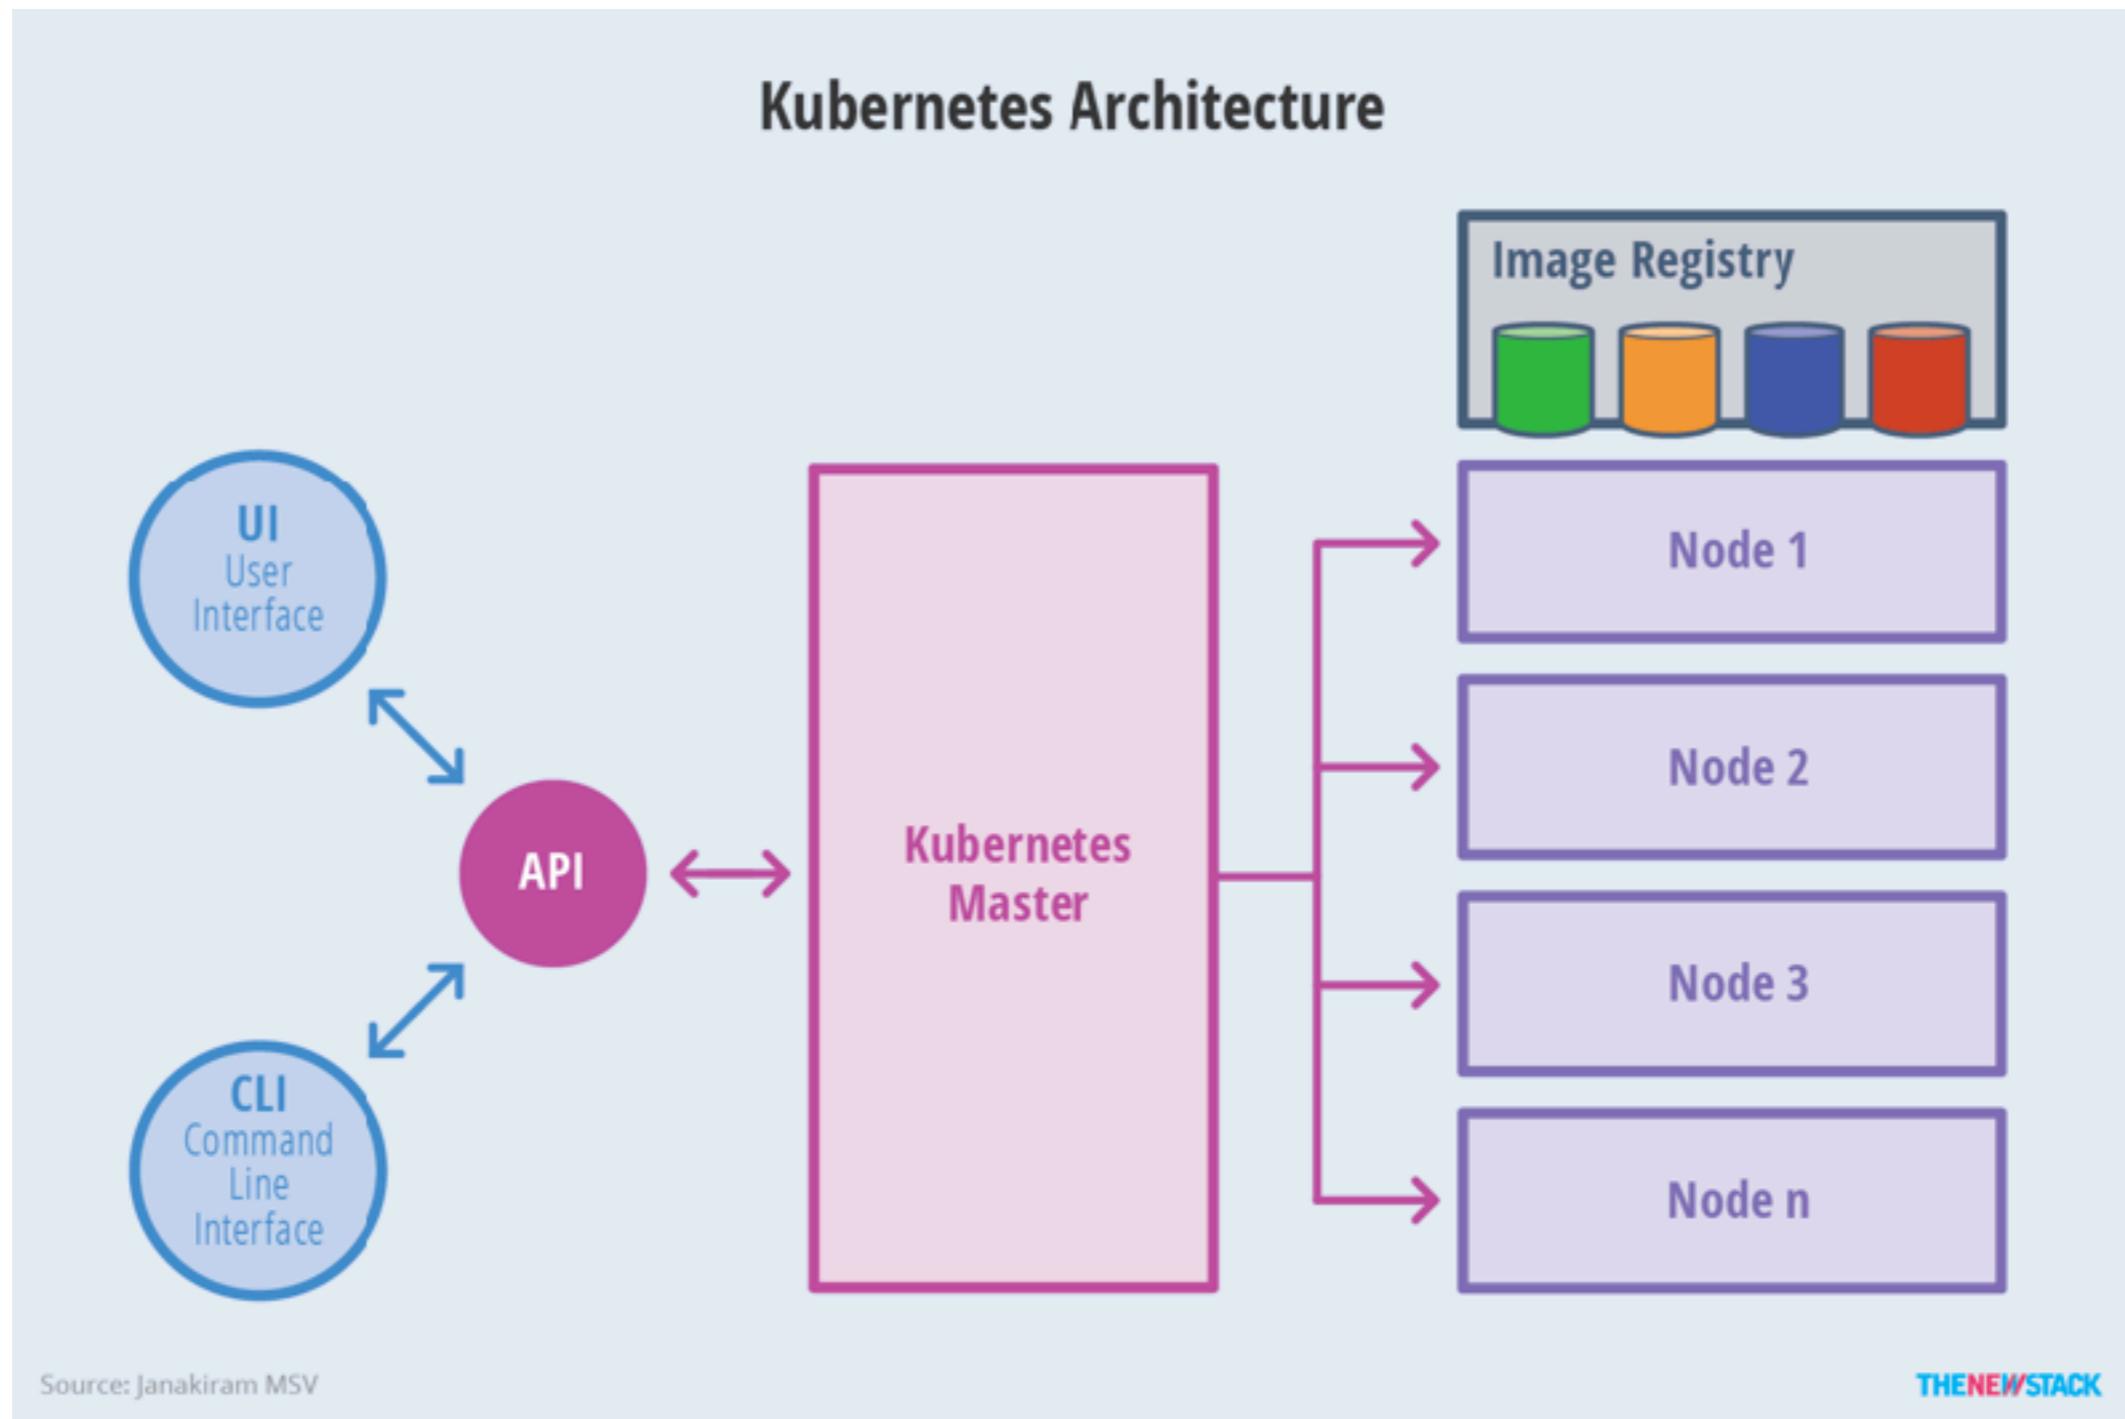
<https://12factor.net/>



# System architecture



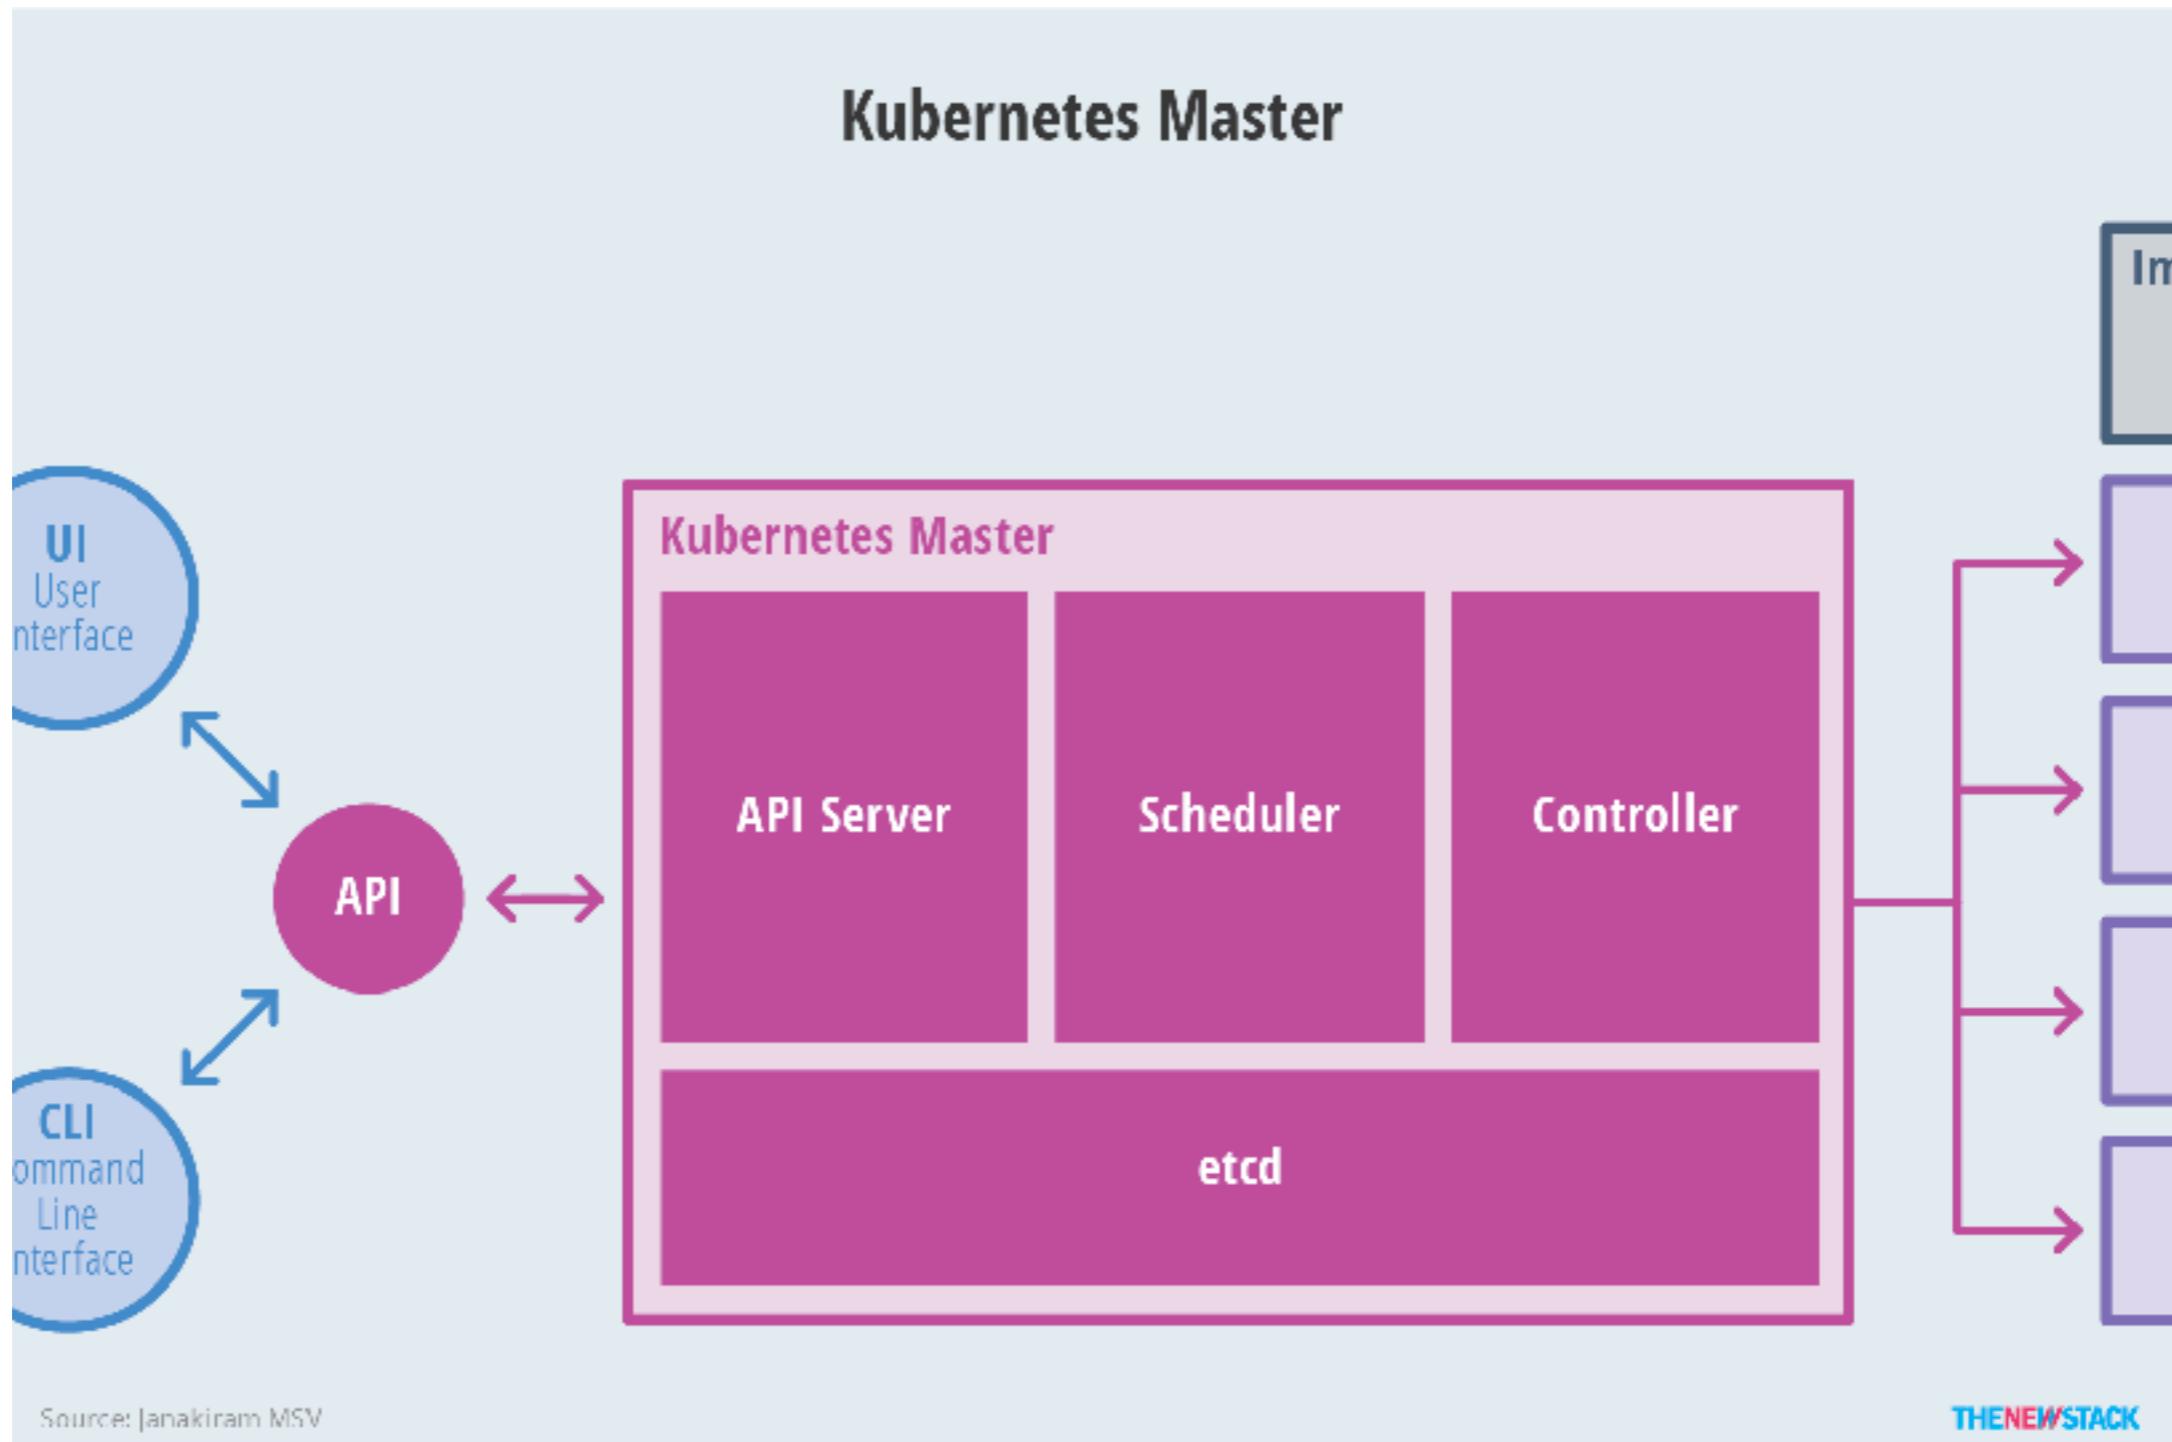
# System architecture



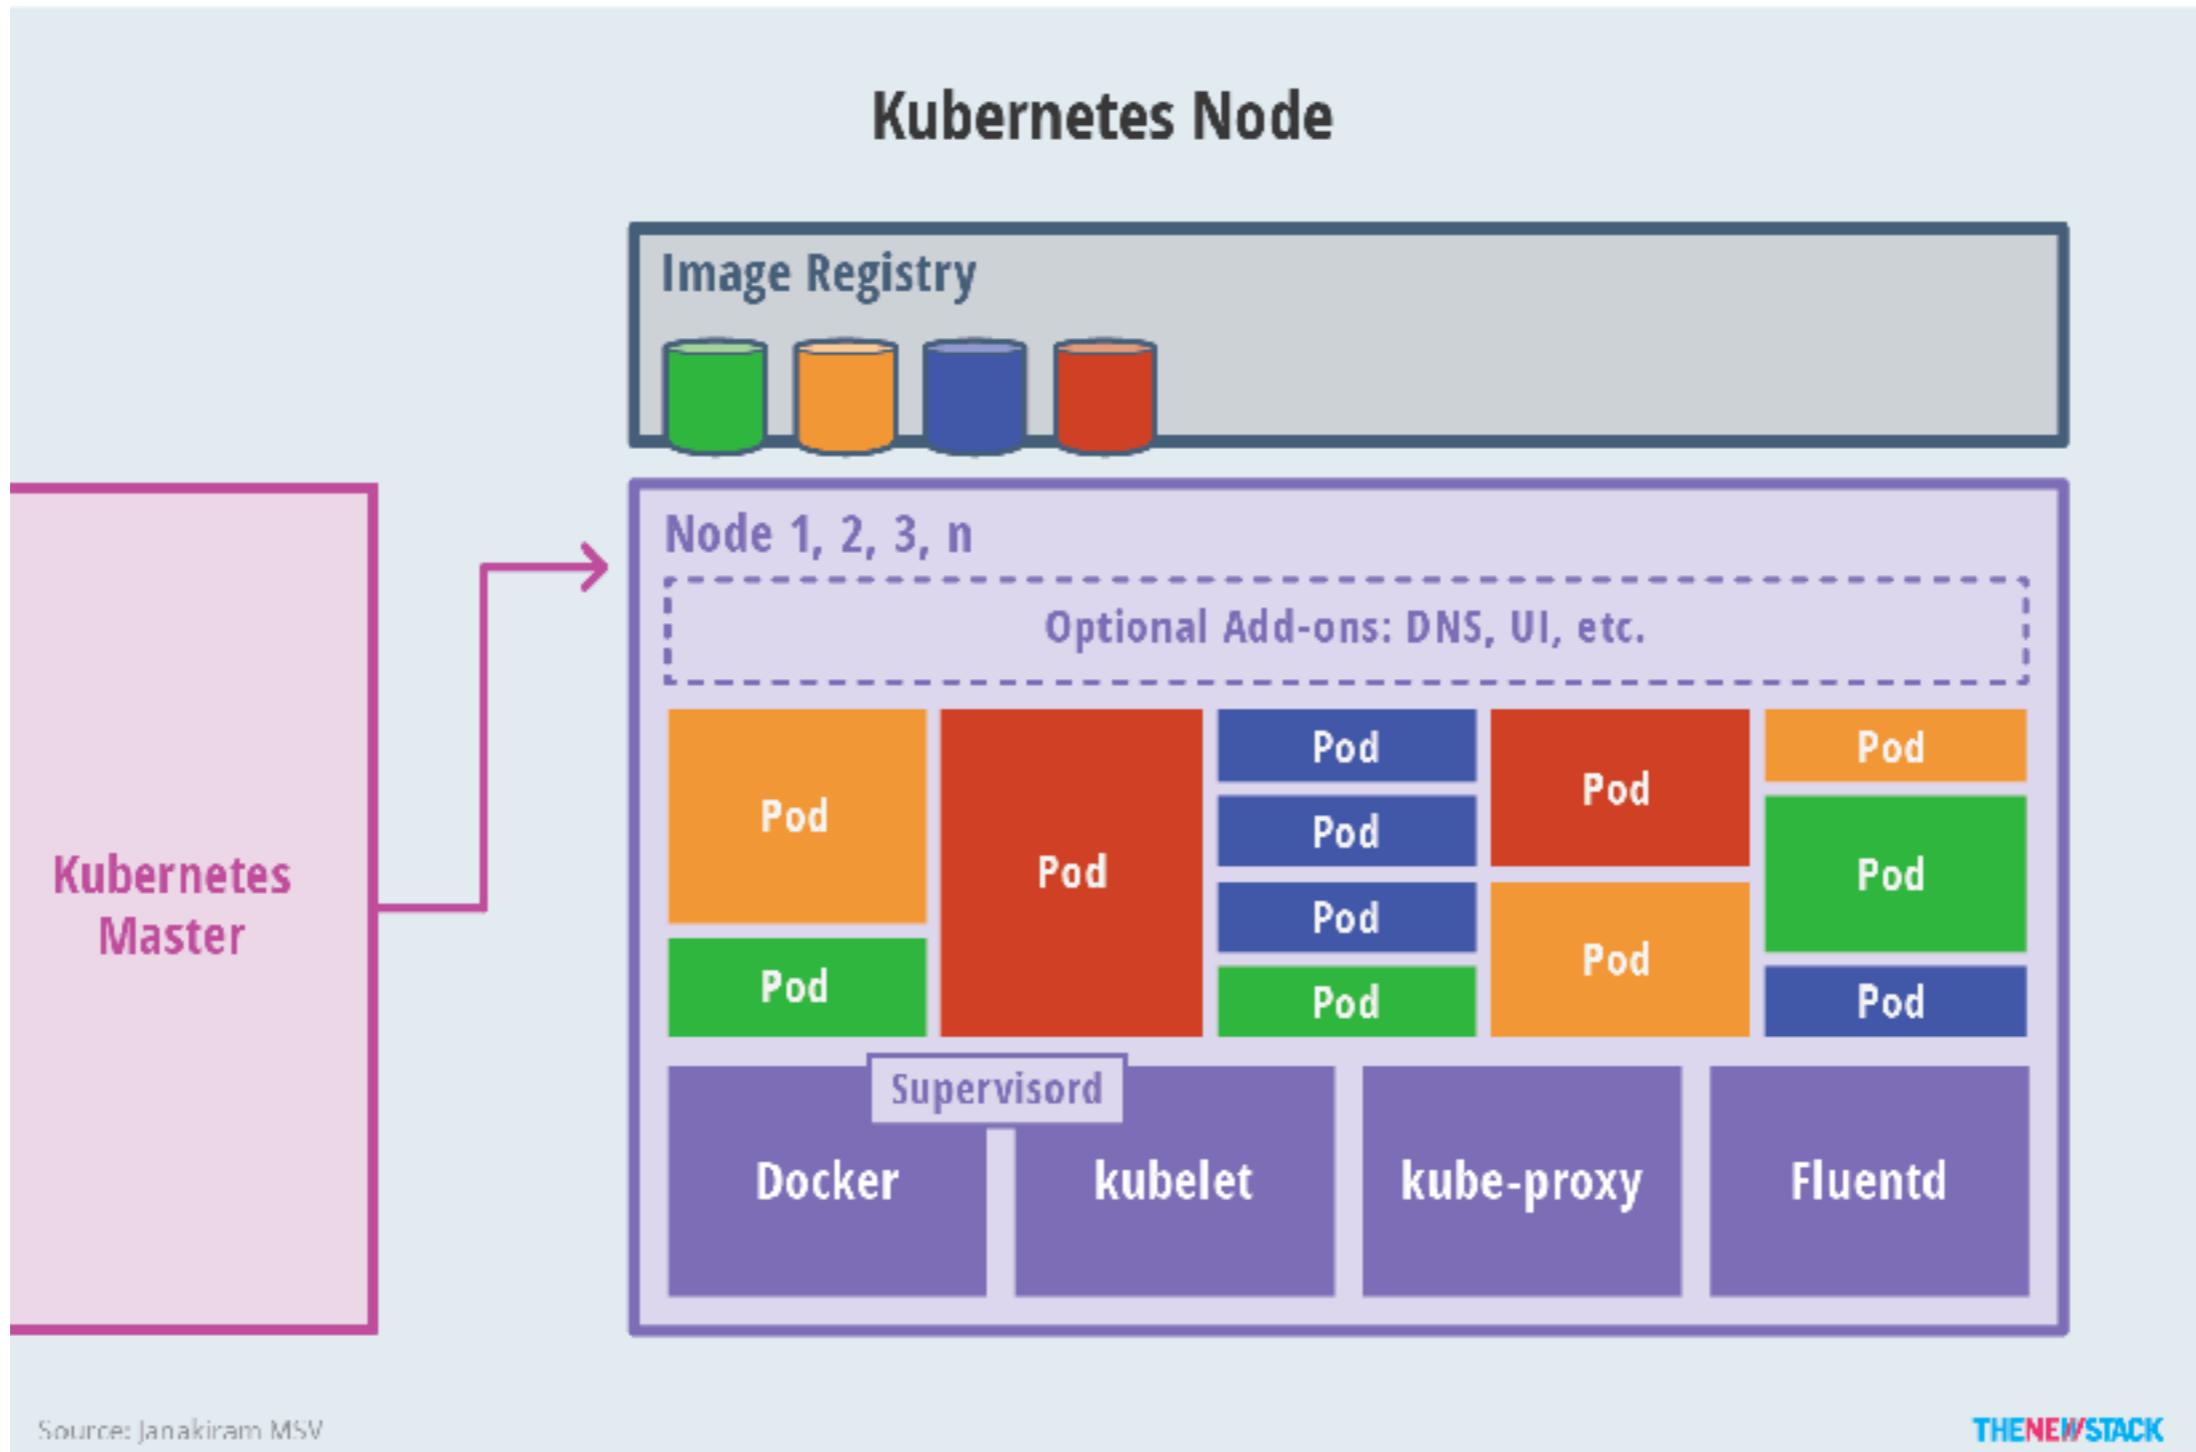
<https://thenewstack.io/kubernetes-an-overview/>



# Master



# Node



<https://thenewstack.io/kubernetes-an-overview/>



# Core concepts of Kubernetes



# Core concepts

Pods vs containers

Services

Replication Controllers

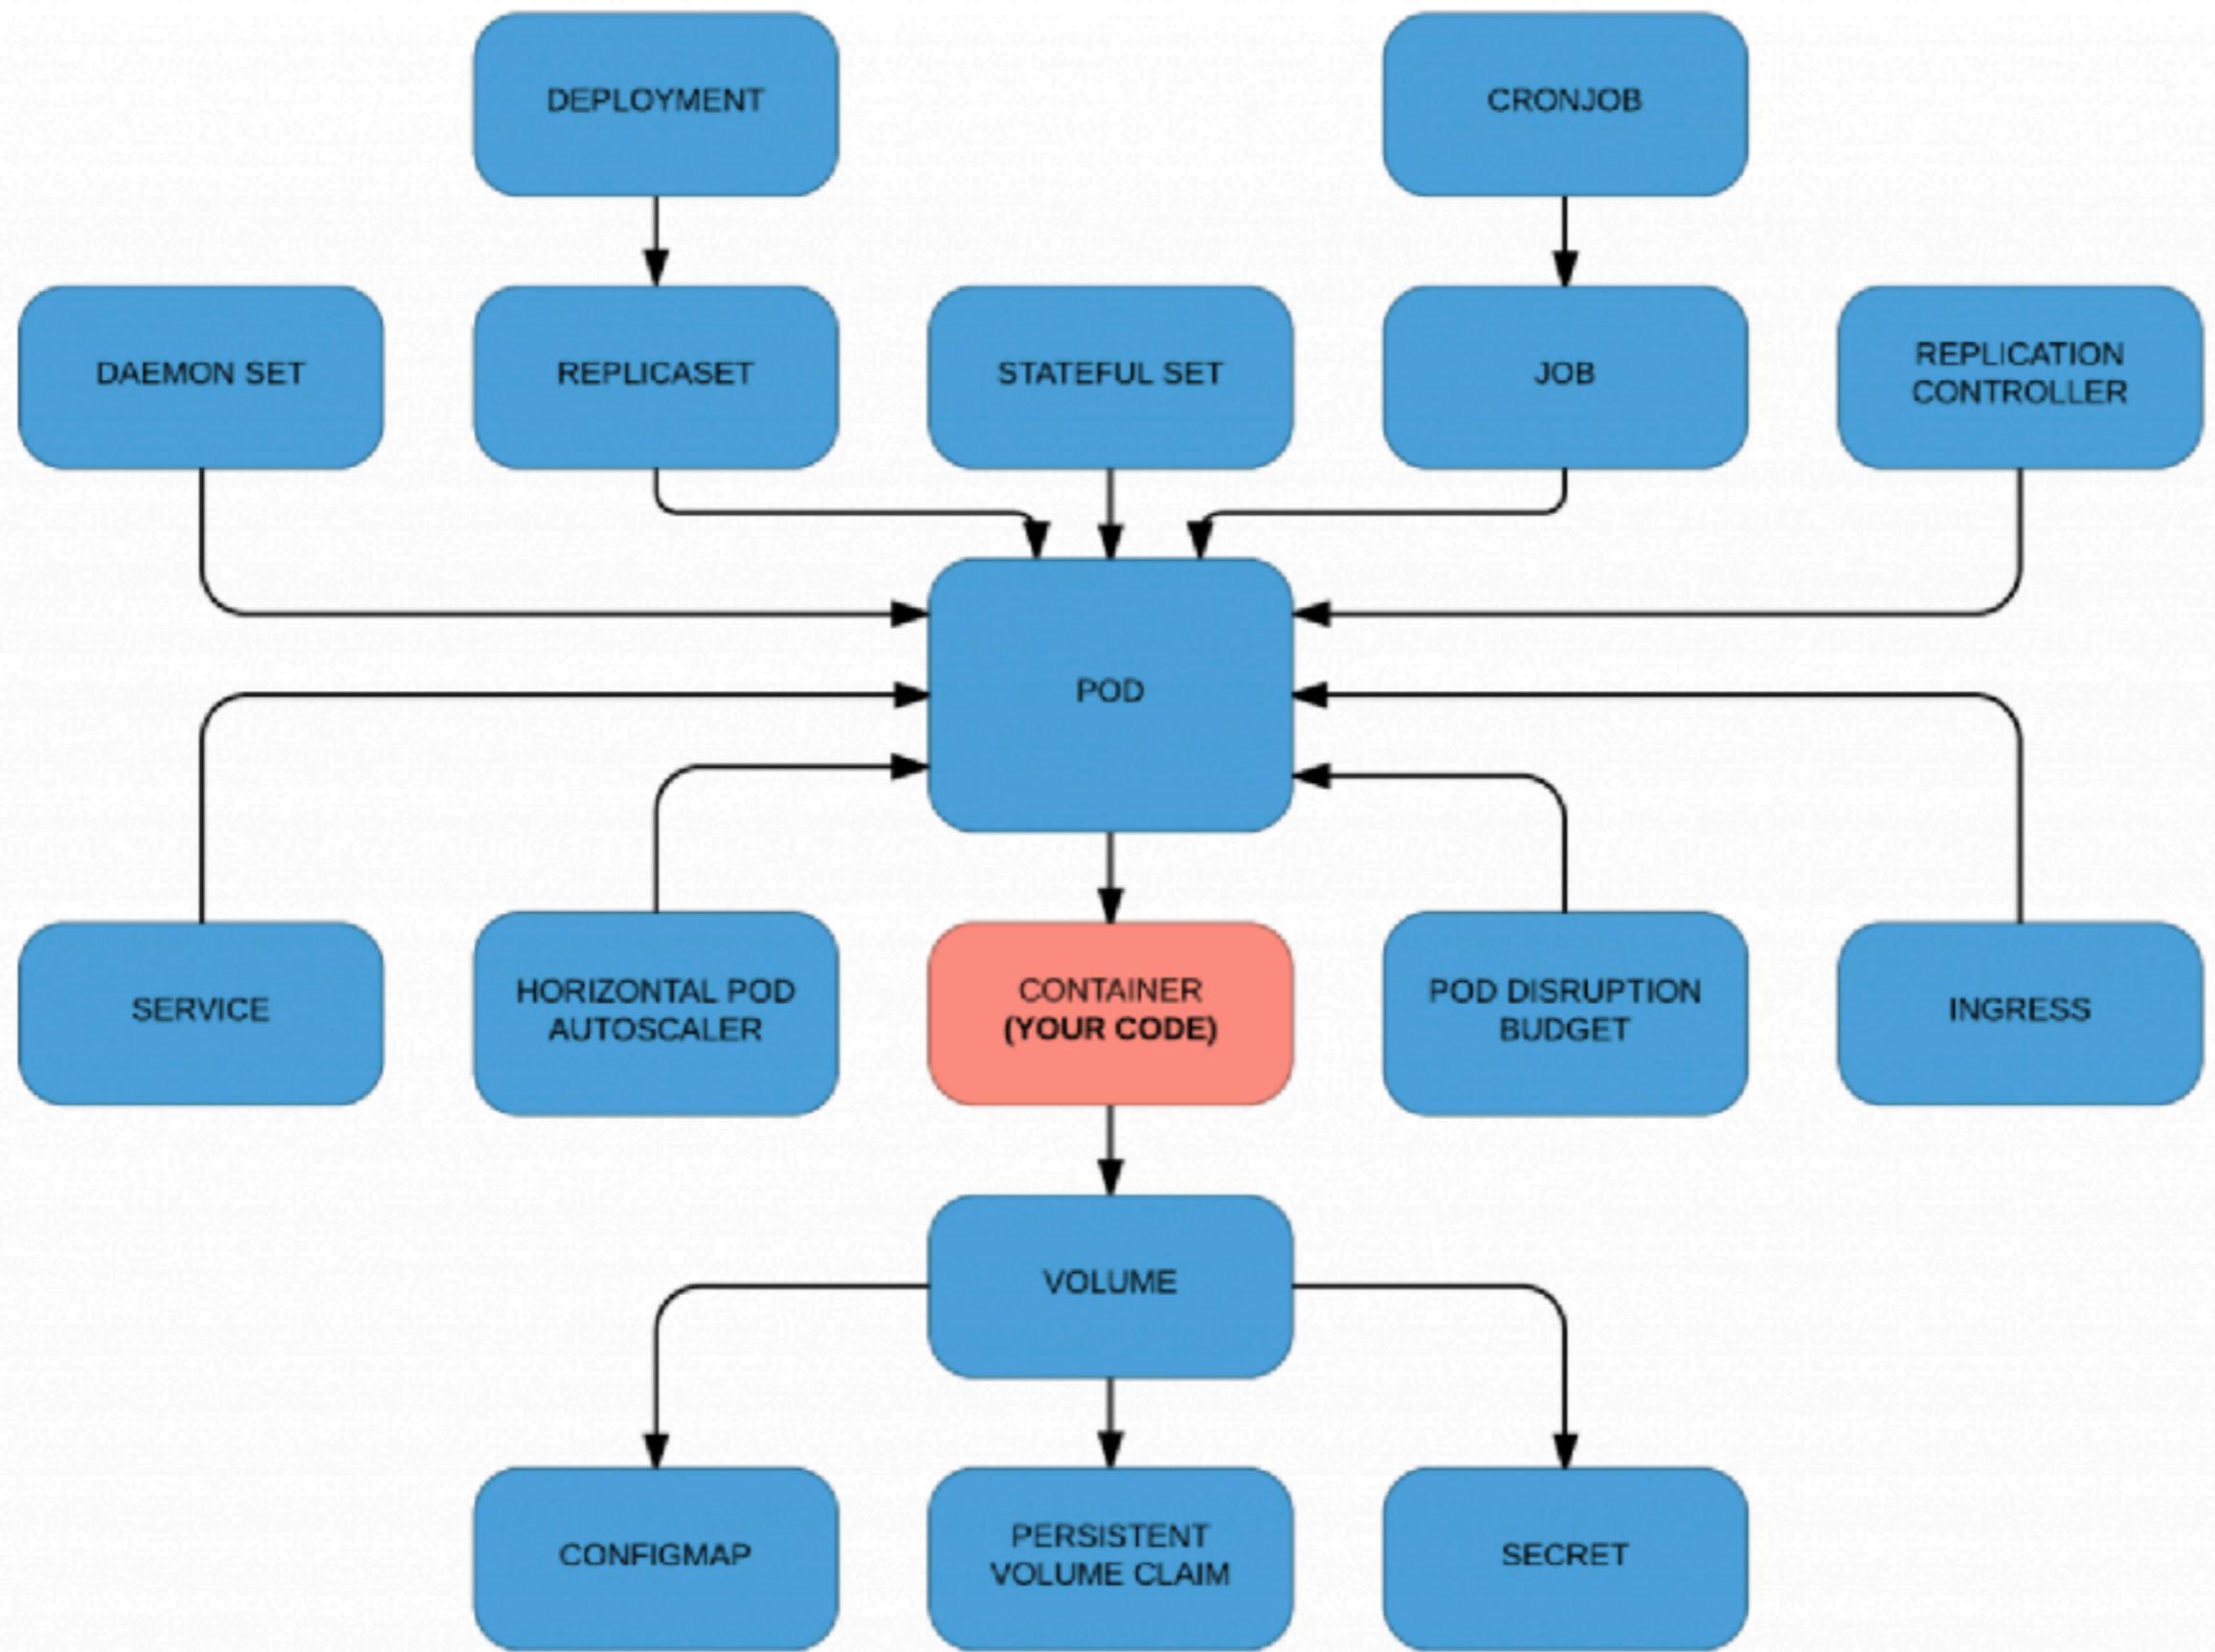
Deployments

ConfigMap and Secret

Volumes

StatefulSets



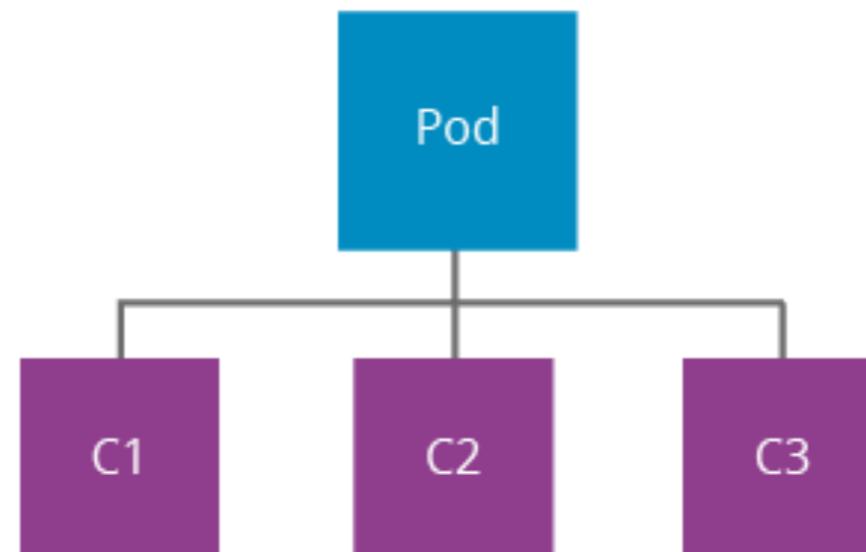


# Pods vs containers



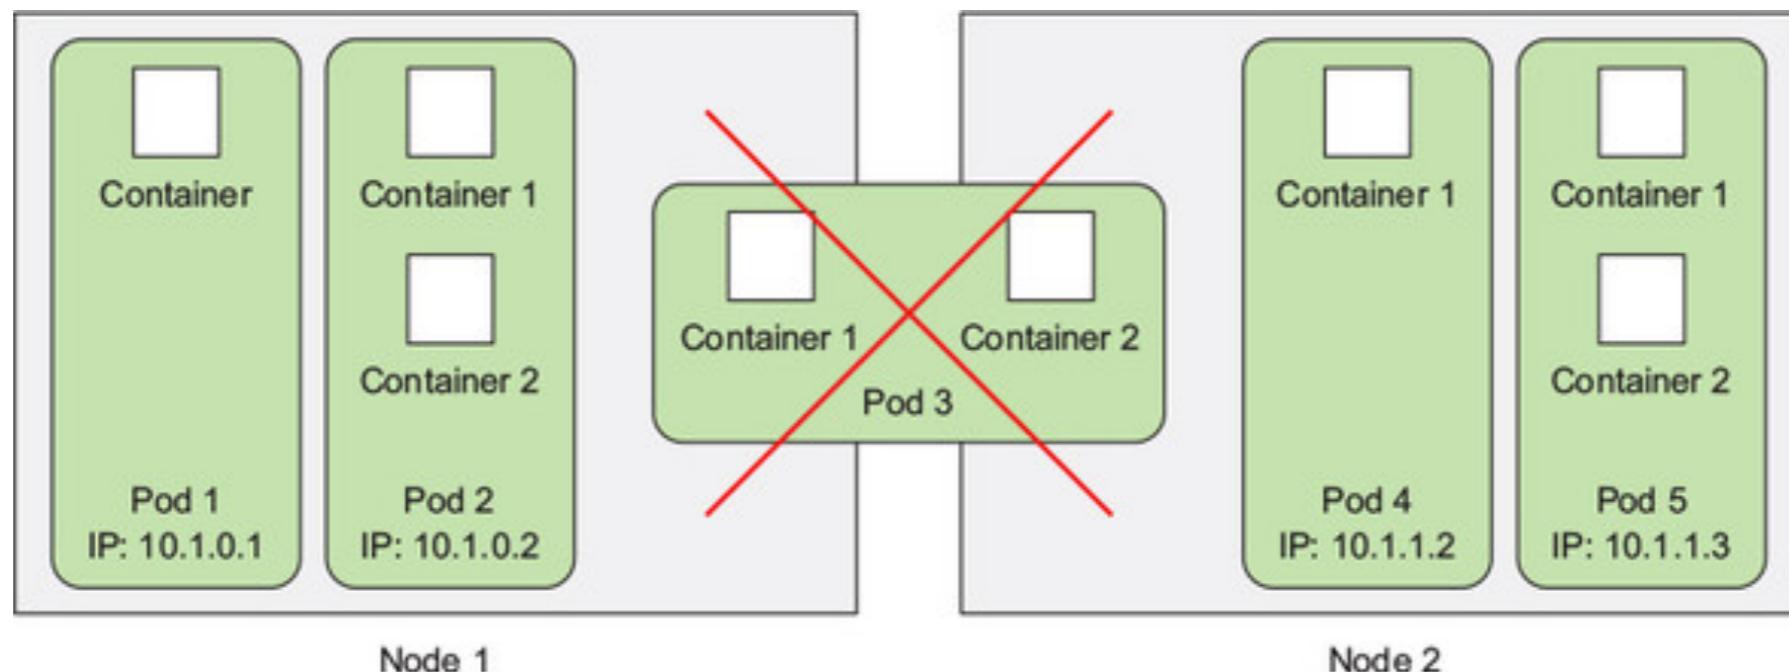
# Pods

Small group of co-located containers  
Optionally shared volume between containers  
Basic deployment unit in Kubernetes



# Pods

1 pods = 1 container  
1 pods = N containers



# All containers in same Pods

Share process ID

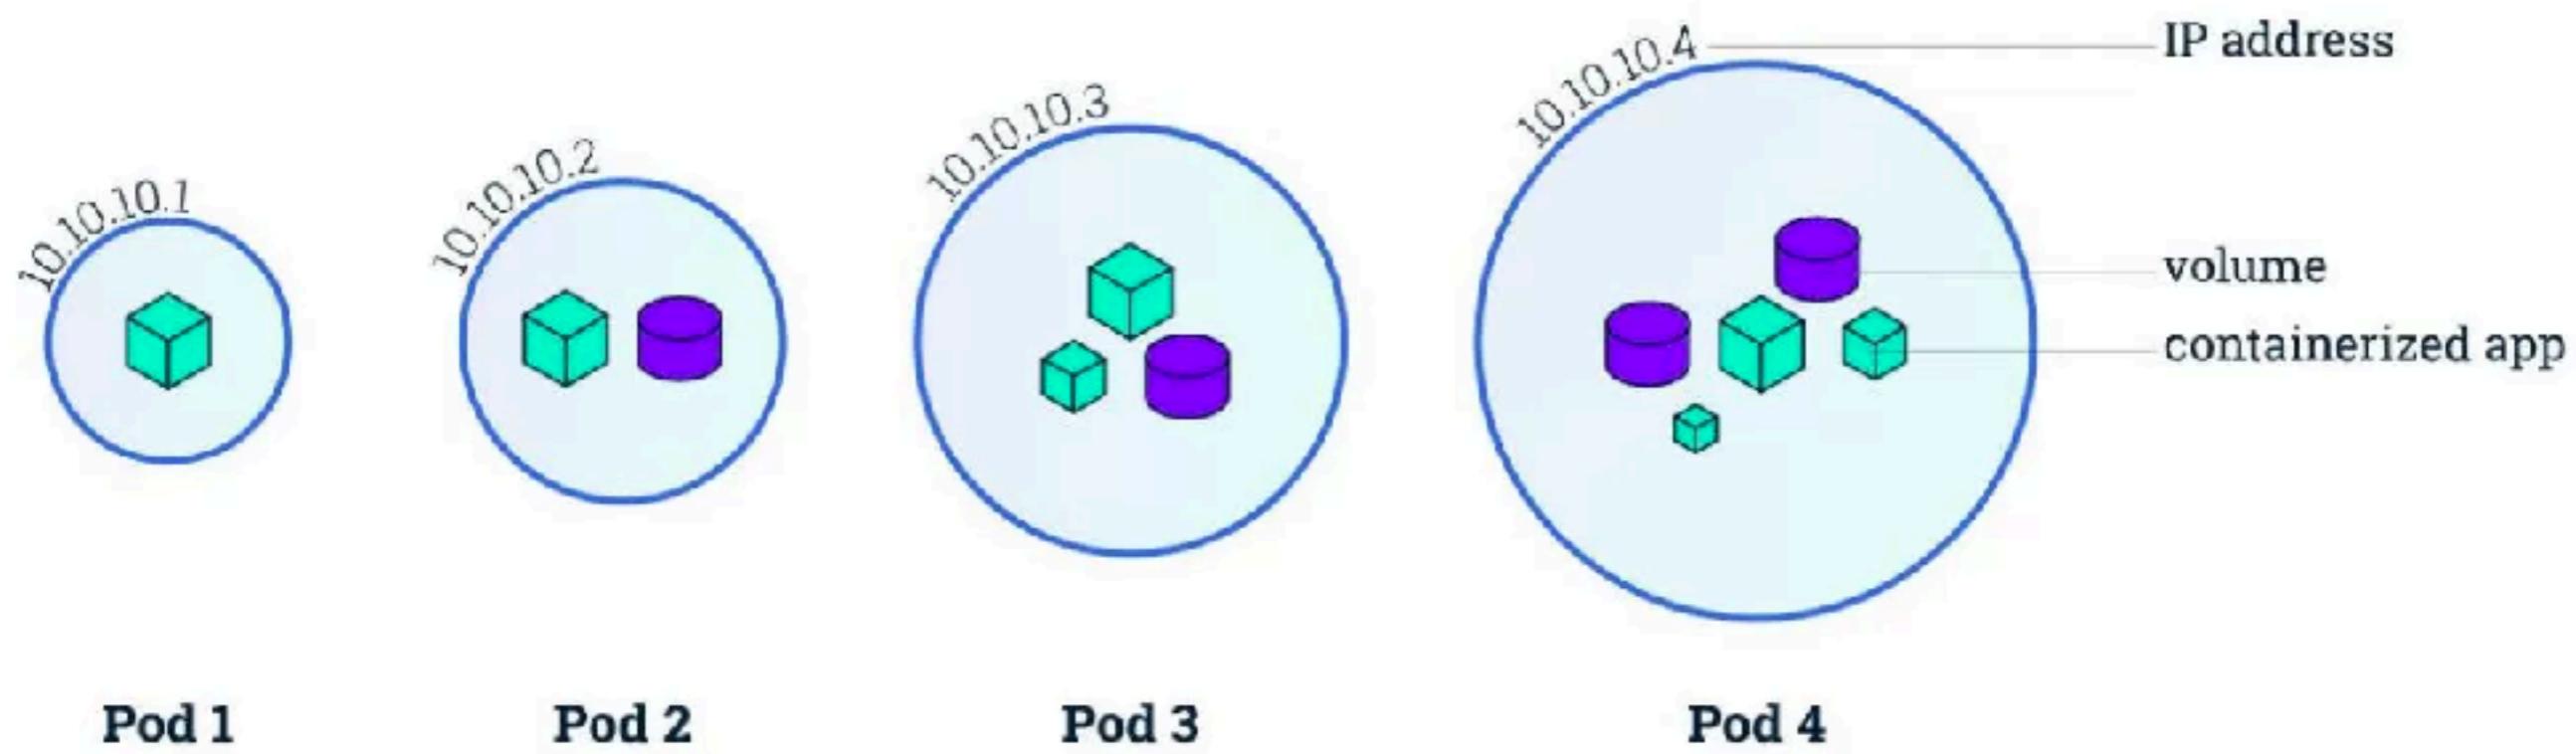
Share network interface

Share Hostname/IP/Port

Share Unix Time Sharing (UTS)

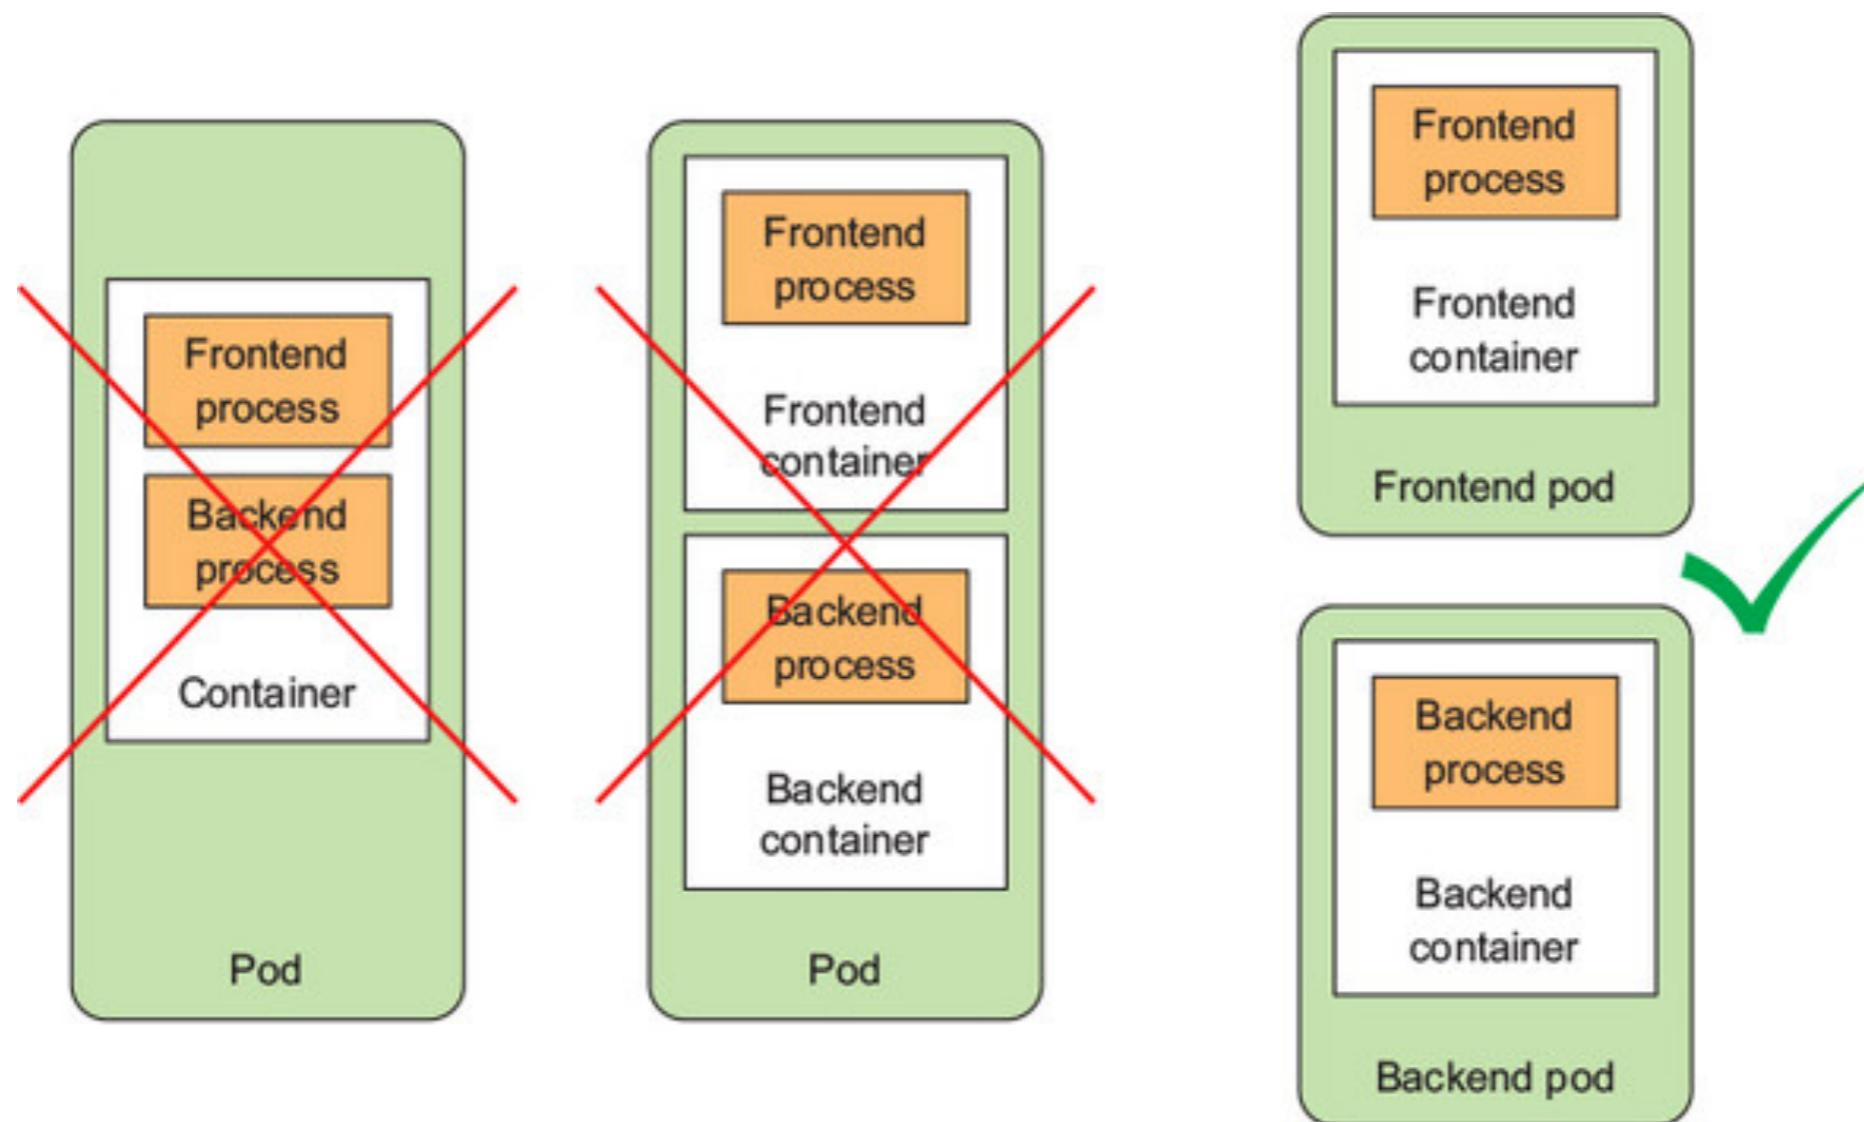


# Pods



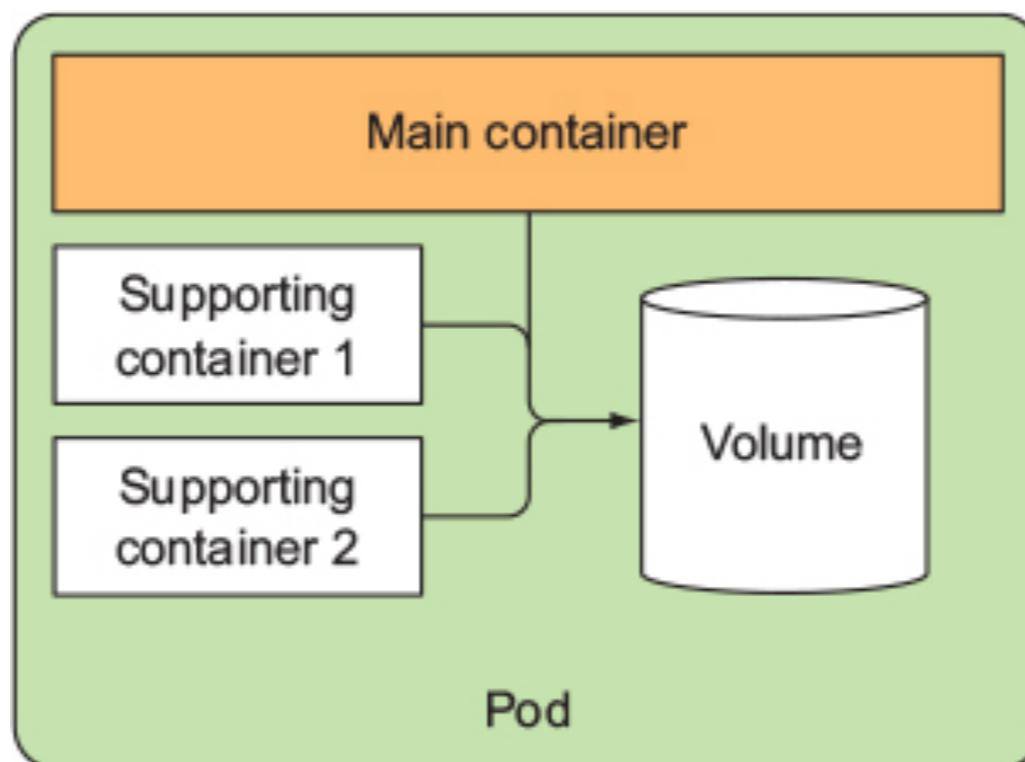
# Organize container across Pods

Split multi-tiers app into multiple pods

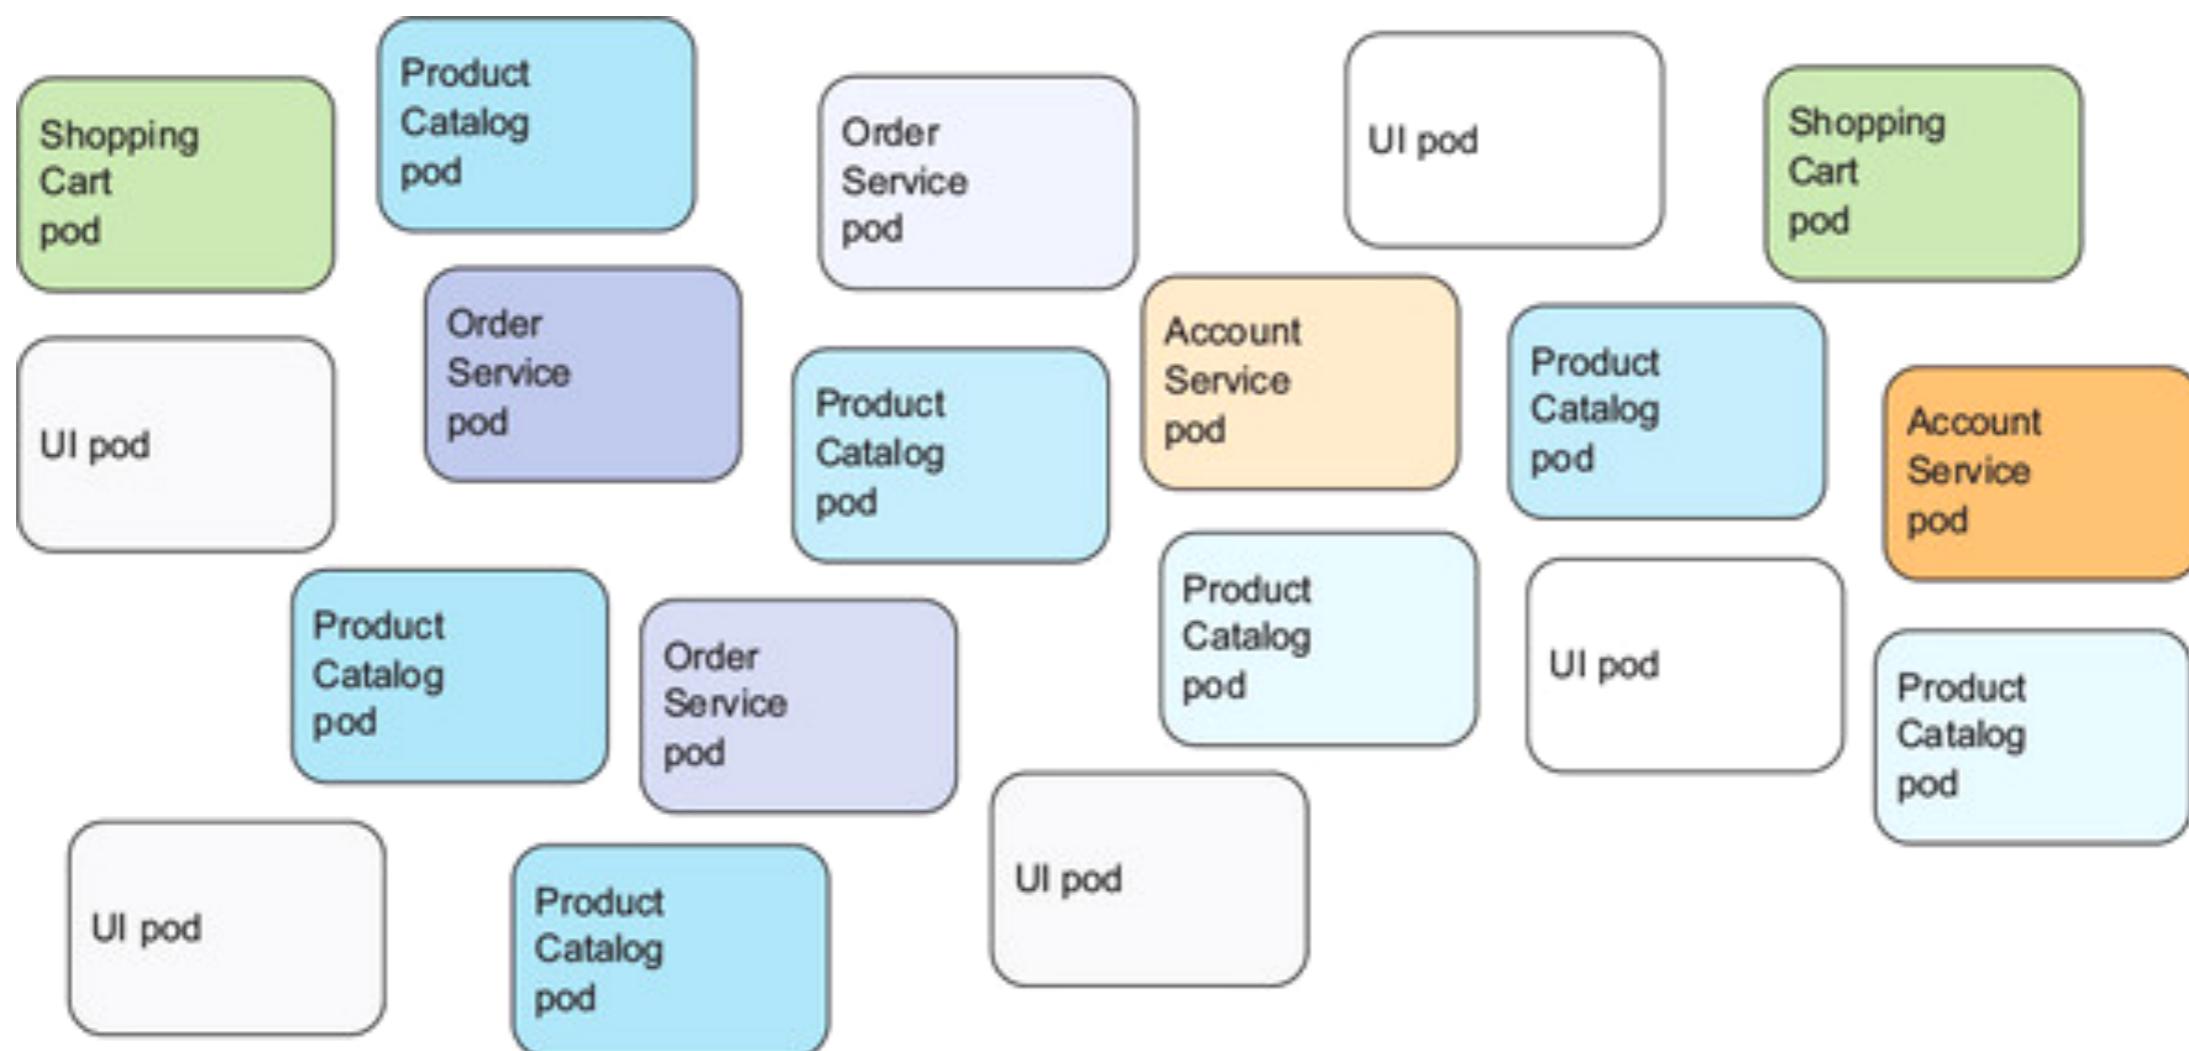


# Organize container across Pods

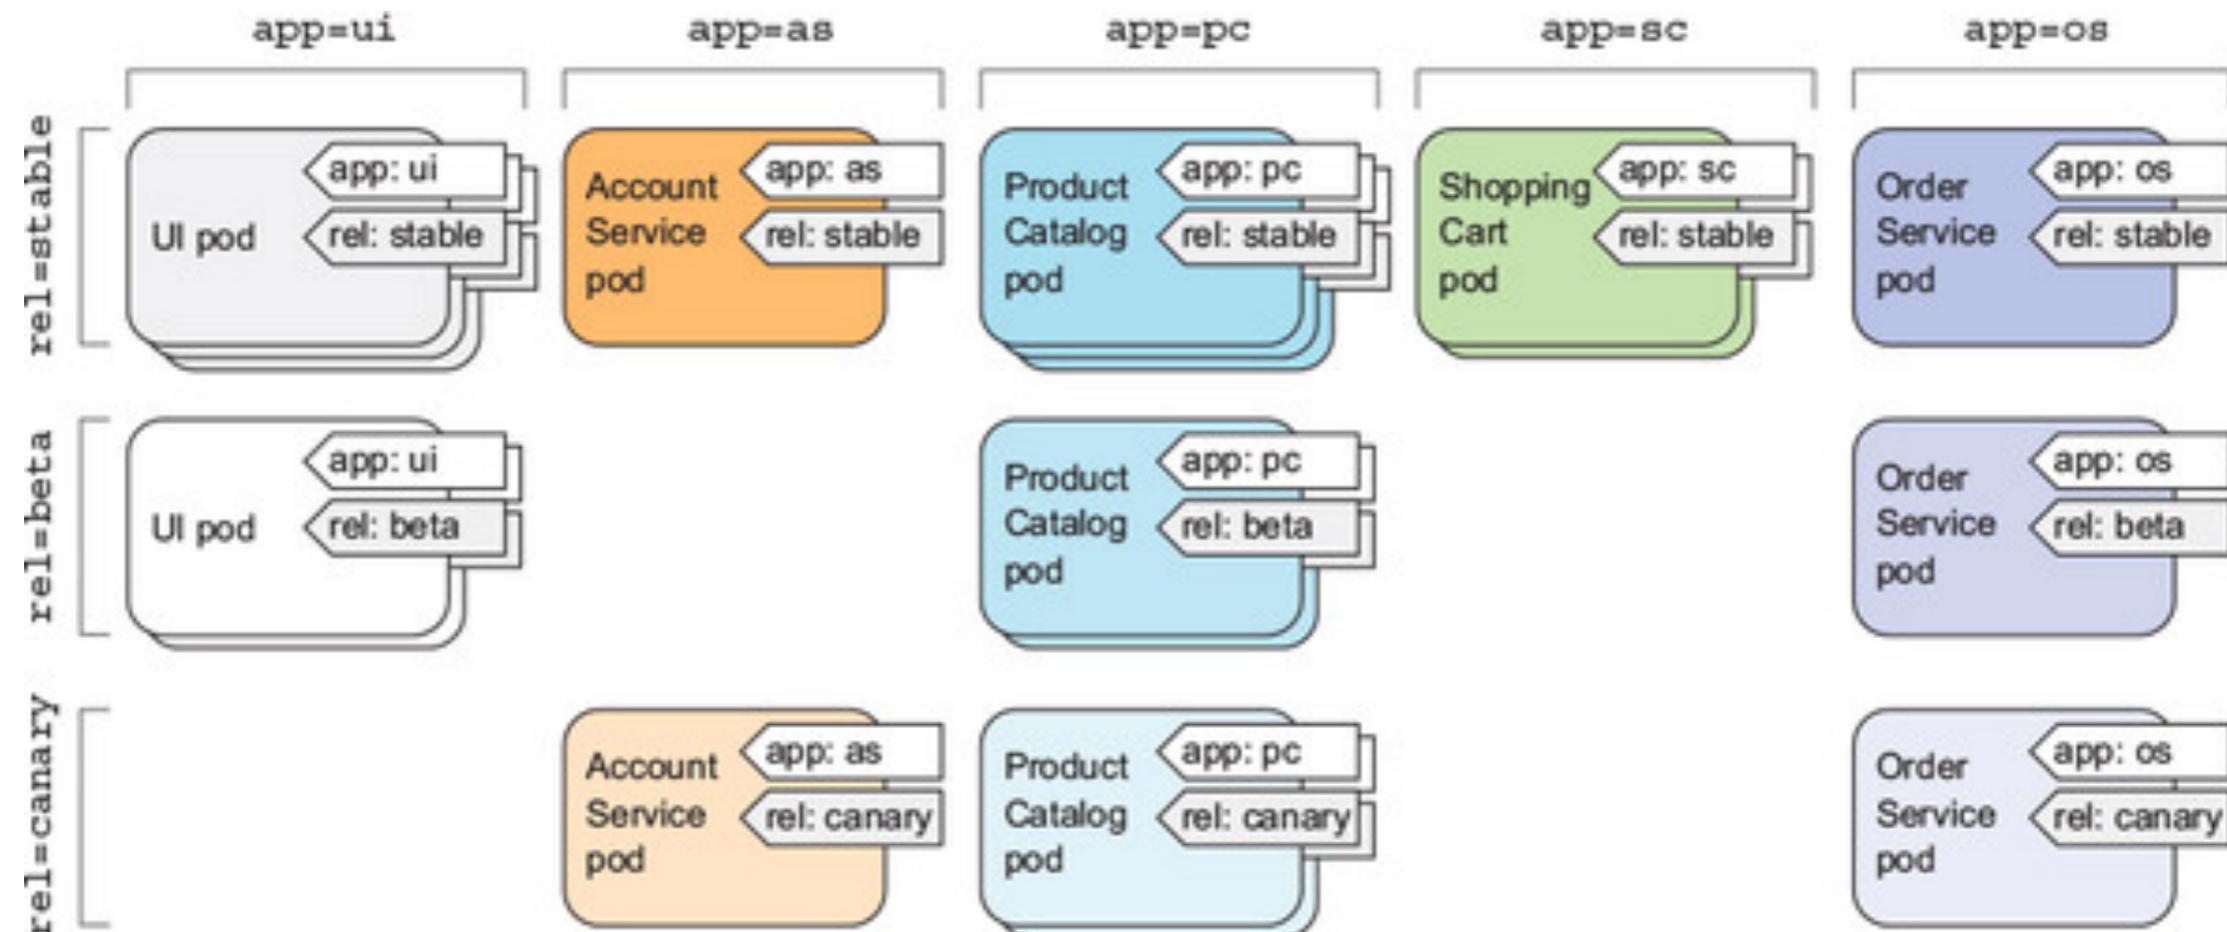
When to use multiple containers in a pods



# Organize pods with Labels

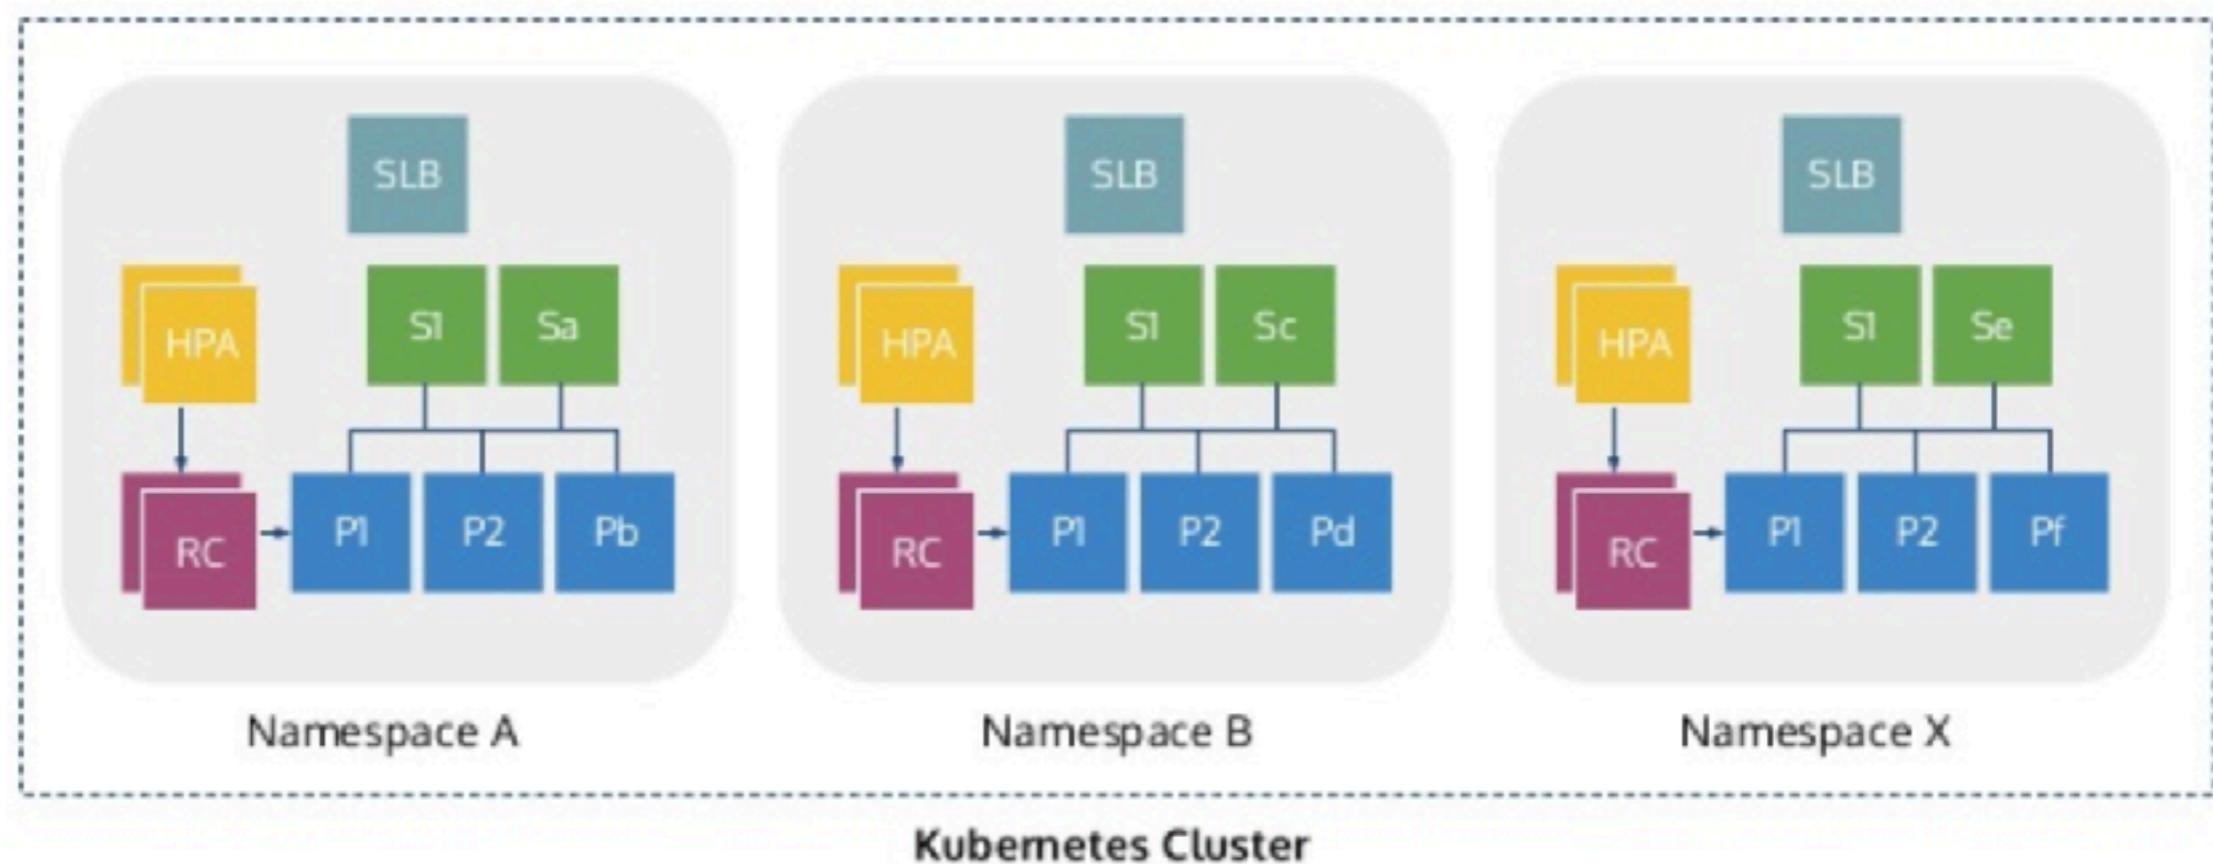


# Organize pods with Labels



# Pods namespaces

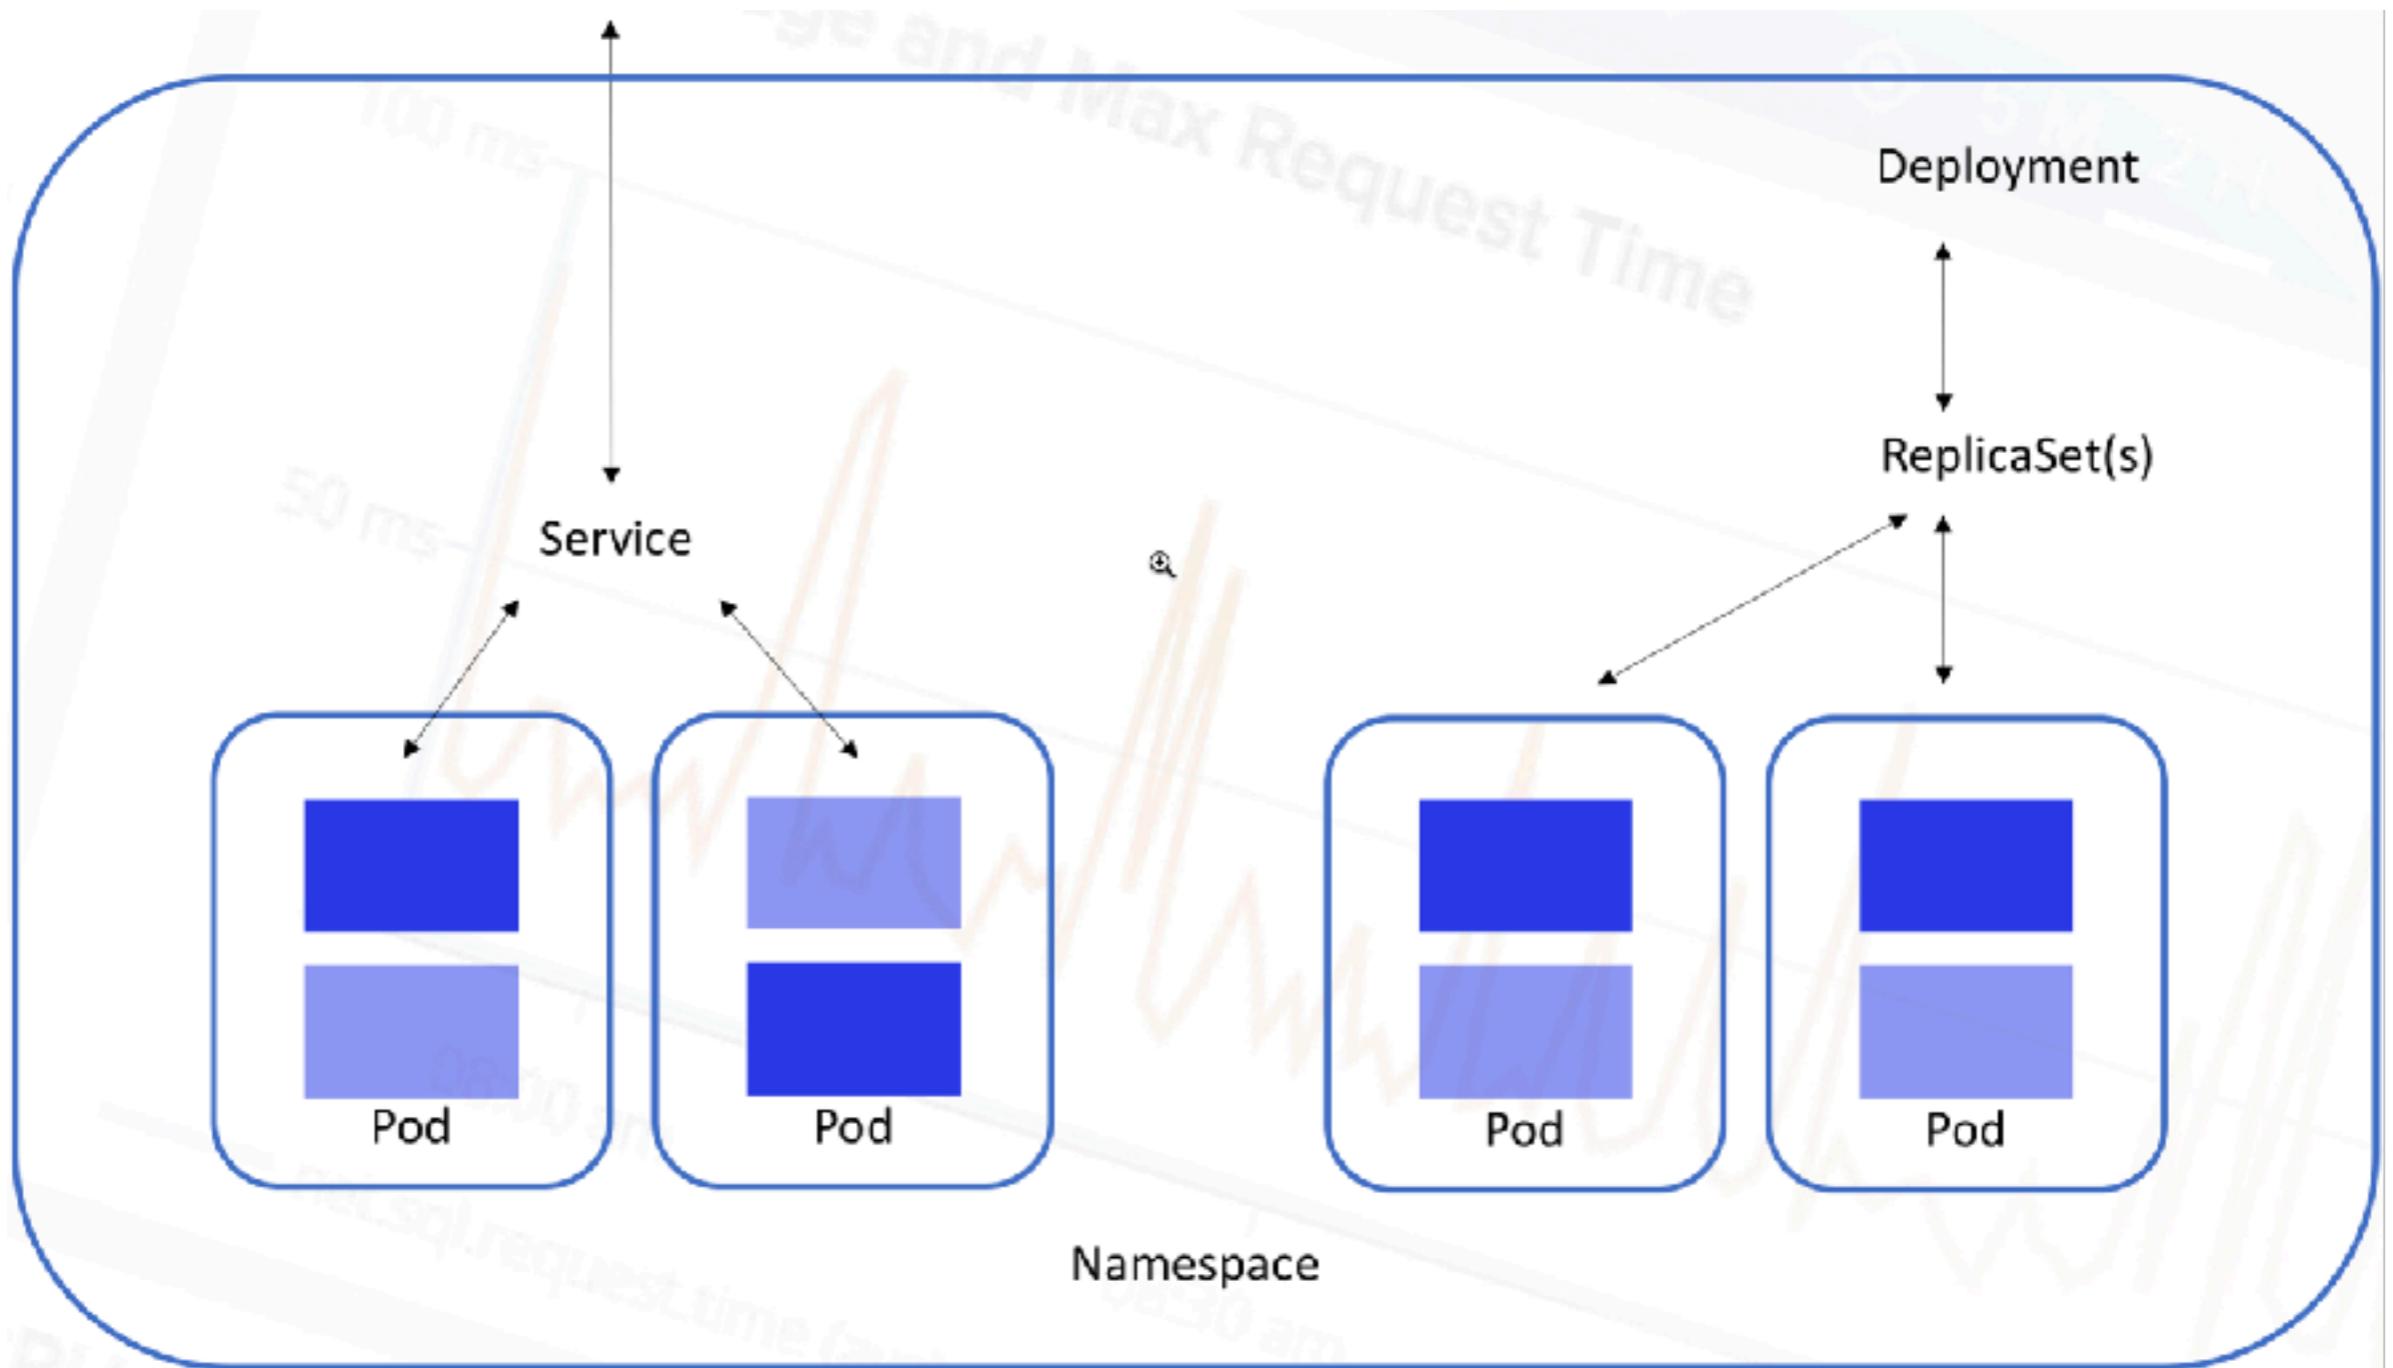
Allow different teams to use the same cluster



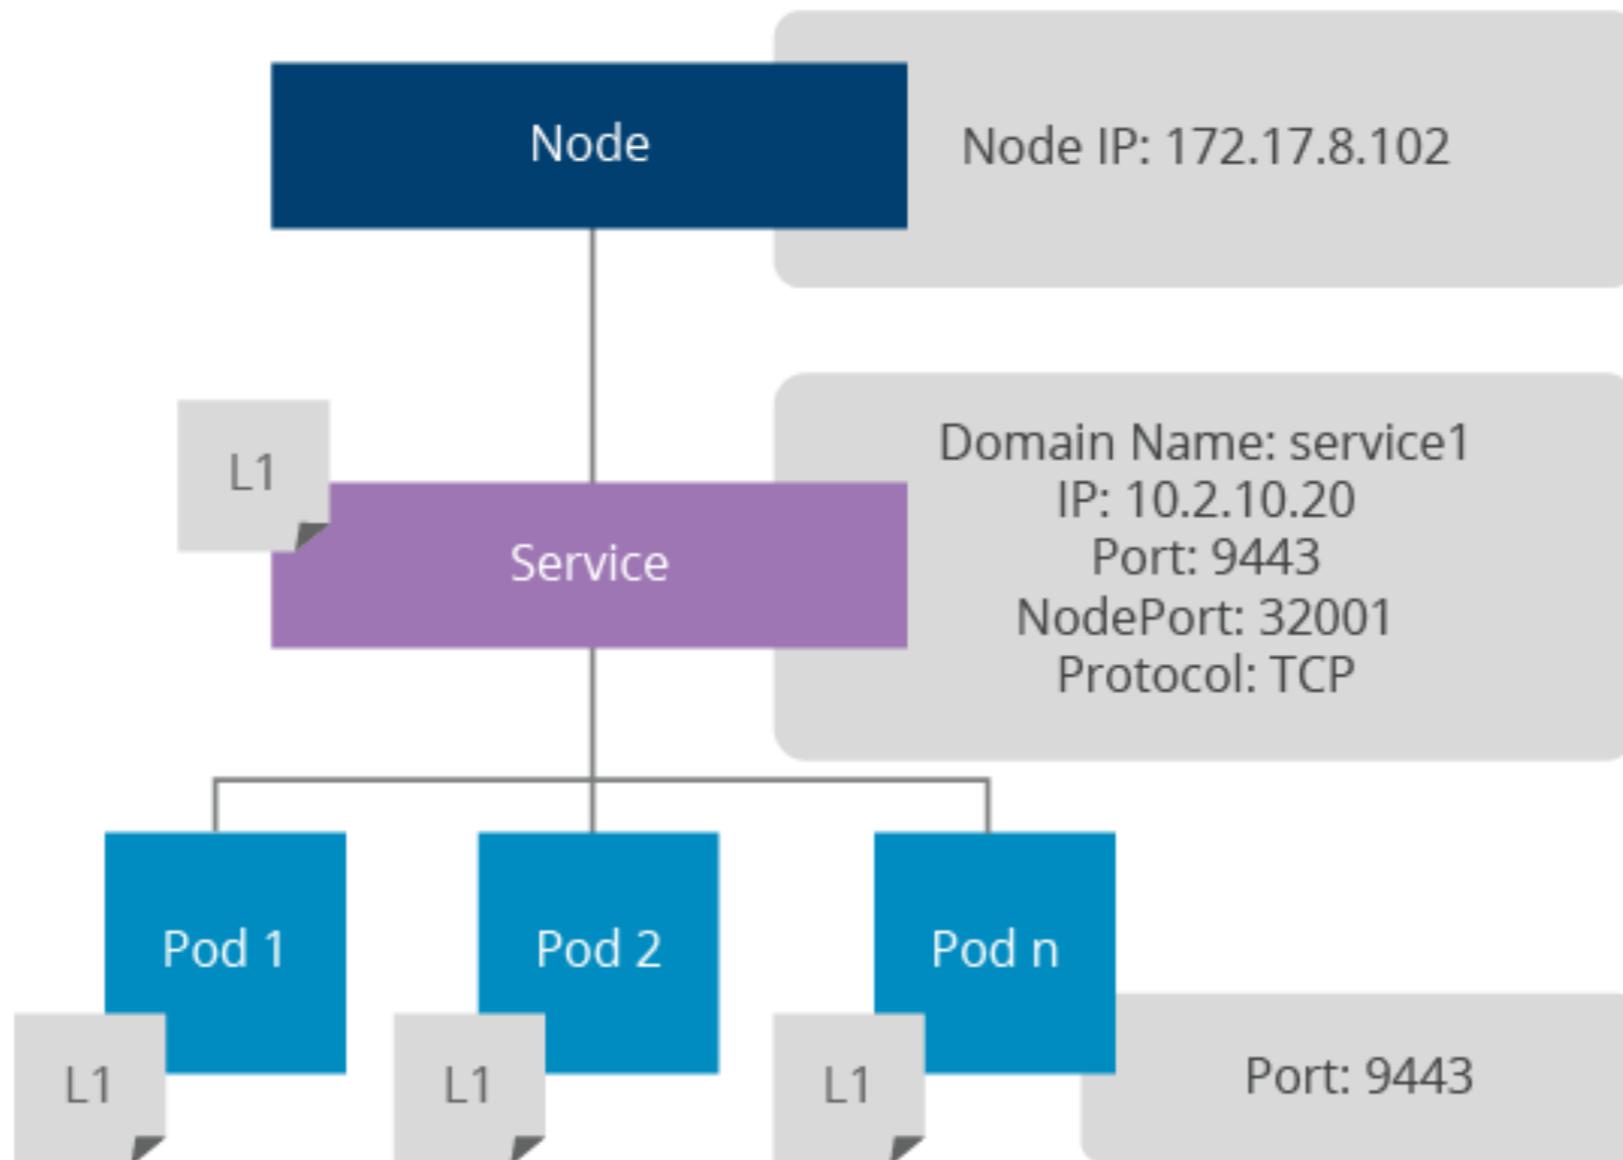
# **Services discovery and Load balancing**



# Services



# Services

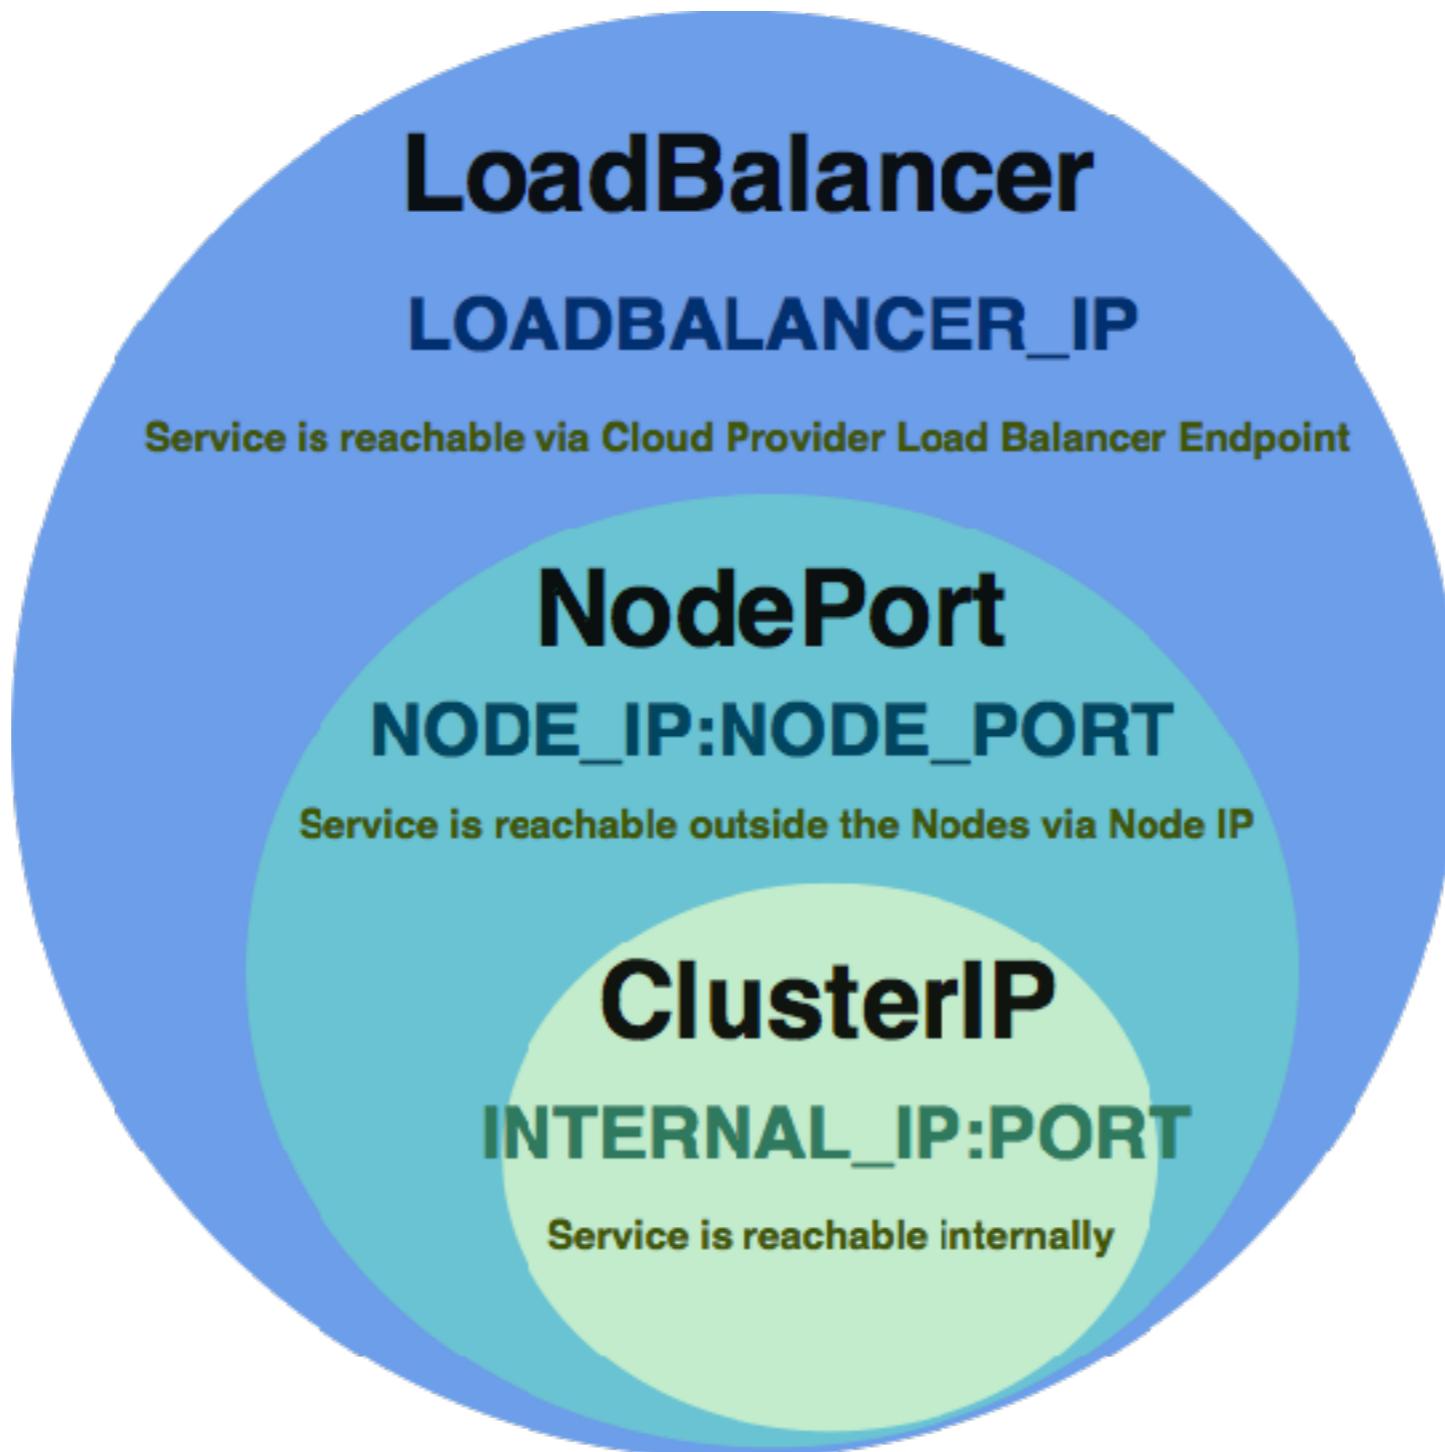


# Services

- Independent from Pods
- Abstraction layer of Pods
- Provide load balance
- Expose access Pods/Load balance
- Find Pods by label selector



# 3 types of services



# Kubernetes Object Management

1. Imperative command
2. Imperative object configuration
3. Declarative object configuration

<https://kubernetes.io/docs/concepts/overview/object-management-kubectl/overview/>



# 1. Imperative command

```
$kubectl run nginx --image nginx
```

```
$kubectl create deployment nginx --image nginx
```



## 2. Imperative object configuration

```
$kubectl create -f nginx.yaml
```

```
$kubectl delete -f nginx.yaml
```



# 3. Declarative object configuration

```
$kubectl apply -f configs/  
$kubectl apply -R -f configs/
```



# Kubernetes Object Management

Management technique	Operates on	Recommended environment	Supported writers	Learning curve
Imperative commands	Live objects	Development projects	1+	Lowest
Imperative object configuration	Individual files	Production projects	1	Moderate
Declarative object configuration	Directories of files	Production projects	1+	Highest

**Warning:** A Kubernetes object should be managed **using only one technique**. Mixing and matching techniques for the same object results in undefined behavior.



# **Workshop :: Pods & Services**

## **1 container per Pods**

**file /02-pod-service/single-pod/  
workshop\_instruction.txt**



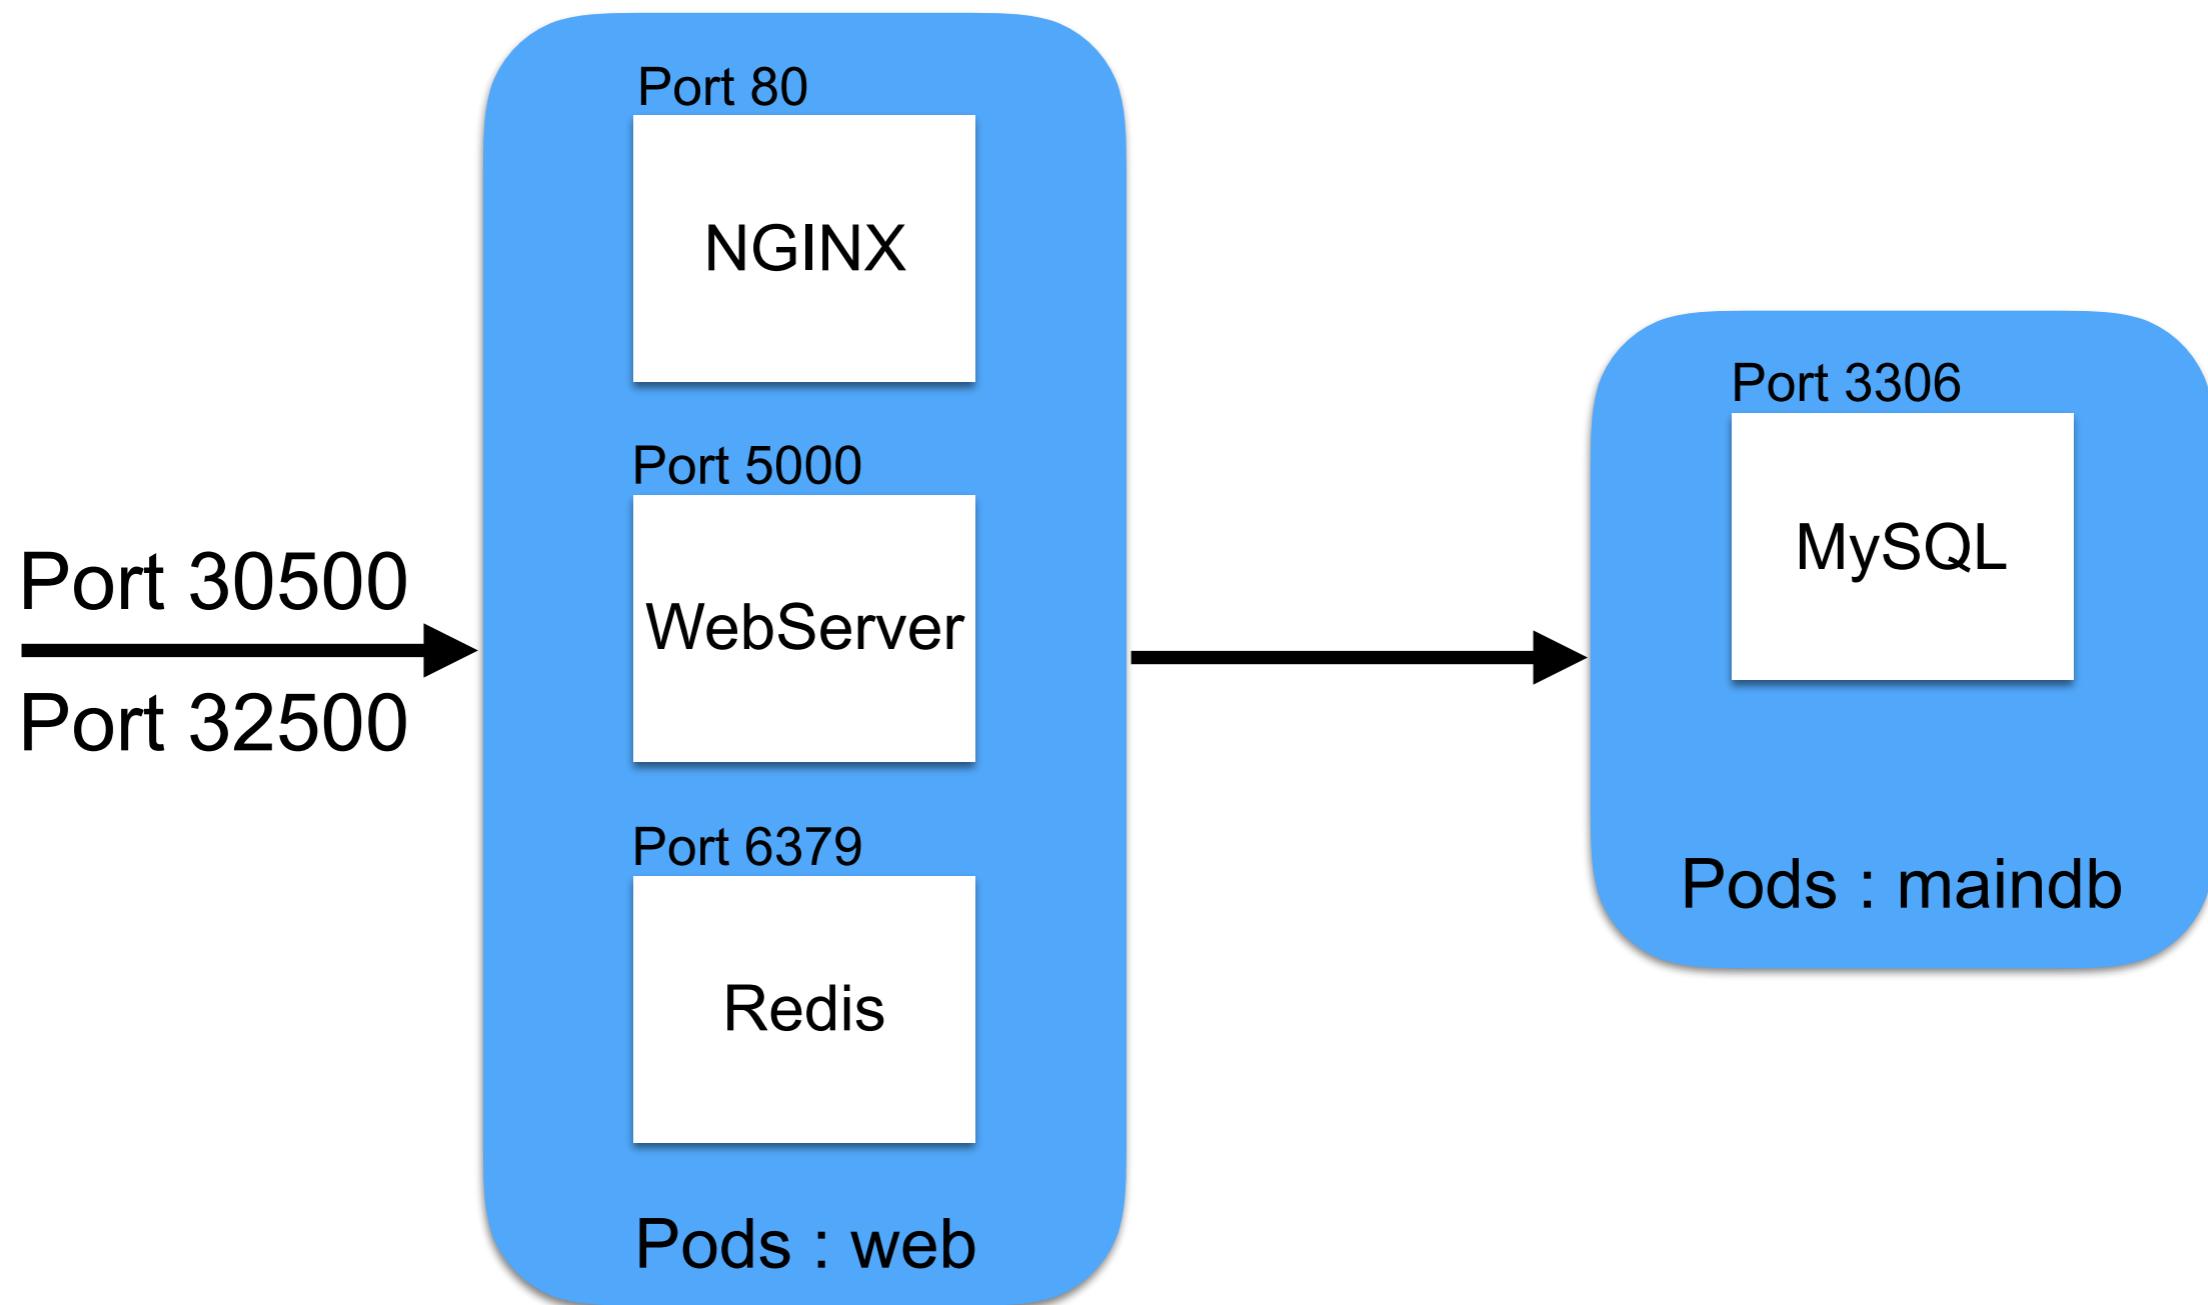
# **Workshop :: Pods & Services**

## **N container per Pods**

**file /01-starter/instruction.txt**



# Workshop



# Replication Controller (RC)



# Replication Controller

Create and maintain Pods

Keep copy of Pods by design

Maintain on cluster level

Auto-healing if Pods crash with any reason

Ensure Pods is up and run with desired number



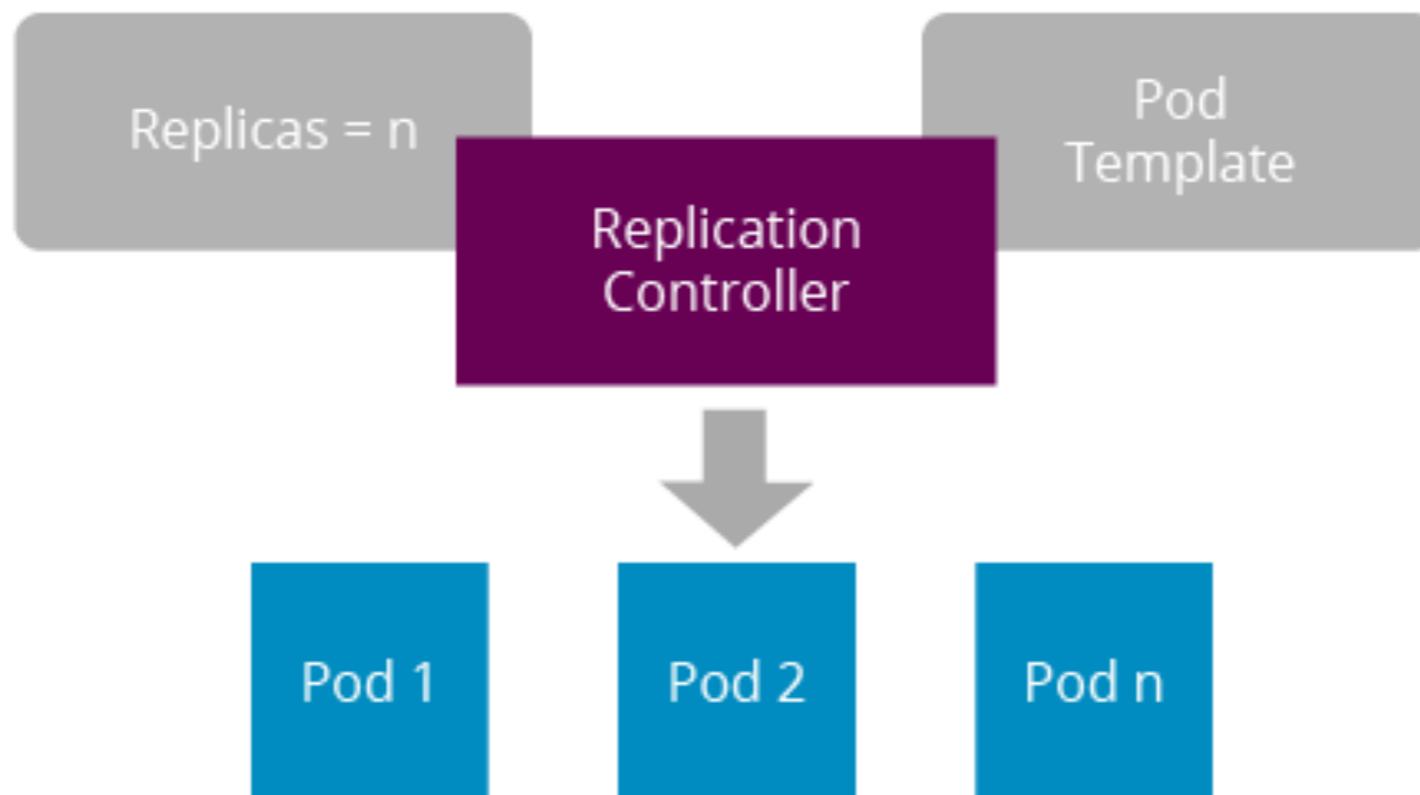
# Replication Controller

Running based on label type  
**Equality-based requirement**

```
selector:  
  name: web  
  version: "1.0"  
  module: WebServer  
  environment: development
```



# Replication Controller



# Scale up replicas of RC

```
$kubectl scale <option>  
--replicas=<number>  
<type or name>
```

```
bash-3.2$ kubectl get rc  
NAME      DESIRED   CURRENT   READY      AGE  
hello     5          5          5          15m  
bash-3.2$ kubectl get pods  
NAME        READY   STATUS    RESTARTS   AGE  
hello-jptv7  1/1     Running   0          43s  
hello-kdbsc  1/1     Running   0          43s  
hello-rn545  1/1     Running   0          15m  
hello-w6t5q  1/1     Running   0          43s  
hello-x6x8k  1/1     Running   0          43s
```



# Detail of RC

## \$kubectl describe rc <name>

### Containers:

```
hello:  
  Image:      somkiat/hello:latest  
  Port:       8080/TCP  
  Host Port:  0/TCP  
  Environment: <none>  
  Mounts:     <none>  
  Volumes:    <none>
```

### Events:

Type	Reason	Age	From	Message
Normal	SuccessfulCreate	16m	replication-controller	Created pod: hello-rn545
Normal	SuccessfulCreate	16m	replication-controller	Created pod: hello-pfwnj
Normal	SuccessfulCreate	16m	replication-controller	Created pod: hello-jfqfl
Normal	SuccessfulCreate	13m	replication-controller	Created pod: hello-6pfqm
Normal	SuccessfulCreate	11m	replication-controller	Created pod: hello-l4lqc
Normal	SuccessfulCreate	11m	replication-controller	Created pod: hello-pfgr2
Normal	SuccessfulDelete	9m	replication-controller	Deleted pod: hello-l4lqc
Normal	SuccessfulDelete	9m	replication-controller	Deleted pod: hello-pfwnj
Normal	SuccessfulDelete	9m	replication-controller	Deleted pod: hello-6pfqm
Normal	SuccessfulDelete	9m	replication-controller	Deleted pod: hello-pfgr2
Normal	SuccessfulCreate	2m	replication-controller	Created pod: hello-kdbsc
Normal	SuccessfulCreate	2m	replication-controller	Created pod: hello-jptv7
Normal	SuccessfulCreate	2m	replication-controller	Created pod: hello-x6x8k

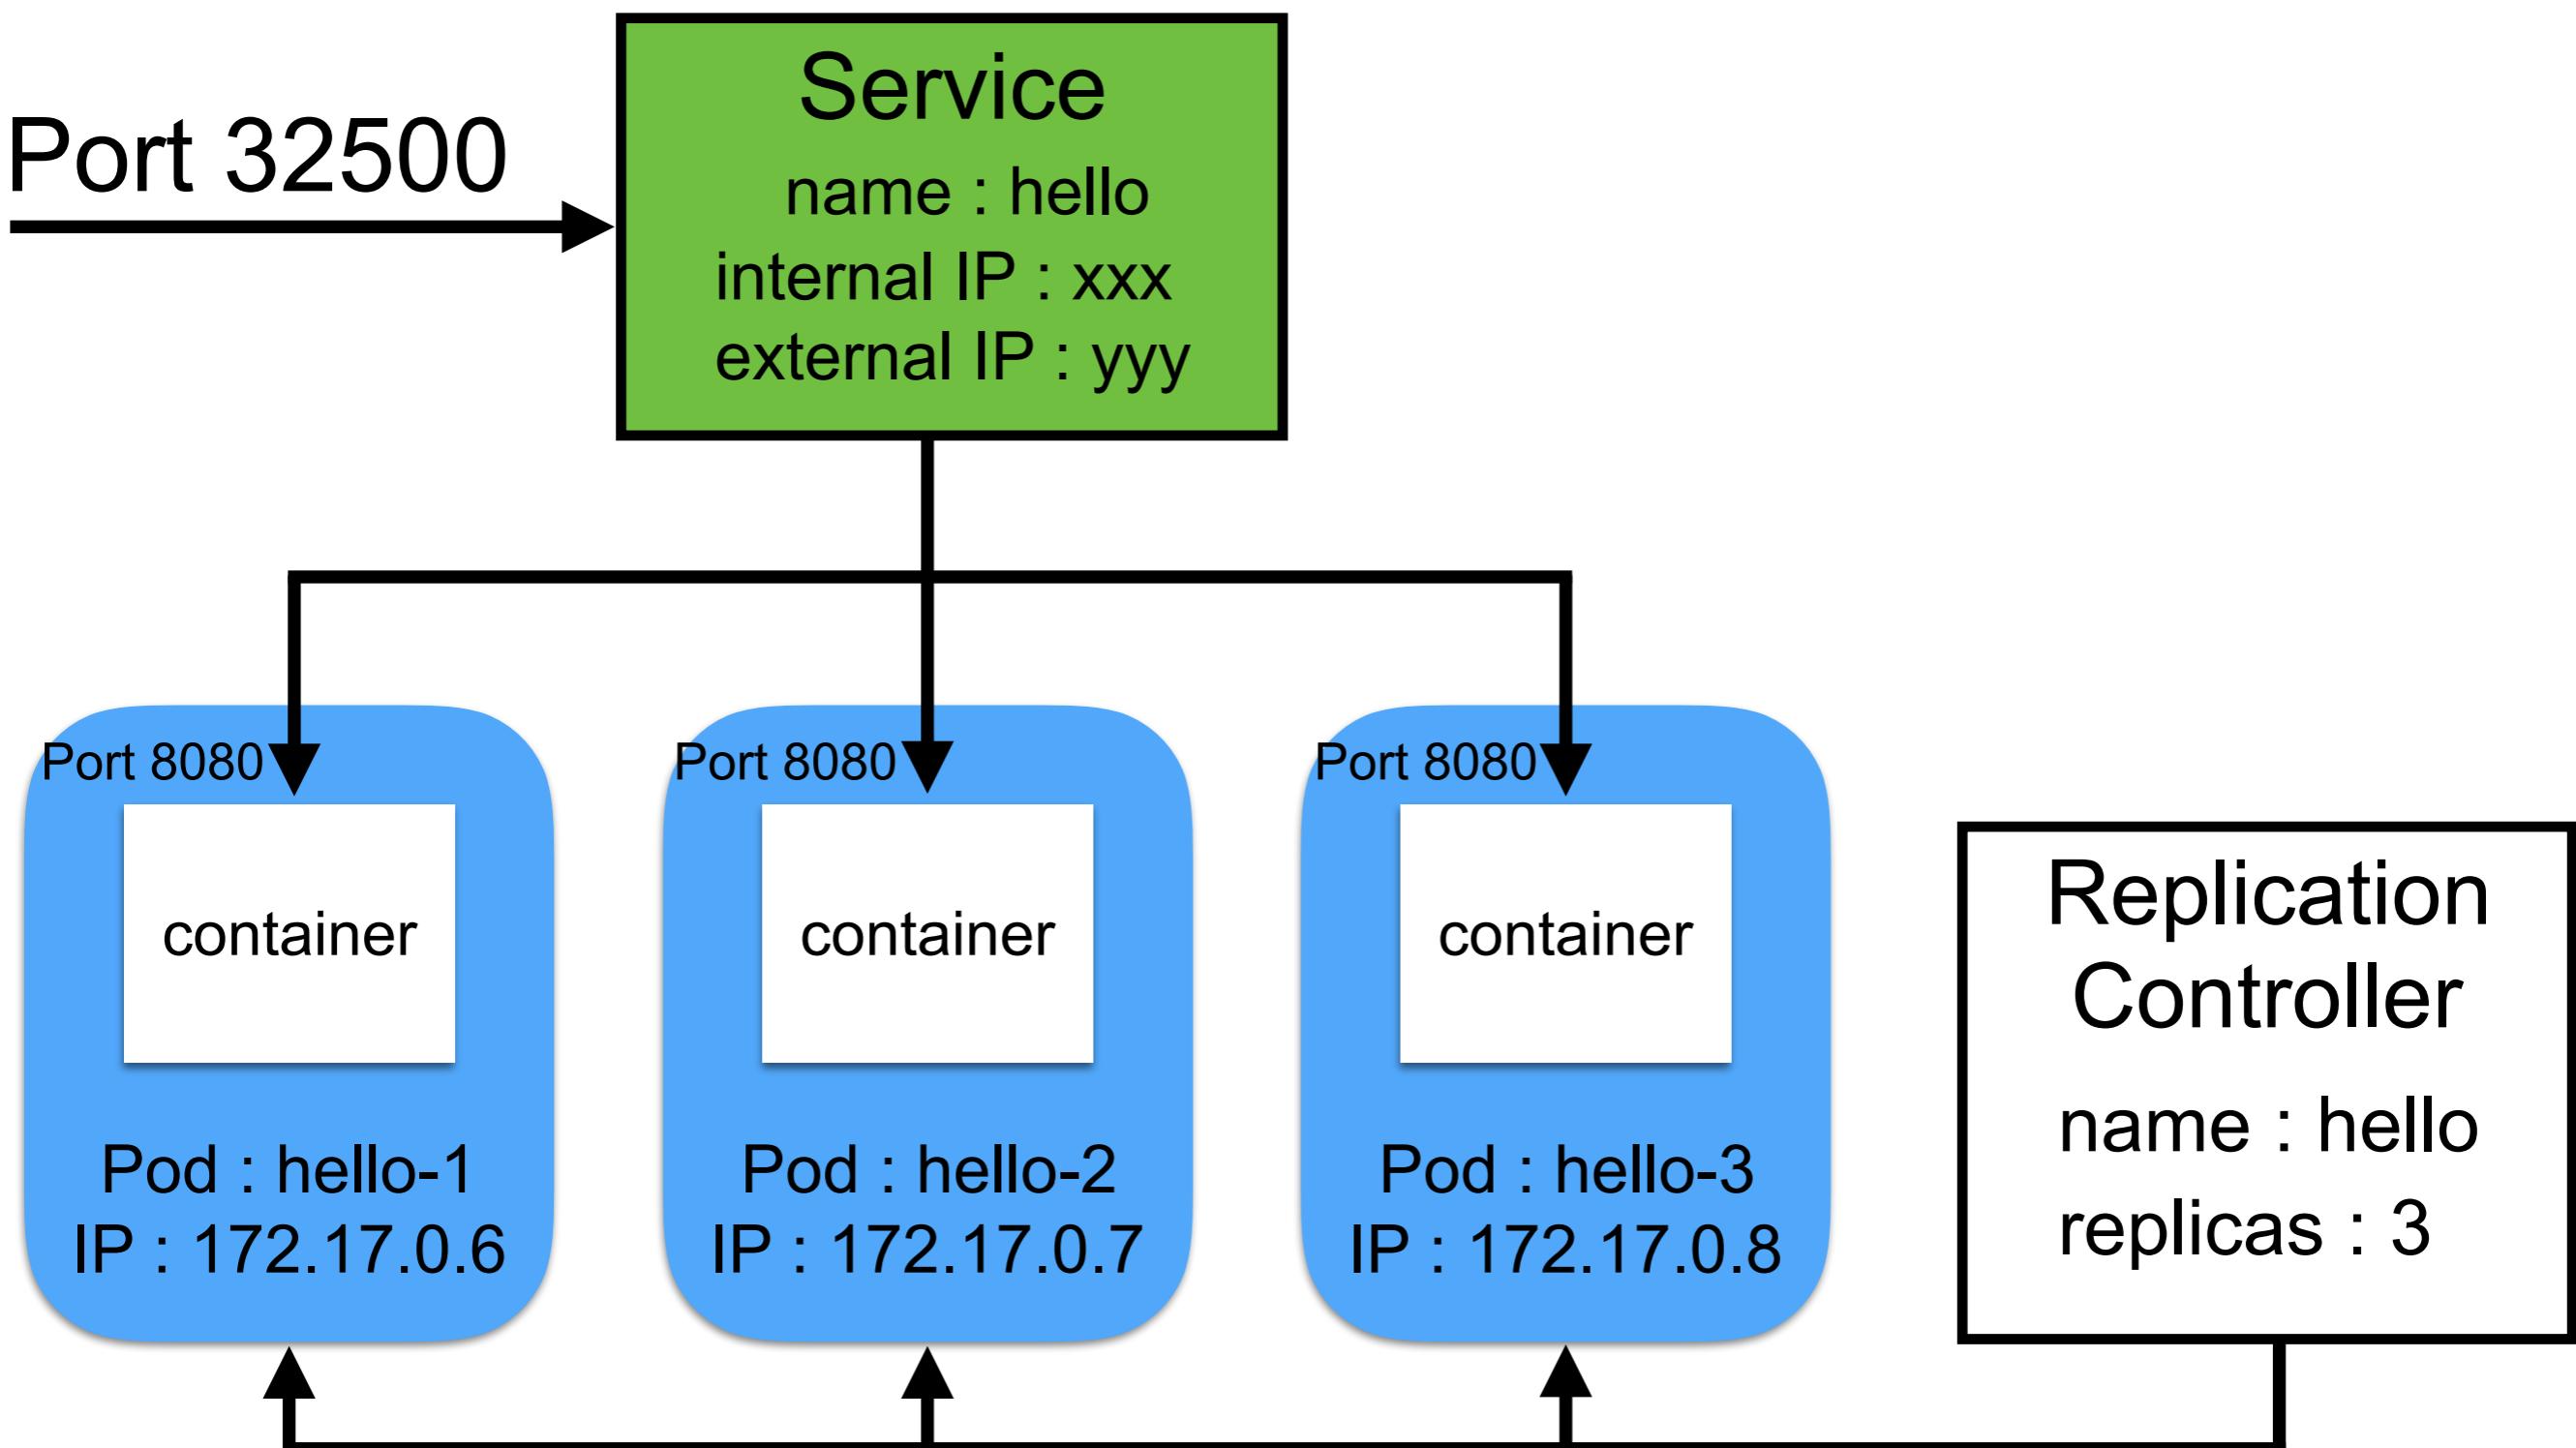


# Workshop Replication Controller

file /03-replication-controller



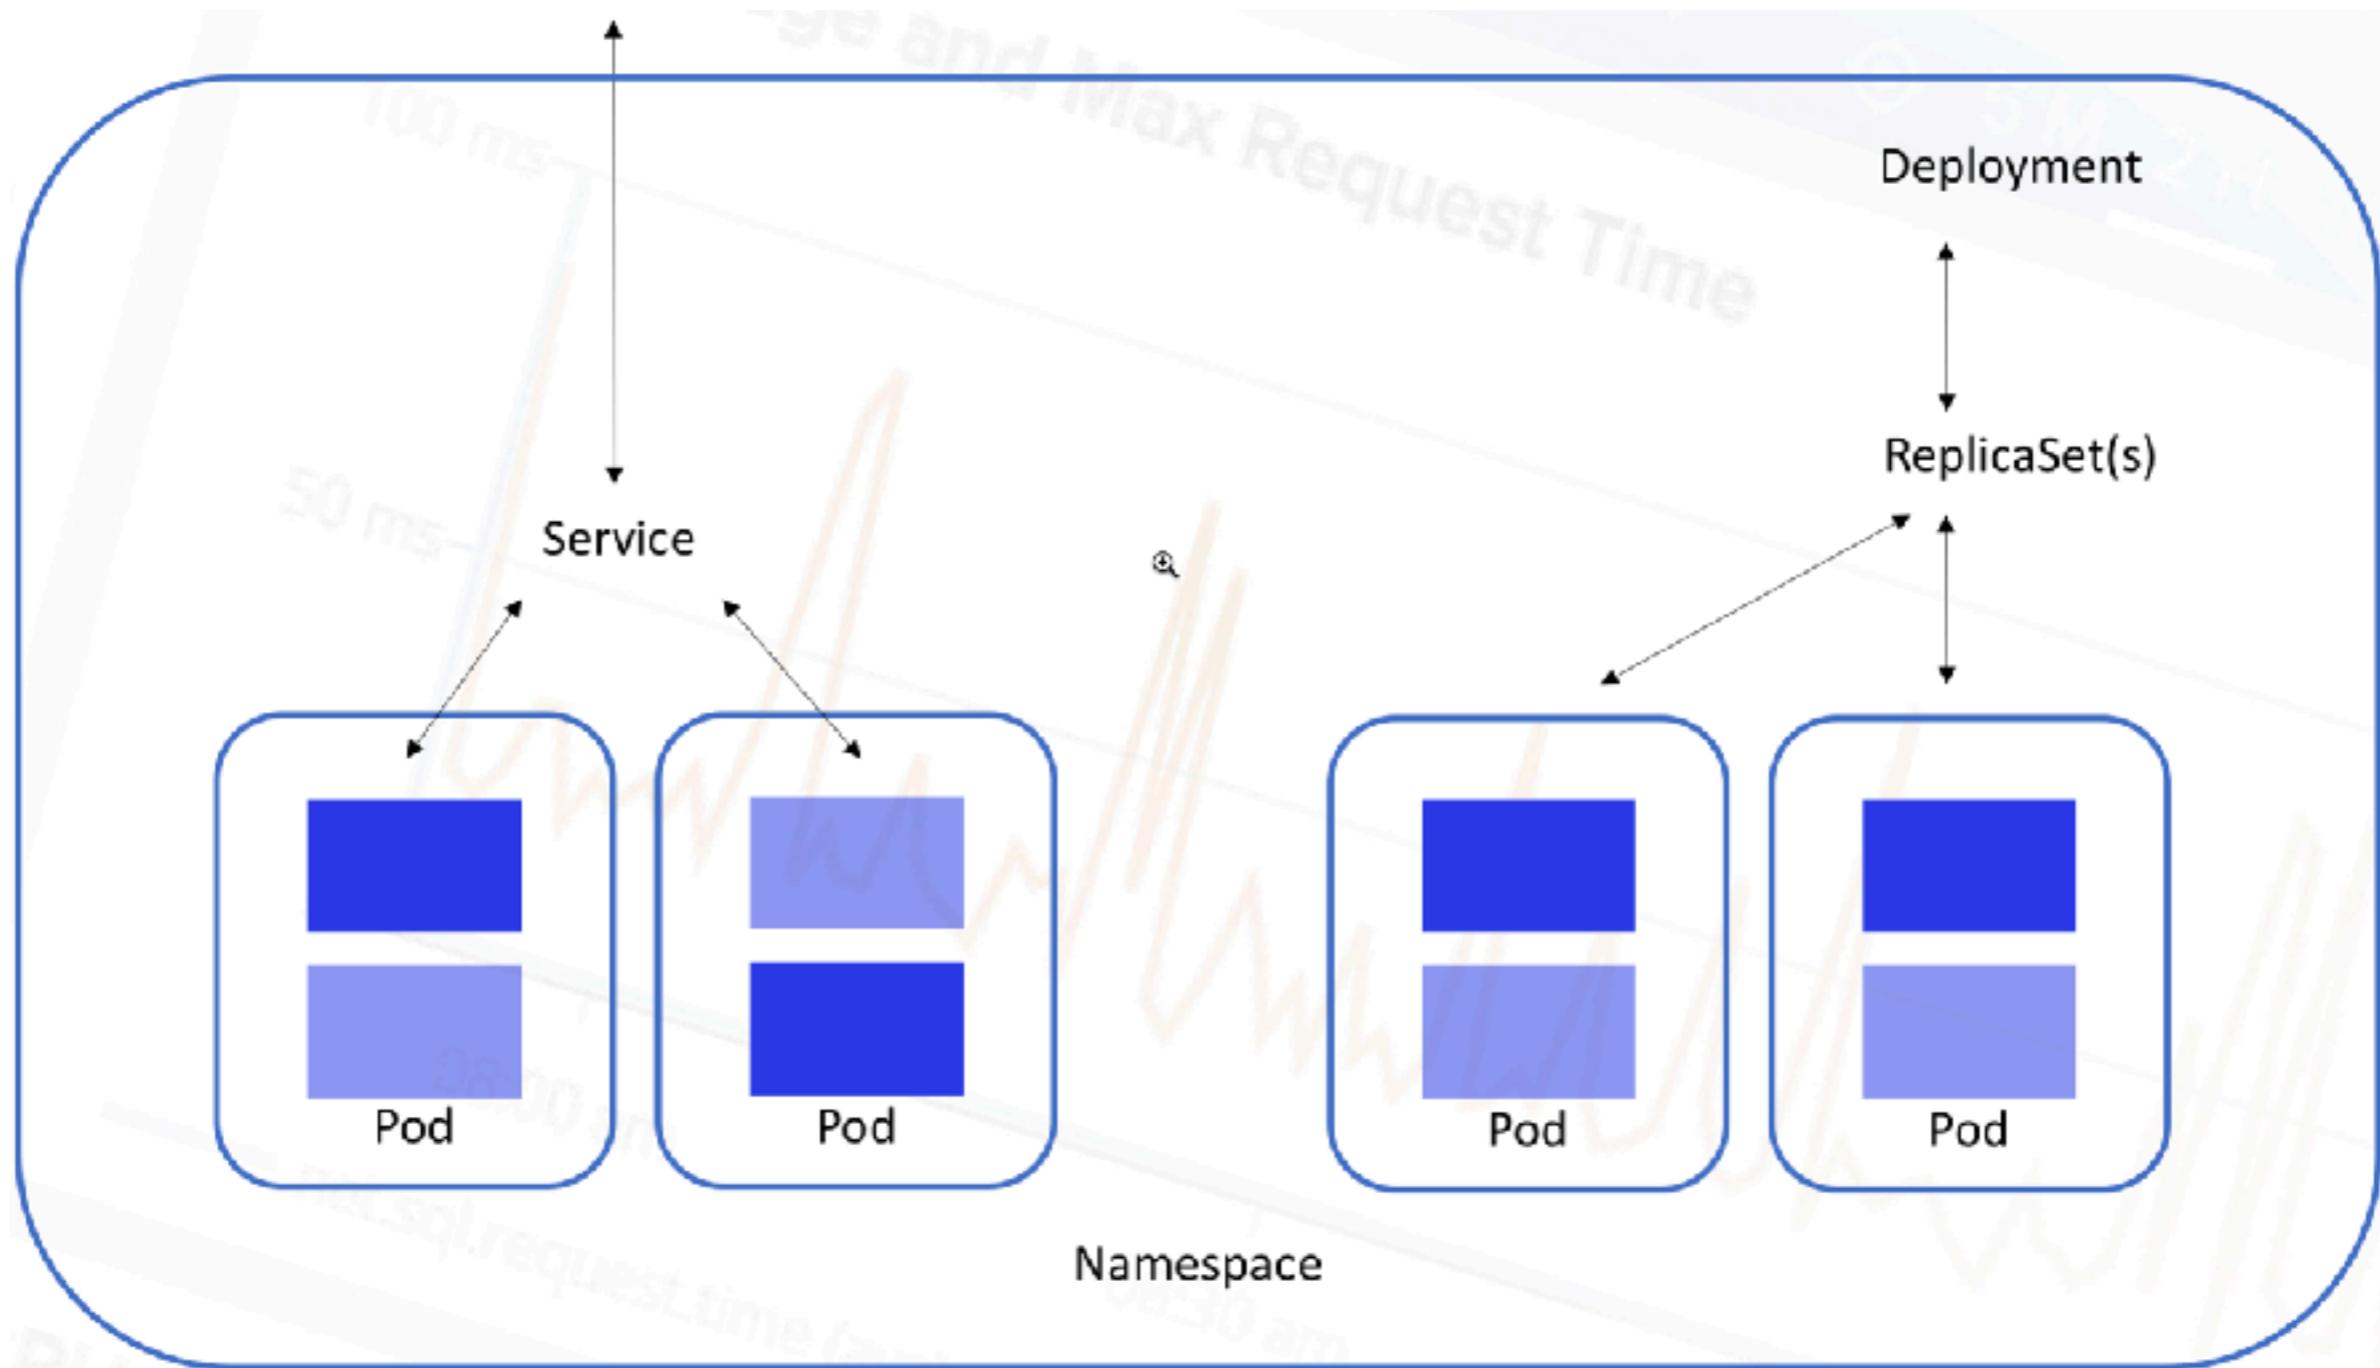
# Scaling the application



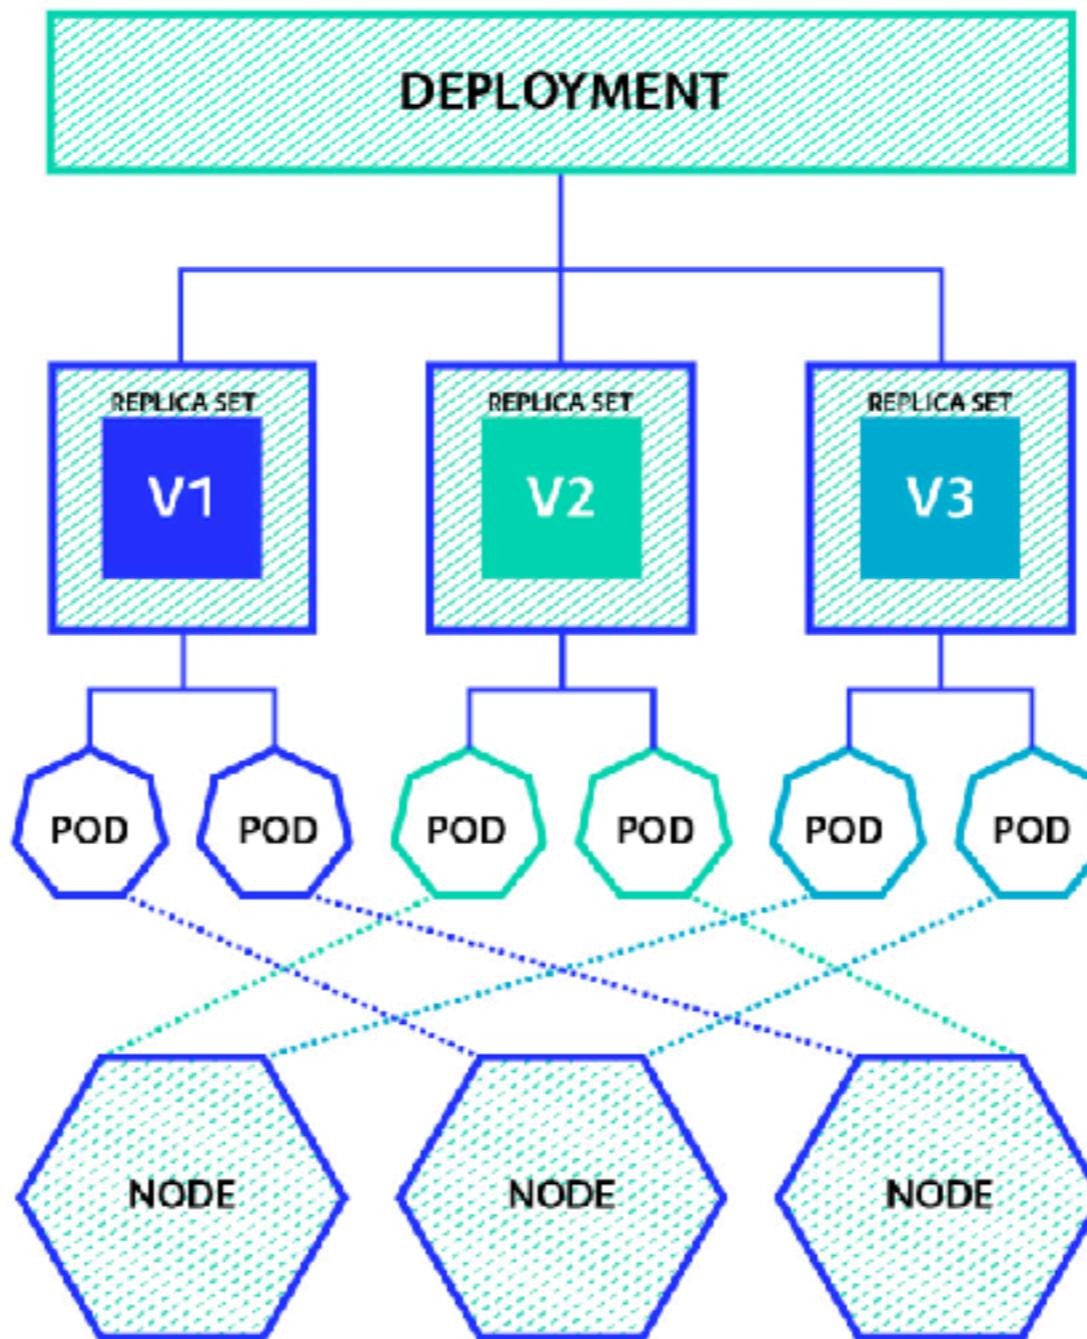
# **Deployment and ReplicaSet (RS)**



# Deployment and ReplicaSet



# Deployment and ReplicaSet



<https://thenewstack.io/kubernetes-deployments-work/>



# Deployment and RS

Next-generation of RC

Provide function to maintain versioning of Pods

Update new version (Rollout)

Revert to old version (Rollback)

Scale a deployment

Pause process

Resume process



# New version from Deployment

Create new RS

Start to scale as desired

Scale down existing RS to 0

Delete existing RS



# Deployment/RS vs RC

RS more dynamic than RC

RS support label with methods:

- Equality-based requirement

- Set-based requirement



# Rollout strategy

**Set online**

**Edit online**

**Modify YAML file and apply**



# Set online

```
$kubectl set image deployment/hello  
hello=somkiat/hello:v2
```

```
Waiting for rollout to finish: 1 out of 3 new replicas have been updated...  
Waiting for rollout to finish: 1 out of 3 new replicas have been updated...  
Waiting for rollout to finish: 1 out of 3 new replicas have been updated...  
Waiting for rollout to finish: 1 out of 3 new replicas have been updated...  
Waiting for rollout to finish: 1 out of 3 new replicas have been updated...  
Waiting for rollout to finish: 1 out of 3 new replicas have been updated...  
Waiting for rollout to finish: 2 out of 3 new replicas have been updated...  
Waiting for rollout to finish: 2 out of 3 new replicas have been updated...  
Waiting for rollout to finish: 2 old replicas are pending termination...  
Waiting for rollout to finish: 1 old replicas are pending termination...  
Waiting for rollout to finish: 1 old replicas are pending termination...  
deployment "hello" successfully rolled out
```



# Edit online

## \$kubectl edit deployment hello

```
# Please edit the object below. Lines beginning with a '#' will be ignored,
# and an empty file will abort the edit. If an error occurs while saving this file,
# it will be reopened with the relevant failures.
#
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  annotations:
    deployment.kubernetes.io/revision: "6"
    kubectl.kubernetes.io/last-applied-configuration: |
      {"apiVersion":"apps/v1","kind":"Deployment","metadata":{"annotations":{}},
      "namespace":"default","spec":{"replicas":3,"revisionHistoryLimit":1,"selector":{"matchLabels":{"app":"web"}}, "template":{"metadata":{"labels":{"app":"web"}}, "spec":{"containers":[{"image":"nginx:1.13.6-alpine", "ports":[{"containerPort":8080, "protocol":"TCP"}]}]}}}
  creationTimestamp: 2018-04-01T17:17:51Z
  generation: 10
  labels:
    app: web
  name: hello
  namespace: default
  resourceVersion: "33405"
  selfLink: /apis/extensions/v1beta1/namespaces/default/deployments/hello
  uid: 9b773859-35d0-11e8-9d36-0800275c6c60
```



# Modify YAML file and apply

```
$kubectl apply -f hello_deployment.yml
```



# Show history of rollout

```
$kubectl rollout history deployment/hello
```

```
deployments "hello"
REVISION  CHANGE-CAUSE
2          <none>
3          <none>
4          <none>
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: hello
  labels:
    app: web
spec:
  replicas: 3
  revisionHistoryLimit: 1
  selector:
    matchLabels:
      app: web
```

spec.revisionHistory = 2 (default)  
spec.revisionHistory = 0 (clean all)



# Try to rollback to revision

```
$kubectl rollout undo deployment/hello --to-revision=2
```



# Try to scale a deployment

\$kubectl scale deployment hello --replicas=5

```
bash-3.2$ kubectl scale deployment hello --replicas=3
deployment.extensions "hello" scaled
bash-3.2$ kubectl get deployment -o wide
NAME      DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
hello      3          3          3           3           21m
CONTAINERS   IMAGES
hello      somkiat/hello
SELECTOR    app=web
bash-3.2$ kubectl scale deployment hello --replicas=5
deployment.extensions "hello" scaled
bash-3.2$ kubectl get deployment -o wide
NAME      DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
hello      5          5          5           5           21m
CONTAINERS   IMAGES
hello      somkiat/hello
SELECTOR    app=web
bash-3.2$ kubectl scale deployment hello --replicas=1
deployment.extensions "hello" scaled
bash-3.2$ kubectl get deployment -o wide
NAME      DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
hello      1          1          1           1           21m
CONTAINERS   IMAGES
hello      somkiat/hello
SELECTOR    app=web
```



# Pause and resume rollout

```
$kubectl rollout pause deployment/hello  
$kubectl rollout resume deployment/hello
```



# **Workshop**

# **Deployment/ReplicaSet**

**file /04-deployment-and-replica-set**



# Volume

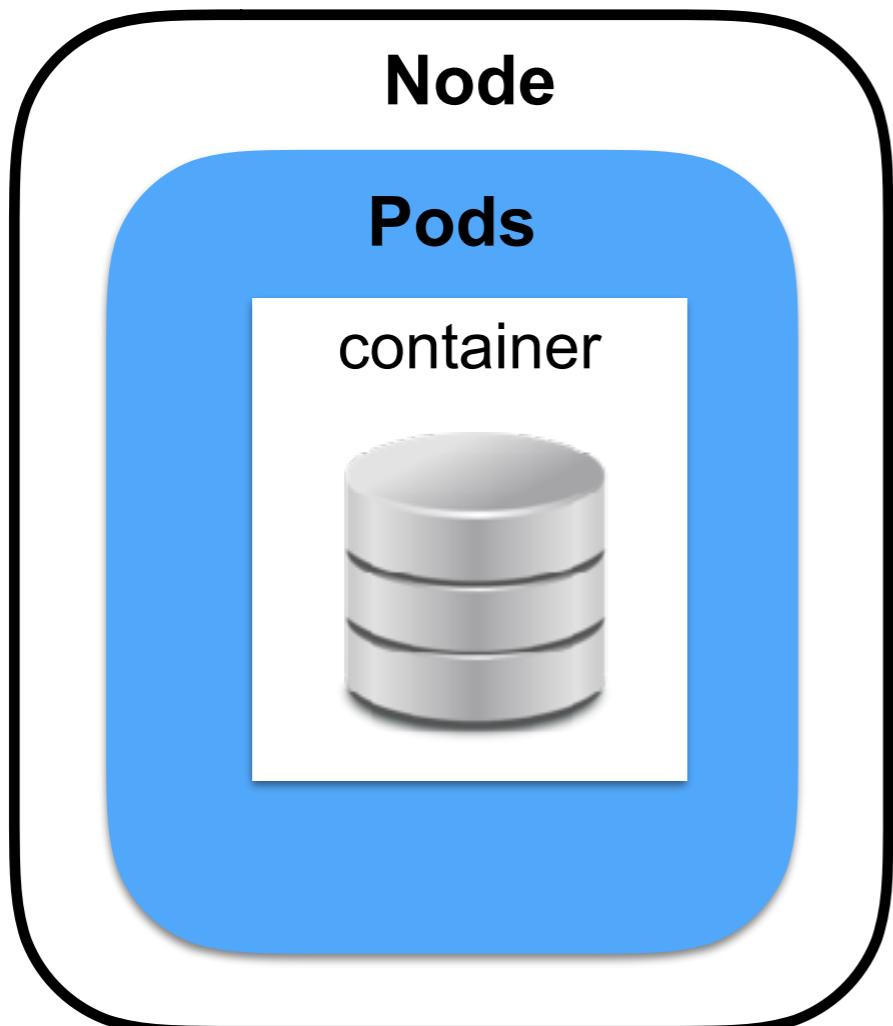
<https://kubernetes.io/docs/concepts/storage/volumes/#resources>



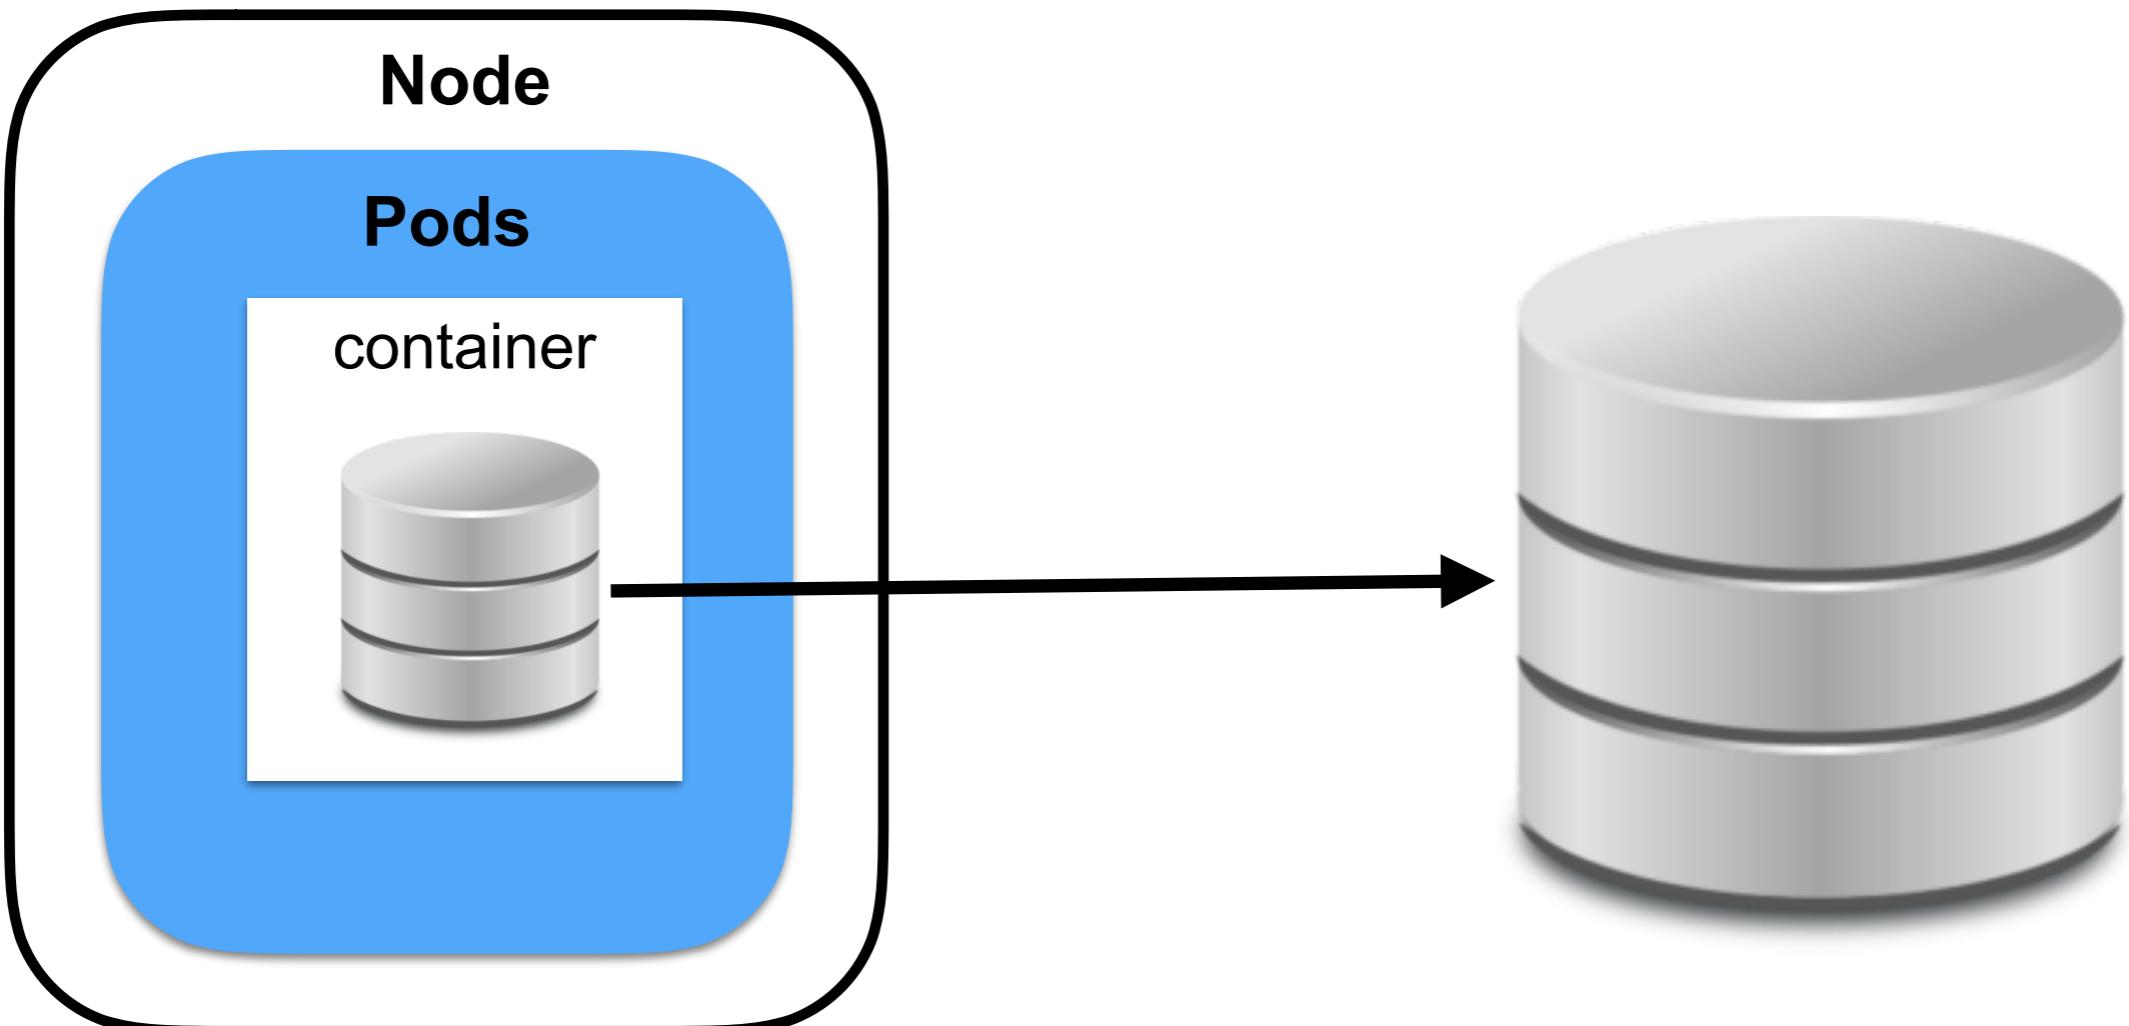
Kubernetes in practice

© 2017 - 2018 Siam Chamnankit Company Limited. All rights reserved.

# Volume



# Volume



# Volume

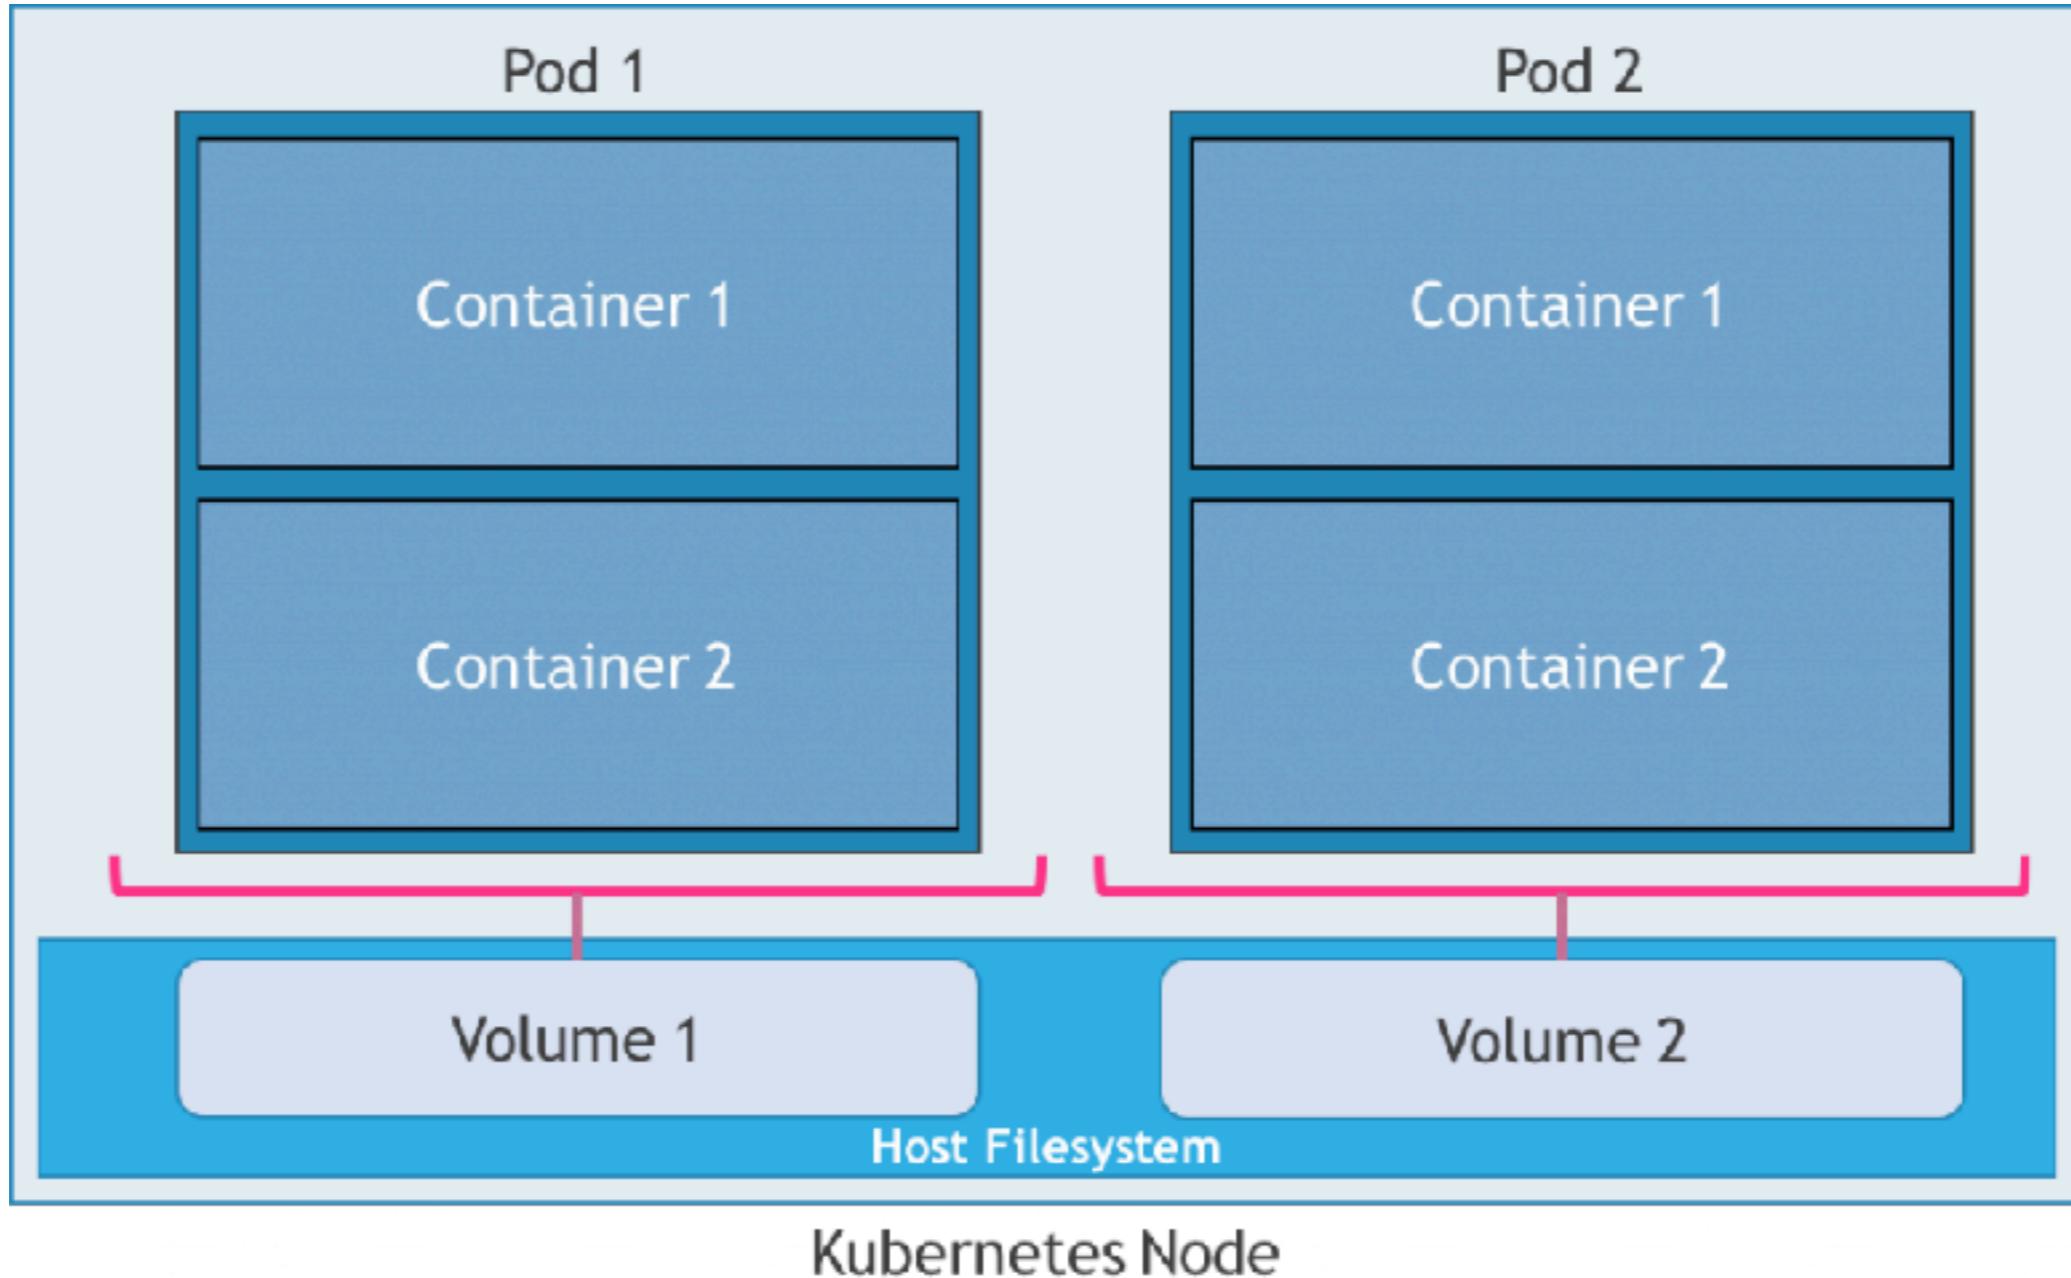
Containers in the same Pods share storage  
(EmptyDir)

- Read and write for all containers in Pods
- Container crash not effect to this storage
- When Pods is deleted then Empty will deleted

Data of Pods/container is **ephemeral**  
all data may loss when Pods is restarted



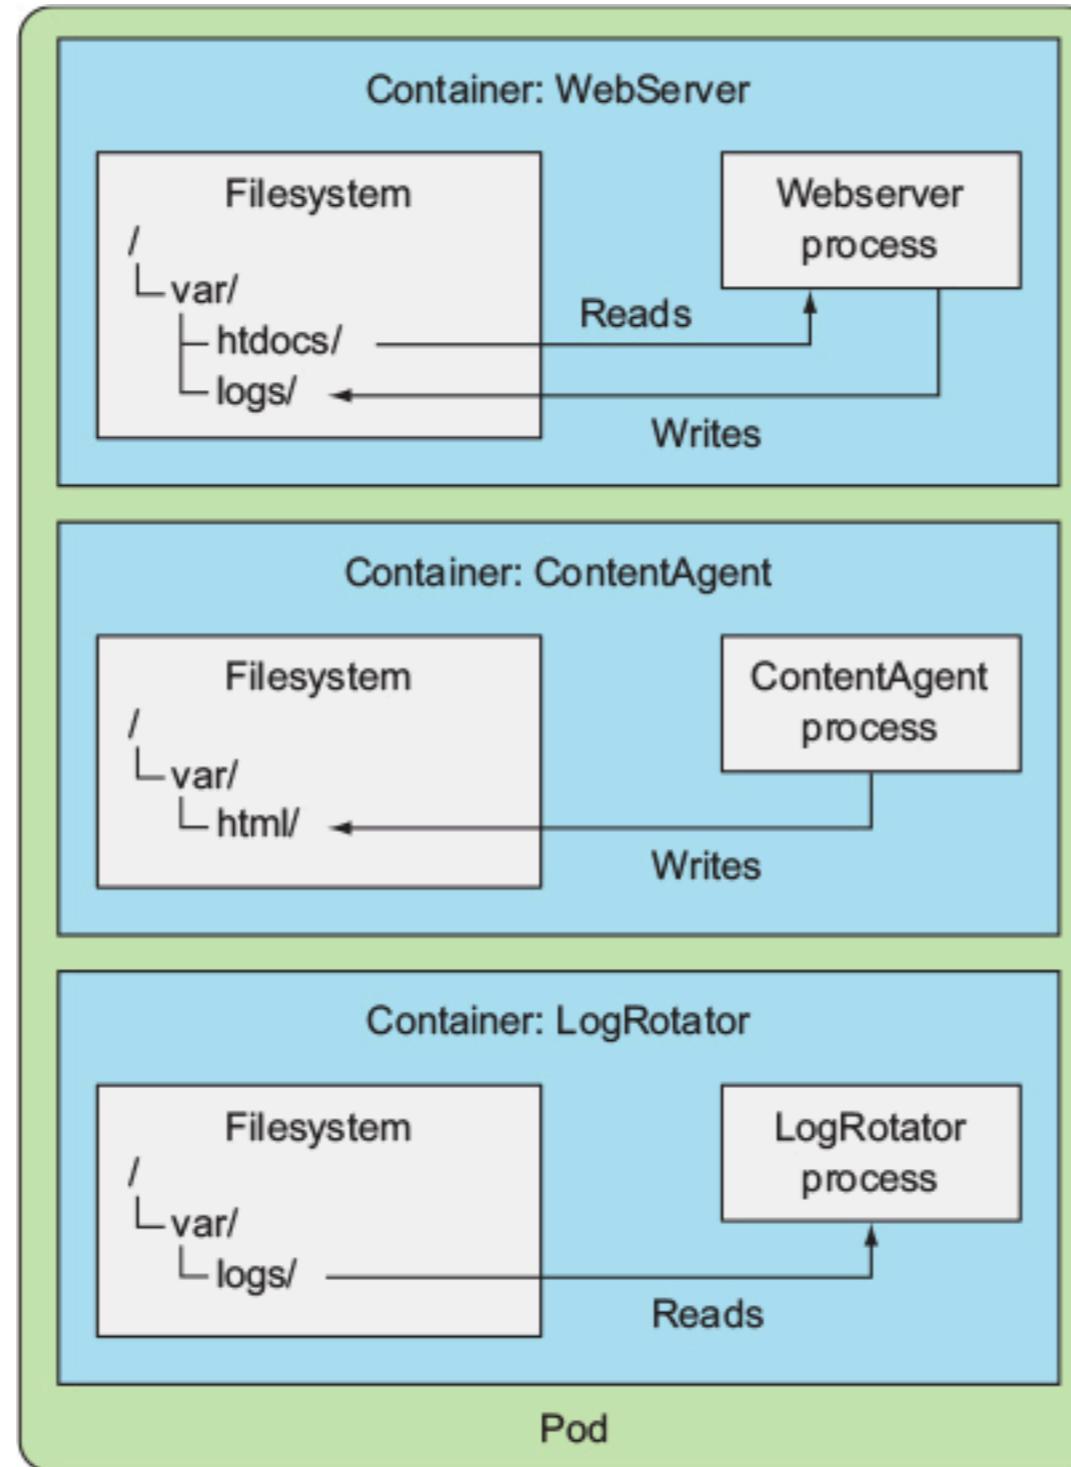
# Containers shared volume



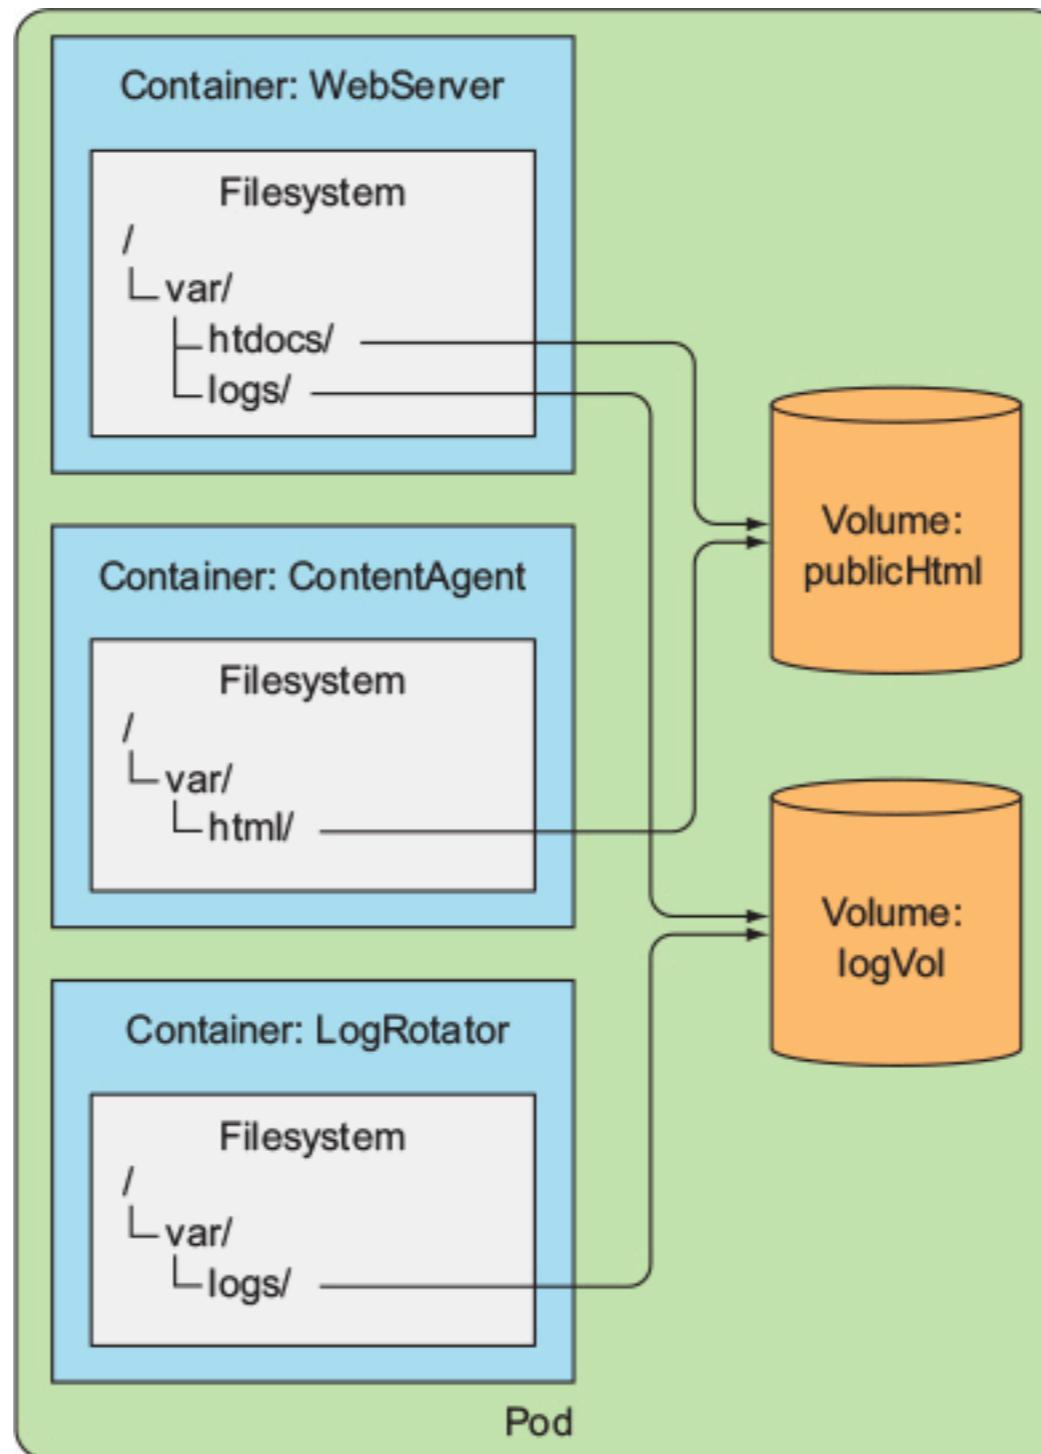
<https://thenewstack.io/strategies-running-stateful-applications-kubernetes-volumes/>



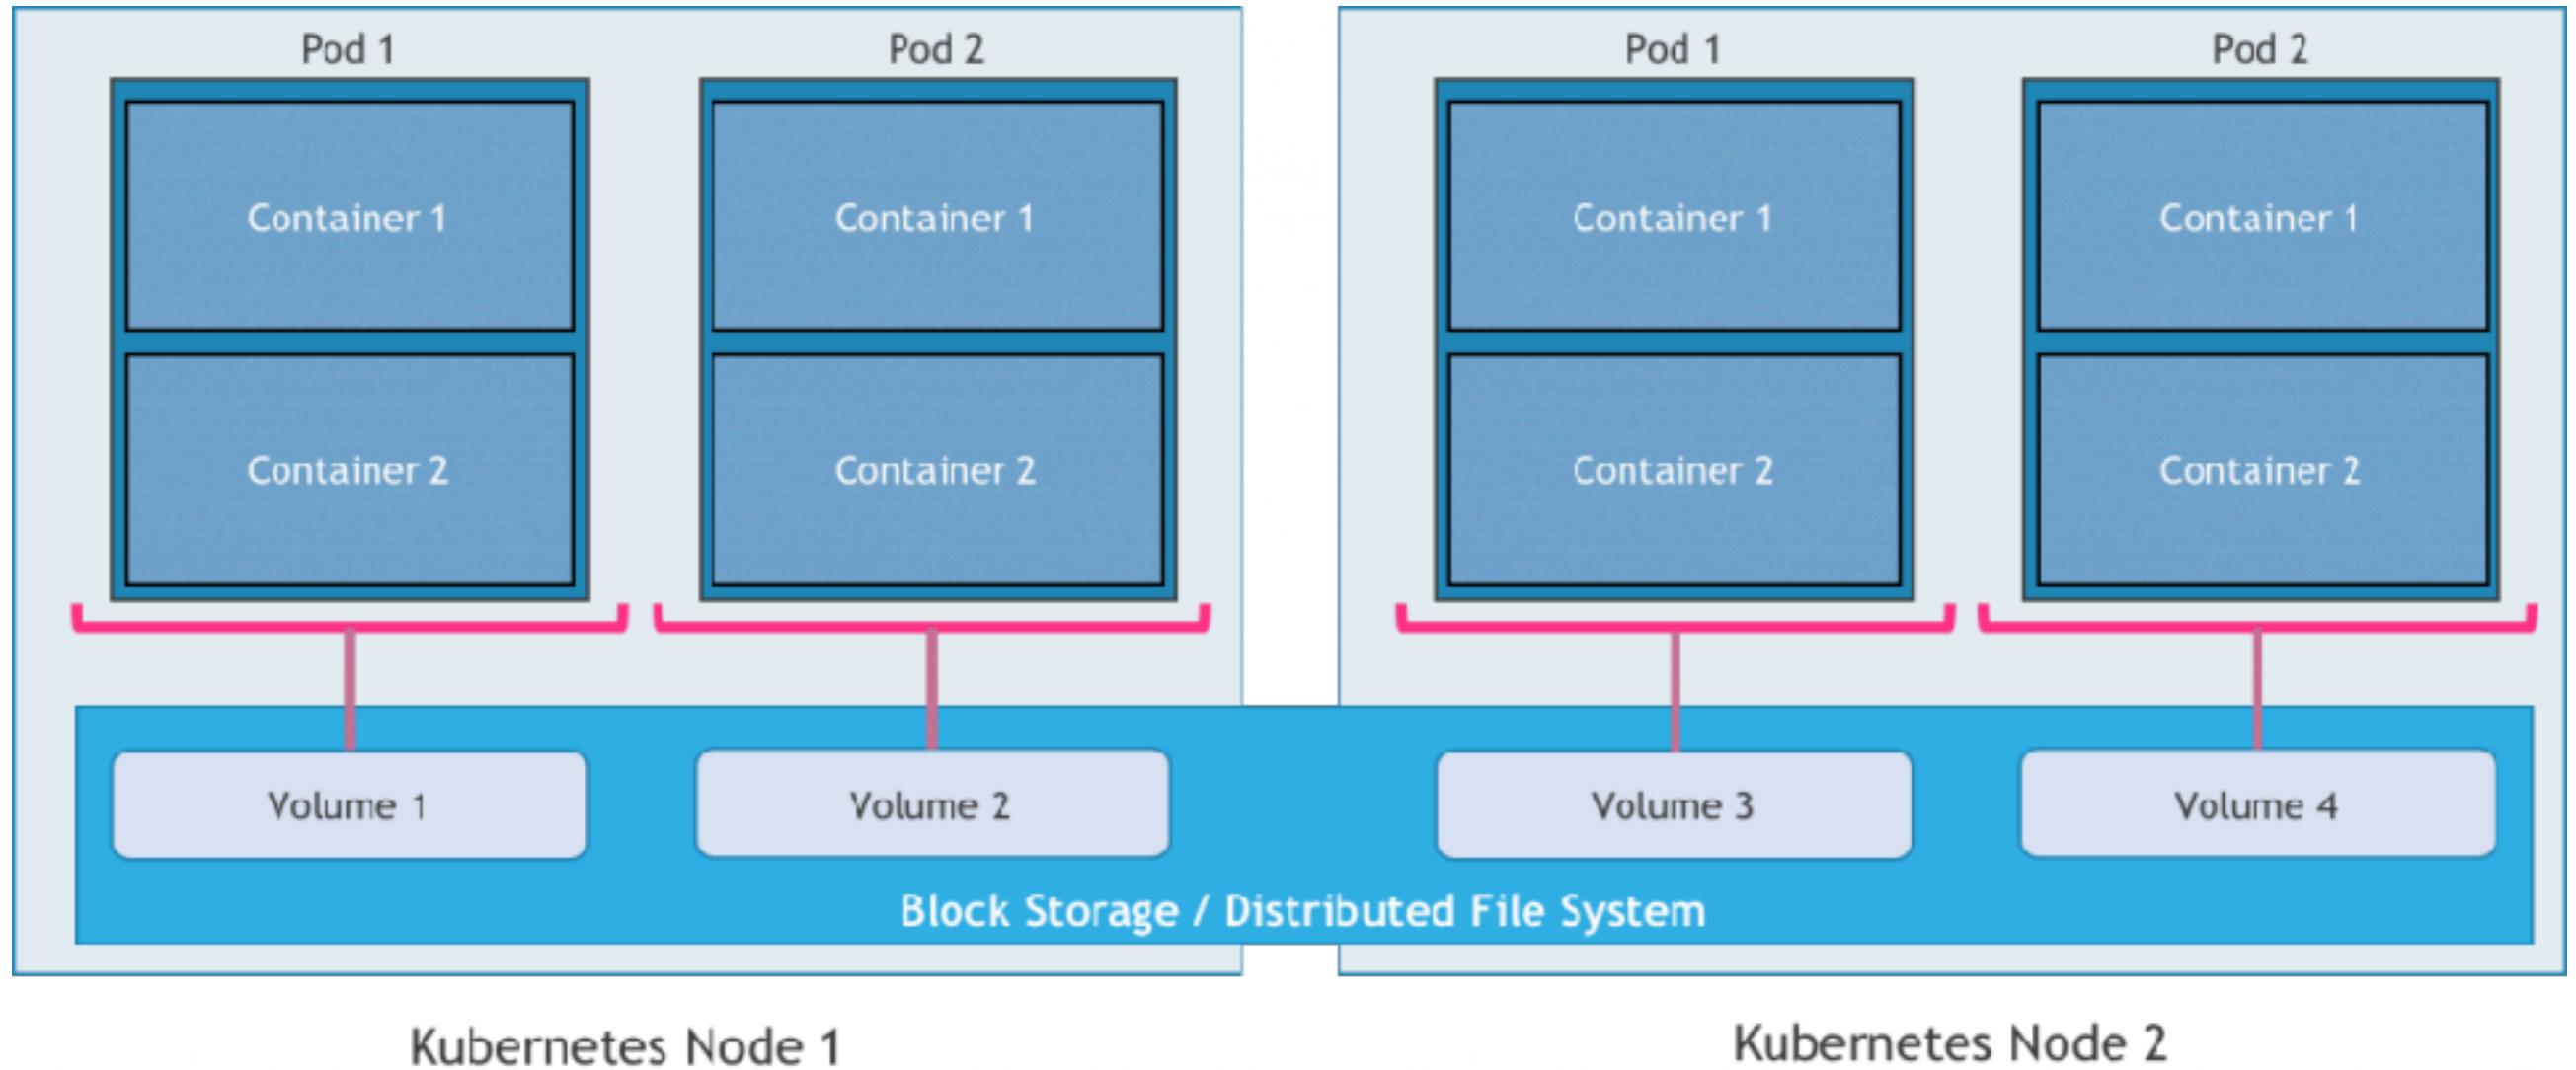
# Without shared storage



# Containers shared volume



# Nodes shared volume



<https://thenewstack.io/strategies-running-stateful-applications-kubernetes-volumes/>



# Volume on Kubernetes

**EmptyDir  
hostPath  
gitRepo  
NFS**

**icePersistentDisk**

**Flocker**

**Persistent Volume Claim (PVC)**

**configMap, secret**

<https://kubernetes.io/docs/concepts/storage/volumes/>



# Kubernetes volumes

EmptyDir

HostPath

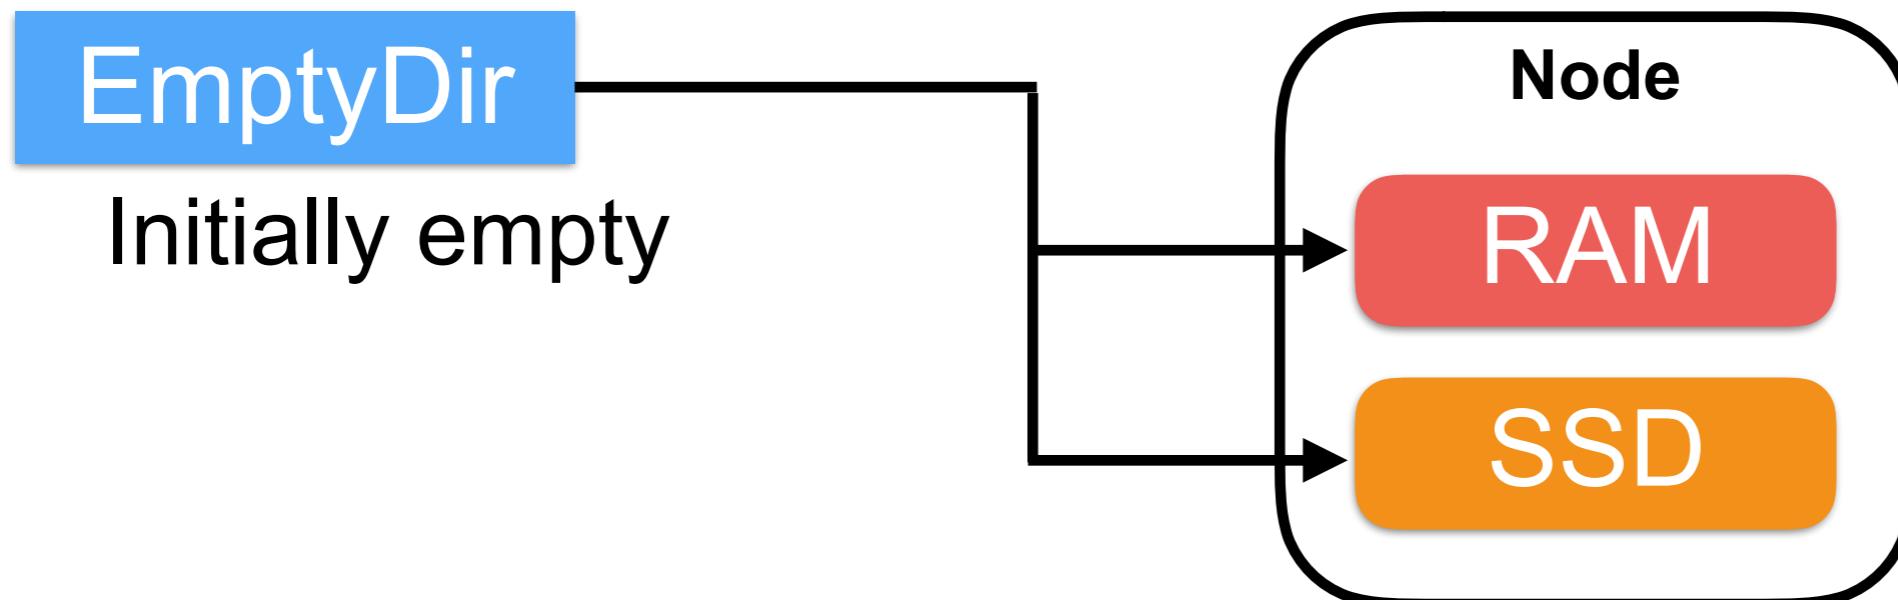
NFS

Cloud Volume

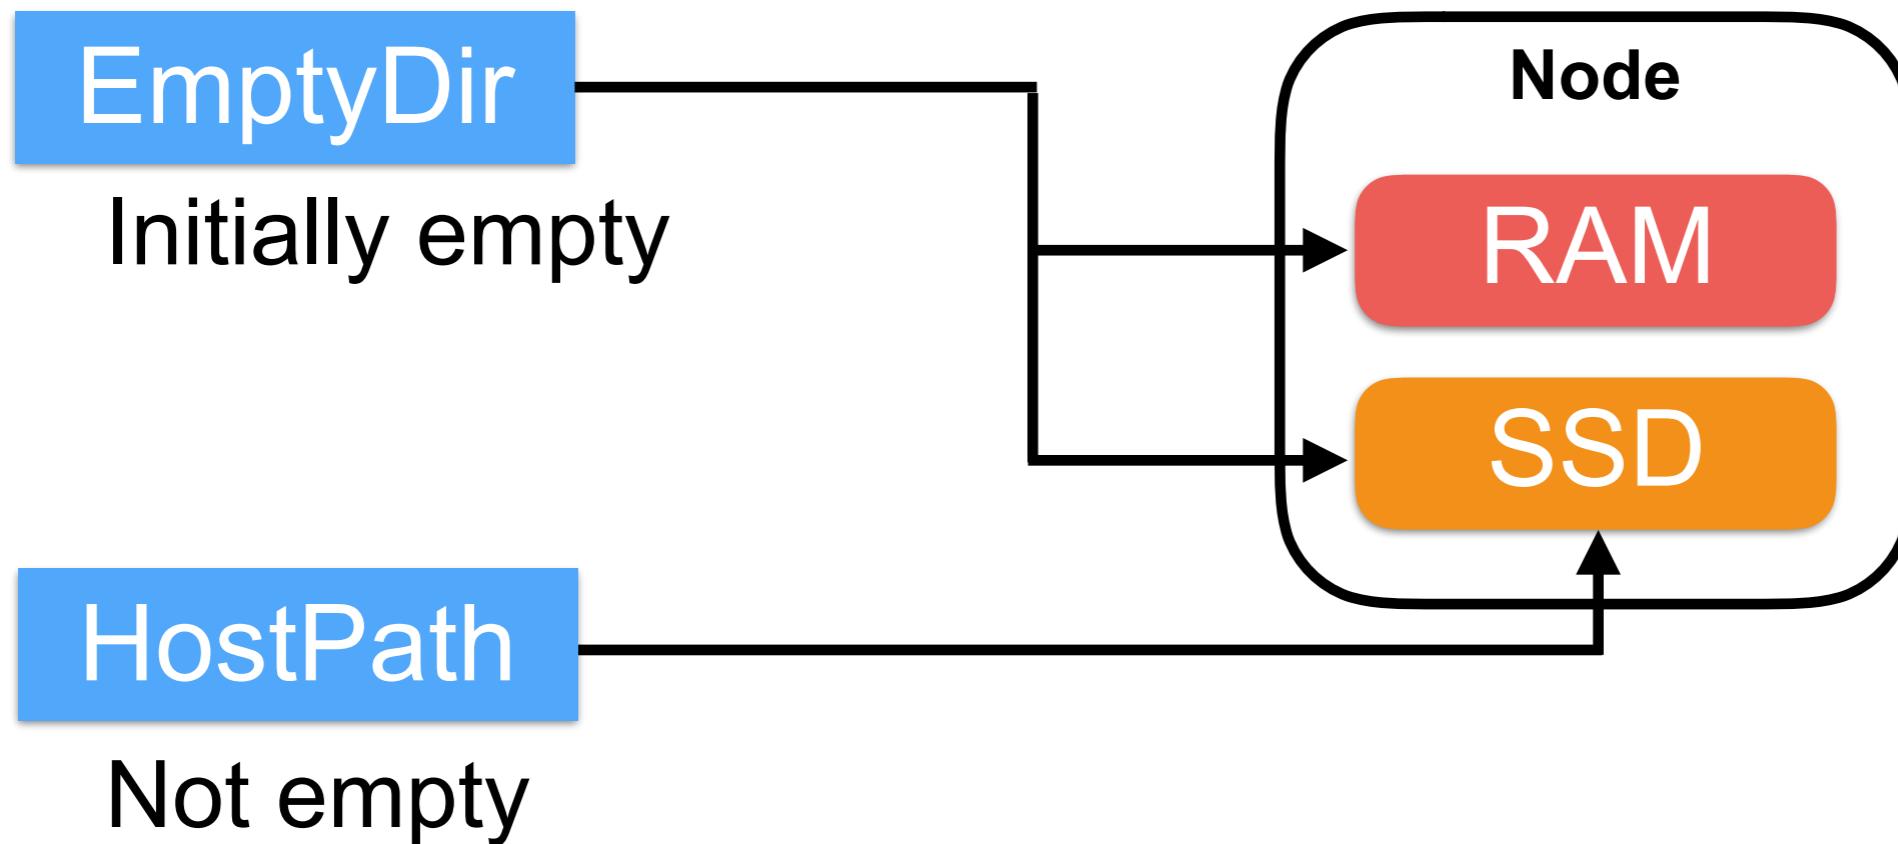
Persistent Volume Claim



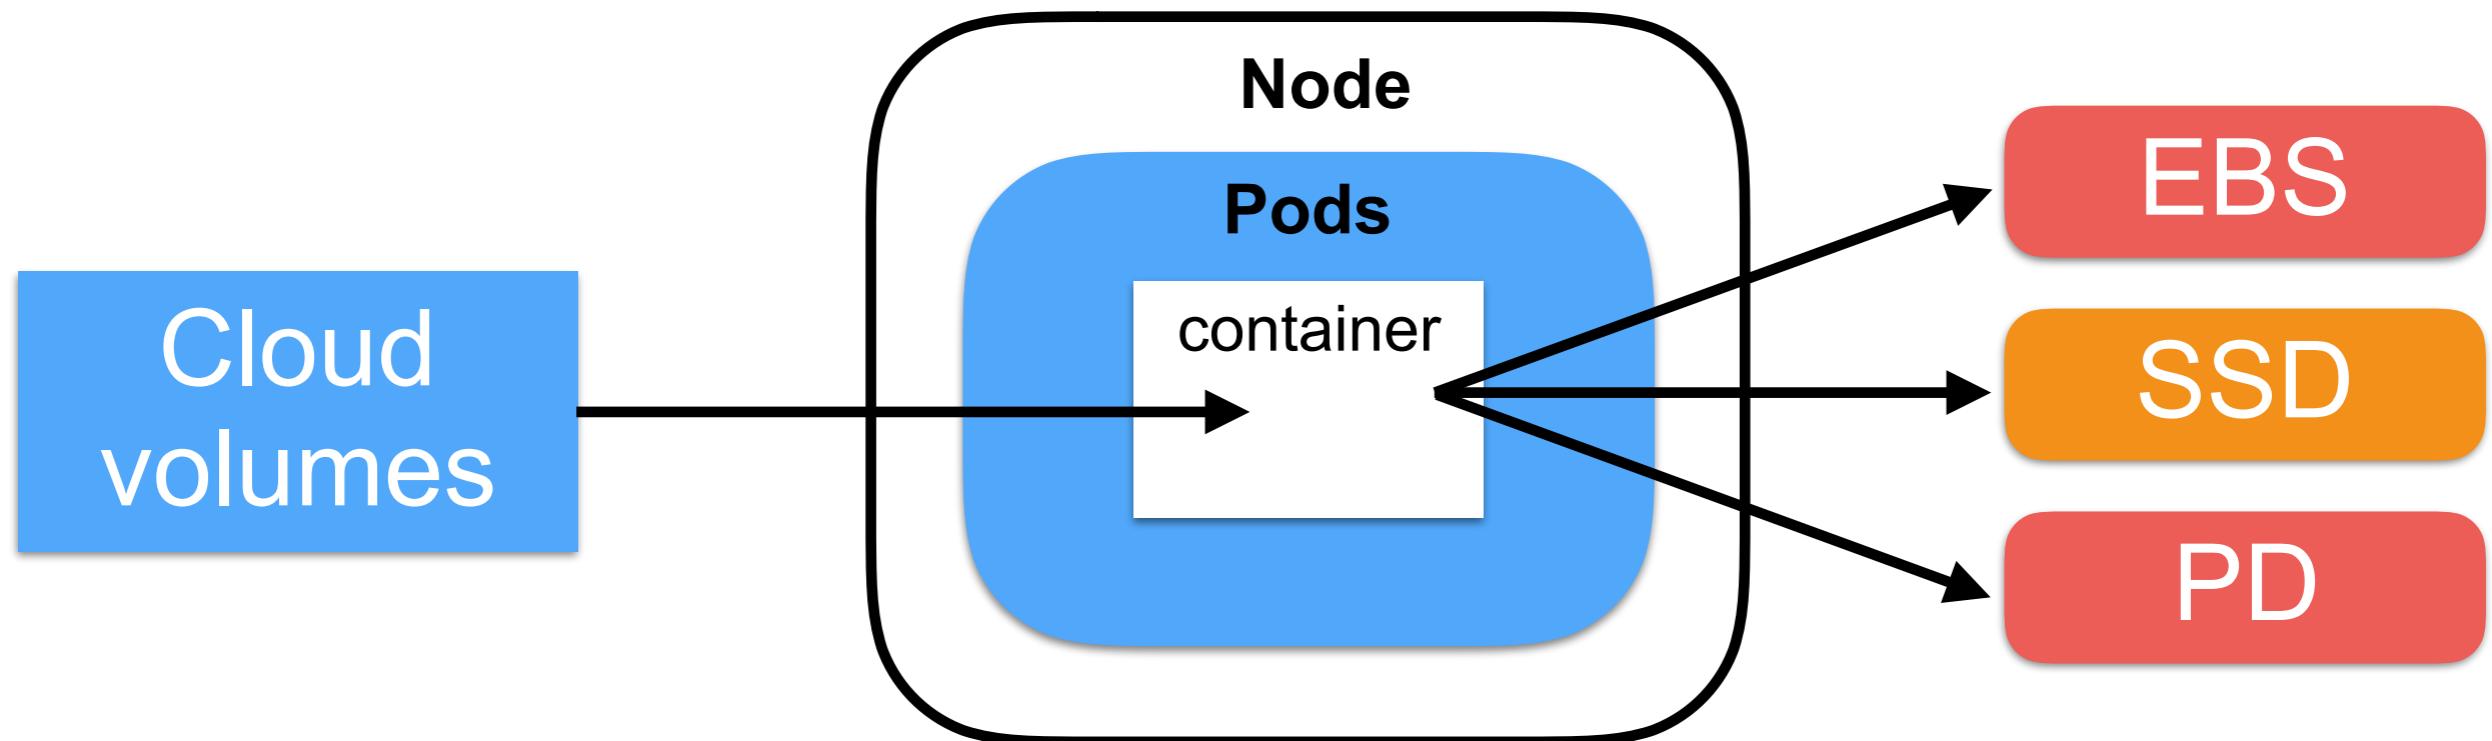
# Kubernetes volumes



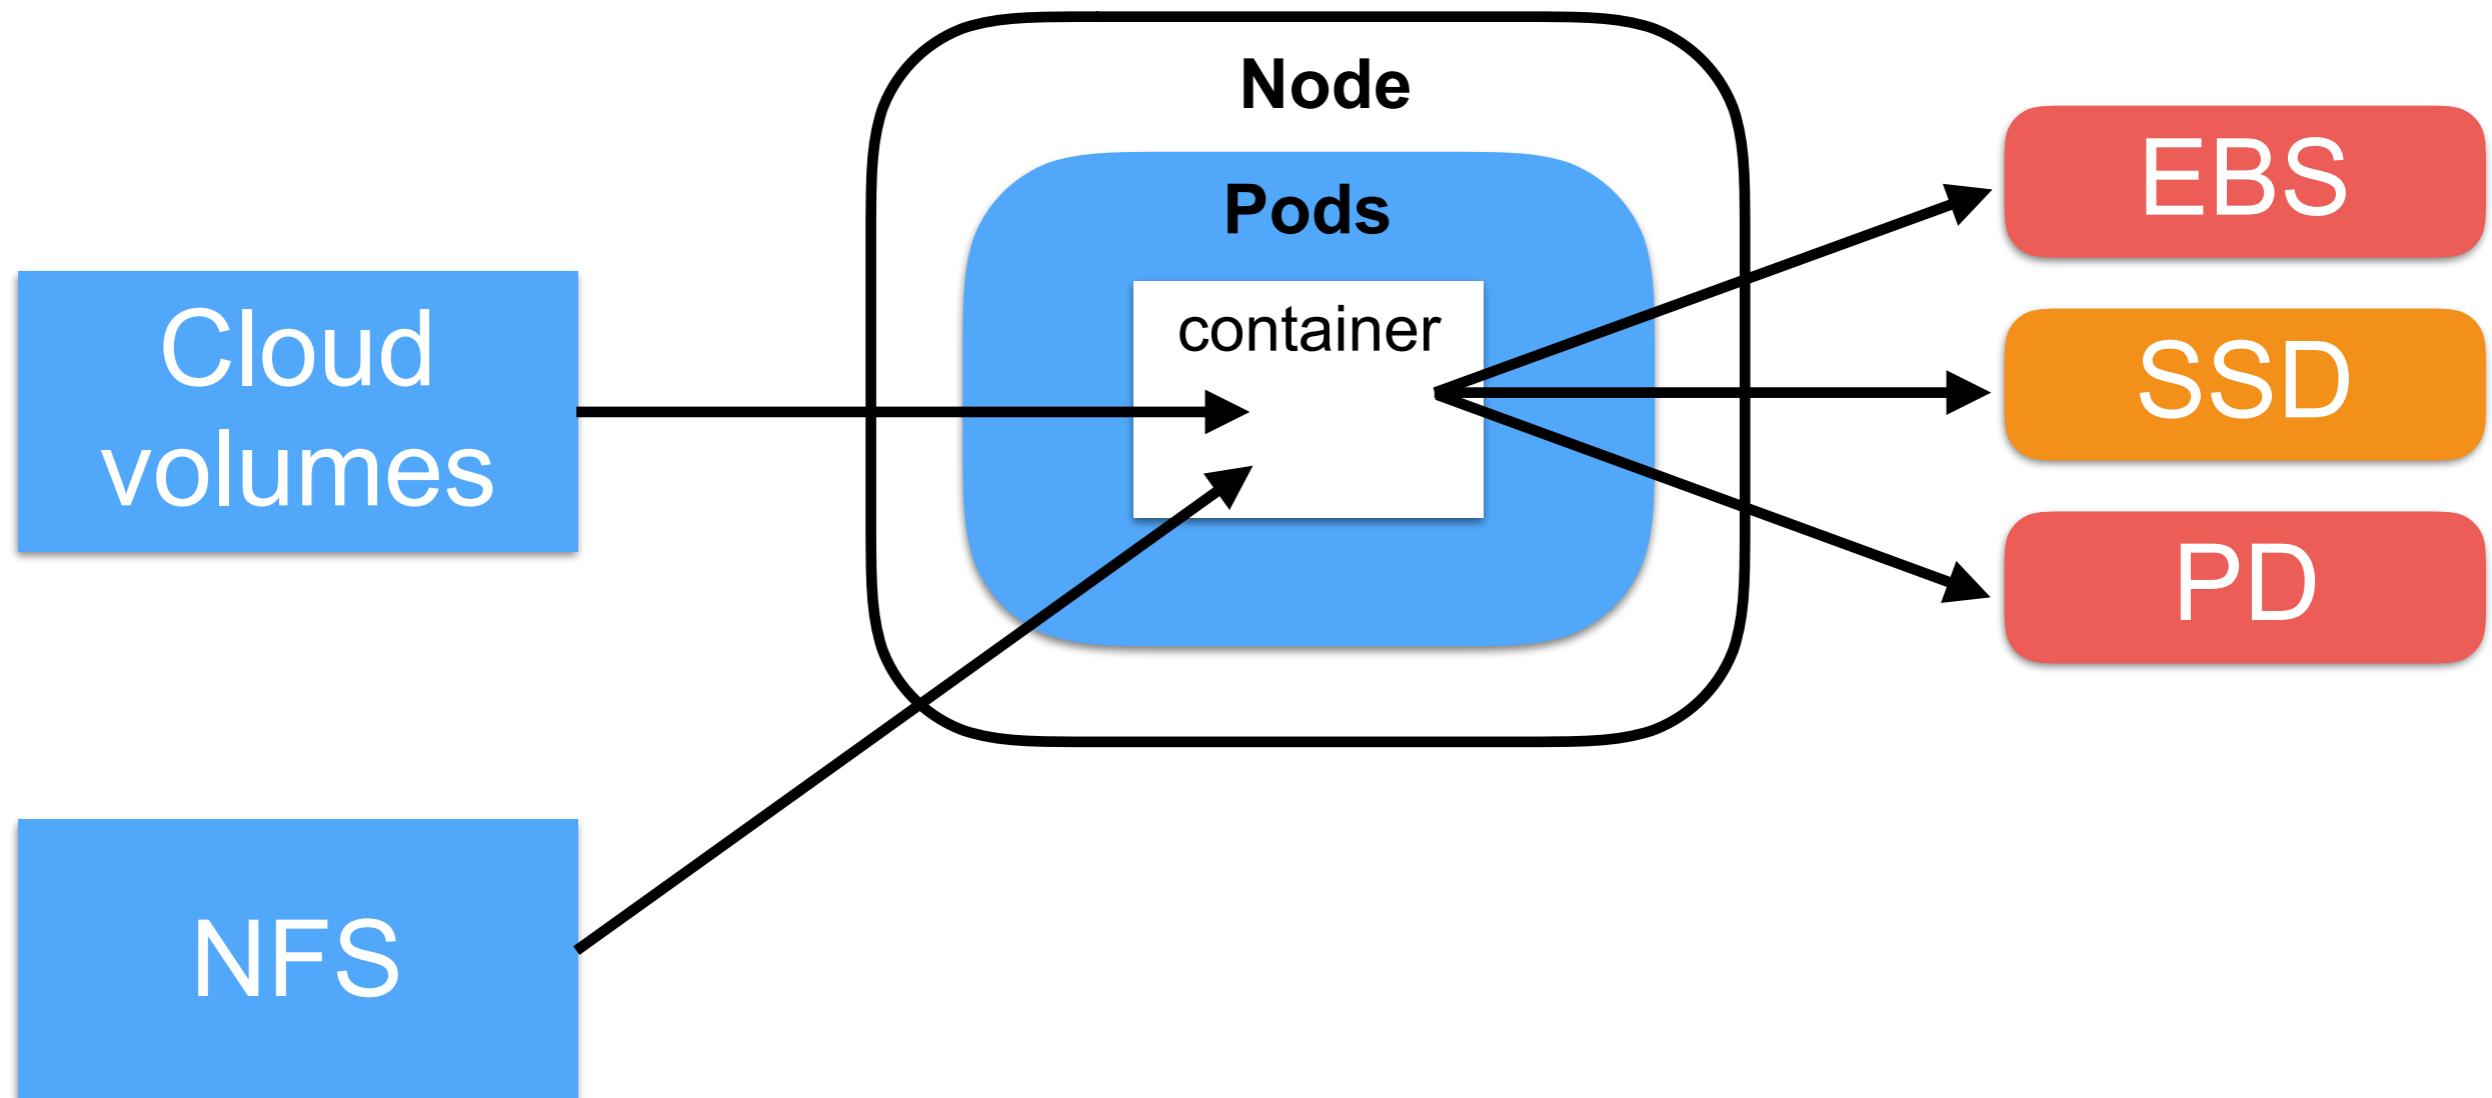
# Kubernetes volumes



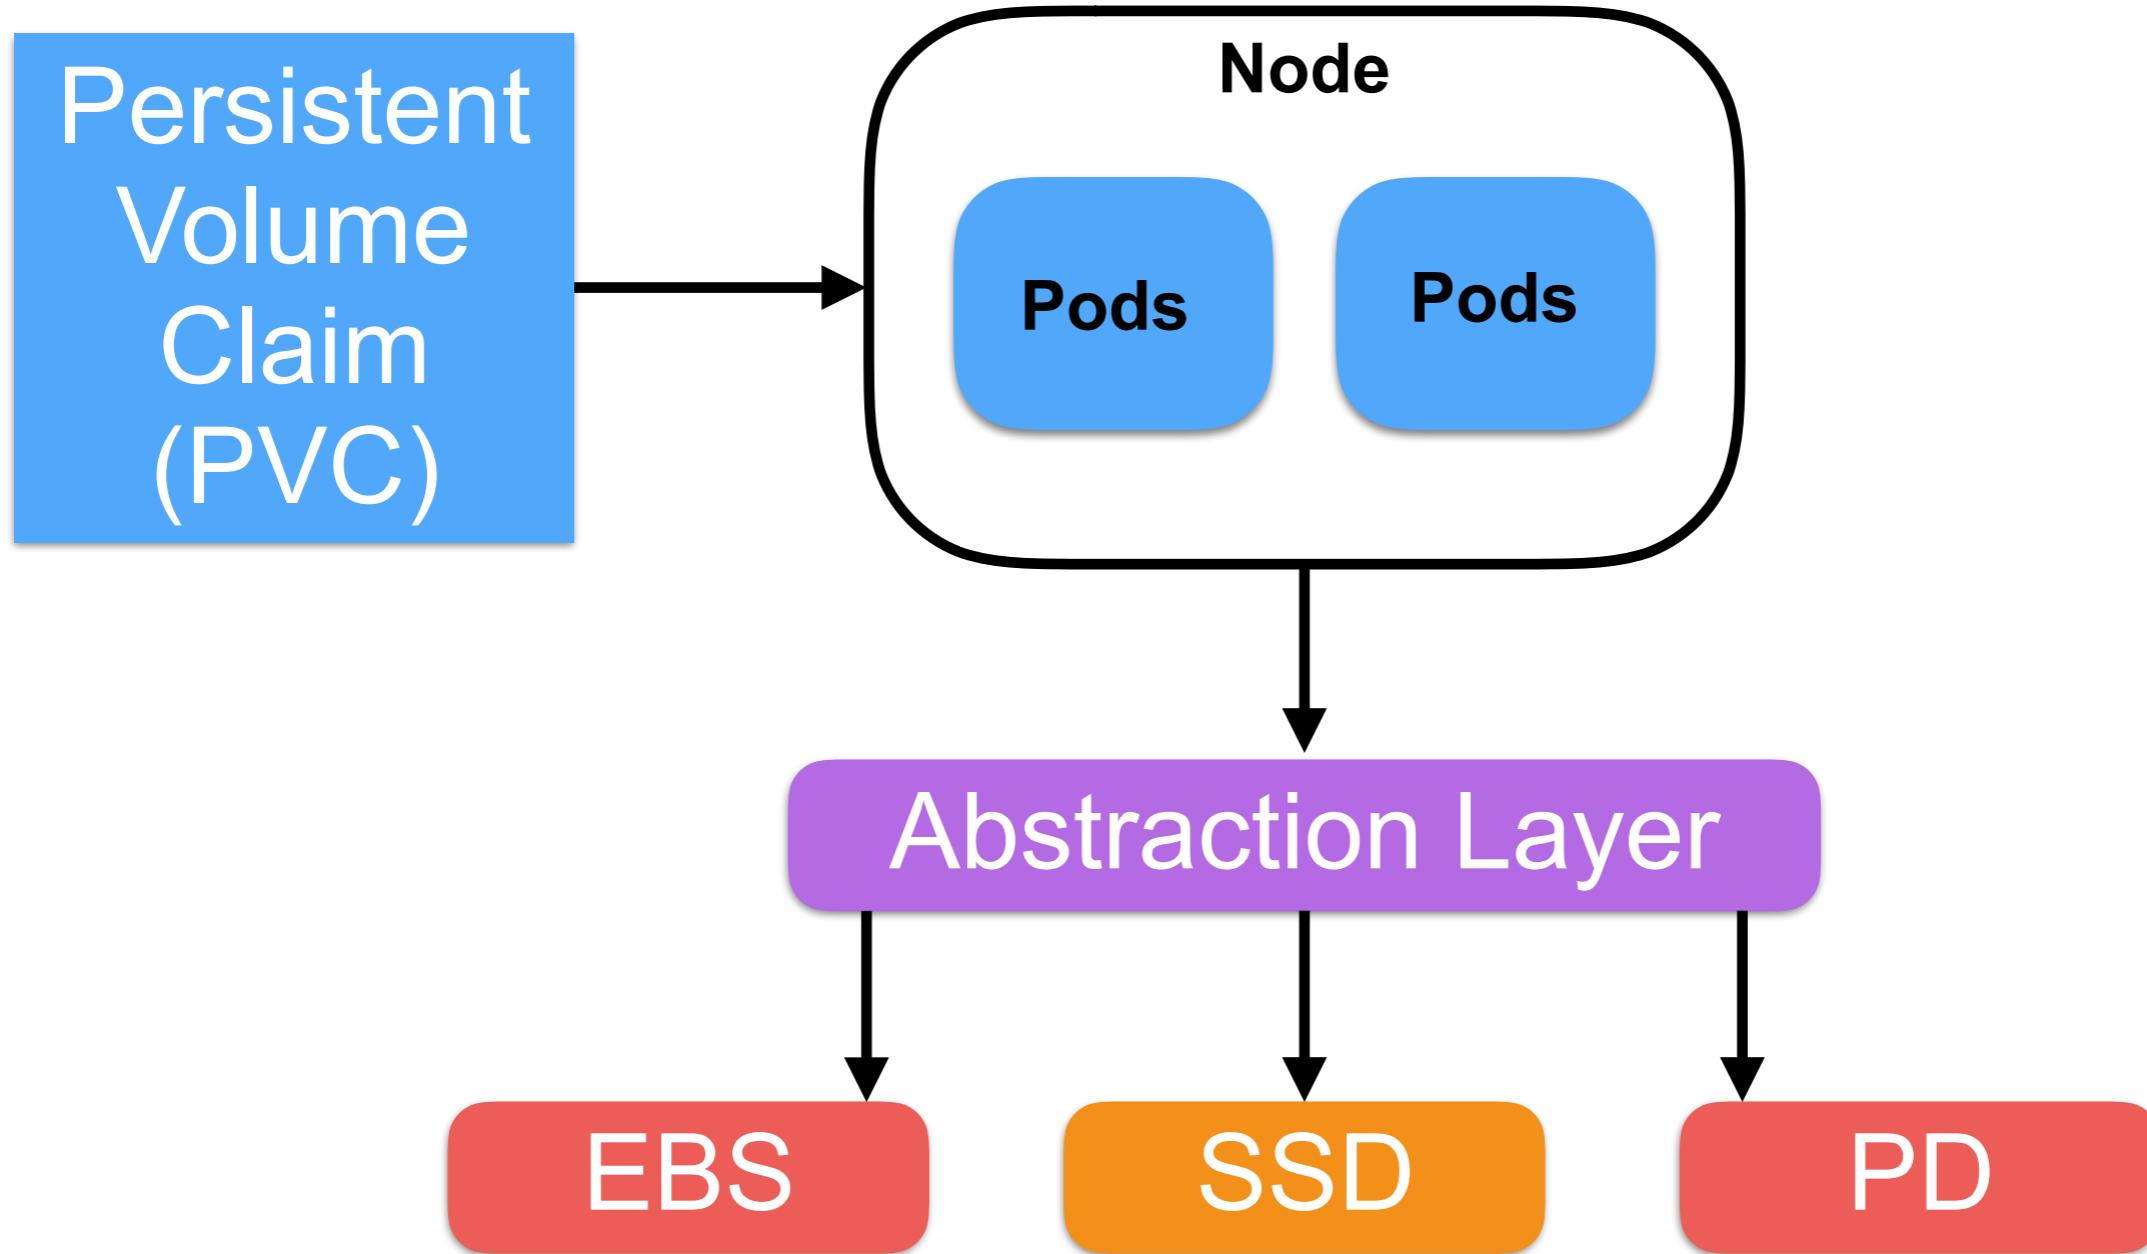
# Kubernetes volumes



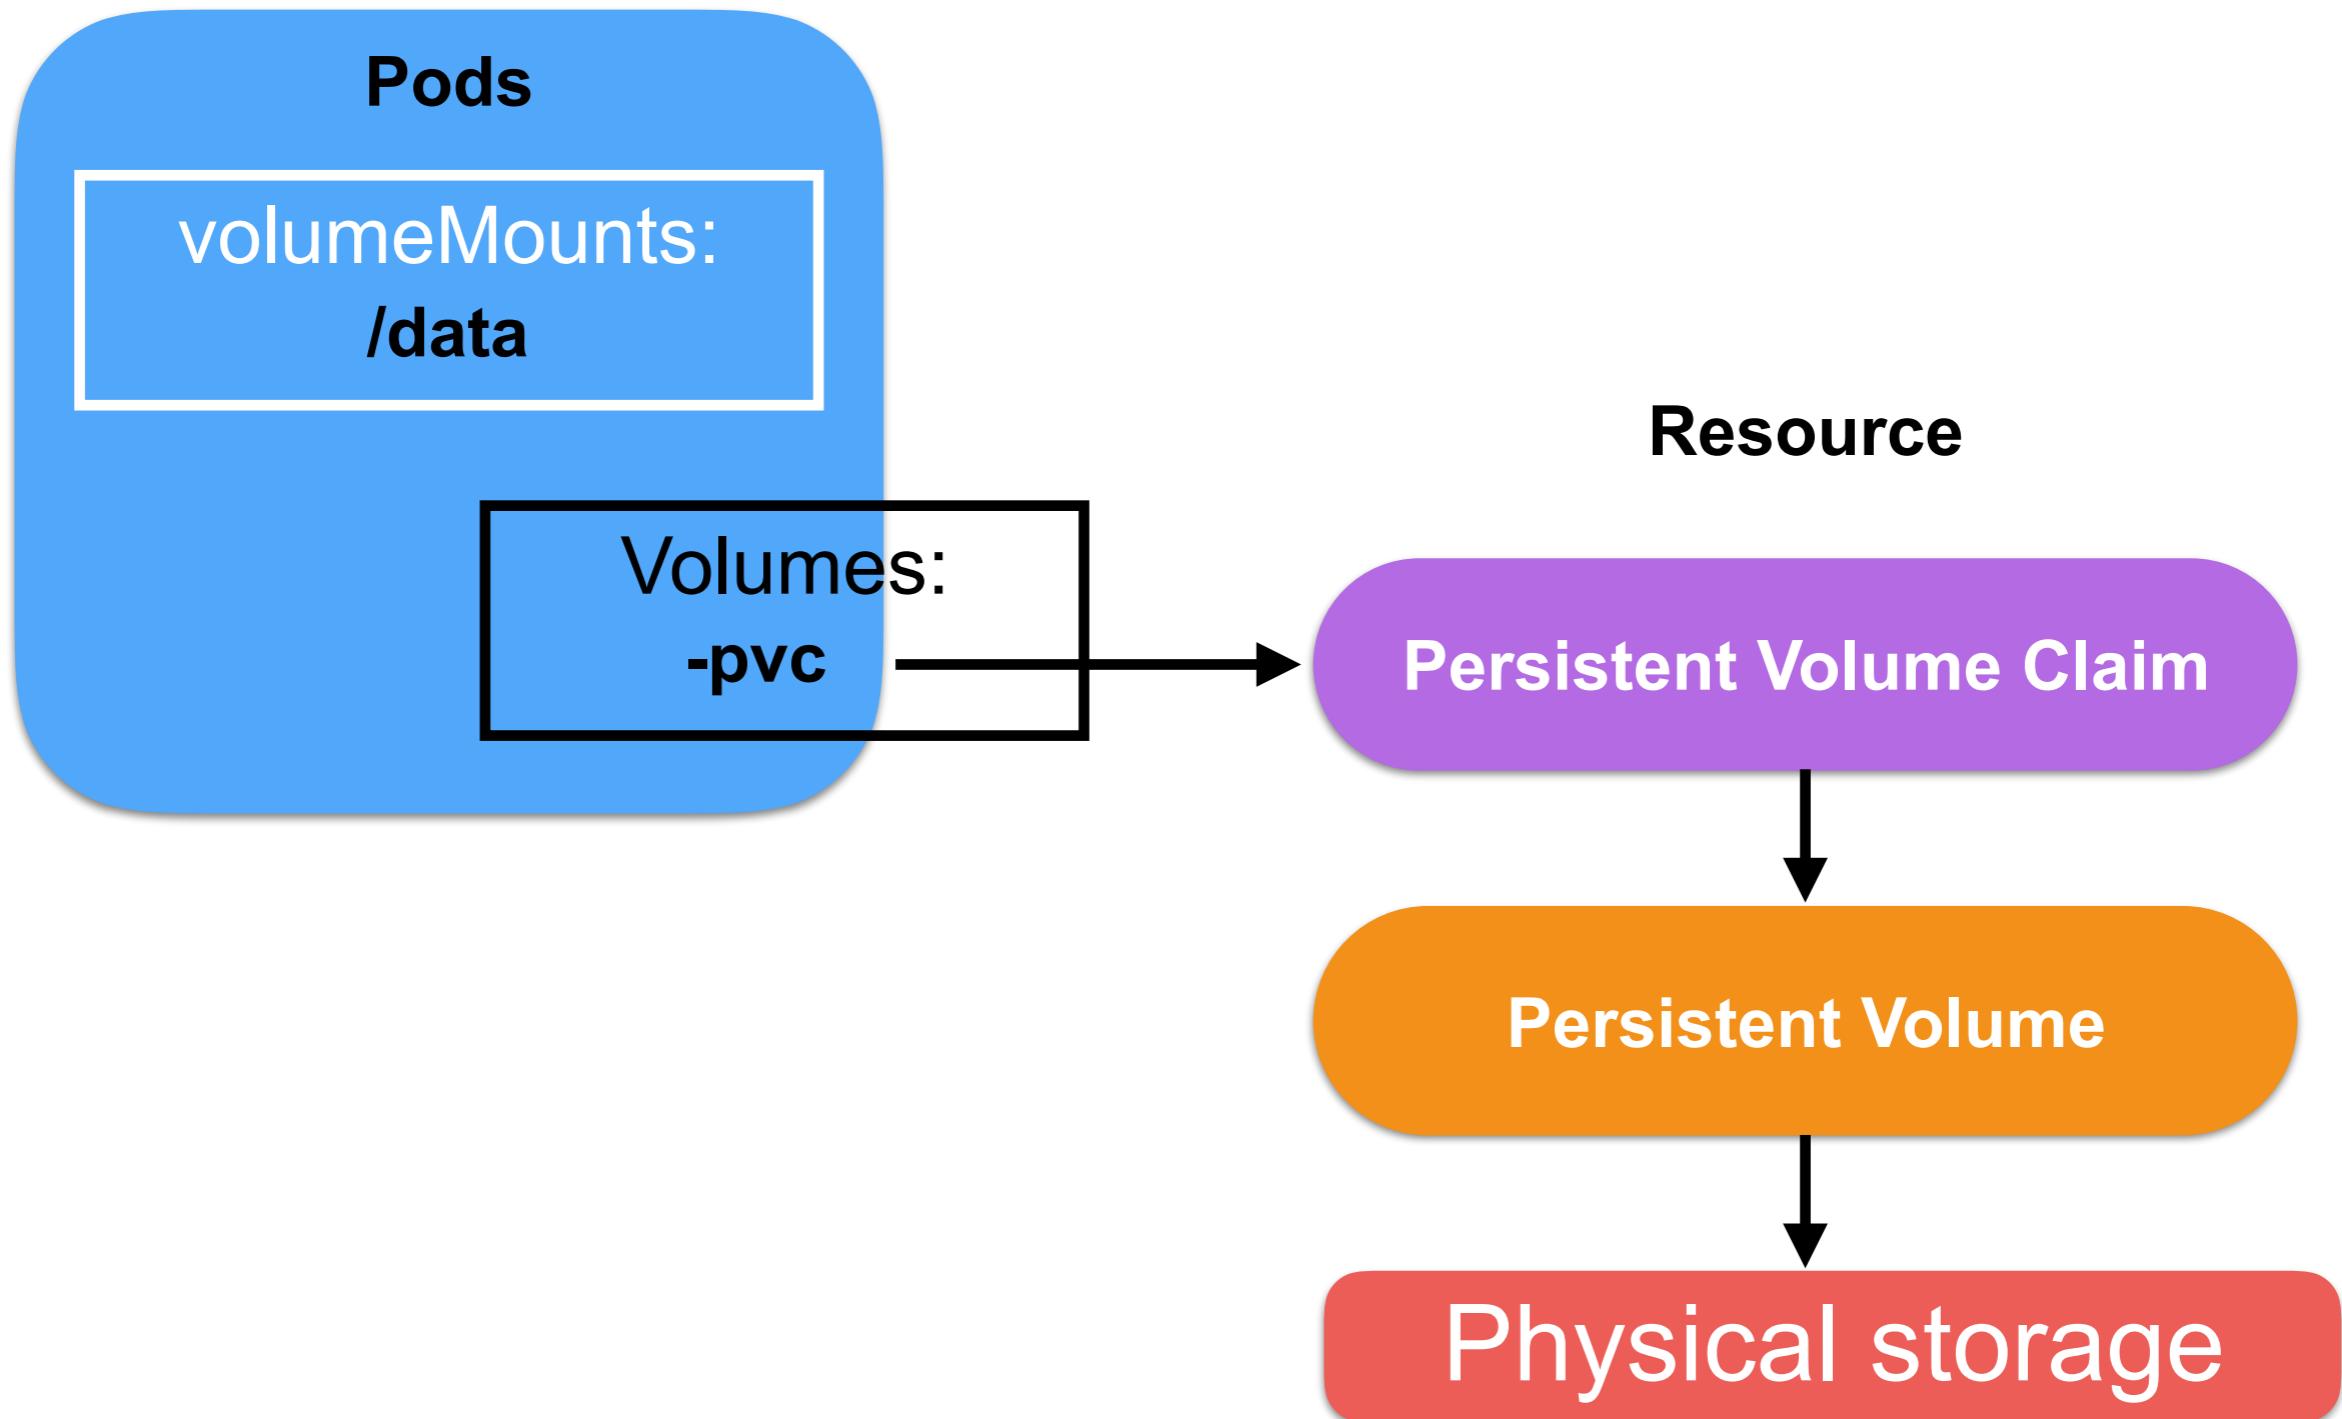
# Kubernetes volumes



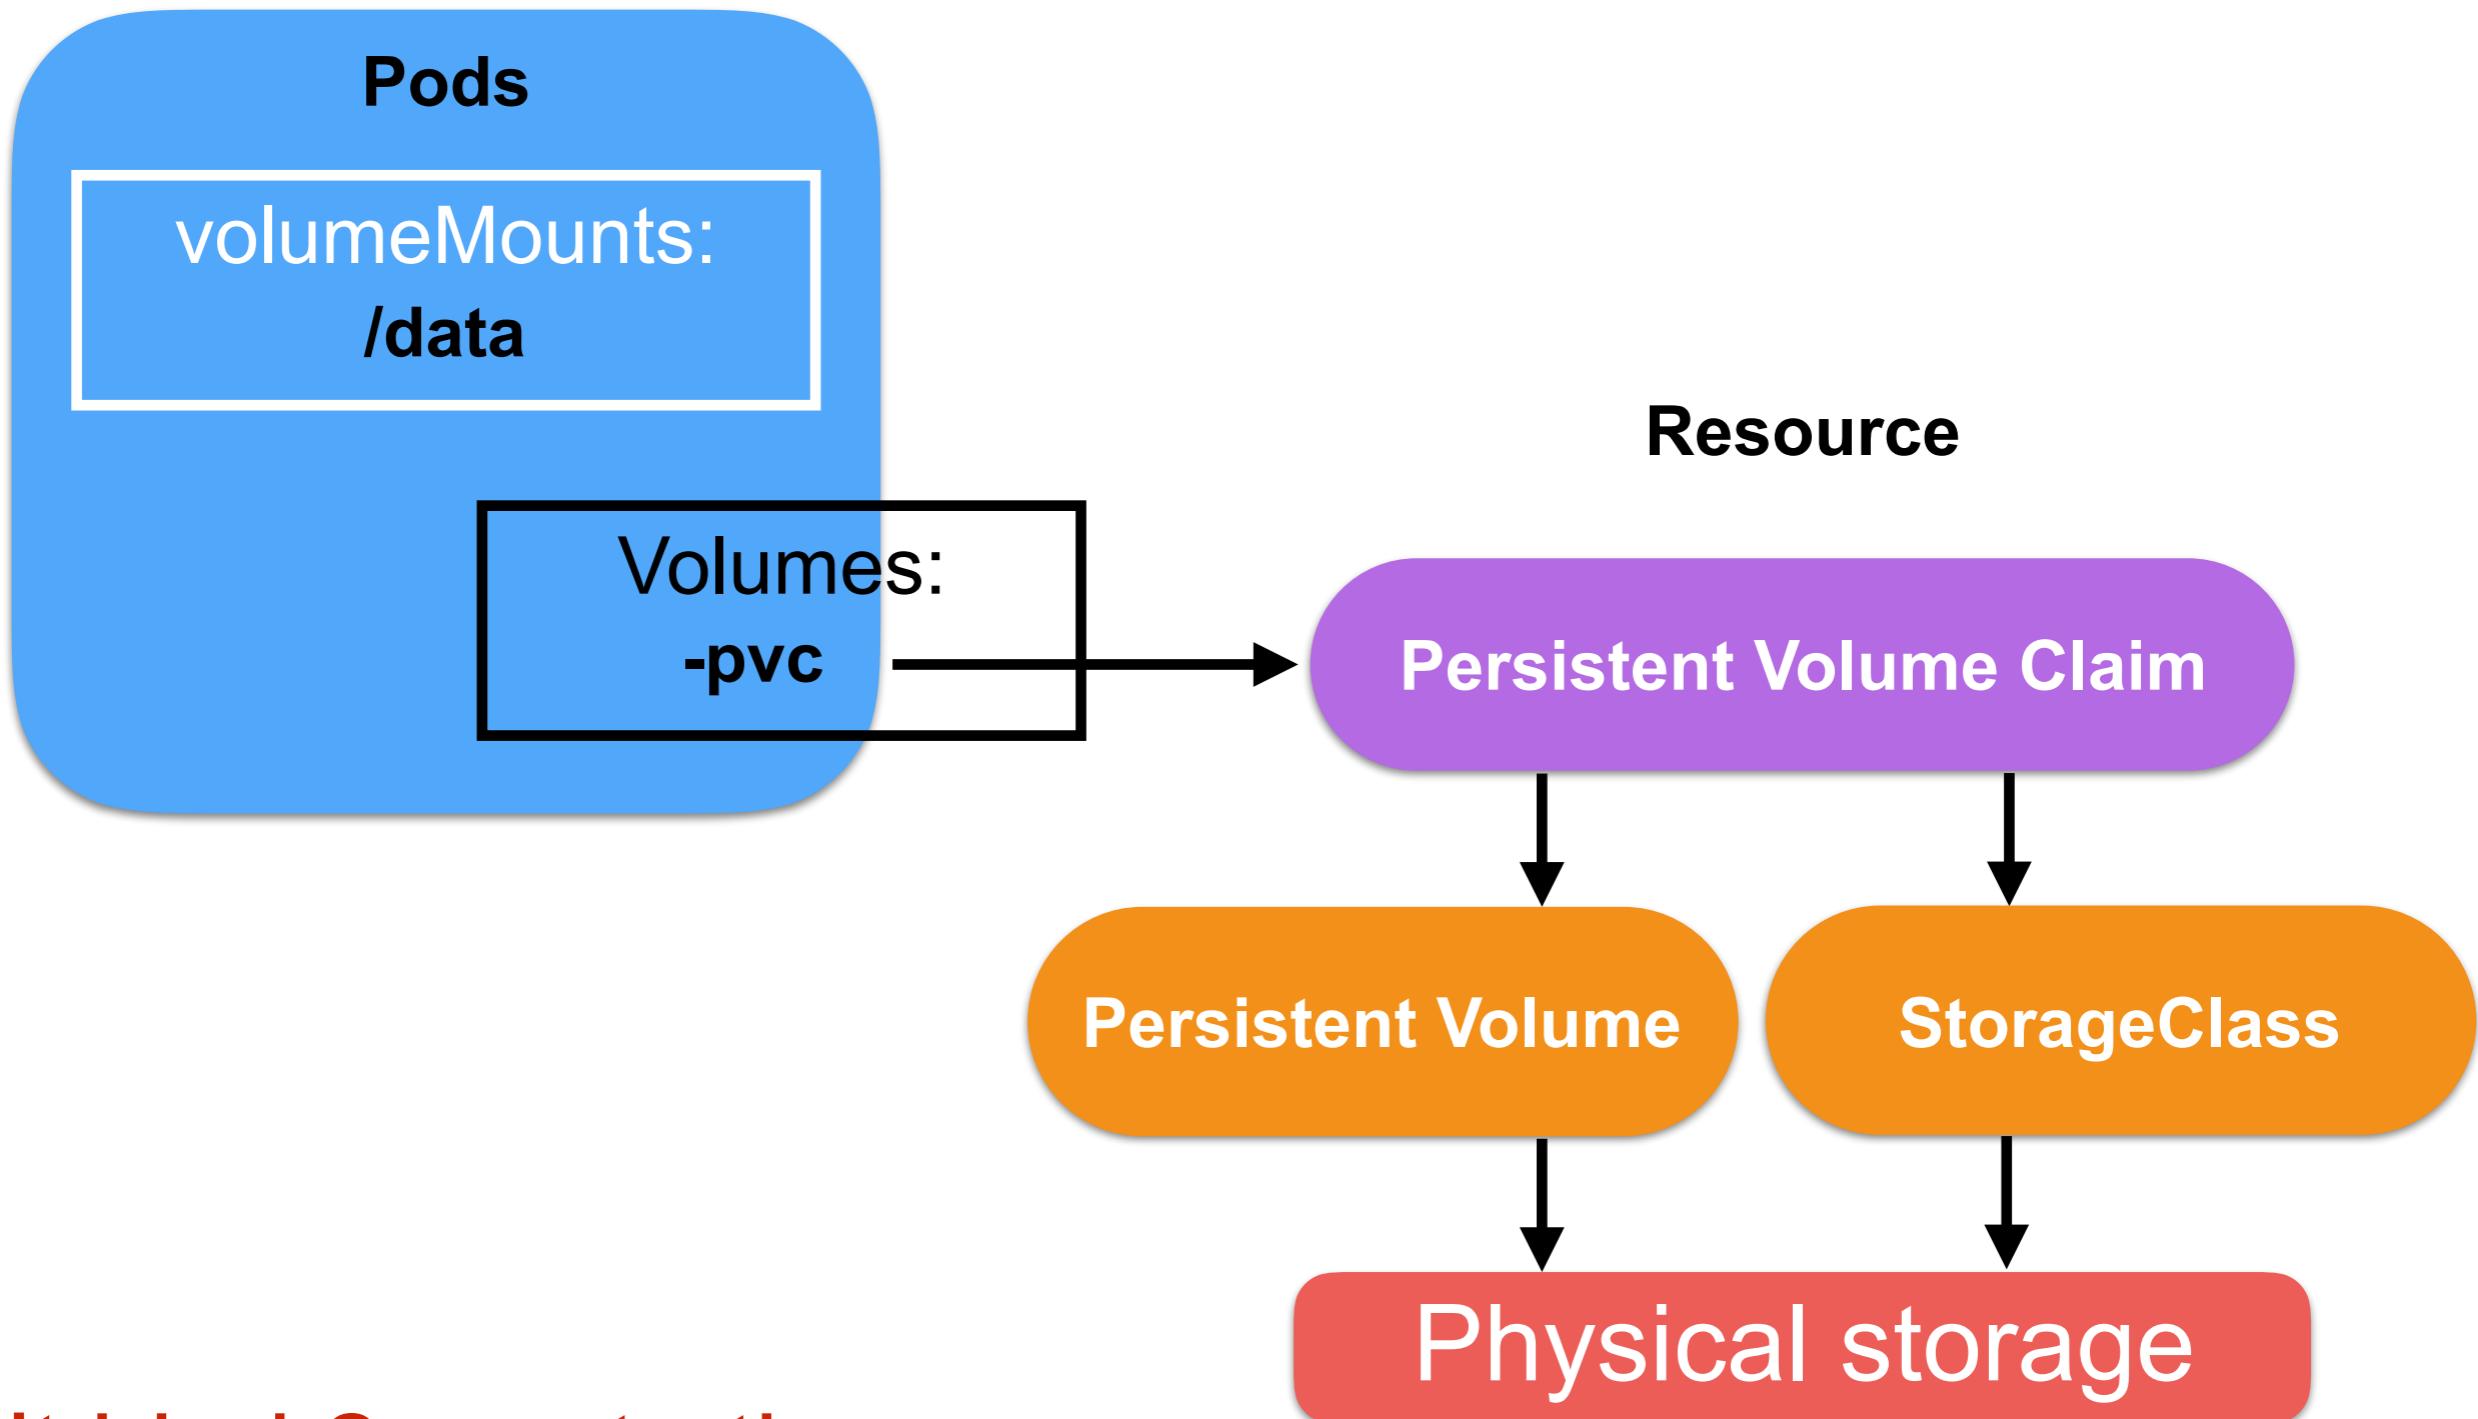
# Kubernetes volumes



# Persistent volume



# Persistent volume



Can't bind 2 pvc to the same pv.



# EmptyDir

Shares volume/storage in same Pods

```
- image: nginx:alpine
  name: web-server
  volumeMounts:
    - name: html
      mountPath: /usr/share/nginx/html
      readOnly: true
  ports:
    - containerPort: 80
      protocol: TCP

  volumes:
    - name: html
      emptyDir:
        medium: Memory
```

